

Join the discussion @ [p2p.wrox.com](http://p2p.wrox.com)



Wrox Programmer to Programmer™



# Beginning C#® 6 Programming with Visual Studio® 2015

Benjamin Perkins, Jacob Vibe Hammer, Jon D. Reid



BEGINNING

# **Visual C#® 2015 Programming**



BEGINNING

# Visual C#® 2015 Programming

Benjamin Perkins

Jacob Vibe Hammer

Jon D. Reid



A Wiley Brand

## Beginning Visual C#® 2015 Programming

Published by

John Wiley & Sons, Inc.  
10475 Crosspoint Boulevard  
Indianapolis, IN 46256  
[www.wiley.com](http://www.wiley.com)

Copyright © 2016 by John Wiley & Sons, Inc., Indianapolis, Indiana

Published simultaneously in Canada

ISBN: 978-1-119-09668-9

ISBN: 978-1-119-09655-9 (ebk)

ISBN: 978-1-119-09656-6 (ebk)

Manufactured in the United States of America

10 9 8 7 6 5 4 3 2 1

No part of this publication may be reproduced, stored in a retrieval system or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, scanning or otherwise, except as permitted under Sections 107 or 108 of the 1976 United States Copyright Act, without either the prior written permission of the Publisher, or authorization through payment of the appropriate per-copy fee to the Copyright Clearance Center, 222 Rosewood Drive, Danvers, MA 01923, (978) 750-8400, fax (978) 646-8600. Requests to the Publisher for permission should be addressed to the Permissions Department, John Wiley & Sons, Inc., 111 River Street, Hoboken, NJ 07030, (201) 748-6011, fax (201) 748-6008, or online at <http://www.wiley.com/go/permissions>.

**Limit of Liability/Disclaimer of Warranty:** The publisher and the author make no representations or warranties with respect to the accuracy or completeness of the contents of this work and specifically disclaim all warranties, including without limitation warranties of fitness for a particular purpose. No warranty may be created or extended by sales or promotional materials. The advice and strategies contained herein may not be suitable for every situation. This work is sold with the understanding that the publisher is not engaged in rendering legal, accounting, or other professional services. If professional assistance is required, the services of a competent professional person should be sought. Neither the publisher nor the author shall be liable for damages arising herefrom. The fact that an organization or Web site is referred to in this work as a citation and/or a potential source of further information does not mean that the author or the publisher endorses the information the organization or Web site may provide or recommendations it may make. Further, readers should be aware that Internet Web sites listed in this work may have changed or disappeared between when this work was written and when it is read.

For general information on our other products and services please contact our Customer Care Department within the United States at (877) 762-2974, outside the United States at (317) 572-3993 or fax (317) 572-4002.

Wiley publishes in a variety of print and electronic formats and by print-on-demand. Some material included with standard print versions of this book may not be included in e-books or in print-on-demand. If this book refers to media such as a CD or DVD that is not included in the version you purchased, you may download this material at <http://book-support.wiley.com>. For more information about Wiley products, visit [www.wiley.com](http://www.wiley.com).

**Library of Congress Control Number:** 2015957031

**Trademarks:** Wiley, the Wiley logo, Wrox, the Wrox logo, Programmer to Programmer, and related trade dress are trademarks or registered trademarks of John Wiley & Sons, Inc. and/or its affiliates, in the United States and other countries, and may not be used without written permission. Visual C# is a registered trademark of Microsoft Corporation. All other trademarks are the property of their respective owners. John Wiley & Sons, Inc., is not associated with any product or vendor mentioned in this book.

# ABOUT THE AUTHORS

**BENJAMIN PERKINS** (MBA, MCSD, ITIL) is currently employed at Microsoft in Munich, Germany, as a Senior Technical Advisor for IIS, ASP.NET, and Azure App Services. He has been working professionally in the IT industry for over two decades. He started computer programming with QBasic at the age of 11 on an Atari 1200XL desktop computer. He takes pleasure in the challenges that troubleshooting technical issues have to offer and savors in the rewards of a well written program. After completing high school he joined the United States Army. After successfully completing his military service, he attended Texas A&M University in College Station, Texas, where he received a Bachelor of Business Administration in Management Information Systems.

His roles in the IT industry have spanned the entire spectrum including programmer, system architect, technical support engineer, team leader, and mid-level management. While employed at Hewlett-Packard, he received numerous awards, degrees, and certifications. He has a passion for technology and customer service and looks forward to troubleshooting and writing more world-class technical solutions.

“My approach is to write code with support in mind, and to write it once correctly and completely so we do not have to come back to it again, except to enhance it.”

Benjamin is married to Andrea and has two wonderful children, Lea and Noa.

**JACOB VIBE HAMMER** is a software architect and developer at Kamstrup, where he is helping the company develop world-class Smart Grid solutions for large public utilities. He started programming just about the time when he was able to spell the word “BASIC” — which, incidentally, is the first programming language he ever used. Since then, he has worked with numerous programming languages and solution architectures; however, since the turn of the century, he has worked primarily with the .NET platform. Today, his programming time is spent working primarily with C# and WPF, as well as toying with NoSQL databases. A Danish citizen, Jacob lives in Aarhus, Denmark, with his wife and two sons.

**JON D. REID** is a Product Solution Manager for IFS Field Service Management ([www.IFSWORLD.com](http://www.IFSWORLD.com)). He has coauthored a number of books, including *Beginning Visual C# 2010*, *Fast Track C#*, *Pro Visual Studio .NET*, and many others.

# ABOUT THE TECHNICAL EDITOR

**JOHN MUELLER** is a freelance author and technical editor. He has writing in his blood, having produced 99 books and more than 600 articles to date. The topics range from networking to artificial intelligence and from database management to heads-down programming. Some of his current books include books on topics such as Python for beginners, Python for data scientists, and MATLAB. He has also written a Java e-learning kit, a book on HTML5 development with JavaScript, and another on CSS3. His technical editing skills have helped more than 63 authors refine the content of their manuscripts. John has provided technical editing services to both *Data Based Advisor* and *Coast Compute* magazines. Be sure to read John's blog at <http://blog.johnmuellerbooks.com/>.

# **ACKNOWLEDGMENTS**

It takes a lot of work to get content into a presentable format for students and IT professionals to read and get value from. The authors indeed have technical knowledge and experiences to share, but without the technical writers, technical reviewers, developers, editors, publishers, graphic designers, the list goes on, providing their valuable input, a book of high quality could not be written. The rate of change occurs too quickly for an individual to perform all these tasks and still publish a book that is valid before the technology becomes stale. This is why authors worked together with a great team to get all the components of the book together quickly. It was done to ensure that the most up to date information gets to the reader while the features are still fresh and current. I would like to thank Kelly Talbot for his great project management and technical review of the content as well as John Mueller for his technical review and suggestions throughout the process. Lastly, I would like to thank all the numerous people behind the scenes who helped get this book together.

# CREDITS

**Senior Acquisitions Editor**  
Kenyon Brown

**Project Editor**  
Kelly Talbot

**Technical Editor**  
John Mueller

**Production Editor**  
Joel Jones

**Copy Editor**  
Kelly Talbot Editing Services

**Manager of Content Development & Assembly**  
Mary Beth Wakefield

**Production Manager**  
Kathleen Wisor

**Marketing Director**  
David Mayhew

**Marketing Manager**  
Carrie Sherrill

**Professional Technology & Strategy Director**  
Barry Pruett

**Business Manager**  
Amy Kries

**Associate Publisher**  
Jim Minatel

**Project Coordinator, Cover**  
Brent Savage

**Proofreader**  
Nancy Bell

**Indexer**  
Jack Lewis

**Cover Designer**  
Wiley

**Cover Image**  
© Nomad\_Soul/Shutterstock

# CONTENTS

## INTRODUCTION

xix

## PART I: THE OOP LANGUAGE

<b>CHAPTER 1: INTRODUCING C#</b>	<b>3</b>
What Is the .NET Framework?	4
What's in the .NET Framework?	4
Writing Applications Using the .NET Framework	5
What Is C#?	8
Applications You Can Write with C#	9
C# in this Book	10
Visual Studio 2015	10
Visual Studio Express 2015 Products	10
Solutions	11
<b>CHAPTER 2: WRITING A C# PROGRAM</b>	<b>13</b>
The Visual Studio 2015 Development Environment	14
Console Applications	17
The Solution Explorer	20
The Properties Window	21
The Error List Window	22
Desktop Applications	22
<b>CHAPTER 3: VARIABLES AND EXPRESSIONS</b>	<b>29</b>
Basic C# Syntax	30
Basic C# Console Application Structure	33
Variables	34
Simple Types	34
Variable Naming	39
Literal Values	39
Expressions	42
Mathematical Operators	42
Assignment Operators	47
Operator Precedence	48
Namespaces	49

<b>CHAPTER 4: FLOW CONTROL</b>	<b>53</b>
<b>Boolean Logic</b>	<b>54</b>
Boolean Bitwise and Assignment Operators	56
Operator Precedence Updated	58
<b>Branching</b>	<b>59</b>
The Ternary Operator	59
The if Statement	59
The switch Statement	63
<b>Looping</b>	<b>66</b>
do Loops	66
while Loops	69
for Loops	71
Interrupting Loops	72
Infinite Loops	73
<b>CHAPTER 5: MORE ABOUT VARIABLES</b>	<b>77</b>
<b>Type Conversion</b>	<b>78</b>
Implicit Conversions	78
Explicit Conversions	80
Explicit Conversions Using the Convert Commands	83
<b>Complex Variable Types</b>	<b>85</b>
Enumerations	85
Structs	89
Arrays	92
<b>String Manipulation</b>	<b>99</b>
<b>CHAPTER 6: FUNCTIONS</b>	<b>107</b>
<b>Defining and Using Functions</b>	<b>108</b>
Return Values	110
Parameters	112
<b>Variable Scope</b>	<b>119</b>
Variable Scope in Other Structures	122
Parameters and Return Values versus Global Data	123
<b>The Main() Function</b>	<b>125</b>
<b>Struct Functions</b>	<b>127</b>
<b>Overloading Functions</b>	<b>128</b>
<b>Using Delegates</b>	<b>130</b>

<b>CHAPTER 7: DEBUGGING AND ERROR HANDLING</b>	<b>135</b>
Debugging in Visual Studio	136
Debugging in Nonbreak (Normal) Mode	136
Debugging in Break Mode	144
Error Handling	153
try...catch...finally	153
Listing and Configuring Exceptions	160
<b>CHAPTER 8: INTRODUCTION TO OBJECT-ORIENTED PROGRAMMING</b>	<b>163</b>
What Is Object-Oriented Programming?	164
What Is an Object?	165
Everything's an Object	168
The Life Cycle of an Object	168
Static and Instance Class Members	169
OOP Techniques	170
Interfaces	171
Inheritance	172
Polymorphism	175
Relationships between Objects	177
Operator Overloading	179
Events	180
Reference Types versus Value Types	180
OOP in Desktop Applications	180
<b>CHAPTER 9: DEFINING CLASSES</b>	<b>187</b>
Class Definitions in C#	188
Interface Definitions	190
System.Object	193
Constructors and Destructors	195
Constructor Execution Sequence	196
OOP Tools in Visual Studio	200
The Class View Window	200
The Object Browser	202
Adding Classes	203
Class Diagrams	204

Class Library Projects	206
Interfaces versus Abstract Classes	209
Struct Types	212
Shallow Copying versus Deep Copying	214
<b>CHAPTER 10: DEFINING CLASS MEMBERS</b>	<b>217</b>
Member Definitions	218
Defining Fields	218
Defining Methods	219
Defining Properties	220
Refactoring Members	225
Automatic Properties	226
Additional Class Member Topics	227
Hiding Base Class Methods	227
Calling Overridden or Hidden Base Class Methods	229
Using Nested Type Definitions	230
Interface Implementation	232
Implementing Interfaces in Classes	233
Partial Class Definitions	235
Partial Method Definitions	237
Example Application	238
Planning the Application	238
Writing the Class Library	239
A Client Application for the Class Library	246
The Call Hierarchy Window	248
<b>CHAPTER 11: COLLECTIONS, COMPARISONS, AND CONVERSIONS</b>	<b>251</b>
Collections	252
Using Collections	253
Defining Collections	258
Indexers	259
Adding a Cards Collection to CardLib	262
Keyed Collections and IDictionary	264
Iterators	266
Iterators and Collections	270
Deep Copying	271
Adding Deep Copying to CardLib	273
Comparisons	275

Type Comparisons	275
Value Comparisons	279
<b>Conversions</b>	<b>295</b>
Overloading Conversion Operators	295
The as Operator	297
<b>CHAPTER 12: GENERICS</b>	<b>301</b>
What Are Generics?	302
Using Generics	303
Nullable Types	303
The System.Collections.Generic Namespace	311
Defining Generic Types	321
Defining Generic Classes	322
Defining Generic Interfaces	332
Defining Generic Methods	333
Defining Generic Delegates	334
Variance	335
Covariance	336
Contravariance	336
<b>CHAPTER 13: ADDITIONAL C# TECHNIQUES</b>	<b>341</b>
The :: Operator and the Global Namespace Qualifier	342
Custom Exceptions	343
Adding Custom Exceptions to CardLib	343
Events	345
What Is an Event?	345
Handling Events	347
Defining Events	350
Expanding and Using CardLib	357
Attributes	365
Reading Attributes	366
Creating Attributes	367
Initializers	368
Object Initializers	368
Collection Initializers	371
Type Inference	374
Anonymous Types	376
Dynamic Lookup	380
The dynamic Type	381

Advanced Method Parameters	384
Optional Parameters	385
Named Parameters	386
Lambda Expressions	391
Anonymous Methods Recap	391
Lambda Expressions for Anonymous Methods	392
Lambda Expression Parameters	396
Lambda Expression Statement Bodies	396
Lambda Expressions as Delegates and Expression Trees	398
Lambda Expressions and Collections	399

**PART II: WINDOWS PROGRAMMING**

<b>CHAPTER 14: BASIC DESKTOP PROGRAMMING</b>	<b>407</b>
XAML	408
Separation of Concerns	409
XAML in Action	409
The Playground	411
WPF Controls	412
Properties	413
Events	417
Control Layout	422
Stack Order	423
Alignment, Margins, Padding, and Dimensions	423
Border	424
Canvas	424
DockPanel	426
StackPanel	428
WrapPanel	429
Grid	430
The Game Client	433
The About Window	433
The Options Window	439
Data Binding	448
Starting a Game with the ListBox Control	453
<b>CHAPTER 15: ADVANCED DESKTOP PROGRAMMING</b>	<b>461</b>
The Main Window	462
The Menu Control	462
Routed Commands with Menus	462

Creating and Styling Controls	466
Styles	467
Templates	467
Value Converters	472
Triggers	473
Animations	475
WPF User Controls	478
Implementing Dependency Properties	478
Putting It All Together	489
Refactoring the Domain Model	489
The View Models	494
Completing the Game	502

## PART III: CLOUD PROGRAMMING

<b>CHAPTER 16: BASIC CLOUD PROGRAMMING</b>	<b>515</b>
The Cloud, Cloud Computing, and the Cloud Optimized Stack	516
Cloud Patterns and Best Practices	519
Using Microsoft Azure C# Libraries to Create a Storage Container	520
Creating an ASP.NET 4.6 Web Site That Uses the Storage Container	530
<b>CHAPTER 17: ADVANCED CLOUD PROGRAMMING AND DEPLOYMENT</b>	<b>539</b>
Creating an ASP.NET Web API	540
Deploying and Consuming an ASP.NET Web API on Microsoft Azure	544
Scaling an ASP.NET Web API on Microsoft Azure	551

## PART IV: DATA ACCESS

<b>CHAPTER 18: FILES</b>	<b>561</b>
File Classes for Input and Output	562
The File and Directory Classes	563
The FileInfo Class	564
The DirectoryInfo Class	566
Path Names and Relative Paths	566

<b>Streams</b>	<b>567</b>
Classes for Using Streams	567
The FileStream Object	568
The StreamWriter Object	575
The StreamReader Object	577
Asynchronous File Access	580
Reading and Writing Compressed Files	580
<b>Monitoring the File System</b>	<b>584</b>
<b>CHAPTER 19: XML AND JSON</b>	<b>593</b>
<b>XML Basics</b>	<b>594</b>
<b>JSON Basics</b>	<b>594</b>
<b>XML Schemas</b>	<b>595</b>
<b>XML Document Object Model</b>	<b>597</b>
The XmlDocument Class	598
TheXmlElement Class	598
Changing the Values of Nodes	603
<b>Converting XML to JSON</b>	<b>609</b>
<b>Searching XML with XPath</b>	<b>611</b>
<b>CHAPTER 20: LINQ</b>	<b>617</b>
<b>LINQ to XML</b>	<b>618</b>
LINQ to XML Functional Constructors	618
Working with XML Fragments	621
<b>LINQ Providers</b>	<b>624</b>
<b>LINQ Query Syntax</b>	<b>624</b>
Declaring a Variable for Results Using the var Keyword	626
Specifying the Data Source: from Clause	627
Specify Condition: where Clause	627
Selecting Items: select Clause	627
Finishing Up: Using the foreach Loop	628
Deferred Query Execution	628
<b>LINQ Method Syntax</b>	<b>628</b>
LINQ Extension Methods	629
Query Syntax versus Method Syntax	629
Lambda Expressions	630
<b>Ordering Query Results</b>	<b>632</b>
<b>Understanding the orderby Clause</b>	<b>633</b>
<b>Querying a Large Data Set</b>	<b>634</b>
<b>Using Aggregate Operators</b>	<b>636</b>

Using the Select Distinct Query	640
Ordering by Multiple Levels	642
Using Group Queries	644
Using Joins	646
<b>CHAPTER 21: DATABASES</b>	<b>651</b>
Using Databases	651
Installing SQL Server Express	652
Entity Framework	652
A Code First Database	653
But Where Is My Database?	660
Navigating Database Relationships	661
Handling Migrations	668
Creating and Querying XML from an Existing Database	669
<b>PART V: ADDITIONAL TECHNIQUES</b>	
<b>CHAPTER 22: WINDOWS COMMUNICATION FOUNDATION</b>	<b>677</b>
What Is WCF?	678
WCF Concepts	679
WCF Communication Protocols	679
Addresses, Endpoints, and Bindings	680
Contracts	682
Message Patterns	683
Behaviors	683
Hosting	683
WCF Programming	684
The WCF Test Client	690
Defining WCF Service Contracts	693
Self-Hosted WCF Services	700
<b>CHAPTER 23: UNIVERSAL APPS</b>	<b>709</b>
Getting Started	709
Universal Apps	710
App Concepts and Design	711
Screen Orientation	711
Menus and Toolbars	711
Tiles and Badges	712

App Lifetime	712
Lock Screen Apps	712
<b>App Development</b>	<b>712</b>
Adaptive Displays	713
Sandboxed Apps	721
Navigation between Pages	725
The CommandBar Control	728
Managing State	729
<b>Common Elements of Windows Store Apps</b>	<b>732</b>
<b>The Windows Store</b>	<b>733</b>
Packaging an App	733
Creating the Package	734
<b>APPENDIX: EXERCISE SOLUTIONS</b>	<b>737</b>
<b>INDEX</b>	<b>781</b>

# INTRODUCTION

**C# IS A RELATIVELY NEW LANGUAGE** that was unveiled to the world when Microsoft announced the first version of its .NET Framework in July 2000. Since then its popularity has rocketed, and it has arguably become the language of choice for desktop, web, and cloud developers who use the .NET Framework. Part of the appeal of C# comes from its clear syntax, which derives from C/C++ but simplifies some things that have previously discouraged some programmers. Despite this simplification, C# has retained the power of C++, and there is now no reason not to move into C#. The language is not difficult and it's a great one to learn elementary programming techniques with. This ease of learning, combined with the capabilities of the .NET Framework, make C# an excellent way to start your programming career.

The latest release of C#, C# 6, which is included with version 4.6 of the .NET Framework, builds on the existing successes and adds even more attractive features. The latest release of Visual Studio (Visual Studio 2015) and the Visual Studio Express/Community 2015 line of development tools also bring many tweaks and improvements to make your life easier and to dramatically increase your productivity.

This book is intended to teach you about all aspects of C# programming, including the language itself, desktop and cloud programming, making use of data sources, and some new and advanced techniques. You'll also learn about the capabilities of Visual Studio 2015 and all the ways that this product can aid your application development.

The book is written in a friendly, mentor-style fashion, with each chapter building on previous ones, and every effort is made to ease you into advanced techniques painlessly. At no point will technical terms appear from nowhere to discourage you from continuing; every concept is introduced and discussed as required. Technical jargon is kept to a minimum; but where it is necessary, it, too, is properly defined and laid out in context.

The authors of this book are all experts in their field and are all enthusiastic in their passion for both the C# language and the .NET Framework. Nowhere will you find a group of people better qualified to take you under their collective wing and nurture your understanding of C# from first principles to advanced techniques. Along with the fundamental knowledge it provides, this book is packed full of helpful hints, tips, exercises, and full-fledged example code (available for download at [p2p.wrox.com](http://p2p.wrox.com)) that you will find yourself returning to repeatedly as your career progresses.

We pass this knowledge on without begrudging it and hope that you will be able to use it to become the best programmer you can be. Good luck, and all the best!

## WHO THIS BOOK IS FOR

This book is for everyone who wants to learn how to program in C# using the .NET Framework. It is for absolute beginners who want to give programming a try by learning a clean, modern, elegant programming language. But it is also for people familiar with other programming languages who want to explore the .NET platform, as well as for existing .NET developers who want to give Microsoft's .NET flagship language a try.

## WHAT THIS BOOK COVERS

The early chapters cover the language itself, assuming no prior programming experience. If you have programmed in other languages before, much of the material in these chapters will be familiar. Many aspects of C# syntax are shared with other languages, and many structures are common to practically all programming languages (such as looping and branching structures). However, even if you are an experienced programmer, you will benefit from looking through these chapters to learn the specifics of how these techniques apply to C#.

If you are new to programming, you should start from the beginning, where you will learn basic programming concepts and become acquainted with both C# and the .NET platform that underpins it. If you are new to the .NET Framework but know how to program, you should read Chapter 1 and then skim through the next few chapters before continuing with the application of the C# language. If you know how to program but haven't encountered an object-oriented programming language before, you should read the chapters from Chapter 8 onward.

Alternatively, if you already know the C# language, you might want to concentrate on the chapters dealing with the most recent .NET Framework and C# language developments, specifically the chapters on collections, generics, and C# language enhancements (Chapters 11 to 13), or skip the first section of the book completely and start with Chapter 14.

The chapters in this book have been written with a dual purpose in mind: They can be read sequentially to provide a complete tutorial in the C# language, and they can be dipped into as required reference material.

In addition to the core material, starting with Chapter 3 each chapter also includes a selection of exercises at the end, which you can work through to ensure that you have understood the material. The exercises range from simple multiple choice or true/false questions to more complex exercises that require you to modify or build applications. The answers to all the exercises are provided in Appendix A. You can also find these exercises as part of the [www.wrox.com](http://www.wrox.com) code downloads for this book at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming).

This book also gives plenty of love and attention to coincide with the release of C# 6 and .NET 4.6. Every chapter received an overhaul, with less relevant material removed, and new material added. All of the code has been tested against the latest version of the development tools used, and all of the screenshots have been retaken in Windows 8.1/10 to provide the most current windows and dialog boxes.

New highlights of this edition include the following:

- Additional and improved code examples for you to try out
- Coverage of everything that's new in C# 6 and .NET 4.6, including how to create Universal Windows Apps
- Examples of programming cloud applications and using Azure SDK to create and access cloud resources

## HOW THIS BOOK IS STRUCTURED

This book is divided into six sections:

- **Introduction** — Purpose and general outline of the book's contents
- **The C# Language** — Covers all aspects of the C# language, from the fundamentals to object-oriented techniques
- **Windows Programming** — How to write and deploy desktop applications with the Windows Presentation Foundation library (WPF)
- **Cloud Programming** — Cloud application development and deployment, including the creation and consumption of a Web API
- **Data Access** — How to use data in your applications, including data stored in files on your hard disk, data stored in XML format, and data in databases
- **Additional Techniques** — An examination of some extra ways to use C# and the .NET Framework, including Windows Communication Foundation (WCF) and Universal Windows Applications

The following sections describe the chapters in the five major parts of this book.

### The C# Language (Chapters 1–13)

Chapter 1 introduces you to C# and how it fits into the .NET landscape. You'll learn the fundamentals of programming in this environment and how Visual Studio 2015 (VS) fits in.

Chapter 2 starts you off with writing C# applications. You'll look at the syntax of C# and put the language to use with sample command-line and Windows applications. These examples demonstrate just how quick and easy it can be to get up and running, and along the way you'll be introduced to the Visual Studio development environment and the basic windows and tools that you'll be using throughout the book.

Next you'll learn more about the basics of the C# language. You'll learn what variables are and how to manipulate them in Chapter 3. You'll enhance the structure of your applications with flow control (looping and branching) in Chapter 4, and you'll see some more advanced variable types

such as arrays in Chapter 5. In Chapter 6 you'll start to encapsulate your code in the form of functions, which makes it much easier to perform repetitive operations and makes your code much more readable.

By the beginning of Chapter 7 you'll have a handle on the fundamentals of the C# language, and you will focus on debugging your applications. This involves looking at outputting trace information as your applications are executed, and at how Visual Studio can be used to trap errors and lead you to solutions for them with its powerful debugging environment.

From Chapter 8 onward you'll learn about object-oriented programming (OOP), starting with a look at what this term means and an answer to the eternal question, "What is an object?" OOP can seem quite difficult at first. The whole of Chapter 8 is devoted to demystifying it and explaining what makes it so great, and you won't actually deal with much C# code until the very end of the chapter.

Everything changes in Chapter 9, when you put theory into practice and start using OOP in your C# applications. This is where the true power of C# lies. You'll start by looking at how to define classes and interfaces, and then move on to class members (including fields, properties, and methods) in Chapter 10. At the end of that chapter you'll start to assemble a card game application, which is developed over several chapters and will help to illustrate OOP.

Once you've learned how OOP works in C#, Chapter 11 moves on to look at common OOP scenarios, including dealing with collections of objects, and comparing and converting objects. Chapter 12 takes a look at a very useful feature of C# that was introduced in .NET 2.0: generics, which enable you to create very flexible classes. Next, Chapter 13 continues the discussion of the C# language and OOP with some additional techniques, notably events, which become very important in, for example, Windows programming. Chapter 13 wraps up the fundamentals by focusing on C# language features that were introduced with versions 3.0, 4, 5, and 6 of the language.

## **Windows Programming (Chapters 14–15)**

Chapter 14 starts by introducing you to what is meant by Windows programming and looks at how this is achieved in Visual Studio. It focuses on WPF as a tool that enables you to build desktop applications in a graphical way and assemble advanced applications with the minimum of effort and time. You'll start with the basics of WPF programming and build up your knowledge in both this chapter and Chapter 15, which demonstrates how you can use the wealth of controls supplied by the .NET Framework in your applications.

## **Cloud Programming (Chapters 16–17)**

Chapter 16 starts by describing what cloud programming is and discusses the cloud optimized stack. The cloud environment is not identical to the way programs have been traditionally coded, so a few cloud programming patterns are discussed and defined. To complete this chapter, you require an Azure account, which is free, so that you can create an App Services Web App, then using the Azure SDK with C#, you create and access a storage account from an ASP.NET 4.6 web application.

In Chapter 17, you learn how to create and deploy an ASP.NET Web API to the cloud and then consume the Web API from a similar ASP.NET 4.6 web application. The chapter ends discussing two of the most valuable features in the cloud, scaling and the optimal utilization of hardware resources.

## Data Access (Chapters 18–21)

Chapter 18 looks at how your applications can save and retrieve data to disk, both as simple text files and as more complex representations of data. You'll also learn how to compress data and how to monitor and act on file system changes.

In Chapter 19 you'll learn about the de facto standard for data exchange — namely, XML — and a rapidly emerging format called JSON. By this point in the book, you'll have touched on XML several times in preceding chapters, but this chapter lays out the ground rules and shows you what all the excitement is about.

The remainder of this part looks at LINQ, which is a query language built in to the latest versions of the .NET Framework. You start in Chapter 20 with a general introduction to LINQ, and then you will use LINQ to access a database and other data in Chapter 21.

## Additional Techniques (Chapters 22–23)

Chapter 22 is an introduction to Windows Communication Foundation (WCF), which provides you with the tools you need for enterprise-level programmatic access to information and capabilities across local networks and the Internet. You will see how you can use WCF to expose complex data and functionality to web and desktop applications in a platform-independent way.

Chapter 23 shows you how you can create Universal Windows Apps, which are new to Windows. This chapter builds on the foundation of Chapters 14 and 15 to show you how to create Windows Apps that can run on all windows platforms.

## WHAT YOU NEED TO USE THIS BOOK

The code and descriptions of C# and the .NET Framework in this book apply to C# 6 and .NET 4.6. You don't need anything other than the Framework to understand this aspect of the book, but many of the examples require a development tool. This book uses Visual Studio 2015 as its primary development tool; however, if you don't have this, you will be able to use the free Visual Studio Express/Community 2015 line of products. For the first part of the book, Visual Studio Express/Community 2012 for Windows Desktop will enable you to create desktop and console applications. For later chapters, you may also use Visual Studio Express/Community 2015 for Windows 10 in order to create Universal Windows Apps, Visual Studio Express/Community 2015 for Cloud to create cloud applications, and SQL Server Express 2014 for applications that access databases. Some functionality is available only in Visual Studio 2015, but this won't stop you from working through any of the examples in this book.

The source code for the samples is available for download from the Wrox website at:

[www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming)

## CONVENTIONS

To help you get the most from the text and keep track of what's happening, we've used a number of conventions throughout the book.

### TRY IT OUT

The *Try It Out* is an exercise you should work through, following the text in the book.

1. They usually consist of a set of steps.
2. Each step has a number.
3. Follow the steps through with your copy of the database.

### How It Works

After each *Try It Out*, the code you've typed will be explained in detail.

**WARNING** Warnings hold important, not-to-be-forgotten information that is directly relevant to the surrounding text.

**NOTE** Notes indicates notes, tips, hints, tricks, or and asides to the current discussion.

As for styles in the text:

- We *highlight* new terms and important words when we introduce them.
- We show keyboard strokes like this: Ctrl+A.
- We show file names, URLs, and code within the text like so: `persistence.properties`.

We present code in two different ways:

We use a monofont type with no highlighting for most code examples.

We use bold to emphasize code that is particularly important in the present context or to show changes from a previous code snippet.

## SOURCE CODE

As you work through the examples in this book, you may choose either to type in all the code manually, or to use the source code files that accompany the book. All the source code used in this book is available for download at [www.wrox.com](http://www.wrox.com). Specifically for this book, the code download is on the Download Code tab at:

[www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming)

You can also search for the book at [www.wrox.com](http://www.wrox.com) by ISBN (the ISBN for this book is 978-1-119-09668-9) to find the code. And a complete list of code downloads for all current Wrox books is available at [www.wrox.com/dynamic/books/download.aspx](http://www.wrox.com/dynamic/books/download.aspx).

Most of the code on [www.wrox.com](http://www.wrox.com) is compressed in a .ZIP, .RAR archive or similar archive format appropriate to the platform. Once you download the code, just decompress it with an appropriate compression tool.

**NOTE** Because many books have similar titles, you may find it easiest to search by ISBN; this book's ISBN is 978-1-119-09668-9.

Alternately, you can go to the main Wrox code download page at [www.wrox.com/dynamic/books/download.aspx](http://www.wrox.com/dynamic/books/download.aspx) to see the code available for this book and all other Wrox books.

## ERRATA

We make every effort to ensure that there are no errors in the text or in the code. However, no one is perfect, and mistakes do occur. If you find an error in one of our books, like a spelling mistake or faulty piece of code, we would be very grateful for your feedback. By sending in errata, you may save another reader hours of frustration, and at the same time, you will be helping us provide even higher quality information.

To find the errata page for this book, go to

[www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming)

And click the Errata link. On this page you can view all errata that has been submitted for this book and posted by Wrox editors.

If you don't spot "your" error on the Book Errata page, go to [www.wrox.com/contact/techsupport.shtml](http://www.wrox.com/contact/techsupport.shtml) and complete the form there to send us the error you have found. We'll check the information and, if appropriate, post a message to the book's errata page and fix the problem in subsequent editions of the book.

## P2P.WROX.COM

For author and peer discussion, join the P2P forums at <http://p2p.wrox.com>. The forums are a Web-based system for you to post messages relating to Wrox books and related technologies and interact with other readers and technology users. The forums offer a subscription feature to e-mail you topics of interest of your choosing when new posts are made to the forums. Wrox authors, editors, other industry experts, and your fellow readers are present on these forums.

At <http://p2p.wrox.com>, you will find a number of different forums that will help you, not only as you read this book, but also as you develop your own applications. To join the forums, just follow these steps:

1. Go to <http://p2p.wrox.com> and click the Register link.
2. Read the terms of use and click Agree.
3. Complete the required information to join, as well as any optional information you wish to provide, and click Submit.
4. You will receive an e-mail with information describing how to verify your account and complete the joining process.

**NOTE** *You can read messages in the forums without joining P2P, but in order to post your own messages, you must join.*

Once you join, you can post new messages and respond to messages other users post. You can read messages at any time on the Web. If you would like to have new messages from a particular forum e-mailed to you, click the Subscribe to this Forum icon by the forum name in the forum listing.

For more information about how to use the Wrox P2P, be sure to read the P2P FAQs for answers to questions about how the forum software works, as well as many common questions specific to P2P and Wrox books. To read the FAQs, click the FAQ link on any P2P page.

# PART I

## The OOP Language

---

- **CHAPTER 1:** Introducing C#
- **CHAPTER 2:** Writing a C# Program
- **CHAPTER 3:** Variables and Expressions
- **CHAPTER 4:** Flow Control
- **CHAPTER 5:** More about Variables
- **CHAPTER 6:** Functions
- **CHAPTER 7:** Debugging and Error Handling
- **CHAPTER 8:** Introduction to Object-Oriented Programming
- **CHAPTER 9:** Defining Classes
- **CHAPTER 10:** Defining Class Members
- **CHAPTER 11:** Collections, Comparisons, and Conversions
- **CHAPTER 12:** Generics
- **CHAPTER 13:** Additional C# Techniques



# 1

## Introducing C#

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Exploring the .NET Framework
- Learning how .NET applications work
- Exploring C# and how it relates to the .NET Framework
- Discovering tools for creating .NET applications with C#

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

The wrox.com code downloads for this chapter are found at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 1 download and individually named according to the names throughout the chapter.

Welcome to the first chapter of the first section of this book. This section provides you with the basic knowledge you need to get up and running with the most recent version of C#. Specifically, this chapter provides an overview of C# and the .NET Framework, including what these technologies are, the motivation for using them, and how they relate to each other.

It begins with a general discussion of the .NET Framework. This technology contains many concepts that are tricky to come to grips with initially. This means that the discussion, by necessity, covers many concepts in a short amount of space. However, a quick look at the basics is essential to understanding how to program in C#. Later in the book, you revisit many of the topics covered here, exploring them in more detail.

After that general introduction, the chapter provides a basic description of C# itself, including its origins and similarities to C++. Finally, you look at the primary tool used throughout this book: Visual Studio 2015 (VS). Visual Studio 2015 is the latest in a long line of development environments that Microsoft has produced, and it includes all sorts of features (including full support for Windows Store applications) that you will learn about throughout this book.

## WHAT IS THE .NET FRAMEWORK?

The .NET Framework (now at version 4.6) is a revolutionary platform created by Microsoft for developing applications. The most interesting thing about this statement is how vague it is — but there are good reasons for this. For a start, note that it doesn't "develop applications on the Windows operating system." Although the Microsoft release of the .NET Framework runs on the Windows and Windows Phone operating systems, it is possible to find alternative versions that will work on other systems. One example of this is Mono, an open-source version of the .NET Framework (including a C# compiler) that runs on several operating systems, including various flavors of Linux and Mac OS; you can read more about it at <http://www.mono-project.com>. There are also variants of Mono that run on iPhone (MonoTouch) and Android (Mono for Android, a.k.a. MonoDroid) smartphones. One of the key motivations behind the .NET Framework is its intended use as a means of integrating disparate operating systems.

In addition, the preceding definition of the .NET Framework includes no restriction on the type of applications that are possible. That's because there is no restriction — the .NET Framework enables the creation of desktop applications, Windows Store applications, cloud/web applications, Web APIs, and pretty much anything else you can think of. Also, with web, cloud and Web API applications it's worth noting that these are, by definition, multi-platform applications, since any system with a web browser can access them.

The .NET Framework has been designed so that it can be used from any language, including C# (the subject of this book) as well as C++, Visual Basic, JScript, and even older languages such as COBOL. For this to work, .NET-specific versions of these languages have also appeared, and more are being released all the time. For a list of languages, see this site <https://msdn.microsoft.com/library/aa292164.aspx>. Not only do all of these have access to the .NET Framework, but they can also communicate with each other. It is possible for C# developers to make use of code written by Visual Basic programmers, and vice versa.

All of this provides an extremely high level of versatility and is part of what makes using the .NET Framework such an attractive prospect.

## What's in the .NET Framework?

The .NET Framework consists primarily of a gigantic library of code that you use from your client languages (such as C#) using object-oriented programming (OOP) techniques. This library is categorized into different modules — you use portions of it depending on the results you want to achieve. For example, one module contains the building blocks for Windows applications, another for network programming, and another for web development. Some modules are divided into more specific submodules, such as a module for building web services within the module for web development.

The intention is for different operating systems to support some or all of these modules, depending on their characteristics. A smartphone, for example, includes support for all the core .NET functionality but is unlikely to require some of the more esoteric modules.

Part of the .NET Framework library defines some basic *types*. A type is a representation of data, and specifying some of the most fundamental of these (such as "a 32-bit signed integer") facilitates

---

interoperability between languages using the .NET Framework. This is called the *Common Type System* (CTS).

As well as supplying this library, the .Net Framework also includes the .NET *Common Language Runtime* (CLR), which is responsible for the execution of all applications developed using the .NET library.

## Writing Applications Using the .NET Framework

Writing an application using the .NET Framework means writing code (using any of the languages that support the Framework) using the .NET code library. In this book you use Visual Studio for your development. Visual Studio is a powerful, integrated development environment that supports C# (as well as managed and unmanaged C++, Visual Basic, and some others). The advantage of this environment is the ease with which .NET features can be integrated into your code. The code that you create will be entirely C# but use the .NET Framework throughout, and you'll make use of the additional tools in Visual Studio where necessary.

In order for C# code to execute, it must be converted into a language that the target operating system understands, known as *native code*. This conversion is called *compiling* code, an act that is performed by a *compiler*. Under the .NET Framework, this is a two-stage process.

### CIL and JIT

When you compile code that uses the .NET Framework library, you don't immediately create operating system-specific native code. Instead, you compile your code into *Common Intermediate Language* (CIL) code. This code isn't specific to any operating system (OS) and isn't specific to C#. Other .NET languages — Visual Basic .NET, for example — also compile to this language as a first stage. This compilation step is carried out by Visual Studio when you develop C# applications.

Obviously, more work is necessary to execute an application. That is the job of a *just-in-time* (JIT) compiler, which compiles CIL into native code that is specific to the OS and machine architecture being targeted. Only at this point can the OS execute the application. The *just-in-time* part of the name reflects the fact that CIL code is compiled only when it is needed. This compilation can happen on the fly while your application is running, although luckily this isn't something that you normally need to worry about as a developer. Unless you are writing extremely advanced code where performance is critical, it's enough to know that this compilation process will churn along merrily in the background, without interfering.

In the past, it was often necessary to compile your code into several applications, each of which targeted a specific operating system and CPU architecture. Typically, this was a form of optimization (to get code to run faster on an AMD chipset, for example), but at times it was critical (for applications to work in both Win9x and WinNT/2000 environments, for example). This is now unnecessary, because JIT compilers (as their name suggests) use CIL code, which is independent of the machine, operating system, and CPU. Several JIT compilers exist, each targeting a different architecture, and the CLR uses the appropriate one to create the native code required.

The beauty of all this is that it requires a lot less work on your part — in fact, you can forget about system-dependent details and concentrate on the more interesting functionality of your code.

**NOTE** You might come across references to Microsoft Intermediate Language (MSIL) or just IL. MSIL was the original name for CIL, and many developers still use this terminology.

## Assemblies

When you compile an application, the CIL code is stored in an *assembly*. Assemblies include both executable application files that you can run directly from Windows without the need for any other programs (these have a .exe file extension) and libraries (which have a .dll extension) for use by other applications.

In addition to containing CIL, assemblies also include *meta* information (that is, information about the information contained in the assembly, also known as *metadata*) and optional *resources* (additional data used by the CIL, such as sound files and pictures). The meta information enables assemblies to be fully self-descriptive. You need no other information to use an assembly, meaning you avoid situations such as failing to add required data to the system registry and so on, which was often a problem when developing with other platforms.

This means that deploying applications is often as simple as copying the files into a directory on a remote computer. Because no additional information is required on the target systems, you can just run an executable file from this directory and (assuming the .NET CLR is installed) you're good to go.

Of course, you won't necessarily want to include everything required to run an application in one place. You might write some code that performs tasks required by multiple applications. In situations like that, it is often useful to place the reusable code in a place accessible to all applications. In the .NET Framework, this is the *global assembly cache* (GAC). Placing code in the GAC is simple — you just place the assembly containing the code in the directory containing this cache.

## Managed Code

The role of the CLR doesn't end after you have compiled your code to CIL and a JIT compiler has compiled that to native code. Code written using the .NET Framework is *managed* when it is executed (a stage usually referred to as *runtime*). This means that the CLR looks after your applications by managing memory, handling security, allowing cross-language debugging, and so on. By contrast, applications that do not run under the control of the CLR are said to be *unmanaged*, and certain languages such as C++ can be used to write such applications, which, for example, access low-level functions of the operating system. However, in C# you can write only code that runs in a managed environment. You will make use of the managed features of the CLR and allow .NET itself to handle any interaction with the operating system.

## Garbage Collection

One of the most important features of managed code is the concept of *garbage collection*. This is the .NET method of making sure that the memory used by an application is freed up completely when the application is no longer in use. Prior to .NET this was mostly the responsibility of programmers, and a few simple errors in code could result in large blocks of memory mysteriously disappearing as a result of being allocated to the wrong place in memory. That usually meant a progressive slowdown of your computer, followed by a system crash.

.NET garbage collection works by periodically inspecting the memory of your computer and removing anything from it that is no longer needed. There is no set time frame for this; it might happen thousands of times a second, once every few seconds, or whenever, but you can rest assured that it will happen.

There are some implications for programmers here. Because this work is done for you at an unpredictable time, applications have to be designed with this in mind. Code that requires a lot of memory to run should tidy itself up, rather than wait for garbage collection to happen, but that isn't as tricky as it sounds.

## Fitting It Together

Before moving on, let's summarize the steps required to create a .NET application as discussed previously:

1. Application code is written using a .NET-compatible language such as C# (see Figure 1-1).
2. That code is compiled into CIL, which is stored in an assembly (see Figure 1-2).




FIGURE 1-1




FIGURE 1-2

3. When this code is executed (either in its own right if it is an executable or when it is used from other code), it must first be compiled into native code using a JIT compiler (see Figure 1-3).




FIGURE 1-3

4. The native code is executed in the context of the managed CLR, along with any other running applications or processes, as shown in Figure 1-4.




FIGURE 1-4

## Linking

Note one additional point concerning this process. The C# code that compiles into CIL in step 2 needn't be contained in a single file. It's possible to split application code across multiple source-code files, which are then compiled together into a single assembly. This extremely useful process is known as *linking*. It is required because it is far easier to work with several smaller files than one enormous one. You can separate logically related code into an individual file so that it can be worked on independently and then practically forgotten about when completed. This also makes it easy to locate specific pieces of code when you need them and enables teams of developers to divide the programming burden into manageable chunks, whereby individuals can "check out" pieces of code to work on without risking damage to otherwise satisfactory sections or sections other people are working on.

## WHAT IS C#?

C#, as mentioned earlier, is one of the languages you can use to create applications that will run in the .NET CLR. It is an evolution of the C and C++ languages and has been created by Microsoft specifically to work with the .NET platform. The C# language has been designed to incorporate many of the best features from other languages, while clearing up their problems.

Developing applications using C# is simpler than using C++, because the language syntax is simpler. Still, C# is a powerful language, and there is little you might want to do in C++ that you can't do in C#. Having said that, those features of C# that parallel the more advanced features of C++, such as directly accessing and manipulating system memory, can be carried out only by using code marked as *unsafe*. This advanced programmatic technique is potentially dangerous (hence its name) because it is possible to overwrite system-critical blocks of memory with potentially catastrophic results. For this reason, and others, this book does not cover that topic.

At times, C# code is slightly more verbose than C++. This is a consequence of C# being a *typesafe* language (unlike C++). In layperson's terms, this means that once some data has been assigned to a type, it cannot subsequently transform itself into another unrelated type. Consequently, strict rules must be adhered to when converting between types, which means you will often need to write more code to carry out the same task in C# than you might write in C++. However, there are benefits to this — the code is more robust, debugging is simpler, and .NET can always track the type of a piece of data at any time. In C#, you therefore might not be able to do things such as “take the region of memory 4 bytes into this data and 10 bytes long and interpret it as X,” but that’s not necessarily a bad thing.

C# is just one of the languages available for .NET development, but it is certainly the best. It has the advantage of being the only language designed from the ground up for the .NET Framework and is the principal language used in versions of .NET that are ported to other operating systems. To keep languages such as the .NET version of Visual Basic as similar as possible to their predecessors yet compliant with the CLR, certain features of the .NET code library are not fully supported, or at least require unusual syntax.

By contrast, C# can make use of every feature that the .NET Framework code library has to offer. Also, each new version of .NET has included additions to the C# language, partly in response to requests from developers, making it even more powerful.

## Applications You Can Write with C#

The .NET Framework has no restrictions on the types of applications that are possible, as discussed earlier. C# uses the framework and therefore has no restrictions on possible applications. However, here are a few of the more common application types:

- **Desktop applications** — Applications, such as Microsoft Office, that have a familiar Windows look and feel about them. This is made simple by using the Windows Presentation Foundation (WPF) module of the .NET Framework, which is a library of *controls* (such as buttons, toolbars, menus, and so on) that you can use to build a Windows user interface (UI).
- **Windows Store applications** — Windows 8 introduced a new type of application, known as a Windows Store application. This type of application is designed primarily for touch devices, and it is usually run full-screen, with a minimum of clutter, and an emphasis on simplicity. You can create these applications in several ways, including using WPF.
- **Cloud/Web applications** — The .NET Framework includes a powerful system named ASP .NET, for generating web content dynamically, enabling personalization, security, and much more. Additionally, these applications can be hosted and accessed in the Cloud, for example on the Microsoft Azure platform.
- **Web APIs** — An ideal framework for building RESTful HTTP services that support a broad variety of clients, including mobile devices and browsers.
- **WCF services** — A way to create versatile distributed applications. Using WCF you can exchange virtually any data over local networks or the Internet, using the same simple syntax regardless of the language used to create a service or the system on which it resides.

Any of these types might also require some form of database access, which can be achieved using the ADO.NET (Active Data Objects .NET) section of the .NET Framework, through the ADO.NET Entity Framework, or through the LINQ (Language Integrated Query) capabilities of C#. Many other resources can be drawn on, such as tools for creating networking components, outputting graphics, performing complex mathematical tasks, and so on.

## C# in this Book

The first part of this book deals with the syntax and usage of the C# language without too much emphasis on the .NET Framework. This is necessary because you can't use the .NET Framework at all without a firm grounding in C# programming. You'll start off even simpler, in fact, and leave the more involved topic of OOP until you've covered the basics. These are taught from first principles, assuming no programming knowledge at all.

After that, you'll be ready to move on to developing more complex (but more useful) applications. Part II tackles cloud based web application programming, and Part III examines data access (for ORM database concepts, filesystem, and XML data) and LINQ. Part IV of this book looks at desktop and Windows Store application programming.

## VISUAL STUDIO 2015

In this book, you use the Visual Studio 2015 development tool for all of your C# programming, from simple command-line applications to more complex project types. A development tool, or integrated development environment (IDE), such as Visual Studio isn't essential for developing C# applications, but it makes things much easier. You can (if you want to) manipulate C# source code files in a basic text editor, such as the ubiquitous Notepad application, and compile code into assemblies using the command-line compiler that is part of the .NET Framework. However, why do this when you have the power of an IDE to help you?

## Visual Studio Express 2015 Products

In addition to Visual Studio 2015, Microsoft also supplies several simpler development tools known as Visual Studio Express or Community 2015 Products. These are freely available at <https://www.visualstudio.com/en-us/downloads/download-visual-studio-vs>.

The various express products enable you to create almost any C# application you might need. They function as slimmed-down versions of Visual Studio and retain the same look and feel. While they offer many of the same features as Visual Studio, some notable feature are absent, although not so many that they would prevent you from using these tools to work through the chapters of this book.

**NOTE** This book was written using the Enterprise version of Visual Studio 2015 because the Express products were not available. At the time of writing, there is an Express product scheduled for release called Visual Studio Express 2015 for Windows Desktop that should be sufficient for following along with the first part of this book. The remainder of the book may also allow you to use Visual Studio Express 2015 for Windows 10 and Visual Studio Express 2015 for Web, but at the time of writing we can't say for certain whether that will hold true.

## Solutions

When you use Visual Studio to develop applications, you do so by creating *solutions*. A solution, in Visual Studio terms, is more than just an application. Solutions contain *projects*, which might be WPF projects, Cloud/Web Application projects, and so on. Because solutions can contain multiple projects, you can group together related code in one place, even if it will eventually compile to multiple assemblies in various places on your hard disk.

This is very useful because it enables you to work on shared code (which might be placed in the GAC) at the same time as applications that use this code. Debugging code is a lot easier when only one development environment is used, because you can step through instructions in multiple code modules.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
.NET Framework fundamentals	The .NET Framework is Microsoft's latest development platform, and is currently in version 4.6. It includes a common type system (CTS) and common language runtime (CLR). .NET Framework applications are written using object-oriented programming (OOP) methodology, and usually contain managed code. Memory management of managed code is handled by the .NET runtime; this includes garbage collection.
.NET Framework applications	Applications written using the .NET Framework are first compiled into CIL. When an application is executed, the CLR uses a JIT to compile this CIL into native code as required. Applications are compiled and different parts are linked together into assemblies that contain the CIL.
C# basics	C# is one of the languages included in the .NET Framework. It is an evolution of previous languages such as C++, and can be used to write any number of applications, including web and desktop applications.
Integrated Development Environments (IDEs)	You can use Visual Studio 2015 to write any type of .NET application using C#. You can also use the free, but less powerful, Express product range to create .NET applications in C#. Both of these IDEs work with solutions, which can consist of multiple projects.

# 2

## Writing a C# Program

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Understanding Visual Studio 2015 basics
- Writing a simple console application
- Writing a simple desktop application

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

The wrox.com code downloads for this chapter are found at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 2 download and individually named according to the names throughout the chapter.

Now that you've spent some time learning what C# is and how it fits into the .NET Framework, it's time to get your hands dirty and write some code. You use Visual Studio 2015 (VS) throughout this book, so the first thing to do is have a look at some of the basics of this development environment.

Visual Studio is an enormous and complicated product, and it can be daunting to first-time users, but using it to create basic applications can be surprisingly simple. As you start to use Visual Studio in this chapter, you will see that you don't need to know a huge amount about it to begin playing with C# code. Later in the book you'll see some of the more complicated operations that Visual Studio can perform, but for now a basic working knowledge is all that is required.

After you've looked at the IDE, you put together two simple applications. You don't need to worry too much about the code in these for now; you just want to prove that things work. By working through the application-creation procedures in these early examples, they will become second nature before too long.

You will learn how to create two basic types of applications in this chapter: a *console application* and a *desktop application*.

The first application you create is a simple console application. Console applications don't use the graphical windows environment, so you won't have to worry about buttons, menus, interaction with the mouse pointer, and so on. Instead, you run the application in a command prompt window and interact with it in a much simpler way.

The second application is a desktop application, which you create using Windows Presentation Foundation (WPF). The look and feel of a desktop application is very familiar to Windows users, and (surprisingly) the application doesn't require much more effort to create. However, the syntax of the code required is more complicated, even though in many cases you don't actually have to worry about details.

You use both types of application in Part III and Part IV of the book, with more emphasis on console applications at the beginning. The additional flexibility of desktop applications isn't necessary when you are learning the C# language, while the simplicity of console applications enables you to concentrate on learning the syntax without worrying about the look and feel of the application.

## THE VISUAL STUDIO 2015 DEVELOPMENT ENVIRONMENT

When Visual Studio is first loaded, it immediately presents you with the option to Sign in to Visual Studio using your Microsoft Account. By doing this, your Visual Studio settings are synced between devices so that you do not have to configure the IDE when using it on multiple workstations. If you do not have a Microsoft Account, follow the process for the creation of one and then use it to sign in. If you do not want to sign in, click the "Not now, maybe later" link, and continue the initial configuration of Visual Studio. At some point, it is recommended that you sign in and get a developer license.

If this is the first time you've run Visual Studio, you will be presented with a list of preferences intended for users who have experience with previous releases of this development environment. The choices you make here affect a number of things, such as the layout of windows, the way that console windows run, and so on. Therefore, choose Visual C# Development Settings from the drop-down; otherwise, you might find that things don't quite work as described in this book. Note that the options available vary depending on the options you chose when installing Visual Studio, but as long as you chose to install C# this option will be available.

If this isn't the first time that you've run Visual Studio, but you chose a different option the first time, don't panic. To reset the settings to Visual C# Development settings, you simply have to import them. To do this, select Tools → Import and Export Settings, and choose the Reset All Settings option, shown in Figure 2-1.

Click Next, and indicate whether you want to save your existing settings before proceeding. If you have customized things, you might want to do this; otherwise, select No and click Next again. From the next dialog box, select Visual C#, shown in Figure 2-2. Again, the available options may vary.




FIGURE 2-1




FIGURE 2-2

Finally, click Finish, then Close to apply the settings.

The Visual Studio environment layout is completely customizable, but the default is fine here. With C# Developer Settings selected, it is arranged as shown in Figure 2-3.




FIGURE 2-3

The main window, which contains a helpful Start Page by default when Visual Studio is started, is where all your code is displayed. This window can contain many documents, each indicated by a tab, so you can easily switch between several files by clicking their filenames. It also has other functions: It can display GUIs that you are designing for your projects, plain-text files, HTML, and various tools that are built into Visual Studio. You will come across all of these in the course of this book.

Above the main window are toolbars and the Visual Studio menu. Several different toolbars can be placed here, with functionality ranging from saving and loading files to building and running projects to debugging controls. Again, you are introduced to these as you need to use them.

Here are brief descriptions of each of the main features that you will use the most:

- The Toolbox window pops up when you click its tab. It provides access to, among other things, the user interface building blocks for desktop applications. Another tab, Server Explorer, can also appear here (selectable via the View ➔ Server Explorer menu option) and includes various additional capabilities, such as Azure subscription details, providing access to data sources, server settings, services, and more.
- The Solution Explorer window displays information about the currently loaded *solution*. A solution, as you learned in the previous chapter, is Visual Studio terminology for one or more

projects along with their configurations. The Solution Explorer window displays various views of the projects in a solution, such as what files they contain and what is contained in those files.

- The Team Explorer window displays information about the current Team Foundation Server or Team Foundation Service connection. This allows you access to source control, bug tracking, build automation, and other functionality. However, this is an advanced subject and is not covered in this book.
- Just below the Solution Explorer window you can display a Properties window, not shown in Figure 2-3 because it appears only when you are working on a project (you can also toggle its display using View  $\Rightarrow$  Properties Window). This window provides a more detailed view of the project's contents, enabling you to perform additional configuration of individual elements. For example, you can use this window to change the appearance of a button in a desktop application.
- Also not shown in the screenshot is another extremely important window: the Error List window, which you can display using View  $\Rightarrow$  Error List. It shows errors, warnings, and other project-related information. The window updates continuously, although some information appears only when a project is compiled.

This might seem like a lot to take in, but it doesn't take long to get comfortable. You start by building the first of your example projects, which involves many of the Visual Studio elements just described.

**NOTE** Visual Studio is capable of displaying many other windows, both informational and functional. Many of these can share screen space with the windows mentioned here, and you can switch between them using tabs, dock them elsewhere, or even detach them and place them on other displays if you have multiple monitors. Several of these windows are used later in the book, and you'll probably discover more yourself when you explore the Visual Studio environment in more detail.

## CONSOLE APPLICATIONS

You use console applications regularly in this book, particularly at the beginning, so the following Try It Out provides a step-by-step guide to creating a simple one.

### TRY IT OUT Creating a Simple Console Application: ConsoleApplication1\Program.cs

1. Create a new console application project by selecting File  $\Rightarrow$  New  $\Rightarrow$  Project, as shown in Figure 2-4.




FIGURE 2-4




FIGURE 2-5

2. Ensure that the Visual C# node is selected in the left pane of the window that appears, and choose the Console Application project type in the middle pane (see Figure 2-5). Change the Location text box to C:\BegVCSharp\Chapter02 (this directory is created automatically if it doesn't already exist). Leave the default text in the Name text box (ConsoleApplication1) and the other settings as they are (refer to Figure 2-5).
3. Click the OK button.

4. Once the project is initialized, add the following lines of code to the file displayed in the main window:

```
namespace ConsoleApplication1
{
    class Program
    {
        static void Main(string[] args)
        {
            // Output text to the screen.
            Console.WriteLine("The first app in Beginning Visual C# 2015!");
            Console.ReadKey();
        }
    }
}
```

5. Select the Debug  $\Rightarrow$  Start Debugging menu item. After a few moments you should see the window shown in Figure 2-6.




FIGURE 2-6

6. Press any key to exit the application (you might need to click on the console window to focus on it first). The display in Figure 2-6 appears only if the Visual C# Developer Settings are applied, as described earlier in this chapter. For example, with Visual Basic Developer Settings applied, an empty console window is displayed, and the application output appears in a window labeled Immediate. In this case, the `Console.ReadKey()` code also fails, and you see an error. If you experience this problem, the best solution for working through the examples in this book is to apply the Visual C# Developer Settings — that way, the results you see match the results shown here.

### *How It Works*

For now, I won't dissect the code used thus far because the focus here is on how to use the development tools to get code up and running. Clearly, Visual Studio does a lot of the work for you and makes the process of compiling and executing code simple. In fact, there are multiple ways to perform even these basic steps — for instance, you can create a new project by using the menu item mentioned earlier, by pressing **Ctrl+Shift+N**, or by clicking the corresponding icon in the toolbar.

Similarly, your code can be compiled and executed in several ways. The process you used in the example — selecting **Debug  $\Rightarrow$  Start Debugging** — also has a keyboard shortcut (**F5**) and a toolbar icon. You can also run code without being in debugging mode using the **Debug  $\Rightarrow$  Start Without Debugging** menu item (or by pressing **Ctrl+F5**), or compile your project without running it (with debugging on or off) using **Build  $\Rightarrow$  Build Solution** or pressing **F6**. Note that you can execute a project without debugging or build a project using toolbar icons, although these icons don't appear on the toolbar by default. After you have compiled your code, you can also execute it simply by running the **.exe** file produced

in Windows Explorer, or from the command prompt. To do this, open a command prompt window, change the directory to C:\BegVCSharp\Chapter02\ConsoleApplication1\ConsoleApplication1\bin\Debug\, type ConsoleApplication1, and press Enter.

**NOTE** In future examples, when you see the instructions “create a new console project” or “execute the code,” you can choose whichever method you want to perform these steps. Unless otherwise stated, all code should be run with debugging enabled. In addition, the terms “start,” “execute,” and “run” are used interchangeably in this book, and discussions following examples always assume that you have exited the application in the example.

Console applications terminate as soon as they finish execution, which can mean that you don’t get a chance to see the results if you run them directly through the IDE. To get around this in the preceding example, the code is told to wait for a key press before terminating, using the following line:

```
Console.ReadKey();
```

You will see this technique used many times in later examples. Now that you’ve created a project, you can take a more detailed look at some of the regions of the development environment.

## The Solution Explorer

By default, the Solution Explorer window is docked in the top-right corner of the screen. As with other windows, you can move it wherever you like, or you can set it to auto-hide by clicking the pin icon. The Solution Explorer window shares space with another useful window called Class View, which you can display using View  $\Rightarrow$  Class View. Figure 2-7 shows both of these windows with all nodes expanded (you can toggle between them by clicking on the tabs at the bottom of the window when the window is docked).




FIGURE 2-7

This Solution Explorer view shows the files that make up the `ConsoleApplication1` project. The file to which you added code, `Program.cs`, is shown along with another code file, `AssemblyInfo.cs`, and several references.

**NOTE** All C# code files have a `.cs` file extension.

You don't have to worry about the `AssemblyInfo.cs` file for the moment. It contains extra information about your project that doesn't concern you yet.

You can use this window to change what code is displayed in the main window by double-clicking `.cs` files; right-clicking them and selecting View Code; or by selecting them and clicking the toolbar button that appears at the top of the window. You can also perform other operations on files here, such as renaming them or deleting them from your project. Other file types can also appear here, such as project resources (resources are files used by the project that might not be C# files, such as bitmap images and sound files). Again, you can manipulate them through the same interface.

You can also expand code items such as `Program.cs` to see what is contained. This overview of your code structure can be a very useful tool; it also enables you to navigate directly to specific parts of your code file, instead of opening the code file and scrolling to the part you want.

The References entry contains a list of the .NET libraries you are using in your project. You'll look at this later; the standard references are fine for now. Class View presents an alternative view of your project by showing the structure of the code you created. You'll come back to this later in the book; for now the Solution Explorer display is appropriate. As you click on files or other icons in these windows, notice that the contents of the Properties window (shown in Figure 2-8) changes.




FIGURE 2-8

## The Properties Window

The Properties window (select `View`  $\leftrightarrow$  `Properties Window` if it isn't already displayed) shows additional information about whatever you select in the window above it. For example, the view shown in Figure 2-8 is displayed when the `Program.cs` file from the project is selected. This window also displays information about other selected items, such as user interface components (as shown in the “Desktop Applications” section of this chapter).

Often, changes you make to entries in the Properties window affect your code directly, adding lines of code or changing what you have in your files. With some projects, you spend as much time manipulating things through this window as making manual code changes.

## The Error List Window

Currently, the Error List window (View  $\leftrightarrow$  Error List) isn't showing anything interesting because there is nothing wrong with the application. However, this is a very useful window indeed. As a test, remove the semicolon from one of the lines of code you added in the previous section. After a moment, you should see a display like the one shown in Figure 2-9.




FIGURE 2-9

In addition, the project will no longer compile.

**NOTE** In Chapter 3, when you start looking at C# syntax, you will learn that semicolons are expected throughout your code — at the end of most lines, in fact.

This window helps you eradicate bugs in your code because it keeps track of what you have to do to compile projects. If you double-click the error shown here, the cursor jumps to the position of the error in your source code (the source file containing the error will be opened if it isn't already open), so you can fix it quickly. Red wavy lines appear at the positions of errors in the code, so you can quickly scan the source code to see where problems lie.

The error location is specified as a line number. By default, line numbers aren't displayed in the Visual Studio text editor, but that is something well worth turning on. To do so, tick the Line numbers check box in the Options dialog box (selected via the Tools  $\leftrightarrow$  Options menu item). It appears in the Text Editor  $\leftrightarrow$  All Languages  $\leftrightarrow$  General category.

You can also change this setting on a per-language basis through the language-specific settings pages in the dialog box. Many other useful options can be found through this dialog box, and you will use several of them later in this book.

## DESKTOP APPLICATIONS

It is often easier to demonstrate code by running it as part of a desktop application than through a console window or via a command prompt. You can do this using user interface building blocks to piece together a user interface.

The following Try It Out shows just the basics of doing this, and you'll see how to get a desktop application up and running without a lot of details about what the application is actually doing. You'll use WPF here, which is Microsoft's recommended technology for creating desktop applications. Later, you take a detailed look at desktop applications and learn much more about what WPF is and what it's capable of.

## TRY IT OUT Creating a Simple Windows Application: WpfApplication1\MainWindow.xaml and WpfApplication1\MainWindow.xaml.cs

1. Create a new project of type WPF Application in the same location as before (`C:\BegVCSharp\Chapter02`), with the default name `WpfApplication1`. If the first project is still open, make sure the Create New Solution option is selected to start a new solution. These settings are shown in Figure 2-10.




FIGURE 2-10

2. Click OK to create the project. You should see a new tab that's split into two panes. The top pane shows an empty window called `MainWindow` and the bottom pane shows some text. This text is actually the code that is used to generate the window, and you'll see it change as you modify the UI.
3. Click the Toolbox tab on the top left of the screen, then double-click the Button entry in the Common WPF Controls section to add a button to the window.
4. Double-click the button that has been added to the window.

5. The C# code in `MainWindow.xaml.cs` should now be displayed. Modify it as follows (only part of the code in the file is shown here for brevity):

```
private void button_Click(object sender, RoutedEventArgs e)
{
    MessageBox.Show("The first desktop app in the book!");
}
```

6. Run the application.  
7. Click the button presented to open a message dialog box, as shown in Figure 2-11.




FIGURE 2-11

8. Click OK, and then exit the application by clicking the X in the top-right corner, as is standard for desktop applications.

### How It Works

Again, it is plain that the IDE has done a lot of work for you and made it simple to create a functional desktop application with little effort. The application you created behaves just like other windows — you can move it around, resize it, minimize it, and so on. You don't have to write the code to do that — it just works. The same is true for the button you added. Simply by double-clicking it, the IDE knew that you wanted to write code to execute when a user clicked the button in the running application. All you had to do was provide that code, getting full button-clicking functionality for free.

Of course, desktop applications aren't limited to plain windows with buttons. Look at the Toolbox window where you found the Button option and you'll see a whole host of user interface building blocks (known as *controls*), some of which might be familiar. You will use most of these at some point in the book, and you'll find that they are all easy to use and save you a lot of time and effort.


The code for your application, in `MainWindow.xaml.cs`, doesn't look much more complicated than the code in the previous section, and the same is true for the code in the other files in the Solution Explorer window. The code in `MainWindow.xaml` (the split-pane view where you added the button) also looks pretty straightforward.

This code is written in XAML, which is the language used to define user interfaces in WPF applications.


Now take a closer look at the button you added to the window. In the top pane of `MainWindow.xaml`, click once on the button to select it. When you do so, the Properties window in the bottom-right corner of the screen shows the properties of the button control (controls have properties much like the files shown in the last example). Ensure that the application isn't currently running, scroll down to the

Content property, which is currently set to `Button`, and change the value to `Click Me`, as shown in Figure 2-12.

The text written on the button in the designer should also reflect this change, as should the XAML code, as shown in Figure 2-13.



**FIGURE 2-12**



**FIGURE 2-13**

There are many properties for this button, ranging from simple formatting of the color and size to more obscure settings such as data binding, which enables you to establish links to data. As briefly mentioned in the previous example, changing properties often results in direct changes to code, and

this is no exception, as you saw with the XAML code change. However, if you switch back to the code view of `MainWindow.xaml.cs`, you won't see any changes there. This is because WPF applications are capable of keeping design aspects of your applications (such as the text on a button) separate from the functionality aspects (such as what happens when you click a button).

**NOTE** Note that it is also possible to use Windows Forms to create desktop applications. WPF is a newer technology that is intended to replace Windows Forms and provides a far more flexible and powerful way to create desktop applications, which is why this book doesn't cover Windows Forms.

---

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>Visual Studio 2015 settings</b>	This book requires the C# development settings option, which you choose when you first run Visual Studio or by resetting the settings.
<b>Console applications</b>	Console applications are simple command-line applications, used in much of this book to illustrate techniques. Create a new console application with the Console Application template that you see when you create a new project in Visual Studio. To run a project in debug mode, use the Debug ➔ Start Debugging menu item, or press F5.
<b>IDE windows</b>	The project contents are shown in the Solution Explorer window. The properties of the selected item are shown in the Properties window. Errors are shown in the Error List window.
<b>Desktop applications</b>	Desktop applications are applications that have the look and feel of standard Windows applications, including the familiar icons to maximize, minimize, and close an application. They are created with the WPF Application template in the New Project dialog box.



# 3

## Variables and Expressions

### WHAT YOU WILL LEARN IN THIS CHAPTER

---

- Understanding basic C# syntax
- Using variables
- Using expressions

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 3 download and individually named according to the names throughout the chapter.

To use C# effectively, it's important to understand what you're actually doing when you create a computer program. Perhaps the most basic description of a computer program is that it is a series of operations that manipulate data. This is true even of the most complicated examples, including vast, multi-featured Windows applications (such as the Microsoft Office Suite). Although this is often completely hidden from users of applications, it is always going on behind the scenes.

To illustrate this further, consider the display unit of your computer. What you see onscreen is often so familiar that it is difficult to imagine it as anything other than a “moving picture.” In fact, what you see is only a representation of some data, which in its raw form is merely a stream of 0s and 1s stashed away somewhere in the computer’s memory. Any onscreen action — moving a mouse pointer, clicking on an icon, typing text into a word processor — results in the shunting around of data in memory.

Of course, simpler situations show this just as well. When using a calculator application, you are supplying data as numbers and performing operations on the numbers in much the same way as you would with paper and pencil — but a lot quicker!

If computer programs are fundamentally performing operations on data, this implies that you need a way to store that data, and some methods to manipulate it. These two functions are provided by *variables* and *expressions*, respectively, and this chapter explores what that means, both in general and specific terms.

First, though, you'll take a look at the basic syntax involved in C# programming, because you need a context in which you can learn about and use variables and expressions in the C# language.

## BASIC C# SYNTAX

The look and feel of C# code is similar to that of C++ and Java. This syntax can look quite confusing at first and it's a lot less like written English than some other languages. However, as you immerse yourself in the world of C# programming, you'll find that the style used is a sensible one, and it is possible to write very readable code without much effort.

Unlike the compilers of some other languages such as Python, C# compilers ignore additional spacing in code, whether it results from spaces, carriage returns, or tab characters (collectively known as *whitespace characters*). This means you have a lot of freedom in the way that you format your code, although conforming to certain rules can help make your code easier to read.

C# code is made up of a series of *statements*, each of which is terminated with a semicolon. Because whitespace is ignored, multiple statements can appear on one line, although for readability it is usual to add carriage returns after semicolons, to avoid multiple statements on one line. It is perfectly acceptable (and quite normal), however, to use statements that span several lines of code.

C# is a *block-structured language*, meaning statements are part of a block of code. These blocks, which are delimited with curly brackets ({ and }), may contain any number of statements, or none at all. Note that the curly bracket characters do not need accompanying semicolons.

For example, a simple block of C# code could take the following form:

```
{  
    <code line 1, statement 1>;  
    <code line 2, statement 2>  
    <code line 3, statement 2>;  
}
```

Here the <code line x, statement y> sections are not actual pieces of C# code; this text is used as a placeholder where C# statements would go. In this case, the second and third lines of code are part of the same statement, because there is no semicolon after the second line. Indenting the third line of code makes it easier to recognize that it is actually a continuation of the second line.

The following simple example uses *indentation* to clarify the C# itself. This is actually standard practice, and in fact Visual Studio automatically does this for you by default. In general, each block of code has its own level of indentation, meaning how far to the right it is. Blocks of code may be

nested inside each other (that is, blocks may contain other blocks), in which case nested blocks will be indented further:

```
{
    <code line 1>;
    {
        <code line 2>;
        <code line 3>;
    }
    <code line 4>;
}
```

In addition, lines of code that are continuations of previous lines are usually indented further as well, as in the third line of code in the first code example.

**NOTE** Look in the Visual Studio Options dialog box (select Tools → Options) to see the rules that Visual Studio uses for formatting your code. There are many of these, in subcategories of the Text Editor → C# → Formatting node. Most of the settings here reflect parts of C# that haven't been covered yet, but you might want to return to these settings later if you want to tweak them to suit your personal style better. For clarity, this book shows all code snippets as they would be formatted by the default settings.

Of course, this style is by no means mandatory. If you don't use it, however, you will quickly find that things can get very confusing as you move through this book!

Comments are something else you often see in C# code. A comment is not, strictly speaking, C# code at all, but it happily cohabits with it. Comments are self-explanatory: They enable you to add descriptive text to your code — in plain English (or French, German, Mongolian, and so on) — which is ignored by the compiler. When you start dealing with lengthy code sections, it's useful to add reminders about exactly what you are doing, such as “this line of code asks the user for a number” or “this code section was written by Bob.”

C# provides two ways of doing this. You can either place markers at the beginning and end of a comment or you can use a marker that means “everything on the rest of this line is a comment.” The latter method is an exception to the rule mentioned previously about C# compilers ignoring carriage returns, but it is a special case.

To indicate comments using the first method, you use /\* characters at the start of the comment and \*/ characters at the end. These may occur on a single line, or on different lines, in which case all lines in between are part of the comment. The only thing you can't type in the body of a comment is \*/, because that is interpreted as the end marker. For example, the following are okay:

```
/* This is a comment */
/* And so. . .
   . . . is this! */
```

The following, however, causes problems:

```
/* Comments often end with "*/" characters */
```

Here, the end of the comment (the characters after "\*/") will be interpreted as C# code, and errors will occur.

The other commenting approach involves starting a comment with //. After that, you can write whatever you like — as long as you keep to one line! The following is okay:

```
// This is a different sort of comment.
```

The following fails, however, because the second line is interpreted as C# code:

```
// So is this,  
// but this bit isn't.
```

This sort of commenting is useful to document statements because both can be placed on a single line:

```
<A statement>; // Explanation of statement
```

It was stated earlier that there are two ways of commenting C# code, but there is a third type of comment in C# — although strictly speaking this is an extension of the // syntax. You can use single-line comments that start with three / symbols instead of two, like this:

```
/// A special comment
```

Under normal circumstances, they are ignored by the compiler — just like other comments — but you can configure Visual Studio to extract the text after these comments and create a specially formatted text file when a project is compiled. You can then use it to create documentation. In order for this documentation to be created, the comments must follow the rules of XML documentation as described here <https://msdn.microsoft.com/library/aa288481.aspx> — a subject not covered in this book but one that is well worth learning about if you have some spare time.

A *very* important point about C# code is that it is *case sensitive*. Unlike some other languages, you must enter code using exactly the right case, because using an uppercase letter instead of a lowercase one will prevent a project from compiling. For example, consider the following line of code, taken from Chapter 2:

```
Console.WriteLine("The first app in Beginning C# Programming!");
```

This code is understood by the C# compiler, as the case of the `Console.WriteLine()` command is correct. However, none of the following lines of code work:

```
console.WriteLine("The first app in Beginning C# Programming!");  
CONSOLE.WRITELINE("The first app in Beginning C# Programming!");  
Console.Writeline("The first app in Beginning C# Programming!");
```

Here, the case used is wrong, so the C# compiler won't know what you want. Luckily, as you will soon discover, Visual Studio is very helpful when it comes to entering code, and most of the time it knows (as much as a program can know) what you are trying to do. As you type, it suggests commands that you might like to use, and it tries to correct case problems.

## BASIC C# CONSOLE APPLICATION STRUCTURE

Here, you'll take a closer look at the console application example from Chapter 2 (`ConsoleApplication1`) and break down the structure a bit. Here's the code:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace ConsoleApplication1
{
    class Program
    {
        static void Main(string[] args)
        {
            // Output text to the screen.
            Console.WriteLine("The first app in Beginning C# Programming!");
            Console.ReadKey();
        }
    }
}
```

You can immediately see that all the syntactic elements discussed in the previous section are present here — semicolons, curly braces, and comments, along with appropriate indentation.

The most important section of code at the moment is the following:

```
static void Main(string[] args)
{
    // Output text to the screen.
    Console.WriteLine("The first app in Beginning C# Programming!");
    Console.ReadKey();
}
```

This is the code that is executed when you run your console application. Well, to be more precise, the code block enclosed in curly braces is executed. The comment line doesn't do anything, as mentioned earlier; it's just there for clarity. The other two code lines output some text to the console window and wait for a response, respectively, although the exact mechanisms of this don't need to concern you for now.

Note how to achieve the code outlining functionality shown in the previous chapter, albeit for a Windows application, since it is such a useful feature. You can do this with the `#region` and `#endregion` keywords, which define the start and end of a region of code that can be expanded and collapsed. For example, you could modify the generated code for `ConsoleApplication1` as follows:

```
#region Using directives
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
#endregion
```

This enables you to collapse this code into a single line and expand it again later should you want to look at the details. The `using` statements contained here, and the `namespace` statement just underneath, are explained at the end of this chapter.

**NOTE** Any keyword that starts with a `#` is actually a preprocessor directive and not, strictly speaking, a C# keyword. Other than the two described here, `#region` and `#endregion`, these can be quite complicated, and they have very specialized uses. This is one subject you might like to investigate yourself after you've worked through this book.

For now, don't worry about the other code in the example, because the purpose of these first few chapters is to explain basic C# syntax, so the exact method of how the application execution gets to the point where `Console.WriteLine()` is called is of no concern. Later, the significance of this additional code is made clear.

## VARIABLES

As mentioned earlier, variables are concerned with the storage of data. Essentially, you can think of variables in computer memory as boxes sitting on a shelf. You can put things in boxes and take them out again, or you can just look inside a box to see if anything is there. The same goes for variables; you place data in them and can take it out or look at it, as required.

Although all data in a computer is effectively the same thing (a series of 0s and 1s), variables come in different flavors, known as *types*. Using the box analogy again, boxes come in different shapes and sizes, so some items fit only in certain boxes. The reasoning behind this type system is that different types of data may require different methods of manipulation, and by restricting variables to individual types you can avoid mixing them up. For example, it wouldn't make much sense to treat the series of 0s and 1s that make up a digital picture as an audio file.

To use variables, you have to *declare* them. This means that you have to assign them a *name* and a *type*. After you have declared variables, you can use them as storage units for the type of data that you declared them to hold.

C# syntax for declaring variables merely specifies the type and variable name:

```
<type> <name>;
```

If you try to use a variable that hasn't been declared, your code won't compile, but in this case the compiler tells you exactly what the problem is, so this isn't really a disastrous error. Trying to use a variable without assigning it a value also causes an error, but, again, the compiler detects this.

## Simple Types

Simple types include types such as numbers and Boolean (true or false) values that make up the fundamental building blocks for your applications. Unlike complex types, simple types cannot have

children or attributes. Most of the simple types available are numeric, which at first glance seems a bit strange — surely, you only need one type to store a number?

The reason for the plethora of numeric types is because of the mechanics of storing numbers as a series of 0s and 1s in the memory of a computer. For integer values, you simply take a number of *bits* (individual digits that can be 0 or 1) and represent your number in binary format. A variable storing N bits enables you to represent any number between 0 and  $(2^N - 1)$ . Any numbers above this value are too big to fit into this variable.

For example, suppose you have a variable that can store two bits. The mapping between integers and the bits representing those integers is therefore as follows:

```
0 = 00
1 = 01
2 = 10
3 = 11
```

In order to store more numbers, you need more bits (three bits enable you to store the numbers from 0 to 7, for example).

The inevitable result of this system is that you would need an infinite number of bits to be able to store every imaginable number, which isn't going to fit in your trusty PC. Even if there were a quantity of bits you could use for every number, it surely wouldn't be efficient to use all these bits for a variable that, for example, was required to store only the numbers between 0 and 10 (because storage would be wasted). Four bits would do the job fine here, enabling you to store many more values in this range in the same space of memory.

Instead, a number of different integer types can be used to store various ranges of numbers, which take up differing amounts of memory (up to 64 bits). These types are shown in Table 3-1.

**NOTE** Each of these types uses one of the standard types defined in the .NET Framework. As discussed in Chapter 1, this use of standard types is what enables language interoperability. The names you use for these types in C# are aliases for the types defined in the framework. Table 3-1 lists the names of these types as they are referred to in the .NET Framework library.

TABLE 3-1: Integer Types

TYPE	ALIAS FOR	ALLOWED VALUES
sbyte	System.SByte	Integer between -128 and 127
byte	System.Byte	Integer between 0 and 255
short	System.Int16	Integer between -32768 and 32767
ushort	System.UInt16	Integer between 0 and 65535

*continues*

**TABLE 3-1** (continued)

TYPE	ALIAS FOR	ALLOWED VALUES
int	System.Int32	Integer between -2147483648 and 2147483647
uint	System.UInt32	Integer between 0 and 4294967295
long	System.Int64	Integer between -9223372036854775808 and 9223372036854775807
ulong	System.UInt64	Integer between 0 and 18446744073709551615

The u characters before some variable names are shorthand for *unsigned*, meaning that you can't store negative numbers in variables of those types, as shown in the Allowed Values column of the preceding table.

Of course, you also need to store *floating-point* values, those that aren't whole numbers. You can use three floating-point variable types: `float`, `double`, and `decimal`. The first two store floating points in the form  $6m \times 2^e$ , where the allowed values for m and e differ for each type. `decimal` uses the alternative form  $6m \times 10^e$ . These three types are shown in Table 3-2, along with their allowed values of m and e, and these limits in real numeric terms.

**TABLE 3-2:** Floating-point Types

TYPE	ALIAS FOR	MIN M	MAX M	MIN E	MAX E	APPROX MIN VALUE	APPROX MAX VALUE
float	System.Single	0	22 <sup>4</sup>	-149	104	$1.5 \times 10^{-45}$	$3.4 \times 10^{38}$
double	System.Double	0	25 <sup>3</sup>	-1075	970	$5.0 \times 10^{-324}$	$1.7 \times 10^{308}$
decimal	System.Decimal	0	29 <sup>6</sup>	-28	0	$1.0 \times 10^{-28}$	$7.9 \times 10^{28}$

In addition to numeric types, three other simple types are available (see Table 3-3).

**TABLE 3-3:** Text and Boolean Types

TYPE	ALIAS FOR	ALLOWED VALUES
char	System.Char	Single Unicode character, stored as an integer between 0 and 65535
bool	System.Boolean	Boolean value, true or false
string	System.String	A sequence of characters

Note that there is no upper limit on the amount of characters making up a `string`, because it can use varying amounts of memory.

The Boolean type `bool` is one of the most commonly used variable types in C#, and indeed similar types are equally prolific in code in other languages. Having a variable that can be either `true` or `false` has important ramifications when it comes to the flow of logic in an application. As a simple example, consider how many questions can be answered with `true` or `false` (or yes and no). Performing comparisons between variable values or validating input are just two of the programmatic uses of Boolean variables that you will examine very soon.

Now that you've seen these types, consider a short example that declares and uses them. In the following Try It Out you use some simple code that declares two variables, assigns them values, and then outputs these values.

### TRY IT OUT Using Simple Type Variables: Ch03Ex01\Program.cs

1. Create a new console application called Ch03Ex01 and save it in the directory `C:\BegVCSharp\Chapter03`.
2. Add the following code to `Program.cs`:

```
static void Main(string[] args)
{
    int myInteger;
    string myString;
    myInteger = 17;
    myString = "\"myInteger\" is";
    Console.WriteLine($"{myString} {myInteger}");
    Console.ReadKey();
}
```

3. Execute the code. The result is shown in Figure 3-1.




FIGURE 3-1

### How It Works

The added code performs three tasks:

- It declares two variables.
- It assigns values to those two variables.
- It outputs the values of the two variables to the console.

Variable declaration occurs in the following code:

```
int myInteger;
string myString;
```

The first line declares a variable of type `int` with a name of `myInteger`, and the second line declares a variable of type `string` called `myString`.

**NOTE** Variable naming is restricted; you can't use just any sequence of characters. You learn about this in the section titled "Variable Naming."

The next two lines of code assign values:

```
myInteger = 17;
myString = "\"myInteger\" is";
```

Here, you assign two fixed values (known as *literal* values in code) to your variables using the `= assignment operator` (the "Expressions" section of this chapter has more details about operators). You assign the integer value `17` to `myInteger`, and you assigned the following string (including the quotes) to `myString`:

```
"myInteger" is
```

When you assign string literal values in this way, double quotation marks are required to enclose the string. Therefore, certain characters might cause problems if they are included in the string itself, such as the double quotation characters, and you must escape some characters by substituting a sequence of other characters (an *escape sequence*) that represents the character(s) you want to use. In this example, you use the sequence `\"` to escape a double quotation mark:

```
myString = "\"myInteger\" is";
```

If you didn't use these escape sequences and tried coding this as follows, you would get a compiler error:

```
myString = ""myInteger" is";
```

Note that assigning string literals is another situation in which you must be careful with line breaks — the C# compiler rejects string literals that span more than one line. If you want to add a line break, then use the escape sequence for a newline character in your string, which is `\n`. For example, consider the following assignment:

```
myString = "This string has a\nline break.";
```

This string would be displayed on two lines in the console view as follows:

```
This string has a
line break.
```

All escape sequences consist of the backslash symbol followed by one of a small set of characters (you'll see the full set later). Because this symbol is used for this purpose, there is also an escape sequence for the backslash symbol itself, which is simply two consecutive backslashes (`\\"\\`).

Getting back to the code, there is one more new line to look at:

```
Console.WriteLine($"{myString} {myInteger}");
```

This is a new feature in C# 6 called *String Interpolation* and looks similar to the simple method of writing text to the console that you saw in the first example, but now you are specifying your variables. It's too soon to dive into the details of this line of code, but suffice it to say that it is the technique you will be using in the first part of this book to output text to the console window.

This method of outputting text to the console is what you use to display output from your code in the examples that follow. Finally, the code includes the line shown in the earlier example for waiting for user input before terminating:

```
Console.ReadKey();
```

Again, the code isn't dissected now, but you will see it frequently in later examples. For now, understand that it pauses code execution until you press a key.

## Variable Naming

As mentioned in the previous section, you can't just choose any sequence of characters as a variable name. This isn't as worrying as it might sound, however, because you're still left with a very flexible naming system.

The basic variable naming rules are as follows:

- The first character of a variable name must be either a letter, an underscore character(\_), or the *at* symbol (@).
- Subsequent characters may be letters, underscore characters, or numbers.

There are also certain keywords that have a specialized meaning to the C# compiler, such as the *using* and *namespace* keywords shown earlier. If you use one of these by mistake, the compiler complains, however, so don't worry about it.

For example, the following variable names are fine:

```
myBigVar  
VAR1  
_test
```

These are not, however:

```
99BottlesOfBeer  
namespace  
It's-All-Over
```

## Literal Values

The previous Try It Out showed two examples of literal values: an integer (17) and a string ("\"myInteger\" is"). The other variable types also have associated literal values, as shown in Table 3-4. Many of these involve *suffixes*, whereby you add a sequence of characters to the end of the literal value to specify the type desired. Some literals have multiple types, determined at compile time by the compiler based on their context (also shown in Table 3-4).

TABLE 3-4: Literal Values

Type(s)	Category	Suffix	Example/Allowed Values
bool	Boolean	None	True or false
int, uint, long, ulong	Integer	None	100
uint, ulong	Integer	u or U	100U
long, ulong	Integer	l or L	100L
ulong	Integer	ul, uL, U1, UL, lu, lU, Lu, or LU	100UL
float	Real	f or F	1.5F
double	Real	None, d, or D	1.5
decimal	Real	m or M	1.5M
char	Character	None	'a', or escape sequence
string	String	None	"a. . .a", may include escape sequences

## String Literals

Earlier in the chapter, you saw a few of the escape sequences you can use in `string` literals. Table 3-5 lists these for reference purposes.

TABLE 3-5: Escape Sequences for String Literals

Escape Sequence	Character Produced	Unicode Value of Character
\'	Single quotation mark	0x0027
\"	Double quotation mark	0x0022
\\\	Backslash	0x005C
\0	Null	0x0000
\a	Alert (causes a beep)	0x0007
\b	Backspace	0x0008
\f	Form feed	0x000C
\n	New line	0x000A

ESCAPE SEQUENCE	CHARACTER PRODUCED	UNICODE VALUE OF CHARACTER
\r	Carriage return	0x000D
\t	Horizontal tab	0x0009
\v	Vertical tab	0x000B

The Unicode Value of Character column of the preceding table shows the hexadecimal values of the characters as they are found in the Unicode character set. As well as the preceding, you can specify any Unicode character using a Unicode escape sequence. These consist of the standard \ character followed by a u and a four-digit hexadecimal value (for example, the four digits after the x in Table 3-5).

This means that the following strings are equivalent:

```
"Benjamin\'s string."
"Benjamin\u0027s string."
```

Obviously, you have more versatility using Unicode escape sequences.

You can also specify strings *verbatim*. This means that all characters contained between two double quotation marks are included in the string, including end-of-line characters and characters that would otherwise need escaping. The only exception to this is the escape sequence for the double quotation mark character, which must be specified to avoid ending the string. To do this, place the @ character before the string:

```
@"Verbatim string literal."
```

This string could just as easily be specified in the normal way, but the following requires the @ character:

```
@"A short list:
item 1
item 2"
```

Verbatim strings are particularly useful in filenames, as these use plenty of backslash characters. Using normal strings, you'd have to use double backslashes all the way along the string:

```
"C:\\\\Temp\\\\MyDir\\\\MyFile.doc"
```

With verbatim string literals you can make this more readable. The following verbatim string is equivalent to the preceding one:

```
@"C:\\Temp\\MyDir\\MyFile.doc"
```

**NOTE** As shown later in the book, strings are reference types. This is in contrast to the other types you've seen in this chapter, which are value types. One consequence of this is that strings can also be assigned the value null, which means that the string variable doesn't reference a string (or anything else, for that matter).

## EXPRESSIONS

C# contains a number of *operators* for this purpose. By combining operators with variables and literal values (together referred to as *operands* when used with operators), you can create *expressions*, which are the basic building blocks of computation.

The operators available range from the simple to the highly complex, some of which you might never encounter outside of mathematical applications. The simple ones include all the basic mathematical operations, such as the + operator to add two operands; the complex ones include manipulations of variable content via the binary representation of this content. There are also logical operators specifically for dealing with Boolean values, and assignment operators such as =.

This chapter focuses on the mathematical and assignment operators, leaving the logical ones for the next chapter, where you examine Boolean logic in the context of controlling program flow.

Operators can be roughly classified into three categories:

- **Unary** — Act on single operands
- **Binary** — Act on two operands
- **Ternary** — Act on three operands

Most operators fall into the binary category, with a few unary ones, and a single ternary one called the *conditional operator* (the conditional operator is a logical one and is discussed in Chapter 4, “Flow Control”). Let’s start by looking at the mathematical operators, which span both the unary and binary categories.

## Mathematical Operators

There are five simple mathematical operators, two of which (+ and -) have both binary and unary forms. Table 3-6 lists each of these operators, along with a short example of its use and the result when it’s used with simple numeric types (integer and floating point).

**TABLE 3-6:** Simple Mathematical Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
+	Binary	<code>var1 = var2 + var3;</code>	<code>var1</code> is assigned the value that is the sum of <code>var2</code> and <code>var3</code> .
-	Binary	<code>var1 = var2 - var3;</code>	<code>var1</code> is assigned the value that is the value of <code>var3</code> subtracted from the value of <code>var2</code> .
*	Binary	<code>var1 = var2 * var3;</code>	<code>var1</code> is assigned the value that is the product of <code>var2</code> and <code>var3</code> .

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
/	Binary	<code>var1 = var2 / var3;</code>	<code>var1</code> is assigned the value that is the result of dividing <code>var2</code> by <code>var3</code> .
%	Binary	<code>var1 = var2 % var3;</code>	<code>var1</code> is assigned the value that is the remainder when <code>var2</code> is divided by <code>var3</code> .
+	Unary	<code>var1 = +var2;</code>	<code>var1</code> is assigned the value of <code>var2</code> .
-	Unary	<code>var1 = -var2;</code>	<code>var1</code> is assigned the value of <code>var2</code> multiplied by <code>-1</code> .

**NOTE** The `+` (unary) operator is slightly odd, as it has no effect on the result. It doesn't force values to be positive, as you might assume — if `var2` is `-1`, then `+var2` is also `-1`. However, it is a universally recognized operator, and as such is included. The most useful fact about this operator is shown later in this book when you look at operator overloading.

The examples use simple numeric types because the result can be unclear when using the other simple types. What would you expect if you added two Boolean values, for example? In this case, nothing, because the compiler complains if you try to use `+` (or any of the other mathematical operators) with `bool` variables. Adding `char` variables is also slightly confusing. Remember that `char` variables are actually stored as numbers, so adding two `char` variables also results in a number (of type `int`, to be precise). This is an example of *implicit conversion*, which you'll learn a lot more about shortly (along with explicit conversion), because it also applies to cases where `var1`, `var2`, and `var3` are of mixed types.

The binary `+` operator *does* make sense when used with string type variables. In this case, the table entry should read as shown in Table 3-7.

TABLE 3-7 The String Concatenation Operator

Operator	Category	Example Expression	Result
+	Binary	<code>var1 = var2 + var3;</code>	<code>var1</code> is assigned the value that is the concatenation of the two strings stored in <code>var2</code> and <code>var3</code> .

None of the other mathematical operators, however, work with strings.

The other two operators you should look at here are the increment and decrement operators, both of which are unary operators that can be used in two ways: either immediately before or immediately after the operand. The results obtained in simple expressions are shown in Table 3-8.

TABLE 3-8: Increment and Decrement Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
<code>++</code>	Unary	<code>var1 = ++var2;</code>	<code>var1</code> is assigned the value of <code>var2 + 1</code> . <code>var2</code> is incremented by 1.
<code>--</code>	Unary	<code>var1 = --var2;</code>	<code>var1</code> is assigned the value of <code>var2 - 1</code> . <code>var2</code> is decremented by 1.
<code>++</code>	Unary	<code>var1 = var2++;</code>	<code>var1</code> is assigned the value of <code>var2</code> . <code>var2</code> is incremented by 1.
<code>--</code>	Unary	<code>var1 = var2--;</code>	<code>var1</code> is assigned the value of <code>var2</code> . <code>var2</code> is decremented by 1.

These operators always result in a change to the value stored in their operand:

- `++` always results in its operand being incremented by one.
- `--` always results in its operand being decremented by one.

The differences between the results stored in `var1` are a consequence of the fact that the placement of the operator determines when it takes effect. Placing one of these operators before its operand means that the operand is affected before any other computation takes place. Placing it after the operand means that the operand is affected after all other computation of the expression is completed.

This merits another example! Consider this code:

```
int var1, var2 = 5, var3 = 6;
var1 = var2++ * --var3;
```

What value will be assigned to `var1`? Before the expression is evaluated, the `--` operator preceding `var3` takes effect, changing its value from 6 to 5. You can ignore the `++` operator that follows `var2`, as it won't take effect until after the calculation is completed, so `var1` will be the product of 5 and 5, or 25.

These simple unary operators come in very handy in a surprising number of situations. They are really just shorthand for expressions such as this:

```
var1 = var1 + 1;
```

This sort of expression has many uses, particularly where *looping* is concerned, as shown in the next chapter. The following Try It Out provides an example demonstrating how to use the mathematical operators, and it introduces a couple of other useful concepts as well. The code prompts you to type in a string and two numbers and then demonstrates the results of performing some calculations.

**TRY IT OUT****Manipulating Variables with Mathematical Operators: Ch03Ex02\Program.cs**

1. Create a new console application called Ch03Ex02 and save it to the directory C:\BegVCSharp\Chapter03.
2. Add the following code to Program.cs:

```
static void Main(string[] args)
{
    double firstNumber, secondNumber;
    string userName;
    Console.WriteLine("Enter your name:");
    userName = Console.ReadLine();
    Console.WriteLine($"Welcome {userName}!");
    Console.WriteLine("Now give me a number:");
    firstNumber = Convert.ToDouble(Console.ReadLine());
    Console.WriteLine("Now give me another number:");
    secondNumber = Convert.ToDouble(Console.ReadLine());
    Console.WriteLine($"The sum of {firstNumber} and {secondNumber} is " +
        $"{firstNumber + secondNumber}.");
    Console.WriteLine($"The result of subtracting {secondNumber} from " +
        $"{firstNumber} is {firstNumber - secondNumber}.");
    Console.WriteLine($"The product of {firstNumber} and {secondNumber} " +
        $"is {firstNumber * secondNumber}.");
    Console.WriteLine($"The result of dividing {firstNumber} by " +
        $"{secondNumber} is {firstNumber / secondNumber}.");
    Console.WriteLine($"The remainder after dividing {firstNumber} by " +
        $"{secondNumber} is {firstNumber % secondNumber}.");
    Console.ReadKey();
}
```

3. Execute the code. The display shown in Figure 3-2 appears.




FIGURE 3-2

4. Enter your name and press Enter. Figure 3-3 shows the display.




FIGURE 3-3

5. Enter a number, press Enter, enter another number, and then press Enter again. Figure 3-4 shows an example result.




FIGURE 3-4

## How It Works

As well as demonstrating the mathematical operators, this code introduces two important concepts that you will often come across:

- User input
- Type conversion

User input uses a syntax similar to the `Console.WriteLine()` command you've already seen — you use `Console.ReadLine()`. This command prompts the user for input, which is stored in a `string` variable:

```
string userName;
Console.WriteLine("Enter your name:");
userName = Console.ReadLine();
Console.WriteLine($"Welcome {userName}!");
```

This code writes the contents of the assigned variable, `userName`, straight to the screen.

You also read in two numbers in this example. This is slightly more involved, because the `Console.ReadLine()` command generates a string, but you want a number. This introduces the topic of *type conversion*, which is covered in more detail in Chapter 5, “More about Variables,” but let’s have a look at the code used in this example.

First, you declare the variables in which you want to store the number input:

```
double firstNumber, secondNumber;
```

Next, you supply a prompt and use the command `Convert.ToDouble()` on a string obtained by `Console.ReadLine()` to convert the string into a double type. You assign this number to the `firstNumber` variable you have declared:

```
Console.WriteLine("Now give me a number:");
firstNumber = Convert.ToDouble(Console.ReadLine());
```

This syntax is remarkably simple, and many other conversions can be performed in a similar way.

The remainder of the code obtains a second number in the same way:

```
Console.WriteLine("Now give me another number:");
secondNumber = Convert.ToDouble(Console.ReadLine());
```

Next, you output the results of adding, subtracting, multiplying, and dividing the two numbers, in addition to displaying the remainder after division, using the remainder (%) operator:

```
Console.WriteLine($"The sum of {firstNumber} and {secondNumber} is " +
    $"{firstNumber + secondNumber}.");
Console.WriteLine($"The result of subtracting {secondNumber} from " +
    $"{firstNumber} is {firstNumber - secondNumber}.");
Console.WriteLine($"The product of {firstNumber} and {secondNumber} " +
    $"is {firstNumber * secondNumber}.");
Console.WriteLine($"The result of dividing {firstNumber} by " +
    $"{secondNumber} is {firstNumber / secondNumber}.");
Console.WriteLine($"The remainder after dividing {firstNumber} by " +
    $"{secondNumber} is {firstNumber % secondNumber}.");
```

Note that you are supplying the expressions, `firstNumber + secondNumber` and so on, as a parameter to the `Console.WriteLine()` statement, without using an intermediate variable:

```
Console.WriteLine($"The sum of {firstNumber} and {secondNumber} is " +
    $"{firstNumber + secondNumber}.");
```

This kind of syntax can make your code very readable, and reduce the number of lines of code you need to write.

## Assignment Operators

So far, you've been using the simple = assignment operator, and it may come as a surprise that any other assignment operators exist at all. There are more, however, and they're quite useful! All of the assignment operators other than = work in a similar way. Like =, they all result in a value being assigned to the variable on their left side based on the operands and operators on their right side.

Table 3-9 describes the operators.

**TABLE 3-9:** Assignment Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
=	Binary	<code>var1 = var2;</code>	<code>var1</code> is assigned the value of <code>var2</code> .
+=	Binary	<code>var1 += var2;</code>	<code>var1</code> is assigned the value that is the sum of <code>var1</code> and <code>var2</code> .
-=	Binary	<code>var1 -= var2;</code>	<code>var1</code> is assigned the value that is the value of <code>var2</code> subtracted from the value of <code>var1</code> .
*=	Binary	<code>var1 *= var2;</code>	<code>var1</code> is assigned the value that is the product of <code>var1</code> and <code>var2</code> .

*continues*

**TABLE 3-9** (continued)

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
/=	Binary	var1 /= var2;	var1 is assigned the value that is the result of dividing var1 by var2.
%=	Binary	var1 %= var2;	var1 is assigned the value that is the remainder when var1 is divided by var2.

As you can see, the additional operators result in var1 being included in the calculation, so code like

```
var1 += var2;
```

has exactly the same result as

```
var1 = var1 + var2;
```

**NOTE** The += operator can also be used with strings, just like +.

Using these operators, especially when employing long variable names, can make code much easier to read.

## Operator Precedence

When an expression is evaluated, each operator is processed in sequence, but this doesn't necessarily mean evaluating these operators from left to right. As a trivial example, consider the following:

```
var1 = var2 + var3;
```

Here, the + operator acts before the = operator. There are other situations where operator precedence isn't so obvious, as shown here:

```
var1 = var2 + var3 * var4;
```

In the preceding example, the \* operator acts first, followed by the + operator, and finally the = operator. This is standard mathematical order, and it provides the same result as you would expect from working out the equivalent algebraic calculation on paper.

Similarly, you can gain control over operator precedence by using parentheses, as shown in this example:

```
var1 = (var2 + var3) * var4;
```

Here, the content of the parentheses is evaluated first, meaning that the + operator acts before the \* operator.

Table 3-10 shows the order of precedence for the operators you've encountered so far. Operators of equal precedence (such as \* and /) are evaluated from left to right.

TABLE 3-10: Operator Precedence

PRECEDENCE	OPERATORS
Highest	<code>++, --</code> (used as prefixes); <code>+, -</code> (unary)
	<code>*, /, %</code>
	<code>+, -</code>
	<code>=, *=, /=, %=, +=, -=</code>
Lowest	<code>++, --</code> (used as postfixes)

**NOTE** You can use parentheses to override this precedence order, as described previously. In addition, note that `++` and `--`, when used as postfixes, only have lowest priority in conceptual terms, as described in Table 3-10. They don't operate on the result of, say, an assignment expression, so you can consider them to have a higher priority than all other operators. However, because they change the value of their operand after expression evaluation, it's easier to think of their precedence as shown in Table 3-10.

## Namespaces

Before moving on, it's worthwhile to consider one more important subject — *namespaces*. These are the .NET way of providing containers for application code, such that code and its contents may be uniquely identified. Namespaces are also used as a means of categorizing items in the .NET Framework. Most of these items are type definitions, such as the simple types in this chapter (`System.Int32` and so on).

C# code, by default, is contained in the *global namespace*. This means that items contained in this code are accessible from other code in the global namespace simply by referring to them by name. You can use the `namespace` keyword, however, to explicitly define the namespace for a block of code enclosed in curly brackets. Names in such a namespace must be *qualified* if they are used from code outside of this namespace.

A qualified name is one that contains all of its hierarchical information, which basically means that if you have code in one namespace that needs to use a name defined in a different namespace, you must include a reference to this namespace. Qualified names use period characters (.) between namespace levels, as shown here:

```
namespace LevelOne
{
    // code in LevelOne namespace
    // name "NameOne" defined
}
// code in global namespace
```

This code defines one namespace, `LevelOne`, and a name in this namespace, `NameOne` (no actual code is shown here to keep the discussion general; instead, a comment appears where the definition would go). Code written inside the `LevelOne` namespace can simply refer to this name using `NameOne` — no classification is necessary. Code in the global namespace, however, must refer to this name using the classified name `LevelOne.NameOne`.

Note one more important point here: The `using` statement doesn't in itself give you access to names in another namespace. Unless the code in a namespace is in some way linked to your project, by being defined in a source file in the project or being defined in some other code linked to the project, you won't have access to the names contained. In addition, if code containing a namespace is linked to your project, then you have access to the names contained in that code, regardless of whether you use `using`. `using` simply makes it easier for you to access these names, and it can shorten otherwise lengthy code to make it more readable.

Going back to the code in `ConsoleApplication1` shown at the beginning of this chapter, the following lines that apply to namespaces appear:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace ConsoleApplication1
{
    ...
}
```

The five lines that start with the `using` keyword are used to declare that the `System`, `System.Collections.Generic`, `System.Linq`, `System.Text`, and `System.Threading.Tasks` namespaces will be used in this C# code and should be accessible from all namespaces in this file without classification. The `System` namespace is the root namespace for .NET Framework applications and contains all the basic functionality you need for console applications. The other four namespaces are very often used in console applications, so they are there just in case. Additionally, notice that a namespace is declared for the application code itself, `ConsoleApplication1` itself.

New to C# 6 is the `using static` keyword. This keyword allows the inclusion of static members directly into the scope of a C# program. For example, both Try It Out code walkthroughs in this chapter have used the `System.Console.WriteLine()` method, which is part of the `System.Console` static class. Notice that in these examples it is required to include the `Console` class combined with the `WriteLine()` method. When the `using static System.Console` namespace is added to the list of included namespaces, accessing the `WriteLine()` method no longer requires the preceding static class name.

All code examples requiring the `System.Console` static class from this point forward include the `using static System.Console` keyword.

## EXERCISES

- 3.1 In the following code, how would you refer to the name `great` from code in the namespace `fabulous`?

```
namespace fabulous
{
    // code in fabulous namespace
}
namespace super
{
    namespace smashing
    {
        // great name defined
    }
}
```

- 3.2 Which of the following is not a legal variable name?

- `myVariableIsGood`
- `99Flake`
- `_floor`
- `time2GetJiggyWidIt`
- `wrox.com`

- 3.3 Is the string `"supercalifragilisticexpialidocious"` too big to fit in a `string` variable? If so, why?

- 3.4 By considering operator precedence, list the steps involved in the computation of the following expression:

```
resultVar += var1 * var2 + var3 % var4 / var5;
```

- 3.5 Write a console application that obtains four `int` values from the user and displays the product. Hint: You may recall that the `Convert.ToDouble()` command was used to convert the input from the console to a `double`; the equivalent command to convert from a `string` to an `int` is `Convert.ToInt32()`.

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Basic C# syntax	C# is a case-sensitive language, and each line of code is terminated with a semicolon. Lines can be indented for ease of reading if they get too long, or to identify nested blocks. You can include non-compiled comments with // or /* ... */ syntax. Blocks of code can be collapsed into regions, also to ease readability.
Variables	Variables are chunks of data that have a name and a type. The .NET Framework defines plenty of simple types, such as numeric and string (text) types for you to use. Variables must be declared and initialized for you to use them. You can assign literal values to variables to initialize them, and variables can be declared and initialized in a single step.
Expressions	Expressions are built from operators and operands, where operators perform operations on operands. There are three types of operators — unary, binary, and ternary — that operate on 1, 2, and 3 operands, respectively. Mathematical operators perform operations on numeric values, and assignment operators place the result of an expression into a variable. Operators have a fixed precedence that determines the order in which they are processed in an expression.
Namespaces	All names defined in a .NET application, including variable names, are contained in a namespace. Namespaces are hierarchical, and you often have to qualify names according to the namespace that contains them in order to access them.

# 4

## Flow Control

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Using Boolean logic
- Branching code
- Looping code

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 4 download and individually named according to the names throughout the chapter.

All of the C# code you've seen so far has had one thing in common. In each case, program execution has proceeded from one line to the next in top-to-bottom order, missing nothing. If all applications worked like this, then you would be very limited in what you could do. This chapter describes two methods for controlling program flow — that is, the order of execution of lines of C# code: *branching* and *looping*. Branching executes code conditionally, depending on the outcome of an evaluation, such as "Execute this code only if the variable `myVal` is less than 10." Looping repeatedly executes the same statements, either a certain number of times or until a test condition has been reached.

Both of these techniques involve the use of *Boolean logic*. In the last chapter, you saw the `bool` type, but didn't actually do much with it. In this chapter, you'll use it a lot, so the chapter begins by discussing what is meant by Boolean logic, and then goes on to cover how you can use it in flow control scenarios.

## BOOLEAN LOGIC

The `bool` type introduced in the previous chapter can hold one of only two values: `true` or `false`. This type is often used to record the result of some operation, so that you can act on this result. In particular, `bool` types are used to store the result of a *comparison*.

**NOTE** As a historical aside, it is the work of the mid-nineteenth-century English mathematician George Boole that forms the basis of Boolean logic.

For instance, consider the situation (mentioned in the chapter introduction) in which you want to execute code based on whether a variable, `myVal`, is less than 10. To do this, you need some indication of whether the statement “`myVal` is less than 10” is `true` or `false` — that is, you need to know the Boolean result of a comparison.

Boolean comparisons require the use of Boolean *comparison operators* (also known as *relational operators*), which are shown in Table 4-1.

TABLE 4-1: Boolean Comparison Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
<code>==</code>	Binary	<code>var1 = var2 == var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is equal to <code>var3</code> , or <code>false</code> otherwise.
<code>!=</code>	Binary	<code>var1 = var2 != var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is not equal to <code>var3</code> , or <code>false</code> otherwise.
<code>&lt;</code>	Binary	<code>var1 = var2 &lt; var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is less than <code>var3</code> , or <code>false</code> otherwise.
<code>&gt;</code>	Binary	<code>var1 = var2 &gt; var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is greater than <code>var3</code> , or <code>false</code> otherwise.
<code>&lt;=</code>	Binary	<code>var1 = var2 &lt;= var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is less than or equal to <code>var3</code> , or <code>false</code> otherwise.
<code>&gt;=</code>	Binary	<code>var1 = var2 &gt;= var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> is greater than or equal to <code>var3</code> , or <code>false</code> otherwise.

In all cases in Table 4-1, `var1` is a `bool` type variable, whereas the types of `var2` and `var3` may vary.

You might use operators such as these on numeric values in code:

```
bool isLessThan10;
isLessThan10 = myVal < 10;
```

The preceding code results in `isLessThan10` being assigned the value `true` if `myVal` stores a value less than 10, or `false` otherwise.

You can also use these comparison operators on other types, such as strings:

```
bool isBenjamin;
isBenjamin = myString == "Benjamin";
```

Here, `isBenjamin` is `true` only if `myString` stores the string `"Benjamin"`.

You can also compare variables with Boolean values:

```
bool isTrue;
isTrue = myBool == true;
```

Here, however, you are limited to the use of the `==` and `!=` operators.

**NOTE** A common code error occurs if you unintentionally assume that because `val1 < val2` is `false`, `val1 > val2` is `true`. If `val1 == val2`, both these statements are `false`.

The `&` and `|` operators also have two similar operators, known as *conditional Boolean* operators, shown in Table 4-2.

**TABLE 4-2:** Conditional Boolean Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
<code>&amp;&amp;</code>	Binary	<code>var1 = var2 &amp;&amp; var3;</code>	<code>var1</code> is assigned the value <code>true</code> if <code>var2</code> and <code>var3</code> are both <code>true</code> , or <code>false</code> otherwise. (Logical AND)
<code>  </code>	Binary	<code>var1 = var2    var3;</code>	<code>var1</code> is assigned the value <code>true</code> if either <code>var2</code> or <code>var3</code> (or both) is <code>true</code> , or <code>false</code> otherwise. (Logical OR)

The result of these operators is exactly the same as `&` and `|`, but there is an important difference in the way this result is obtained, which can result in better performance. Both of these look at the value of their first operands (`var2` in Table 4.2) and, based on the value of this operand, may not need to process the second operands (`var3` in Table 4.2) at all.

If the value of the first operand of the `&&` operator is `false`, then there is no need to consider the value of the second operand, because the result will be `false` regardless. Similarly, the `||` operator returns `true` if its first operand is `true`, regardless of the value of the second operand.

## Boolean Bitwise and Assignment Operators

Boolean comparisons can be combined with assignments by combining Boolean bitwise and assignment operators. These work in the same way as the mathematical assignment operators that were introduced in the preceding chapter (`+=`, `*=`, and so on). The Boolean versions are shown in Table 4-3. When expressions use both the assignment (`=`) and bitwise operators (`&`, `|`, and `^`), the binary representation of the compared quantities are used to compute the outcome, instead of the integer, string, or similar values.

TABLE 4-3: Boolean Assignment Operators

OPERATOR	CATEGORY	EXAMPLE EXPRESSION	RESULT
<code>&amp;=</code>	Binary	<code>var1 &amp;= var2;</code>	<code>var1</code> is assigned the value that is the result of <code>var1 &amp; var2</code> .
<code>  =</code>	Binary	<code>var1  = var2;</code>	<code>var1</code> is assigned the value that is the result of <code>var1   var2</code> .
<code>^=</code>	Binary	<code>var1 ^= var2;</code>	<code>var1</code> is assigned the value that is the result of <code>var1 ^ var2</code> .

For example, the equation `var1 ^= var2` is similar to `var1 = var1 ^ var2` where `var1 = true` and `var2 = false`. When comparing the binary representation of `false` which is `0000` to `true`, which is typically anything other than `0000` (usually `0001`), `var1` is set to `true`.

**NOTE** Note that the `&=` and `| =` assignment operators do not make use of the `&&` and `||` conditional Boolean operators; that is, all operands are processed regardless of the value to the left of the assignment operator.

In the Try It Out that follows, you type in an integer and then the code performs various Boolean evaluations using that integer.

### TRY IT OUT Using Boolean Operators: Ch04Ex01\Program.cs

1. Create a new console application called Ch04Ex01 and save it in the directory `C:\BegVCSharp\Chapter04`.
2. Add the following code to `Program.cs`:

```
static void Main(string[] args)
{
    WriteLine("Enter an integer:");
}
```

```

        int myInt =ToInt32(ReadLine());
        bool isLessThan10 = myInt < 10;
        bool isBetween0And5 = (0 <= myInt) && (myInt <= 5);
        WriteLine($"Integer less than 10? {isLessThan10}");
        WriteLine($"Integer between 0 and 5? {isBetween0And5}");
        WriteLine($"Exactly one of the above is true?
                  {isLessThan10} ^ {isBetween0And5}");
        ReadKey();
    }
}

```

3. Execute the application and enter an integer when prompted. The result is shown in Figure 4-1.




FIGURE 4-1

## How It Works

The first two lines of code prompt for and accept an integer value using techniques you've already seen:

```

WriteLine("Enter an integer:");
int myInt =ToInt32(ReadLine());

```

You use `Convert.ToInt32()` to obtain an integer from the string input, which is simply another conversion command in the same family as the `Convert.ToDouble()` command used previously. Both the `ToInt32()` and `ToDouble()` methods are part of the `System.Convert` static class. As discussed in Chapter 3, since C# 6, it is possible to access the method of a static class directly (in this example `System.Convert`) by including the `using static System.Convert` class to the list of included namespaces. Also note that there is no check to make certain the user has actually entered an integer. If a value other than an integer is provided, for example a string, an exception would occur when trying to perform the conversion. You can handle this using a `try{...}.catch{}` block or by checking if the entered value is an integer before performing the conversion using the `GetType()` method. Both approaches are discussed in later chapters.

Next, two Boolean variables, `isLessThan10` and `isBetween0And5`, are declared and assigned values with logic that matches the description in their names:

```

bool isLessThan10 = myInt < 10;
bool isBetween0And5 = (0 <= myInt) && (myInt <= 5);

```

These variables are used in the next three lines of code, the first two of which output their values, whereas the third performs an operation on them and outputs the result. You work through this code assuming that the user enters 7, as shown in the screenshot.

The first output is the result of the operation `myInt < 10`. If `myInt` is 6, which is less than 10, the result is `true`, which is what you see displayed. Values of `myInt` of 10 or higher result in `false`.

The second output is a more involved calculation: `(0 <= myInt) && (myInt <= 5)`. It uses two comparison operations to determine whether `myInt` is greater than or equal to 0 and less than or equal to 5,

and a Boolean AND operation on the results obtained. With a value of 6, `(0 <= myInt)` returns `true`, and `(myInt <= 5)` returns `false`. The result is then `(true) && (false)`, which is `false`, as you can see from the display.

Finally, you perform a logical exclusive OR on the two Boolean variables `isLessThan10` and `isBetween0And5`. This will return `true` if one of the values is `true` and the other `false`; that is, it returns `true` only if `myInt` is 6, 7, 8, or 9. With a value of 6, as in the example, the result is `true`.

## Operator Precedence Updated

Now that you have a few more operators to consider, Table 3-10: “Operator Precedence” from the previous chapter should be updated to include them. The new order is shown in Table 4-4.

TABLE 4-4: Operator Precedence (Updated)

PRECEDENCE	OPERATORS
Highest	<code>++, --</code> (used as prefixes); <code>()</code> , <code>+, -</code> (unary), <code>!, ~</code>
	<code>*, /, %</code>
	<code>+, -</code>
	<code>&lt;&lt;, &gt;&gt;</code>
	<code>&lt;, &gt;, &lt;=, &gt;=</code>
	<code>==, !=</code>
	<code>&amp;</code>
	<code>^</code>
	<code> </code>
	<code>&amp;&amp;</code>
	<code>  </code>
	<code>=, *=, /=, %=, +=, -=, &lt;=&gt;, &amp;=, ^=,  =</code>
Lowest	<code>++, --</code> (used as suffixes)

This adds quite a few more levels but explicitly defines how expressions such as the following will be evaluated, where the `&&` operator is processed after the `<=` and `>=` operators (in this code `var2` is an `int` value):

```
var1 = var2 <= 4 && var2 >= 2;
```

It doesn't hurt to add parentheses to make expressions such as this one clearer. The compiler knows what order to process operators in, but we humans are prone to forget such things (and you might want to change the order). Writing the previous expression as

```
var1 = (var2 <= 4) && (var2 >= 2);
```

solves this problem by explicitly ordering the computation.

## BRANCHING

Branching is the act of controlling which line of code should be executed next. The line to jump to is controlled by some kind of conditional statement. This conditional statement is based on a comparison between a test value and one or more possible values using Boolean logic.

This section describes three branching techniques available in C#:

- The ternary operator
- The `if` statement
- The `switch` statement

## The Ternary Operator

The simplest way to perform a comparison is to use the *ternary* (or *conditional*) operator mentioned in the last chapter. You've already seen unary operators that work on one operand, and binary operators that work on two operands, so it won't come as a surprise that this operator works on three operands. The syntax is as follows:

```
<test> ? <resultIfTrue>: <resultIfFalse>
```

Here, `<test>` is evaluated to obtain a Boolean value, and the result of the operator is either `<resultIfTrue>` or `<resultIfFalse>` based on this value.

You might use this as follows to test the value of an `int` variable called `myInteger`:

```
string resultString = (myInteger < 10) ? "Less than 10"
                                         : "Greater than or equal to 10";
```

The result of the ternary operator is one of two strings, both of which may be assigned to `resultString`. The choice of which string to assign is made by comparing the value of `myInteger` to 10. In this case, a value of less than 10 results in the first string being assigned, and a value of greater than or equal to 10 results in the second string being assigned. For example, if `myInteger` is 4, then `resultString` will be assigned the string `Less than 10`.

## The if Statement

The `if` statement is a far more versatile and useful way to make decisions. Unlike `? :` statements, `if` statements don't have a result (so you can't use them in assignments); instead, you use the statement to conditionally execute other statements.

The simplest use of an `if` statement is as follows, where `<test>` is evaluated (it must evaluate to a Boolean value for the code to compile) and the line of code that follows the statement is executed if `<test>` evaluates to `true`:

```
if (<test>
    <code executed if <test> is true>;
```

After this code is executed, or if it isn't executed due to `<test>` evaluating to `false`, program execution resumes at the next line of code.

You can also specify additional code using the `else` statement in combination with an `if` statement. This statement is executed if `<test>` evaluates to `false`:

```
if (<test>
    <code executed if <test> is true>;
else
    <code executed if <test> is false>;
```

Both sections of code can span multiple lines using blocks in braces:

```
if (<test>
{
    <code executed if <test> is true>;
}
else
{
    <code executed if <test> is false>;
}
```

As a quick example, you could rewrite the code from the last section that used the ternary operator:

```
string resultString = (myInteger < 10) ? "Less than 10"
                                         : "Greater than or equal to 10";
```

Because the result of the `if` statement cannot be assigned to a variable, you have to assign a value to the variable in a separate step:

```
string resultString;
if (myInteger < 10)
    resultString = "Less than 10";
else
    resultString = "Greater than or equal to 10";
```

Code such as this, although more verbose, is far easier to read and understand than the equivalent ternary form, and enables far more flexibility.

The following Try It Out illustrates the use of the `if` statement.

### TRY IT OUT Using the if Statement: Ch04Ex02\Program.cs

1. Create a new console application called Ch04Ex02 and save it in the directory C:\BegVCSharp\Chapter04.
2. Add the following code to Program.cs:

```

static void Main(string[] args)
{
    string comparison;
    WriteLine("Enter a number:");
    double var1 = ToDouble(ReadLine());
    WriteLine("Enter another number:");
    double var2 = ToDouble(ReadLine());
    if (var1 < var2)
        comparison = "less than";
    else
    {
        if (var1 == var2)
            comparison = "equal to";
        else
            comparison = "greater than";
    }
    WriteLine($"The first number is
              {comparison} the second number.");
    ReadKey();
}

```

3. Execute the code and enter two numbers at the prompts (see Figure 4-2).




FIGURE 4-2

## How It Works

The first section of code is very familiar. It simply obtains two `double` values from user input:

```

string comparison;
WriteLine("Enter a number:");
double var1 = ToDouble(ReadLine());
WriteLine("Enter another number:");
double var2 = ToDouble(ReadLine());

```

Next, you assign a string to the `string` variable `comparison` based on the values obtained for `var1` and `var2`. First, you check whether `var1` is less than `var2`:

```

if (var1 < var2)
    comparison = "less than";

```

If this isn't the case, then `var1` is either greater than or equal to `var2`. In the `else` section of the first comparison, you need to nest a second comparison:

```

else
{
    if (var1 == var2)
        comparison = "equal to";
}

```

The `else` section of this second comparison is reached only if `var1` is greater than `var2`:

```
    else
        comparison = "greater than";
}
```

Finally, you write the value of `comparison` to the console:

```
WriteLine("The first number is {0} the second number.",
    comparison);
```

The nesting used here is just one method of performing these comparisons. You could equally have written this:

```
if (var1 < var2)
    comparison = "less than";
if (var1 == var2)
    comparison = "equal to";
if (var1 > var2)
    comparison = "greater than";
```

The disadvantage to this method is that you are performing three comparisons regardless of the values of `var1` and `var2`. With the first method, you perform only one comparison if `var1 < var2` is true, and two comparisons otherwise (you also perform the `var1 == var2` comparison), resulting in fewer lines of code being executed. The difference in performance here is slight, but it would be significant in applications where speed of execution is crucial.

---

## Checking More Conditions Using `if` Statements

In the preceding example, you checked for three conditions involving the value of `var1`. This covered all possible values for this variable. Sometimes, you might want to check for specific values — for example, if `var1` is equal to 1, 2, 3, or 4, and so on. Using code such as the preceding can result in annoyingly nested code:

```
if (var1 == 1)
{
    // Do something.
}
else
{
    if (var1 == 2)
    {
        // Do something else.
    }
    else
    {
        if (var1 == 3 || var1 == 4)
        {
            // Do something else.
        }
    }
}
```

```

        else
        {
            // Do something else.
        }
    }
}

```

**WARNING** It's a common mistake to write conditions such as `if (var1 == 3 || var1 == 4)` as `if (var1 == 3 || 4)`. Here, owing to operator precedence, the `==` operator is processed first, leaving the `||` operator to operate on a Boolean and a numeric operand, which causes an error.

In these situations, consider using a slightly different indentation scheme and contracting the section of code for the `else` blocks (that is, using a single line of code after the `else` blocks, rather than a block of code). That way, you end up with a structure involving `else if` statements:

```

if (var1 == 1)
{
    // Do something.
}
else if (var1 == 2)
{
    // Do something else.
}
else if (var1 == 3 || var1 == 4)
{
    // Do something else.
}
else
{
    // Do something else.
}

```

These `else if` statements are really two separate statements, and the code is functionally identical to the previous code, but much easier to read. When making multiple comparisons such as this, consider using the `switch` statement as an alternative branching structure.

## The `switch` Statement

The `switch` statement is similar to the `if` statement in that it executes code conditionally based on the value of a test. However, `switch` enables you to test for multiple values of a test variable in one go, rather than just a single condition. This test is limited to discrete values, rather than clauses such as “greater than X,” so its use is slightly different; however, it can be a powerful technique.

The basic structure of a `switch` statement is as follows:

```

switch (<testVar>
{
    case <comparisonVal1>:

```

```
<code to execute if <testVar> == <comparisonVal1> >
break;
case <comparisonVal2>:
<code to execute if <testVar> == <comparisonVal2> >
break;
...
case <comparisonValN>:
<code to execute if <testVar> == <comparisonValN> >
break;
default:
<code to execute if <testVar> != comparisonVals>
break;
}
```

The value in `<testVar>` is compared to each of the `<comparisonValX>` values (specified with `case` statements). If there is a match, then the code supplied for this match is executed. If there is no match, then the code in the `default` section is executed if this block exists.

On completion of the code in each section, you have an additional command, `break`. It is illegal for the flow of execution to reach a second `case` statement after processing one `case` block.

**NOTE** The behavior where the flow of execution is forbidden from flowing from one `case` block to the next is one area in which C# differs from C++. In C++ the processing of `case` statements is allowed to run from one to another.

The `break` statement here simply terminates the `switch` statement, and processing continues on the statement following the structure.

There are alternative methods for preventing flow from one `case` statement to the next in C# code. You can use the `return` statement, which results in termination of the current function, rather than just the `switch` structure (see Chapter 6 for more details about this), or a `goto` statement. `goto` statements (as detailed earlier) work here because `case` statements actually define labels in C# code. Here is an example:

```
switch (<testVar>)
{
    case <comparisonVal1>:
        <code to execute if <testVar> == <comparisonVal1> >
        goto case <comparisonVal2>;
    case <comparisonVal2>:
        <code to execute if <testVar> == <comparisonVal2> >
        break;
    ...
}
```

Here's one exception to the rule that the processing of one `case` statement can't run freely into the next: If you place multiple `case` statements together (*stack* them) before a single block of code, then you are in effect checking for multiple conditions at once. If any of these conditions is met, then the code is executed. Here's an example:

```
switch (<testVar>)
{
    case <comparisonVal1>:
```

```

case <comparisonVal2>:
    <code to execute if <testVar> == <comparisonVal1> or
        <testVar> == <comparisonVal2> >
    break;
...

```

These conditions also apply to the `default` statement. There is no rule stipulating that this statement must be the last in the list of comparisons, and you can stack it with `case` statements if you want. Adding a breakpoint with `break`, or `return`, ensures that a valid execution path exists through the structure in all cases.

The following Try It Out uses a `switch` statement to write different strings to the console, depending on the value you enter for a test string.

### TRY IT OUT Using the switch Statement: Ch04Ex03\Program.cs

1. Create a new console application called Ch04Ex03 and save it to the directory C:\BegVCSharp\Chapter04.
2. Add the following code to Program.cs:

```

static void Main(string[] args)
{
    const string myName = "benjamin";
    const string niceName = "andrea";
    const string sillyName = "floppy";
    string name;
    WriteLine("What is your name?");
    name = ReadLine();
    switch (name.ToLower())
    {
        case myName:
            WriteLine("You have the same name as me!");
            break;
        case niceName:
            WriteLine("My, what a nice name you have!");
            break;
        case sillyName:
            WriteLine("That's a very silly name.");
            break;
    }
    WriteLine($"Hello {name}!");
    ReadKey();
}

```

3. Execute the code and enter a name. The result is shown in Figure 4-3.




FIGURE 4-3

## How It Works

The code sets up three constant strings, accepts a string from the user, and then writes out text to the console based on the string entered. Here, the strings are names.

When you compare the name entered (in the variable `name`) to your constant values, you first force it into lowercase with `name.ToLower()`. This is a standard command that works with all string variables, and it comes in handy when you're not sure what the user entered. Using this technique, the strings `Benjamin`, `benJamin`, `benjamin`, and so on all match the test string `benjamin`.

The `switch` statement itself attempts to match the string entered with the constant values you have defined, and, if successful, writes out a personalized message to greet the user. If no match is made, you offer a generic greeting.

## LOOPING

*Looping* refers to the repeated execution of statements. This technique comes in very handy because it means that you can repeat operations as many times as you want (thousands, even millions, of times) without having to write the same code each time.

As a simple example, consider the following code for calculating the amount of money in a bank account after 10 years, assuming that interest is paid each year and no other money flows into or out of the account:

```
double balance = 1000;
double interestRate = 1.05; // 5% interest/year
balance *= interestRate;
```

Writing the same code 10 times seems a bit wasteful, and what if you wanted to change the duration from 10 years to some other value? You'd have to manually copy the line of code the required amount of times, which would be a bit of a pain! Luckily, you don't have to do this. Instead, you can have a loop that executes the instruction you want the required number of times.

Another important type of loop is one in which you loop until a certain condition is fulfilled. These loops are slightly simpler than the situation detailed previously (although no less useful), so they're a good starting point.

## do Loops

`do` loops operate as follows. The code you have marked out for looping is executed, a Boolean test is performed, and the code executes again if this test evaluates to `true`, and so on. When the test evaluates to `false`, the loop exits.

The structure of a do loop is as follows, where <Test> evaluates to a Boolean value:

```
do
{
    <code to be looped>
} while (<Test>);
```

**NOTE** The semicolon after the while statement is required.

For example, you could use the following to write the numbers from 1 to 10 in a column:

```
int i = 1;
do
{
    WriteLine("{0}", i++);
} while (i <= 10);
```

Here, you use the suffix version of the ++ operator to increment the value of i after it is written to the screen, so you need to check for *i <= 10* to include 10 in the numbers written to the console.

The following Try It Out uses this for a slightly modified version of the code shown earlier, where you calculated the balance in an account after 10 years. Here, you use a loop to calculate how many years it will take to get a specified amount of money in the account, based on a starting amount and a fixed interest rate.

### TRY IT OUT Using do Loops: Ch04Ex04\Program.cs

1. Create a new console application called Ch04Ex04 and save it to the directory C:\BegVCSharp\Chapter04.
2. Add the following code to Program.cs:

```
static void Main(string[] args)
{
    double balance, interestRate, targetBalance;
    WriteLine("What is your current balance?");
    balance = ToDouble(ReadLine());
    WriteLine("What is your current annual interest rate (in %)?");
    interestRate = 1 + ToDouble(ReadLine()) / 100.0;
    WriteLine("What balance would you like to have?");
    targetBalance = ToDouble(ReadLine());
    int totalYears = 0;
    do
    {
        balance *= interestRate;
        ++totalYears;
    }
    while (balance < targetBalance);
    WriteLine($"In {totalYears} year{(totalYears == 1 ? "" : "s")}
            you'll have a balance of {balance}.");
    ReadKey();
}
```

3. Execute the code and enter some values. A sample result is shown in Figure 4-4.

```
file:///C:/BegVCSharp/Chapter04/Ch04Ex04/bin/Debug/Ch...
What is your current balance?
1000
What is your current annual interest rate (in %)?
3.2
What balance would you like to have?
10000
In 74 years you'll have a balance of 10287.2098594564.
```

FIGURE 4-4

## How It Works

This code simply repeats the simple annual calculation of the balance with a fixed interest rate as many times as is necessary for the balance to satisfy the terminating condition. You keep a count of how many years have been accounted for by incrementing a counter variable with each loop cycle:

```
int totalYears = 0;
do
{
    balance *= interestRate;
    ++totalYears;
}
while (balance < targetBalance);
```

You can then use this counter variable as part of the result output:

```
WriteLine($"In {totalYears}
         year{(totalYears == 1 ? "" : "s")}
         you'll have a balance of {balance}.");
```

**NOTE** Perhaps the most common usage of the ?: (ternary) operator is to conditionally format text with the minimum of code. Here, you output an "s" after "year" if totalYears isn't equal to 1.

Unfortunately, this code isn't perfect. Consider what happens when the target balance is less than the current balance. The output will be similar to what is shown in Figure 4-5.

```
file:///C:/BegVCSharp/Chapter04/Ch04Ex04/bin/Debug/Ch...
What is your current balance?
10000
What is your current annual interest rate (in %)?
3.2
What balance would you like to have?
1000
In 1 year you'll have a balance of 10320.
```

FIGURE 4-5

`do` loops always execute at least once. Sometimes, as in this situation, this isn't ideal. Of course, you could add an `if` statement:

```
int totalYears = 0;
if (balance < targetBalance)
{
    do
    {
        balance *= interestRate;
        ++totalYears;
    }
    while (balance < targetBalance);
}
WriteLine($"In {totalYears} year{(totalYears == 1 ? "" : "s")}" +
    $"you'll have a balance of {balance}.");
```

Clearly, this adds unnecessary complexity. A far better solution is to use a `while` loop.

---

## while Loops

`while` loops are very similar to `do` loops, but they have one important difference: The Boolean test in a `while` loop takes place at the start of the loop cycle, not at the end. If the test evaluates to `false`, then the loop cycle is never executed. Instead, program execution jumps straight to the code following the loop.

Here's how `while` loops are specified:

```
while (<Test>)
{
    <code to be looped>
}
```

They can be used in almost the same way as `do` loops:

```
int i = 1;
while (i <= 10)
{
    WriteLine($"{i++}");
}
```

This code has the same result as the `do` loop shown earlier; it outputs the numbers 1 to 10 in a column. The following Try It Out demonstrates how you can modify the last example to use a `while` loop.

### TRY IT OUT Using while Loops: Ch04Ex05\Program.cs

1. Create a new console application called Ch04Ex05 and save it to the directory `C:\BegVCSharp\Chapter04`.
2. Modify the code as follows (use the code from Ch04Ex04 as a starting point, and remember to delete the `while` statement at the end of the original `do` loop):

```

static void Main(string[] args)
{
    double balance, interestRate, targetBalance;
    WriteLine("What is your current balance?");
    balance = ToDouble(ReadLine());
    WriteLine("What is your current annual interest rate (in %)?");
    interestRate = 1 + ToDouble(ReadLine()) / 100.0;
    WriteLine("What balance would you like to have?");
    targetBalance = ToDouble(ReadLine());
    int totalYears = 0;
    while (balance < targetBalance)
    {
        balance *= interestRate;
        ++totalYears;
    }
    WriteLine($"In {totalYears} year{(totalYears == 1 ? "" : "s")}" +
        $"you'll have a balance of {balance}.");
    if (totalYears == 0)
        WriteLine(
            "To be honest, you really didn't need to use this calculator.");
    ReadKey();
}

```

3. Execute the code again, but this time use a target balance that is less than the starting balance, as shown in Figure 4-6.




FIGURE 4-6

## How It Works

This simple change from a do loop to a while loop has solved the problem in the last example. By moving the Boolean test to the beginning, you provide for the circumstance where no looping is required, and you can jump straight to the result.

Of course, other alternatives are possible in this situation. For example, you could check the user input to ensure that the target balance is greater than the starting balance. In that case, you can place the user input section in a loop as follows:

```

WriteLine("What balance would you like to have?");
do
{
    targetBalance = ToDouble(ReadLine());
    if (targetBalance <= balance)
        WriteLine("You must enter an amount greater than " +
            "your current balance!\nPlease enter another value.");
}
while (targetBalance <= balance);

```

This rejects values that don't make sense, so the output looks like Figure 4-7.

```
file:///C:/BegVCSharp/Chapter04/Ch04Ex05/bin/Debug/Ch...
What is your current balance?
10000
What is your current annual interest rate (in %)?
5
What balance would you like to have?
1000
You must enter an amount greater than your current balance!
Please enter another value.
50000
In 52 years you'll have a balance of 51445.128831486.
```

FIGURE 4-7

This *validation* of user input is an important topic when it comes to application design. It is sometimes referred to as a *range check*, and many examples of it appear throughout this book.

## for Loops

The last type of loop to look at in this chapter is the `for` loop. This type of loop executes a set number of times and maintains its own counter. To define a `for` loop you need the following information:

- A starting value to initialize the counter variable
- A condition for continuing the loop, involving the counter variable
- An operation to perform on the counter variable at the end of each loop cycle

For example, if you want a loop with a counter that increments from 1 to 10 in steps of one, then the starting value is 1; the condition is that the counter is less than or equal to 10; and the operation to perform at the end of each cycle is to add 1 to the counter.

This information must be placed into the structure of a `for` loop as follows:

```
for (<initialization>; <condition>; <operation>)
{
    <code to loop>
}
```

This works exactly the same way as the following `while` loop:

```
<initialization>
while (<condition>)
{
    <code to loop>
    <operation>
}
```

Earlier, you used `do` and `while` loops to write out the numbers from 1 to 10. The code that follows shows what is required to do this using a `for` loop:

```
int i;
for (i = 1; i <= 10; ++i)
```

```
{  
    WriteLine($"{{i}}");  
}
```

The counter variable, an integer called `i`, starts with a value of 1 and is incremented by 1 at the end of each cycle. During each cycle, the value of `i` is written to the console.

When the code resumes after the loop, `i` has a value of 11. That's because at the end of the cycle where `i` is equal to 10, `i` is incremented to 11. This happens before the condition `i <= 10` is processed, at which point the loop ends. As with `while` loops, `for` loops execute only if the condition evaluates to `true` before the first cycle, so the code in the loop doesn't necessarily run at all.

As a final note, you can declare the counter variable as part of the `for` statement, rewriting the preceding code as follows:

```
for (int i = 1; i <= 10; ++i)  
{  
    WriteLine($"{{i}}");  
}
```

If you do this, though, the variable `i` won't be accessible from code outside this loop (see the “Variable Scope” section in Chapter 6).

## Interrupting Loops

Sometimes you want finer-grained control over the processing of looping code. C# provides commands to help you here:

- `break` — Causes the loop to end immediately
- `continue` — Causes the current loop cycle to end immediately (execution continues with the next loop cycle)
- `return` — Jumps out of the loop and its containing function (see Chapter 6)

The `break` command simply exits the loop, and execution continues at the first line of code after the loop, as shown in the following example:

```
int i = 1;  
while (i <= 10)  
{  
    if (i == 6)  
        break;  
    WriteLine($"{{i++}}");  
}
```

This code writes out the numbers from 1 to 5 because the `break` command causes the loop to exit when `i` reaches 6.

`continue` only stops the current cycle, not the whole loop, as shown here:

```
int i;  
for (i = 1; i <= 10; i++)
```

```

{
    if ((i % 2) == 0)
        continue;
    WriteLine(i);
}

```

In the preceding example, whenever the remainder of `i` divided by 2 is zero, the `continue` statement stops the execution of the current cycle, so only the numbers 1, 3, 5, 7, and 9 are displayed.

## Infinite Loops

It is possible, through both coding errors and design, to define loops that never end, so-called *infinite loops*. As a very simple example, consider the following:

```

while (true)
{
    // code in loop
}

```

This can be useful, and you can always exit such loops using code such as `break` statements or manually by using the Windows Task Manager. However, when this occurs by accident, it can be annoying. Consider the following loop, which is similar to the `for` loop in the previous section:

```

int i = 1;
while (i <= 10)
{
    if ((i % 2) == 0)
        continue;
    WriteLine($"{i++}");
}

```

Here, `i` isn't incremented until the last line of code in the loop, which occurs after the `continue` statement. If this `continue` statement is reached (which it will be when `i` is 2), the next loop cycle will be using the same value of `i`, continuing the loop, testing the same value of `i`, continuing the loop, and so on. This will cause the application to freeze. Note that it's still possible to quit the frozen application in the normal way, so you won't have to reboot if this happens.

### EXERCISES

- 4.1** If you have two integers stored in variables `var1` and `var2`, what Boolean test can you perform to determine whether one or the other (but not both) is greater than 10?
  
- 4.2** Write an application that includes the logic from Exercise 1, obtains two numbers from the user, and displays them, but rejects any input where both numbers are greater than 10 and asks for two new numbers.

**4.3** What is wrong with the following code?

```
int i;
for (i = 1; i <= 10; i++)
{
    if ((i % 2) = 0)
        continue;
    WriteLine(i);
}
```

---

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Boolean logic	Boolean logic involves using Boolean ( <code>true</code> or <code>false</code> ) values to evaluate conditions. Boolean operators are used to perform comparisons between values and return Boolean results. Some Boolean operators are also used to perform bitwise operations on the underlying bit structure of values, and there are some specialized bitwise operators too.
Branching	You can use Boolean logic to control program flow. The result of an expression that evaluates to a Boolean value can be used to determine whether a block of code is executed. You do this with <code>if</code> statements or the <code>? :</code> (ternary) operator for simple branching, or the <code>switch</code> statement to check multiple conditions simultaneously.
Looping	Looping allows you to execute blocks of code a number of times according to conditions you specify. You can use <code>do</code> and <code>while</code> loops to execute code while a Boolean expression evaluates to <code>true</code> , and <code>for</code> loops to include a counter in your looping code. Loops can be interrupted by <code>cycle</code> (with <code>continue</code> ) or completely (with <code>break</code> ). Some loops end only if you interrupt them; these are called infinite loops.



# 5

## More about Variables

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Performing implicit and explicit conversions between types
- Creating and using enum types
- Creating and using struct types
- Creating and using arrays
- Manipulating string values

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 5 download and individually named according to the names throughout the chapter.

Now that you've seen a bit more of the C# language, you can go back and tackle some of the more involved topics concerning variables.

The first subject you look at in this chapter is *type conversion*, whereby you convert values from one type into another. You've already seen a bit of this, but you look at it formally here. A grasp of this topic gives you a greater understanding of what happens when you mix types in expressions (intentionally or unintentionally), as well as tighter control over the way that data is manipulated. This helps you to streamline your code and avoid nasty surprises.

Then you'll look at a few more types of variables that you can use:

- **Enumerations** — Variable types that have a user-defined discrete set of possible values that can be used in a human-readable way.
- **Structs** — Composite variable types made up of a user-defined set of other variable types.
- **Arrays** — Types that hold multiple variables of one type, allowing index access to the individual value.

These are slightly more complex than the simple types you've been using up to now, but they can make your life much easier. Finally, you'll explore another useful subject concerning strings: basic string manipulation.

## TYPE CONVERSION

Earlier in this book, you saw that all data, regardless of type, is simply a sequence of bits — that is, a sequence of zeros and ones. The meaning of the variable is determined by the way in which this data is interpreted. The simplest example of this is the `char` type. This type represents a character in the Unicode character set using a number. In fact, the number is stored in exactly the same way as a `ushort` — both of them store a number between 0 and 65535.

However, in general, the different types of variables use varying schemes to represent data. This implies that even if it were possible to place the sequence of bits from one variable into a variable of a different type (perhaps they use the same amount of storage, or perhaps the target type has enough storage space to include all the source bits), the results might not be what you expect.

Instead of this one-to-one mapping of bits from one variable into another, you need to use *type conversion* on the data. Type conversion takes two forms:

- **Implicit conversion** — Conversion from type A to type B is possible in all circumstances, and the rules for performing the conversion are simple enough for you to trust in the compiler.
- **Explicit conversion** — Conversion from type A to type B is possible only in certain circumstances or where the rules for conversion are complicated enough to merit additional processing of some kind.

## Implicit Conversions

Implicit conversion requires no work on your part and no additional code. Consider the code shown here:

```
var1 = var2;
```

This assignment may involve an implicit conversion if the type of `var2` can be implicitly converted into the type of `var1`; however, it could just as easily involve two variables with the same type, in which case no implicit conversion is necessary. For example, the values of `ushort` and `char` are effectively interchangeable, because both store a number between 0 and 65535. You can convert values between these types implicitly, as demonstrated by the following code:

```
ushort destinationVar;
char sourceVar = 'a';
destinationVar = sourceVar;
WriteLine($"sourceVar val: {sourceVar}");
WriteLine($"destinationVar val: {destinationVar}");
```

Here, the value stored in `sourceVar` is placed in `destinationVar`. When you output the variables with the two `WriteLine()` commands, you get the following output:

```
sourceVar val: a
destinationVar val: 97
```

Even though the two variables store the same information, they are interpreted in different ways using their type.

There are many implicit conversions of simple types; `bool` and `string` have no implicit conversions, but the numeric types have a few. For reference, Table 5-1 shows the numeric conversions that the compiler can perform implicitly (remember that `chars` are stored as numbers, so `char` counts as a numeric type).

**TABLE 5-1:** Implicit Numeric Conversions

TYPE	CAN SAFELY BE CONVERTED TO
<code>byte</code>	<code>short, ushort, int, uint, long, ulong, float, double, decimal</code>
<code>sbyte</code>	<code>short, int, long, float, double, decimal</code>
<code>short</code>	<code>int, long, float, double, decimal</code>
<code>ushort</code>	<code>int, uint, long, ulong, float, double, decimal</code>
<code>int</code>	<code>long, float, double, decimal</code>
<code>uint</code>	<code>long, ulong, float, double, decimal</code>
<code>long</code>	<code>float, double, decimal</code>
<code>ulong</code>	<code>float, double, decimal</code>
<code>float</code>	<code>double</code>
<code>char</code>	<code>ushort, int, uint, long, ulong, float, double, decimal</code>

Don't worry — you don't need to learn this table by heart, because it's actually quite easy to work out which conversions the compiler can do implicitly. Back in Chapter 3, Tables 3-1, 3-2 and 3-3 showed the range of possible values for every simple numeric type. The implicit conversion rule for these types is this: Any type A whose range of possible values completely fits inside the range of possible values of type B can be implicitly converted into that type.

The reasoning for this is simple. If you try to fit a value into a variable, but that value is outside the range of values the variable can take, then there will be a problem. For example, a `short` type variable is capable of storing values up to 32767, and the maximum value allowed into a `byte` is 255, so there could be problems if you try to convert a `short` value into a `byte` value. If the `short` holds a value between 256 and 32767, then it simply won't fit into a `byte`.

If you know that the value in your `short` type variable is less than 255, then you should be able to convert the value, right? The simple answer is that, of course, you can. The slightly more complex answer is that, of course, you can, but you must use an *explicit* conversion. Performing an explicit conversion is a bit like saying “Okay, I know you’ve warned me about doing this, but I’ll take responsibility for what happens.”

## Explicit Conversions

As the name suggests, an explicit conversion occurs when you explicitly ask the compiler to convert a value from one data type to another. These conversions require extra code, and the format of this code may vary, depending on the exact conversion method. Before you look at any of this explicit conversion code, look at what happens if you *don’t* add any.

For example, the following modification to the code from the last section attempts to convert a `short` value into a `byte`:

```
byte destinationVar;
short sourceVar = 7;
destinationVar = sourceVar;
WriteLine($"sourceVar val: {sourceVar}");
WriteLine($"destinationVar val: {destinationVar}");
```

If you attempt to compile the preceding code, you will receive the following error:

```
Cannot implicitly convert type 'short' to 'byte'. An explicit conversion exists
(are you missing a cast?)
```

To get this code to compile, you need to add the code to perform an explicit conversion. The easiest way to do that in this context is to *cast* the `short` variable into a `byte` (as suggested by the preceding error string). Casting basically means forcing data from one type into another, and it uses the following simple syntax:

```
(<destinationType>) <sourceVar>
```

This will convert the value in `<sourceVar>` into `<destinationType>`.

**NOTE** Casting is only possible in some situations. Types that bear little or no relation to each other are likely not to have casting conversions defined.

You can, therefore, modify your example using this syntax to force the conversion from a `short` to a `byte`:

```
byte destinationVar;
short sourceVar = 7;
destinationVar = (byte)sourceVar;
WriteLine($"sourceVar val: {sourceVar}");
WriteLine($"destinationVar val: {destinationVar}");
```

This results in the following output:

```
sourceVar val: 7
destinationVar val: 7
```

What happens when you try to force a value into an incompatible variable type? For example, you can't fit a large integer into a numeric type that's too small. Modifying your code as follows illustrates this:

```
byte destinationVar;
short sourceVar = 281;
destinationVar = (byte)sourceVar;
WriteLine($"sourceVar val: {sourceVar}");
WriteLine($"destinationVar val: {destinationVar}");
```

This results in the following:

```
sourceVar val: 281
destinationVar val: 25
```

What happened? Well, look at the binary representations of these two numbers, along with the maximum value that can be stored in a byte, which is 255:

```
281 = 100011001
25 = 000011001
255 = 011111111
```

You can see that the leftmost bit of the source data has been lost. This immediately raises a question: How can you tell when this happens? Obviously, there will be times when you will need to explicitly cast one type into another, and it would be nice to know if any data has been lost along the way. Not detecting this could cause serious errors — for example, in an accounting application or an application determining the trajectory of a rocket to the moon.

One way to do this is simply to check the value of the source variable and compare it with the known limits of the destination variable. Another technique is to force the system to pay special attention to the conversion at runtime. Attempting to fit a value into a variable when that value is too big for the type of that variable results in an *overflow*, and this is the situation you want to check for.

Two keywords exist for setting what is called the *overflow checking context* for an expression: `checked` and `unchecked`. You use these in the following way:

```
checked(<expression>)
unchecked(<expression>)
```

You can force overflow checking in the last example:

```
byte destinationVar;
short sourceVar = 281;
destinationVar = checked((byte)sourceVar);
WriteLine($"sourceVar val: {sourceVar}");
WriteLine($"destinationVar val: {destinationVar}");
```

When this code is executed, it will crash with the error message shown in Figure 5-1 (this was compiled in a project called `OverflowCheck`).




FIGURE 5-1

However, if you replace `checked` with `unchecked` in this code, you get the result shown earlier, and no error occurs. That is identical to the default behavior, also shown earlier.

You also can configure your application to behave as if every expression of this type includes the `checked` keyword, unless that expression explicitly uses the `unchecked` keyword (in other words, you can change the default setting for overflow checking). To do this, you modify the properties for your project by right-clicking on it in the Solution Explorer window and selecting the Properties option. Click Build on the left side of the window to bring up the Build settings.

The property you want to change is one of the Advanced settings, so click the Advanced button. In the dialog box that appears, enable the Check for Arithmetic Overflow/Underflow box, as shown in Figure 5-2. By default, this setting is disabled; enabling it provides the `checked` behavior detailed previously.





FIGURE 5-2

## Explicit Conversions Using the Convert Commands

The type of explicit conversion you have been using in many of the Try It Out examples in this book is a bit different from those you have seen so far in this chapter. You have been converting string values into numbers using commands such as `ToDouble()`, which is obviously something that won't work for every possible string.

If, for example, you try to convert a string like `Number` into a double value using `ToDouble()`, you will see the dialog box shown in Figure 5-3 when you execute the code.



**FIGURE 5-3**

As you can see, the operation fails. For this type of conversion to work, the string supplied *must* be a valid representation of a number, and that number must be one that won't cause an overflow. A valid representation of a number is one that contains an optional sign (that is, plus or minus), zero or more digits, an optional period followed by one or more digits, and an optional “e” or “E” followed by an optional sign, one or more digits, and nothing else except spaces (before or after this sequence). Using all of these optional extras, you can recognize strings as complex as `-1.2451e-24` as being a number.

The important thing to note about these conversions is that they are *always* overflow-checked, and the `checked` and `unchecked` keywords and project property settings have no effect.

The next Try It Out is an example that covers many of the conversion types from this section. It declares and initializes a number of variables of different types and then converts between them implicitly and explicitly.

### TRY IT OUT Type Conversions in Practice: Ch05Ex01\Program.cs

1. Create a new console application called Ch05Ex01 and save it in the directory `C:\BegVCSharp\Chapter05`.
2. Add the following code to `Program.cs`:

```
static void Main(string[] args)
{
    short shortResult, shortVal = 4;
```

```

int     integerVal = 67;
long    longResult;
float   floatVal = 10.5F;
double  doubleResult, doubleVal = 99.999;
string  stringResult, stringVal = "17";
bool    boolVal = true;
WriteLine("Variable Conversion Examples\n");
doubleResult = floatVal * shortVal;
WriteLine($"Implicit, -> double: {floatVal} * {shortVal} -> {doubleResult}");
shortResult = (short)floatVal;
WriteLine($"Explicit, -> short: {floatVal} -> {shortResult}");
stringResult = Convert.ToString(boolVal) +
    Convert.ToString(doubleVal);
WriteLine($"Explicit, -> string: \"{boolVal}\" + \"{doubleVal}\" -> " +
    $"{stringResult}");
longResult = integerVal +ToInt64(stringVal);
WriteLine($"Mixed, -> long: {integerVal} + {stringVal} -> {longResult}");
.ReadKey();
}

```

3. Execute the code. The result is shown in Figure 5-4.




FIGURE 5-4

## How It Works

This example contains all of the conversion types you've seen so far — both in simple assignments, as in the short code examples in the preceding discussion, and in expressions. You need to consider both cases because the processing of *every* non-unary operator may result in type conversions, not just assignment operators. For example, the following multiplies a `short` value by a `float` value:

```
shortVal * floatVal
```

In situations such as this, where no explicit conversion is specified, implicit conversion will be used if possible. In this example, the only implicit conversion that makes sense is to convert the `short` into a `float` (as converting a `float` into a `short` requires explicit conversion), so this is the one that will be used.

However, you can override this behavior should you want, as shown here:

```
shortVal * (short)floatVal
```

**NOTE** Interestingly, multiplying two short values together doesn't return a short value. Because the result of this operation is quite likely to exceed 32767 (the maximum value a short can hold), it actually returns an int.

The conversion process can seem complex at first glance, but as long as you break expressions down into parts by taking the operator precedence order into account, you should be able to work things out.

## COMPLEX VARIABLE TYPES

In addition to all the simple variable types, C# also offers three slightly more complex (but very useful) sorts of variables: enumerations (often referred to as enums), structs (occasionally referred to as structures), and arrays.

### Enumerations

Each of the types you've seen so far (with the exception of `string`) has a clearly defined set of allowed values. Admittedly, this set is so large in types such as `double` that it can practically be considered a continuum, but it is a fixed set nevertheless. The simplest example of this is the `bool` type, which can take only one of two values: `true` or `false`.

There are many other circumstances in which you might want to have a variable that can take one of a fixed set of results. For example, you might want to have an `orientation` type that can store one of the values `north`, `south`, `east`, or `west`.

In situations like this, *enumerations* can be very useful. Enumerations do exactly what you want in this `orientation` type: They allow the definition of a type that can take one of a finite set of values that you supply. What you need to do, then, is create your own enumeration type called `orientation` that can take one of the four possible values.

Note that there is an additional step involved here — you don't just declare a variable of a given type; you declare and detail a user-defined type and then declare a variable of this new type.

### Defining Enumerations

You can use the `enum` keyword to define enumerations as follows:

```
enum <typeName>
{
    <value1>,
    <value2>,
    <value3>,
    ...
    <valueN>
}
```

Next, you can declare variables of this new type as follows:

```
<typeName> <varName>;
```

You can assign values using the following:

```
<varName> = <typeName>.<value>;
```

Enumerations have an *underlying type* used for storage. Each of the values that an enumeration type can take is stored as a value of this underlying type, which by default is `int`. You can specify a different underlying type by adding the type to the enumeration declaration:

```
enum <typeName> : <underlyingType>
{
    <value1>,
    <value2>,
    <value3>,
    ...
    <valueN>
}
```

Enumerations can have underlying types of `byte`, `sbyte`, `short`, `ushort`, `int`, `uint`, `long`, and `ulong`.

By default, each value is assigned a corresponding underlying type value automatically according to the order in which it is defined, starting from zero. This means that `<value1>` gets the value 0, `<value2>` gets 1, `<value3>` gets 2, and so on. You can override this assignment by using the `=` operator and specifying actual values for each enumeration value:

```
enum <typeName> : <underlyingType>
{
    <value1> = <actualVal1>,
    <value2> = <actualVal2>,
    <value3> = <actualVal3>,
    ...
    <valueN> = <actualValN>
}
```

In addition, you can specify identical values for multiple enumeration values by using one value as the underlying value of another:

```
enum <typeName> : <underlyingType>
{
    <value1> = <actualVal1>,
    <value2> = <value1>,
    <value3>,
    ...
    <valueN> = <actualValN>
}
```

Any values left unassigned are given an underlying value automatically, whereby the values used are in a sequence starting from 1 greater than the last explicitly declared one. In the preceding code, for example, `<value3>` will get the value `<value1> + 1`.

Note that this can cause problems, with values specified after a definition such as `<value2> = <value1>` being identical to other values. For example, in the following code `<value4>` will have the same value as `<value2>`:

```
enum <typeName> : <underlyingType>
{
    <value1> = <actualVal1>,
    <value2>,
    <value3> = <value1>,
    <value4>,
    ...
    <valueN> = <actualValN>
}
```

Of course, if this is the behavior you want, then this code is fine. Note also that assigning values in a circular fashion will cause an error:

```
enum <typeName> : <underlyingType>
{
    <value1> = <value2>,
    <value2> = <value1>
}
```


The following Try It Out shows an example of all of this. The code defines and then uses an enumeration called `orientation`.

### TRY IT OUT Using an Enumeration: Ch05Ex02\Program.cs

1. Create a new console application called Ch05Ex02 and save it in the directory C:\BegVCSharp\Chapter05.
2. Add the following code to Program.cs:

```
namespace Ch05Ex02
{
    enum orientation : byte
    {
        north = 1,
        south = 2,
        east = 3,
        west = 4
    }
    class Program
    {
        static void Main(string[] args)
        {
            orientation myDirection = orientation.north;
            WriteLine($"myDirection = {myDirection}");
            ReadKey();
        }
    }
}
```

3. Execute the application. You should see the output shown in Figure 5-5.




```
file:///C:/BegVCSharp/Chapter05/Ch05Ex02/bin/Debug/Ch...
myDirection = north
```

FIGURE 5-5

4. Quit the application and modify the code as follows:

```
byte directionByte;
string directionString;
orientation myDirection = orientation.north;
WriteLine($"myDirection = {myDirection}");
directionByte = (byte)myDirection;
directionString = Convert.ToString(myDirection);
WriteLine($"byte equivalent = {directionByte}");
WriteLine($"string equivalent = {directionString}");
ReadKey();
```

5. Execute the application again. The output is shown in Figure 5-6.



```
file:///C:/BegVCSharp/Chapter05/Ch05Ex02/bin/Debug/Ch...
myDirection = north
byte equivalent = 1
string equivalent = north
```

FIGURE 5-6

## How It Works

This code defines and uses an enumeration type called `orientation`. The first thing to notice is that the type definition code is placed in your namespace, `Ch05Ex02`, but not in the same place as the rest of your code. That is because definitions are not executed; that is, at runtime you don't step through the code in a definition as you do the lines of code in your application. Application execution starts in the place you're used to and has access to your new type because it belongs to the same namespace.

The first iteration of the example demonstrates the basic method of creating a variable of your new type, assigning it a value, and outputting it to the screen. Next, you modify the code to show the conversion of enumeration values into other types. Note that you must use explicit conversions here. Even though the underlying type of `orientation` is `byte`, you still have to use the `(byte)` cast to convert the value of `myDirection` into a `byte` type:

```
directionByte = (byte)myDirection;
```

The same explicit casting is necessary in the other direction, too, if you want to convert a `byte` into an `orientation`. For example, you could use the following code to convert a `byte` variable called `myByte` into an `orientation` and assign this value to `myDirection`:

```
myDirection = (orientation)myByte;
```

Of course, you must be careful here because not all permissible values of `byte` type variables map to defined `Orientation` values. The `Orientation` type can store other `byte` values, so you won't get an error straight away, but this may break logic later in the application.

To get the string value of an enumeration value you can use `Convert.ToString()`:

```
directionString = Convert.ToString(myDirection);
```

Using a `(string)` cast won't work because the processing required is more complicated than just placing the data stored in the enumeration variable into a `string` variable. Alternatively, you can use the `ToString()` command of the variable itself. The following code gives you the same result as using `Convert.ToString()`:

```
directionString = myDirection.ToString();
```

Converting a `string` to an enumeration value is also possible, except that here the syntax required is slightly more complex. A special command exists for this sort of conversion, `Enum.Parse()`, which is used in the following way:

```
(enumerationType)Enum.Parse(typeof(enumerationType), enumerationValueString);
```

This uses another operator, `typeof`, which obtains the type of its operand. You could use this for your `Orientation` type as follows:

```
string myString = "north";
orientation myDirection = (orientation)Enum.Parse(typeof(orientation),
                                                       myString);
```

Of course, not all string values will map to an `Orientation` value! If you pass in a value that doesn't map to one of your enumeration values, you will get an error. Like everything else in C#, these values are case sensitive, so you still get an error if your string agrees with a value in everything but case (for example, if `myString` is set to `North` rather than `north`).

## Structs

The *struct* (short for structure) is just that. That is, structs are data structures composed of several pieces of data, possibly of different types. They enable you to define your own types of variables based on this structure. For example, suppose that you want to store the route to a location from a starting point, where the route consists of a direction and a distance in miles. For simplicity, you can assume that the direction is one of the compass points (such that it can be represented using the `Orientation` enumeration from the last section), and that distance in miles can be represented as a `double` type.

You could use two separate variables for this using code you've seen already:

```
orientation myDirection;
double      myDistance;
```

There is nothing wrong with using two variables like this, but it is far simpler (especially where multiple routes are required) to store this information in one place.

## Defining Structs

Structs are defined using the `struct` keyword as follows:

```
struct <typeName>
{
    <memberDeclarations>
}
```

The `<memberDeclarations>` section contains declarations of variables (called the *data members* of the struct) in almost the same format as usual. Each member declaration takes the following form:

```
<accessibility> <type> <name>;
```

To allow the code that calls the struct to access the struct's data members, you use the keyword `public` for `<accessibility>`. For example:

```
struct route
{
    public orientation direction;
    public double      distance;
}
```

Once you have a struct type defined, you use it by defining variables of the new type:

```
route myRoute;
```

In addition, you have access to the data members of this composite variable via the period character:

```
myRoute.direction = orientation.north;
myRoute.distance  = 2.5;
```

This is demonstrated in the following Try It Out, where the `orientation` enumeration from the last Try It Out is used with the `route` struct shown earlier. This struct is then manipulated in code to give you a feel for how structs work.

### TRY IT OUT Using a Struct: Ch05Ex03\Program.cs

1. Create a new console application called Ch05Ex03 and save it in the directory C:\BegVCSharp\Chapter05.
2. Add the following code to Program.cs:

```
namespace Ch05Ex03
{
    enum orientation: byte
    {
        north = 1,
        south = 2,
        east  = 3,
        west  = 4
    }
    struct route
    {
```

```

        public orientation direction;
        public double      distance;
    }
class Program
{
    static void Main(string[] args)
    {
        route myRoute;
        int myDirection = -1;
        double myDistance;
        WriteLine("1) North\n2) South\n3) East\n4) West");
        do
        {
            WriteLine("Select a direction:");
            myDirection =ToInt32(ReadLine());
        }
        while ((myDirection < 1) || (myDirection > 4));
        WriteLine("Input a distance:");
        myDistance = ToDouble(ReadLine());
        myRoute.direction = (orientation)myDirection;
        myRoute.distance = myDistance;
        WriteLine($"myRoute specifies a direction of {myRoute.direction} " +
                 $"and a distance of {myRoute.distance}");
        ReadKey();
    }
}
}

```

3. Execute the code, select a direction by entering a number between 1 and 4, and then enter a distance. The result is shown in Figure 5-7.




FIGURE 5-7

## How It Works

Structs, like enumerations, are declared outside of the main body of the code. You declare your `route` struct just inside the namespace declaration, along with the `orientation` enumeration that it uses:

```

enum orientation: byte
{
    north = 1,
    south = 2,
    east  = 3,
    west  = 4
}

```

```
struct route
{
    public orientation direction;
    public double      distance;
}
```

The main body of the code follows a structure similar to some of the example code you've already seen, requesting input from the user and displaying it. You perform some simple validation of user input by placing the direction selection in a do loop, rejecting any input that isn't an integer between 1 and 4 (with values chosen such that they map onto the enumeration members for easy assignment).

**NOTE** *Input that cannot be interpreted as an integer will result in an error. You'll see why this happens, and what to do about it, later in the book.*

*The interesting point to note is that when you refer to members of route they are treated exactly the same way that variables of the same type as the members are. The assignment is as follows:*

```
myRoute.direction = (orientation)myDirection;
myRoute.distance = myDistance;
```

*You could simply take the input value directly into myRoute.distance with no ill effects as follows:*

```
myRoute.distance = ToDouble(ReadLine());
```

*The extra step allows for more validation, although none is performed in this code. Any access to members of a structure is treated in the same way. Expressions of the form <structVar>.<memberVar> can be said to evaluate to a variable of the type of <memberVar>.*

---

## Arrays

All the types you've seen so far have one thing in common: Each of them stores a single value (or a single set of values in the case of structs). Sometimes, in situations where you want to store a lot of data, this isn't very convenient. You may want to store several values of the same type at the same time, without having to use a different variable for each value.

For example, suppose you want to perform some processing that involves the names of all your friends. You could use simple string variables as follows:

```
string friendName1 = "Todd Anthony";
string friendName2 = "Kevin Holton";
string friendName3 = "Shane Laigle";
```

But this looks like it will require a lot of effort, especially because you need to write different code to process each variable. You couldn't, for example, iterate through this list of strings in a loop.

The alternative is to use an *array*. Arrays are indexed lists of variables stored in a single array type variable. For example, you might have an array called `friendNames` that stores the three names shown in the preceding string variables. You can access individual members of the array by specifying their index in square brackets, as shown here:

```
friendNames[<index>]
```

The index is simply an integer, starting with 0 for the first entry, using 1 for the second, and so on. This means that you can go through the entries using a loop:

```
int i;
for (i = 0; i < 3; i++)
{
    WriteLine($"Name with index of {i}: {friendNames[i]}");
}
```

Arrays have a single *base type* — that is, individual entries in an array are all of the same type. This `friendNames` array has a base type of `string` because it is intended for storing `string` variables. Array entries are often referred to as *elements*.

## Declaring Arrays

Arrays are declared in the following way:

```
<baseType>[] <name>;
```

Here, `<baseType>` may be any variable type, including the enumeration and struct types you've seen in this chapter. Arrays must be initialized before you have access to them. You can't just access or assign values to the array elements like this:

```
int[] myIntArray;
myIntArray[10] = 5;
```

Arrays can be initialized in two ways. You can either specify the complete contents of the array in a literal form or specify the size of the array and use the `new` keyword to initialize all array elements.

Specifying an array using literal values simply involves providing a comma-separated list of element values enclosed in curly braces:

```
int[] myIntArray = { 5, 9, 10, 2, 99 };
```

Here, `myIntArray` has five elements, each with an assigned integer value.

The other method requires the following syntax:

```
int[] myIntArray = new int[5];
```

Here, you use the `new` keyword to explicitly initialize the array, and a constant value to define the size. This method results in all the array members being assigned a default value, which is 0 for numeric types. You can also use nonconstant variables for this initialization:

```
int[] myIntArray = new int[arraySize];
```

In addition, you can combine these two methods of initialization if you want:

```
int[] myIntArray = new int[5] { 5, 9, 10, 2, 99 };
```

With this method the sizes *must* match. You can't, for example, write the following:

```
int[] myIntArray = new int[10] { 5, 9, 10, 2, 99 };
```

Here, the array is defined as having 10 members, but only five are defined, so compilation will fail. A side effect of this is that if you define the size using a variable, then that variable must be a constant:

```
const int arraySize = 5;
int[] myIntArray = new int[arraySize] { 5, 9, 10, 2, 99 };
```

If you omit the `const` keyword, this code will fail.

As with other variable types, there is no need to initialize an array on the same line that you declare it. The following is perfectly legal:

```
int[] myIntArray;
myIntArray = new int[5];
```

In the following Try It Out you create and use an array of strings, using the example from the introduction to this section.

### TRY IT OUT Using an Array: Ch05Ex04\Program.cs

1. Create a new console application called Ch05Ex04 and save it in the directory C:\BegVCSharp\Chapter05.
2. Add the following code to Program.cs:

```
static void Main(string[] args)
{
    string[] friendNames = { "Todd Anthony", "Kevin Holton",
                            "Shane Laigle" };
    int i;
    WriteLine($"Here are {friendNames.Length} of my friends:");
    for (i = 0; i < friendNames.Length; i++)
    {
        WriteLine(friendNames[i]);
    }
    ReadKey();
}
```

3. Execute the code. The result is shown in Figure 5-8.




FIGURE 5-8

## How It Works

This code sets up a `string` array with three values and lists them in the console in a `for` loop. Note that you have access to the number of elements in the array using `friendNames.Length`:


```
WriteLine($"Here are {friendNames.Length} of my friends:");
```

This is a handy way to get the size of an array. Outputting values in a `for` loop is easy to get wrong. For example, try changing `<` to `<=` as follows:

```
for (i = 0; i <= friendNames.Length; i++)
{
    WriteLine(friendNames[i]);
}
```

Compiling and executing the preceding code results in the dialog box shown in Figure 5-9.

Here, the code attempted to access `friendNames[3]`. Remember that array indices start from 0, so the last element is `friendNames[2]`. If you attempt to access elements outside of the array size, the code will fail. It just so happens that there is a more resilient method of accessing all the members of an array: using `foreach` loops.



**FIGURE 5-9**

## foreach Loops

A `foreach` loop enables you to address each element in an array using this simple syntax:

```
foreach (<baseType> <name> in <array>)
{
    // can use <name> for each element
}
```

This loop will cycle through each element, placing it in the variable `<name>` in turn, without danger of accessing illegal elements. You don't have to worry about how many elements are in the array, and you can be sure that you'll get to use each one in the loop. Using this approach, you can modify the code in the last example as follows:

```
static void Main(string[] args)
{
    string[] friendNames = { "Todd Anthony", "Kevin Holton",
                            "Shane Laigle" };
    WriteLine($"Here are {friendNames.Length} of my friends:");
    foreach (string friendName in friendNames)
    {
        WriteLine(friendName);
    }
    ReadKey();
}
```

The output of this code will be exactly the same as that of the previous Try It Out. The main difference between using this method and a standard `for` loop is that `foreach` gives you *read-only* access to the array contents, so you can't change the values of any of the elements. You couldn't, for example, do the following:

```
foreach (string friendName in friendNames)
{
    friendName = "Rupert the bear";
}
```

If you try this, compilation will fail. If you use a simple `for` loop, however, you can assign values to array elements.

## Multidimensional Arrays

A multidimensional array is simply one that uses multiple indices to access its elements. For example, suppose you want to plot the height of a hill against the position measured. You might specify a position using two coordinates, `x` and `y`. You want to use these two coordinates as indices, such that an array called `hillHeight` would store the height at each pair of coordinates. This involves using multidimensional arrays.

A two-dimensional array such as this is declared as follows:

```
<baseType>[,] <name>;
```

Arrays of more dimensions simply require more commas:

```
<baseType>[,,] <name>;
```

This would declare a four-dimensional array. Assigning values also uses a similar syntax, with commas separating sizes. Declaring and initializing the two-dimensional array `hillHeight`, with a base type of `double`, an `x` size of 3, and a `y` size of 4 requires the following:

```
double[,] hillHeight = new double[3,4];
```

Alternatively, you can use literal values for initial assignment. Here, you use nested blocks of curly braces, separated by commas:

```
double[,] hillHeight = { { 1, 2, 3, 4 }, { 2, 3, 4, 5 }, { 3, 4, 5, 6 } };
```

This array has the same dimensions as the previous one — that is, three rows and four columns. By providing literal values, these dimensions are defined implicitly.

To access individual elements of a multidimensional array, you simply specify the indices separated by commas:

```
hillHeight [2,1]
```

You can then manipulate this element just as you can other elements. This expression will access the second element of the third nested array as defined previously (the value will be 4). Remember that you start counting from 0 and that the first number is the nested array. In other words, the first number specifies the pair of curly braces, and the second number specifies the element within that pair of braces. You can represent this array visually, as shown in Figure 5-10.




FIGURE 5-10

The `foreach` loop gives you access to all elements in a multidimensional way, just as with single-dimensional arrays:

```
double[,] hillHeight = { { 1, 2, 3, 4 }, { 2, 3, 4, 5 }, { 3, 4, 5, 6 } };
foreach (double height in hillHeight)
{
    WriteLine("{0}", height);
}
```

The order in which the elements are output is the same as the order used to assign literal values. This sequence is as follows (the element identifiers are shown here rather than the actual values):

```
hillHeight [0,0]
hillHeight [0,1]
hillHeight [0,2]
```

```
hillHeight[0,3]
hillHeight[1,0]
hillHeight[1,1]
hillHeight[1,2]
...
...
```

## Arrays of Arrays

Multidimensional arrays, as discussed in the last section, are said to be *rectangular* because each “row” is the same size. Using the last example, you can have a y coordinate of 0 to 3 for any of the possible x coordinates.

It is also possible to have *jagged* arrays, whereby “rows” may be different sizes. For this, you need an array in which each element is another array. You could also have arrays of arrays of arrays, or even more complex situations. However, all this is possible only if the arrays have the same base type.

The syntax for declaring arrays of arrays involves specifying multiple sets of square brackets in the declaration of the array, as shown here:

```
int[][] jaggedIntArray;
```

Unfortunately, initializing arrays such as this isn’t as simple as initializing multidimensional arrays. You can’t, for example, follow the preceding declaration with this:

```
jaggedIntArray = new int[3][4];
```

Even if you could do this, it wouldn’t be that useful because you can achieve the same effect with simple multidimensional arrays with less effort. Nor can you use code such as this:

```
jaggedIntArray = { { 1, 2, 3 }, { 1 }, { 1, 2 } };
```

You have two options. You can initialize the array that contains other arrays (let’s call these sub-arrays for clarity) and then initialize the sub-arrays in turn:

```
jaggedIntArray = new int[2][];
jaggedIntArray[0] = new int[3];
jaggedIntArray[1] = new int[4];
```

Alternatively, you can use a modified form of the preceding literal assignment:

```
jaggedIntArray = new int[3][] { new int[] { 1, 2, 3 }, new int[] { 1 },
                                new int[] { 1, 2 } };
```

This can be simplified if the array is initialized on the same line as it is declared, as follows:

```
int[][] jaggedIntArray = { new int[] { 1, 2, 3 }, new int[] { 1 },
                           new int[] { 1, 2 } };
```

You can use `foreach` loops with jagged arrays, but you often need to nest these to get to the actual data. For example, suppose you have the following jagged array that contains 10 arrays, each of which contains an array of integers that are divisors of an integer between 1 and 10:

```
int[][] divisors1To10 = { new int[] { 1 },
                           new int[] { 1, 2 },
```

```

        new int[] { 1, 3 },
        new int[] { 1, 2, 4 },
        new int[] { 1, 5 },
        new int[] { 1, 2, 3, 6 },
        new int[] { 1, 7 },
        new int[] { 1, 2, 4, 8 },
        new int[] { 1, 3, 9 },
        new int[] { 1, 2, 5, 10 } };
    
```

The following code will fail:

```

foreach (int divisor in divisors1To10)
{
    WriteLine(divisor);
}
    
```

The failure occurs because the array `divisors1To10` contains `int []` elements, not `int` elements. Instead, you have to loop through every sub-array, as well as through the array itself:

```

foreach (int[] divisorsOfInt in divisors1To10)
{
    foreach(int divisor in divisorsOfInt)
    {
        WriteLine(divisor);
    }
}
    
```

As you can see, the syntax for using jagged arrays can quickly become complex! In most cases, it is easier to use rectangular arrays or a simpler storage method. Nonetheless, there may well be situations in which you are forced to use this method, and a working knowledge can't hurt. An example of this happens when working with XML documents where some elements have sub-children and other do not.

## STRING MANIPULATION

Your use of strings so far has consisted of writing strings to the console, reading strings from the console, and concatenating strings using the `+` operator. In the course of programming more interesting applications, you will discover that manipulating strings is something that you end up doing *a lot*. Therefore, it is worth spending a few pages looking at some of the more common string-manipulation techniques available in C#.

To start with, a `string` type variable can be treated as a read-only array of `char` variables. This means that you can access individual characters using syntax like the following:

```

string myString = "A string";
char myChar = myString[1];
    
```

However, you can't assign individual characters this way. To get a `char` array that you can write to, you can use the following code. This uses the `ToCharArray()` command of the array variable:

```

string myString = "A string";
char[] myChars = myString.ToCharArray();
    
```

Then you can manipulate the `char` array the standard way. You can also use strings in `foreach` loops, as shown here:

```
foreach (char character in myString)
{
    WriteLine($"{character}");
}
```

As with arrays, you can also get the number of elements using `myString.Length`. This gives you the number of characters in the string:

```
string myString = ReadLine();
WriteLine($"You typed {myString.Length} characters.");
```

Other basic string manipulation techniques use commands with a format similar to this `<string>.ToCharArray()` command. Two simple, but useful, ones are `<string>.ToLower()` and `<string>.ToUpper()`. These enable strings to be converted into lowercase and uppercase, respectively. To see why this is useful, consider the situation in which you want to check for a specific response from a user — for example, the string `yes`. If you convert the string entered by the user into lowercase, then you can also check for the strings `YES`, `Yes`, `yes`, and so on — you saw an example of this in the previous chapter:

```
string userResponse = ReadLine();
if (userResponse.ToLower() == "yes")
{
    // Act on response.
}
```

This command, like the others in this section, doesn't actually change the string to which it is applied. Instead, combining this command with a string results in the creation of a new string, which you can compare to another string (as shown here) or assign to another variable. The other variable may be the same one that is being operated on:

```
userResponse = userResponse.ToLower();
```

This is an important point to remember, because just writing

```
userResponse.ToLower();
```

doesn't actually achieve very much!

There are other things you can do to ease the interpretation of user input. What if the user accidentally put an extra space at the beginning or end of the input? In this case, the preceding code won't work. You need to trim the string entered, which you can do using the `<string>.Trim()` command:

```
string userResponse = ReadLine();
userResponse = userResponse.Trim();
if (userResponse.ToLower() == "yes")
{
    // Act on response.
}
```

The preceding code is also able to detect strings like this:

```
" YES"
"Yes "
```

You can also use these commands to remove any other characters, by specifying them in a `char` array, for example:

```
char[] trimChars = { ' ', 'e', 's' };
string userResponse = ReadLine();
userResponse = userResponse.ToLower();
userResponse = userResponse.Trim(trimChars);
if (userResponse == "y")
{
    // Act on response.
}
```

This eliminates any occurrences of spaces, as well as the letters "e" and "s" from the beginning or end of your string. Providing there aren't any other characters in the string, this will result in the detection of strings such as

```
"Yeeeees"
" y"
```

and so on.

You can also use the `<string>.TrimStart()` and `<string>.TrimEnd()` commands, which will trim spaces from the beginning and end of a string, respectively. These can also have `char` arrays specified.

You can use two other string commands to manipulate the spacing of strings: `<string>.PadLeft()` and `<string>.PadRight()`. They enable you to add spaces to the left or right of a string to force it to the desired length. You use them as follows:

```
<string>.PadX(<desiredLength>);
```

Here is an example:

```
myString = "Aligned";
myString = myString.PadLeft(10);
```

This would result in three spaces being added to the left of the word `Aligned` in `myString`. These methods can be helpful when aligning strings in columns, which is particularly useful for positioning strings containing numbers.

As with the trimming commands, you can also use these commands in a second way, by supplying the character to pad the string with. This involves a single `char`, not an array of `chars` as with trimming:

```
myString = "Aligned";
myString = myString.PadLeft(10, '-');
```

This would add three dashes to the start of `myString`.

There are many more of these string-manipulation commands, many of which are only useful in very specific situations. These are discussed as you use them in the forthcoming chapters. Before moving on, though, it is worth looking at one of the features contained in Visual Studio 2015 that you may have noticed over the course of the last few chapters, and especially this one. In the following Try It Out, you examine auto-completion, whereby the IDE tries to help you out by suggesting what code you might like to insert.

### TRY IT OUT Statement Auto-Completion in Visual Studio: Ch05Ex05\Program.cs

1. Create a new console application called Ch05Ex05 and save it in the directory C:\BegVCSharp\Chapter05.
2. Type the following code into `Program.cs`, exactly as written, noting windows that pop up as you do so:

```
static void Main(string[] args)
{
    string myString = "This is a test.";
    char[] separator = {' '};
    string[] myWords;
    myWords = myString.
```

3. As you type the final period, the window shown in Figure 5-11 appears.



FIGURE 5-11

4. Without moving the cursor, type `sp`. The pop-up window changes, and the Tooltip shown in Figure 5-12 appears.



FIGURE 5-12

5. Type the following characters: `(se`. Another pop-up window and Tooltip appears, as shown in Figure 5-13.
6. Then, type these two characters: `) ;`. The code should look as follows, and the pop-up windows should disappear:

```
static void Main(string[] args)
```

```

{
    string myString = "This is a test.";
    char[] separator = {' '};
    string[] myWords;
    myWords = myString.Split(separator);
}

```




FIGURE 5-13

7. Add the following code, noting the windows as they pop up:

```

static void Main(string[] args)
{
    string myString = "This is a test.";
    char[] separator = {' '};
    string[] myWords;
    myWords = myString.Split(separator);
    foreach (string word in myWords)
    {
        WriteLine(${word});
    }
    ReadKey();
}

```

8. Execute the code. The result is shown in Figure 5-14.




FIGURE 5-14

### How It Works

Two main aspects of this code are the new string command used and the use of the auto-completion functionality. The command, `<string>.Split()`, converts a `string` into a `string` array by splitting it at the points specified. These points take the form of a `char` array, which in this case is simply populated by a single element, the space character:

```
char[] separator = {' '};
```

The following code obtains the substrings you get when the string is split at each space — that is, you get an array of individual words:

```
string[] myWords;  
myWords = myString.Split(separator);
```

Next, you loop through the words in this array using `foreach` and write each one to the console:

```
foreach (string word in myWords)  
{  
    WriteLine(${word});  
}
```

**NOTE** Each word obtained has no spaces, either embedded in the word or at either end. The separators are removed when you use `Split()`.

---

## EXERCISES

5.1 Which of the following conversions can't be performed implicitly?

- a. int to short
- b. short to int
- c. bool to string
- d. byte to float

5.2 Show the code for a `color` enumeration based on the `short` type containing the colors of the rainbow plus black and white. Can this enumeration be based on the `byte` type?

5.3 Will the following code compile? Why or why not?

```
string[] blab = new string[5]  
blab[5] = 5th string.
```

5.4 Write a console application that accepts a string from the user and outputs a string with the characters in reverse order.

5.5 Write a console application that accepts a string and replaces all occurrences of the string `no` with `yes`.

5.6 Write a console application that places double quotes around each word in a string.

---

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPT
Type conversion	You can convert values from one type into another, but there are rules that apply when you do so. Implicit conversion happens automatically, but only when all possible values of the source value type are available in the target value type. Explicit conversion is also possible, but you run the risk of values not being assigned as expected, or even causing errors.
Enumerations	Enums, or enumerations, are types that have a discrete set of values, each of which has a name. Enums are defined with the <code>enum</code> keyword. This makes them easy to understand in code because they are very readable. Enums have an underlying numeric type ( <code>int</code> by default), and you can use this property of enum values to convert between enum values and numeric values, or to identify enum values.
Structs	Structs, or structures, are types that contain several different values at the same time. Structs are defined with the <code>struct</code> keyword. The values contained in a struct each have a name and a type; there is no requirement that every value stored in a struct is the same type.
Arrays	An array is a collection of values of the same type. Arrays have a fixed size, or length, which determines how many values they can contain. You can define multidimensional or jagged arrays to hold different amounts and shapes of data. You can also iterate through the values in an array with a <code>foreach</code> loop.



# 6

## Functions

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Defining and using simple functions that don't accept or return any data
- Transferring data to and from functions
- Working with variable scope
- Using command-line arguments with the `Main()` function
- Supplying functions as members of struct types
- Using function overloading
- Using delegates

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 6 download and individually named according to the names throughout the chapter.

All the code you have seen so far has taken the form of a single block, perhaps with some looping to repeat lines of code, and branching to execute statements conditionally. Performing an operation on your data has meant placing the code required right where you want it to work.

This kind of code structure is limited. Often, some tasks — such as finding the highest value in an array, for example — might need to be performed at several points in a program. You can place identical (or nearly identical) sections of code in your application whenever necessary, but this has its own problems. Changing even one minor detail concerning a common task (to correct a code error, for example) can require changes to multiple sections of code, which can be spread throughout the application. Missing one of these can have dramatic consequences and cause the whole application to fail. In addition, the application can get very lengthy.

The solution to this problem is to use *functions*. Functions in C# are a means of providing blocks of code that can be executed at any point in an application.

**NOTE** *Functions of the specific type examined in this chapter are known as methods, but this term has a very specific meaning in .NET programming that will only become clear later in this book. Therefore, for now, the term method will not be used.*

For example, you could have a function that calculates the maximum value in an array. You can use the function from any point in your code, and use the same lines of code in each case. Because you need to supply this code only once, any changes you make to it will affect this calculation wherever it is used. The function can be thought of as containing *reusable* code.

Functions also have the advantage of making your code more readable, as you can use them to group related code together. This way, your application body can be very short, as the inner workings of the code are separated out. This is similar to the way in which you can collapse regions of code together in the IDE using the outline view, and it gives your application a more logical structure.

Functions can also be used to create multipurpose code, enabling them to perform the same operations on varying data. You can supply a function with information to work with in the form of arguments, and you can obtain results from functions in the form of return values. In the preceding example, you could supply an array to search as a argument and obtain the maximum value in the array as a return value. This means that you can use the same function to work with a different array each time. A function definition consists of a name, a return type, and a list of parameters that specify the number and type of arguments that the function requires. The name and parameters of a function (but not its return type) collectively define the *signature* of a function.

## DEFINING AND USING FUNCTIONS

This section describes how you can add functions to your applications and then use (call) them from your code. Starting with the basics, you look at simple functions that don't exchange any data with code that calls them, and then look at more advanced function usage. The following Try It Out gets things moving.

### TRY IT OUT Defining and Using a Basic Function: Ch06Ex01\Program.cs

1. Create a new console application called Ch06Ex01 and save it in the directory C:\BegVCSharp\Chapter06.
2. Add the following code to Program.cs:

```
class Program
{
    static void Write()
```

```

{
    WriteLine("Text output from function.");
}
static void Main(string[] args)
{
    Write();
    ReadKey();
}
}

```

3. Execute the code. The result is shown in Figure 6-1.




FIGURE 6-1

## How It Works

The following four lines of your code define a function called `Write()`:

```

static void Write()
{
    WriteLine("Text output from function.");
}

```

The code contained here simply outputs some text to the console window, but this behavior isn't that important at the moment, because the focus here is on the mechanisms behind function definition and use.

The function definition consists of the following:

- Two keywords: `static` and `void`
- A function name followed by parentheses: `Write()`
- A block of code to execute, enclosed in curly braces

**NOTE** Function names are usually written in *PascalCase*.

The code that defines the `Write()` function looks very similar to some of the other code in your application:

```

static void Main(string[] args)
{
    ...
}

```

That's because all the code you have written so far (apart from type definitions) has been part of a function. This function, `Main()`, is the entry point function for a console application. When a C# application is executed, the entry point function it contains is called; and when that function is completed, the application terminates. All C# executable code must have an entry point.

The only difference between the `Main()` function and your `Write()` function (apart from the lines of code they contain) is that there is some code inside the parentheses after the function name `Main`. This is how you specify parameters, which you see in more detail shortly.

Both `Main()` and `Write()` are defined using the `static` and `void` keywords. The `static` keyword relates to object-oriented concepts, which you come back to later in the book. For now, you only need to remember that all the functions you use in your applications in this section of the book must use this keyword.

In contrast, `void` is much simpler to explain. It's used to indicate that the function does not return a value. Later in this chapter, you'll see the code that you need to use when a function has a return value.

Moving on, the code that calls your function is as follows:

```
Write();
```

You simply type the name of the function followed by empty parentheses. When program execution reaches this point, the code in the `Write()` function runs.

**NOTE** *The parentheses used both in the function definition and where the function is called are mandatory. Try removing them if you like — the code won't compile.*

---

## Return Values

The simplest way to exchange data with a function is to use a return value. Functions that have return values *evaluate* to that value exactly the same way that variables evaluate to the values they contain when you use them in expressions. Just like variables, return values have a type.

For example, you might have a function called `GetString()` whose return value is a string. You could use this in code, such as the following:

```
string myString;  
myString = GetString();
```

Alternatively, you might have a function called `GetVal()` that returns a `double` value, which you could use in a mathematical expression:

```
double myVal;  
double multiplier = 5.3;  
myVal = GetVal() * multiplier;
```

When a function returns a value, you have to modify your function in two ways:

- Specify the type of the return value in the function declaration instead of using the `void` keyword.
- Use the `return` keyword to end the function execution and transfer the return value to the calling code.

In code terms, this looks like the following in a console application function of the type you've been looking at:

```
static <returnType> <FunctionName>()
{
    ...
    return <returnValue>;
}
```

The only limitation here is that `<returnValue>` must be a value that either is of type `<returnType>` or can be implicitly converted to that type. However, `<returnType>` can be any type you want, including the more complicated types you've seen. This might be as simple as the following:

```
static double GetVal()
{
    return 3.2;
}
```

However, return values are usually the result of some processing carried out by the function; the preceding could be achieved just as easily using a `const` variable.

When the `return` statement is reached, program execution returns to the calling code immediately. No lines of code after this statement are executed, although this doesn't mean that `return` statements can only be placed on the last line of a function body. You can use `return` earlier in the code, perhaps after performing some branching logic. Placing `return` in a `for` loop, an `if` block, or any other structure causes the structure to terminate immediately and the function to terminate:

```
static double GetVal()
{
    double checkVal;
    // checkVal assigned a value through some logic (not shown here).
    if (checkVal < 5)
        return 4.7;
    return 3.2;
}
```

Here, one of two values is returned, depending on the value of `checkVal`. The only restriction in this case is that a `return` statement must be processed before reaching the closing `}` of the function. The following is illegal:

```
static double GetVal()
{
    double checkVal;
    // checkVal assigned a value through some logic.
    if (checkVal < 5)
        return 4.7;
}
```

If `checkVal` is  $\geq 5$ , then no `return` statement is met, which isn't allowed. All processing paths must reach a `return` statement. In most cases, the compiler detects this and gives you the error “not all code paths return a value.”

Functions that execute a single line of code can use a feature introduced in C# 6 called expression-bodied methods. The following function pattern uses a `=>` (lambda arrow) to implement this feature.

```
static <returnType> <FunctionName>() => <myVal1 * myVal2>;
```

For example, a `Multiply()` function which prior to C# 6 is written like this:

```
static double Multiply(double myVal1, double myVal2)
{
    return myVal1 * myVal2;
}
```

Can now be written using the `=>` (lambda arrow). The result of the code written here expresses the intent of the method in a much simpler and consolidated way.

```
static double Multiply(double myVal1, double myVal2) => myVal1 * myVal2;
```

## Parameters

When a function needs to accept parameters, you must specify the following:

- A list of the parameters accepted by the function in its definition, along with the types of those parameters
- A matching list of arguments in each function call

**NOTE** Note that careful reading of the C# specification shows a subtle distinction between parameters and arguments. Parameters are defined as part of a function definition, whereas arguments are passed to a function by calling code. However, these terms are often used interchangeably, and nobody seems to get too upset about that.

This involves the following code, where you can have any number of parameters, each with a type and a name:

```
static <returnType> <FunctionName>(<paramType> <paramName>, ...)
{
    ...
    return <returnValue>;
}
```

The parameters are separated using commas, and each of these parameters is accessible from code within the function as a variable. For example, a simple function might take two double parameters and return their product:

```
static double Product(double param1, double param2) => param1 * param2;
```

The following Try It Out provides a more complex example.

### TRY IT OUT Exchanging Data with a Function (Part 1): Ch06Ex02\Program.cs

1. Create a new console application called Ch06Ex02 and save it in the directory C:\BegVCSharp\Chapter06.
2. Add the following code to Program.cs:

```
class Program
{
    static int MaxValue(int[] intArray)
    {
        int maxVal = intArray[0];
        for (int i = 1; i < intArray.Length; i++)
        {
            if (intArray[i] > maxVal)
                maxVal = intArray[i];
        }
        return maxVal;
    }
    static void Main(string[] args)
    {
        int[] myArray = { 1, 8, 3, 6, 2, 5, 9, 3, 0, 2 };
        int maxVal = MaxValue(myArray);
        WriteLine($"The maximum value in myArray is {maxVal}");
        ReadKey();
    }
}
```

3. Execute the code. The result is shown in Figure 6-2.




FIGURE 6-2

### How It Works

This code contains a function that does what the example function at the beginning of this chapter hoped to do. It accepts an array of integers as a parameter and returns the highest number in the array. The function definition is as follows:

```
static int MaxValue(int[] intArray)
{
    int maxVal = intArray[0];
    for (int i = 1; i < intArray.Length; i++)
    {
        if (intArray[i] > maxVal)
            maxVal = intArray[i];
    }
    return maxVal;
}
```

The function, `MaxValue()`, has a single parameter defined, an `int` array called `intArray`. It also has a return type of `int`. The calculation of the maximum value is simple. A local integer variable called `maxVal` is initialized to the first value in the array, and then this value is compared with each of the subsequent elements in the array. If an element contains a higher value than `maxVal`, then this value replaces the current value of `maxVal`. When the loop finishes, `maxVal` contains the highest value in the array, and is returned using the `return` statement.

The code in `Main()` declares and initializes a simple integer array to use with the `MaxValue()` function:

```
int[] myArray = { 1, 8, 3, 6, 2, 5, 9, 3, 0, 2 };
```

The call to `MaxValue()` is used to assign a value to the `int` variable `maxVal`:

```
int maxVal = MaxValue(myArray);
```

Next, you write that value to the screen using `WriteLine()`:

```
WriteLine($"The maximum value in myArray is {maxVal}");
```

---

## Parameter Matching

When you call a function, you must supply arguments that match the parameters as specified in the function definition. This means matching the parameter types, the number of parameters, and the order of the parameters. For example, the function

```
static void MyFunction(string myString, double myDouble)
{
    ...
}
```

can't be called using the following:

```
MyFunction(2.6, "Hello");
```

Here, you are attempting to pass a `double` value as the first argument, and a `string` value as the second argument, which is not the order in which the parameters are defined in the function definition. The code won't compile because the parameter type is wrong. In the “Overloading Functions” section later in this chapter, you'll learn a useful technique for getting around this problem.

## Parameter Arrays

C# enables you to specify one (and only one) special parameter for a function. This parameter, which must be the last parameter in the function definition, is known as a *parameter array*.

Parameter arrays enable you to call functions using a variable amount of parameters, and they are defined using the `params` keyword.

Parameter arrays can be a useful way to simplify your code because you don't have to pass arrays from your calling code. Instead, you pass several arguments of the same type, which are placed in an array you can use from within your function.

The following code is required to define a function that uses a parameter array:

```
static <returnType> <FunctionName>(<p1Type> <p1Name>, ...,
                                         params <type>[] <name>)
{
    ...
    return <returnValue>;
}
```

You can call this function using code like the following:

```
<FunctionName>(<p1>, ..., <val1>, <val2>, ...)
```

`<val1>`, `<val2>`, and so on are values of type `<type>`, which are used to initialize the `<name>` array. The number of arguments that you can specify here is almost limitless; the only restriction is that they must all be of type `<type>`. You can even specify no arguments at all.

The following Try It Out defines and uses a function with a `params` type parameter.

### TRY IT OUT Exchanging Data with a Function (Part 2): Ch06Ex03\Program.cs

1. Create a new console application called Ch06Ex03 and save it in the directory C:\BegVCSharp\Chapter06.
2. Add the following code to Program.cs:

```
class Program
{
    static int SumVals(params int[] vals)
    {
        int sum = 0;
        foreach (int val in vals)
        {
            sum += val;
        }
        return sum;
    }
    static void Main(string[] args)
    {
        int sum = SumVals(1, 5, 2, 9, 8);
        WriteLine($"Summed Values = {sum}");
        ReadKey();
    }
}
```

3. Execute the code. The result is shown in Figure 6-3.




FIGURE 6-3

## How It Works

The function `SumVals()` is defined using the `params` keyword to accept any number of `int` arguments (and no others):

```
static int SumVals(params int[] vals)
{
    ...
}
```

The code in this function simply iterates through the values in the `vals` array and adds the values together, returning the result.

In `Main()`, you call `SumVals()` with five integer arguments:

```
int sum = SumVals(1, 5, 2, 9, 8);
```

You could just as easily call this function with none, one, two, or 100 integer arguments — there is no limit to the number you can specify.

**NOTE** C# includes alternative ways to specify function parameters, including a far more readable way to include optional parameters. You will learn about these methods in Chapter 13, which looks at the C# language.

---

## Reference and Value Parameters

All the functions defined so far in this chapter have had value parameters. That is, when you have used parameters, you have passed a value into a variable used by the function. Any changes made to this variable in the function have no effect on the argument specified in the function call. For example, consider a function that doubles and displays the value of a passed parameter:

```
static void ShowDouble(int val)
{
    val *= 2;
    WriteLine($"val doubled = {0}", val);
}
```

Here, the parameter, `val`, is doubled in this function. If you call it like this,

```
int myNumber = 5;
WriteLine($"myNumber = {myNumber}");
ShowDouble(myNumber);
WriteLine($"myNumber = {myNumber}",);
```

then the text output to the console is as follows:

```
myNumber = 5
val doubled = 10
myNumber = 5
```

Calling `ShowDouble()` with `myNumber` as an argument doesn't affect the value of `myNumber` in `Main()`, even though the parameter it is assigned to, `val`, is doubled.

That's all very well, but if you *want* the value of `myNumber` to change, you have a problem. You could use a function that returns a new value for `myNumber`, like this:

```
static int DoubleNum(int val)
{
    val *= 2;
    return val;
}
```

You could call this function using the following:

```
int myNumber = 5;
WriteLine($"myNumber = {myNumber}");
myNumber = DoubleNum(myNumber);
WriteLine($"myNumber = {myNumber}");
```

However, this code is hardly intuitive and won't cope with changing the values of multiple variables used as arguments (as functions have only one return value).

Instead, you want to pass the parameter by *reference*, which means that the function will work with exactly the same variable as the one used in the function call, not just a variable that has the same value. Any changes made to this variable will, therefore, be reflected in the value of the variable used as an argument. To do this, you simply use the `ref` keyword to specify the parameter:

```
static void ShowDouble(ref int val)
{
    val *= 2;
    WriteLine($"val doubled = {val}");
}
```

Then, specify it again in the function call (this is mandatory):

```
int myNumber = 5;
WriteLine($"myNumber = {myNumber},");
ShowDouble(ref myNumber);
WriteLine($"myNumber = {myNumber}");
```

The text output to the console is now as follows:

```
myNumber = 5
val doubled = 10
myNumber = 10
```

Note two limitations on the variable used as a `ref` parameter. First, the function might result in a change to the value of a reference parameter, so you must use a *nonconstant* variable in the function call. The following is therefore illegal:

```
const int myNumber = 5;
WriteLine($"myNumber = {myNumber},");
ShowDouble(ref myNumber);
WriteLine($"myNumber = {myNumber}");
```

Second, you must use an initialized variable. C# doesn't allow you to assume that a `ref` parameter will be initialized in the function that uses it. The following code is also illegal:

```
int myNumber;
ShowDouble(ref myNumber);
WriteLine("myNumber = {myNumber}");
```

## Out Parameters

In addition to passing values by reference, you can specify that a given parameter is an *out parameter* by using the `out` keyword, which is used in the same way as the `ref` keyword (as a modifier to the parameter in the function definition and in the function call). In effect, this gives you almost exactly the same behavior as a reference parameter, in that the value of the parameter at the end of the function execution is returned to the variable used in the function call. However, there are important differences:

- Whereas it is illegal to use an unassigned variable as a `ref` parameter, you can use an unassigned variable as an `out` parameter.
- An `out` parameter must be treated as an unassigned value by the function that uses it.

This means that while it is permissible in calling code to use an assigned variable as an `out` parameter, the value stored in this variable is lost when the function executes.

As an example, consider an extension to the `MaxValue()` function shown earlier, which returns the maximum value of an array. Modify the function slightly so that you obtain the index of the element with the maximum value within the array. To keep things simple, obtain just the index of the first occurrence of this value when there are multiple elements with the maximum value. To do this, you add an `out` parameter by modifying the function as follows:

```
static int MaxValue(int[] intArray, out int maxIndex)
{
    int maxVal = intArray[0];
    maxIndex = 0;
    for (int i = 1; i < intArray.Length; i++)
    {
        if (intArray[i] > maxVal)
        {
            maxVal = intArray[i];
            maxIndex = i;
        }
    }
    return maxVal;
}
```

You might use the function like this:

```
int[] myArray = { 1, 8, 3, 6, 2, 5, 9, 3, 0, 2 };
int maxIndex;
WriteLine($"The maximum value in myArray is
{MaxValue(myArray, out maxIndex)}");
WriteLine($"The first occurrence of this value is at element
{maxIndex + 1}");
```

That results in the following:

```
The maximum value in myArray is 9
The first occurrence of this value is at element 7
```

You must use the `out` keyword in the function call, just as with the `ref` keyword.

## VARIABLE SCOPE

Throughout the last section, you might have been wondering why exchanging data with functions is necessary. The reason is that variables in C# are accessible only from localized regions of code. A given variable is said to have a *scope* from which it is accessible.

Variable scope is an important subject and one best introduced with an example. The following Try It Out illustrates a situation in which a variable is defined in one scope, and an attempt to use it is made in a different scope.

### TRY IT OUT Variable Scope: using the Ch06Ex01\Program.cs

1. Make the following changes to Ch06Ex01 in `Program.cs` created previously:

```
class Program
{
    static void Write()
    {
        WriteLine($"myString = {myString}");
    }
    static void Main(string[] args)
    {
        string myString = "String defined in Main()";
        Write();
        ReadKey();
    }
}
```

2. Compile the code and note the error and warning that appear in the error list:

```
The name 'myString' does not exist in the current context
The variable 'myString' is assigned but its value is never used
```

### How It Works

What went wrong? Well, the variable `myString` defined in the main body of your application (the `Main()` function) isn't accessible from the `Write()` function.

The reason for this inaccessibility is that variables have a scope within which they are valid. This scope encompasses the code block that they are defined in and any directly nested code blocks. The blocks of code in functions are separate from the blocks of code from which they are called. Inside `Write()`, the name `myString` is undefined, and the `myString` variable defined in `Main()` is *out of scope* — it can be used only from within `Main()`.

In fact, you can have a completely separate variable in `Write()` called `myString`. Try modifying the code as follows:

```
class Program
{
    static void Write()
    {
        string myString = "String defined in Write()";
        WriteLine("Now in Write()");
        WriteLine($"myString = {myString}");
    }
    static void Main(string[] args)
    {
        string myString = "String defined in Main()";
        Write();
        WriteLine("\nNow in Main()");
        WriteLine($"myString = {myString}");
        ReadKey();
    }
}
```

This code does compile, resulting in the output shown in Figure 6-4.




FIGURE 6-4

The operations performed by this code are as follows:

- `Main()` defines and initializes a string variable called `myString`.
- `Main()` transfers control to `Write()`.
- `Write()` defines and initializes a string variable called `myString`, which is a different variable from the `myString` defined in `Main()`.
- `Write()` outputs a string to the console containing the value of `myString` as defined in `Write()`.
- `Write()` transfers control back to `Main()`.
- `Main()` outputs a string to the console containing the value of `myString` as defined in `Main()`.

Variables whose scopes cover a single function in this way are known as *local variables*. It is also possible to have *global variables*, whose scopes cover multiple functions. Modify the code as follows:

```
class Program
{
    static string myString;
    static void Write()
    {
        string myString = "String defined in Write()";
```

```

        WriteLine("Now in Write()");
        WriteLine($"Local myString = {myString}");
        WriteLine($"Global myString = {Program.myString}");
    }
    static void Main(string[] args)
    {
        string myString = "String defined in Main()";
        Program.myString = "Global string";
        Write();
        WriteLine("\nNow in Main()");
        WriteLine($"Local myString = {myString}");
        WriteLine($"Global myString = {Program.myString}");
        ReadKey();
    }
}

```

The result is now as shown in Figure 6-5.

```

file:///C:/BegVCSharp/Chapter06Code/Ch06Ex01/bin/Debug...
Now in Write()
Local myString = String defined in Write()
Global myString = Global String
Now in Main()
myString = String defined in Main()
Global myString = Global String

```

FIGURE 6-5

Here, you have added another variable called `myString`, this time further up the hierarchy of names in the code. The variable is defined as follows:

```
static string myString;
```

Again, the `static` keyword is required. Without going into too much detail, understand that in this type of console application, you must use either the `static` or the `const` keyword for global variables of this form. If you want to modify the value of the global variable, you need to use `static` because `const` prohibits the value of the variable from changing.

To differentiate between this variable and the local variables in `Main()` and `Write()` with the same names, you have to classify the variable name using a fully qualified name, as described in Chapter 3. Here, you refer to the global version as `Program.myString`. This is necessary only when you have global and local variables with the same name; if there were no local `myString` variable, you could simply use `myString` to refer to the global variable, rather than `Program.myString`. When you have a local variable with the same name as a global variable, the global variable is said to be *hidden*.

The value of the global variable is set in `Main()` with

```
Program.myString = "Global string";
```

and accessed in `Write()` with

```
WriteLine($"Global myString = {Program.myString}");
```

You might be wondering why you shouldn't just use this technique to exchange data with functions, rather than the parameter passing shown earlier. There are indeed situations where this is an acceptable way to exchange data, for example if you are writing a single object to be used as a plugin or a short script for use in a larger project. However, there are many scenarios where it isn't a good idea. The most common issue with using global variables has to do with the management of concurrency. For example, a global variable can be written to and read from numerous methods within a class or from different threads. Can you be certain that the value in the global variable contains valid data if numerous threads and methods can write to it? Without some extra synchronization code, the answer is probably not. Additionally, overtime it is possible the actual intent of the global variable is forgotten and used later for some other reason. Therefore, the choice of whether to use global variables depends on the intended use of the function in question.

The problem with using global variables is that they are generally unsuitable for “general-purpose” functions, which are capable of working with whatever data you supply, not just data in a specific global variable. You look at this in more depth a little later.

---

## Variable Scope in Other Structures

One of the points made in the last section has consequences above and beyond variable scope between functions: that the scopes of variables encompass the code blocks in which they are defined and any directly nested code blocks. You can find the code discussed next in the chapter download in `VariableScopeInLoops\Program.cs`. This also applies to other code blocks, such as those in branching and looping structures. Consider the following code:

```
int i;
for (i = 0; i < 10; i++)
{
    string text = "Line " + Convert.ToString(i);
    WriteLine(${text});
}
WriteLine($"Last text output in loop: {text}");
```

Here, the string variable `text` is local to the `for` loop. This code won't compile because the call to `WriteLine()` that occurs outside of this loop attempts to use the variable `text`, which is out of scope outside of the loop. Try modifying the code as follows:

```
int i;
string text;
for (i = 0; i < 10; i++)
{
    text = "Line " + Convert.ToString(i);
    WriteLine(${text});
}
WriteLine($"Last text output in loop: {text}");
```

This code will also fail because variables must be declared and initialized before use, and `text` is only initialized in the `for` loop. The value assigned to `text` is lost when the loop block is exited as it isn't initialized outside the block. However, you can make the following change:

```
int i;
string text = "";
```

```

for (i = 0; i < 10; i++)
{
    text = "Line " + Convert.ToString(i);
    WriteLine($"{{text}}");
}
WriteLine($"Last text output in loop: {{text}}");

```

This time `text` is initialized outside of the loop, and you have access to its value. The result of this simple code is shown in Figure 6-6.

```

Line 0
Line 1
Line 2
Line 3
Line 4
Line 5
Line 6
Line 7
Line 8
Line 9
Last text output in loop: Line 9

```

FIGURE 6-6

The last value assigned to `text` in the loop is accessible from outside the loop. As you can see, this topic requires a bit of effort to come to grips with. It is not immediately obvious why, in light of the earlier example, `text` doesn't retain the empty string it is assigned before the loop in the code after the loop.

The explanation for this behavior is related to memory allocation for the `text` variable, and indeed any variable. Merely declaring a simple variable type doesn't result in very much happening. It is only when values are assigned to the variables that values are allocated a place in memory to be stored. When this allocation takes place inside a loop, the value is essentially defined as a local value and goes out of scope outside of the loop.

Even though the variable itself isn't localized to the loop, the value it contains is. However, assigning a value outside of the loop ensures that the value is local to the main code, and is still in scope inside the loop. This means that the variable doesn't go out of scope before the main code block is exited, so you have access to its value outside of the loop.

Luckily for you, the C# compiler detects variable scope problems, and responding to the error messages it generates certainly helps you to understand the topic of variable scope.

## Parameters and Return Values versus Global Data

Let's take a closer look at exchanging data with functions via global data and via parameters and return values. To recap, consider the following code:

```

class Program
{
    static void ShowDouble(ref int val)
    {
        val *= 2;
    }
}

```

```

        WriteLine($"val doubled = {val}");
    }
    static void Main(string[] args)
    {
        int val = 5;
        WriteLine($"val = {val}");
        ShowDouble(ref val);
        WriteLine($"val = {val}");
    }
}

```

**NOTE** This code is slightly different from the code shown earlier in this chapter, when you used the variable name `myNumber` in `Main()`. This illustrates the fact that local variables can have identical names and yet not interfere with each other.

Now compare it with this code:

```

class Program
{
    static int val;
    static void ShowDouble()
    {
        val *= 2;
        WriteLine($"val doubled = {val}");
    }
    static void Main(string[] args)
    {
        val = 5;
        WriteLine($"val = {val}");
        ShowDouble();
        WriteLine($"val = {val}");
    }
}

```

The results of these `ShowDouble()` functions are identical.

There are no hard-and-fast rules for using one technique rather than another, and both techniques are perfectly valid, but you might want to consider the following guidelines.

To start with, as mentioned when this topic was first introduced, the `ShowDouble()` version that uses the global value only uses the global variable `val`. To use this version, you must use this global variable. This limits the versatility of the function slightly and means that you must continuously copy the global variable value into other variables if you intend to store the results. In addition, global data might be modified by code elsewhere in your application, which could cause unpredictable results (values might change without you realizing it until it's too late).

Of course, it could also be argued that this simplicity actually makes your code more difficult to understand. Explicitly specifying parameters enables you to see at a glance what is changing. If you see a call that reads `FunctionName(val1, out val2)`, you instantly know that `val1` and `val2` are

the important variables to consider and that `val2` will be assigned a new value when the function is completed. Conversely, if this function took no parameters, then you would be unable to make any assumptions about what data it manipulated.

Feel free to use either technique to exchange data. In general, use parameters rather than global data; however, there are certainly cases where global data might be more suitable, and it certainly isn't an error to use that technique.

## THE MAIN() FUNCTION

Now that you've covered most of the simple techniques used in the creation and use of functions, it's time to take a closer look at the `Main()` function.

Earlier, you saw that `Main()` is the entry point for a C# application and that execution of this function encompasses the execution of the application. That is, when execution is initiated, the `Main()` function executes, and when the `Main()` function finishes, execution ends.

The `Main()` function can return either `void` or `int`, and can optionally include a `string[] args` parameter, so you can use any of the following versions:

```
static void Main()
static void Main(string[] args)
static int Main()
static int Main(string[] args)
```

The third and fourth versions return an `int` value, which can be used to signify how the application terminates, and often is used as an indication of an error (although this is by no means mandatory). In general, returning a value of 0 reflects normal termination (that is, the application has completed and can terminate safely).

The optional `args` parameter of `Main()` provides you with a way to obtain information from outside the application, specified at runtime. This information takes the form of *command-line parameters*.

When a console application is executed, any specified command-line parameters are placed in this `args` array. You can then use these parameters in your application. The following Try It Out shows this in action. You can specify any number of command-line arguments, each of which will be output to the console.

### TRY IT OUT Command-Line Arguments: Ch06Ex04\Program.cs

1. Create a new console application called Ch06Ex04 and save it in the directory `C:\BegVCSharp\Chapter06`.
2. Add the following code to `Program.cs`:

```
class Program
{
    static void Main(string[] args)
    {
        WriteLine($"{args.Length} command line arguments were specified:");
        foreach (string arg in args)
```

```

        WriteLine(arg);
        ReadKey();
    }
}

```

3. Open the property pages for the project (right-click on the Ch06Ex04 project name in the Solution Explorer window and select Properties).
4. Select the Debug page and add any command-line arguments you want to the Command Line Arguments setting. Figure 6-7 shows an example.




FIGURE 6-7

5. Run the application. Figure 6-8 shows the output.




FIGURE 6-8

## How It Works

The code used here is very simple:

```

WriteLine(${args.Length} command line arguments were specified:");
foreach (string arg in args)
    WriteLine(arg);

```

You're just using the `args` parameter as you would any other string array. You're not doing anything fancy with the arguments; you're just writing whatever is specified to the screen. You supplied the arguments via the project properties in the IDE. This is a handy way to use the same command-line arguments whenever you run the application from the IDE, rather than type them at a command-line prompt every time. The same result can be obtained by opening a command prompt window in the same directory as the project output (`C:\BegCSharp\Chapter06\Ch06Ex04\Ch06Ex04\bin\Debug`) and typing this:

```
Ch06Ex04 256 myFile.txt "a longer argument"
```

Each argument is separated from the next by spaces. To supply an argument that includes spaces, you can enclose it in double quotation marks, which prevents it from being interpreted as multiple arguments.

## STRUCT FUNCTIONS

The last chapter covered struct types for storing multiple data elements in one place. Structs are actually capable of a lot more than this. For example, they can contain functions as well as data. That might seem a little strange at first, but it is, in fact, very useful. As a simple example, consider the following struct:

```
struct CustomerName
{
    public string firstName, lastName;
}
```

If you have variables of type `CustomerName` and you want to output a full name to the console, you are forced to build the name from its component parts. You might use the following syntax for a `CustomerName` variable called `myCustomer`, for example:

```
CustomerName myCustomer;
myCustomer.firstName = "John";
myCustomer.lastName = "Franklin";
WriteLine($"{myCustomer.firstName} {myCustomer.lastName}");
```

By adding functions to structs, you can simplify this by centralizing the processing of common tasks. For example, you can add a suitable function to the struct type as follows:

```
struct CustomerName
{
    public string firstName, lastName;
    public string Name() => firstName + " " + lastName;
}
```

This looks much like any other function you've seen in this chapter, except that you haven't used the `static` modifier. The reasons for this will become clear later in the book; for now, it is

enough to know that this keyword isn't required for struct functions. You can use this function as follows:

```
CustomerName myCustomer;
myCustomer.firstName = "John";
myCustomer.lastName = "Franklin";
WriteLine(myCustomer.Name());
```

This syntax is much simpler, and much easier to understand, than the previous syntax. The `Name()` function has direct access to the `firstName` and `lastName` struct members. Within the `customerName` struct, they can be thought of as global.

## OVERLOADING FUNCTIONS

Earlier in this chapter, you saw how you must match the signature of a function when you call it. This implies that you need to have separate functions to operate on different types of variables. Function overloading provides you with the capability to create multiple functions with the same name, but each working with different parameter types. For example, earlier you used the following code, which contains a function called `MaxValue()`:

```
class Program
{
    static int MaxValue(int[] intArray)
    {
        int maxVal = intArray[0];
        for (int i = 1; i < intArray.Length; i++)
        {
            if (intArray[i] > maxVal)
                maxVal = intArray[i];
        }
        return maxVal;
    }
    static void Main(string[] args)
    {
        int[] myArray = { 1, 8, 3, 6, 2, 5, 9, 3, 0, 2 };
        int maxVal = MaxValue(myArray);
        WriteLine("The maximum value in myArray is {maxVal}");
        ReadKey();
    }
}
```

This function can be used only with arrays of `int` values. You could provide different named functions for different parameter types, perhaps renaming the preceding function as `IntArray.MaxValue()` and adding functions such as `DoubleArray.MaxValue()` to work with other types. Alternatively, you could just add the following function to your code:

```
static double MaxValue(double[] doubleArray)
{
    double maxVal = doubleArray[0];
    for (int i = 1; i < doubleArray.Length; i++)
    {
```

```

        if (doubleArray[i] > maxVal)
            maxVal = doubleArray[i];
    }
    return maxVal;
}

```

The difference here is that you are using `double` values. The function name, `MaxValue()`, is the same, but (crucially) its *signature* is different. That's because the signature of a function, as shown earlier, includes both the name of the function and its parameters. It would be an error to define two functions with the same signature, but because these two functions have different signatures, this is fine.

**NOTE** *The return type of a function isn't part of its signature, so you can't define two functions that differ only in return type; they would have identical signatures.*

After adding the preceding code, you have two versions of `MaxValue()`, which accept `int` and `double` arrays, returning an `int` or `double` maximum, respectively.

The beauty of this type of code is that you don't have to explicitly specify which of these two functions you want to use. You simply provide an array parameter, and the correct function is executed depending on the type of parameter used.

Note another aspect of the IntelliSense feature in Visual Studio: When you have the two functions shown previously in an application and then proceed to type the name of the function, for example, `Main()`, the IDE shows you the available overloads for that function. For example, if you type

```
double result = MaxValue(
```

the IDE gives you information about both versions of `MaxValue()`, which you can scroll between using the Up and Down arrow keys, as shown in Figure 6-9.




FIGURE 6-9

All aspects of the function signature are included when overloading functions. You might, for example, have two different functions that take parameters by value and by reference, respectively:

```

static void ShowDouble(ref int val)
{
    ...
}
static void ShowDouble(int val)
{
    ...
}

```

Deciding which version to use is based purely on whether the function call contains the `ref` keyword. The following would call the reference version:

```
ShowDouble(ref val);
```

This would call the value version:

```
ShowDouble(val);
```

Alternatively, you could have functions that differ in the number of parameters they require, and so on.

## USING DELEGATES

A *delegate* is a type that enables you to store references to functions. Although this sounds quite involved, the mechanism is surprisingly simple. The most important purpose of delegates will become clear later in the book when you look at events and event handling, but it's useful to briefly consider them here. Delegates are declared much like functions, but with no function body and using the `delegate` keyword. The delegate declaration specifies a return type and parameter list.

After defining a delegate, you can declare a variable with the type of that delegate. You can then initialize the variable as a reference to any function that has the same return type and parameter list as that delegate. Once you have done this, you can call that function by using the delegate variable as if it were a function.

When you have a variable that refers to a function, you can also perform other operations that would be otherwise impossible. For example, you can pass a delegate variable to a function as a parameter, and then that function can use the delegate to call whatever function it refers to, without knowing which function will be called until runtime. The following Try It Out demonstrates using a delegate to access one of two functions.

### TRY IT OUT Using a Delegate to Call a Function: Ch06Ex05\Program.cs

1. Create a new console application called Ch06Ex05 and save it in the directory C:\BegVCSharp\Chapter06.
2. Add the following code to Program.cs:

```
class Program
{
    delegate double ProcessDelegate(double param1, double param2);
    static double Multiply(double param1, double param2) => param1 * param2;
    static double Divide(double param1, double param2) => param1 / param2;

    static void Main(string[] args)
    {
        ProcessDelegate process;
        WriteLine("Enter 2 numbers separated with a comma:");
        string input = ReadLine();
        int commaPos = input.IndexOf(',');
        double param1 = ToDouble(input.Substring(0, commaPos));
        double param2 = ToDouble(input.Substring(commaPos + 1,
```

```
        input.Length - commaPos - 1));
WriteLine("Enter M to multiply or D to divide:");
input = ReadLine();
if (input == "M")
    process = new ProcessDelegate(Multiply);
else
    process = new ProcessDelegate(Divide);
WriteLine($"Result: {process(param1, param2)}");
.ReadKey();
}
```

3. Execute the code and enter the values when prompted. Figure 6-10 shows the result.



**FIGURE 6-10**

## How It Works

This code defines a delegate (`ProcessDelegate`) whose return type and parameters match those of the two functions (`Multiply()` and `Divide()`). Notice that the `Multiply()` and `Divide()` methods use the `=>` (lambda arrow) introduced in C# 6.

```
static double Multiply(double param1, double param2) => param1 * param2;
```

The delegate definition is as follows:

```
delegate double ProcessDelegate(double param1, double param2);
```

The delegate keyword specifies that the definition is for a delegate, rather than a function (the definition appears in the same place that a function definition might). Next, the definition specifies a double return value and two double parameters. The actual names used are arbitrary; you can call the delegate type and parameter names whatever you like. This example uses a delegate called ProcessDelegate and double parameters called param1 and param2.

The code in Main() starts by declaring a variable using the new delegate type:

```
static void Main(string[] args)
{
    ProcessDelegate process;
```

Next, you have some fairly standard C# code that requests two numbers separated by a comma, and then places these numbers in two `double` variables:

```
WriteLine("Enter 2 numbers separated with a comma:");
string input = ReadLine();
int commaPos = input.IndexOf(',');
double param1 = ToDouble(input.Substring(0, commaPos));
```

```
double param2 = ToDouble(input.Substring(commaPos + 1,
    input.Length - commaPos - 1));
```

**NOTE** For demonstration purposes, no user input validation is included here. If this were “real” code, you’d spend much more time ensuring that you had valid values in the local `param1` and `param2` variables.

Next, you ask the user to multiply or divide these numbers:

```
WriteLine("Enter M to multiply or D to divide:");
input = ReadLine();
```

Based on the user’s choice, you initialize the process delegate variable:

```
if (input == "M")
    process = new ProcessDelegate(Multiply);
else
    process = new ProcessDelegate(Divide);
```

To assign a function reference to a delegate variable, you use slightly odd-looking syntax. Much like assigning array values, you can use the `new` keyword to create a new delegate. After this keyword, you specify the delegate type and supply an argument referring to the function you want to use — namely, the `Multiply()` or `Divide()` function. This argument doesn’t match the parameters of the delegate type or the target function; it is a syntax unique to delegate assignment. The argument is simply the name of the function to use, without any parentheses.

In fact, you can use slightly simpler syntax here, if you want:

```
if (input == "M")
    process = Multiply;
else
    process = Divide;
```

The compiler recognizes that the delegate type of the process variable matches the signature of the two functions, and automatically initializes a delegate for you. Which syntax you use is up to you, although some people prefer to use the longhand version, as it is easier to see at a glance what is happening.

Finally, call the chosen function using the delegate. The same syntax works, regardless of which function the delegate refers to:

```
WriteLine($"Result: {process(param1, param2)}");
.ReadKey();
}
```

Here, you treat the delegate variable as if it were a function name. Unlike a function, though, you can also perform additional operations on this variable, such as passing it to a function via a parameter, as shown in this simple example:

```
static void ExecuteFunction(ProcessDelegate process)
    => process(2.2, 3.3);
```

This means that you can control the behavior of functions by passing them function delegates, much like choosing a “snap-in” to use. For example, you might have a function that sorts a string array alphabetically. You can use several techniques to sort lists, with varying performance depending on the characteristics of the list being sorted. By using delegates, you can specify the function to use by passing a sorting algorithm function delegate to a sorting function.

There are many such uses for delegates, but, as mentioned earlier, their most prolific use is in event handling, covered in Chapter 13.

## EXERCISES

- 6.1** The following two functions have errors. What are they?

```
static bool Write()
{
    WriteLine("Text output from function.");
}
static void MyFunction(string label, params int[] args, bool showLabel)
{
    if (showLabel)
        WriteLine(label);
    foreach (int i in args)
        WriteLine("{0}", i);
}
```

- 6.2** Write an application that uses two command-line arguments to place values into a string and an integer variable, respectively. Then display those values.
- 6.3** Create a delegate and use it to impersonate the `ReadLine()` function when asking for user input.
- 6.4** Modify the following struct to include a function that returns the total price of an order:

```
struct order
{
    public string itemName;
    public int unitCount;
    public double unitCost;
}
```

- 6.5** Add another function to the `order` struct that returns a formatted string as follows (as a single line of text, where italic entries enclosed in angle brackets are replaced by appropriate values):

Order Information: <unit count> <item name> items at \$<unit cost> each,  
total cost \$<total cost>

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Defining functions	Functions are defined with a name, zero or more parameters, and a return type. The name and parameters of a function collectively define the signature of the function. It is possible to define multiple functions whose signatures are different even though their names are the same — this is called <i>function overloading</i> . Functions can also be defined within struct types.
Return values and parameters	The return type of a function can be any type, or <code>void</code> if the function does not return a value. Parameters can also be of any type, and consist of a comma-separated list of type and name pairs. A variable number of parameters of a specified type can be specified through a parameter array. Parameters can be specified as <code>ref</code> or <code>out</code> parameters in order to return values to the caller. When calling a function, any arguments specified must match the parameters in the definition both in type and in order and must include matching <code>ref</code> and <code>out</code> keywords if these are used in the parameter definition.
Variable scope	Variables are scoped according to the block of code where they are defined. Blocks of code include methods as well as other structures, such as the body of a loop. It is possible to define multiple, separate variables with the same name at different scope levels.
Command-line parameters	The <code>Main()</code> function in a console application can receive command-line parameters that are passed to the application when it is executed. When executing the application, these parameters are specified by arguments separated by spaces, and longer arguments can be passed in quotes.
Delegates	As well as calling functions directly, it is possible to call them through delegates. Delegates are variables that are defined with a return type and parameter list. A given delegate type can match any method whose return type and parameters match the delegate definition.

# 7

# Debugging and Error Handling

## **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Debugging methods available in the IDE
- Error-handling techniques available in C#

## **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 7 download and individually named according to the names throughout the chapter.

So far this book has covered all the basics of simple programming in C#. Before you move on to object-oriented programming in the next part, you need to look at debugging and error handling in C# code.

Errors in code are something that will always be with you. No matter how good a programmer is, problems will always slip through, and part of being a good programmer is realizing this and being prepared to deal with it. Of course, some problems are minor and don't affect the execution of an application, such as a spelling mistake on a button, but glaring errors are also possible, including those that cause applications to fail completely (usually known as *fatal errors*). Fatal errors include simple errors in code that prevent compilation (syntax errors), or more serious problems that occur only at runtime. Some errors are subtle. Perhaps your application fails to add a record to a database because a requested field is missing, or adds a record with the wrong data in other restricted circumstances. Errors such as these, where application logic is in some way flawed, are known as *semantic errors* or *logic errors*.

Often, you won't know about these subtle errors until a user complains that something isn't working properly. This leaves you with the task of tracing through your code to find out what's happening and fixing it so that it does what it was intended to do. In these situations, the debugging capabilities of Visual Studio are a fantastic help. The first part of this chapter looks at some of the techniques available and applies them to some common problems.

Then, you'll learn the error-handling techniques available in C#. These enable you to take precautions in cases where errors are likely, and to write code that is resilient enough to cope with errors that might otherwise be fatal. The techniques are part of the C# language, rather than a debugging feature, but the IDE provides some tools to help you here too.

## DEBUGGING IN VISUAL STUDIO

Earlier, you learned that you can execute applications in two ways: with debugging enabled or without debugging enabled. By default, when you execute an application from Visual Studio (VS), it executes with debugging enabled. This happens, for example, when you press F5 or click the green Start arrow in the toolbar. To execute an application without debugging enabled, choose Debug ➔ Start Without Debugging, or press Ctrl+F5.

Visual Studio allows you to build applications in numerous configurations, including Debug (the default) and Release. You can switch between these configurations using the Solution Configurations drop-down menu in the Standard toolbar.

When you build an application in debug configuration and execute it in debug mode, more is going on than the execution of your code. Debug builds maintain *symbolic information* about your application, so that the IDE knows exactly what is happening as each line of code is executed. Symbolic information means keeping track of, for example, the names of variables used in uncompiled code, so they can be matched to the values in the compiled machine code application, which won't contain such human-readable information. This information is contained in .pdb files, which you may have seen in your computer's Debug directories.

In the release configuration, application code is optimized, and you cannot perform these operations. However, release builds also run faster; when you have finished developing an application, you will typically supply users with release builds because they won't require the symbolic information that debug builds include.

This section describes debugging techniques you can use to identify and fix areas of code that don't work as expected, a process known as *debugging*. The techniques are grouped into two sections according to how they are used. In general, debugging is performed either by interrupting program execution or by making notes for later analysis. In Visual Studio terms, an application is either running or in break mode — that is, normal execution is halted. You'll look at the nonbreak mode (run-time or normal) techniques first.

### Debugging in Nonbreak (Normal) Mode

One of the commands you've been using throughout this book is the `WriteLine()` function, which outputs text to the console. As you are developing applications, this function comes in handy for getting extra feedback about operations:

```
WriteLine("MyFunc() Function about to be called.");
MyFunc("Do something.");
WriteLine("MyFunc() Function execution completed.");
```

This code snippet shows how you can get extra information concerning a function called `MyFunc()`. This is all very well, but it can make your console output a bit cluttered; and when you develop other types of applications, such as desktop applications, you won't have a console to output information to. As an alternative, you can output text to a separate location — the Output window in the IDE.

Chapter 2, which describes the Error List window, mentions that other windows can also be displayed in the same place. One of these, the Output window, can be very useful for debugging. To display this window, select View  $\Rightarrow$  Output. This window provides information related to compilation and execution of code, including errors encountered during compilation. You can also use this window, shown in Figure 7-1, to display custom diagnostic information by writing to it directly.

**NOTE** The Output window contains a drop-down menu from which different modes can be selected, including Build, Deployment and Debug. These modes display compilation and runtime information, respectively. When you read "writing to the Output window" in this section, it actually means "writing to the debug mode view of the Output window."




FIGURE 7-1

Alternatively, you might want to create a logging file, which has information appended to it when your application is executed. The techniques for doing this are much the same as those for writing text to the Output window, although the process requires an understanding of how to access the filesystem from C# applications. For now, leave that functionality on the back burner because there is plenty you can do without getting bogged down by file-access techniques.

## Outputting Debugging Information

Writing text to the Output window at runtime is easy. You simply replace calls to `WriteLine()` with the required call to write text where you want it. There are two commands you can use to do this:

- `Debug.WriteLine()`
- `Trace.WriteLine()`

These commands function in almost exactly the same way, with one key difference — the first command works in debug builds only; the latter works for release builds as well. In fact, the `Debug.WriteLine()` command won't even be compiled into a release build; it just disappears, which certainly has its advantages (the compiled code will be smaller, for one thing).

**NOTE** Both `Debug.WriteLine()` and `Trace.WriteLine()` methods are contained within the `System.Diagnostics` namespace. The `using static` directive can only be used with static classes, for example `System.Console` which includes the `WriteLine()` method.

These functions don't work exactly like `WriteLine()`. They work with only a single string parameter for the message to output, rather than letting you insert variable values using `{x}` syntax. This means you must use an alternative technique to embed variable values in strings — for example, the `+` concatenation operator. You can also (optionally) supply a second string parameter, which displays a category for the output text. This enables you to see at a glance which output messages are displayed in the Output window, which is useful when similar messages are output from different places in the application.

The general output of these functions is as follows:

```
<category>: <message>
```

For example, the following statement, which has "MyFunc" as the optional category parameter,

```
Debug.WriteLine("Added 1 to i", "MyFunc");
```

would result in the following:

```
MyFunc: Added 1 to i
```

The next Try It Out demonstrates outputting debugging information in this way.

### TRY IT OUT Writing Text to the Output Window: Ch07Ex01\Program.cs

1. Create a new console application called Ch07Ex01 and save it in the directory `C:\BegVCsharp\Chapter07`.
2. Modify the code as follows:

```
using System;
using System.Collections.Generic;
using System.Diagnostics;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using static System.Console;
namespace Ch07Ex01
{
    class Program
    {
```

```
static void Main(string[] args)
{
    int[] testArray = {4, 7, 4, 2, 7, 3, 7, 8, 3, 9, 1, 9};
    int[] maxValIndices;
    int maxVal = Maxima(testArray, out maxValIndices);
    WriteLine($"Maximum value {maxVal} found at element indices:");
    foreach (int index in maxValIndices)
    {
        WriteLine(index);
    }
    ReadKey();
}
static int Maxima(int[] integers, out int[] indices)
{
    Debug.WriteLine("Maximum value search started.");
    indices = new int[1];
    int maxVal = integers[0];
    indices[0] = 0;
    int count = 1;
    Debug.WriteLine(string.Format(
        $"Maximum value initialized to {maxVal}, at element index 0."));
    for (int i = 1; i < integers.Length; i++)
    {
        Debug.WriteLine(string.Format(
            $"Now looking at element at index {i}."));
        if (integers[i] > maxVal)
        {
            maxVal = integers[i];
            count = 1;
            indices = new int[1];
            indices[0] = i;
            Debug.WriteLine(string.Format(
                $"New maximum found. New value is {maxVal}, at
                    element index {i}."));
        }
        else
        {
            if (integers[i] == maxVal)
            {
                count++;
                int[] oldIndices = indices;
                indices = new int[count];
                oldIndices.CopyTo(indices, 0);
                indices[count - 1] = i;
                Debug.WriteLine(string.Format(
                    $"Duplicate maximum found at element index {i}."));
            }
        }
    }
    Trace.WriteLine(string.Format(
        $"Maximum value {maxVal} found, with {count} occurrences."));
    Debug.WriteLine("Maximum value search completed.");
    return maxVal;
}
```

3. Execute the code in debug mode. The result is shown in Figure 7-2.




FIGURE 7-2

4. Terminate the application and check the contents of the Output window (in debug mode). A truncated version of the output is shown here:

```
...
Maximum value search started.
Maximum value initialized to 4, at element index 0.
Now looking at element at index 1.
New maximum found. New value is 7, at element index 1.
Now looking at element at index 2.
Now looking at element at index 3.
Now looking at element at index 4.
Duplicate maximum found at element index 4.
Now looking at element at index 5.
Now looking at element at index 6.
Duplicate maximum found at element index 6.
Now looking at element at index 7.
New maximum found. New value is 8, at element index 7.
Now looking at element at index 8.
Now looking at element at index 9.
New maximum found. New value is 9, at element index 9.
Now looking at element at index 10.
Now looking at element at index 11.
Duplicate maximum found at element index 11.
Maximum value 9 found, with 2 occurrences.
Maximum value search completed.
The thread ##### has exited with code 0 (0x0).
```

5. Change to release mode using the drop-down menu on the Standard toolbar, as shown in Figure 7-3.
6. Run the program again, this time in release mode, and recheck the Output window when execution terminates. The output (again truncated) is as follows:

```
...
Maximum value 9 found, with 2 occurrences.
The thread ##### has exited with code 0 (0x0).
```




FIGURE 7-3

## How It Works

This application is an expanded version of one shown in Chapter 6, using a function to calculate the maximum value in an integer array. This version also returns an array of the indices where maximum values are found in an array, so that the calling code can manipulate these elements.

First, an additional `using` directive appears at the beginning of the code:

```
using System.Diagnostics;
```

This simplifies access to the functions discussed earlier because they are contained in the `System.Diagnostics` namespace. Without this `using` directive, code such as,

```
Debug.WriteLine("Bananas");
```

would need further qualification, and would have to be rewritten as:

```
System.Diagnostics.Debug.WriteLine("Bananas");
```

The code in `Main()` simply initializes a test array of integers called `testArray`; it also declares another integer array called `maxValIndices` to store the index output of `Maxima()` (the function that performs the calculation), and then calls this function. Once the function returns, the code simply outputs the results.

`Maxima()` is slightly more complicated, but it doesn't use much code that you haven't already seen. The search through the array is performed in a similar way to the `MaxVal()` function in Chapter 6, but a record is kept of the indices of maximum values.

Note the function used to keep track of the indices (other than the lines that output debugging information). Rather than return an array that would be large enough to store every index in the source array (needing the same dimensions as the source array), `Maxima()` returns an array just large enough to hold the indices found. It does this by continually recreating arrays of different sizes as the search progresses. This is necessary because arrays can't be resized once they are created.

The search is initialized by assuming that the first element in the source array (called `integers` locally) is the maximum value and that there is only one maximum value in the array. Values can therefore be set for `maxVal` (the return value of the function and the maximum value found) and `indices`, the `out` parameter array that stores the indices of the maximum values found. `maxVal` is assigned the value of the first element in `integers`, and `indices` is assigned a single value, simply `0`, which is the index of the array's first element. You also store the number of maximum values found in a variable called `count`, which enables you to keep track of the `indices` array.

The main body of the function is a loop that cycles through the values in the `integers` array, omitting the first one because it has already been processed. Each value is compared to the current value of `maxVal` and ignored if `maxVal` is greater. If the currently inspected array value is greater than `maxVal`, then `maxVal` and `indices` are changed to reflect this. If the value is equal to `maxVal`, then `count` is incremented and a new array is substituted for `indices`. This new array is one element bigger than the old `indices` array, containing the new index.

The code for this last piece of functionality is as follows:

```
if (integers[i] == maxVal)
{
    count++;
    int[] oldIndices = indices;
    indices = new int[count];
    oldIndices.CopyTo(indices, 0);
```

```
        indices[count - 1] = i;
        Debug.WriteLine(string.Format(
            $"Duplicate maximum found at element index {i}."));
    }
```

This works by backing up the old `indices` array into `oldIndices`, an integer array local to this `if` code block. Note that the values in `oldIndices` are copied into the new `indices` array using the `<array>.CopyTo()` function. This function simply takes a target array and an index to use for the first element to copy to and pastes all values into the target array.

Throughout the code, various pieces of text are output using the `Debug.WriteLine()` and `Trace.WriteLine()` functions. These functions use the `string.Format()` function to embed variable values in strings in the same way as `WriteLine()`. This is slightly more efficient than using the `+ concatenation operator`.

When you run the application in debug mode, you see a complete record of the steps taken in the loop that give you the result. In release mode, you see just the result of the calculation, because no calls to `Debug.WriteLine()` are made in release builds.

---

## Tracepoints

An alternative to writing information to the Output window is to use *tracepoints*. These are a feature of Visual Studio, rather than C#, but they serve the same function as using `Debug.WriteLine()`. Essentially, they enable you to output debugging information without modifying your code.

To demonstrate tracepoints, you can use them to replace the debugging commands in the previous example. (See the `Ch07Ex01TracePoints` file in the downloadable code for this chapter.) The process for adding a tracepoint is as follows:

1. Position the cursor at the line where you want the tracepoint to be inserted. The tracepoint will be processed *before* this line of code is executed.
2. Right-click the line of code and select `Breakpoint`  $\Rightarrow$  `Insert Tracepoint`. Right-click the red circle placed next to the line of code and select the `Settings` menu item.
3. Check the `Actions` checkbox and type the string to be output in the `Message` text box in the `Log a message` section. If you want to output variable values, enclose the variable name in curly braces.
4. Click `OK`. A red diamond appears to the left of the line of code containing a tracepoint, and the line of code itself is shown in red.

As implied by the title of the dialog box for adding tracepoints and the menu selections required for them, tracepoints are a form of breakpoint (and can cause application execution to pause, just like a breakpoint, if desired). You look at breakpoints, which typically serve a more advanced debugging purpose, a little later in the chapter.

Figure 7-4 shows the tracepoint required for line 32 of `Ch07Ex01TracePoints`, where line numbering applies to the code after the existing `Debug.WriteLine()` statements have been removed.



FIGURE 7-4

There is another window that you can use to quickly see the tracepoints in an application. To display this window, select Debug  $\Rightarrow$  Windows  $\Rightarrow$  Breakpoints from the Visual Studio menu. This is a general window for displaying breakpoints (tracepoints, as noted earlier, are a form of breakpoint). You can customize the display to show more tracepoint-specific information by adding the When Hit column from the Columns drop-down in this window. Figure 7-5 shows the display with this column configured and all the tracepoints added to Ch07Ex01TracePoints.

Executing this application in debug mode has the same result as before. You can remove or temporarily disable tracepoints by right-clicking on them in the code window or via the Breakpoints window. In the Breakpoints window, the check box to the left of the tracepoint indicates whether the tracepoint is enabled; disabled tracepoints are unchecked and displayed in the code window as diamond outlines, rather than solid diamonds.




FIGURE 7-5

## Diagnostics Output Versus Tracepoints

Now that you have seen two methods of outputting essentially the same information, consider the pros and cons of each. First, tracepoints have no equivalent to the `Trace` commands; that is, there is no way to output information in a release build using tracepoints. This is because tracepoints are not included in your application. Tracepoints are handled by Visual Studio and, as such, do not exist in the compiled version of your application. You will see tracepoints doing something only when your application is running in the Visual Studio debugger.

The chief disadvantage of tracepoints is also their major advantage, which is that they are stored in Visual Studio. This makes them quick and easy to add to your applications as you need them, but also makes them all too easy to delete. Deleting a tracepoint is as simple as clicking on the red diamond indicating its position, which can be annoying if you are outputting a complicated string of information.

One bonus of tracepoints, though, is the additional information that can be easily added, such as `$FUNCTION` which adds the current function name to the output message. Although this information is available to code written using `Debug` and `Trace` commands, it is trickier to obtain. In summary, use these two methods of outputting debug information as follows:

- **Diagnostics output** — Use when debug output is something you always want to output from an application, particularly when the string you want to output is complex, involving several variables or a lot of information. In addition, `Trace` commands are often the only option should you want output during execution of an application built in release mode.
- **Tracepoints** — Use these when debugging an application to quickly output important information that may help you resolve semantic errors.

## Debugging in Break Mode

The rest of the debugging techniques described in this chapter work in break mode. This mode can be entered in several ways, all of which result in the program pausing in some way.

### Entering Break Mode

The simplest way to enter break mode is to click the Pause button in the IDE while an application is running. This Pause button is found on the Debug toolbar, which you should add to the toolbars that appear by default in Visual Studio. To do this, right-click in the toolbar area and select Debug. Figure 7-6 shows the Debug toolbar that appears.



FIGURE 7-6

The first three buttons on the toolbar allow manual control of breaking. In Figure 7-6, these are grayed out because they don't work with a program that isn't currently executing. The following sections describe the rest of the buttons as needed.



FIGURE 7-7

When an application is running, the toolbar changes to look like Figure 7-7.

The three buttons that were grayed out now enable you to do the following:

- Pause the application and enter break mode.
- Stop the application completely (this doesn't enter break mode; it just quits).
- Restart the application.

Pausing the application is perhaps the simplest way to enter break mode, but it doesn't give you fine-grained control over exactly where to stop. You are likely to stop in a natural pause in the application, perhaps where you request user input. You might also be able to enter break mode during a lengthy operation, or a long loop, but the exact stop point is likely to be fairly random. In general, it is far better to use breakpoints.

## Breakpoints


A *breakpoint* is a marker in your source code that triggers automatic entry into break mode. Breakpoints can be configured to do the following:

- Enter break mode immediately when the breakpoint is reached.
- Enter break mode when the breakpoint is reached if a Boolean expression evaluates to true.
- Enter break mode once the breakpoint is reached a set number of times.
- Enter break mode once the breakpoint is reached and a variable value has changed since the last time the breakpoint was reached.

These features are available only in debug builds. If you compile a release build, all breakpoints are ignored.

There are several ways to add breakpoints. To add simple breakpoints that break when a line is reached, just left-click on the far left of the line of code. Alternatively, you can right-click on the line and select Breakpoint ➔ Insert Breakpoint, select Debug ➔ Toggle Breakpoint from the menu, or press F9.

A breakpoint appears as a red circle next to the line of code, which is highlighted, as shown in Figure 7-8.



The screenshot shows the Visual Studio code editor with the file 'Program.cs' open. The code defines a static void Main method that initializes a test array, calls a Maxima method, and then prints the maximum value and its index. A red circular breakpoint is visible on the far left of line 20, which contains the foreach loop. The code editor interface includes tabs for 'Program.cs', 'Ch07Ex01Tracepoints', and 'Maxima(int[] integers, out int maxValIndices)'. The status bar at the bottom shows '100 %'.

```

13
14
15
16
17
18
19
20
21
22
23
24

```

```

static void Main(string[] args)
{
    int[] testArray = { 4, 7, 4, 2, 7, 3, 7, 8, 3, 9, 1, 9 };
    int[] maxValIndices;
    int maxVal = Maxima(testArray, out maxValIndices);
    WriteLine($"Maximum value {maxVal} found at element indices:");
    foreach (int index in maxValIndices)
    {
        WriteLine(index);
    }
    ReadKey();
}

```

FIGURE 7-8

You can also see information about a file's breakpoints using the Breakpoints window (you saw how to enable this window earlier). You can use the Breakpoints window to disable breakpoints (by removing the tick to the left of a description; a disabled breakpoint shows up as an unfilled red circle), to delete breakpoints, and to edit the properties of breakpoints. You can also add labels to breakpoints, which is a handy way to group selected breakpoints. You can see labels in the Labels column and filter the items shown in this window by label.

The other columns shown in this window, Condition and Hit Count, are only two of the available ones, but they are the most useful. You can edit these by right-clicking a breakpoint (in code or in this window) and selecting Condition or Hit Count.

Selecting Condition opens a dialog box in which you can type any Boolean expression, which may involve any variables in scope at the breakpoint. For example, you could configure a breakpoint that triggers when it is reached and the value of `maxVal` is greater than 4 by entering the expression "`maxVal > 4`" and selecting the `Is true` option. You can also check whether the value of this expression has changed and only trigger the breakpoint then (you might trigger it if `maxVal` changed from 2 to 6 between breakpoint encounters, for example).

Selecting Hit Count opens a dialog box in which you can specify how many times a breakpoint needs to be hit before it is triggered. A drop-down list offers the following options:

- Break always
- Break when the hit count is equal to
- Break when the hit count is a multiple of
- Break when the hit count is greater than or equal to

The option you choose, combined with the value entered in the text box next to the options, determines the behavior of the breakpoint. The hit count is useful in long loops, when you might want to break after, say, the first 5,000 cycles. It would be a pain to break and restart 5,000 times if you couldn't do this!

## Other Ways to Enter Break Mode

There are two more ways to get into break mode. One is to enter it when an unhandled exception is thrown. This subject is covered later in this chapter, when you look at error handling. The other way is to break when an assertion is generated.

*Assertions* are instructions that can interrupt application execution with a user-defined message. They are often used during application development to test whether things are going smoothly. For example, at some point in your application you might require a given variable to have a value less than 10. You can use an assertion to confirm that this is true, interrupting the program if it isn't. When the assertion occurs, you have the option to Abort, which terminates the application; Retry, which causes break mode to be entered; or Ignore, which causes the application to continue as normal.

As with the debug output functions shown earlier, there are two versions of the assertion function:

- `Debug.Assert()`
- `Trace.Assert()`

Again, the debug version is only compiled into debug builds.

These functions take three parameters. The first is a Boolean value, whereby a value of `false` causes the assertion to trigger. The second and third are string parameters to write information both to a pop-up dialog box and the Output window. The preceding example would need a function call such as the following:

```
Debug.Assert(myVar < 10, "myVar is 10 or greater.",
    "Assertion occurred in Main().");
```

Assertions are often useful in the early stages of user adoption of an application. You can distribute release builds of your application containing `Trace.Assert()` functions to keep tabs on things. Should an assertion be triggered, the user will be informed, and this information can be passed on to you. You can then determine what has gone wrong even if you don't know how it went wrong.

You might, for example, provide a brief description of the error in the first string, with instructions as to what to do next as the second string:

```
Trace.Assert(myVar < 10, "Variable out of bounds.",
    "Please contact vendor with the error code KCW001.");
```

Should this assertion occur, the user will see the dialog box shown in Figure 7-9.

Admittedly, this isn't the most user-friendly dialog box in the world, as it contains a lot of information that could confuse users, but if they send you a screenshot of the error, you could quickly track down the problem.

Now it's time to look at what you can actually do after application execution is halted and you are in break mode. In general, you enter break mode to find an error in your code (or to reassure yourself that things are working properly). Once you are in break mode, you can use various techniques, all of which enable you to analyze your code and the exact state of the application at the point in its execution where it is paused.




FIGURE 7-9

## Monitoring Variable Content

Monitoring variable content is just one example of how Visual Studio helps you a great deal by simplifying things. The easiest way to check the value of a variable is to hover the mouse over its name in the source code while in break mode. A tooltip showing information about the variable appears, including the variable's current value.

You can also highlight entire expressions to get information about their results in the same way. For more complex values, such as arrays, you can even expand values in the tooltip to see individual element entries.

It is possible to pin these tooltip windows to the code view, which can be useful if there is a variable you are particularly interested in. Pinned tooltips persist, so they are available even if you stop and restart debugging. You can also add comments to pinned tooltips, move them around, and see the value of the last variable value, even when the application isn't running.

You may have noticed that when you run an application, the layout of the various windows in the IDE changes. By default, the following changes are likely to occur at runtime (this behavior may vary slightly depending on your installation):

- The Properties window disappears, along with some other windows, probably including the Solution Explorer window.
- The Error List window is replaced with two new windows across the bottom of the IDE window.
- Several new tabs appear in the new windows.

The new screen layout is shown in Figure 7-10. This may not match your display exactly, and some of the tabs and windows may not look exactly the same, but the functionality of these windows as described later will be the same, and this display is customizable via the View and Debug ⇔ Windows menus (during break mode), as well as by dragging windows around the screen to reposition them.

The new window that appears in the bottom-left corner is particularly useful for debugging. It enables you to keep tabs on the values of variables in your application when in break mode:

- **Autos** — Variables in use in the current and previous statements (Ctrl+D, A)
- **Locals** — All variables in scope (Ctrl+D, L)
- **Watch N** — Customizable variable and expression display (where N is 1 to 4, found on Debug ⇔ Windows ⇔ Watch)

All these tabs work in more or less the same way, with various additional features depending on their specific function. In general, each tab contains a list of variables, with information on each variable's name, value, and type. More complex variables, such as arrays, may be further examined using the + and - tree expansion/contraction symbols to the left of their names, enabling a tree view of their content. For example, Figure 7-11 shows the Locals tab obtained by placing a breakpoint in the example code. It shows the expanded view for one of the array variables, `maxValIndices`.




FIGURE 7-10

Locals		
Name	Type	Value
args	string[]	{string[0]}
testArray	int[]	{int[12]}
maxValIndices	int[]	{int[2]}
[0]	int	9
[1]	int	11
maxVal	int	9
index	int	9

FIGURE 7-11

You can also edit the content of variables from this view. This effectively bypasses any other variable assignment that might have happened in earlier code. To do this, simply type a new value into the Value column for the variable you want to edit. You might do this to try out some scenarios that would otherwise require code changes, for example.

The Watch window enables you to monitor specific variables, or expressions involving specific variables. To use this window, type the name of a variable or expression into the Name column and view the results. Note that not all variables in an application are in scope all the time, and are labeled as such in a Watch window. For example, Figure 7-12 shows a Watch window with a

few sample variables and expressions in it, obtained when a breakpoint just before the end of the Maxima() function is reached.




FIGURE 7-12

The `testArray` array is local to `Main()`, so you don't see a value here and it is grayed out.

## Stepping through Code

So far, you've learned how to discover what is going on in your applications at the point where break mode is entered. Now it's time to see how you can use the IDE to step through code while remaining in break mode, which enables you to see the exact results of the code being executed. This is an extremely valuable technique for those of us who can't think as fast as computers can.

When Visual Studio enters break mode, a yellow arrow cursor appears to the left of the code view (which may initially appear inside the red circle of a breakpoint if a breakpoint was used to enter break mode) next to the line of code that is about to be executed, as shown in Figure 7-13.




FIGURE 7-13

This shows you what point execution has reached when break mode is entered. At this point, you can execute the program on a line-by-line basis. To do so, you use some of the Debug toolbar buttons shown in Figure 7-14.



FIGURE 7-14

The sixth, seventh, and eighth icons control program flow in break mode. In order, they are as follows:

- **Step Into** — Execute and move to the next statement to execute.
- **Step Over** — Similar to Step Into, but won't enter nested blocks of code, including functions.
- **Step Out** — Run to the end of the code block and resume break mode at the statement that follows.

To look at every single operation carried out by the application, you can use Step Into to follow the instructions sequentially. This includes moving inside functions, such as `Maxima()` in the preceding example. Clicking this icon when the cursor reaches line 17, which is the call to `Maxima()`, results in the cursor moving to the first line inside the `Maxima()` function. Alternatively, clicking Step Over when you reach line 17 moves the cursor straight to line 18, without going through the code in `Maxima()` (although this code is still executed). If you do step into a function that you aren't interested in, you can click Step Out to return to the code that called the function. As you step through code, the values of variables are likely to change. If you keep an eye on the monitoring windows just discussed, you can clearly see this happening.

You can also change which line of code will be executed next by right-clicking on a line of code and selecting Set Next Statement, or by dragging the yellow arrow to a different line of code. This doesn't always work, such as when skipping variable initialization. However, it can be very useful for skipping problematic lines of code to see what will happen, or for repeating the execution of code by moving the arrow backward.

In code that has semantic errors, these techniques may be the most useful ones at your disposal. You can step through code right up to the point where you expect problems to occur, and the errors will be generated as if you were running the program normally. Or you can cause statements to be executed more than once by changing the executing code. Along the way, you can watch the data to see just what is going wrong. Later in this chapter, you'll step through some code to find out what is happening in an example application.

## Immediate and Command Windows

The Command and Immediate windows (found on the Debug Windows menu) enable you to execute commands while an application is running. The Command window enables you to perform Visual Studio operations manually (such as menu and toolbar operations), and the Immediate window enables you to execute additional code besides the source code lines being executed, and to evaluate expressions.

These windows are intrinsically linked. You can even switch between them by entering commands — `immed` to move from the Command window to the Immediate window and `cmd` to move back.

This section concentrates on the Immediate window because the Command window is only really useful for complex operations. The simplest use of this window is to evaluate expressions, a bit like

a one-shot use of the Watch windows. To do this, type an expression and press Return. The information requested will then be displayed. An example is shown in Figure 7-15.



FIGURE 7-15

You can also change variable content here, as demonstrated in Figure 7-16.



FIGURE 7-16

In most cases, you can get the effects you want more easily using the variable monitoring windows shown earlier, but this technique is still handy for tweaking values, and it's good for testing expressions.

## The Call Stack Window

The final window to look at is the Call Stack window, which shows you the way in which the program reached the current location. In simple terms, this means showing the current function along with the function that called it, the function that called that, and so on (that is, a list of nested function calls). The exact points where calls are made are also recorded.

In the earlier example, entering break mode when in `Maxima()`, or moving into this function using code stepping, reveals the information shown in Figure 7-17.




FIGURE 7-17

If you double-click an entry, you are taken to the appropriate location, enabling you to track the way code execution has reached the current point. This window is particularly useful when errors are first detected, because you can see what happened immediately before the error. Where errors occur in commonly used functions, this helps you determine the source of the error.

## ERROR HANDLING

The first part of this chapter explained how to find and correct errors during application development so that they don't occur in release-level code. Sometimes, however, you know that errors are likely to occur and there is no way to be 100 percent sure that they won't. In those situations, it may be preferable to anticipate problems and write code that is robust enough to deal with these errors gracefully, without interrupting execution.

*Error handling* is the term for all techniques of this nature, and this section looks at exceptions and how you can deal with them. An exception is an error generated either in your code or in a function called by your code that occurs at runtime. The definition of error here is more vague than it has been up until now, because exceptions may be generated manually, in functions and so on. For example, you might generate an exception in a function if one of its string parameters doesn't start with the letter "a." Strictly speaking, this isn't an error outside of the context of the function, although the code that calls the function treats it as an error.

You've seen exceptions a few times already in this book. Perhaps the simplest example is attempting to address an array element that is out of range:

```
int[] myArray = { 1, 2, 3, 4 };
int myElem = myArray[4];
```

This outputs the following exception message and then terminates the application:

```
Index was outside the bounds of the array.
```

Exceptions are defined in namespaces, and most have names that make their purpose clear. In this example, the exception generated is called `System.IndexOutOfRangeException`, which makes sense because you have supplied an index that is not in the range of indices permissible in `myArray`. This message appears, and the application terminates, only when the exception is unhandled. In the next section, you'll see exactly what you have to do to handle an exception.

## try...catch...finally

The C# language includes syntax for *structured exception handling* (SEH). Three keywords mark code as being able to handle exceptions, along with instructions specifying what to do when an exception occurs: `try`, `catch`, and `finally`. Each of these has an associated code block and must be used in consecutive lines of code. The basic structure is as follows:

```
try
{
    ...
}
catch (<exceptionType> e) when (filterIsTrue)
{
    <await methodName(e);>
    ...
}
```

```
finally
{
    <await methodName;>
    ...
}
```

Optionally using `await` within either a `catch` or `finally` block was introduced in C# 6. The `await` keyword is used to support advanced asynchronous programming techniques that avoid bottlenecks and can improve the overall performance and responsiveness of an application. Asynchronous programming, utilizing the `async` and `await` keywords, is not discussed in this book; nevertheless, as those keywords do simplify the implementation of this programming technique, it is highly recommended to learn about them.

It is also possible, however, to have a `try` block and a `finally` block with no `catch` block, or a `try` block with multiple `catch` blocks. If one or more `catch` blocks exist, then the `finally` block is optional; otherwise, it is mandatory. The usage of the blocks is as follows:

- `try` — Contains code that might throw exceptions (“throw” is the C# way of saying “generate” or “cause” when talking about exceptions).
- `catch` — Contains code to execute when exceptions are thrown. `catch` blocks can respond only to specific exception types (such as `System.IndexOutOfRangeException`) using `<exceptionType>`, hence the ability to provide multiple `catch` blocks. It is also possible to omit this parameter entirely, to get a general `catch` block that responds to all exceptions. C# 6 introduced a concept called *exception filtering* that is implemented by adding the `when` keyword after the exception type expressions. If that exception type occurs and the filter expression is true, only then will the code in the `catch` block execute.
- `finally` — Contains code that is always executed, either after the `try` block if no exception occurs, after a `catch` block if an exception is handled, or just before an unhandled exception moves “up the call stack.” This phrase means that SEH allows you to nest `try...catch...finally` blocks inside one another, either directly or because of a call to a function within a `try` block. For example, if an exception isn’t handled by any `catch` blocks in the called function, it might be handled by a `catch` block in the calling code. Eventually, if no `catch` blocks are matched, then the application will terminate. The fact that the `finally` block is processed before this happens is the reason for its existence; otherwise, you might just as well place code outside of the `try...catch...finally` structure. This nested functionality is discussed further in the “Notes on Exception Handling” section a little later, so don’t worry if it sounds a little confusing.

Here’s the sequence of events that occurs after an exception occurs in code in a `try` block, also illustrated by Figure 7-18.

- The `try` block terminates at the point where the exception occurred.
- If a `catch` block exists, then a check is made to determine whether the block matches the type of exception that was thrown. If no `catch` block exists, then the `finally` block (which must be present if there are no `catch` blocks) executes.

- If a catch block exists but there is no match, then a check is made for other catch blocks.
- If a catch block matches the exception type and there is an exception filter that results in true, the code within it executes and then the finally block is executed (if it is present).
- If a catch block matches the exception type and there is no exception filter, the code it contains executes, and then the finally block executes if it is present.
- If no catch blocks match the exception type, then the finally block of code executes if it is present.




FIGURE 7-18

**NOTE** If two catch blocks exist that handle the same exception type, only the code within the catch block with an exception filter resulting in true is executed. If a catch block also exists handling the same exception type with no filter exception or a filter exception resulting in false, it is disregarded. Only one catch block code is executed and ordering of the catch block does not affect the execution flow.

The following Try It Out demonstrates handling exceptions. It shows throwing and handling them in several ways so you can see how things work.

### TRY IT OUT Exception Handling: Ch07Ex02\Program.cs

1. Create a new console application called Ch07Ex02 and save it in the directory C:\BegVCSharp\Chapter07.
2. Modify the code as follows (the line number comments shown here will help you match up your code to the discussion afterward, and they are duplicated in the downloadable code for this chapter for your convenience):

```
class Program
{
    static string[] eTypes = { "none", "simple", "index",
                             "nested index", "filter" };
    static void Main(string[] args)
    {
        foreach (string eType in eTypes)
        {
            try
            {
                WriteLine("Main() try block reached.");           // Line 21
                WriteLine($"ThrowException(\"{eType}\") called.");
                ThrowException(eType);
                WriteLine("Main() try block continues.");        // Line 23
            }
            catch (System.IndexOutOfRangeException e) when (eType == "filter")
            {
                WriteLine("Main() FILTERED System.IndexOutOfRangeException" +
                          $"catch block reached. Message:{\n\"{e.Message}\n}");
            }
            catch (System.IndexOutOfRangeException e)           // Line 32
            {
                WriteLine("Main() System.IndexOutOfRangeException catch " +
                          $"block reached. Message:{\n\"{e.Message}\n}");
            }
            catch                                         // Line 36
            {
                WriteLine("Main() general catch block reached.");
            }
            finally
            {
                WriteLine("Main() finally block reached.");
            }
            WriteLine();
        }
        ReadKey();
    }

    static void ThrowException(string exceptionType)
```

```

    WriteLine($"ThrowException(\"{exceptionType}\") reached.");
    switch (exceptionType)
    {
        case "none":
            WriteLine("Not throwing an exception.");
            break;
        case "simple":
            WriteLine("Throwing System.Exception.");
            throw new System.Exception(); // Line 57
        case "index":
            WriteLine("Throwing System.IndexOutOfRangeException.");
            eTypes[5] = "error";
            break;
        case "nested index":
            try // Line 66
            {
                WriteLine("ThrowException(\"nested index\") " +
                    "try block reached.");
                WriteLine("ThrowException(\"index\") called.");
                ThrowException("index"); // Line 71
            }
            catch // Line 73
            {
                WriteLine("ThrowException(\"nested index\") general" +
                    " catch block reached.");
            }
            finally
            {
                WriteLine("ThrowException(\"nested index\") finally" +
                    " block reached.");
            }
            break;
        case "filter":
            try // Line 86
            {
                WriteLine("ThrowException(\"filter\") " +
                    "try block reached.");
                WriteLine("ThrowException(\"index\") called.");
                ThrowException("index"); // Line 91
            }
            catch // Line 93
            {
                WriteLine("ThrowException(\"filter\") general" +
                    " catch block reached.");
            }
            break;
    }
}
}

```

3. Run the application. The result is shown in Figure 7-19.

```
file:///C:/BegVCSharp/Chapter07/Ch07Ex02/bin/Debug/Ch... - x
ThrowException("nested index") called.
ThrowException("nested index") reached.
ThrowException("nested index") try block reached.
ThrowException("index") called.
ThrowException("index") reached.
ThrowException("index") general catch block reached.
ThrowException("nested index") finally block reached.
Main() System.IndexOutOfRangeException catch block reached. Message:
"Index was outside the bounds of the array."
Main() finally block reached.

Main() try block reached.
ThrowException("filter") called.
ThrowException("filter") reached.
ThrowException("filter") try block reached.
ThrowException("index") called.
ThrowException("index") reached.
Throwing System.IndexOutOfRangeException.
ThrowException("filter") general catch block reached.
Main() FILTERED System.IndexOutOfRangeException catch block reached. Message:
"Index was outside the bounds of the array."
Main() finally block reached.
```

FIGURE 7-19

## How It Works

This application has a `try` block in `Main()` that calls a function called `ThrowException()`. This function may throw exceptions, depending on the parameter it is called with:

- `ThrowException ("none")`— Doesn't throw an exception
- `ThrowException ("simple")`— Generates a general exception
- `ThrowException ("index")`— Generates a `System.IndexOutOfRangeException` exception
- `ThrowException ("nested index")`— Contains its own `try` block, which contains code that calls `ThrowException ("index")` to generate a `System.IndexOutOfRangeException` exception
- `ThrowException ("filter")`— Contains its own `try` block, which contains code that calls `ThrowException ("index")` to generate a `System.IndexOutOfRangeException` exception where the exception filter results in `true`.

Each of these string parameters is held in the global `eTypes` array, which is iterated through in the `Main()` function to call `ThrowException()` once with each possible parameter. During this iteration, various messages are written to the console to indicate what is happening. This code gives you an excellent opportunity to use the code-stepping techniques shown earlier in the chapter. By working your way through the code one line at a time, you can see exactly how code execution progresses.

Add a new breakpoint (with the default properties) to line 21 of the code, which reads as follows:

```
WriteLine("Main() try block reached.");
```

**NOTE** Code is referred to by line numbers as they appear in the downloadable version of this code. If you have line numbers turned off, remember that you can turn them back on (select Tools ➔ Options and then change the Line numbers setting in the Text Editor ➔ C# ➔ General options section). Comments are included in the preceding code so that you can follow the text without having the file open in front of you.

Run the application in debug mode. Almost immediately, the program will enter break mode, with the cursor on line 20. If you select the Locals tab in the variable monitoring window, you should see that eType is currently "none". Use the Step Into button to process lines 21 and 22, and confirm that the first line of text has been written to the console. Next, use the Step Into button to step into the ThrowException() function on line 23.

Once in the ThrowException() function, the Locals window changes. eType and args are no longer in scope (they are local to Main()); instead, you see the local exceptionType argument, which is, of course, "none". Keep pressing Step Into and you'll reach the switch statement that checks the value of exceptionType and executes the code that writes out the string Not throwing an exception to the screen. When you execute the break statement (on line 57), you exit the function and resume processing in Main() at line 24. Because no exception was thrown, the try block continues.

Next, processing continues with the finally block. Click Step Into a few more times to complete the finally block and the first cycle of the foreach loop. The next time you reach line 23, ThrowException() is called using a different parameter, "simple".

Continue using Step Into through ThrowException(), and you'll eventually reach line 60:

```
throw new System.Exception();
```

You use the C# throw keyword to generate an exception. This keyword simply needs to be provided with a new-initialized exception as a parameter, and it will throw that exception. Here, you are using another exception from the System namespace, System.Exception.

**NOTE** When you use throw in a case block, no break; statement is necessary.  
throw is enough to end execution of the block.

When you process this statement with Step Into, you find yourself at the general catch block starting on line 36. There was no match with the earlier catch block starting on line 26, so this one is processed instead. Stepping through this code takes you through this block, through the finally block, and back into another loop cycle that calls ThrowException() with a new parameter on line 23. This time the parameter is "index".

Now ThrowException() generates an exception on line 63:

```
eTypes[5] = "error";
```

The eTypes array is global, so you have access to it here. However, here you are attempting to access the sixth element in the array (remember that counting starts at 0), which generates a System.IndexOutOfRangeException exception.

This time there are multiple matched catch blocks in Main(). One has a filter expression of (eType == "filter") on line 26, and the other, on line 32, has no filter expression. The value stored in eType is currently "index" and therefore the filter expression results in false which skips this catch code block.

Stepping into the code takes you to the next catch block, starting at line 32. The WriteLine() call in this block writes out the message stored in the exception using e.Message (you have access to the

exception through the parameter of the `catch` block). Again, stepping through takes you through the `finally` block (but not the second `catch` block, as the exception is already handled) and back into the loop cycle, again calling `ThrowException()` on line 23.

When you reach the `switch` structure in `ThrowException()`, this time you enter a new `try` block, starting on line 66. When you reach line 71, you perform a nested call to `ThrowException()`, this time with the parameter "index". You can use the Step Over button to skip the lines of code that are executed here because you've been through them already. As before, this call generates a `System.IndexOutOfRangeException` exception, but this time it's handled in the nested `try...catch...finally` structure, the one in `ThrowException()`. This structure has no explicit match for this type of exception, so the general `catch` block (starting on line 73) deals with it.

Continue stepping through the code and when you reach the `switch` structure in the `ThrowException()` this time, you enter a new `try` block starting at line 86. When you reach line 91, you perform a nested call to `ThrowException()` same as before. However, this time when the `catch` block that handles the `System.IndexOutOfRangeException` in the `Main()` checks the filter expression of (`eType == "filter"`), the result is true and that `catch` block is executed instead of the `catch` block handling the `System.IndexOutOfRangeException` without the exception filter.

As with the earlier exception handling, you now step through this `catch` block and the associated `finally` block, and reach the end of the function call, but with one crucial difference. Although an exception was thrown, it was also handled — by the code in `ThrowException()`. This means there is no exception left to handle in `Main()`, so you go straight to the `finally` block, at which point the application terminates.

---

## Listing and Configuring Exceptions

The .NET Framework contains a host of exception types, and you are free to throw and handle any of these in your own code. The IDE supplies a dialog box for examining and editing the available exceptions, which can be called up with the `Debug`  $\Rightarrow$  `Exception Settings` menu item (or by pressing `Ctrl+D, E`). Figure 7-20 shows the Exception Settings dialog box.




FIGURE 7-20

Exceptions are listed by category and .NET library namespace. You can see the exceptions in the `System` namespace by expanding the `Common Language Runtime Exceptions` plus sign. The list includes the `System.IndexOutOfRangeException` exception you used earlier.

Each exception may be configured using the check boxes next to the exception type. When checked, the debugger will (break when) Throw, causing a break into the debugger even for exceptions that are handled.

## EXERCISES

- 7.1 “Using `Trace.WriteLine()` is preferable to using `Debug.WriteLine()`, as the `Debug` version works only in debug builds.” Do you agree with this statement? If so, why?
- 7.2 Provide code for a simple application containing a loop that generates an error after 5,000 cycles. Use a breakpoint to enter break mode just before the error is caused on the 5,000th cycle. (Note: A simple way to generate an error is to attempt to access a nonexistent array element, such as `myArray[1000]` in an array with 100 elements.)
- 7.3 “`finally` code blocks execute only if a `catch` block isn’t executed.” True or false?
- 7.4 Given the enumeration data type `Orientation` defined in the following code, write an application that uses structured exception handling (SEH) to cast a `byte`-type variable into an `Orientation`-type variable in a safe way. (Note: You can force exceptions to be thrown using the `checked` keyword, an example of which is shown here. This code should be used in your application.)

```
enum Orientation : byte
{
    North = 1,
    South = 2,
    East = 3,
    West = 4
}
myDirection = checked((Orientation)myByte);
```

Answers to these exercises are in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Error types	Fatal errors cause your application to fail completely, either at compile time (syntax errors) or at runtime. Semantic, or logic, errors are more insidious, and may cause your application to function incorrectly or unpredictably.
Outputting debugging information	You can write code that outputs helpful information to the Output window to aid debugging in the IDE. You do this with the <code>Debug</code> and <code>Trace</code> family of functions, where <code>Debug</code> functions are ignored in release builds. For production applications, you may want to write debugging output to a log file instead. You can also use tracepoints to output debugging information.
Break mode	You can enter break mode (essentially a state where the application is paused) manually, through breakpoints, through assertions, or when unhandled exceptions occur. You can add breakpoints anywhere in your code and you can configure breakpoints to break execution only under specific conditions. When in break mode, you can inspect the content of variables (with the help of various debug information windows) and step through code a line at a time to assist you in determining where the errors are.
Exceptions	Exceptions are errors that occur at runtime and that you can trap and process programmatically to prevent your application from terminating. There are many types of exceptions that can occur when you call functions or manipulate variables. You can also generate exceptions with the <code>throw</code> keyword.
Exception handling	Exceptions that are not handled in your code will cause the application to terminate. You handle exceptions with <code>try</code> , <code>catch</code> , and <code>finally</code> code blocks. <code>try</code> blocks mark out a section of code for which exception handling is enabled. <code>catch</code> blocks consist of code that is executed only if an exception occurs, and can match specific types of exceptions. You can include multiple <code>catch</code> blocks. <code>finally</code> blocks specify code that is executed after exception handling has occurred, or after the <code>try</code> block finishes if no exception occurs. You can include only a single <code>finally</code> block, and if you include any <code>catch</code> blocks, then the <code>finally</code> block is optional.

# 8

# Introduction to Object-Oriented Programming

## **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Understanding object-oriented programming
- Using OOP techniques
- Learning how desktop applications rely on OOP

## **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the [wrox.com](http://www.wrox.com/go/beginningvisualc#2015programming) code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 8 download and individually named according to the names throughout the chapter.

At this point in the book, you've covered all the basics of C# syntax and programming, and have learned how to debug your applications. Already, you can assemble usable console applications. However, to access the real power of the C# language and the .NET Framework, you need to make use of *object-oriented programming* (OOP) techniques. In fact, as you will soon see, you've been using these techniques already, although to keep things simple we haven't focused on this.

This chapter steers away from code temporarily and focuses instead on the principles behind OOP. This leads you back into the C# language because it has a symbiotic relationship with OOP. All of the concepts introduced in this chapter are revisited in later chapters, with illustrative code — so don't panic if you don't grasp everything in the first read-through of this material.

To start with, you'll look at the basics of OOP, which include answering that most fundamental of questions, "What is an *object*?" You will quickly find that a lot of terminology related to

OOP can be confusing at first, but plenty of explanations are provided. You will also see that using OOP requires you to look at programming in a different way.

As well as discussing the general principles of OOP, this chapter looks at an area requiring a thorough understanding of OOP: desktop applications. This type of application relies on the Windows environment, with features such as menus, buttons, and so on. As such, it provides plenty of scope for description, and you will be able to observe OOP points effectively in the Windows environment.

## WHAT IS OBJECT-ORIENTED PROGRAMMING?

Object-oriented programming seeks to address many of the problems with traditional programming techniques. The type of programming you have seen so far is known as *procedural programming*, which often results in so-called monolithic applications, meaning all functionality is contained in a few modules of code (often just one). With OOP techniques, you often use many more modules of code, with each offering specific functionality. Also, each module can be isolated or even completely independent of the others. This modular method of programming gives you much more versatility and provides more opportunity for code reuse.

To illustrate this further, imagine that a high-performance application on your computer is a top-of-the-range race car. Written with traditional programming techniques, this sports car is basically a single unit. If you want to improve this car, then you have to replace the whole unit by sending it back to the manufacturer and getting their expert mechanics to upgrade it, or by buying a new one. If OOP techniques are used, however, you can simply buy a new engine from the manufacturer and follow their instructions to replace it yourself, rather than taking a hacksaw to the bodywork.

In a more traditional application, the flow of execution is often simple and linear. Applications are loaded into memory, begin executing at point A, end at point B, and are then unloaded from memory. Along the way various other entities might be used, such as files on storage media, or the capabilities of a video card, but the main body of the processing occurs in one place. The code along the way is generally concerned with manipulating data through various mathematical and logical means. The methods of manipulation are usually quite simple, using basic types such as integers and Boolean values to build more complex representations of data.

With OOP, things are rarely so linear. Although the same results are achieved, the way of getting there is often very different. OOP techniques are firmly rooted in the structure and meaning of data, and the interaction between that data and other data. This usually means putting more effort into the design stages of a project, but it has the benefit of extensibility. After an agreement is made as to the representation of a specific type of data, that agreement can be worked into later versions of an application, and even entirely new applications. The fact that such an agreement exists can reduce development time dramatically. This explains how the race car example works. The agreement here is how the code for the “engine” is structured, such that new code (for a new engine) can be substituted with ease, rather than requiring a trip back to the manufacturer. It also means that the engine, once created, can be used for other purposes. You could put it in a different car, or use it to power a submarine, for example.

OOP often simplifies things by providing an agreement about the approach to data representation, as well as about the structure and usage of more abstract entities. For example, an agreement can be

made not just on the format of data that should be used to send output to a device such as a printer, but also on the methods of data exchange with that device, including what instructions it understands, and so on. In the race car analogy, the agreement would include how the engine connects to the fuel tank, how it passes drive power to the wheels, and so on.

As the name of the technology suggests, this is achieved using objects.

## What Is an Object?

An *object* is a building block of an OOP application. This building block encapsulates part of the application, which can be a process, a chunk of data, or a more abstract entity.

In the simplest sense, an object can be very similar to a struct type such as those shown earlier in the book, containing members of variable and function types. The variables contained make up the data stored in the object, and the functions contained allow access to the object's functionality. Slightly more complex objects might not maintain any data; instead, they can represent a process by containing only functions. For example, an object representing a printer might be used, which would have functions enabling control over a printer (so you can print a document, a test page, and so on).

Objects in C# are created from types, just like the variables you've seen already. The type of an object is known by a special name in OOP, its *class*. You can use class definitions to *instantiate* objects, which means creating a real, named *instance* of a class. The phrases *instance of a class* and *object* mean the same thing here; but *class* and *object* mean fundamentally different things.

**NOTE** The terms *class* and *object* are often confused, and it is important to understand the distinction. It might help to visualize these terms using the earlier race car analogy. Think of a class as the template for the car, or perhaps the plans used to build the car. The car itself is an instance of those plans, so it could be referred to as an object.

In this chapter, you work with classes and objects using *Unified Modeling Language* (UML) syntax. UML is designed for modeling applications, from the objects that build them to the operations they perform to the use cases that are expected. Here, you use only the basics of this language, which are explained as you go along. UML is a specialized subject to which entire books are devoted, so its more complex aspects are not covered here.




FIGURE 8-1

Figure 8-1 shows a UML representation of your printer class, called `Printer`. The class name is shown in the top section of this box (you learn about the bottom two sections a little later).

Figure 8-2 shows a UML representation of an instance of this `Printer` class called `myPrinter`.

Here, the instance name is shown first in the top section, followed by the name of its class. The two names are separated by a colon.



FIGURE 8-2

## Properties and Fields

Properties and fields provide access to the data contained in an object. This object data differentiates separate objects because it is possible for different objects of the same class to have different values stored in properties and fields.

The various pieces of data contained in an object together make up the *state* of that object. Imagine an object class that represents a cup of coffee, called `CupOfCoffee`. When you instantiate this class (that is, create an object of this class), you must provide it with a state for it to be meaningful. In this case, you might use properties and fields to enable the code that uses this object to set the type of coffee used, whether the coffee contains milk and/or sugar, whether the coffee is instant, and so on. A given coffee cup object would then have a given state, such as “Colombian filter coffee with milk and two sugars.”

Both fields and properties are typed, so you can store information in them as `string` values, as `int` values, and so on. However, properties differ from fields in that they don’t provide direct access to data. Objects can shield users from the nitty-gritty details of their data, which needn’t be represented on a one-to-one basis in the properties that exist. If you used a field for the number of sugars in a `CupOfCoffee` instance, then users could place whatever values they liked in the field, limited only by the limits of the type used to store this information. If, for example, you used an `int` to store this data, then users could use any value between `-2147483648` and `2147483647`, as shown in Chapter 3. Obviously, not all values make sense, particularly the negative ones, and some of the large positive amounts might require an inordinately large cup. If you use a property for this information, you could limit this value to, say, a number between `0` and `2`.


In general, it is better to provide properties rather than fields for state access because you have more control over various behaviors. This choice doesn’t affect code that uses object instances because the syntax for using properties and fields is the same.

Read/write access to properties can also be clearly defined by an object. Certain properties can be read-only, allowing you to see what they are but not change them (at least not directly). This is often a useful technique for reading several pieces of state simultaneously. You might have a read-only property of the `CupOfCoffee` class called `Description`, returning a string representing the state of an instance of this class (such as the string given earlier) when requested. You might be able to assemble the same data by interrogating several properties, but a property such as this one might save you time and effort. You might also have write-only properties that operate in a similar way.

As well as this read/write access for properties, you can also specify a different sort of access permission for both fields and properties, known as *accessibility*. Accessibility determines which code can access these members — that is, whether they are available to all code (`public`), only to code within the class (`private`), or should use a more complex scheme (covered in more detail later in the chapter, when it becomes pertinent). One common practice is to make fields private and provide access to them via public properties. This means that code within the class has direct access to data stored in the field, while the public property shields external users from this data and prevents them from placing invalid content there. Public members are said to be *exposed* by the class.

One way to visualize this is to equate it with variable scope. Private fields and properties, for example, can be thought of as local to the object that possesses them, whereas the scope of public fields and properties also encompasses code external to the object.

In the UML representation of a class, you use the second section to display properties and fields, as shown in Figure 8-3.



**FIGURE 8-3**

This is a representation of the `CupOfCoffee` class, with five members (properties or fields, because no distinction is made in UML) defined as discussed earlier. Each of the entries contains the following information:

- Accessibility — A + symbol is used for a public member, a – symbol is used for a private member. In general, though, private members are not shown in the diagrams in this chapter because this information is internal to the class. No information is provided as to read/write access.
- The member name.
- The type of the member.


A colon is used to separate the member names and types.

## Methods

*Method* is the term used to refer to functions exposed by objects. These can be called in the same way as any other function and can use return values and parameters in the same way — you looked at functions in detail in Chapter 6.

Methods are used to provide access to the object's functionality. Like fields and properties, they can be public or private, restricting access to external code as necessary. They often make use of an object's state to affect their operations, and have access to private members, such as private fields, if required. For example, the `CupOfCoffee` class might define a method called `AddSugar()`, which would provide a more readable syntax for incrementing the amount of sugar than setting the corresponding `Sugar` property.

In UML, class boxes show methods in the third section, as shown in Figure 8-4.



**FIGURE 8-4**

The syntax here is similar to that for fields and properties, except that the type shown at the end is the return type, and method parameters are shown. Each parameter is displayed in UML with one of the following identifiers: `return`, `in`, `out`, or `inout`. These are used to signify the direction of data flow, where `out` and `inout` roughly correspond to the use of the C# keywords `out` and `ref` described in Chapter 6. `in` roughly corresponds to the default C# behavior, where neither the `out` nor `ref` keyword is used and `return` signifies that a value is passed back to the calling method.

## Everything's an Object

At this point, it's time to come clean: You have been using objects, properties, and methods throughout this book. In fact, everything in C# and the .NET Framework is an object! The `Main()` function in a console application is a method of a class. Every variable type you've looked at is a class. Every command you have used has been a property or a method, such as `<String>.Length`, `<String>.ToUpper()`, and so on. (The period character here separates the object instance's name from the property or method name, and methods are shown with `()` at the end to differentiate them from properties.)

Objects really are everywhere, and the syntax to use them is often very simple. It has certainly been simple enough for you to concentrate on some of the more fundamental aspects of C# up until now. From this point on, you'll begin to look at objects in detail. Bear in mind that the concepts introduced here have far-reaching consequences — applying even to that simple little `int` variable you've been happily playing around with.

## The Life Cycle of an Object

Every object has a clearly defined life cycle. Apart from the normal state of “being in use,” this life cycle includes two important stages:

- **Construction** — When an object is first instantiated it needs to be initialized. This initialization is known as *construction* and is carried out by a constructor function, often referred to simply as a *constructor* for convenience.
- **Destruction** — When an object is destroyed, there are often some clean-up tasks to perform, such as freeing memory. This is the job of a destructor function, also known as a *destructor*.

## Constructors

Basic initialization of an object is automatic. For example, you don't have to worry about finding the memory to fit a new object into. However, at times you will want to perform additional tasks during an object's initialization stage, such as initializing the data stored by an object. A constructor is what you use to do this.

All class definitions contain at least one constructor. These constructors can include a *default constructor*, which is a parameter-less method with the same name as the class itself. A class definition might also include several constructor methods with parameters, known as *nondefault constructors*.

These enable code that instantiates an object to do so in many ways, perhaps providing initial values for data stored in the object.

In C#, constructors are called using the `new` keyword. For example, you could instantiate a `CupOfCoffee` object using its default constructor in the following way:

```
CupOfCoffee myCup = new CupOfCoffee();
```

Objects can also be instantiated using nondefault constructors. For example, the `CupOfCoffee` class might have a nondefault constructor that uses a parameter to set the bean type at instantiation:

```
CupOfCoffee myCup = new CupOfCoffee("Blue Mountain");
```

Constructors, like fields, properties, and methods, can be public or private. Code external to a class can't instantiate an object using a private constructor; it must use a public constructor. In this way, you can, for example, force users of your classes to use a nondefault constructor (by making the default constructor private).

Some classes have no public constructors, meaning it is impossible for external code to instantiate them (they are said to be *noncreatable*). However, that doesn't make them completely useless, as you will see shortly.

## Destructors

Destructors are used by the .NET Framework to clean up after objects. In general, you don't have to provide code for a destructor method; instead, the default operation does the work for you. However, you can provide specific instructions if anything important needs to be done before the object instance is deleted.

For example, when a variable goes out of scope, it may not be accessible from your code; however, it might still exist somewhere in your computer's memory. Only when the .NET runtime performs its garbage collection clean-up is the instance completely destroyed.

## Static and Instance Class Members

As well as having members such as properties, methods, and fields that are specific to object instances, it is also possible to have *static* (also known as *shared*, particularly to our Visual Basic brethren) members, which can be methods, properties, or fields. Static members are shared between instances of a class, so they can be thought of as global for objects of a given class. Static properties and fields enable you to access data that is independent of any object instances, and static methods enable you to execute commands related to the class type but not specific to object instances. When using static members, in fact, you don't even need to instantiate an object.

For example, the `Console.WriteLine()` and `Convert.ToString()` methods you have been using are static. At no point do you need to instantiate the `Console` or `Convert` classes (indeed, if you try, you'll find that you can't, as the constructors of these classes aren't publicly accessible, as discussed earlier).

There are many situations such as these where static properties and methods can be used to good effect. For example, you might use a static property to keep track of how many instances of a class have been created. In UML syntax, static members of classes appear with underlining, as shown in Figure 8-5.

## Static Constructors

When using static members in a class, you might want to initialize these members beforehand. You can supply a static member with an initial value as part of its declaration, but sometimes you might want to perform a more complex initialization, or perhaps perform some operations before assigning values or allowing static methods to execute.

You can use a static constructor to perform initialization tasks of this type. A class can have a single static constructor, which must have no access modifiers and cannot have any parameters. A static constructor can never be called directly; instead, it is executed when one of the following occurs:

- An instance of the class containing the static constructor is created.
- A static member of the class containing the static constructor is accessed.

In both cases, the static constructor is called first, before the class is instantiated or static members accessed. No matter how many instances of a class are created, its static constructor will be called only once. To differentiate between static constructors and the constructors described earlier in this chapter, all nonstatic constructors are also known as *instance constructors*.

## Static Classes

Often, you will want to use classes that contain only static members and cannot be used to instantiate objects (such as `Console`). A shorthand way to do this, rather than make the constructors of the class private, is to use a *static class*. A static class can contain only static members and can't have instance constructors, since by implication it can never be instantiated. Static classes can, however, have a static constructor, as described in the preceding section.

**NOTE** If you are completely new to OOP, you might like to take a break before embarking on the remainder of this chapter. It is important to fully grasp the fundamentals before learning about the more complicated aspects of this methodology.

## OOP TECHNIQUES

Now that you know the basics, and what objects are and how they work, you can spend some time looking at some of the other features of objects. This section covers all of the following:

- Interfaces
- Inheritance

- Polymorphism
- Relationships between objects
- Operator overloading
- Events
- Reference versus value types


## Interfaces

An *interface* is a collection of public instance (that is, nonstatic) methods and properties that are grouped together to encapsulate specific functionality. After an interface has been defined, you can implement it in a class. This means that the class will then support all of the properties and members specified by the interface.

Interfaces cannot exist on their own. You can't "instantiate an interface" as you can a class. In addition, interfaces cannot contain any code that implements its members; it just defines the members. The implementation must come from classes that implement the interface.

In the earlier coffee example, you might group together many of the more general-purpose properties and methods into an interface, such as `AddSugar()`, `Milk`, `Sugar`, and `Instant`. You could call this interface something like `IHotDrink` (interface names are normally prefixed with a capital `I`). You could use this interface on other objects, perhaps those of a `CupOfTea` class. You could therefore treat these objects in a similar way, and they can still have their own individual properties (`BeanType` for `CupOfCoffee` and `LeafType` for `CupOfTea`, for example).

Interfaces implemented on objects in UML are shown using *lollipop* syntax. In Figure 8-6, members of `IHotDrink` are split into a separate box using class-like syntax.



**FIGURE 8-6**

A class can support multiple interfaces, and multiple classes can support the same interface. The concept of an interface, therefore, makes life easier for users and other developers. For example, you might have some code that uses an object with a certain interface. Provided that you don't use other properties and methods of this object, it is possible to replace one object with another (code using

the `IHotDrink` interface shown earlier could work with both `CupOfCoffee` and `CupOfTea` instances, for example). In addition, the developer of the object itself could supply you with an updated version of an object, and as long as it supports an interface already in use, it would be easy to use this new version in your code.

Once an interface is published — that is, it has been made available to other developers or end users — it is good practice not to change it. One way of thinking about this is to imagine the interface as a contract between class creators and class consumers. You are effectively saying, “Every class that supports interface `x` will support these methods and properties.” If the interface changes later, perhaps due to an upgrade of the underlying code, this could cause consumers of that interface to run it incorrectly, or even fail. Instead, you should create a new interface that extends the old one, perhaps including a version number, such as `x2`. This has become the standard way of doing things, and you are likely to come across numbered interfaces frequently.

## Disposable Objects

One interface of particular interest is `IDisposable`. An object that supports the `IDisposable` interface must implement the `Dispose()` method — that is, it must provide code for this method. This method can be called when an object is no longer needed (just before it goes out of scope, for example) and should be used to free up any critical resources that might otherwise linger until the destructor method is called on garbage collection. This gives you more control over the resources used by your objects.

C# enables you to use a structure that makes excellent use of this method. The `using` keyword enables you to initialize an object that uses critical resources in a code block, where `Dispose()` is automatically called at the end of the code block:

```
<ClassName> <VariableName> = new <ClassName>();
...
using (<VariableName>
{
    ...
}
```

Alternatively, you can instantiate the object `<VariableName>` as part of the `using` statement:

```
using (<ClassName> <VariableName> = new <ClassName>())
{
    ...
}
```

In both cases, the variable `<VariableName>` will be usable within the `using` code block and will be disposed of automatically at the end (that is, `Dispose()` is called when the code block finishes executing).

## Inheritance

*Inheritance* is one of the most important features of OOP. Any class may *inherit* from another, which means that it will have all the members of the class from which it inherits. In OOP terminology, the class being inherited from (*derived from*) is the *parent* class (also known as the *base* class). Classes in C# can derive only from a single base class directly, although of course that base class can have a base class of its own, and so on.

Inheritance enables you to extend or create more specific classes from a single, more generic base class. For example, consider a class that represents a farm animal (as used by ace octogenarian developer Old MacDonald in his livestock application). This class might be called `Animal` and possess methods such as `EatFood()` or `Breed()`. You could create a derived class called `Cow`, which would support all of these methods but might also supply its own, such as `Moo()` and `SupplyMilk()`. You could also create another derived class, `Chicken`, with `Cluck()` and `LayEgg()` methods.

In UML, you indicate inheritance using arrows, as shown in Figure 8-7.




FIGURE 8-7

**NOTE** In Figure 8-7, the member return types are omitted for clarity.

When using inheritance from a base class, the question of member accessibility becomes an important one. Private members of the base class are not accessible from a derived class, but public members are. However, public members are accessible to both the derived class and external code. Therefore, if you could use only these two levels of accessibility, you couldn't have a member that was accessible both by the base class and the derived class but not external code.

To get around this, there is a third type of accessibility, *protected*, in which only derived classes have access to a member. As far as external code is aware, this is identical to a private member — it doesn't have access in either case.

As well as defining the protection level of a member, you can also define an inheritance behavior for it. Members of a base class can be *virtual*, which means that the member can be overridden by the class that inherits it. Therefore, the derived class can provide an alternative implementation for the member. This alternative implementation doesn't delete the original code, which is still accessible from within the class, but it does shield it from external code. If no alternative is supplied, then any external code that uses the member through the derived class automatically uses the base class implementation of the member.

**NOTE** Virtual members cannot be private because that would cause a paradox — it is impossible to say that a member can be overridden by a derived class at the same time you say that it is inaccessible from the derived class.

In the animals example, you could make `EatFood()` virtual and provide a new implementation for it on any derived class — for example, just on the `Cow` class, as shown in Figure 8-8. This displays the `EatFood()` method on the `Animal` and `Cow` classes to signify that they have their own implementations.

Base classes may also be defined as *abstract* classes. An abstract class can't be instantiated directly; to use it you need to inherit from it. Abstract classes can have abstract members, which have no implementation in the base class, so an implementation must be supplied in the derived class. If `Animal` were an abstract class, then the UML would look as shown in Figure 8-9.

**NOTE** Abstract class names are shown in *italics* (or with a dashed line for their boxes).





FIGURE 8-8

In Figure 8-9, both `EatFood()` and `Breed()` are shown in the derived classes `Chicken` and `Cow`, implying that these methods are either abstract (and, therefore, must be overridden in derived classes) or virtual (and, in this case, have been overridden in `Chicken` and `Cow`). Of course, abstract base classes can provide implementation of members, which is very common. The fact that you can't instantiate an abstract class doesn't mean you can't encapsulate functionality in it.

**FIGURE 8-9**

Finally, a class may be *sealed*. A sealed class cannot be used as a base class, so no derived classes are possible.

C# provides a common base class for all objects called `object` (which is an alias for the `System.Object` class in the .NET Framework). You take a closer look at this class in Chapter 9.

**NOTE** Interfaces, described earlier in this chapter, can also inherit from other interfaces. Unlike classes, interfaces can inherit from multiple base interfaces (in the same way that classes can support multiple interfaces).

## Polymorphism

One consequence of inheritance is that classes deriving from a base class have an overlap in the methods and properties that they expose. Because of this, it is often possible to treat objects instantiated from classes with a base type in common using identical syntax. For example, if a base class called `Animal` has a method called `EatFood()`, then the syntax for calling this method from the derived classes `Cow` and `Chicken` will be similar:

```

Cow myCow = new Cow();
Chicken myChicken = new Chicken();
myCow.EatFood();
myChicken.EatFood();
  
```

Polymorphism takes this a step further. You can assign a variable that is of a derived type to a variable of one of the base types, as shown here:

```
Animal myAnimal = myCow;
```

No casting is required for this. You can then call methods of the base class through this variable:

```
myAnimal.EatFood();
```

This results in the implementation of `EatFood()` in the derived class being called. Note that you can't call methods defined on the derived class in the same way. The following code won't work:

```
myAnimal.Moo();
```

However, you can cast a base type variable into a derived class variable and call the method of the derived class that way:

```
Cow myNewCow = (Cow)myAnimal;  
myNewCow.Moo();
```

This casting causes an exception to be raised if the type of the original variable was anything other than `Cow` or a class derived from `Cow`. There are ways to determine the type of an object, which you'll learn in the next chapter.

Polymorphism is an extremely useful technique for performing tasks with a minimum of code on different objects descending from a single class. It isn't just classes sharing the same parent class that can make use of polymorphism. It is also possible to treat, say, a child and a grandchild class in the same way, as long as there is a common class in their inheritance hierarchy.

As a further note here, remember that in C# all classes derive from the base class `object` at the root of their inheritance hierarchies. It is therefore possible to treat all objects as instances of the class `object`. This is how `WriteLine()` can process an almost infinite number of parameter combinations when building strings. Every parameter after the first is treated as an `object` instance, allowing output from any object to be written to the screen. To do this, the method `ToString()` (a member of `object`) is called. You can override this method to provide an implementation suitable for your class, or simply use the default, which returns the class name (qualified according to any namespaces it is in).

## Interface Polymorphism

Although you can't instantiate interfaces in the same way as objects, you can have a variable of an interface type. You can then use the variable to access methods and properties exposed by this interface on objects that support it.

For example, suppose that instead of an `Animal` base class being used to supply the `EatFood()` method, you place this `EatFood()` method on an interface called `IConsume`. The `Cow` and `Chicken` classes could both support this interface, the only difference being that they are forced to provide an implementation for `EatFood()` because interfaces contain no implementation. You can then access this method using code such as the following:

```
Cow myCow = new Cow();  
Chicken myChicken = new Chicken();  
IConsume consumeInterface;  
consumeInterface = myCow;  
consumeInterface.EatFood();
```

---

```
consumeInterface = myChicken;
consumeInterface.EatFood();
```

This provides a simple way for multiple objects to be called in the same manner, and it doesn't rely on a common base class. For example, this interface could be implemented by a class called `VenusFlyTrap` that derives from `Vegetable` instead of `Animal`:

```
VenusFlyTrap myVenusFlyTrap = new VenusFlyTrap();
IConsume consumeInterface;
consumeInterface = myVenusFlyTrap;
consumeInterface.EatFood();
```

In the preceding code snippets, calling `consumeInterface.EatFood()` results in the `EatFood()` method of the `Cow`, `Chicken`, or `VenusFlyTrap` class being called, depending on which instance has been assigned to the interface type variable.

Note here that derived classes inherit the interfaces supported by their base classes. In the first of the preceding examples, it might be that either `Animal` supports `IConsume` or that both `Cow` and `Chicken` support `IConsume`. Remember that classes with a base class in common do not necessarily have interfaces in common, and vice versa.

## Relationships between Objects

Inheritance is a simple relationship between objects that results in a base class being completely exposed by a derived class, where the derived class can also have some access to the inner workings of its base class (through protected members). There are other situations in which relationships between objects become important.

This section takes a brief look at the following

- **Containment** — One class contains another. This is similar to inheritance but allows the containing class to control access to members of the contained class and even perform additional processing before using members of a contained class.
- **Collections** — One class acts as a container for multiple instances of another class. This is similar to having arrays of objects, but collections have additional functionality, including indexing, sorting, resizing, and more.

### Containment

Containment is simple to achieve by using a member field to hold an object instance. This member field might be public, in which case users of the container object have access to its exposed methods and properties, much like with inheritance. However, you won't have access to the internals of the class via the derived class, as you would with inheritance.

Alternatively, you can make the contained member object a private member. If you do this, then none of its members will be accessible directly by users, even if they are public. Instead, you can provide access to these members using members of the containing class. This means that you have complete control over which members of the contained class to expose, if any, and you can perform additional processing in the containing class members before accessing the contained class members.

For example, a `Cow` class might contain an `Udder` class with the public method `Milk()`. The `Cow` object could call this method as required, perhaps as part of its `SupplyMilk()` method, but these details will not be apparent (or important) to users of the `Cow` object.

Contained classes can be visualized in UML using an association line. For simple containment, you label the ends of the lines with 1s, showing a one-to-one relationship (one `Cow` instance will contain one `Udder` instance). You can also show the contained `Udder` class instance as a private field of the `Cow` class for clarity (see Figure 8-10).




FIGURE 8-10

## Collections

Chapter 5 described how you can use arrays to store multiple variables of the same type. This also works for objects (remember, the variable types you have been using are really objects, so this is no real surprise). Here's an example:

```
Animal[] animals = new Animal[5];
```

A *collection* is basically an array with bells and whistles. Collections are implemented as classes in much the same way as other objects. They are often named in the plural form of the objects they store — for example, a class called `Animals` might contain a collection of `Animal` objects.

The main difference from arrays is that collections usually implement additional functionality, such as `Add()` and `Remove()` methods to add and remove items to and from the collection. There is also usually an `Item` property that returns an object based on its index. More often than not this property is implemented in such a way as to allow more sophisticated access. For example, it would be possible to design `Animals` so that a given `Animal` object could be accessed by its name.

In UML you can visualize this as shown in Figure 8-11. Members are not included in Figure 8-11 because it's the relationship that is being illustrated. The numbers on the ends of the connecting lines show that one `Animals` object will contain zero or more `Animal` objects. You take a more detailed look at collections in Chapter 11.




FIGURE 8-11

## Operator Overloading

Earlier in the book, you saw how operators can be used to manipulate simple variable types. There are times when it is logical to use operators with objects instantiated from your own classes. This is possible because classes can contain instructions regarding how operators should be treated.

For example, you might add a new property to the `Animal` class called `Weight`. You could then compare animal weights using the following:

```

if (cowA.Weight > cowB.Weight)
{
    ...
}
  
```

Using operator overloading, you can provide logic that uses the `Weight` property implicitly in your code, so that you can write code such as the following:

```

if (cowA > cowB)
{
    ...
}
  
```

Here, the greater-than operator (`>`) has been *overloaded*. An overloaded operator is one for which you have written the code to perform the operation involved — this code is added to the class definition of one of the classes that it operates on. In the preceding example, you are using two `Cow` objects, so the operator overload definition is contained in the `Cow` class. You can also overload operators to work with different classes in the same way, where one (or both) of the class definitions contains the code to achieve this.

You can only overload existing C# operators in this way; you can't create new ones. However, you can provide implementations for both unary (single operand) and binary (two operands) usages of operators such as `+` or `>`. You see how to do this in C# in Chapter 13.

## Events

Objects can raise (and consume) *events* as part of their processing. Events are important occurrences that you can act on in other parts of code, similar to (but more powerful than) exceptions. You might, for example, want some specific code to execute when an `Animal` object is added to an `Animals` collection, where that code isn't part of either the `Animals` class or the code that calls the `Add()` method. To do this, you need to add an *event handler* to your code, which is a special kind of function that is called when the event occurs. You also need to configure this handler to listen for the event you are interested in.

You can create *event-driven applications*, which are far more prolific than you might think. For example, bear in mind that Windows-based applications are entirely dependent on events. Every button click or scroll bar drag you perform is achieved through event handling, as the events are triggered by the mouse or keyboard.

Later in this chapter you will see how this works in Windows applications, and there is a more in-depth discussion of events in Chapter 13.

## Reference Types versus Value Types

Data in C# is stored in a variable in one of two ways, depending on the type of the variable. This type will fall into one of two categories: reference or value. The difference is as follows:

- Value types store themselves and their content in one place in memory.
- Reference types hold a reference to somewhere else in memory (called the *heap*) where content is stored.

In fact, you don't have to worry about this too much when using C#. So far, you've used `string` variables (which are reference types) and other simple variables (most of which are value types, such as `int`) in pretty much the same way.

One key difference between value types and reference types is that value types always contain a value, whereas reference types can be `null`, reflecting the fact that they contain no value. It is, however, possible to create a value type that behaves like a reference type in this respect (that is, it can be `null`) by using *nullable types*. These are described in Chapter 12, when you look at the advanced technique of *generic types* (which include nullable types).

The only simple types that are reference types are `string` and `object`, although arrays are implicitly *reference types as well*. Every class you create will be a reference type, which is why this is stressed here.

## OOP IN DESKTOP APPLICATIONS

In Chapter 2, you created a simple desktop application in C# using Windows Presentation Foundation (WPF). WPF desktop applications are heavily dependent on OOP techniques, and this

section takes a look at this to illustrate some of the points made in this chapter. The following Try It Out enables you to work through a simple example.

### TRY IT OUT Objects in Action: Ch08Ex01

1. Create a new WPF application called Ch08Ex01 and save it in the directory C:\BegVCSharp\Chapter08.
2. Add a new Button control using the Toolbox, and position it in the center of MainWindow, as shown in Figure 8-12.




FIGURE 8-12

3. Double-click on the button to add code for a mouse click. Modify the code that appears as follows:

```
private void Button_Click_1(object sender, RoutedEventArgs e)
{
    ((Button)sender).Content = "Clicked!";
    Button newButton = new Button();
    newButton.Content = "New Button! ";
    newButton.Margin = new Thickness(10, 10, 200, 200);
    newButton.Click += newButton_Click;
    ((Grid)((Button)sender).Parent).Children.Add(newButton);
}
private void newButton_Click(object sender, RoutedEventArgs e)
{
    ((Button)sender).Content = "Clicked!!";
}
```

4. Run the application. The window is shown in Figure 8-13.




FIGURE 8-13

5. Click the button marked Button. The display changes (see Figure 8-14).




FIGURE 8-14

6. Click the button marked New Button! The display changes (see Figure 8-15).




FIGURE 8-15

### How It Works

By adding just a few lines of code you've created a desktop application that does something, while at the same time illustrating some OOP techniques in C#. The phrase "everything's an object" is even more true when it comes to desktop applications. From the form that runs to the controls on the form, you need to use OOP techniques all the time. This example highlights some of the concepts you looked at earlier in this chapter to show how everything fits together.

The first thing you do in this application is add a new button to the `MainWindow` window. The button is an object; it's an instance of a class called `Button`, and the window is an instance of a class called `MainWindow`, which is derived from a class called `Window`. Next, by double-clicking the button, you add an event handler to listen for the `Click` event that the `Button` class exposes. The event handler is added to the code for the `MainWindow` object that encapsulates your application, as a private method:

```
private void Button_Click_1(object sender, RoutedEventArgs e)
{
}
```

The code uses the C# keyword `private` as a qualifier. Don't worry too much about that for now; the next chapter explains the C# code required for the OOP techniques covered in this chapter.

The first line of code you add changes the text on the button that is clicked. This makes use of polymorphism, described earlier in the chapter. The `Button` object representing the button that you click is sent to the event handler as an `object` parameter, which you cast into a `Button` type (this is possible

because the `Button` object inherits from `System.Object`, which is the .NET class that `object` is an alias for). You then change the `Content` property of the object to change the text displayed:

```
((Button)sender).Content = "Clicked!";
```

Next, you create a new `Button` object with the `new` keyword (note that namespaces are set up in this project to enable this simple syntax; otherwise, you need to use the fully qualified name of this object, `System.Windows.Controls.Button`):

```
Button newButton = new Button();
```

You also set the `Content` and `Margin` properties of the newly created `Button` object to suitable values for displaying the button. Note that the `Margin` property is of type `Thickness`, so you create a `Thickness` object using a non-default constructor before assigning it to the property:

```
newButton.Content = "New Button!";
newButton.Margin = new Thickness(10, 10, 200, 200);
```

Elsewhere in the code a new event handler is added, which you use to respond to the `Click` event generated by the new button:

```
private void newButton_Click(object sender, RoutedEventArgs e)
{
    ((Button)sender).Content = "Clicked!!";
}
```

You register the event handler as a listener for the `Click` event, using overloaded operator syntax:

```
newButton.Click += newButton_Click;
```

Finally, you add the new button to the window. To do this, you find the parent of the existing button (using its `Parent` property), cast it to the correct type (which is `Grid`), and use the `Add()` method of the `Grid.Children` property to add the button, passing the button as a method parameter:

```
((Grid)((Button)sender).Parent).Children.Add(newButton);
```

This code looks more complicated than it actually is. Once you get the hang of the way that WPF represents the content of a window through a hierarchy of controls (including buttons and containers), this sort of thing will become second nature.

This short example used almost all of the techniques introduced in this chapter. As you can see, OOP programming needn't be complicated — it just requires a different point of view to get right.

## EXERCISES

### 8.1 Which of the following are real levels of accessibility in OOP?

- a. Friend
- b. Public
- c. Secure
- d. Private

- e. Protected
- f. Loose
- g. Wildcard

8.2 "You must call the destructor of an object manually or it will waste memory." True or false?

8.3 Do you need to create an object to call a static method of its class?

8.4 Draw a UML diagram similar to the ones shown in this chapter for the following classes and interface:

- An abstract class called `HotDrink` that has the methods `Drink`, `AddMilk`, and `AddSugar`, and the properties `Milk` and `Sugar`
- An interface called `ICup` that has the methods `Refill` and `Wash`, and the properties `Color` and `Volume`
- A class called `CupOfCoffee` that derives from `HotDrink`, supports the `ICup` interface, and has the additional property `BeanType`
- A class called `CupOfTea` that derives from `HotDrink`, supports the `ICup` interface, and has the additional property `LeafType`

8.5 Write some code for a function that will accept either of the two cup objects in the preceding example as a parameter. The function should call the `AddMilk`, `Drink`, and `Wash` methods for any cup object it is passed.

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Objects and classes	Objects are the building blocks of OOP applications. Classes are type definitions that are used to instantiate objects. Objects can contain data and/or expose operations that other code can use. Data can be made available to external code through properties, and operations can be made available to external code through methods. Both properties and methods are referred to as class members. Properties can allow read access, write access, or both. Class members can be public (available to all code), or private (available only to code inside the class definition). In .NET, everything is an object.
Object life cycle	An object is instantiated by calling one of its constructors. When an object is no longer needed, it is destroyed by executing its destructor. To clean up after an object, it is often necessary to manually dispose of it.
Static and instance members	Instance members are available only on object instances of a class. Static members are available only through the class definition directly, and are not associated with an instance.
Interfaces	Interfaces are a collection of public properties and methods that can be implemented on a class. An instance-typed variable can be assigned a value of any object whose class definition implements that interface. Only the interface-defined members are then available through the variable.
Inheritance	Inheritance is the mechanism through which one class definition can derive from another. A class inherits members from its parent, of which it can have only one. Child classes cannot access private members in its parent, but it is possible to define protected members that are available only within a class or classes that derive from that class. Child classes can override members that are defined as virtual in a parent class. All classes have an inheritance chain that ends in <code>System.Object</code> , which has the alias <code>object</code> in C#.
Polymorphism	All objects instantiated from a derived class can be treated as if they were instances of a parent class.
Object relationships and features	Objects can contain other objects, and can also represent collections of other objects. To manipulate objects in expressions, you often need to define how operators work with objects, through operator overloading. Objects can expose events that are triggered due to some internal process, and client code can respond to events by providing event handlers.

# 9

## Defining Classes

### WHAT YOU WILL LEARN IN THIS CHAPTER

---

- Defining classes and interfaces in C#
- Using the keywords that control accessibility and inheritance
- Discovering the `System.Object` class and its role in class definitions
- Using some helpful tools provided by Visual Studio (VS)
- Defining class libraries
- Knowing the differences and similarities between interfaces and abstract classes
- Exploring struct types
- Understanding important object copying considerations

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 9 download and individually named according to the names throughout the chapter.

In Chapter 8, you looked at the features of object-oriented programming (OOP). In this chapter, you put theory into practice and define classes in C#. You won't go so far as to define class members in this chapter, but you will concentrate on the class definitions themselves.

To begin, you explore the basic class definition syntax, the keywords you can use to determine class accessibility and more, and the way in which you can specify inheritance. You also look at interface definitions because they are similar to class definitions in many ways.

The rest of the chapter covers various related topics that apply when defining classes in C#.

## CLASS DEFINITIONS IN C#

C# uses the `class` keyword to define classes:

```
class MyClass
{
    // Class members.
```

This code defines a class called `MyClass`. Once you have defined a class, you are free to instantiate it anywhere else in your project that has access to the definition. By default, classes are declared as *internal*, meaning that only code in the current project will have access to them. You can specify this explicitly using the `internal` access modifier keyword as follows (although you don't have to):

```
internal class MyClass
{
    // Class members.
```

Alternatively, you can specify that the class is public and should also be accessible to code in other projects. To do so, you use the `public` keyword:

```
public class MyClass
{
    // Class members.
```

In addition to these two access modifier keywords, you can also specify that the class is either *abstract* (cannot be instantiated, only inherited, and can have abstract members) or *sealed* (cannot be inherited). To do this, you use one of the two mutually exclusive keywords, `abstract` or `sealed`. An abstract class is declared as follows:

```
public abstract class MyClass
{
    // Class members, may be abstract.
```

Here, `MyClass` is a public abstract class, while internal abstract classes are also possible.

Sealed classes are declared as follows:

```
public sealed class MyClass
{
    // Class members.
```

As with abstract classes, sealed classes can be public or internal.

Inheritance can also be specified in the class definition. You simply put a colon after the class name, followed by the base class name:

```
public class MyClass : MyBase
{
    // Class members.
```

Only one base class is permitted in C# class definitions; and if you inherit from an abstract class, you must implement all the abstract members inherited (unless the derived class is also abstract).

The compiler does not allow a derived class to be more accessible than its base class. This means that an internal class can inherit from a public base, but a public class can't inherit from an internal base. This code is legal:

```
public class MyBase
{
    // Class members.
}
internal class MyClass : MyBase
{
    // Class members.
}
```

The following code won't compile:

```
internal class MyBase
{
    // Class members.
}
public class MyClass : MyBase
{
    // Class members.
}
```

If no base class is used, the class inherits only from the base class `System.Object` (which has the alias `object` in C#). Ultimately, all classes have `System.Object` at the root of their inheritance hierarchy. You will take a closer look at this fundamental class a little later.

In addition to specifying base classes in this way, you can also specify interfaces supported after the colon character. If a base class is specified, it must be the first thing after the colon, with interfaces specified afterward. If no base class is specified, you specify the interfaces immediately after the colon. Commas must be used to separate the base class name (if there is one) and the interface names from one another.

For example, you could add an interface to `MyClass` as follows:

```
public class MyClass : IMyInterface
{
    // Class members.
}
```

All interface members must be implemented in any class that supports the interface, although you can provide an “empty” implementation (with no functional code) if you don’t want to do anything with a given interface member, and you can implement interface members as abstract in abstract classes.

The following declaration is invalid because the base class `MyBase` isn’t the first entry in the inheritance list:

```
public class MyClass : IMyInterface, MyBase
{
    // Class members.
}
```

The correct way to specify a base class and an interface is as follows:

```
public class MyClass : MyBase, IMyInterface
{
    // Class members.
}
```

Remember that multiple interfaces are possible, so the following is also valid:

```
public class MyClass : MyBase, IMyInterface, IMySecondInterface
{
    // Class members.
}
```

Table 9-1 shows the allowed access modifier combinations for class definitions.

**TABLE 9-1:** Access Modifiers for Class Definitions

MODIFIER	DESCRIPTION
none or internal	Class is accessible only from within the current project
public	Class is accessible from anywhere
abstract or internal abstract	Class is accessible only from within the current project, and cannot be instantiated, only derived from
public abstract	Class is accessible from anywhere, and cannot be instantiated, only derived from
sealed or internal sealed	Class is accessible only from within the current project, and cannot be derived from, only instantiated
public sealed	Class is accessible from anywhere, and cannot be derived from, only instantiated

## Interface Definitions

Interfaces are declared in a similar way to classes, but using the `interface` keyword, rather than `class`:

```
interface IMyInterface
{
    // Interface members.
}
```

The access modifier keywords `public` and `internal` are used in the same way; and as with classes, interfaces are defined as `internal` by default. To make an interface publicly accessible, you must use the `public` keyword:

```
public interface IMyInterface
{
    // Interface members.
}
```

The keywords `abstract` and `sealed` are not allowed because neither modifier makes sense in the context of interfaces (they contain no implementation, so they can't be instantiated directly, and they must be inheritable to be useful).

Interface inheritance is also specified in a similar way to class inheritance. The main difference here is that multiple base interfaces can be used, as shown here:

```
public interface IMyInterface : IMyBaseInterface, IMyBaseInterface2
{
    // Interface members.
}
```

Interfaces are not classes, and thus do not inherit from `System.Object`. However, the members of `System.Object` are available via an interface type variable, purely for convenience. In addition, as already discussed, it is impossible to instantiate an interface in the same way as a class. The following Try It Out provides an example of some class definitions, along with some code that uses them.

### TRY IT OUT Defining Classes: Ch09Ex01\Program.cs

1. Create a new console application called Ch09Ex01 and save it in the directory `C:\BegVCSharp\Chapter09`.
2. Modify the code in `Program.cs` as follows:

```
using static System.Console;
namespace Ch09Ex01
{
    public abstract class MyBase {}
    internal class MyClass : MyBase {}
    public interface IMyBaseInterface {}
    internal interface IMyBaseInterface2 {}
    internal interface IMyInterface : IMyBaseInterface, IMyBaseInterface2 {}
    internal sealed class MyComplexClass : MyClass, IMyInterface {}
    class Program
    {
        static void Main(string[] args)
        {
            MyComplexClass myObj = new MyComplexClass();
            WriteLine(myObj.ToString());
            ReadKey();
        }
    }
}
```

3. Execute the project. Figure 9-1 shows the output.




FIGURE 9-1

## How It Works

This project defines classes and interfaces in the inheritance hierarchy shown in Figure 9-2.




FIGURE 9-2

`Program` is included because it is a class defined in the same way as the other classes, even though it isn't part of the main class hierarchy. The `Main()` method possessed by this class is the entry point for your application.

`MyBase` and `IMyBaseInterface` are public definitions, so they are available from other projects. The other classes and interfaces are internal, and only available in this project.

The code in `Main()` calls the `ToString()` method of `myObj`, an instance of `MyComplexClass`:

```

MyComplexClass myObj = new MyComplexClass();
WriteLine(myObj.ToString());

```

`ToString()` is one of the methods inherited from `System.Object` (not shown in the diagram because members of this class are omitted for clarity) and simply returns the class name of the object as a string, qualified by any relevant namespaces.

This example doesn't do a lot, but you will return to it later in this chapter, where it is used to demonstrate several key concepts and techniques.

## SYSTEM.OBJECT

Because all classes inherit from `System.Object`, all classes have access to the protected and public members of this class. Therefore, it is worthwhile to take a look at what is available there. `System.Object` contains the methods described in Table 9-2.

**TABLE 9-2:** Methods of `System.Object`

METHOD	RETURN TYPE	VIRTUAL	STATIC	DESCRIPTION
<code>Object()</code>	N/A	No	No	Constructor for the <code>System.Object</code> type. Automatically called by constructors of derived types.
<code>~Object()</code> (also known as <code>Finalize()</code> — see the next section)	N/A	No	No	Destructor for the <code>System.Object</code> type. Automatically called by destructors of derived types; cannot be called manually.
<code>Equals(object)</code>	bool	Yes	No	Compares the object for which this method is called with another object and returns <code>true</code> if they are equal. The default implementation checks whether the <code>object</code> parameter refers to the same object (because objects are reference types). This method can be overridden if you want to compare objects in a different way, for example, to compare the state of two objects.
<code>Equals(object, object)</code>	bool	No	Yes	Compares the two objects passed to it and checks whether they are equal. This check is performed using the <code>Equals(object)</code> method. If both objects are null references, then this method returns <code>true</code> .
<code>ReferenceEquals(object, object)</code>	bool	No	Yes	Compares the two objects passed to it and checks whether they are references to the same instance.

*continues*

**TABLE 9-2** (*continued*)

METHOD	RETURN TYPE	VIRTUAL	STATIC	DESCRIPTION
ToString()	string	Yes	No	Returns a string corresponding to the object instance. By default, this is the qualified name of the class type, but this can be overridden to provide an implementation appropriate to the class type.
MemberwiseClone()	object	No	No	Copies the object by creating a new object instance and copying members. This member copying does not result in new instances of these members. Any reference type members of the new object refer to the same objects as the original class. This method is protected, so it can be used only from within the class or from derived classes.
GetType()	System.Type	No	No	Returns the type of the object in the form of a System.Type object.
GetHashCode()	int	Yes	No	Used as a hash function for objects where this is required. A hash function returns a value identifying the object state in some compressed form.

These are the basic methods that must be supported by object types in the .NET Framework, although you might never use some of them (or you might use them only in special circumstances, such as `GetHashCode()`).

`GetType()` is helpful when you are using polymorphism because it enables you to perform different operations with objects depending on their type, rather than the same operation for all objects, as is often the case. For example, if you have a function that accepts an object type parameter (meaning you can pass it just about anything), you might perform additional tasks if certain objects are encountered. Using a combination of `GetType()` and `typeof` (a C# operator that converts a class name into a `System.Type` object), you can perform comparisons such as the following:

```
if (myObj.GetType() == typeof(MyComplexClass))
{
    // myObj is an instance of the class MyComplexClass.
}
```

The `System.Type` object returned is capable of a lot more than that, but only this is covered here. It can also be very useful to override the `ToString()` method, particularly in situations where the contents of an object can be easily represented with a single human-readable string. You see these `System.Object` methods repeatedly in subsequent chapters, so you'll learn more details as necessary.

## CONSTRUCTORS AND DESTRUCTORS

When you define a class in C#, it's often unnecessary to define associated constructors and destructors because the compiler adds them for you when you build your code if you don't supply them. However, you can provide your own, if required, which enables you to initialize and clean up after your objects, respectively.

You can add a simple constructor to a class using the following syntax:

```
class MyClass
{
    public MyClass()
    {
        // Constructor code.
    }
}
```

This constructor has the same name as the class that contains it, has no parameters (making it the default constructor for the class), and is public so that objects of the class can be instantiated using this constructor (refer to Chapter 8 for more information about this).

You can also use a private default constructor, meaning that object instances of this class cannot be created using this constructor (it is *non-creatable* — again, see the discussion in Chapter 8):

```
class MyClass
{
    private MyClass()
    {
        // Constructor code.
    }
}
```

Finally, you can add nondefault constructors to your class in a similar way, simply by providing parameters:

```
class MyClass
{
    public MyClass()
    {
        // Default constructor code.
    }
    public MyClass(int myInt)
    {
        // Nondefault constructor code (uses myInt).
    }
}
```

You can supply an unlimited number of constructors (until you run out of memory or out of distinct sets of parameters, so maybe “almost unlimited” is more appropriate).

Destructors are declared using a slightly different syntax. The destructor used in .NET (and supplied by the `System.Object` class) is called `Finalize()`, but this isn’t the name you use to declare a destructor. Instead of overriding `Finalize()`, you use the following:

```
class MyClass
{
    ~MyClass()
    {
        // Destructor body.
    }
}
```

Thus, the destructor of a class is declared by the class name (just as the constructor is), with the tilde (~) prefix. The code in the destructor is executed when garbage collection occurs, enabling you to free resources. After the destructor is called, implicit calls to the destructors of base classes also occur, including a call to `Finalize()` in the `System.Object` root class. This technique enables the .NET Framework to ensure that this occurs, because overriding `Finalize()` would mean that base class calls would need to be explicitly performed, which is potentially dangerous (you learn how to call base class methods in the next chapter).

## Constructor Execution Sequence

If you perform multiple tasks in the constructors of a class, it can be handy to have this code in one place, which has the same benefits as splitting code into functions, as shown in Chapter 6. You could do this using a method (see Chapter 10), but C# provides a nice alternative. You can configure any constructor to call any other constructor before it executes its own code.

First, though, you need to take a closer look at what happens by default when you instantiate a class instance. Apart from facilitating the centralization of initialization code, as noted previously, this is worth knowing about in its own right. During development, objects often don’t behave quite as you expect them to due to errors during constructor calling — usually a base class somewhere in the inheritance hierarchy of your class that you are not instantiating correctly, or information that is not being properly supplied to base class constructors. Understanding what happens during this phase of an object’s lifecycle can make it much easier to solve this sort of problem.

For a derived class to be instantiated, its base class must be instantiated. For this base class to be instantiated, its own base class must be instantiated, and so on all the way back to `System.Object` (the root of all classes). As a result, whatever constructor you use to instantiate a class, `System.Object.Object()` is always called first.

Regardless of which constructor you use in a derived class (the default constructor or a nondefault constructor), unless you specify otherwise, the default constructor for the base class is used. (You’ll

see how to change this behavior shortly.) Here's a short example illustrating the sequence of execution. Consider the following object hierarchy:

```
public class MyBaseClass
{
    public MyBaseClass()
    {
    }
    public MyBaseClass(int i)
    {
    }
}
public class MyDerivedClass : MyBaseClass
{
    public MyDerivedClass()
    {
    }
    public MyDerivedClass(int i)
    {
    }
    public MyDerivedClass(int i, int j)
    {
    }
}
```

You could instantiate `MyDerivedClass` as follows:

```
MyDerivedClass myObj = new MyDerivedClass();
```

In this case, the following sequence of events will occur:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass()` constructor will execute.
- The `MyDerivedClass.MyDerivedClass()` constructor will execute.

Alternatively, you could use the following:

```
MyDerivedClass myObj = new MyDerivedClass(4);
```

The sequence is as follows:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass()` constructor will execute.
- The `MyDerivedClass.MyDerivedClass(int i)` constructor will execute.

Finally, you could use this:

```
MyDerivedClass myObj = new MyDerivedClass(4, 8);
```

The result is the following sequence:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass()` constructor will execute.
- The `MyDerivedClass.MyDerivedClass(int i, int j)` constructor will execute.

This system works fine most of the time, but sometimes you will want a little more control over the events that occur. For example, in the last instantiation example, you might want to have the following sequence:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass(int i)` constructor will execute.
- The `MyDerivedClass.MyDerivedClass(int i, int j)` constructor will execute.

Using this sequence you could place the code that uses the `int i` parameter in `MyBaseClass(int i)`, which means that the `MyDerivedClass(int i, int j)` constructor would have less work to do — it would only need to process the `int j` parameter. (This assumes that the `int i` parameter has an identical meaning in both scenarios, which might not always be the case; but in practice, with this kind of arrangement, it usually is.) C# allows you to specify this kind of behavior if you want.

To do this, you can use a *constructor initializer*, which consists of code placed after a colon in the method definition. For example, you could specify the base class constructor to use in the definition of the constructor in your derived class, as follows:

```
public class MyDerivedClass : MyBaseClass
{
    ...
    public MyDerivedClass(int i, int j) : base(i)
    {
    }
}
```

The `base` keyword directs the .NET instantiation process to use the base class constructor, which has the specified parameters. Here, you are using a single `int` parameter (the value of which is the value passed to the `MyDerivedClass` constructor as the parameter `i`), so `MyBaseClass(int i)` will be used. Doing this means that `MyBaseClass` will not be called, giving you the sequence of events listed prior to this example — exactly what you want here.

You can also use this keyword to specify literal values for base class constructors, perhaps using the default constructor of `MyDerivedClass` to call a nondefault constructor of `MyBaseClass`:

```
public class MyDerivedClass : MyBaseClass
{
    public MyDerivedClass() : base(5)
    {
    }
    ...
}
```

This gives you the following sequence:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass(int i)` constructor will execute.
- The `MyDerivedClass.MyDerivedClass()` constructor will execute.

As well as this `base` keyword, you can use one more keyword as a constructor initializer: `this`. This keyword instructs the .NET instantiation process to use a nondefault constructor on the current class before the specified constructor is called:

```
public class MyDerivedClass : MyBaseClass
{
    public MyDerivedClass() : this(5, 6)
    {
    }
    ...
    public MyDerivedClass(int i, int j) : base(i)
    {
    }
}
```

Here, using the `MyDerivedClass.MyDerivedClass()` constructor gives you the following sequence:

- The `System.Object.Object()` constructor will execute.
- The `MyBaseClass.MyBaseClass(int i)` constructor will execute.
- The `MyDerivedClass.MyDerivedClass(int i, int j)` constructor will execute.
- The `MyDerivedClass.MyDerivedClass()` constructor will execute.

The only limitation here is that you can specify only a single constructor using a constructor initializer. However, as demonstrated in the last example, this isn't much of a limitation, because you can still construct fairly sophisticated execution sequences.

**NOTE** If you don't specify a constructor initializer for a constructor, the compiler adds one for you: `base()`. This results in the default behavior described earlier in this section.

Be careful not to accidentally create an infinite loop when defining constructors. For example, consider this code:

```
public class MyBaseClass
{
    public MyBaseClass() : this(5)
    {
    }
```

```
public MyBaseClass(int i) : this()  
{  
}  
}
```

Using either one of these constructors requires the other to execute first, which in turn requires the other to execute first, and so on. This code will compile, but if you try to instantiate `MyBaseClass` you will receive a `SystemOverflowException`.


# OOP TOOLS IN VISUAL STUDIO

Because OOP is such a fundamental aspect of the .NET Framework, several tools are provided by Visual Studio to aid development of OOP applications. This section describes some of these.

## The Class View Window

In Chapter 2, you saw that the Solution Explorer window shares space with a window called Class View. This window shows you the class hierarchy of your application and enables you to see at a glance the characteristics of the classes you use. Figure 9-3 shows a view of the example project in the previous Try It Out.

The window is divided into two main sections; the bottom section shows members of types. Note that Figure 9-3 shows the display when all items in the Class View Settings drop-down, at the top of the Class View window, are checked.



**FIGURE 9-3**

Many symbols can be used here, including the ones shown in Table 9-3.

**TABLE 9-3:** Class View Icons

ICON	MEANING	ICON	MEANING	ICON	MEANING
	Project		Property		Event
	Namespace		Field		Delegate
	Class		Struct		Assembly
	Interface		Enumeration		
	Method		Enumeration item		

Some of these are used for type definitions other than classes, such as enumerations and struct types.

Some of the entries can have other symbols placed below them, signifying their access level (no symbol appears for public entries). These are listed in Table 9-4.

**TABLE 9-4:** Additional Class View Icons

ICON	MEANING	ICON	MEANING	ICON	MEANING
	Private		Protected		Internal

No symbols are used to denote abstract, sealed, or virtual entries.

As well as being able to look at this information here, you can also access the relevant code for many of these items. Double-clicking on an item, or right-clicking and selecting Go To Definition, takes you straight to the code in your project that defines the item, if it is available. If the code isn't available, such as code in an inaccessible base type (for example, `System.Object`), you instead have the option to select Browse Definition, which will take you to the Object Browser view (described in the next section).

One other entry that appears in Figure 9-3 is Project References. This enables you to see which assemblies are referenced by your projects, which in this case includes (among others) the core .NET types in `mscorlib` and `System`, data access types in `System.Data`, and XML manipulation types in `System.Xml`. The references here can be expanded, showing you the namespaces and types contained within these assemblies.

You can find occurrences of types and members in your code by right-clicking on an item and selecting Find All References; a list of search results displays in the Find Symbol Results window, which

appears at the bottom of the screen as a tabbed window in the Error List display area. You can also rename items using the Class View window. If you do this, you're given the option to rename references to the item wherever it occurs in your code. This means you have no excuse for spelling mistakes in class names because you can change them as often as you like!

In addition, you can navigate through your code with a view called *Call Hierarchy*, which is accessible from the Class View window through the View Call Hierarchy right-click menu option. This functionality is extremely useful for looking at how class members interact with each other, and you'll look at it in the next chapter.

## The Object Browser

The Object Browser is an expanded version of the Class View window, enabling you to view other classes available to your project, and even external classes. It is entered either automatically (for example, in the situation noted in the last section) or manually via View  $\Rightarrow$  Object Browser. The view appears in the main window, and you can browse it in the same way as the Class View window.

This window provides the same information as Class View but also shows you more of the .NET types. When an item is selected, you also get information about it in a third window, as shown in Figure 9-4.




FIGURE 9-4

Here, the `.ReadKey()` method of the `Console` class has been selected. (`Console` is found in the `System` namespace in the `mscorlib` assembly.) The information window in the bottom-right corner shows you the method signature, the class to which the method belongs, and a summary of the method function. This information can be useful when you are exploring the .NET types, or if you are just refreshing your memory about what a particular class can do.

Additionally, you can make use of this information window in types that you create. Make the following change to the code created previously in Ch09Ex01:

```
/// <summary>
/// This class contains my program!
/// </summary>
class Program
{
    static void Main(string[] args)
    {
        MyComplexClass myObj = new MyComplexClass();
        WriteLine(myObj.ToString());
        ReadKey();
    }
}
```

Return to the Object Browser. The change is reflected in the information window. This is an example of XML documentation, a subject not covered in this book but well worth learning about when you have a spare moment.

**NOTE** If you made this code change manually, then you noticed that simply typing the three slashes (///) causes the IDE to add most of the rest of the code for you. It automatically analyzes the code to which you are applying XML documentation and builds the basic XML documentation — more evidence, should you need any, that Visual Studio is a great tool to work with!

## Adding Classes

Visual Studio contains tools that can speed up some common tasks, and some of these are applicable to OOP. One of these tools, the Add New Item Wizard, enables you to add new classes to your project with a minimum amount of typing.

This tool is accessible through the Project  $\Rightarrow$  Add New Item menu item or by right-clicking on your project in the Solution Explorer window and selecting the appropriate item. Either way, a dialog box appears, enabling you to choose the item to add. To add a class, select the Class item in the templates window, as shown in Figure 9-5, provide a filename for the file that will contain the class, and click Add. The class created is named according to the filename you provided.




FIGURE 9-5

In the Try It Out earlier in this chapter, you added class definitions manually to your `Program.cs` file. Often, keeping classes in separate files makes it easier to keep track of your classes. Entering the information in the Add New Item dialog box when the Ch09Ex01 project is open results in the following code being generated in `MyNewClass.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch09Ex01
{
    class MyNewClass
    {
    }
}
```

This class, `MyNewClass`, is defined in the same namespace as your entry point class, `Program`, so you can use it from code just as if it were defined in the same file. As shown in the code, the class generated for you contains no constructor. Recall that if a class definition doesn't include a constructor, then the compiler adds a default constructor when you compile your code.

## Class Diagrams

One powerful feature of Visual Studio that you haven't looked at yet is the capability to generate class diagrams from code and use them to modify projects. The class diagram editor in Visual Studio enables you to generate UML-like diagrams of your code with ease. You'll see this in action in the following Try It Out when you generate a class diagram for the Ch09Ex01 project you created earlier.

## TRY IT OUT Generating a Class Diagram

1. Open the Ch09Ex01 project created earlier in this chapter.
2. In the Solution Explorer window, right-click the Ch09Ex01 project and then select View  $\Rightarrow$  View Class Diagram menu item.
3. A class diagram appears, called ClassDiagram1.cd.
4. Click the IMyInterface lollipop and, using the Properties window, change its Position property to Right.
5. Right-click MyBase and select Show Base Type from the context menu.
6. Move the objects in the drawing around by dragging them to achieve a more pleasing layout. At this point, the diagram should look a little like Figure 9-6.




FIGURE 9-6

## How It Works

With very little effort, you have created a class diagram not unlike the UML diagram presented in Figure 9-2 (without the color, of course). The following features are evident:

- Classes are shown as blue boxes, including their name and type.
- Interfaces are shown as green boxes, including their name and type.
- Inheritance is shown with arrows with white heads (and in some cases, text inside class boxes).

- Classes implementing interfaces have lollipops.
- Abstract classes are shown with a dotted outline and italicized name.
- Sealed classes are shown with a thick black outline.

Clicking on an object shows you additional information in a Class Details window at the bottom of the screen (right-click an object and select Class Details if this window doesn't appear). Here, you can see (and modify) class members. You can also modify class details in the Properties window.

From the Toolbox, you can add new items such as classes, interfaces, and enums to the diagram, and define relationships between objects in the diagram. When you do this, the code for the new items is automatically generated for you.

---

## CLASS LIBRARY PROJECTS

As well as placing classes in separate files within your project, you can also place them in completely separate projects. A project that contains nothing but classes (along with other relevant type definitions, but no entry point) is called a *class library*.

Class library projects compile into .dll assemblies, and you can access their contents by adding references to them from other projects (which might be part of the same solution, but don't have to be). This extends the encapsulation that objects provide because class libraries can be revised and updated without touching the projects that use them. That means you can easily upgrade services provided by classes (which might affect multiple consumer applications).

The following Try It Out provides an example of a class library project and a separate project that makes use of the classes that it contains.

### TRY IT OUT Using a Class Library: Ch09ClassLib and Ch09Ex02\Program.cs

1. Create a new project of type Class Library called Ch09ClassLib and save it in the directory C:\BegVCSharp\Chapter09, as shown in Figure 9-7.
2. Rename the file `Class1.cs` to `MyExternalClass.cs` (by right-clicking on the file in the Solution Explorer window and selecting Rename). Click Yes on the dialog box that appears.
3. The code in `MyExternalClass.cs` automatically changes to reflect the class name change:

```
public class MyExternalClass
{
}
```

4. Add a new class to the project, using the filename `MyInternalClass.cs`.
5. Modify the code to make the class `MyInternalClass` explicitly internal:  

```
internal class MyInternalClass
{
}
```
6. Compile the project (this project has no entry point, so you can't run it as normal — instead, you can build it by selecting Build ⇔ Build Solution).




FIGURE 9-7

7. Create a new console application project called Ch09Ex02 and save it in the directory C:\BegVCSharp\Chapter09.
8. Select Project ➤ Add Reference, or select the same option after right-clicking References in the Solution Explorer window.
9. Click the Browse tab, navigate to C:\BegVCSharp\Chapter09\Chapter09\Ch09ClassLib\bin\Debug\, and double-click on Ch09ClassLib.dll.
10. When the operation completes, confirm that a reference was added in the Solution Explorer window, as shown in Figure 9-8.




FIGURE 9-8

11. Open the Object Browser window and examine the new reference to see what objects it contains (see Figure 9-9).




FIGURE 9-9

12. Modify the code in `Program.cs` as follows:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using static System.Console;
using Ch09ClassLib;
namespace Ch09Ex02
{
    class Program
    {
        static void Main(string[] args)
        {
            MyExternalClass myObj = new MyExternalClass();
            WriteLine(myObj.ToString());
            ReadKey();
        }
    }
}
```

13. Run the application. The result is shown in Figure 9-10.




FIGURE 9-10

### How It Works

This example created two projects: a class library project and a console application project. The class library project, Ch09ClassLib, contains two classes: `MyExternalClass`, which is publicly accessible, and

`MyInternalClass`, which is internally accessible. Note that this class was implicitly internal by default when you created it, as it had no access modifier. It is good practice to be explicit about accessibility, though, because it makes your code more readable, which is why you add the `internal` keyword. The console application project, Ch09Ex02, contains simple code that makes use of the class library project.

**NOTE** When an application uses classes defined in an external library, you can call that application a client application of the library. Code that uses a class that you define is often similarly referred to as client code.

To use the classes in Ch09ClassLib, you added a reference to `Ch09ClassLib.dll` to the console application. For the purposes of this example, you simply point at the output file for the class library, although it would be just as easy to copy this file to a location local to Ch09Ex02, enabling you to continue development of the class library without affecting the console application. To replace the old assembly version with the new one, simply copy the newly generated DLL file over the old one.

After adding the reference, you took a look at the available classes using the Object Browser. Because the `MyInternalClass` is internal, you can't see it in this display — it isn't accessible to external projects. However, `MyExternalClass` is accessible, and it's the one you use in the console application.

You could replace the code in the console application with code attempting to use the internal class as follows:

```
static void Main(string[] args)
{
    MyInternalClass myObj = new MyInternalClass();
    WriteLine(myObj.ToString());
    ReadKey();
}
```

If you attempt to compile this code, you receive the following compilation error:

```
'Ch09ClassLib.MyInternalClass'
    is inaccessible due to its protection level
```

This technique of making use of classes in external assemblies is key to programming with C# and the .NET Framework. It is, in fact, exactly what you are doing when you use any of the classes in the .NET Framework because they are treated the same way.

## INTERFACES VERSUS ABSTRACT CLASSES

This chapter has demonstrated how you can create both interfaces and abstract classes (without members for now — you get to them in Chapter 10). The two types are similar in a number of ways, so it would be useful to know how to determine when you should use one technique or the other.

First the similarities: Both abstract classes and interfaces can contain members that can be inherited by a derived class. Neither interfaces nor abstract classes can be directly instantiated, but it is possible to declare variables of these types. If you do, you can use polymorphism to assign objects that

inherit from these types to variables of these types. In both cases, you can then use the members of these types through these variables, although you don't have direct access to the other members of the derived object.

Now the differences: Derived classes can only inherit from a single base class, which means that only a single abstract class can be inherited directly (although it is possible for a chain of inheritance to include multiple abstract classes). Conversely, classes can use as many interfaces as they want, but this doesn't make a massive difference — similar results can be achieved either way. It's just that the interface way of doing things is slightly different.

Abstract classes can possess both *abstract members* (these have no code body and must be implemented in the derived class unless the derived class is itself abstract) and *non-abstract members* (these possess a code body, and can be virtual so that they can be overridden in the derived class). *Interface members*, conversely, must be implemented on the class that uses the interface — they do not possess code bodies. Moreover, interface members are by definition public (because they are intended for external use), but members of abstract classes can also be private (as long as they aren't abstract), protected, internal, or protected internal (where protected internal members are accessible only from code within the application or from a derived class). In addition, interfaces can't contain fields, constructors, destructors, static members, or constants.

**NOTE** Abstract classes are intended for use as the base class for families of objects that share certain central characteristics, such as a common purpose and structure. Interfaces are intended for use by classes that might differ on a far more fundamental level, but can still do some of the same things.

For example, consider a family of objects representing trains. The base class, `Train`, contains the core definition of a train, such as wheel gauge and engine type (which could be steam, diesel, and so on). However, this class is abstract because there is no such thing as a “generic” train. To create an “actual” train, you add characteristics specific to that train. For example, you derive classes such as `PassengerTrain`, `FreightTrain`, and `424DoubleBogey`, as shown in Figure 9-11.




FIGURE 9-11

A family of car objects might be defined in the same way, with an abstract base class of `Car` and derived classes such as `Compact`, `SUV`, and `PickUp`. `Car` and `Train` might even derive from a common base class, such as `Vehicle`. This is shown in Figure 9-12.




FIGURE 9-12

Some of the classes lower in the hierarchy can share characteristics because of their purpose, not just because of what they are derived from. For example, PassengerTrain, Compact, SUV, and Pickup are all capable of carrying passengers, so they might possess an `IPassengerCarrier` interface. FreightTrain and Pickup can carry heavy loads, so they might both have an `IHeavyLoadCarrier` interface as well. This is illustrated in Figure 9-13.




FIGURE 9-13

By breaking down an object system in this way before going about assigning specifics, you can clearly see which situations should use abstract classes rather than interfaces, and vice versa. The result of this example couldn't be achieved using only interfaces or only abstract inheritance.

## STRUCT TYPES

Chapter 8 noted that structs and classes are very similar but that structs are value types and classes are reference types. What does this actually mean to you? Well, the easiest way of looking at this is with an example, such as the following Try It Out.

### TRY IT OUT Classes versus Structs: Ch09Ex03\Program.cs

1. Create a new console application project called Ch09Ex03 and save it in the directory C:\BegVCSharp\Chapter09.
2. Modify the code as follows:

```
namespace Ch09Ex03
{
    class MyClass
    {
        public int val;
    }
    struct myStruct
    {
        public int val;
    }
    class Program
    {
        static void Main(string[] args)
        {
            MyClass objectA = new MyClass();
            MyClass objectB = objectA;
            objectA.val = 10;
            objectB.val = 20;
            myStruct structA = new myStruct();
            myStruct structB = structA;
            structA.val = 30;
            structB.val = 40;
            WriteLine("objectA.val = {objectA.val}");
            WriteLine("objectB.val = {objectB.val}");
            WriteLine("structA.val = {structA.val}");
            WriteLine("structB.val = {structB.val}");
            ReadKey();
        }
    }
}
```

3. Run the application. Figure 9-14 shows the output.




FIGURE 9-14

### How It Works

This application contains two type definitions: one for a struct called `myStruct`, which has a single public `int` field called `val`, and one for a class called `MyClass` that contains an identical field (you look at class members such as fields in Chapter 10; for now just understand that the syntax is the same here). Next, you perform the same operations on instances of both of these types:

1. Declare a variable of the type.
2. Create a new instance of the type in this variable.
3. Declare a second variable of the type.
4. Assign the first variable to the second variable.
5. Assign a value to the `val` field in the instance in the first variable.
6. Assign a value to the `val` field in the instance in the second variable.
7. Display the values of the `val` fields for both variables.

Although you are performing the same operations on variables of both types, the outcome is different. When you display the values of the `val` field, both object types have the same value, whereas the struct types have different values. What has happened?

Objects are *reference* types. When you assign an object to a variable you are actually assigning that variable with a *pointer* to the object to which it refers. A pointer, in real code terms, is an address in memory. In this case, the address is the point in memory where the object is found. When you assign the first object reference to the second variable of type `MyClass` with the following line, you are actually copying this address:

```
MyClass objectB = objectA;
```

This means that both variables contain pointers to the same object.

Structs are *value* types. Instead of the variable holding a pointer to the struct, the variable contains the struct itself. When you assign the first struct to the second variable of type `myStruct` with the following line, you are actually copying all the information from one struct to the other:

```
myStruct structB = structA;
```

You saw behavior like this earlier in this book for simple variable types such as `int`. The upshot is that the two struct type variables contain different structs. The entire technique of using pointers is hidden from you in managed C# code, making your code much simpler. It is possible to access lower-level operations such as pointer manipulation in C# using unsafe code, but that is an advanced topic not covered here.

---

## SHALLOW COPYING VERSUS DEEP COPYING

Copying objects from one variable to another by value instead of by reference (that is, copying them in the same way as structs) can be quite complex. Because a single object can contain references to many other objects, such as field members and so on, a lot of processing can be involved. Simply copying each member from one object to another might not work because some of these members might be reference types in their own right.

The .NET Framework takes this into account. You can create a simple copy of an object where each member is copied to the new object by using the method `MemberwiseClone()`, inherited from `System.Object`. This is a protected method, but it would be easy to define a public method on an object that called this method. This copying method is known as a *shallow copy*, in that it doesn't take reference type members into account. This means that reference members in the new object refer to the same objects as equivalent members in the source object, which isn't ideal in many cases. If you want to create new instances of the members in question by copying the values across (rather than the references), you need to perform a *deep copy*.

There is an interface you can implement that enables you to deep copy in a standard way: `ICloneable`. If you use this interface, then you must implement the single method it contains, `Clone()`. This method returns a value of type `System.Object`. You can use whatever processing you want to obtain this object, by implementing the method body however you choose. That means you can implement a deep copy if you want to, although the exact behavior isn't mandatory, so you could perform a shallow copy if desired. There are no rules or restrictions on what you actually return from this method, so many people recommend avoiding it. Instead, they recommend implementing your own deep-copy method. You take a closer look at this interface in Chapter 11.

## EXERCISES

### 9.1 What is wrong with the following code?

```
public sealed class MyClass
{
    // Class members.
}
public class myDerivedClass : MyClass
{
    // Class members.
}
```

---

- 9.2** How would you define a non-creatable class?
- 
- 9.3** Why are non-creatable classes still useful? How do you make use of their capabilities?
- 
- 9.4** Write code in a class library project called Vehicles that implements the `Vehicle` family of objects discussed earlier in this chapter. There are nine objects and two interfaces that require implementation.
- 
- 9.5** Create a console application project, Traffic, that references `Vehicles.dll` (created in Question 4). Include a function called `AddPassenger` that accepts any object with the `IPassengerCarrier` interface. To prove that the code works, call this function using instances of each object that supports this interface, calling the `ToString` method inherited from `System.Object` on each one and writing the result to the screen.
- 

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>Class and interface definitions</b>	Classes are defined with the <code>class</code> keyword, and interfaces with the <code>interface</code> keyword. You can use the <code>public</code> and <code>internal</code> keywords to define class and interface accessibility, and classes can be defined as <code>abstract</code> or <code>sealed</code> to control inheritance. Parent classes and interfaces are specified in a comma-separated list after a colon following the class or interface name. Only a single parent class can be specified in a class definition, and it must be the first item in the list.
<b>Constructors and destructors</b>	Classes come ready-equipped with a default constructor and destructor implementation, and you rarely have to provide your own destructor. You can define constructors with an accessibility, the name of the class, and any required parameters. Constructors of base classes are executed before those of derived classes, and you can control the execution sequence within a class with the <code>this</code> and <code>base</code> constructor initializer keywords.
<b>Class libraries</b>	You can create class library projects that only contain class definitions. These projects cannot be executed directly; they must be accessed through client code in an executable application. Visual Studio provides various tools for creating, modifying, and examining classes.
<b>Class families</b>	Classes can be grouped into families that exhibit common behavior or that share common characteristics. You can do this by inheriting from a shared base class (which can be abstract), or by implementing interfaces.
<b>Struct definitions</b>	A struct is defined in a very similar way to a class, but remember that structs are value types whereas classes are reference types.
<b>Copying objects</b>	When you make a copy of an object, you must be careful to copy any objects that it might contain, rather than simply copying the references to those objects. Copying references is referred to as shallow copying, whereas a full copy is referred to as a deep copy. You can use the <code>ICloneable</code> interface as a framework for providing deep-copy capabilities in a class definition.

# 10

## Defining Class Members

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Defining class members
- Controlling class member inheritance
- Defining nested classes
- Implementing interfaces
- Using partial class definitions
- Using the Call Hierarchy window

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 10 download and individually named according to the names throughout the chapter.

This chapter continues exploring class definitions in C# by looking at how you define field, property, and method class members. You start by examining the code required for each of these types, and learn how to generate the structure of this code. You also learn how to modify members quickly by editing their properties.

After covering the basics of member definition, you'll learn some advanced techniques involving members: hiding base class members, calling overridden base class members, nested type definitions, and partial class definitions.

Finally, you put theory into practice by creating a class library that you can build on and use in later chapters.

## MEMBER DEFINITIONS

Within a class definition, you provide definitions for all members of the class, including fields, methods, and properties. All members have their own accessibility levels, defined in all cases by one of the following keywords:

- `public` — Members are accessible from any code.
- `private` — Members are accessible only from code that is part of the class (the default if no keyword is used).
- `internal` — Members are accessible only from code within the assembly (project) where they are defined.
- `protected` — Members are accessible only from code that is part of either the class or a derived class.

The last two of these can be combined, so `protected internal` members are also possible. These are only accessible from code-derived classes within the project (more accurately, the assembly).

Fields, methods, and properties can also be declared using the keyword `static`, which means that they are static members owned by the class, rather than by object instances, as discussed in Chapter 8.

## Defining Fields

Fields are defined using standard variable declaration format (with optional initialization), along with the modifiers discussed previously:

```
class MyClass
{
    public int MyInt;
}
```

**NOTE** Public fields in the .NET Framework are named using PascalCasing, rather than camelCasing, and that's the casing methodology used here. That's why the field in this example is called `MyInt` instead of `myInt`. This is only a suggested casing scheme, but it makes a lot of sense. There is no recommendation for private fields, which are usually named using camelCasing.

Fields can also use the keyword `readonly`, meaning the field can be assigned a value only during constructor execution or by initial assignment:

```
class MyClass
{
    public readonly int MyInt = 17;
}
```

As noted in the chapter introduction, fields can be declared as static using the `static` keyword:

```
class MyClass
{
    public static int MyInt;
}
```

Static fields are accessed via the class that defines them (`MyClass.MyInt` in the preceding example), not through object instances of that class. You can use the keyword `const` to create a constant value. `const` members are static by definition, so you don't need to use the `static` modifier (in fact, it is an error to do so).

## Defining Methods

Methods use standard function format, along with accessibility and optional `static` modifiers, as shown in this example:

```
class MyClass
{
    public string GetString() => return "Here is a string.";
}
```

**NOTE** Like public fields, public methods in the .NET Framework are named using PascalCasing.

Remember that if you use the `static` keyword, then this method is accessible only through the class, not the object instance. You can also use the following keywords with method definitions:

- `virtual` — The method can be overridden.
- `abstract` — The method must be overridden in non-abstract derived classes (only permitted in abstract classes).
- `override` — The method overrides a base class method (it must be used if a method is being overridden).
- `extern` — The method definition is found elsewhere.

Here's an example of a method override:

```
public class MyBaseClass
{
    public virtual void DoSomething()
    {
        // Base implementation.
    }
}
public class MyDerivedClass : MyBaseClass
{
    public override void DoSomething()
    {
        // Derived class implementation, overrides base implementation.
    }
}
```

If `override` is used, then `sealed` can also be used to specify that no further modifications can be made to this method in derived classes — that is, the method can't be overridden by derived classes. Here is an example:

```
public class MyDerivedClass : MyBaseClass
{
    public override sealed void DoSomething()
    {
        // Derived class implementation, overrides base implementation.
    }
}
```

Using `extern` enables you to provide the implementation of a method externally to the project, but this is an advanced topic not covered here.

## Defining Properties

Properties are defined in a similar way to fields, but there's more to them. Properties, as already discussed, are more involved than fields in that they can perform additional processing before modifying state — and, indeed, might not modify state at all. They achieve this by possessing two function-like blocks: one for getting the value of the property and one for setting the value of the property.

These blocks, also known as *accessors*, are defined using `get` and `set` keywords respectively, and can be used to control the access level of the property. You can omit one or the other of these blocks to create read-only or write-only properties (where omitting the `get` block gives you write-only access, and omitting the `set` block gives you read-only access). Of course, that only applies to external code because code elsewhere within the class will have access to the same data that these code blocks have. You can also include accessibility modifiers on accessors — making a `get` block `public` while the `set` block is `protected`, for example. You must include at least one of these blocks to obtain a valid property (and, let's face it, a property you can't read or change wouldn't be very useful).

The basic structure of a property consists of the standard access modifying keyword (`public`, `private`, and so on), followed by a type name, the property name, and one or both of the `get` and `set` blocks that contain the property processing:

```
public int MyIntProp
{
    get
    {
        // Property get code.
    }
    set
    {
        // Property set code.
    }
}
```

**NOTE** Public properties in .NET are also named using *PascalCasing*, rather than *camelCasing*; as with fields and methods, *PascalCasing* is used here.

The first line of the definition is the bit that is very similar to a field definition. The difference is that there is no semicolon at the end of the line; instead, you have a code block containing nested `get` and `set` blocks.

`get` blocks must have a return value of the type of the property. Simple properties are often associated with a single private field controlling access to that field, in which case the `get` block can return the field's value directly:

```
// Field used by property.
private int myInt;
// Property.
public int MyIntProp
{
    get { return myInt; }
    set { // Property set code. }
}
```

Code external to the class cannot access this `myInt` field directly due to its accessibility level (it is private). Instead, external code must use the property to access the field. The `set` function assigns a value to the field similarly. Here, you can use the keyword `value` to refer to the value received from the user of the property:

```
// Field used by property.
private int myInt;
// Property.
public int MyIntProp
{
    get { return myInt; }
    set { myInt = value; }
}
```

`value` equates to a value of the same type as the property, so if the property uses the same type as the field, then you never have to worry about casting in situations like this.

This simple property does little more than shield direct access to the `myInt` field. The real power of properties is apparent when you exert a little more control over the proceedings. For example, you might implement your `set` block as follows:

```
set
{
    if (value >= 0 && value <= 10)
        myInt = value;
}
```

Here, you modify `myInt` only if the value assigned to the property is between 0 and 10. In situations like this, you have an important design choice to make. What should you do if an invalid value is used? You have four options:

- Do nothing (as in the preceding code).
- Assign a default value to the field.
- Continue as if nothing went wrong but log the event for future analysis.
- Throw an exception.

In general, the last two options are preferable. Deciding between them depends on how the class will be used and how much control should be assigned to the users of the class. Exception throwing gives users a fair amount of control and lets them know what is going on so that they can respond appropriately. You can use one of the standard exceptions in the `System` namespace for this:

```
set
{
    if (value >= 0 && value <= 10)
        myInt = value;
    else
        throw (new ArgumentException("MyIntProp", value,
            "MyIntProp must be assigned a value between 0 and 10."));
}
```

This can be handled using `try...catch...finally` logic in the code that uses the property, as you saw in Chapter 7.

Logging data, perhaps to a text file, can be useful, such as in production code where problems really shouldn't occur. It enables developers to check on performance and perhaps debug existing code if necessary.

Properties can use the `virtual`, `override`, and `abstract` keywords just like methods, something that isn't possible with fields. Finally, as mentioned earlier, accessors can have their own accessibilities, as shown here:

```
// Field used by property.
private int myInt;
// Property.
public int MyIntProp
{
    get { return myInt; }
    protected set { myInt = value; }
}
```

Here, only code within the class or derived classes can use the `set` accessor.

The accessibilities that are permitted for accessors depend on the accessibility of the property, and it is forbidden to make an accessor more accessible than the property to which it belongs. This means that a `private` property cannot contain any accessibility modifiers for its accessors, whereas `public` properties can use all modifiers on their accessors.

C# 6 introduced a feature called expression based properties. Similar to the expression based method discussed previously in Chapter 6, this feature reduces the extent of the property to a single line of code. For example, properties that return a one-line mathematical computation on a value can use the lambda arrow followed by the equation.

```
// Field used by property.
private int myDoubledInt = 5;
// Property.
public int MyDoubledIntProp => (myDoubledInt * 2);
```

The following Try It Out enables you to experiment with defining and using fields, methods, and properties.

**TRY IT OUT** Using Fields, Methods, and Properties: Ch10Ex01

1. Create a new console application called Ch10Ex01 and save it in the directory C:\BegVCSharp\Chapter10.
2. Add a new class called MyClass, using the Add Class shortcut, which will cause the new class to be defined in a new file called MyClass.cs.
3. Modify the code in MyClass.cs as follows:

```
public class MyClass
{
    public readonly string Name;
    private int intValue;
    public int Val
    {
        get { return intValue; }
        set {
            if (value >= 0 && value <= 10)
                intValue = value;
            else
                throw (new ArgumentOutOfRangeException("Val", value,
                    "Val must be assigned a value between 0 and 10."));
        }
    }
    public override string ToString() => "Name: " + Name + "\nVal: " + Val;
}
private MyClass() : this("Default Name") { }
public MyClass(string newName)
{
    Name = newName;
    intValue = 0;
}
private int myDoubledInt;
public int myDoubledIntProp => (myDoubledInt * 2);

}
```

4. Modify the code in Program.cs as follows:

```
using static System.Console;
static void Main(string[] args)
{
    WriteLine("Creating object myObj...");
    MyClass myObj = new MyClass("My Object");
    WriteLine("myObj created.");
    for (int i = -1; i <= 0; i++)
    {
        try
        {
            WriteLine($"\\Attempting to assign {i} to myObj.Val...");
            myObj.Val = i;
            WriteLine($"Value {myObj.Val} assigned to myObj.Val.");
        }
        catch (Exception e)
```

```

    {
        WriteLine($"Exception {e.GetType().FullName} thrown.");
        WriteLine($"Message:\n\"{e.Message}\"");
    }
}
WriteLine("\nOutputting myObj.ToString()....");
WriteLine(myObj.ToString());
WriteLine("myObj.ToString() Output.");

WriteLine("\nmyDoubledIntProp = 5....");
WriteLine($"Getting myDoubledIntProp of 5 is {myObj.myDoubledIntProp}");

.ReadKey();
}

```

5. Run the application. The result is shown in Figure 10-1.




FIGURE 10-1

## How It Works

The code in `Main()` creates and uses an instance of the `MyClass` class defined in `MyClass.cs`. The code must instantiate this class by using a nondefault constructor because the default constructor of `MyClass` is private:

```
private MyClass() : this("Default Name") {}
```

Using `this("Default Name")` ensures that `Name` gets a value if this constructor is ever called, which is possible if this class is used to derive a new class. This is necessary because not assigning a value to the `Name` field could be a source of errors later.

The nondefault constructor used assigns values to the `readonly` field `Name` (you can only do this by assignment in the field declaration or in a constructor) and the private field `intval`.

Next, `Main()` attempts two assignments to the `val` property of `myObj` (the instance of `MyClass`). A `for` loop is used to assign the values `-1` and `0` in two cycles, and a `try...catch` structure is used to check for any exception thrown. When `-1` is assigned to the property, an exception of type `System.ArgumentOutOfRangeException` is thrown, and code in the `catch` block outputs information about

the exception to the console window. In the next loop cycle, the value 0 is successfully assigned to the `val` property, and through that property to the private `intval` field.

Use the overridden `ToString()` method to output a formatted string representing the contents of the object:

```
public override string ToString() => "Name: " + Name + "\nVal: " + Val;
```

This method must be declared using the `override` keyword, because it is overriding the virtual `ToString()` method of the base `System.Object` class. The code here uses the property `val` directly, rather than the private field `intval`. There is no reason why you shouldn't use properties from within classes in this way, although there may be a small performance hit (so small that you are unlikely to notice it). Of course, using the property also gives you the validation inherent in property use, which may be beneficial for code within the class as well.

Finally, you created and set a read-only property called `myDoubledInt` in `MyClass.cs` to 5. By using the expression based property feature to return the value multiplied by 2:

```
public int MyDoubledIntProp => (myDoubledInt * 2);
```

when property is accessed using `myObj.myDoubledIntProp` the output is 2 times 5 which is 10, as expected.

## Refactoring Members

One technique that comes in handy when adding properties is the capability to generate a property from a field. This is an example of *refactoring*, which simply means modifying your code using a tool, rather than by hand. This can be accomplished by right-clicking a member in a class diagram or in code view.

For example, if the `MyClass` class contained the field,

```
public string myString;
```

you could right-click on the field and select Quick Actions.... That would bring up the dialog box shown in Figure 10-2.




FIGURE 10-2

Accepting the default options modifies the code for `MyClass` as follows:

```
public string myString;
public string MyString
{
    get { return myString; }
    set { myString = value; }
}
private string myString;
```

Here, the accessibility of the `myString` field has been changed to `private`, and a public property called `MyString` has been created and automatically linked to `myString`. Clearly, reducing the time required to monotonously create properties for fields is a big plus!

## Automatic Properties

Properties are the preferred way to access the state of an object because they shield external code from the implementation of data storage within the object. They also give you greater control over how internal data is accessed, as you have seen several times in this chapter's code. However, you'll typically define properties in a very standard way — that is, you will have a private member that is accessed directly through a public property. The code for this is almost invariably similar to the code in the previous section, which was autogenerated by the Visual Studio refactoring tool.

Refactoring certainly speeds things up when it comes to typing, but C# has another trick up its sleeve: automatic properties. With an *automatic property*, you declare a property with a simplified syntax and the C# compiler fills in the blanks for you. Specifically, the compiler declares a private field that is used for storage, and uses that field in the `get` and `set` blocks of your property — without you having to worry about the details.

Use the following code structure to define an automatic property:

```
public int MyIntProp
{
    get;
    set;
}
```

You can even define an automatic property on a single line of code to save space, without making the property much less readable:

```
public int MyIntProp { get; set; }
```

You define the accessibility, type, and name of the property in the usual way, but you don't provide any implementation for the `get` or `set` block. Instead, the compiler provides the implementations of these blocks (and the underlying field).

**TIP** You can create an automatically implemented property template by using the `prop` code snippet within Visual Studio. Type in "prop" then press the TAB key twice and the following, `public int MyProperty {get; set;}` is created for you.

When you use an automatic property, you only have access to its data through the property, not through its underlying private field. This is because you can't access the private field without knowing its name, which is defined during compilation. However, that's not really a limitation because using the property name directly is fine. The only limitation of automatic properties is that they must include both a `get` and a `set` accessor — you cannot define read- or write-only properties in this way. However, you can change the accessibility of these accessors. For example, this means you can create an externally read-only property as follows:

```
public int MyIntProp { get; private set; }
```

Here you can access the value of `MyIntProp` only from code in the class definition.

C# 6 introduced two new concepts pertaining to automatic properties referred to as getter-only auto-properties and initializers for auto-properties. Prior to C# 6, automatic properties required setters, which limited the utilization of immutable data types. The simple definition of an immutable data type is that it does not change state once it is created, the most famous immutable type being `System.String`. There are many benefits for using immutable data types, such as the simplification of concurrent programming and the synchronization of threads.

Concurrent programming and synchronization of threads are advanced topics and not discussed further in this book; however, it is important to know about the getter-only auto-properties. They are created by using the following syntax, notice that a setter is no longer required:

```
public int MyIntProp { get; }
```

The initialization feature for auto-properties is implemented by the following which is similar to the way fields are declared:

```
public int MyIntProp { get; } = 9;
```

## ADDITIONAL CLASS MEMBER TOPICS

Now you're ready to look at some more advanced member topics. This section tackles the following:

- Hiding base class methods
- Calling overridden or hidden base class methods
- Using nested type definitions

### Hiding Base Class Methods

When you inherit a (non-abstract) member from a base class, you also inherit an implementation. If the inherited member is virtual, then you can override this implementation with the `override` keyword. Regardless of whether the inherited member is virtual, you can, if you want, *hide* the implementation. This is useful when, for example, a public inherited member doesn't work quite as you want it to.

You can do this simply by using code such as the following:

```
public class MyBaseClass
{
    public void DoSomething()
    {
        // Base implementation.
    }
}
public class MyDerivedClass : MyBaseClass
{
    public void DoSomething()
    {
        // Derived class implementation, hides base implementation.
    }
}
```

Although this code works fine, it generates a warning that you are hiding a base class member. That warning gives you the chance to correct it if you have accidentally hidden a member that you want to use. If you really do want to hide the member, you can use the `new` keyword to explicitly indicate that this is what you want to do:

```
public class MyDerivedClass : MyBaseClass
{
    new public void DoSomething()
    {
        // Derived class implementation, hides base implementation.
    }
}
```

This works in exactly the same way but won't show a warning. At this point, it's worthwhile to note the difference between hiding and overriding base class members. Consider the following code:

```
public class MyBaseClass
{
    public virtual void DoSomething() => WriteLine("Base imp");
}
public class MyDerivedClass : MyBaseClass
{
    public override void DoSomething() => WriteLine("Derived imp");
}
```

Here, the overriding method replaces the implementation in the base class, such that the following code uses the new version even though it does so through the base class type (using polymorphism):

```
MyDerivedClass myObj = new MyDerivedClass();
MyBaseClass myBaseObj;
myBaseObj = myObj;
myBaseObj.DoSomething();
```

This results in the following output:

```
Derived imp
```

Alternatively, you could hide the base class method:

```
public class MyBaseClass
{
    public virtual void DoSomething() => WriteLine("Base imp");
}
public class MyDerivedClass : MyBaseClass
{
    new public void DoSomething() => WriteLine("Derived imp");
}
```

The base class method needn't be virtual for this to work, but the effect is exactly the same and the preceding code only requires changes to one line. The result for a virtual or nonvirtual base class method is as follows:

```
Base imp
```

Although the base implementation is hidden, you still have access to it through the base class.

## Calling Overridden or Hidden Base Class Methods

Whether you override or hide a member, you still have access to the base class member from the derived class. There are many situations in which this can be useful, such as the following:

- When you want to hide an inherited public member from users of a derived class but still want access to its functionality from within the class
- When you want to add to the implementation of an inherited virtual member rather than simply replace it with a new overridden implementation

To achieve this, you use the `base` keyword, which refers to the implementation of the base class contained within a derived class (in a similar way to its use in controlling constructors, as shown in the last chapter):

```
public class MyBaseClass
{
    public virtual void DoSomething()
    {
        // Base implementation.
    }
}
public class MyDerivedClass : MyBaseClass
{
    public override void DoSomething()
    {
        // Derived class implementation, extends base class implementation.
        base.DoSomething();
        // More derived class implementation.
    }
}
```

This code executes the version of `DoSomething()` contained in `MyBaseClass`, the base class of `MyDerivedClass`, from within the version of `DoSomething()` contained in `MyDerivedClass`. As `base` works using object instances, it is an error to use it from within a static member.

## The this Keyword

As well as using `base` in the last chapter, you also used the `this` keyword. As with `base`, `this` can be used from within class members, and, like `base`, `this` refers to an object instance, although it is the current object instance (which means you can't use this keyword in static members because static members are not part of an object instance).

The most useful function of the `this` keyword is the capability to pass a reference to the current object instance to a method, as shown in this example:

```
public void doSomething()
{
    MyTargetClass myObj = new MyTargetClass();
    myObj.DoSomethingWith(this);
}
```

Here, the `MyTargetClass` instance that is instantiated (`myObj`) has a method called `DoSomethingWith()`, which takes a single parameter of a type compatible with the class containing the preceding method. This parameter type might be of this class type, a class type from which this class derives, an interface implemented by the class, or (of course) `System.Object`.

Another common use of the `this` keyword is to use it to qualify local type members, for example:

```
public class MyClass
{
    private int someData;
    public int SomeData
    {
        get
        {
            return this.someData;
        }
    }
}
```

Many developers like this syntax, which can be used with any member type, as it is clear at a glance that you are referring to a member rather than a local variable.

## Using Nested Type Definitions

You can define types such as classes in namespaces, and you can also define them inside other classes. Then you can use the full range of accessibility modifiers for the definition, rather than just `public` and `internal`, and you can use the `new` keyword to hide a type definition inherited from a base class. For example, the following code defining `MyClass` also defines a nested class called `MyNestedClass`:

```
public class MyClass
{
    public class MyNestedClass
    {
        public int NestedClassField;
    }
}
```

To instantiate `MyNestedClass` from outside `MyClass`, you must qualify the name, as shown here:

```
MyClass.MyNestedClass myObj = new MyClass.MyNestedClass();
```

However, you might not be able to do this, for example if the nested class is declared as `private`. One reason for the existence of this feature is to define classes that are private to the containing class so that no other code in the namespace has access to them. Another reason is that nested classes have access to private and protected members of their containing class. The next Try it Out examines this feature.

### TRY IT OUT Using Nested Classes: Ch10Ex02

1. Create a new console application called `Ch10Ex02` and save it in the directory `C:\BegVCSharp\Chapter10`.
2. Modify the code in `Program.cs` as follows:

```
namespace Ch10Ex02
{
    public class ClassA
    {
        private int state = -1;
        public int State
        {
            get { return state; }
        }
        public class ClassB
        {
            public void SetPrivateState(ClassA target, int newState)
            {
                target.state = newState;
            }
        }
    }
    class Program
    {
        static void Main(string[] args)
        {
            ClassA myObject = new ClassA();
            WriteLine($"myObject.State = {myObject.State}");
            ClassA.ClassB myOtherObject = new ClassA.ClassB();
            myOtherObject.SetPrivateState(myObject, 999);
            WriteLine($"myObject.State = {myObject.State}");
            ReadKey();
        }
    }
}
```

3. Run the application. The result is shown in Figure 10-3.




FIGURE 10-3

### How It Works

The code in `Main()` creates and uses an instance of `ClassA`, which has a read-only property called `state`. Next, the code creates an instance of the nested class `ClassA.ClassB`. This class has access to the backing field for `ClassA.State`, which is the `ClassA.state` field, even though the field is private. Because of this, the nested class method `SetPrivateState()` can change the value of the read-only `state` property of `ClassA`.

It is important to reiterate that this is possible only because `ClassB` is defined as a nested class of `ClassA`. If you were to move the definition of `ClassB` outside of `ClassA`, then the code wouldn't compile due to this error:

```
'Ch10Ex02.ClassA.state' is inaccessible due to its protection level.
```

Being able to expose the internal state of your classes to nested classes can be extremely useful in some circumstances. However, most of the time it's enough simply to manipulate the internal state through methods that your class exposes.

---

## INTERFACE IMPLEMENTATION

This section takes a closer look at how you go about defining and implementing interfaces. In the last chapter, you learned that interfaces are defined in a similar way as classes, using code such as the following:

```
interface IMyInterface
{
    // Interface members.
}
```

Interface members are defined like class members except for a few important differences:

- No access modifiers (`public`, `private`, `protected`, or `internal`) are allowed — all interface members are implicitly public.
- Interface members can't contain code bodies.
- Interfaces can't define field members.

- Interface members can't be defined using the keywords `static`, `virtual`, `abstract`, or `sealed`.
- Type definition members are forbidden.

You can, however, define members using the `new` keyword if you want to hide members inherited from base interfaces:

```
interface IMyBaseInterface
{
    void DoSomething();
}
interface IMyDerivedInterface : IMyBaseInterface
{
    new void DoSomething();
}
```

This works exactly the same way as hiding inherited class members.

Properties defined in interfaces define either or both of the access blocks — `get` and `set` — which are permitted for the property, as shown here:

```
interface IMyInterface
{
    int MyInt { get; set; }
}
```

Here the `int` property `MyInt` has both `get` and `set` accessors. Either of these can be omitted for a property with more restricted access.

**NOTE** This syntax is similar to automatic properties, but remember that automatic properties are defined for classes, not interfaces, and that automatic properties must have both `get` and `set` accessors.

Interfaces do not specify how the property data should be stored. Interfaces cannot specify fields, for example, that might be used to store property data. Finally, interfaces, like classes, can be defined as members of classes (but not as members of other interfaces because interfaces cannot contain type definitions).

## Implementing Interfaces in Classes

A class that implements an interface must contain implementations for all members of that interface, which must match the signatures specified (including matching the specified `get` and `set` blocks), and must be public, as shown here:

```
public interface IMyInterface
{
    void DoSomething();
    void DoSomethingElse();
}
```

```
public class MyClass : IMyInterface
{
    public void DoSomething() {}
    public void DoSomethingElse() {}
}
```

It is possible to implement interface members using the keyword `virtual` or `abstract`, but not `static` or `const`. Interface members can also be implemented on base classes:

```
public interface IMyInterface
{
    void DoSomething();
    void DoSomethingElse();
}
public class MyBaseClass
{
    public void DoSomething() {}
}
public class MyDerivedClass : MyBaseClass, IMyInterface
{
    public void DoSomethingElse() {}
}
```

Inheriting from a base class that implements a given interface means that the interface is implicitly supported by the derived class. Here's an example:

```
public interface IMyInterface
{
    void DoSomething();
    void DoSomethingElse();
}
public class MyBaseClass : IMyInterface
{
    public virtual void DoSomething() {}
    public virtual void DoSomethingElse() {}
}
public class MyDerivedClass : MyBaseClass
{
    public override void DoSomething() {}
}
```

Clearly, it is useful to define implementations in base classes as `virtual` so that derived classes can replace the implementation, rather than hide it. If you were to hide a base class member using the `new` keyword, rather than `override` it in this way, the method `IMyInterface.DoSomething()` would always refer to the base class version even if the derived class were being accessed via the interface.

## Explicit Interface Member Implementation

Interface members can also be implemented *explicitly* by a class. If you do that, the member can only be accessed through the interface, not the class. *Implicit* members, which you used in the code in the last section, can be accessed either way.

For example, if the class `MyClass` implemented the `DoSomething()` method of `IMyInterface` implicitly, as in the preceding example, then the following code would be valid:

```
MyClass myObj = new MyClass();
myObj.DoSomething();
```

This would also be valid:

```
MyClass myObj = new MyClass();
IMyInterface myInt = myObj;
myInt.DoSomething();
```

Alternatively, if `MyDerivedClass` implements `DoSomething()` explicitly, then only the latter technique is permitted. The code for doing that is as follows:

```
public class MyClass : IMyInterface
{
    void IMyInterface.DoSomething() {}
    public void DoSomethingElse() {}
}
```

Here, `DoSomething()` is implemented explicitly, and `DoSomethingElse()` implicitly. Only the latter is accessible directly through an object instance of `MyClass`.

## Additional Property Accessors

Earlier you learned that if you implement an interface with a property, you must implement matching get/set accessors. That isn't strictly true — it is possible to add a get block to a property in a class in which the interface defining that property only contains a set block, and vice versa. However, this is possible only if you implement the interface implicitly. Also, in most cases you will want to add the accessor with an accessibility modifier that is more restrictive than the accessibility modifier on the accessor defined in the interface. Because the accessor defined by the interface is, by definition, public, this means that you would add nonpublic accessors. Here's an example:

```
public interface IMyInterface
{
    int MyIntProperty { get; }
}
public class MyBaseClass : IMyInterface
{
    public int MyIntProperty { get; protected set; }
}
```

If you define the additional accessor as public, then code with access to the class implementing the interface can access it. However, code that has access only to the interface won't be able to access it.

## PARTIAL CLASS DEFINITIONS

When you create classes with a lot of members of one type or another, the code can get quite confusing, and code files can get very long. One technique that can help, which you've looked at in earlier

chapters, is to use code outlining. By defining regions in code, you can collapse and expand sections to make the code easier to read. For example, you might have a class defined as follows:

```
public class MyClass
{
    #region Fields
    private int myInt;
    #endregion
    #region Constructor
    public MyClass() { myInt = 99; }
    #endregion
    #region Properties
    public int MyInt
    {
        get { return myInt; }
        set { myInt = value; }
    }
    #endregion
    #region Methods
    public void DoSomething()
    {
        // Do something..
    }
    #endregion
}
```

Here, you can expand and contract fields, properties, the constructor, and methods for the class, enabling you to focus only on what you are interested in. It is even possible to nest regions this way, so some regions are visible only when the region that contains them is expanded.

An alternative to using regions is to use *partial class definitions*. Put simply, you use partial class definitions to split the definition of a class across multiple files. You can, for example, put the fields, properties, and constructor in one file, and the methods in another. To do that, you just use the `partial` keyword with the class in each file that contains part of the definition, as follows:

```
public partial class MyClass { ... }
```

If you use partial class definitions, the `partial` keyword must appear in this position in every file containing part of the definition.

For example, a WPF window in a class called `MainWindow` has code stored in both `MainWindow.xaml.cs` and `MainWindow.g.i.cs` (visible if Show All Files is selected in the Solution Explorer window if you drill down into `obj\Debug` folder). This enables you to concentrate on the functionality of your forms, without worrying about your code being cluttered with information that doesn't really interest you.

One final note about partial classes: Interfaces applied to one partial class part apply to the whole class, meaning that the definition,

```
public partial class MyClass : IMyInterface1 { ... }
public partial class MyClass : IMyInterface2 { ... }
```

is equivalent to:

```
public class MyClass : IMyInterface1, IMyInterface2 { ... }
```

Partial class definitions can include a base class in a single partial class definition, or more than one partial class definition. If a base class is specified in more than one definition, though, it must be the *same* base class; recall that classes in C# can inherit only from a single base class.

## PARTIAL METHOD DEFINITIONS

Partial classes can also define partial methods. Partial methods are defined in one partial class definition without a method body, and implemented in another partial class definition. In both places, the `partial` keyword is used:

```
public partial class MyClass
{
    partial void MyPartialMethod();
}
public partial class MyClass
{
    partial void MyPartialMethod()
    {
        // Method implementation
    }
}
```

Partial methods can also be static, but they are always private and can't have a return value. Any parameters they use can't be `out` parameters, although they can be `ref` parameters. They also can't use the `virtual`, `abstract`, `override`, `new`, `sealed`, or `extern` modifiers.

Given these limitations, it is not immediately obvious what purpose partial methods fulfill. In fact, they are important when it comes to code compilation, rather than usage. Consider the following code:

```
public partial class MyClass
{
    partial void DoSomethingElse();
    public void DoSomething()
    {
        WriteLine("DoSomething() execution started.");
        DoSomethingElse();
        WriteLine("DoSomething() execution finished.");
    }
}
public partial class MyClass
{
    partial void DoSomethingElse() =>
        WriteLine("DoSomethingElse() called.");
}
```

Here, the partial method `DoSomethingElse()` is defined and called in the first partial class definition, and implemented in the second. The output, when `DoSomething()` is called from a console application, is what you might expect:

```
DoSomething() execution started.  
DoSomethingElse() called.  
DoSomething() execution finished.
```

If you were to remove the second partial class definition or partial method implementation entirely (or comment out the code), the output would be as follows:

```
DoSomething() execution started.  
DoSomething() execution finished.
```

You might assume that what is happening here is that when the call to `DoSomethingElse()` is made, the runtime discovers that the method has no implementation and therefore continues executing the next line of code. What actually happens is a little subtler. When you compile code that contains a partial method definition without an implementation, the compiler actually removes the method entirely. It also removes any calls to the method. When you execute the code, no check is made for an implementation because there is no call to check. This results in a slight — but nevertheless significant — improvement in performance.

As with partial classes, partial methods are useful when it comes to customizing autogenerated or designer-created code. The designer may declare partial methods that you can choose to implement or not depending on the situation. If you don't implement them, you incur no performance hit because effectively the method does not exist in the compiled code.

Consider at this point why partial methods can't have a return type. If you can answer that to your own satisfaction, you can be sure that you fully understand this topic — so that is left as an exercise for you.

## EXAMPLE APPLICATION

To illustrate some of the techniques you've been using so far, in this section you'll develop a class module that you can build on and make use of in subsequent chapters. The class module contains two classes:

- `Card` — Representing a standard playing card, with a suit of club, diamond, heart, or spade, and a rank that lies between ace and king
- `Deck` — Representing a full deck of 52 cards, with access to cards by position in the deck and the capability to shuffle the deck

You'll also develop a simple client to ensure that things are working, but you won't use the deck in a full card game application — yet.

## Planning the Application

The class library for this application, `Ch10CardLib`, will contain your classes. Before you get down to any code, though, you should plan the required structure and functionality of your classes.

## The Card Class

The `Card` class is basically a container for two read-only fields: `suit` and `rank`. The reason for making the fields read-only is that it doesn't make sense to have a "blank" card, and cards shouldn't be able to change once they have been created. To facilitate this, you'll make the default constructor private, and provide an alternative constructor that builds a card from a supplied suit and rank.

Other than that, the `Card` class will override the `ToString()` method of `System.Object`, so that you can easily obtain a human-readable string representing the card. To make things a little simpler, you'll provide enumerations for the two fields `suit` and `rank`.

The `Card` class is shown in Figure 10-4.




FIGURE 10-4

## The Deck Class

The `Deck` class will maintain 52 `Card` objects. You can use a simple array type for this. The array won't be directly accessible because access to the `Card` object is achieved through a `GetCard()` method, which returns the `Card` object with the given index. This class should also expose a `Shuffle()` method to rearrange the cards in the array. The `Deck` class is shown in Figure 10-5.

## Writing the Class Library

For the purposes of this example, it is assumed that you are familiar enough with the IDE to bypass the standard Try It Out format, so the steps aren't listed explicitly, as they are the same steps you've used many times. The important thing here is a detailed look at the code. Nonetheless, several pointers are included to ensure that you don't run into any problems along the way.

Both your classes and your enumerations will be contained in a class library project called `Ch10CardLib`. This project will contain four `.cs` files: `Card.cs`, which contains the `Card` class definition, `Deck.cs`, which contains the `Deck` class definition, and the `Suit.cs` and `Rank.cs` files containing enumerations.

You can put together a lot of this code using the Visual Studio class diagram tool.

**NOTE** If you'd prefer not to use the class diagram tool, don't worry. Each of the following sections also includes the code generated by the class diagram, so you'll be able to follow along just fine.

To get started, you need to do the following:

1. Create a new class library project called `Ch10CardLib` and save it in the directory `C:\BegVCSharp\Chapter10`.

2. Remove `Class1.cs` from the project.
3. Open the class diagram for the project using the Solution Explorer window (right-click the project and then click View  $\Rightarrow$  View Class Diagram). The class diagram should be blank to start with because the project contains no classes.

## Adding the Suit and Rank Enumerations

You can add an enumeration to the class diagram by dragging an `Enum` from the Toolbox into the diagram, and then filling in the New `Enum` dialog box that appears. For example, for the `Suit` enumeration, fill out the dialog box as shown in Figure 10-6.




FIGURE 10-6


Next, add the members of the enumeration using the Class Details window. Figure 10-7 shows the values that are required.

Name	Value	Summary	Hide
Club			<input checked="" type="checkbox"/>
Diamond			<input type="checkbox"/>
Heart			<input type="checkbox"/>
Spade			<input type="checkbox"/>
<add member>			

FIGURE 10-7

Add the `Rank` enumeration from the Toolbox in the same way. The values required are shown in Figure 10-8.

**NOTE** The value entry for the first member, `Ace`, is set to 1 so that the underlying storage of the `Enum` matches the rank of the card, such that `Six` is stored as 6, for example.

**FIGURE 10-8**

You can find the code generated for these two enumerations in the code files, `Suit.cs` and `Rank.cs`. First, you can find the full code for this example in `Ch10CardLib` folder/`Suit.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
namespace Ch10CardLib
{
    public enum Suit
    {
        Club,
        Diamond,
        Heart,
        Spade,
    }
}
```

And you can find the full code for this example in `Ch10CardLib` folder/`Rank.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
namespace Ch10CardLib
{
    public enum Rank
    {
        Ace = 1,
        Deuce,
        Three,
        Four,
        Five,
        Six,
        Seven,
        Eight,
```

Nine,  
Ten,  
Jack,  
Queen,  
King,  
}


Alternatively, you can add this code manually by adding `Suit.cs` and `Rank.cs` code files and then entering the code. Note that the extra commas added by the code generator after the last enumeration member do not prevent compilation and do not result in an additional “empty” member being created — although they are a little messy.

## Adding the Card Class

To add the `Card` class, you'll use a mix of the class designer and code editor. Adding a class in the class designer is much like adding an enumeration — you drag the appropriate entry from the Toolbox into the diagram. In this case, you drag a `Class` into the diagram and name the new class `Card`.

Use the Class Details window to add the fields rank and suit, and then use the Properties window to set the Constant Kind of the field to readonly. You also need to add two constructors — a private default constructor, and a public constructor that takes two parameters, newSuit and newRank, of types Suit and Rank, respectively. Finally, you override `ToString()`, which requires you to change the Inheritance Modifier in the Properties window to override.

Figure 10-9 shows the Class Details window and the Card class with all the information entered. (You can find this code in Ch10CardLib\Card.cs.)



**FIGURE 10-9**

Next, modify the code for the class in `Card.cs` as follows (or add the code shown to a new class called `Card` in the `ch10CardLib` namespace):

```

public class Card
{
    public readonly Suit suit;
    public readonly Rank rank;
    public Card(Suit newSuit, Rank newRank)
    {
        suit = newSuit;
        rank = newRank;
    }
    private Card()
    {
    }
    public override string ToString()
    {
        return "The " + rank + " of " + suit + "s";
    }
}

```

The overridden `ToString()` method writes the string representation of the enumeration value stored to the returned string, and the nondefault constructor initializes the values of the `suit` and `rank` fields.

## Adding the Deck Class

The `Deck` class needs the following members defined using the class diagram:

- A private field called `cards`, of type `Card[]`
- A public default constructor
- A public method called `GetCard()`, which takes one `int` parameter called `cardNum` and returns an object of type `Card`
- A public method called `Shuffle()`, which takes no parameters and returns `void`

When these are added, the Class Details window for the `Deck` class will appear as shown in Figure 10-10.




FIGURE 10-10

To make things clearer in the diagram, you can show the relationships among the members and types you have added. In the class diagram, right-click on each of the following in turn and select Show as Association from the menu:

- cards in Deck
- suit in Card
- rank in Card

When you have finished, the diagram should look like Figure 10-11.




FIGURE 10-11

Next, modify the code in `Deck.cs` (if you aren't using the class designer, you must add this class first with the code shown here). You can find this code in `Ch10CardLib\Deck.cs`. First you implement the constructor, which simply creates and assigns 52 cards in the `cards` field. You iterate through all combinations of the two enumerations, using each to create a card. This results in `cards` initially containing an ordered list of cards:

```

using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch10CardLib
{
    public class Deck
    {
        private Card[] cards;
        public Deck()
        {
            cards = new Card[52];
            for (int suitVal = 0; suitVal < 4; suitVal++)
                for (int rankVal = 0; rankVal < 13; rankVal++)
                    cards[suitVal * 13 + rankVal] = new Card((Rank)rankVal, (Suit)suitVal);
        }
    }
}
  
```

```

    {
        for (int rankVal = 1; rankVal < 14; rankVal++)
        {
            cards[suitVal * 13 + rankVal - 1] = new Card((Suit)suitVal,
                                              (Rank)rankVal);
        }
    }
}

```

Next, implement the `GetCard()` method, which either returns the `Card` object with the requested index or throws an exception as shown earlier:

```

public Card GetCard(int cardNum)
{
    if (cardNum >= 0 && cardNum <= 51)
        return cards[cardNum];
    else
        throw (new System.ArgumentOutOfRangeException("cardNum", cardNum,
                                                       "Value must be between 0 and 51."));
}

```

Finally, you implement the `Shuffle()` method. This method works by creating a temporary card array and copying cards from the existing `cards` array into this array at random. The main body of this function is a loop that counts from 0 to 51. On each cycle, you generate a random number between 0 and 51, using an instance of the `System.Random` class from the .NET Framework. Once instantiated, an object of this class generates a random number between 0 and x, using the method `Next(x)`. When you have a random number, you simply use that as the index of the `Card` object in your temporary array in which to copy a card from the `cards` array.

To keep a record of assigned cards, you also have an array of `bool` variables, and assign these to `true` as each card is copied. As you are generating random numbers, you check against this array to see whether you have already copied a card to the location in the temporary array specified by the random number. If so, you simply generate another.

This isn't the most efficient way of doing things because many random numbers will be generated before finding a vacant slot into which a card can be copied. However, it works, it's very simple, and C# code executes so quickly you will hardly notice a delay. The code is as follows:

```

public void Shuffle()
{
    Card[] newDeck = new Card[52];
    bool[] assigned = new bool[52];
    Random sourceGen = new Random();
    for (int i = 0; i < 52; i++)
    {
        int destCard = 0;
        bool foundCard = false;
        while (foundCard == false)
        {
            destCard = sourceGen.Next(52);
            if (assigned[destCard] == false)
                foundCard = true;
        }
    }
}

```

```
        assigned[destCard] = true;
        newDeck[destCard] = cards[i];
    }
    newDeck.CopyTo(cards, 0);
}
}
```


The last line of this method uses the `CopyTo()` method of the `System.Array` class (used whenever you create an array) to copy each of the cards in `newDeck` back into `cards`. This means you are using the same set of `Card` objects in the same `cards` object, rather than creating any new instances. If you had instead used `cards = newDeck`, then you would be replacing the object instance referred to by `cards` with another. This could cause problems if code elsewhere were retaining a reference to the original `cards` instance — which wouldn't be shuffled!

That completes the class library code.

# A Client Application for the Class Library

To keep things simple, you can add a client console application to the solution containing the class library. To do so, simply right-click on the solution in Solution Explorer and select Add ➔ New Project. The new project is called Ch10CardClient.

To use the class library you have created from this new console application project, add a reference to your Ch10CardLib class library project. You can do that through the Projects tab of the Reference Manager dialog box, as shown in Figure 10-12.



**FIGURE 10-12**

Select the project and click OK to add the reference.

Because this new project is the second one you've created, you also need to specify that it is the startup project for the solution, meaning the one that is executed when you click Run. To do so, simply right-click on the project name in the Solution Explorer window and select the Set as StartUp Project menu option.

Next, add the code that uses your new classes. That doesn't require anything particularly special, so the following code will do (you can find this code in Ch10CardClient\Program.cs):

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using static System.Console;
using Ch10CardLib;
namespace Ch10CardClient
{
    class Program
    {
        static void Main(string[] args)
        {
            Deck myDeck = new Deck();
            myDeck.Shuffle();
            for (int i = 0; i < 52; i++)
            {
                Card tempCard = myDeck.GetCard(i);
                Write(tempCard.ToString());
                if (i != 51)
                    Write(", ");
                else
                    WriteLine();
            }
            ReadKey();
        }
    }
}
```

Figure 10-13 shows the result you'll get if you run this application.




FIGURE 10-13

This is a random arrangement of the 52 playing cards in the deck. You'll continue to develop and use this class library in later chapters.

## THE CALL HIERARCHY WINDOW

Now is a good time to take a quick look at another feature of Visual Studio: the Call Hierarchy window. This window enables you to interrogate code to find out where your methods are called from and how they relate to other methods. The best way to illustrate this is with an example.

Open the example application from the previous section, and open the `Deck.cs` code file. Find the `Shuffle()` method, right-click on it, and select the View Call Hierarchy menu item. The window that appears is shown in Figure 10-14 (which has some regions expanded).




FIGURE 10-14

Starting from the `Shuffle()` method, you can drill into the tree view in the window to find all the code that calls the method, and all the calls that the method makes. For example, the highlighted method, `Next(int)`, is called from `Shuffle()`, so it appears in the Calls From 'Shuffle' section. When you click on a call you can see the line of code that makes the call on the right, along with its location. You can double-click on the location to navigate instantly to the line of code that is referred to.

You can also drill into methods further down the hierarchy — in Figure 10-19 this has been done for `Main()`, and the display shows calls from and to the `Main()` method.

This window is very useful when you are debugging or refactoring code, as it enables you to see at a glance how different pieces of code are related.

### EXERCISES

- 10.1** Write code that defines a base class, `MyClass`, with the virtual method `GetString()`. This method should return the string stored in the protected field `myString`, accessible through the write-only public property `ContainedString`.

- 10.2 Derive a class, `MyDerivedClass`, from `MyClass`. Override the `GetString()` method to return the string from the base class, using the base implementation of the method, but add the text "(output from derived class)" to the returned string.
- 10.3 Partial method definitions must use the `void` return type. Provide a reason why this is so.
- 10.4 Write a class called `MyCopyableClass` that is capable of returning a copy of itself using the method `GetCopy()`. This method should use the `MemberwiseClone()` method inherited from `System.Object`. Add a simple property to the class, and write client code that uses the class to confirm that everything is working.
- 10.5 Write a console client for the `Ch10CardLib` library that draws five cards at one time from a shuffled `Deck` object. If all five cards are the same suit, then the client should display the card names onscreen along with the text `Flush!`; otherwise, it should quit after 50 cards with the text `No flush`.

---

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Member definitions	You can define field, method, and property members in a class. Fields are defined with an accessibility, name, and type. Methods are defined with an accessibility, return type, name, and parameters. Properties are defined with an accessibility, name, and a <code>get</code> and/or <code>set</code> accessor. Individual property accessors can have their own accessibility, which must be less accessible than the property as a whole.
Member hiding and overrides	Properties and methods can be defined as <code>abstract</code> or <code>virtual</code> in base classes to define inheritance. Derived classes must implement abstract members, and can override virtual members, with the <code>override</code> keyword. They can also provide new implementations with the <code>new</code> keyword, and prevent further overrides of <code>virtual</code> members with the <code>sealed</code> keyword. Base implementations can be called with the <code>base</code> keyword.
Interface implementation	A class that implements an interface must implement all of the members defined by that interface as public. You can implement interfaces implicitly or explicitly, where explicit implementations are only available through an interface reference.
Partial definitions	You can split class definitions across multiple code files with the <code>partial</code> keyword. You can also create partial methods with the <code>partial</code> keyword. Partial methods have certain restrictions, including no return value or <code>out</code> parameters, and are not compiled if no implementation is provided.

# 11

## Collections, Comparisons, and Conversions

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Defining and using collections
- Learning the types of collections that are available
- Comparing types and using the `is` operator
- Comparing values and overloading operators
- Defining and using conversions
- Using the `as` operator

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 11 download and individually named according to the names throughout the chapter.

You've covered all the basic OOP techniques in C# now, but there are some more advanced techniques that are worth becoming familiar with. These techniques relate to certain problems that you must solve regularly when you are writing code. Learning about them will make it much easier to progress and allow you to concentrate on other, potentially more important aspects of your applications. In this chapter, you look at the following:

- **Collections** — Collections enable you to maintain groups of objects. Unlike arrays, which you've used in earlier chapters, collections can include more advanced functionality, such as controlling access to the objects they contain, searching and sorting, and

more. You'll learn how to use and create collection classes and learn about some powerful techniques for getting the most out of them.

- **Comparisons** — When dealing with objects, you often want to make comparisons between them. This is especially important in collections, because it is how sorting is achieved. You'll look at how to compare objects in a number of ways, including operator overloading, and how to use the `IComparable` and `IComparer` interface to sort collections.
- **Conversions** — Earlier chapters showed you how to cast objects from one type into another. In this chapter, you'll learn how to customize type conversions to suit your needs.

## COLLECTIONS

In Chapter 5, you learned how to use arrays to create variable types that contain a number of objects or values. Arrays, however, have their limitations. The biggest limitation is that once arrays have been created, they have a fixed size, so you can't add new items to the end of an existing array without creating a new one. This often means that the syntax used to manipulate arrays can become overly complicated. OOP techniques enable you to create classes that perform much of this manipulation internally, simplifying the code that uses lists of items or arrays.

Arrays in C# are implemented as instances of the `System.Array` class and are just one type of what are known as *collection classes*. Collection classes in general are used for maintaining lists of objects, and they may expose more functionality than simple arrays. Much of this functionality comes through implementing interfaces from the `System.Collections` namespace, thus standardizing collection syntax. This namespace also contains some other interesting things, such as classes that implement these interfaces in ways other than `System.Array`.

Because the collection's functionality (including basic functions such as accessing collection items by using `[index]` syntax) is available through interfaces, you aren't limited to using basic collection classes such as `System.Array`. Instead, you can create your own customized collection classes. These can be made more specific to the objects you want to enumerate (that is, the objects you want to maintain collections of). One advantage of doing this, as you will see, is that custom collection classes can be *strongly typed*. That is, when you extract items from the collection, you don't need to cast them into the correct type. Another advantage is the capability to expose specialized methods. For example, you can provide a quick way to obtain subsets of items. In the deck of cards example, you could add a method to obtain all `Card` items of a particular suit.

Several interfaces in the `System.Collections` namespace provide basic collection functionality:

- `IEnumerable` — Provides the capability to loop through items in a collection
- `ICollection` — Provides the capability to obtain the number of items in a collection and copy items into a simple array type (inherits from `IEnumerable`)
- `IList` — Provides a list of items for a collection along with the capabilities for accessing these items, and some other basic capabilities related to lists of items (inherits from `IEnumerable` and `ICollection`)
- `IDictionary` — Similar to `IList`, but provides a list of items accessible via a key value, rather than an index (inherits from `IEnumerable` and `ICollection`)

The `System.Array` class implements `IList`, `ICollection`, and `IEnumerable`. However, it doesn't support some of the more advanced features of `IList`, and it represents a list of items by using a fixed size.

## Using Collections

One of the classes in the `Systems.Collections` namespace, `System.Collections.ArrayList`, also implements `IList`, `ICollection`, and `IEnumerable`, but does so in a more sophisticated way than `System.Array`. Whereas arrays are fixed in size (you can't add or remove elements), this class can be used to represent a variable-length list of items. To give you more of a feel for what is possible with such a highly advanced collection, the following Try It Out uses this class, as well as a simple array.

### TRY IT OUT Arrays versus More Advanced Collections: Ch11Ex01

1. Create a new console application called Ch11Ex01 and save it in the directory `C:\BegVCSharp\Chapter11`.
2. Add three new classes, `Animal`, `Cow`, and `Chicken`, to the project by right-clicking on the project in the Solution Explorer window and selecting `Add`  $\Rightarrow$  `Class` for each.
3. Modify the code in `Animal.cs` as follows:

```
namespace Ch11Ex01
{
    public abstract class Animal
    {
        protected string name;
        public string Name
        {
            get { return name; }
            set { name = value; }
        }
        public Animal()
        {
            name = "The animal with no name";
        }
        public Animal(string newName)
        {
            name = newName;
        }
        public void Feed() => WriteLine($"{name} has been fed.");
    }
}
```

4. Modify the code in `Cow.cs` as follows:

```
namespace Ch11Ex01
{
    public class Cow : Animal
    {
        public void Milk() => WriteLine($"{name} has been milked.");
        public Cow(string newName) : base(newName) {}
    }
}
```

5. Modify the code in Chicken.cs as follows:


```
namespace Ch11Ex01
{
    public class Chicken : Animal
    {
        public void LayEgg() => WriteLine($"'{name}' has laid an egg.");
        public Chicken(string newName) : base(newName) {}
    }
}
```

6. Modify the code in Program.cs as follows:

```
using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using static System.Console;
namespace Ch11Ex01
{
    class Program
    {
        static void Main(string[] args)
        {
            WriteLine("Create an Array type collection of Animal " +
                "objects and use it:");
            Animal[] animalArray = new Animal[2];
            Cow myCow1 = new Cow("Lea");
            animalArray[0] = myCow1;
            animalArray[1] = new Chicken("Noa");
            foreach (Animal myAnimal in animalArray)
            {
                WriteLine($"New {myAnimal.ToString()} object added to Array" +
                    $" collection, Name = {myAnimal.Name}");
            }
            WriteLine($"Array collection contains {animalArray.Length} objects.");
            animalArray[0].Feed();
            ((Chicken)animalArray[1]).LayEgg();
            WriteLine();
            WriteLine("Create an ArrayList type collection of Animal " +
                "objects and use it:");
            ArrayList animalArrayList = new ArrayList();
            Cow myCow2 = new Cow("Rual");
            animalArrayList.Add(myCow2);
            animalArrayList.Add(new Chicken("Andrea"));
            foreach (Animal myAnimal in animalArrayList)
            {
                WriteLine($"New {myAnimal.ToString()} object added to ArrayList " +
                    $" collection, Name = {myAnimal.Name}");
            }
            WriteLine($"ArrayList collection contains {animalArrayList.Count} " +
                $" objects.");
            ((Animal)animalArrayList[0]).Feed();
```

```
((Chicken)animalArrayList[1]).LayEgg();
WriteLine();
WriteLine("Additional manipulation of ArrayList:");
animalArrayList.RemoveAt(0);
((Animal)animalArrayList[0]).Feed();
animalArrayList.AddRange(animalArray);
((Chicken)animalArrayList[2]).LayEgg();
WriteLine($"The animal called {myCow1.Name} is at " +
    $"index {animalArrayList.IndexOf(myCow1)}.");
myCow1.Name = "Mary";
WriteLine("The animal is now " +
    $" called {((Animal)animalArrayList[1]).Name }.");
.ReadKey();
}
}
```

7. Run the application. The result is shown in Figure 11-1.



**FIGURE 11-1**

## How It Works

This example creates two collections of objects: the first uses the `System.Array` class (that is, a simple array), and the second uses the `System.Collections.ArrayList` class. Both collections are of `Animal` objects, which are defined in `Animal.cs`. The `Animal` class is abstract, so it can't be instantiated, although you can have items in your collection that are instances of the `Cow` and `Chicken` classes, which are derived from `Animal`. You achieve this by using polymorphism, discussed in Chapter 8.

Once created in the `Main()` method in `Class1.cs`, these arrays are manipulated to show their characteristics and capabilities. Several of the operations performed apply to both `Array` and `ArrayList` collections, although their syntax differs slightly. Some, however, are possible only by using the more advanced `ArrayList` type.

You'll learn the similar operations first, comparing the code and results for both types of collection. First, collection creation. With simple arrays you must initialize the array with a fixed size in order to use it. You do this to an array called `animalArray` by using the standard syntax shown in Chapter 5:

```
Animal[] animalArray = new Animal[2];
```

ArrayList collections, conversely, don't need a size to be initialized, so you can create your list (called `animalArrayList`) as follows:

```
ArrayList animalArrayList = new ArrayList();
```

You can use two other constructors with this class. The first copies the contents of an existing collection to the new instance by specifying the existing collection as a parameter; the other sets the capacity of the collection, also via a parameter. This capacity, specified as an `int` value, sets the initial number of items that can be contained in the collection. This is not an absolute capacity, however, because it is doubled automatically if the number of items in the collection ever exceeds this value.

With arrays of reference types (such as the `Animal` and `Animal`-derived objects), simply initializing the array with a size doesn't initialize the items it contains. To use a given entry, that entry needs to be initialized, which means that you need to assign initialized objects to the items:

```
Cow myCow1 = new Cow("Lea");
animalArray[0] = myCow1;
animalArray[1] = new Chicken("Noa");
```

The preceding code does this in two ways: once by assignment using an existing `Cow` object, and once by assignment through the creation of a new `Chicken` object. The main difference here is that the former method creates a reference to the object in the array — a fact that you make use of later in the code.

With the `ArrayList` collection, there are no existing items, not even `null`-referenced ones. This means you can't assign new instances to indices in the same way. Instead, you use the `Add()` method of the `ArrayList` object to add new items:

```
Cow myCow2 = new Cow("Rual");
animalArrayList.Add(myCow2);
animalArrayList.Add(new Chicken("Andrea"));
```

Apart from the slightly different syntax, you can add new or existing objects to the collection in the same way. Once you have added items in this way, you can overwrite them by using syntax identical to that for arrays:

```
animalArrayList[0] = new Cow("Alma");
```

You won't do that in this example, though.

Chapter 5 showed how the `foreach` structure can be used to iterate through an array. This is possible because the `System.Array` class implements the `IEnumerable` interface, and the only method on this interface, `GetEnumerator()`, allows you to loop through items in the collection. You'll look at this in more depth a little later in the chapter. In your code, you write out information about each `Animal` object in the array:

```
foreach (Animal myAnimal in animalArray)
{
    WriteLine($"New {myAnimal.ToString()} object added to Array " +
        $"collection, Name = {myAnimal.Name}");
}
```

The `ArrayList` object you use also supports the `IEnumerable` interface and can be used with `foreach`. In this case, the syntax is identical:

```
foreach (Animal myAnimal in animalArrayList)
{
    WriteLine($"New {myAnimal.ToString()} object added to ArrayList " +
              $"collection, Name = {myAnimal.Name}");
}
```

Next, you use the array's `Length` property to output to the screen the number of items in the array:

```
WriteLine($"Array collection contains {animalArrayList.Length} objects.");
```

You can achieve the same thing with the `ArrayList` collection, except that you use the `Count` property that is part of the `ICollection` interface:

```
WriteLine($"ArrayList collection contains {animalArrayList.Count} objects.");
```

Collections — whether simple arrays or more complex collections — aren't very useful unless they provide access to the items that belong to them. Simple arrays are strongly typed — that is, they allow direct access to the type of the items they contain. This means you can call the methods of the item directly:

```
animalArray[0].Feed();
```

The type of the array is the abstract type `Animal`; therefore, you can't call methods supplied by derived classes directly. Instead you must use casting:

```
((Chicken)animalArray[1]).LayEgg();
```

The `ArrayList` collection is a collection of `System.Object` objects (you have assigned `Animal` objects via polymorphism). This means that you must use casting for all items:

```
((Animal)animalArrayList[0]).Feed();
((Chicken)animalArrayList[1]).LayEgg();
```

The remainder of the code looks at some of the `ArrayList` collection's capabilities that go beyond those of the `ArrayList` collection. First, you can remove items by using the `Remove()` and `RemoveAt()` methods, part of the `IList` interface implementation in the `ArrayList` class. These methods remove items from an array based on an item reference or index, respectively. This example uses the latter method to remove the list's first item, the `Cow` object with a `Name` property of `Hayley`:

```
animalArrayList.RemoveAt(0);
```

Alternatively, you could use

```
animalArrayList.Remove(myCow2);
```

because you already have a local reference to this object — you added an existing reference to the array via `Add()`, rather than create a new object. Either way, the only item left in the collection is the `Chicken` object, which you access as follows:

```
((Animal)animalArrayList[0]).Feed();
```

Any modifications to items in the `ArrayList` object resulting in `N` items being left in the array will be executed in such a way as to maintain indices from `0` to `N-1`. For example, removing the item with the index `0` results in all other items being shifted one place in the array, so you access the `Chicken` object with the index `0`, not `1`. You no longer have an item with an index of `1` (because you only had two items in the first place), so an exception would be thrown if you tried the following:

```
((Animal)animalArrayList[1]).Feed();
```

`ArrayList` collections enable you to add several items at once with the `AddRange()` method. This method accepts any object with the `ICollection` interface, which includes the `animalArray` array created earlier in the code:

```
animalArrayList.AddRange(animalArray);
```

To check that this works, you can attempt to access the third item in the collection, which is the second item in `animalArray`:

```
((Chicken)animalArrayList[2]).LayEgg();
```

The `AddRange()` method isn't part of any of the interfaces exposed by `ArrayList`. This method is specific to the `ArrayList` class and demonstrates the fact that you can exhibit customized behavior in your collection classes, beyond what is required by the interfaces you have looked at. This class exposes other interesting methods too, such as `InsertRange()`, for inserting an array of objects at any point in the list, and methods for tasks such as sorting and reordering the array.

Finally, you make use of the fact that you can have multiple references to the same object. Using the `IndexOf()` method (part of the `IList` interface), you can see that `myCow1` (an object originally added to `animalArray`) is now not only part of the `animalArrayList` collection, but also its index:

```
WriteLine($"The animal called {myCow1.Name} is at index " +  
    $"{animalArrayList.IndexOf(myCow1)}");
```

As an extension of this, the next two lines of code rename the object via the object reference and display the new name via the collection reference:

```
myCow1.Name = "Mary";  
WriteLine($"The animal is now called {((Animal)animalArrayList[1]).Name}.");
```

---

## Defining Collections

Now that you know what is possible using more advanced collection classes, it's time to learn how to create your own strongly typed collection. One way of doing this is to implement the required methods manually, but this can be a time-consuming and complex process. Alternatively, you can derive your collection from a class, such as `System.Collections.CollectionBase`, an abstract class that supplies much of the implementation of a collection for you. This option is strongly recommended.

The `CollectionBase` class exposes the interfaces `IEnumerable`, `ICollection`, and `IList` but provides only some of the required implementation — notably, the `Clear()` and `RemoveAt()` methods

of `IList` and the `Count` property of `ICollection`. You need to implement everything else yourself if you want the functionality provided.

To facilitate this, `CollectionBase` provides two protected properties that enable access to the stored objects themselves. You can use `List`, which gives you access to the items through an `IList` interface, and `InnerList`, which is the `ArrayList` object used to store items.

For example, the basics of a collection class to store `Animal` objects could be defined as follows (you'll see a fuller implementation shortly):

```
public class Animals : CollectionBase
{
    public void Add(Animal newAnimal)
    {
        List.Add(newAnimal);
    }
    public void Remove(Animal oldAnimal)
    {
        List.Remove(oldAnimal);
    }
    public Animals() {}
}
```

Here, `Add()` and `Remove()` have been implemented as strongly typed methods that use the standard `Add()` method of the `IList` interface used to access the items. The methods exposed will now only work with `Animal` classes or classes derived from `Animal`, unlike the `ArrayList` implementations shown earlier, which work with any object.

The `CollectionBase` class enables you to use the `foreach` syntax with your derived collections. For example, you can use code such as this:

```
WriteLine("Using custom collection class Animals:");
Animals animalCollection = new Animals();
animalCollection.Add(new Cow("Lea"));
foreach (Animal myAnimal in animalCollection)
{
    WriteLine($"New { myAnimal.ToString() } object added to custom " +
              $"collection, Name = {myAnimal.Name}");
}
```

You can't, however, do the following:

```
animalCollection[0].Feed();
```

To access items via their indices in this way, you need to use an indexer.

## Indexers

An *indexer* is a special kind of property that you can add to a class to provide array-like access. In fact, you can provide more complex access via an indexer, because you can define and use complex parameter types with the square bracket syntax as you want. Implementing a simple numeric index for items, however, is the most common usage.

You can add an indexer to the `Animals` collection of `Animal` objects as follows:

```
public class Animals : CollectionBase
{
    ...
    public Animal this[int animalIndex]
    {
        get { return (Animal)List[animalIndex]; }
        Set { List[animalIndex] = value; }
    }
}
```

The `this` keyword is used along with parameters in square brackets, but otherwise the indexer looks much like any other property. This syntax is logical, because you access the indexer by using the name of the object followed by the index parameter(s) in square brackets (for example, `MyAnimals[0]`).

The indexer code uses an indexer on the `List` property (that is, on the `IList` interface that provides access to the `ArrayList` in `CollectionBase` that stores your items):

```
return (Animal)List[animalIndex];
```

Explicit casting is necessary here, as the `IList.List` property returns a `System.Object` object. The important point to note here is that you define a type for this indexer. This is the type that will be obtained when you access an item by using this indexer. This strong typing means that you can write code such as

```
animalCollection[0].Feed();
```

rather than:

```
((Animal)animalCollection[0]).Feed();
```

This is another handy feature of strongly typed custom collections. In the following Try It Out, you expand the previous Try It Out to put this into action.

### TRY IT OUT Implementing an Animals Collection: Ch11Ex02

1. Create a new console application called `Ch11Ex02` and save it in the directory `C:\BegVCSharp\Chapter11`.
2. Right-click on the project name in the Solution Explorer window and select Add  $\leftrightarrow$  Existing Item.
3. Select the `Animal.cs`, `Cow.cs`, and `Chicken.cs` files from the `C:\BegVCSharp\Chapter11\Ch11Ex01\Ch11Ex01` directory, and click Add.
4. Modify the namespace declaration in the three files you added as follows:  

```
namespace Ch11Ex02
```
5. Add a new class called `Animals`.

6. Modify the code in Animals.cs as follows:

```
using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch11Ex02
{
    public class Animals : CollectionBase
    {
        public void Add(Animal newAnimal)
        {
            List.Add(newAnimal);
        }
        public void Remove(Animal newAnimal)
        {
            List.Remove(newAnimal);
        }
        public Animal this[int animalIndex]
        {
            get { return (Animal)List[animalIndex]; }
            set { List[animalIndex] = value; }
        }
    }
}
```

7. Modify Program.cs as follows:

```
static void Main(string[] args)
{
    Animals animalCollection = new Animals();
    animalCollection.Add(new Cow("Donna"));
    animalCollection.Add(new Chicken("Kevin"));
    foreach (Animal myAnimal in animalCollection)
    {
        myAnimal.Feed();
    }
    ReadKey();
}
```

8. Execute the application. The result is shown in Figure 11-2.




FIGURE 11-2

## How It Works

This example uses code detailed in the last section to implement a strongly typed collection of `Animal` objects in a class called `Animals`. The code in `Main()` simply instantiates an `Animals` object called `animalCollection`, adds two items (an instance of `Cow` and `Chicken`), and uses a `foreach` loop to call the `Feed()` method that both objects inherit from their base class, `Animal`.

## Adding a Cards Collection to CardLib

In the last chapter, you created a class library project called `Ch10CardLib` that contained a `Card` class representing a playing card, and a `Deck` class representing a deck of cards — that is, a collection of `Card` classes. This collection was implemented as a simple array.

In this chapter, you'll add a new class to this library, renamed `Ch11CardLib`. This new class, `Cards`, will be a custom collection of `Card` objects, giving you all the benefits described earlier in this chapter. Create a new class library called `Ch11CardLib` in the `C:\BegVCSharp\Chapter11` directory. Next, delete the autogenerated `Class1.cs` file; select `Project` → `Add Existing Item`; select the `Card.cs`, `Deck.cs`, `Suit.cs`, and `Rank.cs` files from the `C:\BegVCSharp\Chapter10\Ch10CardLib\Ch10CardLib` directory; and add the files to your project. As with the previous version of this project, introduced in Chapter 10, these changes are presented without using the standard Try It Out format. Should you want to jump straight to the code, feel free to open the version of this project included in the downloadable code for this chapter.

**NOTE** Don't forget that when copying the source files from `Ch10CardLib` to `Ch11CardLib`, you must change the namespace declarations to refer to `Ch11CardLib`. This also applies to the `Ch10CardClient` console application that you will use for testing.

The downloadable code for this chapter includes a `Ch11CardLib` folder that contains all the code you need for the various expansions to the `Ch11CardLib` project. Because of this, you may notice some extra code that isn't included in this example, but this won't affect how it works at this stage. Often you will find that code is commented out; however, when you reach the relevant example, you can uncomment the section you want to experiment with.

If you decide to create this project yourself, add a new class called `Cards` and modify the code in `Cards.cs` as follows:

```
using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch11CardLib
```

```

{
    public class Cards : CollectionBase
    {
        public void Add(Card newCard)
        {
            List.Add(newCard);
        }
        public void Remove(Card oldCard)
        {
            List.Remove(oldCard);
        }
        public Card this[int cardIndex]
        {
            get { return (Card)List[cardIndex]; }
            set { List[cardIndex] = value; }
        }
        /// <summary>
        /// Utility method for copying card instances into another Cards
        /// instance—used in Deck.Shuffle(). This implementation assumes that
        /// source and target collections are the same size.
        /// </summary>
        public void CopyTo(Cards targetCards)
        {
            for (int index = 0; index < this.Count; index++)
            {
                targetCards[index] = this[index];
            }
        }
        /// <summary>
        /// Check to see if the Cards collection contains a particular card.
        /// This calls the Contains() method of the ArrayList for the
        collection,
        /// which you access through the InnerList property.
        /// </summary>
        public bool Contains(Card card) => InnerList.Contains(card);
    }
}

```

Next, modify `Deck.cs` to use this new collection, rather than an array:

```

using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
namespace Ch11CardLib
{
    public class Deck
    {
        private Cards cards = new Cards();
        public Deck()
        {
            // Line of code removed here
            for (int suitVal = 0; suitVal < 4; suitVal++)
            {
                for (int rankVal = 1; rankVal < 14; rankVal++)

```

```
        {
            cards.Add(new Card((Suit)suitVal, (Rank)rankVal));
        }
    }
}
public Card GetCard(int cardNum)
{
    if (cardNum >= 0 && cardNum <= 51)
        return cards[cardNum];
    else
        throw (new System.ArgumentOutOfRangeException("cardNum", cardNum,
            "Value must be between 0 and 51."));
}
public void Shuffle()
{
    Cards newDeck = new Cards();
    bool[] assigned = new bool[52];
    Random sourceGen = new Random();
    for (int i = 0; i < 52; i++)
    {
        int sourceCard = 0;
        bool foundCard = false;
        while (foundCard == false)
        {
            sourceCard = sourceGen.Next(52);
            if (assigned[sourceCard] == false)
                foundCard = true;
        }
        assigned[sourceCard] = true;
        newDeck.Add(cards[sourceCard]);
    }
    newDeck.CopyTo(cards);
}
}
```

Not many changes are necessary here. Most of them involve changing the shuffling logic to allow for the fact that cards are added to the beginning of the new `Cards` collection `newDeck` from a random index in `cards`, rather than to a random index in `newDeck` from a sequential position in `cards`.

The client console application for the `Ch10CardLib` solution, `Ch10CardClient`, can be used with this new library with the same result as before, as the method signatures of `Deck` are unchanged. Clients of this class library can now make use of the `Cards` collection class, however, rather than rely on arrays of `Card` objects — for example, to define hands of cards in a card game application.

## Keyed Collections and `IDictionary`

Instead of implementing the `IList` interface, it is also possible for collections to implement the similar `IDictionary` interface, which allows items to be indexed via a key value (such as a string name), rather than an index. This is also achieved using an indexer, although here the indexer parameter

used is a key associated with a stored item, rather than an `int` index, which can make the collection a lot more user-friendly.

As with indexed collections, there is a base class you can use to simplify implementation of the `IDictionary` interface: `DictionaryBase`. This class also implements `IEnumerable` and `ICollection`, providing the basic collection-manipulation capabilities that are the same for any collection.

`DictionaryBase`, like `CollectionBase`, implements some (but not all) of the members obtained through its supported interfaces. Like `CollectionBase`, the `Clear` and `Count` members are implemented, although `RemoveAt()` isn't because it's a method on the `IList` interface and doesn't appear on the `IDictionary` interface. `IDictionary` does, however, have a `Remove()` method, which is one of the methods you should implement in a custom collection class based on `DictionaryBase`.

The following code shows an alternative version of the `Animals` class, this time derived from `DictionaryBase`. Implementations are included for `Add()`, `Remove()`, and a key-accessed indexer:

```
public class Animals : DictionaryBase
{
    public void Add(string newID, Animal newAnimal)
    {
        Dictionary.Add(newID, newAnimal);
    }
    public void Remove(string animalID)
    {
        Dictionary.Remove(animalID);
    }
    public Animals() {}
    public Animal this[string animalID]
    {
        get { return (Animal)Dictionary[animalID]; }
        set { Dictionary[animalID] = value; }
    }
}
```

The differences in these members are as follows:

- `Add()` — Takes two parameters, a key and a value, to store together. The dictionary collection has a member called `Dictionary` inherited from `DictionaryBase`, which is an `IDictionary` interface. This interface has its own `Add()` method, which takes two object parameters. Your implementation takes a string value as a key and an `Animal` object as the data to store alongside this key.
- `Remove()` — Takes a key parameter, rather than an object reference. The item with the key value specified is removed.
- `Indexer` — Uses a string key value, rather than an index, which is used to access the stored item via the `Dictionary` inherited member. Again, casting is necessary here.

One other difference between collections based on `DictionaryBase` and collections based on `CollectionBase` is that `foreach` works slightly differently. The collection from the last section allowed you to extract `Animal` objects directly from the collection. Using `foreach` with the `DictionaryBase` derived class gives you `DictionaryEntry` structs, another type defined in the `System.Collections` namespace. To get to the `Animal` objects themselves, you must use the `Value` member of this struct, or you can use the `Key` member of the struct to get the associated key. To get code equivalent to the earlier

```
foreach (Animal myAnimal in animalCollection)
{
    WriteLine($"New {myAnimal.ToString()} object added to custom " +
              $"collection, Name = {myAnimal.Name}");
}
```

you need the following:

```
foreach (DictionaryEntry myEntry in animalCollection)
{
    WriteLine($"New {myEntry.Value.ToString()} object added to " +
              $"custom collection, Name = {((Animal)myEntry.Value).Name}");
}
```

It is possible to override this behavior so that you can access `Animal` objects directly through `foreach`. There are several ways to do this, the simplest being to implement an iterator.

## Iterators

Earlier in this chapter, you saw that the `IEnumerable` interface enables you to use `foreach` loops. It's often beneficial to use your classes in `foreach` loops, not just collection classes such as those shown in previous sections.

However, overriding this behavior, or providing your own custom implementation of it, is not always simple. To illustrate this, it's necessary to take a detailed look at `foreach` loops. The following steps show you what actually happens in a `foreach` loop iterating through a collection called `collectionObject`:

1. `collectionObject.GetEnumerator()` is called, which returns an `IEnumerator` reference. This method is available through implementation of the `IEnumerable` interface, although this is optional.
2. The `MoveNext()` method of the returned `IEnumerator` interface is called.
3. If `MoveNext()` returns `true`, then the `Current` property of the `IEnumerator` interface is used to get a reference to an object, which is used in the `foreach` loop.
4. The preceding two steps repeat until `MoveNext()` returns `false`, at which point the loop terminates.

To enable this behavior in your classes, you must override several methods, keep track of indices, maintain the `Current` property, and so on. This can be a lot of work to achieve very little.

A simpler alternative is to use an iterator. Effectively, using iterators generates a lot of the code for you behind the scenes and hooks it all up correctly. Moreover, the syntax for using iterators is much easier to get a grip on.

A good definition of an iterator is a block of code that supplies all the values to be used in a `foreach` block in sequence. Typically, this block of code is a method, although you can also use property accessors and other blocks of code as iterators. To keep things simple, you'll just look at methods here.

Whatever the block of code is, its return type is restricted. Perhaps contrary to expectations, this return type isn't the same as the type of object being enumerated. For example, in a class that represents a collection of `Animal` objects, the return type of the iterator block can't be `Animal`. Two possible return types are the interface types mentioned earlier, `IEnumerable` or `IEnumerator`. You use these types as follows:

- To iterate over a class, use a method called `GetEnumerator()` with a return type of `IEnumerator`.
- To iterate over a class member, such as a method, use `IEnumerable`.

Within an iterator block, you select the values to be used in the `foreach` loop by using the `yield` keyword. The syntax for doing this is as follows:

```
yield return <value>;
```

That information is all you need to build a very simple example, as follows (you can find this code in `SimpleIterators\Program.cs`):

```
public static IEnumerable SimpleList()
{
    yield return "string 1";
    yield return "string 2";
    yield return "string 3";
}
static void Main(string[] args)
{
    foreach (string item in SimpleList())
        WriteLine(item);
    ReadKey();
}
```

Here, the static method `SimpleList()` is the iterator block. Because it is a method, you use a return type of `IEnumerable`. `SimpleList()` uses the `yield` keyword to supply three values to the `foreach` block that uses it, each of which is written to the screen. The result is shown in Figure 11-3.




FIGURE 11-3

Obviously, this iterator isn't a particularly useful one, but it does show how this works in action and how simple the implementation can be. Looking at the code, you might wonder how the code knows to return `string` type items. In fact, it doesn't; it returns `object` type values. As you know, `object` is the base class for all types, so you can return anything from the `yield` statements.

However, the compiler is intelligent enough that you can interpret the returned values as whatever type you want in the context of the `foreach` loop. Here, the code asks for `string` type values, so those are the values you get to work. Should you change one of the `yield` lines so that it returns, say, an integer, you would get a bad cast exception in the `foreach` loop.

One more thing about iterators. It is possible to interrupt the return of information to the `foreach` loop by using the following statement:

```
yield break;
```

When this statement is encountered in an iterator, the iterator processing terminates immediately, as does the `foreach` loop using it.

Now it's time for a more complicated — and useful! — example. In this Try It Out, you'll implement an iterator that obtains prime numbers.

### TRY IT OUT Implementing an Iterator: Ch11Ex03

1. Create a new console application called Ch11Ex03 and save it in the directory `C:\BegVCSharp\Chapter11`.
2. Add a new class called `Primes` and modify the code in `Primes.cs` as follows:


```
using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch11Ex03
{
    public class Primes
    {
        private long min;
        private long max;
        public Primes() : this(2, 100) {}
        public Primes(long minimum, long maximum)
        {
            if (minimum < 2)
                min = 2;
            else
                min = minimum;
            max = maximum;
        }
        public IEnumerator GetEnumerator()
        {
            for (long possiblePrime = min; possiblePrime <= max; possiblePrime++)
                yield return possiblePrime;
        }
    }
}
```

```
    {
        bool isPrime = true;
        for (long possibleFactor = 2; possibleFactor <=
            (long)Math.Floor(Math.Sqrt(possiblePrime)); possibleFactor++)
        {
            long remainderAfterDivision = possiblePrime % possibleFactor;
            if (remainderAfterDivision == 0)
            {
                isPrime = false;
                break;
            }
        }
        if (isPrime)
        {
            yield return possiblePrime;
        }
    }
}
```

3. Modify the code in Program.cs as follows:

```
static void Main(string[] args)
{
    Primes primesFrom2To1000 = new Primes(2, 1000);
    foreach (long i in primesFrom2To1000)
        Write($"'{i}' ");
    ReadKey();
}
```

4. Execute the application. The result is shown in Figure 11-4.



**FIGURE 11-4**

## How It Works

This example consists of a class that enables you to enumerate over a collection of prime numbers between an upper and lower limit. The class that encapsulates the prime numbers uses an iterator to provide this functionality.

The code for `Primes` starts off with the basics: two fields to hold the maximum and minimum values to search between, and constructors to set these values. Note that the minimum value is restricted — it

can't be less than 2. This makes sense, because 2 is the lowest prime number. The interesting code is all in the `GetEnumerator()` method. The method signature fulfills the rules for an iterator block in that it returns an `IEnumerator` type:

```
public I IEnumerator GetEnumerator()
{
```

To extract prime numbers between limits, you need to test each number in turn, so you start with a `for` loop:

```
for (long possiblePrime = min; possiblePrime <= max; possiblePrime++)
{
```

Because you don't know whether a number is prime, you first assume that it is and then check to see if it isn't. That means checking whether any number between 2 and the square root of the number to be tested is a factor. If this is `true`, then the number isn't prime, so you move on to the next one. If the number is indeed prime, then you pass it to the `foreach` loop using `yield`:

```
bool isPrime = true;
for (long possibleFactor = 2; possibleFactor <=
    (long)Math.Floor(Math.Sqrt(possiblePrime)); possibleFactor++)
{
    long remainderAfterDivision = possiblePrime % possibleFactor;
    if (remainderAfterDivision == 0)
    {
        isPrime = false;
        break;
    }
}
if (isPrime)
{
    yield return possiblePrime;
}
```

An interesting fact reveals itself through this code if you set the minimum and maximum limits to very big numbers. When you execute the application, the results appear one at a time, with pauses in between, rather than all at once. This is evidence that the iterator code returns results one at a time, despite the fact that there is no obvious place where the code terminates between `yield` calls. Behind the scenes, calling `yield` does interrupt the code, which resumes when another value is requested — that is, when the `foreach` loop using the iterator begins a new cycle.

---

## Iterators and Collections

Earlier you were promised an explanation of how iterators can be used to iterate over the objects stored in a dictionary-type collection without having to deal with `DictionaryItem` objects. In the downloadable code for this chapter, you will find the code for the next project in the `DictionaryAnimals` folder. Recall the collection class `Animals`:

```
public class Animals : DictionaryBase
{
    public void Add(string newID, Animal newAnimal)
```

```

    {
        Dictionary.Add(newID, newAnimal);
    }
    public void Remove(string animalID)
    {
        Dictionary.Remove(animalID);
    }
    public Animal this[string animalID]
    {
        get { return (Animal)Dictionary[animalID]; }
        set { Dictionary[animalID] = value; }
    }
}
}

```

You can add this simple iterator to the code to get the desired behavior:

```

public new IEnumerator GetEnumerator()
{
    foreach (object animal in Dictionary.Values)
        yield return (Animal)animal;
}

```

Now you can use the following code to iterate through the `Animal` objects in the collection:

```

foreach (Animal myAnimal in animalCollection)
{
    WriteLine($"New {myAnimal.ToString()} object added to " +
              $" custom collection, Name = {myAnimal.Name}");
}

```

## Deep Copying

Chapter 9 described how you can perform shallow copying with the `System.Object.MemberwiseClone()` protected method, by using a method like the `GetCopy()` one shown here:

```

public class Cloner
{
    public int Val;
    public Cloner(int newVal)
    {
        Val = newVal;
    }
    public object GetCopy() => MemberwiseClone();
}

```

Suppose you have fields that are reference types, rather than value types (for example, objects):

```

public class Content
{
    public int Val;
}
public class Cloner
{
    public Content MyContent = new Content();
}

```

```
public Cloner(int newVal)
{
    MyContent.Val = newVal;
}
public object GetCopy() => MemberwiseClone();
}
```

In this case, the shallow copy obtained through `GetCopy()` has a field that refers to the same object as the original object. The following code, which uses this `Cloner` class, illustrates the consequences of shallow copying reference types:

```
Cloner mySource = new Cloner(5);
Cloner myTarget = (Cloner)mySource.GetCopy();
WriteLine($"myTarget.MyContent.Val = {myTarget.MyContent.Val}");
mySource.MyContent.Val = 2;
WriteLine($"myTarget.MyContent.Val = {myTarget.MyContent.Val}");
```

The fourth line, which assigns a value to `mySource.MyContent.Val`, the `Val` public field of the `MyContent` public field of the original object, also changes the value of `myTarget.MyContent.Val`. That's because `mySource.MyContent` refers to the same object instance as `myTarget.MyContent`. The output of the preceding code is as follows:

```
myTarget.MyContent.Val = 5
myTarget.MyContent.Val = 2
```

To get around this, you need to perform a deep copy. You could just modify the `GetCopy()` method used previously to do this, but it is preferable to use the standard .NET Framework way of doing things: implement the `ICloneable` interface, which has the single method `Clone()`. This method takes no parameters and returns an `object` type result, giving it a signature identical to the `GetCopy()` method used earlier.

To modify the preceding classes, try using the following deep copy code:

```
public class Content
{
    public int Val;
}
public class Cloner : ICloneable
{
    public Content MyContent = new Content();
    public Cloner(int newVal)
    {
        MyContent.Val = newVal;
    }
    public object Clone()
    {
        Cloner clonedCloner = new Cloner(MyContent.Val);
        return clonedCloner;
    }
}
```

This created a new `cloner` object by using the `val` field of the `Content` object contained in the original `Cloner` object (`MyContent`). This field is a value type, so no deeper copying is necessary.

Using code similar to that just shown to test the shallow copy — but using `Clone()` instead of `GetCopy()` — gives you the following result:

```
myTarget.MyContent.Val = 5
myTarget.MyContent.Val = 5
```

This time, the contained objects are independent. Note that sometimes calls to `Clone()` are made recursively, in more complex object systems. For example, if the `MyContent` field of the `Cloner` class also required deep copying, then you might need the following:

```
public class Cloner : ICloneable
{
    public Content MyContent = new Content();
    ...
    public object Clone()
    {
        Cloner clonedCloner = new Cloner();
        clonedCloner.MyContent = MyContent.Clone();
        return clonedCloner;
    }
}
```

You're calling the default constructor here to simplify the syntax of creating a new `Cloner` object. For this code to work, you would also need to implement `ICloneable` on the `Content` class.

## Adding Deep Copying to CardLib

You can put this into practice by implementing the capability to copy `Card`, `Cards`, and `Deck` objects by using the `ICloneable` interface. This might be useful in some card games, where you might not necessarily want two decks with references to the same set of `Card` objects, although you might conceivably want to set up one deck to have the same card order as another.

Implementing cloning functionality for the `Card` class in `Ch11CardLib` is simple because shallow copying is sufficient (`Card` contains only value-type data, in the form of fields). Begin by making the following changes to the class definition:

```
public class Card : ICloneable
{
    public object Clone() => MemberwiseClone();
```

This implementation of `ICloneable` is just a shallow copy. There is no rule determining what should happen in the `Clone()` method, and this is sufficient for your purposes.

Next, implement `ICloneable` on the `Cards` collection class. This is slightly more complicated because it involves cloning every `Card` object in the original collection — so you need to make a deep copy:

```
public class Cards : CollectionBase, ICloneable
{
```

```

public object Clone()
{
    Cards newCards = new Cards();
    foreach (Card sourceCard in List)
    {
        newCards.Add((Card)sourceCard.Clone());
    }
    return newCards;
}

```

Finally, implement `ICloneable` on the `Deck` class. Note a slight problem here: The `Deck` class in `Ch11CardLib` has no way to modify the cards it contains, short of shuffling them. There is no way, for example, to modify a `Deck` instance to have a given card order. To get around this, define a new private constructor for the `Deck` class that allows a specific `Cards` collection to be passed in when the `Deck` object is instantiated. Here's the code to implement cloning in this class:

```

public class Deck : ICloneable
{
    public object Clone()
    {
        Deck newDeck = new Deck(cards.Clone() as Cards);
        return newDeck;
    }
    private Deck(Cards newCards)
    {
        cards = newCards;
    }
}

```

Again, you can test this with some simple client code. As before, place this code within the `Main()` method of a client project for testing (you can find this code in `Ch11CardClient\Program.cs` in the chapter's online download):

```

Deck deck1 = new Deck();
Deck deck2 = (Deck)deck1.Clone();
WriteLine($"The first card in the original deck is: {deck1.GetCard(0)}");
WriteLine($"The first card in the cloned deck is: {deck2.GetCard(0)}");
deck1.Shuffle();
WriteLine("Original deck shuffled.");
WriteLine($"The first card in the original deck is: {deck1.GetCard(0)}");
WriteLine($"The first card in the cloned deck is: {deck2.GetCard(0)}");
.ReadKey();

```

The output will be similar to what is shown in Figure 11-5.




FIGURE 11-5

## COMPARISONS

This section covers two types of comparisons between objects:

- Type comparisons
- Value comparisons

Type comparisons — that is, determining what an object is, or what it inherits from — are important in all areas of C# programming. Often when you pass an object — to a method, for example — what happens next depends on the type of the object. You've seen this in passing in this and earlier chapters, but here you will see some more useful techniques.

Value comparisons are also something you've seen a lot of, at least with simple types. When it comes to comparing values of objects, things get a little more complicated. You have to define what is meant by a comparison for a start, and what operators such as `>` mean in the context of your classes. This is especially important in collections, for which you might want to sort objects according to some condition, perhaps alphabetically or according to a more complicated algorithm.

### Type Comparisons

When comparing objects, you often need to know their type, which enables you to determine whether a value comparison is possible. In Chapter 9 you saw the `GetType()` method, which all classes inherit from `System.Object`, and how this method can be used in combination with the `typeof()` operator to determine (and take action depending on) object types:

```
if (myObj.GetType() == typeof(MyComplexClass))
{
    // myObj is an instance of the class MyComplexClass.
}
```

You've also seen how the default implementation of `ToString()`, also inherited from `System.Object`, will get you a string representation of an object's type. You can compare these strings too, although that's a rather messy way to accomplish this.

This section demonstrates a handy shorthand way of doing things: the `is` operator. This operator allows for much more readable code and, as you will see, has the advantage of examining base classes. Before looking at the `is` operator, though, you need to be aware of what often happens behind the scenes when dealing with value types (as opposed to reference types): *boxing* and *unboxing*.

### Boxing and Unboxing

In Chapter 8, you learned the difference between reference types and value types, which was illustrated in Chapter 9 by comparing structs (which are value types) with classes (which are reference types). Boxing is the act of converting a value type into the `System.Object` type or to an interface type that is implemented by the value type. Unboxing is the opposite conversion.

For example, suppose you have the following struct type:

```
struct MyStruct
```

```
{  
    public int Val;  
}
```

You can box a struct of this type by placing it into an object-type variable:

```
MyStruct valType1 = new MyStruct();  
valType1.Val = 5;  
object refType = valType1;
```

Here, you create a new variable (`valType1`) of type `MyStruct`, assign a value to the `Val` member of this struct, and then box it into an object-type variable (`refType`).

The object created by boxing a variable in this way contains a reference to a copy of the value-type variable, not a reference to the original value-type variable. You can verify this by modifying the original struct's contents and then unboxing the struct contained in the object into a new variable and examining its contents:

```
valType1.Val = 6;  
MyStruct valType2 = (MyStruct)refType;  
WriteLine($"valType2.Val = {valType2.Val}");
```

This code gives you the following output:

```
valType2.Val = 5
```

When you assign a reference type to an object, however, you get a different behavior. You can see this by changing `MyStruct` into a class (ignoring the fact that the name of this class isn't appropriate now):

```
class MyStruct  
{  
    public int Val;  
}
```

With no changes to the client code shown previously (again ignoring the misnamed variables), you get the following output:

```
valType2.Val = 6
```

You can also box value types into interface types, so long as they implement that interface. For example, suppose the `MyStruct` type implements the `IMyInterface` interface as follows:

```
interface IMyInterface {}  
struct MyStruct : IMyInterface  
{  
    public int Val;  
}
```

You can then box the struct into an `IMyInterface` type as follows:

```
MyStruct valType1 = new MyStruct();  
IMyInterface refType = valType1;
```

You can unbox it by using the normal casting syntax:

```
MyStruct ValType2 = (MyStruct)refType;
```

As shown in these examples, boxing is performed without your intervention — that is, you don't have to write any code to make it possible. Unboxing a value requires an explicit conversion, however, and requires you to make a cast (boxing is implicit and doesn't have this requirement).

You might be wondering why you would actually want to do this. There are two very good reasons why boxing is extremely useful. First, it enables you to use value types in collections (such as `ArrayList`) where the items are of type `object`. Second, it's the internal mechanism that enables you to call `object` methods on value types, such as `ints` and `structs`.

It is worth noting that unboxing is necessary before access to the value type contents is possible.

## The `is` Operator

Despite its name, the `is` operator isn't a way to determine whether an object is a certain type. Instead, the `is` operator enables you to check whether an object either is or *can be converted into* a given type. If this is the case, then the operator evaluates to `true`.

Earlier examples showed a `Cow` and a `Chicken` class, both of which inherit from `Animal`. Using the `is` operator to compare objects with the `Animal` type will return `true` for objects of all three of these types, not just `Animal`. This is something you'd have a hard time achieving with the `GetType()` method and `typeof()` operator shown previously.

The `is` operator has the following syntax:

```
<operand> is <type>
```

The possible results of this expression are as follows:

- If `<type>` is a class type, then the result is `true` if `<operand>` is of that type, if it inherits from that type, or if it can be boxed into that type.
- If `<type>` is an interface type, then the result is `true` if `<operand>` is of that type or it is a type that implements the interface.
- If `<type>` is a value type, then the result is `true` if `<operand>` is of that type or it is a type that can be unboxed into that type.

The following Try It Out shows how this works in practice.


### TRY IT OUT Using the `is` Operator: Ch11Ex04\Program.cs

1. Create a new console application called Ch11Ex04 in the directory `C:\BegVCSharp\Chapter11`.
2. Modify the code in `Program.cs` as follows:

```
namespace Ch11Ex04
{
    class Checker
    {
        public void Check(object param1)
        {
            if (param1 is ClassA)
                WriteLine("Variable can be converted to ClassA.");
            else
                WriteLine("Variable can't be converted to ClassA.");
        }
    }
}
```

```
if (param1 is IMyInterface)
    WriteLine("Variable can be converted to IMyInterface.");
else
    WriteLine("Variable can't be converted to IMyInterface.");
if (param1 is MyStruct)
    WriteLine("Variable can be converted to MyStruct.");
else
    WriteLine("Variable can't be converted to MyStruct.");
}
}
interface IMyInterface {}
class ClassA : IMyInterface {}
class ClassB : IMyInterface {}
class ClassC {}
class ClassD : ClassA {}
struct MyStruct : IMyInterface {}
class Program
{
    static void Main(string[] args)
    {
        Checker check = new Checker();
        ClassA try1 = new ClassA();
        ClassB try2 = new ClassB();
        ClassC try3 = new ClassC();
        ClassD try4 = new ClassD();
        MyStruct try5 = new MyStruct();
        object try6 = try5;
        WriteLine("\nAnalyzing ClassA type variable:");
        check.Check(try1);
        WriteLine("\nAnalyzing ClassB type variable:");
        check.Check(try2);
        WriteLine("\nAnalyzing ClassC type variable:");
        check.Check(try3);
        WriteLine("\nAnalyzing ClassD type variable:");
        check.Check(try4);
        WriteLine("\nAnalyzing MyStruct type variable:");
        check.Check(try5);
        WriteLine("\nAnalyzing boxed MyStruct type variable:");
        check.Check(try6);
        ReadKey();
    }
}
}
```

3. Execute the code. The result is shown in Figure 11-6.



```

file:///C:/BegVCSharp/Chapter11/Ch11Ex04/bin/Debug/Ch...
Analyzing ClassA type variable:
Variable can be converted to ClassA.
Variable can be converted to IMyInterface.
Variable can't be converted to MyStruct.

Analyzing ClassB type variable:
Variable can't be converted to ClassA.
Variable can be converted to IMyInterface.
Variable can't be converted to MyStruct.

Analyzing ClassC type variable:
Variable can't be converted to ClassA.
Variable can't be converted to IMyInterface.
Variable can't be converted to MyStruct.

Analyzing ClassD type variable:
Variable can be converted to ClassA.
Variable can be converted to IMyInterface.
Variable can't be converted to MyStruct.

Analyzing MyStruct type variable:
Variable can't be converted to ClassA.
Variable can be converted to IMyInterface.
Variable can be converted to MyStruct.

Analyzing boxed MyStruct type variable:
Variable can't be converted to ClassA.
Variable can be converted to IMyInterface.
Variable can be converted to MyStruct.

```

**FIGURE 11-6**

## How It Works

This example illustrates the various results possible when using the `is` operator. Three classes, an interface, and a structure are defined and used as parameters to a method of a class that uses the `is` operator to determine whether they can be converted into the `ClassA` type, the interface type, and the struct type.

Only the `ClassA` and `ClassD` (which inherits from `ClassA`) types are compatible with `ClassA`. Types that don't inherit from a class are not compatible with that class.

The `ClassA`, `ClassB`, and `MyStruct` types all implement `IMyInterface`, so these are all compatible with the `IMyInterface` type. `ClassD` inherits from `ClassA`, so it too is compatible. Therefore, only `ClassC` is incompatible.

Finally, only variables of type `MyStruct` itself and boxed variables of that type are compatible with `MyStruct`, because you can't convert reference types to value types (although, of course, you can unbox previously boxed variables).

## Value Comparisons

Consider two `Person` objects representing people, each with an integer `Age` property. You might want to compare them to see which person is older. You can simply use the following code:

```

if (person1.Age > person2.Age)
{
    ...
}
```

This works fine, but there are alternatives. You might prefer to use syntax such as the following:

```
if (person1 > person2)
{
    ...
}
```

This is possible using *operator overloading*, which you'll look at in this section. This is a powerful technique, but it should be used judiciously. In the preceding code, it is not immediately obvious that ages are being compared — it could be height, weight, IQ, or just general “greatness.”

Another option is to use the `IComparable` and `IComparer` interfaces, which enable you to define how objects will be compared to each other in a standard way. This technique is supported by the various collection classes in the .NET Framework, making it an excellent way to sort objects in a collection.

## Operator Overloading

*Operator overloading* enables you to use standard operators, such as `+`, `>`, and so on, with classes that you design. This is called “overloading” because you are supplying your own implementations for these operators when used with specific parameter types, in much the same way that you overload methods by supplying different parameters for methods with the same name.

Operator overloading is useful because you can perform whatever processing you want in the implementation of the operator overload, which might not be as simple as, for example, `+`, meaning “add these two operands together.” Later, you’ll see a good example of this in a further upgrade of the `CardLib` library, whereby you’ll provide implementations for comparison operators that compare two cards to see which would beat the other in a trick (one round of card game play).

Because a trick in many card games depends on the suits of the cards involved, this isn’t as straightforward as comparing the numbers on the cards. If the second card laid down is a different suit from the first, then the first card wins regardless of its rank. You can implement this by considering the order of the two operands. You can also take a trump suit into account, whereby trumps beat other suits even if that isn’t the first suit laid down. This means that calculating that `card1 > card2` is `true` (that is, `card1` will beat `card2` if `card1` is laid down first), doesn’t necessarily imply that `card2 > card1` is `false`. If neither `card1` nor `card2` are trumps and they belong to different suits, then both of these comparisons will be `true`.

To start with, though, here’s a look at the basic syntax for operator overloading. Operators can be overloaded by adding operator type members (which must be static) to a class. Some operators have multiple uses (such as `-`, which has unary and binary capabilities); therefore, you also specify how many operands you are dealing with and the types of these operands. In general, you will have operands that are the same type as the class in which the operator is defined, although it’s possible to define operators that work on mixed types, as you’ll see shortly.

As an example, consider the simple type `AddClass1`, defined as follows:

```
public class AddClass1
{
    public int val;
}
```

This is just a wrapper around an `int` value but it illustrates the principles. With this class, code such as the following will fail to compile:

```
AddClass1 op1 = new AddClass1();
op1.val = 5;
AddClass1 op2 = new AddClass1();
op2.val = 5;
AddClass1 op3 = op1 + op2;
```

The error you get informs you that the `+` operator cannot be applied to operands of the `AddClass1` type. This is because you haven't defined an operation to perform yet. Code such as the following works, but it won't give you the result you might want:

```
AddClass1 op1 = new AddClass1();
op1.val = 5;
AddClass1 op2 = new AddClass1();
op2.val = 5;
bool op3 = op1 == op2;
```

Here, `op1` and `op2` are compared by using the `==` binary operator to determine whether they refer to the same object, not to verify whether their values are equal. `op3` will be `false` in the preceding code, even though `op1.val` and `op2.val` are identical.

To overload the `+` operator, use the following code:

```
public class AddClass1
{
    public int val;
    public static AddClass1 operator +(AddClass1 op1, AddClass1 op2)
    {
        AddClass1 returnVal = new AddClass1();
        returnVal.val = op1.val + op2.val;
        return returnVal;
    }
}
```

As you can see, operator overloads look much like standard `static` method declarations, except that they use the keyword `operator` and the operator itself, rather than a method name. You can now successfully use the `+` operator with this class, as in the previous example:

```
AddClass1 op3 = op1 + op2;
```

Overloading all binary operators fits the same pattern. Unary operators look similar but have only one parameter:

```
public class AddClass1
{
    public int val;
    public static AddClass1 operator +(AddClass1 op1, AddClass1 op2)
    {
        AddClass1 returnVal = new AddClass1();
        returnVal.val = op1.val + op2.val;
        return returnVal;
    }
}
```

```
public static AddClass1 operator -(AddClass1 op1)
{
    AddClass1 returnVal = new AddClass1();
    returnVal.val = -op1.val;
    return returnVal;
}
```

Both these operators work on operands of the same type as the class and have return values that are also of that type. Consider, however, the following class definitions:

```
public class AddClass1
{
    public int val;
    public static AddClass3 operator +(AddClass1 op1, AddClass2 op2)
    {
        AddClass3 returnVal = new AddClass3();
        returnVal.val = op1.val + op2.val;
        return returnVal;
    }
}
public class AddClass2
{
    public int val;
}
public class AddClass3
{
    public int val;
}
```

This will allow the following code:

```
AddClass1 op1 = new AddClass1();
op1.val = 5;
AddClass2 op2 = new AddClass2();
op2.val = 5;
AddClass3 op3 = op1 + op2;
```

When appropriate, you can mix types in this way. Note, however, that if you added the same operator to `AddClass2`, then the preceding code would fail because it would be ambiguous as to which operator to use. You should, therefore, take care not to add operators with the same signature to more than one class.

In addition, if you mix types, then the operands must be supplied in the same order as the parameters to the operator overload. If you attempt to use your overloaded operator with the operands in the wrong order, the operation will fail. For example, you can't use the operator like,

```
AddClass3 op3 = op2 + op1;
```

unless, of course, you supply another overload with the parameters reversed:

```
public static AddClass3 operator +(AddClass2 op1, AddClass1 op2)
{
    AddClass3 returnVal = new AddClass3();
```

```

        returnVal.val = op1.val + op2.val;
        return returnVal;
    }
}

```

The following operators can be overloaded:

- **Unary operators** — +, -, !, ~, ++, --, true, false
- **Binary operators** — +, -, \*, /, %, &, |, ^, <<, >>
- **Comparison operators** — ==, !=, <, >, <=, >=

**NOTE** If you overload the true and false operators, then you can use classes in Boolean expressions, such as if(op1) {}.

You can't overload assignment operators, such as +=, but these operators use their simple counterparts, such as +, so you don't have to worry about that. Overloading + means that += will function as expected. The = operator can't be overloaded because it has such a fundamental usage, but this operator is related to the user-defined conversion operators, which you'll look at in the next section.

You also can't overload && and ||, but these operators use the & and | operators to perform their calculations, so overloading these is enough.

Some operators, such as < and >, must be overloaded in pairs. That is, you can't overload < unless you also overload >. In many cases, you can simply call other operators from these to reduce the code required (and the errors that might occur), as shown in this example:

```

public class AddClass1
{
    public int val;
    public static bool operator >=(AddClass1 op1, AddClass1 op2)
        => (op1.val >= op2.val);
    public static bool operator <(AddClass1 op1, AddClass1 op2)
        => !(op1 >= op2);
    // Also need implementations for <= and > operators.
}

```

In more complex operator definitions, this can reduce the lines of code. It also means that you have less code to change if you later decide to modify the implementation of these operators.

The same applies to == and !=, but with these operators it is often worth overriding `Object.Equals()` and `Object.GetHashCode()`, because both of these functions can also be used to compare objects. By overriding these methods, you ensure that whatever technique users of the class use, they get the same result. This isn't essential, but it's worth adding for completeness. It requires the following nonstatic override methods:

```

public class AddClass1
{
    public int val;
    public static bool operator ==(AddClass1 op1, AddClass1 op2)

```

```
        => (op1.val == op2.val);
    public static bool operator !=(AddClass1 op1, AddClass1 op2)
        => !(op1 == op2);
    public override bool Equals(object op1) => val == ((AddClass1)op1).val;
    public override int GetHashCode() => val;
}
```

`GetHashCode()` is used to obtain a unique `int` value for an object instance based on its state. Here, using `val` is fine, because it is also an `int` value.

Note that `Equals()` uses an `object` type parameter. You need to use this signature or you will be overloading this method, rather than overriding it, and the default implementation will still be accessible to users of the class. Instead, you must use casting to get the required result. It is often worth checking the object type using the `is` operator discussed earlier, in code such as this:

```
public override bool Equals(object op1)
{
    if (op1 is AddClass1)
    {
        return val == ((AddClass1)op1).val;
    }
    else
    {
        throw new ArgumentException(
            "Cannot compare AddClass1 objects with objects of type " +
            op1.GetType().ToString());
    }
}
```

In this code, an exception is thrown if the operand passed to `Equals` is of the wrong type or cannot be converted into the correct type. Of course, this behavior might not be what you want. You might want to be able to compare objects of one type with objects of another type, in which case more branching would be necessary. Alternatively, you might want to restrict comparisons to those in which both objects are of exactly the same type, which would require the following change to the first `if` statement:

```
if (op1.GetType() == typeof(AddClass1))
```

## Adding Operator Overloads to CardLib

Now you'll upgrade your `Ch11CardLib` project again, adding operator overloading to the `Card` class. Again, you can find the code for the classes that follow in the `Ch11CardLib` folder of this chapter's code download. First, though, you'll add the extra fields to the `Card` class that allow for trump suits and an option to place aces high. You make these static, because when they are set, they apply to all `Card` objects:

```
public class Card
{
    /// <summary>
    /// Flag for trump usage. If true, trumps are valued higher
    /// than cards of other suits.
    /// </summary>
```

```

public static bool useTrumps = false;
/// <summary>
/// Trump suit to use if useTrumps is true.
/// </summary>
public static Suit trump = Suit.Club;
/// <summary>
/// Flag that determines whether aces are higher than kings or lower
/// than deuces.
/// </summary>
public static bool isAceHigh = true;

```

These rules apply to all `Card` objects in every `Deck` in an application. It's not possible to have two decks of cards with cards contained in each that obey different rules. That's fine for this class library, however, as you can safely assume that if a single application wants to use separate rules, then it could maintain these itself, perhaps setting the static members of `Card` whenever decks are switched.

Because you have done this, it is worth adding a few more constructors to the `Deck` class to initialize decks with different characteristics:

```

/// <summary>
/// Nondefault constructor. Allows aces to be set high.
/// </summary>
public Deck(bool isAceHigh) : this()
{
    Card.isAceHigh = isAceHigh;
}
/// <summary>
/// Nondefault constructor. Allows a trump suit to be used.
/// </summary>
public Deck(bool useTrumps, Suit trump) : this()
{
    Card.useTrumps = useTrumps;
    Card.trump = trump;
}
/// <summary>
/// Nondefault constructor. Allows aces to be set high and a trump suit
/// to be used.
/// </summary>
public Deck(bool isAceHigh, bool useTrumps, Suit trump) : this()
{
    Card.isAceHigh = isAceHigh;
    Card.useTrumps = useTrumps;
    Card.trump = trump;
}

```

Each of these constructors is defined by using the `: this()` syntax shown in Chapter 9, so in all cases the default constructor is called before the nondefault one, initializing the deck.

**NOTE** The null condition operator (?.) implemented in the == and > operator overload method is discussed in more detail in Chapter 12. The ?. in this code segment, card1?.suit, of the public static bool operator == method checks if the card1 object is null before attempting to retrieve the value stored in suit. This is important when you implement the method in later chapters.

Now add your operator overloads (and suggested overrides) to the Card class:

```

public static bool operator ==(Card card1, Card card2)
    => card1?.suit == card2?.suit) && (card1?.rank == card2?.rank);
public static bool operator !=(Card card1, Card card2)
    => !(card1 == card2);
public override bool Equals(object card) => this == (Card)card;
public override int GetHashCode()
    => return 13 * (int)suit + (int)rank;
public static bool operator >(Card card1, Card card2)
{
    if (card1.suit == card2.suit)
    {
        if (isAceHigh)
        {
            if (card1.rank == Rank.Ace)
            {
                if (card2.rank == Rank.Ace)
                    return false;
                else
                    return true;
            }
            else
            {
                if (card2.rank == Rank.Ace)
                    return false;
                else
                    return (card1.rank > card2?.rank);
            }
        }
        else
        {
            return (card1.rank > card2.rank);
        }
    }
    else
    {
        if (useTrumps && (card2.suit == Card.trump))
            return false;
        else
            return true;
    }
}

```

```

public static bool operator <(Card card1, Card card2)
    => !(card1 >= card2);
public static bool operator >=(Card card1, Card card2)
{
    if (card1.suit == card2.suit)
    {
        if (isAceHigh)
        {
            if (card1.rank == Rank.Ace)
            {
                return true;
            }
            else
            {
                if (card2.rank == Rank.Ace)
                    return false;
                else
                    return (card1.rank >= card2.rank);
            }
        }
        else
        {
            return (card1.rank >= card2.rank);
        }
    }
    else
    {
        if (useTrumps && (card2.suit == Card.trump))
            return false;
        else
            return true;
    }
}
public static bool operator <=(Card card1, Card card2)
    => !(card1 > card2);

```

There's not much to note here, except perhaps the slightly lengthy code for the `>` and `>=` overloaded operators. If you step through the code for `>`, you can see how it works and why these steps are necessary.

You are comparing two cards, `card1` and `card2`, where `card1` is assumed to be the first one laid down on the table. As discussed earlier, this becomes important when you are using trump cards, because a trump will beat a non-trump even if the non-trump has a higher rank. Of course, if the suits of the two cards are identical, then whether the suit is the trump suit or not is irrelevant, so this is the first comparison you make:

```

public static bool operator >(Card card1, Card card2)
{
    if (card1.suit == card2.suit)
    {

```

If the static `isAceHigh` flag is `true`, then you can't compare the cards' ranks directly via their value in the `Rank` enumeration, because the rank of ace has a value of `1` in this enumeration, which is less than that of all other ranks. Instead, use the following steps:

- If the first card is an ace, then check whether the second card is also an ace. If it is, then the first card won't beat the second. If the second card isn't an ace, then the first card wins:

```
if (isAceHigh)
{
    if (card1.rank == Rank.Ace)
    {
        if (card2.rank == Rank.Ace)
            return false;
        else
            return true;
    }
}
```

- If the first card isn't an ace, then you also need to check whether the second one is. If it is, then the second card wins; otherwise, you can compare the rank values because you know that aces aren't an issue:

```
else
{
    if (card2.rank == Rank.Ace)
        return false;
    else
        return (card1.rank > card2?.rank);
}
}
```

- If aces aren't high, then you just compare the rank values:

```
else
{
    return (card1.rank > card2.rank);
}
```

The remainder of the code concerns the case where the suits of `card1` and `card2` are different. Here, the static `useTrumps` flag is important. If this flag is `true` and `card2` is of the trump suit, then you can say definitively that `card1` isn't a trump (because the two cards have different suits); and trumps always win, so `card2` is the higher card:

```
else
{
    if (useTrumps && (card2.suit == Card.trump))
        return false;
```

If `card2` isn't a trump (or `useTrumps` is `false`), then `card1` wins, because it was the first card laid down:

```
else
    return true;
}
}
```

Only one other operator (`>=`) uses code similar to this, and the other operators are very simple, so there's no need to go into more detail about them.

The following simple client code tests these operators. Simply place it in the `Main()` method of a client project to test it, like the client code shown earlier in the `CardLib` examples (you can find this code in `Ch11CardClient\Program.cs`):

```
Card.isAceHigh = true;
WriteLine("Aces are high.");
Card.useTrumps = true;
Card.trump = Suit.Club;
WriteLine("Clubs are trumps.");
Card card1, card2, card3, card4, card5;
card1 = new Card(Suit.Club, Rank.Five);
card2 = new Card(Suit.Club, Rank.Five);
card3 = new Card(Suit.Club, Rank.Ace);
card4 = new Card(Suit.Heart, Rank.Ten);
card5 = new Card(Suit.Diamond, Rank.Ace);
WriteLine($"{card1.ToString()} == {card2.ToString()} ? {card1 == card2}");
WriteLine($"{card1.ToString()} != {card3.ToString()} ? {card1 != card3}");
WriteLine($"{card1.ToString()}.Equals({card4.ToString()}) ? " +
    $" { card1.Equals(card4) }";
WriteLine($"Card.Equals({card3.ToString()}, {card4.ToString()}) ? " +
    $" { Card.Equals(card3, card4) }";
WriteLine($"{card1.ToString()} > {card2.ToString()} ? {card1 > card2}");
WriteLine($"{card1.ToString()} <= {card3.ToString()} ? {card1 <= card3}");
WriteLine($"{card1.ToString()} > {card4.ToString()} ? {card1 > card4}");
WriteLine($"{card4.ToString()} > {card1.ToString()} ? {card4 > card1}");
WriteLine($"{card5.ToString()} > {card4.ToString()} ? {card5 > card4}");
WriteLine($"{card4.ToString()} > {card5.ToString()} ? {card4 > card5}");
.ReadKey();
```

The results are as shown in Figure 11-7.

In each case, the operators are applied taking the specified rules into account. This is particularly apparent in the last four lines of output, demonstrating how trump cards always beat non-trumps.

```
file:///C:/BegVCSharp/Chapter11/Ch11CardClient/bin/Debug...
Aces are high.
Clubs are trumps.
The Five of Clubs == The Ace of Clubs ? True
The Five of Clubs != The Ace of Clubs ? True
The Five of Clubs.Equals(The Ten of Hearts) ? False
Card.Equals(The Ace of Clubs, The Ten of Hearts) ? False
The Five of Clubs > The Ace of Clubs ? False
The Five of Clubs <= The Ace of Clubs ? True
The Five of Clubs > The Ten of Hearts ? True
The Ten of Hearts > The Five of Clubs ? False
The Ace of Diamonds > The Ten of Hearts ? True
The Ten of Hearts > The Ace of Diamonds ? True
```

FIGURE 11-7

## The `IComparable` and `IComparer` Interfaces

The `IComparable` and `IComparer` interfaces are the standard way to compare objects in the .NET Framework. The difference between the interfaces is as follows:

- `IComparable` is implemented in the class of the object to be compared and allows comparisons between that object and another object.
- `IComparer` is implemented in a separate class, which allows comparisons between any two objects.

Typically, you give a class default comparison code by using `IComparable`, and nondefault comparisons using other classes.

`IComparable` exposes the single method `CompareTo()`, which accepts an object. You could, for example, implement it in a way that enables you to pass a `Person` object to it and determine whether that person is older or younger than the current person. In fact, this method returns an `int`, so you could also determine how much older or younger the second person is:

```
if (person1.CompareTo(person2) == 0)
{
    WriteLine("Same age");
}
else if (person1.CompareTo(person2) > 0)
{
    WriteLine("person 1 is Older");
}
else
{
    WriteLine("person1 is Younger");
}
```

`IComparer` exposes the single method `Compare()`, which accepts two objects and returns an integer result just like `CompareTo()`. With an object supporting `IComparer`, you could use code like the following:

```
if (personComparer.Compare(person1, person2) == 0)
{
    WriteLine("Same age");
}
else if (personComparer.Compare(person1, person2) > 0)
{
    WriteLine("person 1 is Older");
}
else
{
    WriteLine("person1 is Younger");
}
```

In both cases, the parameters supplied to the methods are of the type `System.Object`. This means that you can compare one object to another object of any other type, so you usually have to perform some type comparison before returning a result, and maybe even throw exceptions if the wrong types are used.

The .NET Framework includes a default implementation of the `IComparer` interface on a class called `Comparer`, found in the `System.Collections` namespace. This class is capable of performing culture-specific comparisons between simple types, as well as any type that supports the `IComparable` interface. You can use it, for example, with the following code:

```
string firstString = "First String";
string secondString = "Second String";
WriteLine($"Comparing '{firstString}' and '{secondString}', " +
    $"result: {Comparer.Default.Compare(firstString, secondString)}");
int firstNumber = 35;
int secondNumber = 23;
WriteLine($"Comparing '{firstNumber}' and '{secondNumber}', " +
    $"result: {Comparer.Default.Compare(firstNumber, secondNumber)}");
```

This uses the `Comparer.Default` static member to obtain an instance of the `Comparer` class, and then uses the `Compare()` method to compare first two strings, and then two integers.

The result is as follows:

```
Comparing 'First String' and 'Second String', result: -1
Comparing '35' and '23', result: 1
```

Because F comes before S in the alphabet, it is deemed “less than” S, so the result of the first comparison is `-1`. Similarly, 35 is greater than 23, hence the result of `1`. Note that the results do not indicate the magnitude of the difference.

When using `Comparer`, you must use types that can be compared. Attempting to compare `firstString` with `firstNumber`, for instance, will generate an exception.

Here are a few more points about the behavior of this class:

- Objects passed to `Comparer.Compare()` are checked to determine whether they support `IComparable`. If they do, then that implementation is used.
- Null values are allowed, and are interpreted as being “less than” any other object.
- Strings are processed according to the current culture. To process strings according to a different culture (or language), the `Comparer` class must be instantiated using its constructor, which enables you to pass a `System.Globalization.CultureInfo` object specifying the culture to use.
- Strings are processed in a case-sensitive way. To process them in a non-case-sensitive way, you need to use the `CaseInsensitiveComparer` class, which otherwise works exactly the same.

## Sorting Collections

Many collection classes allow sorting, either by default comparisons between objects or by custom methods. `ArrayList` is one example. It contains the method `Sort()`, which can be used without parameters, in which case default comparisons are used, or it can be passed an `IComparer` interface to use to compare pairs of objects.

When you have an `ArrayList` filled with simple types, such as integers or strings, the default comparer is fine. For your own classes, you must either implement `IComparable` in your class definition or create a separate class supporting `IComparer` to use for comparisons.

Note that some classes in the `System.Collections` namespace, including `CollectionBase`, don't expose a method for sorting. If you want to sort a collection you have derived from this class, then you have to do a bit more work and sort the internal `List` collection yourself.

The following Try It Out shows how to use a default and nondefault comparer to sort a list.

### TRY IT OUT Sorting a List: Ch11Ex05

1. Create a new console application called `Ch11Ex05` in the directory `C:\BegVCSharp\Chapter11`.
2. Add a new class called `Person` and modify the code in `Person.cs` as follows:

```
namespace Ch11Ex05
{
    public class Person : IComparable
    {
        public string Name;
        public int Age;
        public Person(string name, int age)
        {
            Name = name;
            Age = age;
        }
        public int CompareTo(object obj)
        {
            if (obj is Person)
            {
                Person otherPerson = obj as Person;
                return this.Age - otherPerson.Age;
            }
            else
            {
                throw new ArgumentException(
                    "Object to compare to is not a Person object.");
            }
        }
    }
}
```

3. Add another new class called `PersonComparerName` and modify the code as follows:

```
using System;
using System.Collections;
```

```

using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch11Ex05
{
    public class PersonComparerName : IComparer
    {
        public static IComparer Default = new PersonComparerName();
        public int Compare(object x, object y)
        {
            if (x is Person && y is Person)
            {
                return Comparer.Default.Compare(
                    ((Person)x).Name, ((Person)y).Name);
            }
            else
            {
                throw new ArgumentException(
                    "One or both objects to compare are not Person objects.");
            }
        }
    }
}

```

4. Modify the code in Program.cs as follows:


```

using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using static System.Console;
namespace Ch11Ex05
{
    class Program
    {
        static void Main(string[] args)
        {
            ArrayList list = new ArrayList();
            list.Add(new Person("Rual", 30));
            list.Add(new Person("Donna", 25));
            list.Add(new Person("Mary", 27));
            list.Add(new Person("Ben", 44));
            WriteLine("Unsorted people:");
            for (int i = 0; i < list.Count; i++)
            {
                WriteLine($"{(list[i] as Person).Name} ({(list[i] as Person).Age})");
            }
            WriteLine();
            WriteLine(
                "People sorted with default comparer (by age):");
            list.Sort();
        }
    }
}

```

```
        for (int i = 0; i < list.Count; i++)
    {
        WriteLine($"{{(list[i] as Person).Name }} ({{(list[i] as Person).Age }})");
    }
    WriteLine();
    WriteLine(
        "People sorted with nondefault comparer (by name):");
    list.Sort(PersonComparerName.Default);
    for (int i = 0; i < list.Count; i++)
    {
        WriteLine($"{{(list[i] as Person).Name }} ({{(list[i] as Person).Age }})");
    }
    ReadKey();
}
}
```

5. Execute the code. The result is shown in Figure 11-8.



**FIGURE 11-8**

## *How It Works*

An `ArrayList` containing `Person` objects is sorted in two different ways here. By calling the `ArrayList.Sort()` method with no parameters, the default comparison is used, which is the `CompareTo()` method in the `Person` class (because this class implements `IComparable`):

```
public int CompareTo(object obj)
{
    if (obj is Person)
    {
        Person otherPerson = obj as Person;
        return this.Age - otherPerson.Age;
    }
    else
    {
        throw new ArgumentException(
            "Object to compare to is not a Person object.");
    }
}
```

---

This method first checks whether its argument can be compared to a `Person` object — that is, whether the object can be converted into a `Person` object. If there is a problem, then an exception is thrown. Otherwise, the `Age` properties of the two `Person` objects are compared.

Next, a nondefault comparison sort is performed using the `PersonComparerName` class, which implements `IComparer`. This class has a public static field for ease of use:

```
public static IComparer Default = new PersonComparerName();
```

This enables you to get an instance using `PersonComparerName.Default`, just like the `Comparer` class shown earlier. The `CompareTo()` method of this class is as follows:

```
public int Compare(object x, object y)
{
    if (x is Person && y is Person)
    {
        return Comparer.Default.Compare(
            ((Person)x).Name, ((Person)y).Name);
    }
    else
    {
        throw new ArgumentException(
            "One or both objects to compare are not Person objects.");
    }
}
```

Again, arguments are first checked to determine whether they are `Person` objects. If they aren't, then an exception is thrown. If they are, then the default `Comparer` object is used to compare the two string `Name` fields of the `Person` objects.

---

## CONVERSIONS

Thus far, you have used casting whenever you have needed to convert one type into another, but this isn't the only way to do things. Just as an `int` can be converted into a `long` or a `double` implicitly as part of a calculation, you can define how classes you have created can be converted into other classes (either implicitly or explicitly). To do this, you overload conversion operators, much like other operators were overloaded earlier in this chapter. You'll see how in the first part of this section. You'll also see another useful operator, the `as` operator, which in general is preferable to casting when using reference types.

### Overloading Conversion Operators

As well as overloading mathematical operators, as shown earlier, you can define both implicit and explicit conversions between types. This is necessary if you want to convert between types that

aren't related — if there is no inheritance relationship between them and no shared interfaces, for example.

Suppose you define an implicit conversion between `ConvClass1` and `ConvClass2`. This means that you can write code such as the following:

```
ConvClass1 op1 = new ConvClass1();
ConvClass2 op2 = op1;
```

Alternatively, you can define an explicit conversion:

```
ConvClass1 op1 = new ConvClass1();
ConvClass2 op2 = (ConvClass2)op1;
```

As an example, consider the following code:

```
public class ConvClass1
{
    public int val;
    public static implicit operator ConvClass2(ConvClass1 op1)
    {
        ConvClass2 returnVal = new ConvClass2();
        returnVal.val = op1.val;
        return returnVal;
    }
}
public class ConvClass2
{
    public double val;
    public static explicit operator ConvClass1(ConvClass2 op1)
    {
        ConvClass1 returnVal = new ConvClass1();
        checked {returnVal.val = (int)op1.val;};
        return returnVal;
    }
}
```

Here, `ConvClass1` contains an `int` value and `ConvClass2` contains a `double` value. Because `int` values can be converted into `double` values implicitly, you can define an implicit conversion between `ConvClass1` and `ConvClass2`. The reverse is not true, however, and you should define the conversion operator between `ConvClass2` and `ConvClass1` as explicit.

You specify this using the `implicit` and `explicit` keywords as shown. With these classes, the following code is fine:

```
ConvClass1 op1 = new ConvClass1();
op1.val = 3;
ConvClass2 op2 = op1;
```

A conversion in the other direction, however, requires the following explicit casting conversion:

```
ConvClass2 op1 = new ConvClass2();
op1.val = 3e15;
ConvClass1 op2 = (ConvClass1)op1;
```

Because you have used the `checked` keyword in your explicit conversion, you will get an exception in the preceding code, as the `val` property of `op1` is too large to fit into the `val` property of `op2`.

## The as Operator

The `as` operator converts a type into a specified reference type, using the following syntax:

```
<operand> as <type>
```

This is possible only in certain circumstances:

- If `<operand>` is of type `<type>`
- If `<operand>` can be implicitly converted to type `<type>`
- If `<operand>` can be boxed into type `<type>`

If no conversion from `<operand>` to `<type>` is possible, then the result of the expression will be `null`.

Conversion from a base class to a derived class is possible by using an explicit conversion, but it won't always work. Consider the two classes `ClassA` and `ClassD` from an earlier example, where `ClassD` inherits from `ClassA`:

```
class ClassA : IMyInterface {}
class ClassD : ClassA {}
```

The following code uses the `as` operator to convert from a `ClassA` instance stored in `obj1` into the `ClassD` type:

```
ClassA obj1 = new ClassA();
ClassD obj2 = obj1 as ClassD;
```

This will result in `obj2` being `null`.

However, it is possible to store `ClassD` instances in `ClassA`-type variables by using polymorphism. The following code illustrates this, using the `as` operator to convert from a `ClassA`-type variable containing a `ClassD`-type instance into the `ClassD` type:

```
ClassD obj1 = new ClassD();
ClassA obj2 = obj1;
ClassD obj3 = obj2 as ClassD;
```

This time the result is that `obj3` ends up containing a reference to the same object as `obj1`, not `null`.

This functionality makes the `as` operator very useful, because the following code (which uses simple casting) results in an exception being thrown:

```
ClassA obj1 = new ClassA();
ClassD obj2 = (ClassD)obj1;
```

The `as` equivalent of this code results in a `null` value being assigned to `obj2` — no exception is thrown. This means that code such as the following (using two of the classes developed earlier in this chapter, `Animal` and a class derived from `Animal` called `Cow`) is very common in C# applications:

```
public void MilkCow(Animal myAnimal)
{
    Cow myCow = myAnimal as Cow;
    if (myCow != null)
    {
        myCow.Milk();
    }
    else
    {
        WriteLine($"{myAnimal.Name} isn't a cow, and so can't be milked.");
    }
}
```

This is much simpler than checking for exceptions!

## EXERCISES

- 11.1** Create a collection class called `People` that is a collection of the following `Person` class. The items in the collection should be accessible via a string indexer that is the name of the person, identical to the `Person.Name` property.

```
public class Person
{
    private string name;
    private int age;
    public string Name
    {
        get { return name; }
        set { name = value; }
    }
    public int Age
    {
        get { return age; }
        set { age = value; }
    }
}
```

- 11.2** Extend the `Person` class from the preceding exercise so that the `>`, `<`, `>=`, and `<=` operators are overloaded, and compare the `Age` properties of `Person` instances.
- 11.3** Add a `GetOldest()` method to the `People` class that returns an array of `Person` objects with the greatest `Age` property (one or more objects, as multiple items can have the same value for this property), using the overloaded operators defined in Exercise 11.2.
- 11.4** Implement the `ICloneable` interface on the `People` class to provide deep copying capability.

---

**11.5** Add an iterator to the `People` class that enables you to get the ages of all members in a `foreach` loop as follows:

```
foreach (int age in myPeople.Ages)
{
    // Display ages.
```

---

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

Key Concept	Description
Defining collections	Collections are classes that can contain instances of other classes. You can define a collection by deriving from <code>CollectionBase</code> , or implement collection interfaces such as <code>IEnumerable</code> , <code>ICollection</code> , and <code>IList</code> yourself. Typically, you will define an indexer for your collection in order to use <code>collection[index]</code> syntax to access members.
Dictionaries	You can also define keyed collections, or dictionaries, whereby each item has an associated key. In this case, the key can be used to identify an item, rather than using the item's index. You can define a dictionary by implementing <code>IDictionary</code> or by deriving a class from <code>DictionaryBase</code> .
Iterators	You can implement an iterator to control how looping code obtains values in its loop cycles. To iterate over a class, implement a method called <code>GetEnumerator()</code> with a return type of <code>IEnumerator</code> . To iterate over a class member, such as a method, use a return type of <code>IEnumerable</code> . In iterator code blocks, return values with the <code>yield</code> keyword.
Type comparisons	You can use the <code>GetType()</code> method to obtain the type of an object, or the <code>typeof()</code> operator to get the type of a class. You can then compare these type values. You can also use the <code>is</code> operator to determine whether an object is compatible with a certain class type.
Value comparisons	If you want to make classes whose instances can be compared using standard C# operators, you must overload those operators in the class definition. For other types of value comparison, you can use classes that implement the <code>IComparable</code> or <code>IComparer</code> interfaces. These interfaces are particularly useful for sorting collections.
The <code>as</code> operator	You can use the <code>as</code> operator to convert a value to a reference type. If no conversion is possible, the <code>as</code> operator returns a <code>null</code> value.

# 12

## Generics

---

### WHAT YOU WILL LEARN IN THIS CHAPTER:

---

- Discovering generics
- Using generic classes provided by the .NET Framework
- Defining your own generics
- Learning how variance works with generics

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 12 download and individually named according to the names throughout the chapter.

This chapter begins by looking at what generics are. You learn about generics in fairly abstract terms at first, because learning the concepts behind generics is crucial to being able to use them effectively.

Next, you see some of the generic types in the .NET Framework in action. This will help you understand their functionality and power, as well as the new syntax required in your code. You'll then move on to define your own generic types, including generic classes, interfaces, methods, and delegates. You also learn additional techniques for further customizing generic types: the `default` keyword and type constraints.

Finally, you'll look at covariance and contravariance, two forms of variance that were introduced in C# 4 and that allow greater flexibility when using generic classes.

## WHAT ARE GENERICS?

To best illustrate what generics are, and why they are so useful, recall the collection classes from the previous chapter. You saw how basic collections can be contained in classes such as `ArrayList`, but that such collections suffer from being untyped. This requires that you cast object items into whatever type of objects you actually stored in the collection. Because anything that inherits from `System.Object` (that is, practically anything) can be stored in an `ArrayList`, you need to be careful. Assuming that certain types are all that is contained in a collection can lead to exceptions being thrown, and code logic breaking down. You learned some techniques to deal with this, including the code required to check the type of an object.

However, you discovered that a much better solution is to use a strongly typed collection class initially. By deriving from `CollectionBase` and providing your own methods for adding, removing, and otherwise accessing members of the collection, you learned how you could restrict collection members to those derived from a certain base type or supporting a certain interface. This is where you encounter a problem. Every time you create a new class that needs to be held in a collection, you must do one of the following:

- Use a collection class you've already made that can contain items of the new type.
- Create a new collection class that can hold items of the new type, implementing all the required methods.

Typically, with a new type you need extra functionality, so more often than not, you need a new collection class anyway. Therefore, making collection classes can take up a fair amount of your time!

*Generic classes*, conversely, make coding a lot simpler. A generic class is built around whatever type, or types, you supply during instantiation, enabling you to strongly type an object with hardly any effort at all. In the context of collections, creating a “collection of type T objects” is as simple as saying it aloud — and achievable in a single line of code. Instead of code such as,

```
CollectionClass items = new CollectionClass();
items.Add(new ItemClass());
```

you can use this:

```
CollectionClass<ItemClass> items = new CollectionClass<ItemClass>();
items.Add(new ItemClass());
```

The angle bracket syntax is the way you pass type parameters to generic types. In the preceding code, read `CollectionClass<ItemClass>` as `CollectionClass` of `ItemClass`. You will, of course, examine this syntax in more detail later in the chapter.

There's more to the subject of generics than just collections, but they are particularly suited to this area, as you will see later in the chapter when you look at the `System.Collections.Generic` namespace. By creating a generic class, you can generate methods that have a signature that can be strongly typed to any type you want, even catering to the fact that a type can be a value or reference type, and deal with individual cases as they occur. You can even allow only a subset of types to be

used, by restricting the types used to instantiate a generic class to those that support a given interface or are derived from a certain type. Moreover, you're not restricted to generic classes — you can create generic interfaces, generic methods (which can be defined on nongeneric classes), and even generic delegates. All this adds a great deal of flexibility to your code, and judicious use of generics can eliminate hours of development time.

**NOTE** If you are familiar with C++, this is one difference between C++ templates and C# generic classes. In C++ the compiler detects where you used a specific type of template — for example, A of B — and compiles the code necessary to create this type. In C# everything happens at runtime.

You're probably wondering how all this is possible. Usually, when you create a class, it is compiled into a type that you can then use in your code. You might think that when you create a generic class, it would have to be compiled into a plethora of types, so that you could instantiate it. Fortunately, that's not the case — and given the infinite amount of classes possible in .NET, that's just as well. Behind the scenes, the .NET runtime allows generic classes to be dynamically generated as and when you need them. A given generic class A of B won't exist until you ask for it by instantiating it.

## USING GENERICS

Before you look at how to create your own generic types, it's worth looking at the ones supplied by the .NET Framework. These include the types in the `System.Collections.Generic` namespace, a namespace that you've seen several times in your code because it is included by default in console applications. You haven't yet used any of the types in this namespace, but that's about to change. This section looks at the types in this namespace and how you can use them to create strongly typed collections and improve the functionality of your existing collections.

First, though, you'll look at another simpler generic type that gets around a minor issue with value types: *nullable types*.

### Nullable Types

In earlier chapters, you saw that one of the ways in which value types (which include most of the basic types such as `int` and `double` as well as all structs) differ from reference types (`string` and any class) is that they must contain a value. They can exist in an unassigned state, just after they are declared and before a value is assigned, but you can't make use of the value type in that state in any way. Conversely, reference types can be `null`.

There are times, and they crop up more often than you might think (particularly when you work with databases), when it is useful to have a value type that can be `null`. Generics give you a way to do this using the `System.Nullable<T>` type, as shown in this example:

```
System.Nullable<int> nullableInt;
```

This code declares a variable called `nullableInt`, which can have any value that an `int` variable can, plus the value `null`. This enables you to write code such as the following:

```
nullableInt = null;
```

If `nullableInt` were an `int` type variable, then the preceding code wouldn't compile.

The preceding assignment is equivalent to the following:

```
nullableInt = new System.Nullable<int>();
```

As with any other variable, you can't just use it before some kind of initialization, whether to `null` (through either syntax shown previously) or by assigning a value.

You can test nullable types to determine whether they are `null`, just like you test reference types:

```
if (nullableInt == null)
{
    ...
}
```

Alternatively, you can use the `HasValue` property:

```
if (nullableInt.HasValue)
{
    ...
}
```

This wouldn't work for reference types, even one with a `HasValue` property of its own, because having a `null`-valued reference type variable means that no object exists through which to access this property, and an exception would be thrown.

You can also look at the value of a nullable type by using the `Value` property. If `HasValue` is `true`, then you are guaranteed a non-`null` value for `Value`; but if `HasValue` is `false` — that is, `null` has been assigned to the variable — then accessing `Value` will result in an exception of type `System.InvalidOperationException`.

Note that nullable types are so useful that they have resulted in a modification of C# syntax. Rather than use the syntax shown previously to declare a nullable type variable, you can instead use the following:

```
int? nullableInt;
```

`int?` is simply a shorthand for `System.Nullable<int>` but is much more readable. In subsequent sections, you'll use this syntax.

## Operators and Nullable Types

With simple types, such as `int`, you can use operators such as `+`, `-`, and so on to work with values. With nullable type equivalents, there is no difference: The values contained in nullable types are implicitly converted to the required type and the appropriate operators are used. This also applies to structs with operators that you have supplied:

```
int? op1 = 5;
int? result = op1 * 2;
```

Note that here the `result` variable is also of type `int?`. The following code will not compile:

```
int? op1 = 5;
int result = op1 * 2;
```

To get this to work you must perform an explicit conversion or access the value through the `Value` property, which requires code such as,

```
int? op1 = 5;
int result = (int)op1 * 2;
```

or:

```
int? op1 = 5;
int result = op1.Value * 2;
```

This works fine as long as `op1` has a value — if it is `null`, then you will get an exception of type `System.InvalidOperationException`.

This raises the obvious question: What happens when one or both values in an operator evaluation that involves two nullable values are `null`, such as `op1` in the following code?

```
int? op1 = null;
int? op2 = 5;
int? result = op1 * op2;
```

The answer is that for all simple nullable types other than `bool?`, the result of the operation is `null`, which you can interpret as “unable to compute.” For structs you can define your own operators to deal with this situation (as shown later in this chapter), and for `bool?` there are operators defined for `&` and `|` that might result in non-`null` return values. The results in the table make perfect sense logically — if there is enough information to work out the answer of the computation without needing to know the value of one of the operands, then it doesn’t matter if that operand is `null`.

## The `??` Operator

To further reduce the amount of code you need in order to deal with nullable types, and to make it easier to deal with variables that can be `null`, you can use the `??` operator. Known as the *null coalescing operator*, it is a binary operator that enables you to supply an alternative value to use for expressions that might evaluate to `null`. The operator evaluates to its first operand if the first operand is not `null`, or to its second operand if the first operand is `null`. Functionally, the following two expressions are equivalent:

```
op1 ?? op2
op1 == null ? op2 : op1
```

In this code, `op1` can be any nullable expression, including a reference type and, importantly, a nullable type. This means that you can use the `??` operator to provide default values to use if a nullable type is `null`, as shown here:

```
int? op1 = null;
int result = op1 * 2 ?? 5;
```

Because in this example `op1` is `null`, `op1 * 2` will also be `null`. However, the `??` operator detects this and assigns the value 5 to `result`. Importantly, note here that no explicit conversion is required

to put the result in the `int` type variable `result`. The `??` operator handles this conversion for you. Alternatively, you can pass the result of a `??` evaluation into an `int?` with no problems:

```
int? result = op1 * 2 ?? 5;
```

This behavior makes the `??` operator a versatile one to use when dealing with nullable variables, and a handy way to supply defaults without using either a block of code in an `if` structure or the often confusing ternary operator.

## The `?.` Operator

This operator, often referred to as the *Elvis operator* or the *null condition operator*, helps to overcome code ambiguity caused by burdensome null checking. For example, if you wanted to get the count of orders for a given customer, you would need to check for null before setting the count value:

```
int count = 0;
if (customer.orders != null)
{
    count = customer.orders.Count();
}
```

If you were to simply write this code and there were no orders existing for the customer (i.e. it's `null`), a `System.ArgumentNullException` is thrown:

```
int count = customer.orders.Count();
```

Using the `?.` operator results in the `int? count` being set to `null` instead of an exception happening.

```
int? count = customer.orders?.Count();
```

Combining the null coalescing operator `??` discussed in the previous section with the null condition operator `?.` makes it possible to set a default value when the result is `null`.

```
int? count = customer.orders?.Count() ?? 0;
```

Another use of the null conditional operator is to trigger events. Events are discussed in detail in Chapter 13. The most common way to trigger an event is by using this code pattern:

```
var onChanged = OnChanged;
if (onChanged != null)
{
    onChanged(this, args);
}
```

This pattern is not thread safe because someone might unsubscribe the last event handler just after the null check is done. When that happens an exception is thrown and the application crashes. Avoid this by using the null condition operator as shown here:

```
OnChanged?.Invoke(this, args);
```

**NOTE** If you utilize operator overload methods (for example, the == ) without checking for nulls, you receive a System.NullReferenceException.

As mentioned in Chapter 11, use the ?. operator to check for nulls with the == operator overload in the C:\BegVCSharp\Chapter12\Ch12CardLib\Card.cs class to prevent an exception from being thrown when using the method. For example:

```
public static bool operator ==(Card card1, Card card2)
    => (card1?.suit == card2?.suit) && (card1?.rank == card2?.rank);
```

By including the null condition operator in the statement, you are effectively expressing that if the object to the left is not null, (in this case card1 or card2), then retrieve what is to the right. If the object on the left is null (i.e. card1 or card2), then terminate the access chain and return null.

## Working with Nullable Types

Use the following Try It Out to experiment with a nullable vector type.

### TRY IT OUT Nullable Types: Ch12Ex01

1. Create a new console application project called Ch12Ex01 and save it in the directory C:\BegVCSharp\Chapter12.
2. Add a new class called Vector in the file Vector.cs.
3. Modify the code in Vector.cs as follows:

```
using static System.Math;
public class Vector
{
    public double? R = null;
    public double? Theta = null;
    public double? ThetaRadians
    {
        // Convert degrees to radians.
        get { return (Theta * Math.PI / 180.0); }
    }
    public Vector(double? r, double? theta)
    {
        // Normalize.
        if (r < 0)
        {
            r = -r;
            theta += 180;
        }
        theta = theta % 360;
        // Assign fields.
        R = r;
    }
}
```

```
        Theta = theta;
    }
    public static Vector operator +(Vector op1, Vector op2)
    {
        try
        {
            // Get (x, y) coordinates for new vector.
            double newX = op1.R.Value * Sin(op1.ThetaRadians.Value)
                + op2.R.Value * Sin(op2.ThetaRadians.Value);
            double newY = op1.R.Value * Cos(op1.ThetaRadians.Value)
                + op2.R.Value * Cos(op2.ThetaRadians.Value);
            // Convert to (r, theta).
            double newR = Sqrt(newX * newX + newY * newY);
            double newTheta = Atan2(newX, newY) * 180.0 / PI;
            // Return result.
            return new Vector(newR, newTheta);
        }
        catch
        {
            // Return "null" vector.
            return new Vector(null, null);
        }
    }
    public static Vector operator -(Vector op1) => new Vector(-op1.R, op1.Theta);
    public static Vector operator -(Vector op1, Vector op2) => op1 + (-op2);
    public override string ToString()
    {
        // Get string representation of coordinates.
        string rString = R.HasValue ? R.ToString() : "null";
        string thetaString = Theta.HasValue ? Theta.ToString() : "null";
        // Return (r, theta) string.
        return string.Format($"{rString}, {thetaString}");
    }
}
```

4. Modify the code in Program.cs as follows:

```
class Program
{
    static void Main(string[] args)
    {
        Vector v1 = GetVector("vector1");
        Vector v2 = GetVector("vector1");
        WriteLine($"{v1} + {v2} = {v1 + v2}");
        WriteLine($"{v1} - {v2} = {v1 - v2}");
        ReadKey();
    }
    static Vector GetVector(string name)
    {
        WriteLine($"Input {name} magnitude:");
    }
}
```

```

        double? r = GetNullableDouble();
        WriteLine($"Input {name} angle (in degrees):");
        double? theta = GetNullableDouble();
        return new Vector(r, theta);
    }
    static double? GetNullableDouble()
    {
        double? result;
        string userInput = ReadLine();
        try
        {
            result = double.Parse(userInput);
        }
        catch
        {
            result = null;
        }
        return result;
    }
}

```

5. Execute the application and enter values for two vectors. The sample output is shown in Figure 12-1.




FIGURE 12-1

6. Execute the application again, but this time skip at least one of the four values. The sample output is shown in Figure 12-2.




FIGURE 12-2

## How It Works

This example created a class called `Vector` that represents a vector with polar coordinates (that is, with a magnitude and an angle), as shown in Figure 12-3.

The coordinates  $r$  and  $\theta$  are represented in code by the public fields `R` and `Theta`, where `Theta` is expressed in degrees. `ThetaRadians` is supplied to obtain the value of `Theta` in radians — this is necessary because the `Math` class uses radians in its static methods. Both `R` and `Theta` are of type `double?`, so they can be null:

```
public class Vector
{
    public double? R = null;
    public double? Theta = null;
    public double? ThetaRadians
    {
        get
        {
            // Convert degrees to radians.
            return (Theta * PI / 180.0);
        }
    }
}
```

The constructor for `Vector` normalizes the initial values of `R` and `Theta` and then assigns the public fields:

```
public Vector(double? r, double? theta)
{
    // Normalize.
    if (r < 0)
    {
        r = -r;
        theta += 180;
    }
    theta = theta % 360;
    // Assign fields.
    R = r;
    Theta = theta;
}
```

The main functionality of the `Vector` class is to add and subtract vectors using operator overloading, which requires some fairly basic trigonometry not covered here. You might consider taking a look at this site <http://www.onlinemathlearning.com/basic-trigonometry.html>, or search for other resources on the Internet. The important point about the code is that if an exception is thrown when obtaining the `value` property of `R` or `ThetaRadians` — that is, if either is `null` — then a “null” vector is returned:

```
public static Vector operator +(Vector op1, Vector op2)
{
    try
    {
        // Get (x, y) coordinates for new vector.
        ...
    }
```




FIGURE 12-3

```

        catch
    {
        // Return "null" vector.
        return new Vector(null, null);
    }
}

```

If either of the coordinates making up a vector is `null`, then the vector is invalid, which is signified here by a `Vector` class with `null` values for both `R` and `Theta`. The rest of the code in the `Vector` class overrides the other operators required to extend the addition functionality to include subtraction, and overrides `ToString()` to obtain a string representation of a `Vector` object.

The code in `Program.cs` tests the `Vector` class by enabling the user to initialize two vectors, and then adds and subtracts them to and from one another. Should the user omit a value, it will be interpreted as `null`, and the rules mentioned previously apply.

---

## The System.Collections.Generic Namespace

In practically every application used so far in this book, you have seen the following namespaces:

```

using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;

```

The `System` namespace contains most of the basic types used in .NET applications. The `System.Text` namespace includes types relating to string processing and encoding. You'll look at the `System.Linq` namespace later in this book. The `System.Threading.Tasks` namespace contains types that help you to write asynchronous code, which isn't covered in this book. But what about `System.Collections.Generic`, and why is it included by default in console applications?

The answer is that this namespace contains generic types for dealing with collections, and it is likely to be used so often that it is configured with a `using` statement, ready for you to use without qualification.

You'll now look at these types, which are guaranteed to make your life easier. They make it possible for you to create strongly typed collection classes with hardly any effort. Table 12-1 lists two types from the `System.Collections.Generic` namespace that are covered in this section. More of the types in this namespace are covered later in this chapter.

**TABLE 12-1:** Generic Collection Type

TYPE	DESCRIPTION
<code>List&lt;T&gt;</code>	Collection of type <code>T</code> objects
<code>Dictionary&lt;K, V&gt;</code>	Collection of items of type <code>V</code> , associated with keys of type <code>K</code>

This section also describes various interfaces and delegates used with these classes.

## List<T>

Rather than derive a class from `CollectionBase` and implement the required methods as you did in the last chapter, it can be quicker and easier simply to use the `List<T>` generic collection type. An added bonus here is that many of the methods you normally have to implement, such as `Add()`, are implemented for you.

Creating a collection of type `T` objects requires the following code:

```
List<T> myCollection = new List<T>();
```

That's it. You don't have to define any classes, implement any methods, or do anything else. You can also set a starting list of items in the collection by passing a `List<T>` object to the constructor. `List<T>` also has an `Item` property, enabling array-like access:

```
T itemAtIndex2 = myCollectionOfT[2];
```

This class supports several other methods, but that's plenty to get you started. The following Try It Out demonstrates how to use `List<T>` in practice.

### TRY IT OUT Using List<T>: Ch12Ex02

1. Create a new console application called `Ch12Ex02` and save it in the directory `C:\BegVCSharp\Chapter12`.
2. Right-click on the project name in the Solution Explorer window and select the `Add`  $\Rightarrow$  `Existing Item` option.
3. Select the `Animal.cs`, `Cow.cs`, and `Chicken.cs` files from the `C:\BegVCSharp\Chapter11\Ch11Ex01\Ch11Ex01` directory and click `Add`.
4. Modify the namespace declaration in the three files you added as follows:

```
namespace Ch12Ex02
```

5. Modify `Program.cs` as follows:

```
static void Main(string[] args)
{
    List<Animal> animalCollection = new List<Animal>();
    animalCollection.Add(new Cow("Rual"));
    animalCollection.Add(new Chicken("Donna"));
    foreach (Animal myAnimal in animalCollection)
    {
        myAnimal.Feed();
    }
    ReadKey();
}
```

- 
6. Execute the application. The result is exactly the same as the result for Ch11Ex02 in the last chapter.

### How It Works

There are only two differences between this example and Ch11Ex02. The first is that the line of code

```
Animals animalCollection = new Animals();
```

has been replaced with:

```
List<Animal> animalCollection = new List<Animal>();
```

The second, and more crucial, difference is that there is no longer an `Animals` collection class in the project. All that hard work you did earlier to create this class was achieved in a single line of code by using a generic collection class.

An alternative way to get the same result is to leave the code in `Program.cs` as it was in the last chapter, and use the following definition of `Animals`:

```
public class Animals : List<Animal> {}
```

Doing this has the advantage that the code in `Program.cs` is slightly easier to read, plus you can add members to the `Animals` class as you see fit.

---

## Sorting and Searching Generic Lists

Sorting a generic list is much the same as sorting any other list. The last chapter described how you can use the `IComparer` and `IComparable` interfaces to compare two objects and thereby sort a list of that type of object. The only difference here is that you can use the generic interfaces `IComparer<T>` and `IComparable<T>`, which expose slightly different, type-specific methods. Table 12-2 explains these differences.

**TABLE 12-2:** Sorting with Generic Types

GENERIC METHOD	NONGENERIC METHOD	DIFFERENCE
<code>Int IComparable&lt;T&gt;</code> <code>.CompareTo(T otherObj)</code>	<code>int IComparable</code> <code>.CompareTo(object otherObj)</code>	Strongly typed in generic versions.
<code>Bool IComparable&lt;T&gt;</code> <code>.Equals(T otherObj)</code>	N/A	Doesn't exist on a nongeneric interface; can use inherited <code>object.Equals()</code> instead.

*continues*

TABLE 12-2 (continued)

GENERIC METHOD	NONGENERIC METHOD	DIFFERENCE
<code>Int IComparer&lt;T&gt;.Compare(T objectA, T objectB)</code>	<code>int IComparer.Compare(object objectA, object objectB)</code>	Strongly typed in generic versions.
<code>Bool IComparer&lt;T&gt;.Equals(T objectA, T objectB)</code>	N/A	Doesn't exist on a nongeneric interface; can use inherited <code>object.Equals()</code> instead.
<code>Int IComparer&lt;T&gt;.GetHashCode(T objectA)</code>	N/A	Doesn't exist on a nongeneric interface; can use inherited <code>object.GetHashCode()</code> instead.

To sort a `List<T>`, you can supply an `IComparable<T>` interface on the type to be sorted, or supply an `IComparer<T>` interface. Alternatively, you can supply a *generic delegate* as a sorting method. From the perspective of seeing how the code works, this is far more interesting because implementing the interfaces described here takes no more effort than implementing their nongeneric cousins.

In general terms, all you need to sort a list is a method that compares two objects of type `T`; and to search, all you need is a method that checks an object of type `T` to determine whether it meets certain criteria. It is a simple matter to define such methods, and to aid you there are two generic delegate types that you can use:

- `Comparison<T>` — A delegate type for a method used for sorting, with the following return type and parameters:  
`int method(T objectA, T objectB)`
- `Predicate<T>` — A delegate type for a method used for searching, with the following return type and parameters:  
`bool method(T targetObject)`

You can define any number of such methods, and use them to “snap-in” to the searching and sorting methods of `List<T>`. The next Try It Out illustrates this technique.

### TRY IT OUT Sorting and Searching `List<T>`: Ch12Ex03

1. Create a new console application called `Ch12Ex03` and save it in the directory `C:\BegVCSharp\Chapter12`.
2. Right-click on the project name in the Solution Explorer window and select the Add Existing Item option.
3. Select the `Vector.cs` file from the `C:\BegVCSharp\Chapter12\Ch12Ex01\Ch12Ex01` directory and click Add.

- 
4. Modify the namespace declaration in the file you added as follows:

```
namespace Ch12Ex03
```

5. Add a new class called `Vectors`.

6. Modify `Vectors.cs` as follows:

```
public class Vectors : List<Vector>
{
    public Vectors()
    {
    }
    public Vectors(IEnumerable<Vector> initialItems)
    {
        foreach (Vector vector in initialItems)
        {
            Add(vector);
        }
    }
    public string Sum()
    {
        StringBuilder sb = new StringBuilder();
        Vector currentPoint = new Vector(0.0, 0.0);
        sb.Append("origin");
        foreach (Vector vector in this)
        {
            sb.AppendFormat($" + {vector}");
            currentPoint += vector;
        }
        sb.AppendFormat($" = {currentPoint}");
        return sb.ToString();
    }
}
```

7. Add a new class called `VectorDelegates`.

8. Modify `VectorDelegates.cs` as follows:

```
public static class VectorDelegates
{
    public static int Compare(Vector x, Vector y)
    {
        if (x.R > y.R)
        {
            return 1;
        }
        else if (x.R < y.R)
        {
            return -1;
        }
        return 0;
    }
}
```

```

public static bool TopRightQuadrant(Vector target)
{
    if (target.Theta >= 0.0 && target.Theta <= 90.0)
    {
        return true;
    }
    else
    {
        return false;
    }
}

```

9. Modify Program.cs as follows:

```

static void Main(string[] args)
{
    Vectors route = new Vectors();
    route.Add(new Vector(2.0, 90.0));
    route.Add(new Vector(1.0, 180.0));
    route.Add(new Vector(0.5, 45.0));
    route.Add(new Vector(2.5, 315.0));
    WriteLine(route.Sum());
    Comparison<Vector> sorter = new Comparison<Vector>(
        VectorDelegates.Compare);
    route.Sort(sorter);
    WriteLine(route.Sum());
    Predicate<Vector> searcher =
        new Predicate<Vector>(VectorDelegates.TopRightQuadrant);
    Vectors topRightQuadrantRoute = new Vectors(route.FindAll(searcher));
    WriteLine(topRightQuadrantRoute.Sum());
    ReadKey();
}

```

10. Execute the application. The result is shown in Figure 12-4.




FIGURE 12-4

### How It Works

In this example, you created a collection class, `Vectors`, for the `Vector` class created in Ch12Ex01. You could just use a variable of type `List<Vector>`, but because you want additional functionality you use a new class, `Vectors`, and derive from `List<Vector>`, which enables you to add whatever additional members you want.

One member, `Sum()`, returns a string listing each vector in turn, along with the result of summing them all together (using the overloaded `+` operator from the original `Vector` class). Because each vector can be thought of as a direction and a distance, this effectively constitutes a route with an endpoint:

```
public string Sum()
{
    StringBuilder sb = new StringBuilder();
    Vector currentPoint = new Vector(0.0, 0.0);
    sb.Append("origin");
    foreach (Vector vector in this)
    {
        sb.AppendFormat(" + {vector}");
        currentPoint += vector;
    }
    sb.AppendFormat(" = {currentPoint}");
    return sb.ToString();
}
```

This method uses the handy `StringBuilder` class, found in the `System.Text` namespace, to build the response string. This class has members such as `Append()` and `AppendFormat()` (used here), which make it easy to assemble a string — the performance is better than concatenating individual strings. You use the `ToString()` method of this class to obtain the resultant string.

You also create two methods to be used as delegates, as static members of `VectorDelegates`. `Compare()` is used for comparison (sorting), and `TopRightQuadrant()` for searching. You'll look at these as you review the code in `Program.cs`.

The code in `Main()` starts with the initialization of a `Vectors` collection, to which are added several `Vector` objects (you can find this code in `Ch12Ex03\Program.cs`):

```
Vectors route = new Vectors();
route.Add(new Vector(2.0, 90.0));
route.Add(new Vector(1.0, 180.0));
route.Add(new Vector(0.5, 45.0));
route.Add(new Vector(2.5, 315.0));
```

The `Vectors.Sum()` method is used to write out the items in the collection as noted earlier, this time in their initial order:

```
WriteLine(route.Sum());
```

Next, you create the first of your delegates, `sorter`. This delegate is of type `Comparison<Vector>` and, therefore, can be assigned a method with the following return type and parameters:

```
int method(Vector objectA, Vector objectB)
```

This matches `VectorDelegates.Compare()`, which is the method you assign to the delegate:

```
Comparison<Vector> sorter = new Comparison<Vector>(
    VectorDelegates.Compare);
```

Compare() compares the magnitudes of two vectors as follows:

```
public static int Compare(Vector x, Vector y)
{
    if (x.R > y.R)
    {
        return 1;
    }
    else if (x.R < y.R)
    {
        return -1;
    }
    return 0;
}
```

This enables you to order the vectors by magnitude:

```
route.Sort(sorter);
WriteLine(route.Sum());
```

The output of the application gives the result you'd expect — the result of the summation is the same because the endpoint of following the “vector route” is the same regardless of the order in which you carry out the individual steps.

Next, you obtain a subset of the vectors in the collection by searching. This uses VectorDelegates.TopRightQuadrant():

```
public static bool TopRightQuadrant(Vector target)
{
    if (target.Theta >= 0.0 && target.Theta <= 90.0)
    {
        return true;
    }
    else
    {
        return false;
    }
}
```

This method returns true if its Vector argument has a value of Theta between 0 and 90 degrees — that is, if it points up and/or right in a diagram of the sort shown earlier.

In the Main() method, you use this method via a delegate of type Predicate<Vector> as follows:

```
Predicate<Vector> searcher =
    new Predicate<Vector>(VectorDelegates.TopRightQuadrant);
Vectors topRightQuadrantRoute = new Vectors(route.FindAll(searcher));
WriteLine(topRightQuadrantRoute.Sum());
```

This requires the constructor defined in Vectors:

```
public Vectors(IEnumerable<Vector> initialItems)
{
    foreach (Vector vector in initialItems)
    {
```

```

        Add(vector);
    }
}

```

Here, you initialize a new `Vectors` collection using an interface of `IEnumerable<Vector>`, which is necessary because `List<Vector>.FindAll()` returns a `List<Vector>` instance, not a `Vectors` instance.

The result of the searching is that only a subset of `vector` objects is returned, so (again, as you'd expect) the result of the summation is different. The use of these generic delegate types to sort and search generic collections can take a little while to get used to, but the result is code that is streamlined and efficient, and which has a highly logical structure. It is well worth investing the time to learn the techniques presented in this section.

As an aside to this example, note that the code,

```

Comparison<Vector> sorter = new Comparison<Vector>(
    VectorDelegates.Compare);
route.Sort(sorter);

```

can be simplified to the following:

```
route.Sort(VectorDelegates.Compare);
```

This removes the necessity to implicitly reference the `Comparison<Vector>` type. In fact, an instance of this type is still created, but it is created implicitly. The `Sort()` method obviously needs an instance of this type to work, but the compiler realizes this and creates one for you from the method that you supply. In this situation, the reference to `VectorDelegates.Compare()` (without the parentheses) is referred to as a *method group*. There are many situations in which you can use method groups to implicitly create delegates in this way, which can make your code more readable.

## Dictionary<K, V>

The `Dictionary<K, V>` type enables you to define a collection of key-value pairs. Unlike the other generic collection types you've looked at in this chapter, this class requires instantiating two types: the types for both the key and the value that represent each item in the collection.

Once a `Dictionary<K, V>` object is instantiated, you can perform much the same operations on it as you can on a class that inherits from `DictionaryBase`, but with type-safe methods and properties already in place. You can, for example, add key-value pairs using a strongly typed `Add()` method:

```

Dictionary<string, int> things = new Dictionary<string, int>();
things.Add("Green Things", 29);
things.Add("Blue Things", 94);
things.Add("Yellow Things", 34);
things.Add("Red Things", 52);
things.Add("Brown Things", 27);

```

You can iterate through keys and values in the collection by using the `Keys` and `Values` properties:

```
foreach (string key in things.Keys)
{
    WriteLine(key);
}
foreach (int value in things.Values)
{
    WriteLine(value);
}
```

In addition, you can iterate through items in the collection by obtaining each as a `KeyValuePair<K, V>` instance, much like you can with the `DictionaryEntry` objects shown in the last chapter:

```
foreach (KeyValuePair<string, int> thing in things)
{
    WriteLine($"{thing.Key} = {thing.Value}");
}
```

One point to note about `Dictionary<K, V>` is that the key for each item must be unique. Attempting to add an item with an identical key will cause an `ArgumentException` exception to be thrown.

Because of this, `Dictionary<K, V>` allows you to pass an `IComparer<K>` interface to its constructor. This might be necessary if you use your own classes as keys and they don't support an `IComparable` or `IComparable<K>` interface, or if you want to compare objects using a nondefault process. For instance, in the preceding example, you could use a case-insensitive method to compare string keys:

```
Dictionary<string, int> things =
    new Dictionary<string, int>(StringComparer.CurrentCultureIgnoreCase);
```

Now you'll get an exception if you use keys such as this:

```
things.Add("Green Things", 29);
things.Add("Green things", 94);
```

You can also pass an initial capacity (with an `int`) or a set of items (with an `IDictionary<K, V>` interface) to the constructor.

A feature introduced in C# 6 called *index initializers* supports the initialization of indices inside the object initializer:

```
var zahlen = new Dictionary<int, string>()
{
    [1] = "eins",
    [2] = "zwei"
};
```

Index initializers can be streamlined as in many cases there is no need for a temporary variable as show previously via `var zahlen`. Using *expression-bodied methods*, the above example leads to a cascading effect of simplification:

```
public ZObject ToGerman() => new ZObject() { [1] = "eins", [2] = "zwei" };
```

## Modifying CardLib to Use a Generic Collection Class

One simple modification you can make to the `CardLib` project you've been building over recent chapters is to change the `Cards` collection class to use a generic collection class, thus saving many lines of code. The required modification to the class definition for `Cards` is as follows (you can find this code in `Ch12CardLib\Cards.cs`):

```
public class Cards : List<Card>, ICloneable { ... }
```

You can also remove all the methods of `Cards` except `Clone()`, which is required for `ICloneable`, and `CopyTo()`, because the version of `CopyTo()` supplied by `List<Card>` works with an array of `Card` objects, not a `Cards` collection. `Clone()` requires a minor modification because the `List<T>` class does not define a `List` property to use:

```
public object Clone()
{
    Cards newCards = new Cards();
    foreach (Card sourceCard in this)
    {
        newCards.Add((Card)sourceCard.Clone());
    }
    return newCards;
}
```

Rather than show the code here for what is a very simple modification, the updated version of `CardLib`, called `Ch12CardLib`, is included in the downloadable code for this chapter, along with the client code from the last chapter.

## DEFINING GENERIC TYPES

You've now learned enough about generics to create your own. You've seen plenty of code involving generic types and have had plenty of practice using generic syntax. This section looks at defining the following:

- Generic classes
- Generic interfaces
- Generic methods
- Generic delegates

You'll also look at the following more advanced techniques for dealing with the issues that come up when defining generic types:

- The `default` keyword
- Constraining types
- Inheriting from generic classes
- Generic operators

## Defining Generic Classes

To create a generic class, merely include the angle bracket syntax in the class definition:

```
class MyGenericClass<T> { ... }
```

Here, T can be any identifier you like, following the usual C# naming rules, such as not starting with a number and so on. Typically, though, you can just use T. A generic class can have any number of type parameters in its definition, separated by commas:

```
class MyGenericClass<T1, T2, T3> { ... }
```

Once these types are defined, you can use them in the class definition just like any other type. You can use them as types for member variables, return types for members such as properties or methods, and parameter types for method arguments:

```
class MyGenericClass<T1, T2, T3>
{
    private T1 innerT1Object;
    public MyGenericClass(T1 item)
    {
        innerT1Object = item;
    }
    public T1 InnerT1Object
    {
        get { return innerT1Object; }
    }
}
```

Here, an object of type T1 can be passed to the constructor, and read-only access is permitted to this object via the property InnerT1Object. Note that you can make practically no assumptions as to what the types supplied to the class are. The following code, for example, will not compile:

```
class MyGenericClass<T1, T2, T3>
{
    private T1 innerT1Object;
    public MyGenericClass()
    {
        innerT1Object = new T1();
    }
    public T1 InnerT1Object
    {
        get { return innerT1Object; }
    }
}
```

Because you don't know what T1 is, you can't use any of its constructors — it might not even have any, or it might have no publicly accessible default constructor. Without more complicated code involving the techniques shown later in this section, you can make only the following assumption about T1: you can treat it as a type that either inherits from or can be boxed into System.Object.

Obviously, this means that you can't really do anything very interesting with instances of this type, or any of the other types supplied to the generic class `MyGenericClass`. Without using *reflection*, which is an advanced technique used to examine types at runtime (and not covered in this chapter), you're limited to code that's no more complicated than the following:

```
public string GetAllTypesAsString()
{
    return "T1 = " + typeof(T1).ToString()
        + ", T2 = " + typeof(T2).ToString()
        + ", T3 = " + typeof(T3).ToString();
}
```

There is a bit more that you can do, particularly in terms of collections, because dealing with groups of objects is a pretty simple process and doesn't need any assumptions about the object types — which is one good reason why the generic collection classes you've seen in this chapter exist.

Another limitation that you need to be aware of is that using the operator `==` or `!=` is permitted only when comparing a value of a type supplied to a generic type to `null`. That is, the following code works fine:

```
public bool Compare(T1 op1, T1 op2)
{
    if (op1 != null && op2 != null)
    {
        return true;
    }
    else
    {
        return false;
    }
}
```

Here, if `T1` is a value type, then it is always assumed to be non-`null`, so in the preceding code `Compare` will always return `true`. However, attempting to compare the two arguments `op1` and `op2` fails to compile:

```
public bool Compare(T1 op1, T1 op2)
{
    if (op1 == op2)
    {
        return true;
    }
    else
    {
        return false;
    }
}
```

That's because this code assumes that `T1` supports the `==` operator. In short, to do anything interesting with generics, you need to know a bit more about the types used in the class.

## The default Keyword

One of the most basic things you might want to know about types used to create generic class instances is whether they are reference types or value types. Without knowing this, you can't even assign null values with code such as this:

```
public MyGenericClass()
{
    innerT1Object = null;
}
```

If T1 is a value type, then innerT1Object can't have the value null, so this code won't compile. Luckily, this problem has been addressed, resulting in a new use for the `default` keyword (which you've seen being used in `switch` structures earlier in the book). This is used as follows:

```
public MyGenericClass()
{
    innerT1Object = default(T1);
}
```

The result of this is that innerT1Object is assigned a value of null if it is a reference type, or a default value if it is a value type. This default value is 0 for numeric types, while structs have each of their members initialized to 0 or null in the same way. The `default` keyword gets you a bit further in terms of doing a little more with the types you are forced to use, but to truly get ahead, you need to constrain the types that are supplied.

## Constraining Types

The types you have used with generic classes until now are known as *unbounded* types because no restrictions are placed on what they can be. By *constraining* types, it is possible to restrict the types that can be used to instantiate a generic class. There are a number of ways to do this. For example, it's possible to restrict a type to one that inherits from a certain type. Referring back to the `Animal`, `Cow`, and `Chicken` classes used earlier, you could restrict a type to one that was or inherited from `Animal`, so this code would be fine:

```
MyGenericClass<Cow> = new MyGenericClass<Cow>();
```

The following, however, would fail to compile:

```
MyGenericClass<string> = new MyGenericClass<string>();
```

In your class definitions this is achieved using the `where` keyword:

```
class MyGenericClass<T> where T : constraint { ... }
```

Here, `constraint` defines what the constraint is. You can supply a number of constraints in this way by separating them with commas:

```
class MyGenericClass<T> where T : constraint1, constraint2 { ... }
```

You can define constraints on any or all of the types required by the generic class by using multiple `where` statements:

```
class MyGenericClass<T1, T2> where T1 : constraint1 where T2 : constraint2
{ ... }
```

Any constraints that you use must appear after the inheritance specifiers:

```
class MyGenericClass<T1, T2> : MyBaseClass, IMyInterface
    where T1 : constraint1 where T2 : constraint2 { ... }
```

The available constraints are shown in Table 12-3.

TABLE 12-3: Generic Type Constraints

CONSTRAINT	DEFINITION	EXAMPLE USAGE
<code>struct</code>	Type must be a value type.	In a class that requires value types to function — for example, where a member variable of type <code>T</code> being <code>0</code> means something
<code>class</code>	Type must be a reference type.	In a class that requires reference types to function — for example, where a member variable of type <code>T</code> being <code>null</code> means something
<code>base-class</code>	Type must be, or inherit from, <code>base-class</code> . You can supply any class name as this constraint.	In a class that requires certain baseline functionality inherited from <code>base-class</code> in order to function
<code>interface</code>	Type must be, or implement, <code>interface</code> .	In a class that requires certain baseline functionality exposed by <code>interface</code> in order to function
<code>new()</code>	Type must have a public, parameterless constructor.	In a class where you need to be able to instantiate variables of type <code>T</code> , perhaps in a constructor

**NOTE** If `new()` is used as a constraint, it must be the last constraint specified for a type.

It is possible to use one type parameter as a constraint on another through the base-class constraint as follows:

```
class MyGenericClass<T1, T2> where T2 : T1 { ... }
```

Here, T2 must be the same type as T1 or inherit from T1. This is known as a *naked type constraint*, meaning that one generic type parameter is used as a constraint on another.

Circular type constraints, as shown here, are forbidden:

```
class MyGenericClass<T1, T2> where T2 : T1 where T1 : T2 { ... }
```

This code will not compile. In the following Try It Out, you'll define and use a generic class that uses the Animal family of classes shown in earlier chapters.

### TRY IT OUT Defining a Generic Class: Ch12Ex04

1. Create a new console application called Ch12Ex04 and save it in the directory C:\BegVCSharp\Chapter12.
2. Right-click on the project name in the Solution Explorer window and select the Add Existing Item option.
3. Select the Animal.cs, Cow.cs, and Chicken.cs files from the C:\BegVCSharp\Chapter12\Ch12Ex02\Ch12Ex02 directory and click Add.
4. Modify the namespace declaration in the file you have added as follows:

```
namespace Ch12Ex04
```

5. Modify Animal.cs as follows:

```
public abstract class Animal
{
    ...
    public abstract void MakeANoise();
}
```

6. Modify Chicken.cs as follows:

```
public class Chicken : Animal
{
    ...
    public override void MakeANoise()
    {
        WriteLine($"{name} says 'cluck!');");
    }
}
```

7. Modify Cow.cs as follows:

```
public class Cow : Animal
{
    ...
    public override void MakeANoise()
    {
        WriteLine($"{name} says 'moo!');");
    }
}
```

8. Add a new class called SuperCow and modify the code in SuperCow.cs as follows:

```
public class SuperCow : Cow
{
    public void Fly()
    {
        WriteLine($"{name} is flying!");
    }
    public SuperCow(string newName) : base(newName)
    {
    }
    public override void MakeANoise()
    {
        WriteLine(
            $"{name} says 'here I come to save the day!'");
    }
}
```

9. Add a new class called Farm and modify the code in Farm.cs as follows:

```
using System;
using System.Collections;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
namespace Ch12Ex04
{
    public class Farm<T> : IEnumerable<T>
        where T : Animal
    {
        private List<T> animals = new List<T>();
        public List<T> Animals
        {
            get { return animals; }
        }
        public IEnumerator<T> GetEnumerator() => animals.GetEnumerator();
        IEnumerable.GetEnumerator() => animals.GetEnumerator();
        public void MakeNoises()
        {
            foreach (T animal in animals)
            {
                animal.MakeANoise();
            }
        }
        public void FeedTheAnimals()
        {
            foreach (T animal in animals)
            {
                animal.Feed();
            }
        }
        public Farm<Cow> GetCows()
        {
            Farm<Cow> cowFarm = new Farm<Cow>();
```

```

        foreach (T animal in animals)
        {
            if (animal is Cow)
            {
                cowFarm.Animals.Add(animal as Cow);
            }
        }
        return cowFarm;
    }
}
}

```

10. Modify Program.cs as follows:

```

static void Main(string[] args)
{
    Farm<Animal> farm = new Farm<Animal>();
    farm.Animals.Add(new Cow("Rual"));
    farm.Animals.Add(new Chicken("Donna"));
    farm.Animals.Add(new Chicken("Mary"));
    farm.Animals.Add(new SuperCow("Ben"));
    farm.MakeNoises();
    Farm<Cow> dairyFarm = farm.GetCows();
    dairyFarm.FeedTheAnimals();
    foreach (Cow cow in dairyFarm)
    {
        if (cow is SuperCow)
        {
            (cow as SuperCow).Fly();
        }
    }
    ReadKey();
}

```

11. Execute the application. The result is shown in Figure 12-5.




FIGURE 12-5

### How It Works

In this example, you created a generic class called `Farm<T>`, which, rather than inheriting from a generic list class, exposes a generic list class as a public property. The type of this list is determined by the type parameter `T` that is passed to `Farm<T>` and is constrained to be, or inherit from, `Animal`:

```

public class Farm<T> : IEnumerable<T>
    where T : Animal

```

```

{
    private List<T> animals = new List<T>();
    public List<T> Animals
    {
        get { return animals; }
    }
}

```

`Farm<T>` also implements `IEnumerable<T>`, where `T` is passed into this generic interface and is therefore also constrained in the same way. You implement this interface to make it possible to iterate through the items contained in `Farm<T>` without needing to explicitly iterate over `Farm<T>.Animals`. This is simple to achieve: you simply return the enumerator exposed by `Animals`, which is a `List<T>` class that also implements `IEnumerable<T>`:

```
public IEnumerator<T> GetEnumerator() => animals.GetEnumerator();
```

Because `IEnumerable<T>` inherits from `IEnumerable`, you also need to implement `IEnumerable`.`GetEnumerator()`:

```
IEnumerator IEnumerable.GetEnumerator() => animals.GetEnumerator();
```

Next, `Farm<T>` includes two methods that make use of methods of the abstract `Animal` class:

```

public void MakeNoises()
{
    foreach (T animal in animals)
    {
        animal.MakeANoise();
    }
}
public void FeedTheAnimals()
{
    foreach (T animal in animals)
    {
        animal.Feed();
    }
}

```

Because `T` is constrained to `Animal`, this code compiles fine — you are guaranteed to have access to the `MakeANoise()` and `Feed()` methods, whatever type `T` actually is.

The next method, `GetCows()`, is more interesting. This method simply extracts all the items in the collection that are of type `Cow` (or that inherit from `Cow`, such as the new `SuperCow` class):

```

public Farm<Cow> GetCows()
{
    Farm<Cow> cowFarm = new Farm<Cow>();
    foreach (T animal in animals)
    {
        if (animal is Cow)
        {
            cowFarm.Animals.Add(animal as Cow);
        }
    }
    return cowFarm;
}

```

What is interesting here is that this method seems a bit wasteful. If you wanted other methods of the same sort, such as `GetChickens()` and so on, you'd need to implement them explicitly too. In a system with many more types, you'd need many more methods. A far better solution is to use a *generic method*, which you'll implement a little later in the chapter.

The client code in `Program.cs` simply tests the various methods of `Farm` and doesn't contain much you haven't already seen, so there's no need to examine this code in any greater detail — despite the flying cow.

---

## Inheriting from Generic Classes

The `Farm<T>` class in the preceding example, as well as several other classes you've seen in this chapter, inherit from a generic type. In the case of `Farm<T>`, this type was an interface: `IEnumerable<T>`. Here, the constraint on `T` supplied by `Farm<T>` resulted in an additional constraint on `T` used in `IEnumerable<T>`. This can be a useful technique for constraining otherwise unbounded types. However, you do need to follow some rules.

First, you can't "unconstrain" types that are constrained in a type from which you are inheriting. In other words, a type `T` that is used in a type you are inheriting from must be constrained at least as much as it is in that type. For example, the following code is fine:

```
class SuperFarm<T> : Farm<T>
    where T : SuperCow {}
```

This works because `T` is constrained to `Animal` in `Farm<T>`, and constraining it to `SuperCow` is constraining `T` to a subset of these values. However, the following won't compile:

```
class SuperFarm<T> : Farm<T>
    where T : struct{}
```

Here, you can say definitively that the type `T` supplied to `SuperFarm<T>` cannot be converted into a `T` usable by `Farm<T>`, so the code won't compile.

Even situations in which the constraint is a superset have the same problem:

```
class SuperFarm<T> : Farm<T>
    where T : class{}
```

Even though types such as `Animal` would be allowed by `SuperFarm<T>`, other types that satisfy the class constraint won't be allowed in `Farm<T>`. Again, compilation will fail. This rule applies to all the constraint types shown earlier in this chapter.

Also note that if you inherit from a generic type, then you must supply all the required type information, either in the form of other generic type parameters, as shown, or explicitly. This also applies to nongeneric classes that inherit from generic types, as you've seen elsewhere. Here's an example:

```
public class Cards : List<Card>, ICloneable{}
```

This is fine, but attempting the following will fail:

```
public class Cards : List<T>, ICloneable{}
```

Here, no information is supplied for T, so no compilation is possible.

**NOTE** If you supply a parameter to a generic type, as in `List<Card>`, then you can refer to the type as closed. Similarly, inheriting from `List<T>` is inheriting from an open generic type.

## Generic Operators

Operator overrides are implemented in C# just like other methods and can be implemented in generic classes. For example, you could define the following implicit conversion operator in `Farm<T>`:

```
public static implicit operator List<Animal>(Farm<T> farm)
{
    List<Animal> result = new List<Animal>();
    foreach (T animal in farm)
    {
        result.Add(animal);
    }
    return result;
}
```

This allows the `Animal` objects in a `Farm<T>` to be accessed directly as a `List<Animal>` should you require it. This comes in handy if you want to add two `Farm<T>` instances together, such as with the following operators:

```
public static Farm<T> operator +(Farm<T> farm1, List<T> farm2)
{
    Farm<T> result = new Farm<T>();
    foreach (T animal in farm1)
    {
        result.Animals.Add(animal);
    }
    foreach (T animal in farm2)
    {
        if (!result.Animals.Contains(animal))
        {
            result.Animals.Add(animal);
        }
    }
    return result;
}
public static Farm<T> operator +(List<T> farm1, Farm<T> farm2)
=> farm2 + farm1;
```

You could then add instances of `Farm<Animal>` and `Farm<Cow>` as follows:

```
Farm<Animal> newFarm = farm + dairyFarm;
```

In this code, `dairyFarm` (an instance of `Farm<Cow>`) is implicitly converted into `List<Animal>`, which is usable by the overloaded `+` operator in `Farm<T>`.

You might think that this could be achieved simply by using the following:

```
public static Farm<T> operator +(Farm<T> farm1, Farm<T> farm2) { ... }
```

However, because `Farm<Cow>` cannot be converted into `Farm<Animal>`, the summation will fail. To take this a step further, you could solve this using the following conversion operator:

```
public static implicit operator Farm<Animal>(Farm<T> farm)
{
    Farm <Animal> result = new Farm <Animal>();
    foreach (T animal in farm)
    {
        result.Animals.Add(animal);
    }
    return result;
}
```

With this operator, instances of `Farm<T>`, such as `Farm<Cow>`, can be converted into instances of `Farm<Animal>`, solving the problem. You can use either of the methods shown, although the latter is preferable for its simplicity.

## Generic Structs

You learned in earlier chapters that structs are essentially the same as classes, barring some minor differences and the fact that a struct is a value type, not a reference type. Because this is the case, *generic structs* can be created in the same way as generic classes, as shown here:

```
public struct MyStruct<T1, T2>
{
    public T1 item1;
    public T2 item2;
}
```

## Defining Generic Interfaces

You've now seen several generic interfaces in use — namely, those in the `System.Collections.Generic` namespace such as `IEnumerable<T>` used in the last example. Defining a generic interface involves the same techniques as defining a generic class:

```
interface MyFarmingInterface<T>
    where T : Animal
{
    bool AttemptToBreed(T animal1, T animal2);
    T OldestInHerd { get; }
}
```

Here, the generic parameter `T` is used as the type of the two arguments of `AttemptToBreed()` and the type of the `OldestInHerd` property.

The same inheritance rules apply as for classes. If you inherit from a base generic interface, you must obey the rules, such as keeping the constraints of the base interface generic type parameters.

## Defining Generic Methods

The previous Try It Out used a method called `GetCows()`, and in the discussion of the example it was stated that you could make a more general form of this method using a *generic method*. In this section you'll see how this is possible. A generic method is one in which the return and/or parameter types are determined by a generic type parameter or parameters:

```
public T GetDefault<T>() => default(T);
```

This trivial example uses the `default` keyword you looked at earlier in the chapter to return a default value for a type `T`. This method is called as follows:

```
int myDefaultInt = GetDefault<int>();
```

The type parameter `T` is provided at the time the method is called.

This `T` is quite separate from the types used to supply generic type parameters to classes. In fact, generic methods can be implemented by nongeneric classes:

```
public class Defaulter
{
    public T GetDefault<T>() => default(T);
}
```

If the class is generic, though, then you must use different identifiers for generic method types. The following code won't compile:

```
public class Defaulter<T>
{
    public T GetDefault<T>() => default(T);
}
```

The type `T` used by either the method or the class must be renamed.

Constraints can be used by generic method parameters in the same way that they are for classes, and in this case you can make use of any class type parameters:

```
public class Defaulter<T1>
{
    public T2 GetDefault<T2>()
        where T2 : T1
    {
        return default(T2);
    }
}
```

Here, the type `T2` supplied to the method must be the same as, or inherit from, `T1` supplied to the class. This is a common way to constrain generic methods.

In the `Farm<T>` class shown earlier, you could include the following method (included, but commented out, in the downloadable code for Ch12Ex04):

```
public Farm<U> GetSpecies<U>() where U : T
{
```

```
    Farm<U> speciesFarm = new Farm<U>();
    foreach (T animal in animals)
    {
        if (animal is U)
        {
            speciesFarm.Animals.Add(animal as U);
        }
    }
    return speciesFarm;
}
```

This can replace `GetCows()` and any other methods of the same type. The generic type parameter used here, `U`, is constrained by `T`, which is in turn constrained by the `Farm<T>` class to `Animal`. This enables you to treat instances of `T` as instances of `Animal`, should you want to do so.

In the client code for `Ch12Ex04`, in `Program.cs`, using this new method requires one modification:

```
Farm<Cow> dairyFarm = farm.GetSpecies<Cow>();
```

In a similar vein, you could write:

```
Farm<Chicken> poultryFarm = farm.GetSpecies<Chicken>();
```

You can take this same approach with any class that inherits from `Animal`.

Note here that having generic type parameters on a method changes the signature of the method. This means you can have several overloads of a method differing only in generic type parameters, as shown in this example:

```
public void ProcessT<T>(T op1){ ... }
public void ProcessT<T, U>(T op1){ ... }
```

Which method should be used is determined by the amount of generic type parameters specified when the method is called.

## Defining Generic Delegates

The last generic type to consider is the *generic delegate*. You saw these delegates in action earlier in the chapter when you learned how to sort and search generic lists. You used the `Comparison<T>` and `Predicate<T>` delegates, respectively, for this.

Chapter 6 described how to define delegates using the parameters and return type of a method, the `delegate` keyword, and a name for the delegate:

```
public delegate int MyDelegate(int op1, int op2);
```

To define a generic delegate, you simply declare and use one or more generic type parameters:

```
public delegate T1 MyDelegate<T1, T2>(T2 op1, T2 op2) where T1: T2;
```

As you can see, constraints can be applied here too. You'll learn a lot more about delegates in the next chapter, including how you can use them in a common C# programming technique — events.

## VARIANCE

Variance is the collective term for *covariance* and *contravariance*, two concepts that were introduced in .NET 4. In fact, they have been around longer than that (they were available in .NET 2.0), but until .NET 4 it was very difficult to implement them, as this required custom compilation procedures.

The easiest way to grasp what these terms mean is to compare them with polymorphism. Polymorphism, as you will recall, is what enables you to put objects of a derived type into variables of a base type, for example:

```
Cow myCow = new Cow("Geronimo");
Animal myAnimal = myCow;
```

Here, an object of type `Cow` has been placed into a variable of type `Animal` — which is possible because `Cow` derives from `Animal`.

However, the same cannot be said for interfaces. That is to say, the following code will not work:

```
IMethaneProducer<Cow> cowMethaneProducer = myCow;
IMethaneProducer<Animal> animalMethaneProducer = cowMethaneProducer;
```

The first line of code is fine, assuming that `Cow` supports the interface `IMethaneProducer<Cow>`. However, the second line of code presupposes a relationship between the two interface types that doesn't exist, so there is no way of converting one into the other. Or is there? There certainly isn't a way using the techniques you've seen so far in this chapter, as all the type parameters for generic types have been *invariant*. However, it is possible to define variant type parameters on generic interfaces and generic delegates that cater to exactly the situation illustrated in the previous code.

To make the previous code work, the type parameter `T` for the `IMethaneProducer<T>` interface must be *covariant*. Having a covariant type parameter effectively sets up an inheritance relationship between `IMethaneProducer<Cow>` and `IMethaneProducer<Animal>`, so that variables of one type can hold values of the other, just like with polymorphism (although a little more complicated).

To round off this introduction to variance, you need to look at the other kind, *contravariance*. This is similar but works in the other direction. Rather than being able to place a generic interface value into a variable that includes a base type as in covariance, contravariance enables you to place that interface into a variable that uses a derived type, for example:

```
IGrassMuncher<Cow> cowGrassMuncher = myCow;
IGrassMuncher<SuperCow> superCowGrassMuncher = cowGrassMuncher;
```

At first glance this seems a little odd, as you couldn't do the same with polymorphism. However, this is a useful technique in certain circumstances, as you will see in the section called, "Contravariance."

In the next two sections, you look at how to implement variance in generic types and how the .NET Framework uses variance to make your life easier.

**NOTE** All of the code in this section is included in a demo project called VarianceDemo if you want to work through it as you go along.

## Covariance

To define a generic type parameter as covariant, you use the `out` keyword in the type definition, as shown in the following example:

```
public interface IMethaneProducer<out T>{ ... }
```

For interface definitions, covariant type parameters can be used only as return values of methods or property get accessors.

A good example of how this is useful is found in the .NET Framework, in the `IEnumerable<T>` interface that you've used previously. The item type `T` in this interface is defined as being covariant. This means that you can put an object that supports, say, `IEnumerable<Cow>` into a variable of type `IEnumerable<Animal>`.

This enables the following code:

```
static void Main(string[] args)
{
    List<Cow> cows = new List<Cow>();
    cows.Add(new Cow("Geronimo"));
    cows.Add(new SuperCow("Tonto"));
    ListAnimals(cows);
    ReadKey();
}
static void ListAnimals(IEnumerable<Animal> animals)
{
    foreach (Animal animal in animals)
    {
        WriteLine(animal.ToString());
    }
}
```

Here the `cows` variable is of type `List<Cow>`, which supports the `IEnumerable<Cow>` interface. This variable can, through covariance, be passed to a method that expects a parameter of type `IEnumerable<Animal>`. Recalling what you know about how `foreach` loops work, you know that the `GetEnumerator()` method is used to get an enumerator of `IEnumerator<T>`, and the `Current` property of that enumerator is used to access items. `IEnumerator<T>` also defines its type parameter as covariant, which means that it's okay to use it as the get accessor of a parameter, and everything works perfectly.

## Contravariance

To define a generic type parameter as contravariant, you use the `in` keyword in the type definition:

```
public interface IGrassMuncher<in T>{ ... }
```

For interface definitions, contravariant type parameters can be used only as method parameters, not as return types.

Again, the best way to understand this is to look at an example of how contravariance is used in the .NET Framework. One interface that has a contravariant type parameter, again one that you've already used, is `IComparer<T>`. You might implement this interface for animals as follows:

```
public class AnimalNameLengthComparer : IComparer<Animal>
{
    public int Compare(Animal x, Animal y)
        => x.Name.Length.CompareTo(y.Name.Length);
}
```

This comparer compares animals by name length, so you could use it to sort, for example, an instance of `List<Animal>`. However, through contravariance, you can also use it to sort an instance of `List<Cow>`, even though the `List<Cow>.Sort()` method expects an instance of `IComparer<Cow>`:

```
List<Cow> cows = new List<Cow>();
cows.Add(new Cow("Geronimo"));
cows.Add(new SuperCow("Tonto"));
cows.Add(new Cow("Gerald"));
cows.Add(new Cow("Phil"));
cows.Sort(new AnimalNameLengthComparer());
```

In most circumstances, contravariance is something that simply happens — and it's been worked into the .NET Framework to help with just this sort of operation. The good thing about both types of variance in .NET 4 and above, though, is that you can now implement them with the techniques shown in this section whenever you need them.

## EXERCISES

**12.1** Which of the following can be generic?

- a. Classes
- b. Methods
- c. Properties
- d. Operator overloads
- e. Structs
- f. Enumerations

**12.2** Extend the `Vector` class in `Ch12Ex01` such that the `*` operator returns the dot product of two vectors.

**NOTE** The dot product of two vectors is defined as the product of their magnitudes multiplied by the cosine of the angle between them.

**12.3** What is wrong with the following code? Fix it.

```
public class Instantiator<T>
{
    public T instance;
    public Instantiator()
    {
        instance = new T();
    }
}
```

**12.4** What is wrong with the following code? Fix it.

```
public class StringGetter<T>
{
    public string GetString<T>(T item) => item.ToString();
}
```

**12.5** Create a generic class called `ShortList<T>` that implements `IList<T>` and consists of a collection of items with a maximum size. This maximum size should be an integer that can be supplied to the constructor of `ShortList<T>` or defaults to 10. The constructor should also be able to take an initial list of items via an `IEnumerable<T>` parameter. The class should function exactly like `List<T>` but throw an exception of type `IndexOutOfRangeException` if an attempt is made to add too many items to the collection, or if the `IEnumerable<T>` passed to the constructor contains too many items.

**12.6** Will the following code compile? If not, why not?

```
public interface IMethaneProducer<out T>
{
    void BelchAt(T target);
}
```

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU HAVE LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Using generic types	Generic types require one or more type parameters to work. You can use a generic type as the type of a variable by passing the type parameters you require when you declare a variable. You do this by enclosing a comma-separated list of type names in angle brackets.
Nullable types	Nullable types are types that can take any value of a specified value type or the value <code>null</code> . You can use the syntax <code>Nullable&lt;T&gt;</code> or <code>T?</code> to declare a nullable type variable.
The <code>??</code> operator	The null coalescing operator returns either the value of its first operand, or, if the first operand is <code>null</code> , its second operand.
Generic collections	Generic collections are extremely useful as they come with strong typing built-in. You can use <code>List&lt;T&gt;</code> , <code>Collection&lt;T&gt;</code> , and <code>Dictionary&lt;K, V&gt;</code> among other collection types. These also expose generic interfaces. To sort and search generic collections, you use the <code>IComparer&lt;T&gt;</code> and <code>IComparable&lt;T&gt;</code> interfaces.
Defining generic classes	You define a generic type much like any other type, with the addition of generic type parameters where you specify the type name. As with using generic types, you specify these as a comma-separated list enclosed in angle brackets. You can use the generic type parameters in your code anywhere you'd use a type name, for example, in method return values and parameters.
Generic type parameter constraints	In order to use generic type parameters more effectively in your generic type code, you can constrain the types that can be supplied when the type is used. You can constrain type parameters by base class, supported interface, whether they must be value or reference types, and whether they support parameterless constructors. Without such constraints, you must use the <code>default</code> keyword to instantiate a variable of a generic type.
Other generic types	As well as classes, you can define generic interfaces, delegates, and methods.
Variance	Variance is a concept similar to polymorphism, but applied to type parameters. It allows you to use one generic type in place of another, where those generic types vary only in the generic type parameters used. Covariance allows conversion between two types where the target type has a type parameter that is a base class of the type parameter of the source type. Contravariance allows conversion where this relationship is inverted. Covariant type parameters are defined with the <code>out</code> parameter, and can only be used as return types and property get accessor types. Contravariant type parameters are defined with the <code>in</code> parameter and can only be used as method parameters.



# 13

## Additional C# Techniques

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Discovering the :: operator
- Understanding the global namespace qualifier
- Creating custom exceptions
- Using events
- Using anonymous methods
- Using C# attributes
- Working with initializers
- Using the var type and type inference
- Working with anonymous types
- Using the dynamic type
- Using named and optional method parameters
- Working with lambda expressions

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 13 download and individually named according to the names throughout the chapter.

In this chapter, you continue exploring the C# language by looking at a few bits and pieces that haven't quite fit in elsewhere. Anders Hejlsberg (the inventor of C#) and others at Microsoft continue to update and refine the language. At the time of this writing, the most recent changes are part of version 6 of the C# language, which is released as part of the Visual Studio 2015 product line, along with .NET 4.6. At this point in the book, you might be

wondering what else could be needed; indeed, previous versions of C# lack little in terms of functionality. However, this doesn't mean that it isn't possible to make some aspects of C# programming easier, or that the relationships between C# and other technologies can't be streamlined.

You also make some final modifications to the CardLib code that you've been building in the last few chapters, and even use CardLib to create a card game.

## THE :: OPERATOR AND THE GLOBAL NAMESPACE QUALIFIER

The :: operator provides an alternative way to access types in namespaces. This might be necessary if you want to use a namespace alias and there is ambiguity between the alias and the actual namespace hierarchy. If that's the case, then the namespace hierarchy is given priority over the namespace alias. To see what this means, consider the following code:

```
using MyNamespaceAlias = MyRootNamespace.MyNestedNamespace;
namespace MyRootNamespace
{
    namespace MyNamespaceAlias
    {
        public class MyClass {}
    }
    namespace MyNestedNamespace
    {
        public class MyClass {}
    }
}
```

Code in MyRootNamespace might use the following to refer to a class:

```
MyNamespaceAlias.MyClass
```

The class referred to by this code is the `MyRootNamespace.MyNamespaceAlias.MyClass` class, not the `MyRootNamespace.MyNestedNamespace.MyClass` class. That is, the namespace `MyRootNamespace.MyNamespaceAlias` has hidden the alias defined by the `using` statement, which refers to `MyRootNamespace.MyNestedNamespace`. You can still access the `MyRootNamespace.MyNestedNamespace` namespace and the class contained within, but it requires different syntax:

```
MyNestedNamespace.MyClass
```

Alternatively, you can use the :: operator:

```
MyNamespaceAlias::MyClass
```

Using this operator forces the compiler to use the alias defined by the `using` statement, and therefore the code refers to `MyRootNamespace.MyNestedNamespace.MyClass`.

You can also use the keyword `global` with the :: operator, which is essentially an alias to the top-level, root namespace. This can be useful to make it clearer which namespace you are referring to, as shown here:

```
global::System.Collections.Generic.List<int>
```

This is the class you'd expect it to be, the generic `List<T>` collection class. It definitely isn't the class defined with the following code:

```
namespace MyRootNamespace
{
    namespace System
    {
        namespace Collections
        {
            namespace Generic
            {
                class List<T> {}
            }
        }
    }
}
```

Of course, you should avoid giving your namespaces names that already exist as .NET namespaces, although similar problems can arise in large projects, particularly if you are working as part of a large team. Using the `::` operator and the `global` keyword might be the only way you can access the types you want.

## CUSTOM EXCEPTIONS

Chapter 7 covered exceptions and explained how you can use `try...catch...finally` blocks to act on them. You also saw several standard .NET exceptions, including the base class for exceptions, `System.Exception`. Sometimes it's useful to derive your own exception classes from this base class for use in your applications, instead of using the standard exceptions. This enables you to be more specific with the information you send to whatever code catches the exception, and it enables catching code to be more specific about which exceptions it handles. For example, you might add a new property to your exception class that permits access to some underlying information, making it possible for the exception's receiver to make the required changes, or just provide more information about the exception's cause.

**NOTE** Two fundamental exception classes exist in the `System` namespace and derive from `Exception`: `ApplicationException` and `SystemException`. `SystemException` is used as the base class for exceptions that are predefined by the .NET Framework. `ApplicationException` was provided for developers to derive their own exception classes, but more recent best practice dictates that you should not derive your exceptions from this class; you should use `Exception` instead.

## Adding Custom Exceptions to CardLib

How to use custom exceptions is, once again, best illustrated by upgrading the CardLib project. The `Deck.GetCard()` method currently throws a standard .NET exception if an attempt is made

to access a card with an index less than 0 or greater than 51, but you'll modify that to use a custom exception.

First, you need to create a new class library project called Ch13CardLib, save it in the BegVCSharp\Chapter13 directory, and copy the classes from Ch12CardLib as before, changing the namespace to Ch13CardLib as applicable. Next, define the exception. You do this with a new class defined in a new class file called CardOutOfRangeException.cs, which you can add to the Ch13CardLib project with Project → Add Class (you can find this code in Ch13CardLib\CardOutOfRangeException.cs):

```
public class CardOutOfRangeException : Exception
{
    private Cards deckContents;
    public Cards DeckContents
    {
        get { return deckContents; }
    }
    public CardOutOfRangeException(Cards sourceDeckContents)
        : base("There are only 52 cards in the deck.")
    {
        deckContents = sourceDeckContents;
    }
}
```

An instance of the Cards class is required for the constructor of this class. It allows access to this Cards object through a DeckContents property and supplies a suitable error message to the base Exception constructor so that it is available through the Message property of the class.

Next, add code to throw this exception to Deck.cs, replacing the old standard exception (you can find this code in Ch13CardLib\Deck.cs):

```
public Card GetCard(int cardNum)
{
    if (cardNum >= 0 && cardNum <= 51)
        return cards[cardNum];
    else
        throw new CardOutOfRangeException(cards.Clone() as Cards);
}
```

The DeckContents property is initialized with a deep copy of the current contents of the Deck object, in the form of a Cards object. This means that you see the contents at the point where the exception was thrown, so subsequent modification to the deck contents won't "lose" this information.

To test this, use the following client code (you can find this code in Ch13CardClient\Program.cs):

```
Deck deck1 = new Deck();
try
{
    Card myCard = deck1.GetCard(60);
}
```

```

        catch (CardOutOfRangeException e)
    {
        WriteLine(e.Message);
        WriteLine(e.DeckContents[0]);
    }
    ReadKey();
}

```

This code results in the output shown in Figure 13-1.



**FIGURE 13-1**

Here, the catching code has written the exception `Message` property to the screen. You also displayed the first card in the `Cards` object obtained through `DeckContents`, just to prove that you can access the `Cards` collection through your custom exception object.

## EVENTS

This section covers one of the most frequently used OOP techniques in .NET: *events*. You start, as usual, with the basics — looking at what events actually are. After that, you'll see some simple events in action and learn what you can do with them. Then, you learn how you can create and use events of your own.

In this chapter you'll complete your `CardLib` class library by adding an event. Finally, because this is the last port of call before arriving at some advanced topics, you'll have a bit of fun creating a card game application that uses this class library.

### What Is an Event?

Events are similar to exceptions in that they are *raised* (thrown) by objects, and you can supply code that acts on them. However, there are several important differences, the most important of which is that there is no equivalent to the `try...catch` structure for handling events. Instead, you must *subscribe* to them. Subscribing to an event means supplying code that will be executed when an event is raised, in the form of an *event handler*.

Many handlers can be subscribed to a single event, all of which are called when the event is raised. This can include event handlers that are part of the class of the object that raises the event, but event handlers are just as likely to be found in other classes.

Event handlers themselves are simply methods. The only restriction on an event handler method is that it must match the return type and parameters required by the event. This restriction is part of the definition of an event and is specified by a *delegate*.

**NOTE** The fact that delegates are used in events is one of the reasons why delegates are so useful. This is why some space was devoted to them in Chapter 6. You might want to review that material to refresh your memory about delegates and how you use them.

The basic sequence of processing is as follows: First, an application creates an object that can raise an event. For example, suppose an instant messaging application creates an object that represents a connection to a remote user. That connection object might raise an event when a message arrives through the connection from the remote user (see Figure 13-2).




FIGURE 13-2

Next, the application subscribes to the event. Your instant messaging application would do this by defining a method that could be used with the delegate type specified by the event, passing a reference to this method to the event. The event handler method might be a method on another object, such as an object representing a display device to show instant messages when they arrive (see Figure 13-3).




FIGURE 13-3

When the event is raised, the subscriber is notified. When an instant message arrives through the connection object, the event handler method on the display device object is called. Because you are using a standard method, the object that raises the event can pass any relevant information via parameters, making events very versatile. In the example case, one parameter might be the text of the instant message, which the event handler could display on the display device object. This is shown in Figure 13-4.




FIGURE 13-4

## Handling Events

As previously discussed, to handle an event you need to subscribe to it by providing an event handler method whose return type and parameters match that of the delegate specified for use with the event. The following example uses a simple timer object to raise events, which results in a handler method being called.

### TRY IT OUT Handling Events: Ch13Ex01\Program.cs

1. Create a new console application called Ch13Ex01 and save it in the directory C:\BegVCSharp\Chapter13.
2. Modify the code in Program.cs as follows:


```

using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.Timers;
using static System.Console;
namespace Ch13Ex01
{
    class Program
    {
        static int counter = 0;

        static string displayString =
            "This string will appear one letter at a time. ";
        static void Main(string[] args)
        {
            Timer myTimer = new Timer(100);
            myTimer.Elapsed += new ElapsedEventHandler(WriteChar);
        }
        static void WriteChar(object source, ElapsedEventArgs args)
        {
            if (counter < displayString.Length)
                Console.Write(displayString[counter++]);
            else
                Console.WriteLine();
            if (counter < displayString.Length)
                myTimer.Interval = 100;
            else
                myTimer.Enabled = false;
        }
    }
}
  
```

```
        myTimer.Start();
        System.Threading.Thread.Sleep(200);
        ReadKey();
    }
    static void WriteChar(object source, ElapsedEventArgs e)
    {
        Write(displayString[counter++ % displayString.Length]);
    }
}
```

3. Run the application (once it is running, pressing a key will terminate the application). The result, after a short period, is shown in Figure 13-5.



**FIGURE 13-5**

## How It Works

The object you are using to raise events is an instance of the `System.Timers.Timer` class. This object is initialized with a time period (in milliseconds). When the `Timer` object is started using its `Start()` method, a stream of events is raised, spaced out in time according to the specified time period. `Main()` initializes a `Timer` object with a timer period of 100 milliseconds, so it will raise events 10 times a second when started:

```
static void Main(string[] args)
{
    Timer myTimer = new Timer(100);
```

The `Timer` object possesses an event called `Elapsed`, and the event handler required by this event must match the return type and parameters of the `System.Timers.ElapsedEventHandler` delegate type, which is one of the standard delegates defined in the .NET Framework. This delegate specifies the following return type and parameters:

```
void <MethodName>(object source, ElapsedEventArgs e);
```

The Timer object sends a reference to itself in the first parameter and an instance of an ElapsedEventArgs object in its second parameter. It is safe to ignore these parameters for now; you'll take a look at them a little later.

In your code you have a suitable method:

```
static void WriteChar(object source, ElapsedEventArgs e)
{
    Write(displayString[counter++ % displayString.Length]);
}
```

This method uses the two static fields of `Program`, `counter` and `displayString`, to display a single character. Every time the method is called, a different character is displayed.

The next task is to hook this handler up to the event — to subscribe to it. To do this, you use the `+=` operator to add a handler to the event in the form of a new delegate instance initialized with your event handler method:

```
static void Main(string[] args)
{
    Timer myTimer = new Timer(100);
    myTimer.Elapsed += new ElapsedEventHandler(WriteChar);
```

This command (which uses slightly strange-looking syntax, specific to delegates) adds a handler to the list that will be called when the `Elapsed` event is raised. You can add as many handlers as you like to this list as long as they all meet the criteria required. Each handler is called in turn when the event is raised.

All that remains for `Main()` to do is start the timer running:

```
myTimer.Start();
```

You don't want the application terminating before you have handled any events, so you put the `Main()` method on hold. The simplest way to do this is to request user input, as this command won't finish processing until the user has pressed a key:

```
ReadKey();
```

Although processing in `Main()` effectively ceases here, processing in the `Timer` object continues. When it raises events it calls the `WriteChar()` method, which runs concurrently with the `ReadLine()` statement. The `System.Threading.Thread.Sleep(200)` statement is included to give the timer the opportunity to start sending messages to the console application.

Note that the syntax for adding an event handler can be simplified slightly using the method group concept introduced in the previous chapter, as follows:

```
myTimer.Elapsed += WriteChar;
```

The end result is exactly the same, but you do not have to explicitly specify the delegate type; it is inferred by the compiler from the context in which you use it. However, many programmers dislike this syntax because it reduces readability — it is no longer possible to tell at a glance what delegate type you are using. Feel free to use this syntax if you prefer, but in this chapter all the delegates you use will be referenced explicitly to make things clearer.

## Defining Events

Now it's time to define and use your own events. The following Try It Out implements an example version of the instant messaging scenario introduced earlier in this chapter, creating a `Connection` object that raises events that are handled by a `Display` object.

### TRY IT OUT Defining Events: Ch13Ex02

1. Create a new console application called Ch13Ex02 and save it in the directory `C:\BegVCSharp\Chapter13`.
2. Add a new class called `Connection` and modify `Connection.cs` as follows:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.Timers;
using static System.Console;
namespace Ch13Ex02
{
    public delegate void MessageHandler(string messageText);
    public class Connection
    {
        public event MessageHandler MessageArrived;
        private Timer pollTimer;
        public Connection()
        {
            pollTimer = new Timer(100);
            pollTimer.Elapsed += new ElapsedEventHandler(CheckForMessage);
        }
        public void Connect() => pollTimer.Start();
        public void Disconnect() => pollTimer.Stop();
        private static Random random = new Random();
        private void CheckForMessage(object source, ElapsedEventArgs e)
        {
            WriteLine("Checking for new messages.");
            if ((random.Next(9) == 0) && (MessageArrived != null))
            {
                MessageArrived("Hello Mami!");
            }
        }
    }
}
```

3. Add a new class called `Display` and modify `Display.cs` as follows:

```
namespace Ch13Ex02
{
    public class Display
    {
```

```

        public void DisplayMessage(string message)
            => WriteLine($"Message arrived: {message}");
    }
}

```

4. Modify the code in Program.cs as follows:

```

static void Main(string[] args)
{
    Connection myConnection = new Connection();
    Display myDisplay = new Display();
    myConnection.MessageArrived +=
        new MessageHandler (myDisplay.DisplayMessage);
    myConnection.Connect();
    ReadKey();
}

```

5. Run the application. The result is shown in Figure 13-6.




FIGURE 13-6

## How It Works

The Connection class does most of the work in this application. Instances of this class make use of a Timer object much like the one shown in the first example of this chapter, initializing it in the class constructor and providing access to its state (enabled or disabled) via Connect() and Disconnect():

```

public class Connection
{
    private Timer pollTimer;
    public Connection()
    {
        pollTimer = new Timer(100);
        pollTimer.Elapsed += new ElapsedEventHandler(CheckForMessage);
    }
}

```

```
    public void Connect() => pollTimer.Start();
    public void Disconnect() => pollTimer.Stop();
    ...
}
```

Also in the constructor, you register an event handler for the `Elapsed` event, just as you did in the first example. The handler method, `CheckForMessage()`, raises an event on average once every 10 times it is called. You will look at the code for this, but first it would be useful to look at the event definition itself.

Before you define an event, you must define a delegate type to use with the event — that is, a delegate type that specifies the return type and parameters to which an event handling method must conform. You do this using standard delegate syntax, defining it as public inside the `Ch13Ex02` namespace to make the type available to external code:

```
namespace Ch13Ex02
{
    public delegate void MessageHandler(string messageText);
```

This delegate type, called `MessageHandler` here, is a `void` method that has a single `string` parameter. You can use this parameter to pass an instant message received by the `Connection` object to the `Display` object. Once a delegate has been defined (or a suitable existing delegate has been located), you can define the event itself, as a member of the `Connection` class:

```
public class Connection
{
    public event MessageHandler MessageArrived;
```

You simply name the event (here it is `MessageArrived`) and declare it by using the `event` keyword and specifying the delegate type to use (the `MessageHandler` delegate type defined earlier). After you have declared an event in this way, you can raise it simply by calling it by name as if it were a method with the return type and parameters specified by the delegate. For example, you could raise this event using the following:

```
MessageArrived("This is a message.");
```

If the delegate had been defined without any parameters, then you could simply use the following:

```
MessageArrived();
```

Alternatively, you could define more parameters, which would require more code to raise the event. The `CheckForMessage()` method looks like this:

```
private static Random random = new Random();
private void CheckForMessage(object source, ElapsedEventArgs e)
{
    WriteLine("Checking for new messages.");
    if ((random.Next(9) == 0) && (MessageArrived != null))
    {
        MessageArrived("Hello Mami!");
    }
}
```

You use an instance of the `Random` class shown in earlier chapters to generate a random number between 0 and 9, and raise an event if the number generated is 0, which should happen 10 percent of

the time. This simulates polling the connection to determine whether a message has arrived, which won't be the case every time you check. To separate the timer from the instance of `Connection`, you use a private static instance of the `Random` class.

Note that you supply additional logic. You raise an event only if the expression `MessageArrived != null` evaluates to `true`. This expression, which again uses the delegate syntax in a slightly unusual way, means “Does the event have any subscribers?” If there are no subscribers, then `MessageArrived` evaluates to `null`, and there is no point in raising the event.

The class that will subscribe to the event is called `Display` and contains the single method, `DisplayMessage()`, defined as follows:

```
public class Display
{
    public void DisplayMessage(string message)
        => WriteLine($"Message arrived: {message}");
}
```

This method matches the delegate type (and is `public`, which is a requirement of event handlers in classes other than the class that generates the event), so you can use it to respond to the `MessageArrived` event.

All that is left now is for the code in `Main()` to initialize instances of the `Connection` and `Display` classes, hook them up, and start things going. The code required here is similar to the first example:

```
static void Main(string[] args)
{
    Connection myConnection = new Connection();
    Display myDisplay = new Display();
    myConnection.MessageArrived +=
        new MessageHandler(myDisplay.DisplayMessage);
    myConnection.Connect();
    System.Threading.Thread.Sleep(200);
    ReadKey();
}
```

Again, you call `ReadKey()` to pause the processing of `Main()` once you have started things moving with the `Connect()` method of the `Connection` object and inserted a short delay.

## Multipurpose Event Handlers

The delegate you saw earlier, for the `Timer.Elapsed` event, contained two parameters that are of a type often seen in event handlers:

- `object source` — A reference to the object that raised the event
- `ElapsedEventArgs e` — Parameters sent by the event

The reason the `object` type parameter is used in this event, and indeed in many other events, is that you often need to use a single event handler for several identical events generated by different objects and still tell which object generated the event.

To explain and illustrate this concept, the next Try It Out extends the last example a little.

**TRY IT OUT** Using a Multipurpose Event Handler: Ch13Ex03

1. Create a new console application called Ch13Ex03 and save it in the directory C:\BegVCSharp\Chapter13.
2. Copy the code across for Program.cs, Connection.cs, and Display.cs from Ch13Ex02, making sure that you change the namespaces in each file from Ch13Ex02 to Ch13Ex03.
3. Add a new class called MessageArrivedEventArgs and modify MessageArrivedEventArgs.cs as follows:

```
namespace Ch13Ex03
{
    public class MessageArrivedEventArgs : EventArgs
    {
        private string message;
        public string Message
        {
            get { return message; }
        }
        public MessageArrivedEventArgs()
        {
            message = "No message sent.";
        }
        public MessageArrivedEventArgs(string newMessage)
        {
            message = newMessage;
        }
    }
}
```

4. Modify Connection.cs as follows:

```
namespace Ch13Ex03
{
    // delegate definition removed
    public class Connection
    {
        public event EventHandler<MessageArrivedEventArgs> MessageArrived;
        public string Name { get; set; }
        ...
        private void CheckForMessage(object source, EventArgs e)
        {
            WriteLine("Checking for new messages.");
            if ((random.Next(9) == 0) && (MessageArrived != null))
            {
                MessageArrived(this, new MessageArrivedEventArgs("Hello Mami!"));
            }
        }
        ...
    }
}
```

5. Modify Display.cs as follows:

```
public void DisplayMessage(object source, MessageArrivedEventArgs e)
{
    WriteLine($"Message arrived from: {((Connection)source).Name}");
    WriteLine($"Message Text: {e.Message}");
}
```

6. Modify Program.cs as follows:

```
static void Main(string[] args)
{
    Connection myConnection1 = new Connection();
    myConnection1.Name = "First connection.";
    Connection myConnection2 = new Connection();
    myConnection2.Name = "Second connection.";
    Display myDisplay = new Display();
    myConnection1.MessageArrived += myDisplay.DisplayMessage;
    myConnection2.MessageArrived += myDisplay.DisplayMessage;
    myConnection1.Connect();
    myConnection2.Connect();
    System.Threading.Thread.Sleep(200);
    ReadKey();
}
```

7. Run the application. The result is shown in Figure 13-7.




FIGURE 13-7

## How It Works

By sending a reference to the object that raises an event as one of the event handler parameters, you can customize the response of the handler to individual objects. The reference gives you access to the source object, including its properties.

By sending parameters that are contained in a class that inherits from `System.EventArgs` (as `ElapsedEventArgs` does), you can supply whatever additional information is necessary as parameters (such as the `Message` parameter on the `MessageArrivedEventArgs` class).

In addition, these parameters benefit from polymorphism. You could define a handler for the `MessageArrived` event such as this:

```
public void DisplayMessage(object source, EventArgs e)
{
    WriteLine($"Message arrived from: {((Connection)source).Name}");
    WriteLine($"Message Text: {((MessageArrivedEventArgs)e).Message}");
}
```

The application will execute exactly as it did before, but the `DisplayMessage()` method is now more versatile (in theory at least — more implementation is needed to make this production quality). This same handler could work with other events, such as the `Timer.Elapsed`, although you'd have to modify the internals of the handler a bit more such that the parameters sent when this event is raised are handled properly. (Casting them to `Connection` and `MessageArrivedEventArgs` objects in this way will cause an exception; you should use the `as` operator instead and check for `null` values.)

---

## The EventHandler and Generic EventHandler<T> Types

In most cases, you will follow the pattern outlined in the previous section and use event handlers with a `void` return type and two parameters. The first parameter will be of type `object`, and will be the event source. The second parameter will be of a type that derives from `System.EventArgs`, and will contain any event arguments. As this is so common, .NET provides two delegate types to make it easier to define events: `EventHandler` and `EventHandler<T>`. Both of these are delegates that use the standard event handler pattern. The generic version enables you to specify the type of event argument you want to use.

In the previous Try It Out, you saw this in action as you used the generic `EventHandler<T>` delegate type as follows:

```
public class Connection
{
    public event EventHandler<MessageArrivedEventArgs> MessageArrived;
    ...
}
```

This is obviously a good thing to do because it simplifies your code. In general, it is best practice to use these delegate types whenever you define an event. Note that if you have an event that doesn't need event argument data, you can still use the `EventHandler` delegate type. You can simply pass `EventArgs.Empty` as the argument value.

## Return Values and Event Handlers

All the event handlers you've seen so far have had a return type of `void`. It is possible to provide a return type for an event, but this can lead to problems because a given event can result in several event handlers being called. If all of these handlers return a value, then it can be unclear which value was actually returned.

The system deals with this by allowing you access to only the last value returned by an event handler. That will be the value returned by the last event handler to subscribe to an event. Although this

functionality might be of use in some situations, it is recommended that you use `void` type event handlers, and avoid `out` type parameters (which would lead to the same ambiguity regarding the source of the value returned by the parameter).

## Anonymous Methods

Instead of defining event handler methods, you can choose to use *anonymous methods*. An anonymous method doesn't actually exist as a method in the traditional sense — that is, it isn't a method on any particular class. Instead, an anonymous method is created purely for use as a target for a delegate.

To create an anonymous method, you need the following code:

```
delegate (parameters)
{
    // Anonymous method code.
};
```

`parameters` is a list of parameters matching those of the delegate type you are instantiating, as used by the anonymous method code:

```
delegate (Connection source, MessageArrivedEventArgs e)
{
    // Anonymous method code matching MessageHandler event in Ch13Ex03.
};
```

For example, you could use this code to completely bypass the `Display.DisplayMessage()` method in Ch13Ex03:

```
myConnection1.MessageArrived += 
    delegate (Connection source, MessageArrivedEventArgs e)
    {
        WriteLine($"Message arrived from: {source.Name}");
        WriteLine($"Message Text: {e.Message}");
    };
}
```

An interesting point about anonymous methods is that they are effectively local to the code block that contains them, and they have access to local variables in this scope. If you use such a variable, then it becomes an *outer* variable. Outer variables are not disposed of when they go out of scope as other local variables are; instead, they live on until the anonymous methods that use them are destroyed. This might be some time later than you expect, so it's definitely something to be careful about. If an outer variable takes up a large amount of memory, or if it uses resources that are expensive in other ways (for example, resources that are limited in number), then this could cause memory or performance problems.

## EXPANDING AND USING CARDLIB

Now that you've had a look at defining and using events, you can use them in Ch13CardLib. The event you'll add to your library will be generated when the last `Card` object in a `Deck` object is obtained by using `GetCard`, and it will be called `LastCardDrawn`. The event enables subscribers to

reshuffle the deck automatically, cutting down on the processing necessary by a client. The event will use the `EventHandler` delegate type and will pass as its source a reference to the `Deck` object, such that the `Shuffle()` method will be accessible from wherever the handler is. Add the following code to `Deck.cs` (you can find this code in `Ch13CardLib\Deck.cs`) to define and raise the event:

```
namespace Ch13CardLib
{
    public class Deck : ICloneable
    {
        public event EventHandler LastCardDrawn;
        ...
        public Card GetCard(int cardNum)
        {
            if (cardNum >= 0 && cardNum <= 51)
            {
                if ((cardNum == 51) && (LastCardDrawn != null))
                    LastCardDrawn(this, EventArgs.Empty);
                return cards[cardNum];
            }
            else
                throw new CardOutOfRangeException((Cards)cards.Clone());
        }
        ...
    }
}
```

This is all the code required to add the event to the `Deck` class definition.

After spending all this time developing the `CardLib` library, it would be a shame not to use it. Before finishing this section on OOP in C# and the .NET Framework, it's time to have a little fun and write the basics of a card game application that uses the familiar playing card classes.

As in previous chapters, you'll add a client console application to the `Ch13CardLib` solution, add a reference to the `Ch13CardLib` project, and make it the startup project. This application will be called `Ch13CardClient`.

To begin, you'll create a new class called `Player` in a new file in `Ch13CardClient`, `Player.cs`. You can find this code in `Ch13CardClient\Player.cs` in this chapter's online download. This class will contain two automatic properties: `Name` (a `string`) and `PlayHand` (of type `Cards`). Both of these properties have private set accessors, but despite this the `PlayHand` provides write-access to its contents, enabling you to modify the cards in the player's hand.

You'll also hide the default constructor by making it private, and supply a public nondefault constructor that accepts an initial value for the `Name` property of `Player` instances.

Finally, you'll provide a `bool` type method called `HasWon()`, which returns `true` if all the cards in the player's hand are the same suit (a simple winning condition, but that doesn't matter too much).

Here's the code for `Player.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
```

```

using System.Text;
using System.Threading.Tasks;
using Ch13CardLib;
namespace Ch13CardClient
{
    public class Player
    {
        public string Name { get; private set; }
        public Cards PlayHand { get; private set; }
        private Player()
        {
        }
        public Player(string name)
        {
            Name = name;
            PlayHand = new Cards();
        }
        public bool HasWon()
        {
            bool won = true;
            Suit match = PlayHand[0].suit;
            for (int i = 1; i < PlayHand.Count; i++)
            {
                won &= PlayHand[i].suit == match;
            }
            return won;
        }
    }
}

```

Next, define a class that will handle the card game itself, called Game. This class is found in the file Game.cs of the Ch13CardClient project. The class has four private member fields:

- playDeck — A Deck type variable containing the deck of cards to use
- currentCard — An int value used as a pointer to the next card in the deck to draw
- players — An array of Player objects representing the players of the game
- discardedCards — A Cards collection for the cards that have been discarded by players but not shuffled back into the deck

The default constructor for the class initializes and shuffles the Deck stored in playDeck, sets the currentCard pointer variable to 0 (the first card in playDeck), and wires up an event handler called Reshuffle() to the playDeck.LastCardDrawn event. The handler simply shuffles the deck, initializes the discardedCards collection, and resets currentCard to 0, ready to read cards from the new deck.

The Game class also contains two utility methods: SetPlayers() for setting the players for the game (as an array of Player objects) and DealHands() for dealing hands to the players (seven cards each). The allowed number of players is restricted to between two and seven to ensure that there are enough cards to go around.

Finally, there is a `PlayGame()` method that contains the game logic itself. You'll come back to this method shortly, after you've looked at the code in `Program.cs`. The rest of the code in `Game.cs` is as follows (you can find this code in `Ch13CardClient\Game.cs`):

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using Ch13CardLib;
using static System.Console;
namespace Ch13CardClient
{
    public class Game
    {
        private int currentCard;
        private Deck playDeck;
        private Player[] players;
        private Cards discardedCards;
        public Game()
        {
            currentCard = 0;
            playDeck = new Deck(true);
            playDeck.LastCardDrawn += Reshuffle;
            playDeck.Shuffle();
            discardedCards = new Cards();
        }
        private void Reshuffle(object source, EventArgs args)
        {
            WriteLine("Discarded cards reshuffled into deck.");
            ((Deck)source).Shuffle();
            discardedCards.Clear();
            currentCard = 0;
        }
        public void SetPlayers(Player[] newPlayers)
        {
            if (newPlayers.Length > 7)
                throw new ArgumentException(
                    "A maximum of 7 players may play this game.");
            if (newPlayers.Length < 2)
                throw new ArgumentException(
                    "A minimum of 2 players may play this game.");
            players = newPlayers;
        }
        private void DealHands()
        {
            for (int p = 0; p < players.Length; p++)
            {
                for (int c = 0; c < 7; c++)
                {
                    players[p].PlayHand.Add(playDeck.GetCard(currentCard++));
                }
            }
        }
    }
}
```

```

        public int PlayGame()
        {
            // Code to follow.
        }
    }
}

```

`Program.cs` contains the `Main()` method, which initializes and runs the game. This method performs the following steps:

1. An introduction is displayed.
2. The user is prompted for a number of players between 2 and 7.
3. An array of `Player` objects is set up accordingly.
4. Each player is prompted for a name, used to initialize one `Player` object in the array.
5. A `Game` object is created and players are assigned using the `SetPlayers()` method.
6. The game is started by using the `PlayGame()` method.
7. The `int` return value of `PlayGame()` is used to display a winning message (the value returned is the index of the winning player in the array of `Player` objects).

The code for this follows, with comments added for clarity (you can find this code in `Ch13CardClient\Program.cs`):

```

static void Main(string[] args)
{
    // Display introduction.
    WriteLine("BenjaminCards: a new and exciting card game.");
    WriteLine("To win you must have 7 cards of the same suit in" +
              " your hand.");
    WriteLine();
    // Prompt for number of players.
    bool inputOK = false;
    int choice = -1;
    do
    {
        WriteLine("How many players (2-7)?");
        string input = ReadLine();
        try
        {
            // Attempt to convert input into a valid number of players.
            choice = Convert.ToInt32(input);
            if ((choice >= 2) && (choice <= 7))
                inputOK = true;
        }
        catch
        {
            // Ignore failed conversions, just continue prompting.
        }
    } while (inputOK == false);
    // Initialize array of Player objects.
}

```

```
Player[] players = new Player[choice];
// Get player names.
for (int p = 0; p < players.Length; p++)
{
    WriteLine($"Player {p + 1}, enter your name:");
    string playerName = ReadLine();
    players[p] = new Player(playerName);
}
// Start game.
Game newGame = new Game();
newGame.SetPlayers(players);
int whoWon = newGame.PlayGame();
// Display winning player.
WriteLine($"{players[whoWon].Name} has won the game!");
.ReadKey();
}
```

Now you come to `PlayGame()`, the main body of the application. Space limitations preclude us from providing a lot of detail about this method, but the code is commented to make it more comprehensible. None of the code is complicated; there's just quite a bit of it.

Play proceeds with each player viewing his or her cards and an upturned card on the table. They can either pick up this card or draw a new one from the deck. After drawing a card, each player must discard one, replacing the card on the table with another one if it has been picked up, or placing the discarded card on top of the one on the table (also adding the discarded card to the `discardedCards` collection).

As you consider this code, bear in mind how the `Card` objects are manipulated. The reason why these objects are defined as reference types, rather than value types (using a struct), should now be clear. A given `Card` object can appear to exist in several places at once because references can be held by the `Deck` object, the hand fields of the `Player` objects, the `discardedCards` collection, and the `playCard` object (the card currently on the table). This makes it easy to keep track of the cards and is used in particular in the code that draws a new card from the deck. The card is accepted only if it isn't in any player's hand or in the `discardedCards` collection.

The code is as follows:

```
public int PlayGame()
{
    // Only play if players exist.
    if (players == null)
        return -1;
    // Deal initial hands.
    DealHands();
    // Initialize game vars, including an initial card to place on the
    // table: playCard.
    bool GameWon = false;
    int currentPlayer;
    Card playCard = playDeck.GetCard(currentCard++);
    discardedCards.Add(playCard);
    // Main game loop, continues until GameWon == true.
    do
    {
```

```
// Loop through players in each game round.
for (currentPlayer = 0; currentPlayer < players.Length;
    currentPlayer++)
{
    //Write out current player, player hand, and the card on the
    // table.
    WriteLine($"{players[currentPlayer].Name}'s turn.");
    WriteLine("Current hand:");
    foreach (Card card in players[currentPlayer].PlayHand)
    {
        WriteLine(card);
    }
    WriteLine($"Card in play: {playCard}");
    // Prompt player to pick up card on table or draw a new one.
    bool inputOK = false;
    do
    {
        WriteLine("Press T to take card in play or D to draw:");
        string input = ReadLine();
        if (input.ToLower() == "t")
        {
            // Add card from table to player hand.
            WriteLine("Drawn: {playCard}");
            // Remove from discarded cards if possible (if deck
            // is reshuffled it won't be there any more)
            if (discardedCards.Contains(playCard))
            {
                discardedCards.Remove(playCard);
            }
            players[currentPlayer].PlayHand.Add(playCard);
            inputOK = true;
        }
        if (input.ToLower() == "d")
        {
            // Add new card from deck to player hand.
            Card newCard;
            // Only add card if it isn't already in a player hand
            // or in the discard pile
            bool cardIsAvailable;
            do
            {
                newCard = playDeck.GetCard(currentCard++);
                // Check if card is in discard pile
                cardIsAvailable = !discardedCards.Contains(newCard);
                if (cardIsAvailable)
                {
                    // Loop through all player hands to see if newCard
                    // is already in a hand.
                    foreach (Player testPlayer in players)
                    {
                        if (testPlayer.PlayHand.Contains(newCard))
                        {
                            cardIsAvailable = false;
                            break;
                        }
                    }
                }
            }
        }
    }
}
```

```
        }
    }
} while (!cardIsAvailable);
// Add the card found to player hand.
WriteLine($"Drawn: {newCard}");
players[currentPlayer].PlayHand.Add(newCard);
inputOK = true;
}
} while (inputOK == false);
// Display new hand with cards numbered.
WriteLine("New hand:");
for (int i = 0; i < players[currentPlayer].PlayHand.Count;
i++)
{
    WriteLine($"{i + 1}: " +
              $"{ players[currentPlayer].PlayHand[i]}");
}
// Prompt player for a card to discard.
inputOK = false;
int choice = -1;
do
{
    WriteLine("Choose card to discard:");
    string input = ReadLine();
    try
    {
        // Attempt to convert input into a valid card number.
        choice = Convert.ToInt32(input);
        if ((choice > 0) && (choice <= 8))
            inputOK = true;
    }
    catch
    {
        // Ignore failed conversions, just continue prompting.
    }
} while (inputOK == false);
// Place reference to removed card in playCard (place the card
// on the table), then remove card from player hand and add
// to discarded card pile.
playCard = players[currentPlayer].PlayHand[choice - 1];
players[currentPlayer].PlayHand.RemoveAt(choice - 1);
discardedCards.Add(playCard);
WriteLine($"Discarding: {playCard}");
// Space out text for players
WriteLine();
// Check to see if player has won the game, and exit the
player
// loop if so.
GameWon = players[currentPlayer].HasWon();
if (GameWon == true)
    break;
}
} while (GameWon == false);
// End game, noting the winning player.
return currentPlayer;
}
```

Figure 13-8 shows a game in progress.




FIGURE 13-8

As a final exercise, have a close look at the code in `Player.HasWon()`. Can you think of a way that you could make this code more efficient, perhaps without having to examine every card in the player's hand every time this method is called?

## ATTRIBUTES

This section takes a brief look at a useful way to provide additional information to code that consumes types that you create: *attributes*. Attributes give you a way to mark sections of code with information that can be read externally and used in any number of ways to affect how your types are used. This is often referred to as *decorating* the code. You can find the code for this section in `CustomAttributes\Program.cs` in this chapter's online download.

For example, let's say you create a class with a really simple method. In fact, it's so simple that you really aren't that interested in stepping through it. Unfortunately — and to your considerable annoyance — you keep doing precisely that as you debug the code in your application. In this situation, it's possible to add an attribute to the method that tells Visual Studio not to step into the code when you debug it; instead, Visual Studio should step through it and on to the next statement. The code for this is as follows:

```
[DebuggerStepThrough]
public void DullMethod() { ... }
```

The attribute in this code is `[DebuggerStepThrough]`. All attributes are added in this way, by enclosing the name of the attribute in square brackets just before the target to which they apply. You can add multiple attributes to a single target either by separating them with commas or by enclosing each one in square brackets.

The attribute used in the preceding code is actually implemented in a class called `DebuggerStepThroughAttribute`, and is found in the `System.Diagnostics` namespace, so you need a `using` statement for that namespace if you want to use this attribute. You can refer to this attribute either by its full name or, as in the code you saw, with an abbreviated name that doesn't include the suffix `Attribute`.

When you add an attribute in this way, the compiler creates an instance of the attribute class and associates it with the class method. Some attributes are customizable through constructor parameters or properties, and these can be specified when you add the attribute, for example:

```
[DoesInterestingThings(1000, WhatDoesItDo = "voodoo")]
public class DecoratedClass {}
```

This attribute is passing a value of 1000 to the constructor of `DoesInterestingThingsAttribute` and setting the value of a property called `WhatDoesItDo` to the string "voodoo".

## Reading Attributes

In order to read attribute values, you have to use a technique called *reflection*. This is a fairly advanced technique that allows you to dynamically inspect type information at runtime, even to the point where you can create objects and call methods without knowing what those objects are. This book doesn't cover this technique in detail, but you do need to know some basics in order to use attributes.

Essentially, reflection involves using information stored in `Type` objects (which you've seen in several places in this book) along with types in the `System.Reflection` namespace to work with type information. You've already seen a quick way to get type information from a class with the `typeof` operator, and from an object instance using the `GetType()` method. Using reflection you can proceed to interrogate member information from the `Type` object. You can then obtain attribute information from the class or its various members.

The simplest way to do this — and the only way you'll see in this book — is to use the `Type` `.GetCustomAttributes()` method. This method takes up to two parameters and returns an array of object instances, each of which is an attribute instance. First, you can optionally pass the type

or types of attributes you are interested in (any other attributes will be ignored). If you omit this parameter, then all attributes will be returned. Second, you must pass a Boolean value indicating whether to look just at the class or at the class and all classes that the class derives from.

For example, the following code would list the attributes of a class called `DecoratedClass`:

```
Type classType = typeof(DecoratedClass);
object[] customAttributes = classType.GetCustomAttributes(true);
foreach (object customAttribute in customAttributes)
{
    WriteLine($"Attribute of type {customAttribute} found.");
}
```

Once you have found attributes in this way, you can take whatever action is appropriate for the attribute. This is exactly what Visual Studio does when it encounters the `DebuggerStepThroughAttribute` attribute discussed earlier.

## Creating Attributes

You can create your own attributes simply by deriving from the `System.Attribute` class. Sometimes, you don't need to do anything else, as no additional information is required if your code is interested only in the presence or absence of your attribute. However, you can supply nondefault constructors and/or writeable properties if you want the attribute to be customizable.

You also need to decide two things about your attribute: what type of target it can be applied to (class, property, and so on) and whether it can be applied more than once to the same target. You specify this information through an attribute that you apply to your attribute (this has a certain Zen feeling of correctness to it!) called `AttributeUsageAttribute`. This attribute has a constructor parameter of type `AttributeTargets`, which is an enum that allows you to combine its values with the `|` operator. It also has a Boolean property called `AllowMultiple` that specifies whether the attribute can be applied more than once.

For example, the following code specifies an attribute that can be applied (once) to a class or property:

```
[AttributeUsage(AttributeTargets.Class | AttributeTargets.Method,
               AllowMultiple = false)]
class DoesInterestingThingsAttribute : Attribute
{
    public DoesInterestingThingsAttribute(int howManyTimes)
    {
        HowManyTimes = howManyTimes;
    }
    public string WhatDoesItDo { get; set; }
    public int HowManyTimes { get; private set; }
}
```

This attribute, `DoesInterestingThingsAttribute`, can be used as in the earlier code snippet:

```
[DoesInterestingThings(1000, WhatDoesItDo = "voodoo")]
public class DecoratedClass {}
```

And by modifying the code in the previous section, you can gain access to the properties of the attribute:

```
Type classType = typeof(DecoratedClass);
object[] customAttributes = classType.GetCustomAttributes(true);
foreach (object customAttribute in customAttributes)
{
    WriteLine($"Attribute of type {customAttribute} found.");
    DoesInterestingThingsAttribute interestingAttribute =
        customAttribute as DoesInterestingThingsAttribute;
    if (interestingAttribute != null)
    {
        WriteLine($"This class does {interestingAttribute.WhatDoesItDo} x " +
                 $" {interestingAttribute.HowManyTimes}!");
    }
}
```

Putting everything in this section together and using this code would give you the result shown in Figure 13-9.



FIGURE 13-9

Attributes can be extremely useful and crop up all over .NET applications — and WPF and Windows Store applications in particular. You will encounter them repeatedly throughout the remainder of this book.

## INITIALIZERS

Up to now you have learned to instantiate and initialize objects in various ways. Invariably, that has required you either to add code to class definitions to enable initialization or to instantiate and initialize objects with separate statements. You have also learned how to create collection classes of various types, including generic collection classes. Again, you might have noticed that there was no easy way to combine the creation of a collection with adding items to the collection.

Object initializers provide a way to simplify your code by enabling you to combine instantiation and initialization of objects. Collection initializers give you a simple, elegant syntax to create and populate collections in a single step. This section explains how to use both of these features.

### Object Initializers

Consider the following simple class definition:

```
public class Curry
{
    public string MainIngredient { get; set; }
```

```

    public string Style { get; set; }
    public int Spiciness { get; set; }
}

```

This class has three properties that are defined using the automatic property syntax shown in Chapter 10. If you want to instantiate and initialize an object instance of this class, you must execute several statements:

```

Curry tastyCurry = new Curry();
tastyCurry.MainIngredient = "panir tikka";
tastyCurry.Style = "jalfrezi";
tastyCurry.Spiciness = 8;

```

This code uses the default, parameter-less constructor that is supplied by the C# compiler if you don't include a constructor in your class definition. To simplify this initialization, you can supply an appropriate nondefault constructor:

```

public class Curry
{
    public Curry(string mainIngredient, string style,
                int spiciness)
    {
        MainIngredient = mainIngredient;
        Style = style;
        Spiciness = spiciness;
    }
    ...
}

```

That enables you to write code that combines instantiation with initialization:

```
Curry tastyCurry = new Curry("panir tikka", "jalfrezi", 8);
```

This works fine, although it forces code that uses this class to use this constructor, which would prevent the previous code, which used a parameter-less constructor, from working. Often, particularly when classes must be serializable, it is necessary to provide a parameter-less constructor:

```

public class Curry
{
    public Curry() {}
    ...
}

```

Now you have a situation where you can instantiate and initialize the `Curry` class any way you like. However, you have added several lines of code to the initial class definition that don't do anything much other than provide the basic plumbing required for this flexibility.

Enter *object initializers*, which are a way to instantiate and initialize objects without having to add code (such as the constructors detailed here) to a class. When you instantiate an object, you supply values for publicly accessible properties or fields using a name/value pair for each property you want to initialize. The syntax for this is as follows:

```

<ClassName> <variableName> = new <ClassName>
{
    <propertyOrField1> = <value1>,

```

```
<propertyOrField2> = <value2>,
...
<propertyOrFieldN> = <valueN>
};
```

For example, you could rewrite the code shown earlier, which instantiates and initializes an object of type `Curry`, as follows:

```
Curry tastyCurry = new Curry
{
    MainIngredient = "panir tikka",
    Style = "jalfrezi",
    Spiciness = 8
};
```

Often you can put code like that on a single line without seriously degrading readability.

When you use an object initializer, you don't have to explicitly call a constructor of the class. If you omit the constructor parentheses (as in the previous code), the default parameter-less constructor is called automatically. This happens before any parameter values are set by the initializer, which enables you to provide default values for parameters in the default constructor if desired. Alternatively, you can call a specific constructor. Again, this constructor is called first, so any initialization of public properties that takes place in the constructor might be overridden by values that you provide in the initializer. You must have access to the constructor that you use (or the default one if you aren't explicit) in order for object initializers to work.

If one of the properties you want to initialize with an object initializer is more complex than the simple types used in this example, then you might find yourself using a *nested object initializer*. That simply means using the exact same syntax you've already seen:

```
Curry tastyCurry = new Curry
{
    MainIngredient = "panir tikka",
    Style = "jalfrezi",
    Spiciness = 8,
    Origin = new Restaurant
    {
        Name = "King's Balti",
        Location = "York Road",
        Rating = 5
    }
};
```

Here, a property called `Origin` of type `Restaurant` (not shown here) is initialized. The code initializes three properties of the `Origin` property — `Name`, `Location`, and `Rating` — with values of type `string`, `string`, and `int`, respectively. This initialization uses a nested object initializer.

Note that object initializers are not a replacement for nondefault constructors. The fact that you can use object initializers to set property and field values when you instantiate an object does not mean that you will always know what state needs initializing. With constructors you can specify exactly which values are required for an object to function and then execute code in response to those values immediately.

Also, in the previous example there is another (admittedly quite subtle) difference between using a nested object initializer and using constructors. This difference is the order in which objects get created. With a nested initializer, the top level object (`Curry`) gets created first. Next, the nested object (`Restaurant`) is created and assigned to the property `Origin`. If you used a constructor, you would reverse this construction order and pass the `Restaurant` instance to the constructor of `Curry`. In this simple example, there is no practical difference, but in some circumstances this might be significant.

## Collection Initializers

Chapter 5 described how arrays can be initialized with values using the following syntax:

```
int[] myIntArray = new int[5] { 5, 9, 10, 2, 99 };
```

This is a quick and easy way to combine the instantiation and initialization of an array. Collection initializers simply extend this syntax to collections:

```
List<int> myIntCollection = new List<int> { 5, 9, 10, 2, 99 };
```

By combining object and collection initializers, it is possible to configure collections with simple and elegant code. Rather than code like this:

```
List<Curry> curries = new List<Curry>();
curries.Add(new Curry("Chicken", "Pathia", 6));
curries.Add(new Curry("Vegetable", "Korma", 3));
curries.Add(new Curry("Prawn", "Vindaloo", 9));
```

You can use the following:

```
List<Curry> moreCurries = new List<Curry>
{
    new Curry
    {
        MainIngredient = "Chicken",
        Style = "Pathia",
        Spiciness = 6
    },
    new Curry
    {
        MainIngredient = "Vegetable",
        Style = "Korma",
        Spiciness = 3
    },
    new Curry
    {
        MainIngredient = "Prawn",
        Style = "Vindaloo",
        Spiciness = 9
    }
};
```

This works very well for types that are primarily used for data representation, and as such, collection initializers are a great accompaniment to the LINQ technology described later in the book.

The following Try It Out illustrates how you can use object and collection initializers.

### TRY IT OUT Using Initializers: Ch13Ex04

1. Create a new console application called Ch13Ex04 and save it in the directory C:\BegVCSharp\Chapter13.
2. Right-click on the project name in the Solution Explorer window and select the Add Existing Item option.
3. Select the Animal.cs, Cow.cs, Chicken.cs, SuperCow.cs, and Farm.cs files from the C:\BegVCSharp\Chapter12\Ch12Ex04\Ch12Ex04 directory, and click Add.
4. Modify the namespace declaration in the file you have added as follows:

```
namespace Ch13Ex04
```

5. Remove the constructors from the Cow, Chicken, and SuperCow classes.
6. Modify the code in Program.cs as follows:

```
static void Main(string[] args)
{
    Farm<Animal> farm = new Farm<Animal>
    {
        new Cow { Name="Lea" },
        new Chicken { Name="Noa" },
        new Chicken(),
        new SuperCow { Name="Andrea" }
    };
    farm.MakeNoises();
    ReadKey();
}
```

7. Build the application. You should receive the build errors shown in Figure 13-10 because there is no Add(T animal) method definition in the Farm class. It is added in the next step.




FIGURE 13-10

8. Add the following code to Farm.cs:

```
public class Farm<T> : IEnumerable<T> where T : Animal
{
    public void Add(T animal) => animals.Add(animal);
    ...
}
```

9. Run the application. The result is shown in Figure 13-11.




FIGURE 13-11

## How It Works

This example combines object and collection initializers to create and populate a collection of objects in a single step. It uses the farmyard collection of objects that you have seen in previous chapters, although two modifications are necessary for initializers to be used with these classes.

First, you remove the constructors from the classes derived from the base `Animal` class. You can remove these constructors because they set the animal's `Name` property, which you will do with object initializers instead. Alternatively, you could have added default constructors. In either case, when using default constructors, the `Name` property is initialized according to the default constructor in the base class, which has code as follows:

```
public Animal()
{
    name = "The animal with no name";
}
```

However, when an object initializer is used with a class that derives from `Animal`, recall that any properties set by the initializer are set after the object is instantiated, and therefore after this base class constructor is executed. If a value for the `Name` property is supplied as part of an object initializer, it will override this default value. In the example code, the `Name` property is set for all but one of the items added to the collection.

Second, you add an `Add()` method to the `Farm` class. This is in response to a series of compiler errors of the following form:

```
'Ch13Ex04.Farm<Ch13Ex04.Animal>' does not contain a definition for 'Add'
```

This error exposes part of the underlying functionality of collection initializers. Behind the scenes, the compiler calls the `Add()` method of a collection for each item that you supply in a collection initializer. The `Farm` class exposes a collection of `Animal` objects through a property called `Animals`. The compiler cannot guess that this is the property you want to populate (through `Animals.Add()`), so the code fails. To correct this problem, you add an `Add()` method to the class, which is initialized through the object initializer.

Alternatively, you could modify the code in the example to provide a nested initializer for the `Animals` property as follows:

```
static void Main(string[] args)
{
    Farm<Animal> farm = new Farm<Animal>
    {
        Animals =
        {
            new Cow { Name="Lea" },
            new Chicken { Name="Noa" },
            new Chicken(),
            new SuperCow { Name="Andrea" }
        }
    };
    farm.MakeNoises();
    ReadKey();
}
```

With this code there is no need to provide an `Add()` method for the `Farm` class. This alternative technique is appropriate when you have a class that contains multiple collections. In this case, there is no obvious candidate for a collection to add to with an `Add()` method of the containing class.

---

## TYPE INFERENCE

Earlier in this book you saw how C# is a *strongly typed* language, which means that every variable has a fixed type and can be used only in code that takes that type into account. In every code example you've seen so far, you have declared variables in one of two ways:

```
<type> <varName>;
<type> <varName> = <value>;
```

The following code shows at a glance what type of variable `<varName>` is:

```
int myInt = 5;
WriteLine(myInt);
```

You can also see that the IDE is aware of the variable type simply by hovering the mouse pointer over the variable identifier, as shown in Figure 13-12.



FIGURE 13-12

C# 3 introduced the new keyword `var`, which you can use as an alternative for `type` in the preceding code:

```
var <varName> = <value>;
```

In this code, the variable `<varName>` is *implicitly typed* to the type of `<value>`. Note that there is no type called `var`. In the code:

```
var myVar = 5;
```

`myVar` is a variable of type `int`, not of type `var`. Again, as shown in Figure 13-13, the IDE is aware of this.

This is an extremely important point. When you use `var` you are not declaring a variable with no type, or even a type that can change. If that were the case, C# would no longer be a strongly typed language. All you are doing is relying on the compiler to determine the type of the variable.

```
var myInt = 5;
WriteLine(myInt);
```

(local variable) int myInt

FIGURE 13-13

**NOTE** *The introduction of dynamic types in .NET 4 stretched the definition of C# being a strongly typed language, as you will see in the section “Dynamic Lookup” later in this chapter.*

If the compiler is unable to determine the type of variable declared using `var`, then your code will not compile. Therefore, you can't declare a variable using `var` without initializing the variable at the same time. If you do this, there is no value that the compiler can use to determine the type of the variable. The following code, therefore, will not compile:

```
var myVar;
```

The `var` keyword can also be used to infer the type of an array through the array initializer:

```
var myArray = new[] { 4, 5, 2 };
```

In this code, the type `myArray` is implicitly `int []`. When you implicitly type an array in this way, the array elements used in the initializer must be one of the following:

- All the same type
- All the same reference type or `null`
- All elements that can be implicitly converted to a single type

If the last of these rules is applied, then the type that elements can be converted to is referred to as the *best* type for the array elements. If there is any ambiguity as to what this best type might be — that is, if there are two or more types that all the elements can be implicitly converted to — your code will not compile. Instead, you receive the error indicating that no best type is available, as in the following code:

```
var myArray = new[] { 4, "not an int", 2 };
```

Note also that numeric values are never interpreted as nullable types, so the following code will not compile:

```
var myArray = new[] { 4, null, 2 };
```

You can, however, use a standard array initializer to make this work:

```
var myArray = new int?[] { 4, null, 2 };
```

A final point: The identifier `var` is not a forbidden identifier to use for a class name. This means, for example, that if your code has a class called `var` in scope (in the same namespace or in a referenced namespace), then you cannot use implicit typing with the `var` keyword.

In itself, type inference is not particularly useful because in the code you've seen in this section it only serves to complicate things. Using `var` makes it more difficult to see at a glance the type of a given variable. However, as you will see later in this chapter, the concept of inferred types is important because it underlies other techniques. The next subject, anonymous types, is one for which inferred types are essential.

## ANONYMOUS TYPES

After programming for a while you might find, especially in database applications, that you spend a lot of time creating simple, dull classes for data representation. It is not unusual to have families of classes that do absolutely nothing other than expose properties. The `Curry` class shown earlier in this chapter is a perfect example:

```
public class Curry
{
    public string MainIngredient { get; set; }
    public string Style { get; set; }
    public int Spiciness { get; set; }
}
```

This class doesn't do anything — it merely stores structured data. In database or spreadsheet terms, you could think of this class as representing a row in a table. A collection class that was capable of holding instances of this class would be a representation of multiple rows in a table or spreadsheet.

This is a perfectly acceptable use of classes, but writing the code for these classes can become monotonous, and any modifications to the underlying data schema requires you to add, remove, or modify the code that defines the classes.

*Anonymous types* are a way to simplify this programming model. The idea behind anonymous types is that rather than define these simple data storage types, you can instead use the C# compiler to automatically create types based on the data that you want to store in them.

The preceding `Curry` type can be instantiated as follows:

```
Curry curry = new Curry
{
    MainIngredient = "Lamb",
    Style = "Dhansak",
    Spiciness = 5
};
```

Alternatively, you could use an anonymous type, as in the following code:

```
var curry = new
{
    MainIngredient = "Lamb",
    Style = "Dhansak",
    Spiciness = 5
};
```

There are two differences here. First, the var keyword is used. That's because anonymous types do not have an identifier that you can use. Internally they do have an identifier, as you will see in a moment, but it is not available to you in your code. Second, no type name is specified after the new keyword. That's how the compiler knows you want to use an anonymous type.

The IDE detects the anonymous type definition and updates IntelliSense accordingly. With the preceding declaration, you can see the anonymous type, as shown in Figure 13-14.



FIGURE 13-14

Here, internally, the type of the variable `curry` is '`a`'. Obviously, you can't use this type in your code — it's not even a legal identifier name. The '`'` is simply the symbol used to denote an anonymous type in IntelliSense. IntelliSense also enables you to inspect the members of the anonymous type, as shown in Figure 13-15.




FIGURE 13-15

Note that the properties shown here are defined as *read-only* properties. This means that if you want to be able to change the values of properties in your data storage objects, you cannot use anonymous types.

The other members of anonymous types are implemented, as shown in the following Try It Out.

### TRY IT OUT Using Anonymous Types: Ch13Ex05\Program.cs

1. Create a new console application called Ch13Ex05 and save it in the directory C:\BegVCSharp\Chapter13.
2. Modify the code in Program.cs as follows:

```
static void Main(string[] args)
{
    var curries = new[]
    {
        new { MainIngredient = "Lamb", Style = "Dhansak", Spiciness = 5 },
        new { MainIngredient = "Lamb", Style = "Dhansak", Spiciness = 5 },
        new { MainIngredient = "Chicken", Style = "Dhansak", Spiciness = 5 }
    };
    WriteLine(curries[0].ToString());
    WriteLine(curries[0].GetHashCode());
    WriteLine(curries[1].GetHashCode());
    WriteLine(curries[2].GetHashCode());
    WriteLine(curries[0].Equals(curries[1]));
    WriteLine(curries[0].Equals(curries[2]));
    WriteLine(curries[0] == curries[1]);
    WriteLine(curries[0] == curries[2]);
    ReadKey();
}
```

3. Run the application. The result is shown in Figure 13-16.




FIGURE 13-16

### How It Works

In this example you create an array of anonymous type objects that you then proceed to use to perform tests of the members supplied by anonymous types. The code to create the array of anonymously typed objects is as follows:

```
var curries = new[]
{
    new { MainIngredient = "Lamb", Style = "Dhansak", Spiciness = 5 },
    ...
};
```

This uses an array that is implicitly typed to an anonymous type, using a combination of syntax from this section and the “Type Inference” section earlier in this chapter. The result is that the `curries` variable contains three instances of an anonymous type.

The first thing the code does after creating this array is output the result of calling `Tostring()` on the anonymous type:

```
WriteLine(curries[0].ToString());
```

This results in the following output:

```
{ MainIngredient = Lamb, Style = Dhansak, Spiciness = 5 }
```

The implementation of `Tostring()` in an anonymous type outputs the values of each property defined for the type.

The code next calls `GetHashCode()` on each of the array’s three objects:

```
WriteLine(curries[0].GetHashCode());
WriteLine(curries[1].GetHashCode());
WriteLine(curries[2].GetHashCode());
```

When implemented, `GetHashCode()` should return a unique integer for an object based on the object’s state. The first two objects in the array have the same property values, and therefore the same state.

The result of these calls is the same integer for each of these objects, but a different integer for the third object. The output is as follows:

```
1789653062
1789653062
2116426892
```

Next, the `Equals()` method is called to compare the first object with the second object, and then to compare the first object with the third object:

```
WriteLine(curries[0].Equals(curries[1]));
WriteLine(curries[0].Equals(curries[2]));
```

The result is as follows:

```
True
False
```

The implementation of `Equals()` in anonymous types compares the state of objects. The result is `true` where every property of one object contains the same value as the comparable property on another object.

That is not what happens when you use the `==` operator, however. The `==` operator, as shown in previous chapters, compares object references. The last section of code performs the same comparisons as the previous section of code but uses `==` instead of `Equals()`:

```
WriteLine(curries[0] == curries[1]);
WriteLine(curries[0] == curries[2]);
```

Each entry in the `curries` array refers to a different instance of the anonymous type, so the result is `false` in both cases. The output is as expected:

```
False  
False
```

Interestingly, when you create instances of the anonymous types in this example, the compiler notices that the parameters are the same and creates three instances of the *same* anonymous type — not three separate anonymous types. However, this doesn't mean that when you instantiate an object from an anonymous type the compiler looks for a type to match it with. Even if you have defined a class elsewhere that has matching properties, if you use anonymous type syntax, then an anonymous type will be created (or reused as in this example).

---

## DYNAMIC LOOKUP

The `var` keyword, as described earlier, is not in itself a type, and so doesn't break the “strongly typed” methodology of C#. From C# 4 onward, though, things have become a little less fixed. C# 4 introduced the concept of *dynamic variables*, which, as their name suggests, are variables that do not have a fixed type.

The main motivation for this is that there are many situations where you will want to use C# to manipulate objects created by another language. This includes interoperability with older technologies such as the Component Object Model (COM), as well as dealing with dynamic languages such as JavaScript, Python, and Ruby. Without going into too much implementation detail, using C# to access methods and properties of objects created by these languages has, in the past, involved awkward syntax. For example, say you had code that obtained an object from JavaScript with a method called `Add()` that added two numbers. Without dynamic lookup, your code to call this method might look something like the following:

```
ScriptObject jsObj = SomeMethodThatGetsTheObject();  
int sum = Convert.ToInt32(jsObj.Invoke("Add", 2, 3));
```

The `ScriptObject` type (not covered in depth here) provides a way to access a JavaScript object, but even this is unable to give you the capability to do the following:

```
int sum = jsObj.Add(2, 3);
```

Dynamic lookup changes everything — enabling you to write code just like the preceding. However, as you will see in the following sections, this power comes at a price.

Another situation in which dynamic lookup can assist you is when you are dealing with a C# object whose type you do not know. This might sound like an odd situation, but it happens more often than you might think. It is also an important capability when writing generic code that can deal with whatever input it receives. The “old” way to deal with this situation is called *reflection*, which involves using type information to access types and members. The syntax for using reflection to access type members such as methods is quite similar to the syntax used to access the JavaScript object, as shown in the preceding code. In other words, it's messy.

Under the hood, dynamic lookup is supported by the Dynamic Language Runtime (DLR). This is part of .NET 4.5, just as the CLR is. An exact description of the DLR and how it makes interoperability easier is beyond the scope of this book; here you’re more interested in how to use it in C#.

## The dynamic Type

C# 4 introduced the `dynamic` keyword, which you can use to define variables, as in this example:

```
dynamic myDynamicVar;
```

Unlike the `var` keyword introduced earlier, there really is a `dynamic` type, so there is no need to initialize the value of `myDynamicVar` when it is declared.

**NOTE** *Unusually, the dynamic type exists only at compile time; at runtime the System.Object type is used instead. This is a minor implementation detail but one that is worth remembering, as it might clarify some of the discussion that follows.*

Once you have a dynamic variable, you can proceed to access its members (the code to obtain a value for the variable is not shown here):

```
myDynamicVar.DoSomething("With this!");
```

Regardless of the value that `myDynamicVar` contains, this code will compile. However, if the requested member does not exist, you will get an exception when this code is executed, of type `RuntimeBinderException`.

In effect, what you are doing with code like this is providing a “recipe” that should be applied at runtime. The value of `myDynamicVar` will be examined, and a method called `DoSomething()` with a single string parameter will be located and called at the point where it is required.

This is best illustrated with an example.

**WARNING** *The following example is for illustrative purposes only! In general, you should use dynamic types only when they are your only option — for example, when you are dealing with non-.NET objects.*

### TRY IT OUT Using Dynamic Types: Ch13Ex06\Program.cs

1. Create a new console application called Ch13Ex06 and save it in the directory `C:\BegVCSharp\Chapter13`.
2. Modify the code in `Program.cs` as follows:

```
using System;
```

```

using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using Microsoft.CSharp.RuntimeBinder;
namespace Ch13Ex06
{
    class MyClass1
    {
        public int Add(int var1, int var2) => var1 + var2;
    }
    class MyClass2 {}
    class Program
    {
        static int callCount = 0;
        static dynamic GetValue()
        {
            if (callCount++ == 0)
            {
                return new MyClass1();
            }
            return new MyClass2();
        }
        static void Main(string[] args)
        {
            try
            {
                dynamic firstResult = GetValue();
                dynamic secondResult = GetValue();
                WriteLine($"firstResult is: {firstResult.ToString()}");
                WriteLine($"secondResult is: {secondResult.ToString()}");
                WriteLine($"firstResult call: {firstResult.Add(2, 3)}");
                WriteLine($"secondResult call: {secondResult.Add(2, 3)}");
            }
            catch (RuntimeBinderException ex)
            {
                WriteLine(ex.Message);
            }
            ReadKey();
        }
    }
}

```

3. Run the application. The result is shown in Figure 13-17.




FIGURE 13-17

## How It Works

In this example you use a method that returns one of two types of objects to obtain a dynamic value, and then attempts to use the object obtained. The code compiles without any trouble, but an exception is thrown (and handled) when an attempt is made to access a non-existent method.

To begin, you add a `using` statement for the namespace that contains the `RuntimeBinderException` exception:

```
using Microsoft.CSharp.RuntimeBinder;
```

Next, you define two classes, `MyClass1` and `MyClass2`, where `MyClass1` has an `Add()` method and `MyClass2` has no members:

```
class MyClass1
{
    public int Add(int var1, int var2) => var1 + var2;
}
class MyClass2
{
}
```

You also add a field (`callCount`) and a method (`GetValue()`) to the `Program` class to provide a way to obtain an instance of one of these classes:

```
static int callCount = 0;
static dynamic GetValue()
{
    if (callCount++ == 0)
    {
        return new MyClass1();
    }
    return new MyClass2();
}
```

A simple call counter is used so that this method returns an instance of `MyClass1` the first time it is called, and instances of `MyClass2` thereafter. Note that the `dynamic` keyword can be used as a return type for a method.

Next, the code in `Main()` calls the `GetValue()` method twice and then attempts to call `GetString()` and `Add()` on both values returned in turn. This code is placed in a `try...catch` block to trap any exceptions of type `RuntimeBinderException` that occur:

```
static void Main(string[] args)
{
    try
    {
        dynamic firstResult = GetValue();
        dynamic secondResult = GetValue();
        WriteLine($"firstResult is: {firstResult.ToString()}");
        WriteLine($"secondResult is: {secondResult.ToString()}");
```

```
        WriteLine($"firstResult call: {firstResult.Add(2, 3)}");
        WriteLine($"secondResult call: {secondResult.Add(2, 3)}");
    }
    catch (RuntimeBinderException ex)
    {
        WriteLine(ex.Message);
    }
    ReadKey();
}
```

Sure enough, an exception is thrown when `secondResult.Add()` is called, as no such method exists on `MyClass2`. The exception message tells you exactly that.

The `dynamic` keyword can also be used in other places where a type name is required, such as for method parameters. You could rewrite the `Add()` method as follows:

```
public int Add(dynamic var1, dynamic var2) => var1 + var2;
```

This would have no effect on the result. In this case, at runtime the values passed to `var1` and `var2` are inspected to determine whether a compatible operator definition for `+` exists. In the case of two `int` values being passed, such an operator does exist. If incompatible values are used, a `RuntimeBinderException` exception is thrown. For example, if you try,

```
WriteLine("firstResult call: {0}", firstResult.Add("2", 3));
```

the exception message will be as follows:

```
Cannot implicitly convert type 'string' to 'int'
```

---

The lesson to be learned here is that dynamic types are very powerful, but there's a warning to learn too. These sorts of exceptions are entirely avoidable if you use strong typing instead of dynamic typing. For most C# code that you write, avoid the `dynamic` keyword. However, if a situation arises where you need to use it, use it and love it — and spare a thought for those poor programmers of the past who didn't have this powerful tool at their disposal.

## ADVANCED METHOD PARAMETERS

C# 4 extended what is possible when defining and using method parameters. This is primarily in response to a specific problem that arises when using interfaces defined externally, such as the Microsoft Office programming model. Here, certain methods expose a vast number of parameters, many of which are not required for every call. In the past, this has meant that a way to specify missing parameters has been necessary, or that a lot of nulls appear in code:

```
RemoteCall(var1, var2, null, null, null, null, null);
```

In this code it is not at all obvious what the `null` values refer to, or why they have been omitted.

Perhaps, in an ideal world, there would be multiple overloads of this `RemoteCall()` method, including one that only required two parameters as follows:

```
RemoteCall(var1, var2);
```

However, this would require many more methods with alternative combinations of parameters, which in itself would cause more problems (more code to maintain, increased code complexity, and so on).

Languages such as Visual Basic have dealt with this situation in a different way, by allowing named and optional parameters. From version 4, this became possible in C#, demonstrating one way in which the evolution of all .NET languages is converging.

In the following sections you will see how to use these parameter types.

## Optional Parameters

Often when you call a method, you pass in the same value for a particular parameter. This can be a Boolean value, for example, which might control a nonessential part of the method's operation. To be more specific, consider the following method definition:

```
public List<string> GetWords(string sentence, bool capitalizeWords)
{
    ...
}
```

Regardless of the value passed into the `capitalizeWords` parameter, this method will return a list of `string` values, each of which is a word from the input sentence. Depending on how this method was used, you might occasionally want to capitalize the list of words returned (perhaps you are formatting a heading such as the one for this section, "Optional Parameters"). In most cases, though, you might not want to do this, so most calls would be as follows:

```
List<string> words = GetWords(sentence, false);
```

To make this the "default" behavior, you might declare a second method as follows:

```
public List<string> GetWords(string sentence) => GetWords(sentence, false);
```

This method calls into the second method, passing a value of `false` for `capitalizeWords`.

There is nothing wrong with doing this, but you can probably imagine how complicated this would become in a situation where many more parameters were used.

An alternative is to make the `capitalizeWords` parameter an *optional parameter*. This involves defining the parameter as optional in the method definition by providing a default value that will be used if none is supplied, as follows:

```
public List<string> GetWords(string sentence, bool capitalizeWords = false)
{
    ...
}
```

If you were to define a method in this way, then you could supply either one or two parameters, where the second parameter is required only if you want `capitalizeWords` to be `true`.

## Optional Parameter Values

As described in the previous section, a method definition defines an optional parameter with syntax as follows:

```
<parameterType> <parameterName> = <defaultValue>
```

There are restrictions on what you can use for the `<defaultValue>` default value. Default values must be literal values, constant values, or default value type values. The following, therefore, will not compile:

```
public bool CapitalizationDefault;
public List<string> GetWords(string sentence,
    bool capitalizeWords = CapitalizationDefault)
{
    ...
}
```

In order to make this work, the `CapitalizationDefault` value must be defined as a constant:

```
public const bool CapitalizationDefault = false;
```

Whether it makes sense to do this depends on the situation; in most cases you will probably be better off providing a literal value as in the previous section.

## The `OptionalAttribute` Attribute

As an alternative to the syntax described in the previous sections, you can define optional parameters using the `OptionalAttribute` attribute as follows:

```
[Optional] <parameterType> <parameterName>
```

This attribute is found in the `System.Runtime.InteropServices` namespace. Note that if you use this syntax there is no way to provide a default value for the parameter.

## Optional Parameter Order

When you use optional values, they must appear at the end of the list of parameters for a method. No parameters without default values can appear after any parameters with default values.

The following code, therefore, is illegal:

```
public List<string> GetWords(bool capitalizeWords = false, string sentence)
{
    ...
}
```

Here, `sentence` is a required parameter, and must therefore appear before the optional `capitalizedWords` parameter.

## Named Parameters

When you use optional parameters, you might find yourself in a situation where a particular method has several optional parameters. It's not beyond the realm of the imagination, then, to conceive of

a situation where you want to pass a value to, say, only the third optional parameter. With just the syntax from the previous section there is no way to do this without supplying values for the first and second optional parameters.

C# 4 also introduced *named parameters* that enable you to specify whichever parameters you want. This doesn't require you to do anything in particular in your method definition; it is a technique that you use when you are calling a method. The syntax is as follows:

```
MyMethod(
    <param1Name>: <param1Value>,
    ...
    <paramNName>: <paramNValue>);
```

The names of parameters are the names of the variables used in the method definition.

You can specify any number of parameters you like in this way, as long as the named parameters exist, and you can do so in any order. Named parameters can be optional as well.

You can, if you want, use named parameters for only some of the parameters in a method call. This is particularly useful when you have several optional parameters in a method signature, but some required parameters. You might specify the required parameters first, then finish off with named optional parameters. For example:

```
MyMethod(
    requiredParameter1Value,
    optionalParameter5: optionalParameter5Value);
```

If you mix named and positional parameters, though, note that you must include all positional parameters first, before the named parameters. However, you can use a different order if you prefer as long as you use named parameters throughout, as in this example:

```
MyMethod(
    optionalParameter5: optionalParameter5Value,
    requiredParameter1: requiredParameter1Value);
```

If you do this you must include values for all required parameters.

In the following Try It Out, you will see how you can use named and optional parameters.

## TRY IT OUT Using Named and Optional Parameters: Ch13Ex07

1. Create a new console application called Ch13Ex07 and save it in the directory C:\BegVCSharp\Chapter13.
2. Add a class called WordProcessor to the project and modify its code as follows:

```
public static class WordProcessor
{
    public static List<string> GetWords(
        string sentence,
        bool capitalizeWords = false,
        bool reverseOrder = false,
        bool reverseWords = false)
```

```
        List<string> words = new List<string>(sentence.Split(' '));
        if (capitalizeWords)
            words = CapitalizeWords(words);
        if (reverseOrder)
            words = ReverseOrder(words);
        if (reverseWords)
            words = ReverseWords(words);
        return words;
    }
    private static List<string> CapitalizeWords(List<string> words)
    {
        List<string> capitalizedWords = new List<string>();
        foreach (string word in words)
        {
            if (word.Length == 0)
                continue;
            if (word.Length == 1)
                capitalizedWords.Add(
                    word[0].ToString().ToUpper());
            else
                capitalizedWords.Add(
                    word[0].ToString().ToUpper()
                    + word.Substring(1));
        }
        return capitalizedWords;
    }
    private static List<string> ReverseOrder(List<string> words)
    {
        List<string> reversedWords = new List<string>();
        for (int wordIndex = words.Count - 1;
            wordIndex >= 0; wordIndex--)
            reversedWords.Add(words[wordIndex]);
        return reversedWords;
    }
    private static List<string> ReverseWords(List<string> words)
    {
        List<string> reversedWords = new List<string>();
        foreach (string word in words)
            reversedWords.Add(ReverseWord(word));
        return reversedWords;
    }
    private static string ReverseWord(string word)
    {
        StringBuilder sb = new StringBuilder();
        for (int characterIndex = word.Length - 1;
            characterIndex >= 0; characterIndex--)
            sb.Append(word[characterIndex]);
        return sb.ToString();
    }
}
```

3. Modify the code in Program.cs as follows:

```
static void Main(string[] args)
{
```

```

string sentence = "his gaze against the sweeping bars has "
    + "grown so weary";
List<string> words;
words = WordProcessor.GetWords(sentence);
WriteLine("Original sentence:");
foreach (string word in words)
{
    Write(word);
    Write(' ');
}
WriteLine('\n');
words = WordProcessor.GetWords(
    sentence,
    reverseWords: true,
    capitalizeWords: true);
WriteLine("Capitalized sentence with reversed words:");
foreach (string word in words)
{
    Write(word);
    Write(' ');
}
.ReadKey();
}

```

4. Run the application. The result is shown in Figure 13-18.




FIGURE 13-18

## How It Works

In this example you have created a utility class that performs some simple string manipulation, and used that class to modify a string. The single public method exposed by the class contains one required parameter and three optional ones:

```

public static List<string> GetWords(
    string sentence,
    bool capitalizeWords = false,
    bool reverseOrder = false,
    bool reverseWords = false)
{
    ...
}

```

This method returns a collection of `string` values, each of which is a word from the original input. Depending on which (if any) of the three optional parameters are specified, additional transformations can be made — on the string collection as a whole or on individual word values.

**NOTE** You won't look at the functionality of the WordProcessor class in any more depth here; you are free to browse the code at your leisure. Along the way you might like to think about how this code might be improved. For example, should the word 'twas be capitalized as 'Twas? How would you go about making that change?

When this method is called, only two of the available optional parameters are used; the third parameter (`reverseOrder`) will have its default value of `false`:

```
words = WordProcessor.GetWords(
    sentence,
    reverseWords: true,
    capitalizeWords: true);
```

Also, note that the two parameters specified are placed in a different order from how they are defined.

As a final point to note, IntelliSense can be quite handy when dealing with methods that have optional parameters. When entering the code for this Try It Out, you might have noticed the tooltip for the `GetWords()` method, shown in Figure 13-19 (you can also see this tooltip by hovering the mouse pointer over the method call as shown).




FIGURE 13-19

This is a very useful tooltip, as it shows not only the names of available parameters, but also the default values for optional parameters, so you can tell at a glance if you need to override a particular default.

## LAMBDA EXPRESSIONS

*Lambda expressions* are a construct introduced in C# 3 that you can use to simplify certain aspects of C# programming, in particular when combined with LINQ. They can be difficult to grasp at first, mainly because they are so flexible in their usage. Lambda expressions are extremely useful when combined with other C# language features, such as anonymous methods. Without looking at LINQ, a subject left until later in the book, anonymous methods are the best entry point for examining this subject. Start with a quick refresher.

## Anonymous Methods Recap

Previously in this chapter you learned about anonymous methods — methods that you supply inline, where a delegate type variable would otherwise be required. When you add an event handler to an event, the sequence of events is as follows:

1. Define an event handler method whose return type and parameters match those of the delegate required for the event to which you want to subscribe.
2. Declare a variable of the delegate type used for the event.
3. Initialize the delegate variable to an instance of the delegate type that refers to the event handler method.
4. Add the delegate variable to the list of subscribers for the event.

In practice, things are a bit simpler than this because you typically won't bother with a variable to store the delegate — you will just use an instance of the delegate when you subscribe to the event.

This was the case when you previously used the following code:

```
Timer myTimer = new Timer(100);
myTimer.Elapsed += new ElapsedEventHandler(WriteChar);
```

This code subscribes to the `Elapsed` event of a `Timer` object. This event uses the `ElapsedEventHandler` delegate type, which is instantiated using a method identifier, `WriteChar`. The result here is that when the `Timer` raises the `Elapsed` event, the `WriteChar()` method is called. The parameters passed to `WriteChar()` depend on the parameter types defined by the `ElapsedEventHandler` delegate and the values passed by the code in `Timer` that raises the event.

In fact, the C# compiler can achieve the same result with even less code through method group syntax:

```
myTimer.Elapsed += WriteChar;
```

The C# compiler knows the delegate type required by the `Elapsed` event, so it can fill in the blanks. However, you should use this syntax with care because it can make it harder to read your code and know exactly what is happening. When you use an anonymous method, the sequence of events shown earlier is reduced to a single step:

1. Use an inline, anonymous method that matches the return type and the parameters of the delegate required by an event to subscribe to that event.

The inline, anonymous method is defined by using the `delegate` keyword:

```
myTimer.Elapsed +=  
    delegate(object source, ElapsedEventArgs e)  
    {  
        WriteLine("Event handler called after {0} milliseconds.",  
            (source as Timer).Interval);  
    };
```

This code works just as well as using the event handler separately. The main difference is that the anonymous method used here is effectively hidden from the rest of your code. You cannot, for example, reuse this event handler elsewhere in your application. In addition, the syntax used here is, for want of a better description, a little clunky. The `delegate` keyword is immediately confusing because it is effectively being overloaded — you use it both for anonymous methods and for defining delegate types.

## Lambda Expressions for Anonymous Methods

This brings you to lambda expressions. Lambda expressions are a way to simplify the syntax of anonymous methods. In fact, they are more than that, but this section will keep things simple for now. Using a lambda expression, you can rewrite the code at the end of the previous section as follows:

```
myTimer.Elapsed += (source, e) => WriteLine("Event handler called after " +  
    $"{(source as Timer).Interval} milliseconds.");
```

At first glance this looks...well, a little baffling (unless you are familiar with so-called functional programming languages such as Lisp or Haskell, that is). However, if you look closer you can see, or at least infer, how this works and how it relates to the anonymous method that it replaces. The lambda expression is made up of three parts:

- A list of (untyped) parameters in parentheses
- The `=>` operator
- A C# statement

The types of the parameters are inferred from the context, using the same logic shown in the section “Anonymous Types” earlier in this chapter. The `=>` operator simply separates the parameter list from the expression body. The expression body is executed when the lambda expression is called.

The compiler takes this lambda expression and creates an anonymous method that works exactly the same way as the anonymous method in the previous section. In fact, it will be compiled into the same or similar Common Intermediate Language (CIL) code.

The following Try It Out clarifies what occurs in lambda expressions.

### TRY IT OUT Using Simple Lambda Expressions: Ch13Ex08\Program.cs

1. Create a new console application called Ch13Ex08 and save it in the directory C:\BegVCSharp\Chapter13.

2. Modify the code in Program.cs as follows:

```
namespace Ch13Ex08
{
    delegate int TwoIntegerOperationDelegate(int paramA, int paramB);
    class Program
    {
        static void PerformOperations(TwoIntegerOperationDelegate del)
        {
            for (int paramAVal = 1; paramAVal <= 5; paramAVal++)
            {
                for (int paramBVal = 1; paramBVal <= 5; paramBVal++)
                {
                    int delegateCallResult = del(paramAVal, paramBVal);
                    Write($"f({paramAVal}, {paramBVal})={delegateCallResult}");
                    if (paramBVal != 5)
                    {
                        Write(", ");
                    }
                }
                WriteLine();
            }
        }
        static void Main(string[] args)
        {
            WriteLine("f(a, b) = a + b:");
            PerformOperations((paramA, paramB) => paramA + paramB);
            WriteLine();
            WriteLine("f(a, b) = a * b:");
            PerformOperations((paramA, paramB) => paramA * paramB);
            WriteLine();
            WriteLine("f(a, b) = (a - b) % b:");
            PerformOperations((paramA, paramB) => (paramA - paramB)
                % paramB);
            ReadKey();
        }
    }
}
```

3. Run the application. The result is shown in Figure 13-20.

```

file:///C:/BegVCSharp/Chapter13/Ch13Ex08/Ch13Ex08/bin/...
f<a, b> = a + b:
f<1,1>=2, f<1,2>=3, f<1,3>=4, f<1,4>=5, f<1,5>=6
f<2,1>=3, f<2,2>=4, f<2,3>=5, f<2,4>=6, f<2,5>=7
f<3,1>=4, f<3,2>=5, f<3,3>=6, f<3,4>=7, f<3,5>=8
f<4,1>=5, f<4,2>=6, f<4,3>=7, f<4,4>=8, f<4,5>=9
f<5,1>=6, f<5,2>=7, f<5,3>=8, f<5,4>=9, f<5,5>=10

f<a, -b> = a * (-b):
f<1, -1>=-1, f<1, -2>=-2, f<1, -3>=-3, f<1, -4>=-4, f<1, -5>=-5
f<2, -1>=-2, f<2, -2>=-4, f<2, -3>=-6, f<2, -4>=-8, f<2, -5>=-10
f<3, -1>=-3, f<3, -2>=-6, f<3, -3>=-9, f<3, -4>=-12, f<3, -5>=-15
f<4, -1>=-4, f<4, -2>=-8, f<4, -3>=-12, f<4, -4>=-16, f<4, -5>=-20
f<5, -1>=-5, f<5, -2>=-10, f<5, -3>=-15, f<5, -4>=-20, f<5, -5>=-25

```

FIGURE 13-20

## How It Works

This example uses lambda expressions to generate functions that can be used to return the result of performing specific processing on two input parameters. Those functions then operate on 25 pairs of values and output the results to the console.

You start by defining a delegate type called `TwoIntegerOperationDelegate` to represent a method that takes two `int` parameters and returns an `int` result:

```
delegate int TwoIntegerOperationDelegate(int paramA, int paramB);
```

This delegate type is used later when you define your lambda expressions. These lambda expressions compile into methods whose return type and parameter types match this delegate type, as you will see shortly.

Next, you add a method called `PerformOperations()`, which takes a single parameter of type `TwoIntegerOperationDelegate`:

```
static void PerformOperations(TwoIntegerOperationDelegate del)
{
```

The idea behind this method is that you can pass it a delegate instance (or an anonymous method or lambda expression, because these constructs compile to delegate instances) and the method will call the method represented by the delegate instance with an assortment of values:

```
for (int paramAVal = 1; paramAVal <= 5; paramAVal++)
{
    for (int paramBVal = 1; paramBVal <= 5; paramBVal++)
    {
        int delegateCallResult = del(paramAVal, paramBVal);
```

The parameters and results are then output to the console:

```
Write($"f({paramAVal}, " +
      $"{{paramBVal}})={{delegateCallResult}}");
if (paramBVal != 5)
{
    Write(", ");
```

```
        }  
    }  
    }  
    WriteLine();
```

In the `Main()` method you create three lambda expressions and use them to call `PerformOperations()` in turn. The first of these calls is as follows:

```
WriteLine("f(a, b) = a + b:");
PerformOperations((paramA, paramB) => paramA + paramB);
```

The lambda expression used here is as follows:

(paramA, paramB) => paramA + paramB

Again, this breaks down into three parts:

1. A parameter definition section. Here there are two parameters, `paramA` and `paramB`. These parameters are untyped, meaning the compiler can infer the types of these parameters according to the context. In this case the compiler can determine that the `PerformOperations()` method call requires a delegate of type `TwoIntegerOperationDelegate`. This delegate type has two `int` parameters, so by inference both `paramA` and `paramB` are typed as `int` variables.
  2. The `=>` operator. This separates the lambda expression parameters from the lambda expression body.
  3. The expression body. This specifies a simple operation, which is the summation of `paramA` and `paramB`. Notice that there is no need to specify that this is a return value. The compiler knows that in order to create a method that can be used with `TwoIntegerOperationDelegate`, the method must have a return type of `int`. Because the operation specified, `paramA + paramB`, evaluates to an `int`, and no additional information is supplied, the compiler infers that the result of this expression should be the return type of the method.

In longhand then, you can expand the code that uses this lambda expression to the following code that uses an anonymous method:

```
WriteLine("f(a, b) = a + b:");
PerformOperations(delegate(int paramA, int paramB)
{
    return paramA + paramB;
});
```

The remaining code performs operations using two different lambda expressions in the same way:

```
    WriteLine();
    WriteLine("f(a, b) = a * b:");
    PerformOperations((paramA, paramB) => paramA * paramB);
    WriteLine();
    WriteLine("f(a, b) = (a - b) % b:");
    PerformOperations((paramA, paramB) => (paramA - paramB)
        % paramB);
    ReadKey();
```

The last lambda expression involves more calculations but is no more complicated than the others. The syntax for lambda expressions enables you to perform far more complicated operations, as you will see shortly.

## Lambda Expression Parameters

In the code you have seen so far, the lambda expressions have used type inference to determine the types of the parameters passed. In fact, this is not mandatory; you can define types if you want. For example, you could use the following lambda expression:

```
(int paramA, int paramB) => paramA + paramB
```

This has the advantage of making your code more readable, although you lose out in both brevity and flexibility. You could use the implicitly typed lambda expressions from the previous Try It Out for delegate types that used other numeric types, such as `long` variables.

Note that you cannot use implicit and explicit parameter types in the same lambda expression. The following lambda expressions will not compile because `paramA` is explicitly typed and `paramB` is implicitly typed:

```
(int paramA, paramB) => paramA + paramB
```

Parameter lists in lambda expressions always consist of a comma-separated list of either all implicitly typed parameters or all explicitly typed parameters. If you have only one implicitly typed parameter, then you can omit the parentheses; otherwise, they are required as part of the parameter list, as shown earlier. For example, you could have the following as a single-parameter, implicitly typed lambda expression:

```
param1 => param1 * param1
```

You can also define lambda expressions that have no parameters. This is denoted by using empty parentheses, `()`:

```
() => Math.PI
```

This could be used where a delegate requiring no parameters but returning a `double` value is required.

## Lambda Expression Statement Bodies

In all the code that you have seen so far, a single expression has been used in the statement body of lambda expressions. You have also seen how this single expression has been interpreted as the return value of the lambda expression, which is, for example, how you can use the expression `paramA + paramB` as the statement body for a lambda expression for a delegate with a return type of `int` (assuming both `paramA` and `paramB` are implicitly or explicitly typed to `int` values, as they were in the example code).

An earlier example showed how a delegate with a `void` return type was less fussy about the code used in the statement body:

```
myTimer.Elapsed += (source, e) => WriteLine("Event handler called after " +
    $"{(source as Timer).Interval} milliseconds.");
```

Here, the statement doesn't evaluate to anything, so it is simply executed without any return value being used anywhere.

Given that lambda expressions can be visualized as an extension of the anonymous method syntax, you might not be surprised to learn that you can also include multiple statements as a lambda expression statement body. To do so, you simply provide a block of code enclosed in curly braces, much like any other situation in C# where you must supply multiple lines of code:

```
(param1, param2) =>
{
    // Multiple statements ahoy!
}
```

If you use a lambda expression in combination with a delegate type that has a non-`void` return type, then you must return a value with the `return` keyword, just like any other method:

```
(param1, param2) =>
{
    // Multiple statements ahoy!
    return returnValue;
}
```

For example, earlier you saw how you could rewrite the following code from the Try It Out,

```
PerformOperations((paramA, paramB) => paramA + paramB);
```

as:

```
PerformOperations(delegate(int paramA, int paramB)
{
    return paramA + paramB;
});
```

Alternatively, you could rewrite the code as follows:

```
PerformOperations((paramA, paramB) =>
{
    return paramA + paramB;
});
```

This is more in keeping with the original code because it maintains implicit typing of the `paramA` and `paramB` parameters.

For the most part, lambda expressions are at their most useful — and certainly their most elegant — when used with single expressions. To be honest, if you require multiple statements, your code

might read much better if you define a separate, non-anonymous method to use instead of a lambda expression; that also makes your code more reusable.

## Lambda Expressions as Delegates and Expression Trees

You have already seen some of the differences between lambda expressions and anonymous methods where lambda methods have more flexibility — for example, implicitly typed parameters. At this point it is worth noting another key difference, although the implications of this will not become apparent until later in the book when you learn about LINQ.

You can interpret a lambda expression in two ways. The first way, which you have seen throughout this chapter, is as a delegate. That is, you can assign a lambda expression to a delegate type variable, as you did in the previous Try It Out.

In general terms, you can represent a lambda expression with up to eight parameters as one of the following generic types, all defined in the `System` namespace:

- `Action` for lambda expressions with no parameters and a return type of `void`
- `Action<>` for lambda expressions with up to eight parameters and a return type of `void`
- `Func<>` for lambda expressions with up to eight parameters and a return type that is not `void`

`Action<>` has up to eight generic type parameters, one for each parameter, and `Func<>` has up to nine generic type parameters, used for up to eight parameters and the return type. In `Func<>`, the return type is always the last in the list.

For example, the following lambda expression, which you saw earlier:

```
(int paramA, int paramB) => paramA + paramB
```

This expression can be represented as a delegate of type `Func<int, int, int>` because it has two parameters and a return type all of type `int`. Note that you can use these generic delegate types instead of defining your own in many circumstances. For example, you can use them instead of the `TwoIntegerOperationDelegate` delegate you defined in the previous Try It Out.

The second way to interpret a lambda expression is as an *expression tree*. An expression tree is an abstract representation of a lambda expression; and as such, it cannot be executed directly. Instead, you can use the expression tree to analyze the lambda expression programmatically and perform actions in response to the lambda expression.

This is, obviously, a complicated subject. However, expression trees are critical to the LINQ functionality you will learn about later in this book. To give a more concrete example, the LINQ framework includes a generic class called `Expression<>`, which you can use to encapsulate a lambda expression. One of the ways in which this class is used is to take a lambda expression that you have written in C# and convert it into an equivalent SQL script representation for executing directly in a database.

You don't need to know any more about that at this point. When you encounter this functionality later in the book, you will be better equipped to understand what is going on, as you now have a thorough grounding in the key concepts that the C# language provides.

## Lambda Expressions and Collections

Now that you have learned about the `Func<>` generic delegate, you can understand some of the extension methods that the `System.Linq` namespace provides for array types (which you might have seen popping up in IntelliSense at various points during your coding). For example, there is an extension method called `Aggregate()`, which is defined with three overloads as follows:

```
public static TSource Aggregate<TSource>(
    this IEnumerable<TSource> source,
    Func<TSource, TSource, TSource> func);
public static TAccumulate Aggregate<TSource, TAccumulate>(
    this IEnumerable<TSource> source,
    TAccumulate seed,
    Func<TAccumulate, TSource, TAccumulate> func);
public static TResult Aggregate<TSource, TAccumulate,
    Aggregate<TSource, TAccumulate, TResult>() TResult>(
    this IEnumerable<TSource> source,
    TAccumulate seed,
    Func<TAccumulate, TSource, TAccumulate> func,
    Func<TAccumulate, TResult> resultSelector);
```

As with the extension method shown earlier, this looks at first glance to be impenetrable, but if you break it down you can work it out easily enough. The IntelliSense for this function tells you that it does the following:

Applies an accumulator function over a sequence.

This means that an accumulator function (which you can supply in the form of a lambda expression) will be applied to each element in a collection from beginning to end. This accumulator function must have two parameters and one return value. One input is the current element; the other input is either a seed value, the first value in the collection, or the result of the previous evaluation.

In the simplest of the three overloads, there is only one generic type specification, which can be inferred from the type of the instance parameter. For example, in the following code the generic type specification will be `int` (the accumulator function is left blank for now):

```
int[] myIntArray = { 2, 6, 3 };
int result = myIntArray.Aggregate(...);
```

This is equivalent to the following:

```
int[] myIntArray = { 2, 6, 3 };
int result = myIntArray.Aggregate<int>(...);
```

The lambda expression that is required here can be deduced from the extension method specification. Because the type `TSource` is `int` in this code, you must supply a lambda expression for the delegate `Func<int, int, int>`. For example, you could use one you've seen before:

```
int[] myIntArray = { 2, 6, 3 };
int result = myIntArray.Aggregate((paramA, paramB) => paramA + paramB);
```

This call results in the lambda expression being called twice, first with `paramA = 2` and `paramB = 6`, and once with `paramA = 8` (the result of the first calculation) and `paramB = 3`. The final result

assigned to the variable `result` will be the `int` value 11 — the summation of all the elements in the array.

The other two overloads of the `Aggregate()` extension method are similar, but enable you to perform slightly more complicated processing. This is illustrated in the following short Try It Out.

### TRY IT OUT Using Lambda Expressions with Collections: Ch13Ex09\Program.cs

1. Create a new console application called Ch13Ex09 and save it in the directory `C:\BegVCSharp\Chapter13`.
2. Modify the code in `Program.cs` as follows:

```
static void Main(string[] args)
{
    string[] curries = { "pathia", "jalfrezi", "korma" };
    WriteLine(curries.Aggregate(
        (a, b) => a + " " + b));
    WriteLine(curries.Aggregate<string, int>(
        0,
        (a, b) => a + b.Length));
    WriteLine(curries.Aggregate<string, string, string>(
        "Some curries:",
        (a, b) => a + " " + b,
        a => a));
    WriteLine(curries.Aggregate<string, string, int>(
        "Some curries:",
        (a, b) => a + " " + b,
        a => a.Length));
    ReadKey();
}
```

3. Run the application. The result is shown in Figure 13-21.




FIGURE 13-21

### How It Works

In this example you experimented with each of the overloads of the `Aggregate()` extension method, using a string array with three elements as source data.

First, a simple concatenation is performed:

```
WriteLine(curries.Aggregate((a, b) => a + " " + b));
```

The first pair of elements is concatenated into a string using simple syntax. This is far from the best way to concatenate strings — ideally you would use `string.Concat()` or `string.Format()` to optimize performance — but here it provides a very simple way to see what is going on. After this first concatenation, the result is passed back into the lambda expression along with the third element in the array, in much the same way as you saw `int` values being summed earlier. The result is a concatenation of the entire array, with spaces separating entries. You can achieve this effect in a simpler way using the `string.Join()` method, but the remainder of the overloads illustrated in this example provide additional functionality that `string.Join()` doesn't.

The second overload of the `Aggregate()` function, which has the two generic type parameters `TSource` and `TAccumulate`, is used. In this case the lambda expression must be of the form `Func<TAccumulate, TSource, TAccumulate>`. In addition, a seed value of type `TAccumulate` must be specified. This seed value is used in the first call to the lambda expression, along with the first array element. Subsequent calls take the accumulator result of previous calls to the expression. The code used is as follows:

```
WriteLine(curries.Aggregate<string, int>(
    0,
    (a, b) => a + b.Length));
```

The accumulator (and, by implication, the return value) is of type `int`. The accumulator value is initially set to the seed value of `0`, and with each call to the lambda expression it is summed with the length of an element in the array. The final result is the sum of the lengths of each element in the array.

Next you come to the last overload of `Aggregate()`. This takes three generic type parameters and differs from the previous version only in that the return value can be a different type from both the type of the elements in the array and the accumulator value. First, this overload is used to concatenate the string elements with a seed string:

```
WriteLine(curries.Aggregate<string, string, string>(
    "Some curries:",
    (a, b) => a + " " + b,
    a => a));
```

The final parameter of this method, `resultSelector`, must be specified even if (as in this example) the accumulator value is simply copied to the result. This parameter is a lambda expression of type `Func<TAccumulate, TResult>`.

In the final section of code, the same version of `Aggregate()` is used again, but this time with an `int` return value. Here, `resultSelector` is supplied with a lambda expression that returns the length of the accumulator string:

```
WriteLine(curries.Aggregate<string, string, int>(
    "Some curries:",
    (a, b) => a + " " + b,
    a => a.Length));
```

---

This example hasn't done anything spectacular, but it demonstrates how you can use more complicated extension methods that involve generic type parameters, collections, and seemingly complex syntax. You'll see more of this later in the book.

**EXERCISES**

- 13.1. Write the code for an event handler that uses the general-purpose `(object sender, EventArgs e)` syntax that will accept either the `Timer.Elapsed` event or the `Connection.MessageArrived` event from the code shown earlier in this chapter. The handler should output a string specifying which type of event has been received, along with the `Message` property of the `MessageArrivedEventArgs` parameter or the `SignalTime` property of the `ElapsedEventArgs` parameter, depending on which event occurs.
- 13.2. Modify the card game example to check for the more interesting winning condition of the popular card game, rummy. This means that a player wins the game if his or her hand contains two “sets” of cards, one of which consists of three cards and one of which consists of four cards. A set is defined as either a sequence of cards of the same suit (such as 3H, 4H, 5H, 6H) or several cards of the same rank (such as 2H, 2D, 2S).
- 13.3 Why can't you use an object initializer with the following class? After modifying this class to enable the use of an object initializer, give an example of the code you would use to instantiate and initialize this class in one step:

```
public class Giraffe
{
    public Giraffe(double neckLength, string name)
    {
        NeckLength = neckLength;
        Name = name;
    }
    public double NeckLength {get; set;}
    public string Name {get; set;}
}
```

- 13.4 True or false: If you declare a variable of type `var`, you will then be able to use it to hold any object type.
- 13.5 When you use anonymous types, how can you compare two instances to determine whether they contain the same data?
- 13.6 Try to correct the following code for an extension method, which contains an error:

```
public string ToAcronym(this string inputString)
{
    inputString = inputString.Trim();
    if (inputString == "")
    {
        return "";
    }
    string[] inputStringAsArray = inputString.Split(' ');
    StringBuilder sb = new StringBuilder();
    for (int i = 0; i < inputStringAsArray.Length; i++)
    {
        if (inputStringAsArray[i].Length > 0)
        {
            sb.AppendFormat("{0}",
                inputStringAsArray[i].Substring(
                    0, 1).ToUpper());
        }
    }
    return sb.ToString();
}
```

```

        }
    }
    return sb.ToString();
}

```

**13.7** How would you ensure that the extension method in Question 4 was available to your client code?

**13.8** Rewrite the `ToAcronym()` method shown here as a single statement. The code should ensure that strings including multiple spaces between words do not cause errors. Hint: You will require the `? :` tertiary operator, the `string.Aggregate<string, string>()` extension method, and a lambda expression to achieve this.

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Namespace qualification	To avoid ambiguity in namespace qualification, you can use the <code>::</code> operator to force the compiler to use aliases that you have created. You can also use the <code>global</code> namespace as an alias for the top-level namespace.
Custom exceptions	You can create your own exception classes by deriving from the <code>root Exception</code> class. This is helpful because it gives you greater control over catching specific exceptions, and allows you to customize the data that is contained in an exception in order to deal with it effectively.
Event handling	Many classes expose events that are raised when certain triggers occur in their code. You can write handlers for these events to execute code at the point where they are raised. This two-way communication is a great mechanism for responsive code, and prevents you from having to write what would otherwise be complex, convoluted code that might poll an object for changes.
Event definitions	You can define your own event types, which involves creating a named event and a delegate type for any handlers for the event. You can use the standard delegate type with no return type and custom event arguments that derive from <code>System.EventArgs</code> to allow for multipurpose event handlers. You can also use the <code>EventHandler</code> and <code>EventHandler&lt;T&gt;</code> delegate types to define events with simpler code.
Anonymous methods	Often, to make your code more readable, you can use an anonymous method instead of a full event handler method. This means defining the code to execute when an event is raised in-line at the point where you add the event handler. You achieve this with the <code>delegate</code> keyword.
Attributes	Occasionally, either because the framework you are using demands it or because you choose to, you will make use of attributes in your code. You can add attributes to classes, methods and other members using <code>[AttributeName]</code> syntax, and you can create your own attributes by deriving from <code>System.Attribute</code> . You can read attribute values through reflection.

Topic	Key Concepts
Initializers	<p>You can use initializers to initialize an object or collection at the same time as creating it. Both types of initializers consist of a block of code surrounded by curly brackets. Object initializers allow you to set property values by providing a comma-separated list of property name/value pairs. Collection initializers simply require a comma-separated list of values. When you use an object initializer, you can also use a nondefault constructor.</p>
Type inference	<p>The <code>var</code> keyword allows you to omit the type of a variable when you declare it. However, this is possible only if the type can be determined at compile time. Using <code>var</code> does not break the strong typing methodology of C# as a variable declared with <code>var</code> has one and only one possible type.</p>
Anonymous types	<p>For many simple types used to structure data storage, defining a type is not necessary. Instead, you can use an anonymous type, whose members are inferred from usage. You define an anonymous type with object initializer syntax, and every property you set is defined as a read-only property.</p>
Dynamic lookup	<p>Use the <code>dynamic</code> keyword to define a dynamic type variable that can hold any value. You can then access members of the contained value with normal property or method syntax, and these are only checked at runtime. If, at runtime, you attempt to access a nonexistent member, an exception is thrown. This dynamic typing greatly simplifies the syntax required to access non-.NET types, or .NET types whose type information is not available at compile time. However, dynamic types must be used with caution as you lose compile time code checking. You can control the behavior of dynamic lookup by implementing the <code>IDynamicMetaObjectProvider</code> interface.</p>
Optional method parameters	<p>Often, you can define a method with lots of parameters, many of which are only rarely used. Instead of forcing client code to specify values for rarely used parameters, you might provide multiple method overloads. Alternatively, you can define these parameters as optional (and provide default values for parameters that are not specified). Client code that calls your method can then specify only as many parameters as are required.</p>
Named method parameters	<p>Client code can specify method parameter values by position or by name (or a mix of the two where positional parameters are specified first). Named parameters can be specified in any order. This is particularly useful when combined with optional parameters.</p>
Lambda expressions	<p>Lambda expressions are essentially a shorthand way of defining anonymous methods, although they have additional capabilities such as implicit typing. You define a lambda expression with a comma-separated list of parameters (or empty parentheses for no parameters), the <code>=&gt;</code> operator, and an expression. The expression can be a block of code enclosed in curly brackets. Lambda expressions with up to eight parameters and an optional return type can be represented with the <code>Action</code>, <code>Action&lt;&gt;</code>, and <code>Func&lt;&gt;</code> delegate types. Many LINQ extension methods that can be used with collections use lambda expression parameters.</p>

## PART II

# Windows Programming

---

- CHAPTER 14: Basic Desktop Programming
- CHAPTER 15: Advanced Desktop Programming



# 14

## Basic Desktop Programming

### WHAT YOU WILL LEARN IN THIS CHAPTER

---

- Using the WPF designer
- Using controls for displaying information to the user, such as the `Label` and `TextBlock` controls
- Using controls for triggering events, such as the `Button` control
- Using the controls that enable users of your application to enter text, such as the `TextBox` control
- Using controls that enable you to inform users of the current state of the application and allow the user to change that state, such as the `RadioButton` and `CheckButton` controls
- Using controls that enable you to display lists of information, such as the `ListBox` and `ComboBox` controls
- Using panels to lay out your user interfaces

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 14 download and individually named according to the names throughout the chapter.

The first part of this book has concerned itself with the ins and outs of C#, but now it is time to move away from the details of the programming language and into the world of the graphical user interface (GUI).

Over the past 10 years, Visual Studio has provided the Windows developers with a couple of choices for creating user interfaces: Windows Forms, which is a basic tool for creating applications that target classic Windows, and Windows Presentation Foundations (WPF),

which provide a wider range of application types and attempts to solve a number of problems with Windows Forms. WPF is technically platform-independent, and some of its flexibility can be seen in the fact that a subset of WPF called Silverlight is used to create interactive web applications. In this and the next chapter you are going to learn how to use WPF to create Windows applications, and in Chapter 23 you will build on this knowledge when you create Universal Windows Apps.

At the heart of the development of most graphical Windows applications is the Window Designer. You create a user interface by dragging and dropping controls from a Toolbox to your window, placing them where you want them to appear when you run the application. With WPF this is only partly true, as the user interface is in fact written entirely in another language called Extensible Application Markup Language (XAML, pronounced *zammel*). Visual Studio allows you to do both and as you get more comfortable with WPF, you are likely going to combine dragging and dropping controls with writing raw XAML.

In this chapter, you work with the Visual Studio WPF designer to create a number of windows for the card game that you wrote in previous chapters. You learn to use some of the many controls that ship with Visual Studio that cover a wide range of functionality. Through the design capabilities of Visual Studio, developing user interfaces and handling user interaction is very straightforward — and fun! Presenting all of Visual Studio’s controls is impossible within the scope of this book, so this chapter looks at some of the most commonly used controls, ranging from labels and text boxes to menu bars and layout panels.

## XAML

XAML is a language that uses XML syntax and enables controls to be added to a user interface in a declarative, hierarchical way. That is to say, you can add controls in the form of XML elements, and specify control properties with XML attributes. You can also have controls that contain other controls, which is essential for both layout and functionality.

**NOTE** XML is covered in detail in Chapter 19. If you want a quick introduction to the basics of XML at this point, it might be a good idea to skip forward and read the first few pages of that chapter.

XAML is designed with today’s powerful graphics cards in mind, and as such it enables you to use all the advanced capabilities that these graphics cards offer through DirectX. The following lists some of these capabilities:

- Floating-point coordinates and vector graphics to provide layout that can be scaled, rotated, and otherwise transformed with no loss of quality
- 2D and 3D capabilities for advanced rendering
- Advanced font processing and rendering
- Solid, gradient, and texture fills with optional transparency for UI objects

- Animation storyboarding that can be used in all manner of situations, including user-triggered events such as mouse clicks on buttons
- Reusable resources that you can use to dynamically style controls

## Separation of Concerns

One problem that exists with maintaining Windows applications that has been written over the years is that they very often mix the code that generates the user interface and the code that executes based on users' actions. This makes it difficult for multiple developers and designers to work on the same project. WPF solves this in two ways. First, by using XAML to describe the GUI rather than C#, the GUI becomes platform independent, and you can in fact render XAML without any code whatsoever. Second, this means that it feels natural to place the C# code in a different file than you place the GUI code. Visual Studio utilizes something called *code-behind* files, which are C# files that are dynamically linked to the XAML files.

Because the GUI is separated from the code, it is possible to create tailor-made applications for designing the GUI, and this is exactly what Microsoft has done. The design tool Blend for Visual Studio is the favored tool used by designers when creating GUIs for WPF. This tool can load the same projects as Visual Studio, but where Visual Studio targets the developer more than the designer, the opposite is true in Expression Blend. This means that on large projects with designers and developers, everyone can work together on the same project, using their preferred tool without fear of inadvertently influencing the others.

## XAML in Action

As stated, XAML is XML, which means that as long as the files are fairly small, it is possible to see immediately what it is describing. Take a look at this small example and see if you can tell what it does:

```
<Window x:Class="Ch14Ex01.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:WpfApplication1"
    mc:Ignorable="d"
    Title="Hello World" Height="350" Width="525">
    <Grid>
        <Button Content="Hello World"
            HorizontalAlignment="Left"
            Margin="220,151,0,0"
            VerticalAlignment="Top"
            Width="75"/>
    </Grid>
</Window>
```

The XAML in this example creates a window with a single button on it. Both the window and the button display the text "Hello World". XML allows you to place tags inside other tags as long as

you close them properly. When an element is placed inside another in XAML, this element becomes the content of the enclosing element, meaning that the `Button` could also have been written like this:

```
<Button HorizontalAlignment="Left"
        Margin="220,151,0,0"
        VerticalAlignment="Top"
        Width="75">
    Hello World
</Button>
```

Here, the `Content` property of the `Button` has been removed and the text is now a child node of the `Button` control. Content can be just about anything in XAML, which is also demonstrated in this example: The `Button` element is the content of the `Grid` element, which is itself the content of the `Window` element.

Most, if not all, controls can have content, and there are very few limits to what you can do to change the appearance of the built-in controls. Chapter 15 explores this in more detail.

## Namespaces

The `Window` element of the previous example is the root element of the XAML file. This element usually includes a number of namespace declarations. By default, the Visual Studio designer includes two namespaces that you should be aware of: `http://schemas.microsoft.com/winfx/2006/xaml/presentation` and `http://schemas.microsoft.com/winfx/2006/xaml`. The first one is the default namespace of WPF and declares a lot of controls that you are going to use to create user interfaces. The second one declares the XAML language itself. Namespaces don't have to be declared on the root tag, but doing so ensures that their content can be easily accessed throughout the XAML file, so there is rarely any need to move the declarations.

**NOTE** The namespaces looks like they might be URLs, but this is deceiving. In fact they are what is known as Uniform Resource Identifiers (URIs). A URI can be any string as long as it uniquely identifies a resource. Microsoft has chosen to specify the URIs in a form that is normally used for URLs, but in this case you will not get a result if you type them into your browser.

When you create a new window in Visual Studio, the presentation namespace is always declared as the default and the language namespace as `xmlns:x`. As seen with the `Window`, `Button`, and `Grid` tags, this ensures that you don't have to prefix the controls you add to the window, but the language elements you specify must be prefixed with an `x`.

The last namespace that you will see quite often is the system namespace: `xmlns:sys="clr-namespace:System;assembly=mscorlib"`. This namespace allows you to use the built-in types of the .NET Framework in your XAML. By doing this, the markup you write can explicitly declare the types of elements you are creating. For example, it is possible to declare an array in markup and state that the members of the array are strings:

```
<Window.Resources>
    <ResourceDictionary>
```

---

```

<x:Array Type="sys:String" x:Key="localArray">
    <sys:String>Benjamin Perkins</sys:String>
    <sys:String>Jacob Vibe Hammer</sys:String>
    <sys:String>Job D. Reid</sys:String>
</x:Array>
</ResourceDictionary>
</Window.Resources>

```

## Code-Behind Files

Although XAML is a powerful way to declare user interfaces, it is not a programming language. Whenever you want to do more than presentation, you need C#. It is possible to embed C# code directly into XAML, but mixing code and markup is never recommended and you will not see it done in this book. What you will see quite a lot is the use of code-behind files. These files are normal C# files that have the same name as the XAML file, plus a .cs extension. Although you can call them whatever you like, it's best to stick to the naming convention. Visual Studio creates code-behind files automatically when you create a new window in your application, because it expects you to add code to the window. It also adds the `x:Class` property to the `Window` tag in the XAML:

```
<Window x:Class="Ch14Ex01.MainWindow"
```

This tells the compiler that it can find the code for this window in, not a file, but the class `Ch14Ex01.MainWindow`. Because you can specify only the fully qualified class name, and not the assembly in which the class is found, it is not possible to put the code-behind file somewhere outside of the project in which the XAML is defined. Visual Studio puts the code-behind files in the same directory as the XAML files so you never have to worry about this while working in Visual Studio.

## THE PLAYGROUND

Now you know enough about how WPF is constructed to start getting your hands dirty, so it's time to look at the editor. Start by creating a new WPF project by selecting `File`  $\Rightarrow$  `New`  $\Rightarrow$  `Project`. From the New Project dialog box, navigate to the Clasic Desktop node under `Visual C#  $\Rightarrow$  Windows` and select the project template `WPF Application`. To be able to reuse this example with the next examples, name the project `Ch14Ex01`.

Visual Studio now displays an empty window and a number of panels around it. The greater part of the screen is divided in two sections. The upper section, known as the Design View, displays a WYSIWYG (What You See Is What You Get) representation of the window you are designing and the lower section, known as the XAML View, displays a textual representation of the same window.

To the right of the Design View, you see the Solution Explorer that you have seen in previous projects and a Properties panel that displays information about the current selection in the Design and XAML Views. It is worth noting that the selection in the Properties panel, XAML View, and Design View are always in sync, so if you move the cursor in the XAML View you will see the selection change in the other two.

Collapsed to the left of the Design View are a number of panels, one of which is the Toolbox. This chapter shows you how to use many of the controls from the Toolbox panel to create dialog boxes for the card game, so expand it and pin it open by clicking the pin in the top-right corner. While you

are at it, expand the Common WPF Controls node in the panel as well. You will be using most of the controls shown here in this chapter.

## WPF Controls

Controls combine prepackaged code and a GUI that can be reused to create more complex applications. They can define how they draw themselves by default and a set of standard behaviors. Some controls, such as the `Label`, `Button`, and `TextBox` controls are easily recognizable and have been used in Windows applications for about 20 years. Others, such as `Canvas` and `StackPanel`, don't display anything and simply help you create the GUI.

Out-of-the-box controls look exactly as you would expect a control to look in a standard Windows application and use the current Windows Theme to draw themselves. All of this is highly customizable and with only a few clicks you can completely change how a control is displayed. This customization is done using properties that are defined on the controls. WPF uses normal properties that you have seen before and adds a new type of property called a *dependency property*. These are examined in detail in Chapter 15, but for now it is enough to know that many of the properties of WPF do more than just get and set a value; for one, they are able to notify observers of changes.

Besides defining how something looks on the screen, controls also define standard behavior, such as the ability to click on a button and select something in a list. You can change what happens when a user performs an action on a control by “handling” the events that the control defines. When and how you implement the event handler will vary from application to application and from control to control, but generally speaking you will always handle the `Click` event for a button; for a `ListBox` control, you often have to react when the user changes the selection and so the `SelectionChanged` event should be handled. On other controls, such as the `Label` or `TextBlock` controls, you will rarely implement any event.

**WARNING** Although users are often happy when you take the time to provide a more interesting user interface than the standard Windows display, you must be careful when changing the standard behavior of controls. Imagine that you change a `Button` control to work only when users right-click it. Your users will think that your application is broken when nothing happens when they left-click on the button. In fact, even if there are fantastic reasons for changing the button like this, it is likely that you should be using another type of control instead of changing the behavior of the `Button` control.

You can add controls to a window in a number of ways, but the most common way is to drag and drop them from the Toolbox onto the Design View or the XAML View.

## TRY IT OUT Adding Controls to a Window

As you work your way through this chapter, you will add controls to the Design View by dragging them from the Toolbox panel or by typing the XAML manually.

1. Start by dragging a Button control from the Toolbox onto the Design View. Notice how the text in the XAML View is updated to reflect the change you made.
2. Now drag another Button, but this time drop it in the XAML View below the first Button, but above the `</Grid>` tag.

### How It Works

The result you see in the Design View might be somewhat surprising — the second button expands to fill the entire window. When you drop a control onto the Design View, Visual Studio will try to set properties and insert child elements to allow the controls to display themselves in a standard way. This does not happen when you drag controls into the XAML View, where only the tag that is used to define the control is inserted.

There are times when you want to position a control at a specific position on your window and it is difficult to drop it at exactly the right position. When this happens, you might want to drop the control directly in the XAML View or type it manually.

**NOTE** If you want the behavior of the Design View when you drop a control, but can't hit the right spot, just drop it anywhere and then cut and paste the XAML that was generated for you into the correct position.

## Properties

As mentioned, all controls have a number of properties that are used to manipulate the behavior of the control. Some of these are easy to understand such as `height` and `width`, whereas others are less obvious such as `RenderTransform`. All of them can be set using the Properties panel, directly in XAML, or by manipulating the control on the Design View.

**NOTE** Visual Studio will create a default namespace for your classes when you create a new project. That namespace is subsequently used when you add new classes or windows to your project. You can change the namespace by double-clicking Properties in the Solution Explorer. If you find that your classes get a different namespace than given in the examples, it can be helpful to change the default namespace to the namespace from the book. The change will only affect new classes, not anything already in the project.

## TRY IT OUT Manipulating Properties: Ch14Ex01\MainWindow.xaml

Return to the previous example and follow these steps. As you change the properties, notice how your changes affect the XAML and Design Views. You are going to change the window to look like Figure 14-1.

1. Start by selecting the second `Button` control in Design View; this is the button that is currently filling the entire window.
2. You can change the name of the control in the Properties panel at the very top. Change it to `rotatedButton`.
3. Under the Common node, change the Content to `2nd Button`.
4. Under Layout, change width to 75 and height to 22.
5. Expand the Text node and change the text to bold by clicking the B icon.
6. Select the first button and drag it to a position above the second button. Visual Studio will assist with the positioning by snapping the control.




FIGURE 14-1

7. Select the second button again, and hover the mouse pointer over the top-left corner of it. The pointer changes to a quarter-circle with arrows on both ends. Drag down until the button is tilted down.
8. The XAML code for the window should now look like this:

```
<Window x:Class="Ch14Ex01.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:Ch14Ex01"
    mc:Ignorable="d"
    Title="MainWindow" Height="350" Width="525">
<Grid>
    <Button x:Name="button" Content="Button" HorizontalAlignment="Left"
        Margin="218,113,0,0" VerticalAlignment="Top" Width="75"/>
```

```

<Button x:Name="rotatedButton" Content="2nd Button" Width="75"
    Height="22" FontWeight="Bold" Margin="218,138,224,159"
    RenderTransformOrigin="0.5,0.5" >
    <Button.RenderTransform>
        <TransformGroup>
            <ScaleTransform/>
            <SkewTransform/>
            <RotateTransform Angle="-23.896"/>
            <TranslateTransform/>
        </TransformGroup>
    </Button.RenderTransform>
</Button>
</Grid>
</Window>

```

- Run the application by pressing F5. Try to resize the window. Notice that the second button moves with the window, whereas the first button stays fixed.

### How It Works

Any change that you apply in any of the three views is reflected in the other views, but some things are easier to do in certain views. Changing something trivial like the text displayed on a button can be done quickly in XAML View, but adding the information needed to perform a render transformation is much quicker from Design View.

In this exercise, you began by changing the name of the button, which added the `x:Name` property to the button. The name of a control must be unique within the scope of the namespace, so you can use the name for only one control.

Next you changed the `Content` property, set the `Height` and `Width` of the control, and then changed the font to bold. Doing so changed the way the control displayed itself within the window. It used to fill all the space of its container, but now you have limited it to a specific size.

Then you dragged the first button to a specific position on the Design View. As you see later in this chapter, this action will not always yield the same results but is dependent on the container in which the control is placed. In this case, with the `Grid` container, the control can be dragged to a specific position. The action sets the `Margin` property on the control. Two other properties should be mentioned here: `HorizontalAlignment="Left"` and `VerticalAlignment="Top"`. With these two properties set, the margin becomes relative to the top-left corner of the window and thus the control is pushed to the position you placed it in the grid. If you compare the first and second buttons at this point, you will notice that the second control has none of these properties set. By omitting the alignment properties as well as the margin properties, the control is placed at the center of the container, even at runtime. This means that the first button with the margin and alignments set is fixed when the window resizes, but the second button always stays centered.

Finally, you performed a little bit of a party trick. By dragging the control when the `Rotate` mouse pointer is displayed, you can rotate the control. This is a standard feature of XAML and WPF and can be applied to all controls, although there are a few controls that fail to change their content when the control itself is rotated. This includes controls that rely on Windows Forms or old Windows controls to display content.

The animations that you can do in WPF are covered in Chapter 15, but from the XML that was generated when you dragged the cursor, you can see that you can perform some advanced animation simply by manipulating these properties.

---

## Dependency Properties

For the most part, normal .NET properties are simple getters and setters, which is fine for most cases. However, when you are working with a dynamic user interface that can and should change when properties change, you have to write a lot of code in these get and set methods that will be repeated many times. A *dependency property* is a property that is registered with the WPF property system in such a way as to allow extended functionality. This extended functionality includes, but is not limited to, automatic property change notifications. Specifically, dependency properties have the following features:

- You can use styles to change the values of dependency properties.
- You can set the value of a dependency property by using resources or by data binding.
- You can change dependency property values in an animation.
- You can set dependency properties hierarchically in XAML — that is, a value for a dependency property that you set on a parent element can be used to set the default value for the same dependency property of its child elements.
- You can configure notifications for property value changes using a well-defined coding pattern.
- You can configure sets of related properties so that they all update in response to a change to one of them. This is known as *coercion*. The changed property is said to *coerce* the values of the other properties.
- You can apply metadata to a dependency property to specify other behavior characteristics. For example, you might specify that if a given property changes, then it might be necessary to rearrange the user interface.

In practice, because of the way in which dependency properties are implemented, you might not notice much of a difference compared to ordinary properties. However, when you create your own controls, you will quickly find that a lot of functionality suddenly disappears when you use ordinary .NET properties.

Chapter 15 shows how you can implement new dependency properties.

## Attached Properties

An *attached property* is a property that is made available to each child object of an instance of the class that defines the property. For example, as you will see later in this chapter, the `Grid` control that you used in the previous examples allows you to define columns and rows for ordering the child controls of the `Grid`. Each child control can then use the attached properties `Column` and `Row` to specify where it belongs in the grid:

```
<Grid HorizontalAlignment="Left" Height="167" VerticalAlignment="Top" Width="290">
<Button Content="Button" HorizontalAlignment="Left" Margin="10,10,0,0">
```

```

VerticalAlignment="Top" Width="75" Grid.Column="0" Grid.Row="0"
Height="22" />
...
</Grid>
```

Here, the attached property is referred to using the name of the parent element, a period, and the name of the attached property.

In WPF, attached properties serve a variety of uses. You will see a lot of attached properties shortly, when you look at how to position controls in the “Control Layout” section. You will learn how container controls define attached properties that enable child controls to define, for example, which edges of the container to dock to.

## Events

In Chapter 13, you learned what events are and how to use them. This section covers particular kinds of events — specifically, the events generated by WPF controls — and introduces *routed events*, which are usually associated with user actions. For example, when the user clicks a button, that button generates an event indicating what just happened to it. Handling the event is the means by which the programmer can provide some functionality for that button.

Many of the events you handle are common to most of the controls that you work with in this book. This includes events such as `LostFocus` and `MouseEnter`. This is because the events themselves are inherited from base classes such as `Control` or `ContentControl`. Other events such as the `CalendarOpened` event of the `DatePicker` are more specific and only found on specialized controls. Some of the most used events are listed in Table 14-1.

TABLE 14-1: Common Control Events

EVENT	DESCRIPTION
Click	Occurs when a control is clicked. In some cases, this event also occurs when a user presses the Enter key.
Drop	Occurs when a drag-and-drop operation is completed — in other words, when an object has been dragged over the control, and the user releases the mouse button.
DragEnter	Occurs when an object being dragged enters the bounds of the control.
DragLeave	Occurs when an object being dragged leaves the bounds of the control.
DragOver	Occurs when an object has been dragged over the control.
KeyDown	Occurs when a key is pressed while the control has focus. This event always occurs before <code>KeyPress</code> and <code>KeyUp</code> .
KeyUp	Occurs when a key is released while a control has focus. This event always occurs after <code>KeyDown</code> event.

*continues*


**TABLE 14-1 (continued)**

EVENT	DESCRIPTION
GotFocus	Occurs when a control receives focus. Do not use this event to perform validation of controls. Use <code>Validating</code> and <code>Validated</code> instead.
LostFocus	Occurs when a control loses focus. Do not use this event to perform validation of controls. Use <code>Validating</code> and <code>Validated</code> instead.
MouseDoubleClick	Occurs when a control is double-clicked.
MouseDown	Occurs when the mouse pointer is over a control and a mouse button is pressed. This is not the same as a <code>Click</code> event because <code>MouseDown</code> occurs as soon as the button is pressed and before it is released.
MouseMove	Occurs continually as the mouse travels over the control.
MouseUp	Occurs when the mouse pointer is over a control and a mouse button is released.

You will see many of these events in the examples in this chapter.

## Handling Events

There are two basic ways to add a handler for an event. One way is to use the Events list in the Properties window, shown in Figure 14-2, which is displayed when you click the lightning bolt button.

**FIGURE 14-2**

To add a handler for a particular event, either type the name of the event and press Return, or double-click to the right of the event name in the Events list. This causes the event to be added to the XAML tag. The method signature to handle the event is added to the C# code-behind file.

```

<Button x:Name="rotatedButton" Content="2nd Button" Width="75"
    Height="22" FontWeight="Bold" Margin="218,138,224,159"
    RenderTransformOrigin="0.5,0.5"
    Click="rotatedButton_Click">

    ...
</Button>
private void rotatedButton_Click(object sender, RoutedEventArgs e)
{
}

```

You can also type the name of the event directly in XAML and add the name of the handler there. If you do this, you can right-click on the event and chose Navigate to Event Handler. This will add the event handler to the code-behind file.

## Routed Events

WPF uses events that are called *routed events*. A standard .NET event is handled by the code that has explicitly subscribed to it and it is sent only to those subscribers. Routed events are different in that they can send the event to all controls in the hierarchy in which the control participates.

A routed event can travel up and down the hierarchy of the control on which the event occurred. So, if you right-click a button, the `MouseRightButtonDown` event will first be sent to the button itself, then to the parent of the control — in the case of the earlier example, the `Grid` control. If this doesn't handle it, then the event is finally sent to the window. If, on the other hand you don't want the event to travel further up the hierarchy, then you simply set the `RoutedEventArgs.Handled` property to `true`, and no additional calls will be made at that point. When an event travels up the control hierarchy like this, it is called a *bubbling event*.

Routed events can also travel in the other direction, that is, from the root element to the control on which the action was performed. This is called a *tunneling event* and by convention all events like this are prefixed with the word `Preview` and always occur before their bubbling counterparts. An example of this is the `PreviewMouseRightButtonDown` event.

Finally, a routed event can behave exactly like a normal .NET event and only be sent to the control on which the action was made.

## Routed Commands

Routed commands serve much the same purpose as events in that they cause some code to execute. Where Events are bound directly to a single element in the XAML and a handler in the code, Routed Commands are more sophisticated.

The key difference between events and commands is in their use. An event should be used whenever you have a piece of code that has to respond to a user action that happens in only one place in your application. An example of such an event could be when the user clicks OK in a window to save and close it. A command can be used when you have code that will be executed to respond to actions that happen in many locations. An example of this is when the content of an application is saved. There is often a menu with a Save command that can be selected, as well as a toolbar button for the same purpose. It is possible to use event handlers to do this, but it would mean implementing the same code in many locations — a command allows you to write the code just once.

When you create a command, you must also implement code that can respond to the question, “Should this code be available to the user at the moment?” This means that when a command is associated with a button, that button can ask the command if it can execute and set its state accordingly.

A command is much more complicated to implement than an event, so you are not going to see them in use until Chapter 15, where they will be used with menu items.

### TRY IT OUT Routed Events: Ch14Ex01\MainWindow.xaml

This example builds on the example from earlier in the chapter. If you added the rows and columns earlier, you should remove them to match the XAML in this example.

1. Select the button `rotatedButton` and add the event `KeyDown`. You can do this through the Properties panel or by typing the XAML directly. Name it `rotatedButton_KeyDown`.
2. Select the `Grid` by clicking on the tag it in the XAML View, and add the same event to it. Name it `Grid_KeyDown`.
3. Select the `Window` tag in the XAML View and add the event again. Name it `Window_KeyDown`.
4. Repeat Steps 1 through 3, but replace the event with `PreviewKeyDown` and change the name of the event to reflect that it is the Preview handler. The XAML should look like this:

```
<Window x:Class="Ch14Ex01.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:Ch14Ex01"
    mc:Ignorable="d"
    Title="MainWindow" Height="350" Width="525" KeyDown="Window_KeyDown"
    PreviewKeyDown="Window_PreviewKeyDown">
    <Grid KeyDown="Grid_KeyDown" PreviewKeyDown="Grid_PreviewKeyDown">      <Button
        x:Name="button" Content="Button" HorizontalAlignment="Left"
        Margin="27,4,0,0" VerticalAlignment="Top" Width="75" Grid.Column="0"
        Grid.Row="0"/>
        <Button x:Name="rotatedButton" Content="2nd Button" Width="75" Height="22"
            FontWeight="Bold" RenderTransformOrigin="0.5, 0.5"
            KeyDown="rotatedButton_KeyDown"
            PreviewKeyDown="rotatedButton_PreviewKeyDown" Grid.Column="1"
            Grid.Row="1" >
            <Button.RenderTransform>
                <TransformGroup>
                    <ScaleTransform/>
                    <SkewTransform/>
                    <RotateTransform Angle="-23.896"/>
                    <TranslateTransform/>
                </TransformGroup>
            </Button.RenderTransform>
        </Button>
    </Grid>
</Window>
```

5. If you typed the XAML directly, right-click each of the events and add the event handler to the code-behind by selecting the Navigate to Event Handler menu item.
6. Add this code to the event handlers:

```

private void Grid_KeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("Grid handler, bubbling up");
}
private void Grid_PreviewKeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("Grid handler, tunneling down");
}
private void rotatedButton_KeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("rotatedButton handler, bubbling up");
}
private void rotatedButton_PreviewKeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("rotatedButton handler, tunneling down");
}
private void Window_KeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("Window handler, bubbling up");
}
private void Window_PreviewKeyDown(object sender, KeyEventArgs e)
{
    MessageBox.Show("Window handler, tunneling down");
}

```

7. Run the application by pressing F5.
8. Select the rotated button by clicking it and pressing any key except Return, Tab, Escape, the space-bar, or the arrow keys. Observe the events being executed in turn.
9. Stop the application.
10. Go to the `Grid_PreviewKeyDown` event handler and add this line below the `MessageBox` line:  
`e.Handled = true;`
11. Repeat Steps 7 and 8.

### *How It Works*

The `KeyDown` and `PreviewKeyDown` events demonstrate bubbling and tunneling events. When you press a key with `rotatedButton` selected, you see each of the event handlers executing, one after another.

First the `Preview` events execute, starting with the handler on `Window`, then the `Grid`, and finally the `rotatedButton`. Then the `KeyDown` events execute, but in the opposite order, starting with the event handler on the `rotatedButton` and finishing with the handler on `Window`.

If you use any of the keys explicitly stated not to use in Step 8, you will notice that the `Preview` event is fired only on `Window`. This is because these are not considered input keys and are ignored by the grid and buttons.

Then you added this line:

```
e.Handled = true;
```

This changed the behavior dramatically. By setting the `Handled` property of the `RoutedEventArgs` you not only caused the execution of the tunneling events, but also of the bubbling events. This is generally true for all events like this. If you stop the execution of either the `Preview` or the “normal” version of the event handlers, you stop them both.

---

## Control Types

As stated, WPF has a lot of controls to choose from. Two types of interest are the Content and Items controls. Content controls, such as the `Button` control, have a `Content` property that can be set to any other control. This means that you can determine how the control is displayed, but you can specify only a single control directly in the content. That being said, you can specify an Items control, which is a control that allows you to insert multiple controls as content. An example of an Items control is the `Grid` control. When you are creating user interfaces, you are continually combining these two control types.

In addition to Content and Items controls, there are a number of other types of controls that don’t allow you to use other controls as their content. One example of this is the `Image` control, which is used to display an image. Changing that behavior defeats the purpose of the control.

## CONTROL LAYOUT

So far in this chapter you have used the `Grid` element to lay out a few controls, primarily because that is the control supplied by default when you create a new WPF application. However, you haven’t yet examined the full capabilities of this class, nor have you learned about the other layout containers that you can use to achieve alternative layouts. This section looks at control layout in more detail, as it is a fundamental concept of WPF.

All content layout controls derive from the abstract `Panel` class. This class simply defines a container that can contain a collection of objects that derive from `UIElement`. All WPF controls derive from `UIElement`. You cannot use the `Panel` class directly for control layout, but you can derive from it if you want to. Alternatively, you can use one of the following layout controls that derive from `Panel`:

- `Canvas` — This control enables you to position child controls any way you see fit. It doesn’t place any restrictions on child control positioning, but nor does it provide any assistance in positioning.
- `DockPanel` — This control enables you to dock child controls against one of its four edges. The last child control fills the remaining space.
- `Grid` — This control enables flexible positioning of child controls. You can divide the layout of this control into rows and columns, which enables you to align controls in a grid layout.

- `StackPanel` — This control positions its child controls in a sequential horizontal or vertical layout.
- `WrapPanel` — This control positions its child controls in a sequential horizontal or vertical layout as `StackPanel`, but rather than a single row or column of controls, this control wraps its children into multiple rows or columns according to the space available.

You'll look at how to use these controls in more detail shortly. First, however, there are a few basic concepts to understand:

- How controls appear in stack order
- How to use alignment, margins, and padding to position controls and their content
- How to use the `Border` control

## Stack Order

When a container control contains multiple child controls, they are drawn in a specific stack order. You might be familiar with this concept from drawing packages. The best way to think of stack order is to imagine that each control is contained in a plate of glass, and the container contains a stack of these plates of glass. The appearance of the container, therefore, is what you would see if you looked down from the top through these layers of glass. The controls contained by the container overlap, so what you see is determined by the order of the glass plates. If a control is higher up the stack, then it will be the control that you see in the overlap area. Controls lower down may be partially or completely hidden by controls above them.

This also affects hit testing when you click on a window with the mouse. The target control will always be the one that is uppermost in the stack when considering overlapping controls. The stack order of controls is determined by the order in which they appear in the list of children for a container. The first child in a container is placed on the lowest layer in the stack, and the last child on the topmost layer. The children between the first and last child are placed on increasingly higher layers. The stack order of controls has additional implications for some of the layout controls that you can use in WPF, as you will see shortly.

## Alignment, Margins, Padding, and Dimensions

Earlier examples used the `Margin`, `HorizontalAlignment`, and `VerticalAlignment` properties to position controls in a `Grid` container, but without going into much detail about their use. You have also seen how you can use `Height` and `Width` to specify dimensions. These properties, along with `Padding`, which you haven't looked at yet, are useful for all of the layout controls (or most of them, as you will see), but in different ways. Different layout controls can also set default values for these properties. You'll see a lot of this by example in subsequent sections, but before doing that, it is worth covering the basics.

The two alignment properties, `HorizontalAlignment` and `VerticalAlignment`, determine how the control is aligned. `HorizontalAlignment` can be set to `Left`, `Right`, `Center`, or `Stretch`. `Left` and `Right` tend to position controls to the left or right edges of the container, `Center` positions controls

in the middle, and Stretch changes the width of the control so that its edges reach to the sides of the container. VerticalAlignment is similar, and has the values Top, Bottom, Center, or Stretch.

Margin and Padding specify the space to leave blank around the edges of controls and inside the edges of controls, respectively. Earlier examples used Margin to position controls relative to the edges of a window. This worked because with HorizontalAlignment set to Left and VerticalAlignment set to Top, the control is positioned tight against the top-left corner, and Margin inserted a gap around the edge of the control. Padding is used similarly, but spaces out the content of a control from its edges. This is particularly useful for Border, as you will see in the next section. Both Padding and Margin can be specified in four parts (in the form leftAmount, topAmount, rightAmount, bottomAmount) or as a single value (a Thickness value).

Later, you will see how Height and Width are often controlled by other properties. For example, with HorizontalAlignment set to Stretch, the Width property of a control changes as the width of its container changes.

## Border

The Border control is a very simple, and very useful, container control. It holds a single child, not multiple children like the more complicated controls you'll look at in a moment. This child will be sized to completely fill the Border control. This might not seem particularly useful, but remember that you can use the Margin and Padding properties to position the Border within its container, and the content of the Border within the edges of the Border. You can also set, for example, the Background property of a Border so that it is visible. You will see this control in action shortly.

## Canvas

The Canvas control, as previously noted, provides complete freedom over control positioning. Another thing about Canvas is that the HorizontalAlignment and VerticalAlignment properties used with a child element will have no effect whatsoever over the positioning of those elements.

You can use Margin to position elements in a Canvas as it was done in earlier examples, but a better way is to use the Canvas.Left, Canvas.Top, Canvas.Right, and Canvas.Bottom attached properties that the Canvas class exposes:

```
<Canvas...>
  <Button Canvas.Top="10" Canvas.Left="10"...>Button1</Button>
</Canvas>
```

The preceding code positions a Button so that its top edge is 10 pixels from the top edge of the Canvas, and its left edge is 10 pixels from the left edge of the Canvas. Note that the Top and Left properties take precedence over Bottom and Right. For example, if you specify both Top and Bottom, then the Bottom property is ignored.

Figure 14-3 shows two Rectangle controls positioned in a Canvas control, with the window resized to two sizes.

**NOTE** All of the example layouts in this section can be found in the `LayoutExamples` project in the downloadable code for this chapter. See the "Wrox.com Code Downloads for this Chapter" section at the beginning of this chapter for information on how to download this chapter's code.

One `Rectangle` is positioned relative to the top-left corner, and one is positioned relative to the bottom-right corner. As you resize the window, these relative positions are maintained. You can also see the importance of the stacking order of the `Rectangle` controls. The bottom-right `Rectangle` is higher up in the stacking order, so when they overlap this is the control that you see.




FIGURE 14-3

The code for this example is as follows (you can find it in the downloaded code at `LayoutExamples\Canvas.xaml`):

```
<Window x:Class="LayoutExamples.Canvas"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:LayoutExamples"
    mc:Ignorable="d"
    Title="Canvas" Height="300" Width="300">
    <Canvas Background="AliceBlue">
        <Rectangle Canvas.Left="50" Canvas.Top="50" Height="40" Width="100"
        Stroke="Black" Fill="Chocolate" />
        <Rectangle Canvas.Right="50" Canvas.Bottom="50" Height="40" Width="100"
        Stroke="Black" Fill="Bisque" />
    </Canvas>
</Window>
```

## DockPanel

The `DockPanel` control, as its name suggests, enables you to dock controls to one of its edges. This sort of layout should be familiar to you, even if you've never stopped to notice it before. It is how, for example, the `Ribbon` control in Word remains at the top of the Word window, or how the various windows in Visual Studio are positioned. In Visual Studio you can also change the docking of windows by dragging them around.

`DockPanel` has a single attached property that child controls can use to specify the edge to which controls dock: `DockPanel.Dock`. You can set this property to `Left`, `Top`, `Right`, or `Bottom`.

The stack order of controls in a `DockPanel` is extremely important, as every time you dock a control to an edge you also reduce the available space of subsequent child controls. For example, you might dock a toolbar to the top of a `DockPanel` and then a second toolbar to the left of the `DockPanel`.

The first control would stretch across the entire top of the `DockPanel` display area, but the second control would only stretch from the bottom of the first toolbar to the bottom of the `DockPanel` along the left edge.

The last child control you specify will (usually) fill the area that remains after all the previous children have been positioned. (You can control this behavior, which is why this statement is qualified.)

When you position a control in a `DockPanel`, the area occupied by the control might be smaller than the area of the `DockPanel` that is reserved for the control. For example, if you dock a `Button` with a `Width` of 100, a `Height` of 50, and a `HorizontalAlignment` of `Left` to the top of a `DockPanel`, then there will be space to the right of the `Button` that isn't used by other docked children. In addition, if the `Button` control has a `Margin` of 20, then a total of 90 pixels at the top of the `DockPanel` will be reserved (the height of the control plus the top and bottom margins). You need to take this behavior into account when you use `DockPanel` for layout; otherwise, you can end up with unexpected results.

Figure 14-4 shows a sample `DockPanel` layout.




FIGURE 14-4

The code for this layout is as follows (you can find it in the downloadable code at `LayoutExamples\DockPanels.xaml`):

```

<Window x:Class="LayoutExamples.DockPanels"
        xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
        xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
        xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
        xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
        xmlns:local="clr-namespace:LayoutExamples"
        mc:Ignorable="d"
        Title="DockPanels" Height="300" Width="300">
    <DockPanel Background="AliceBlue">
        <Border DockPanel.Dock="Top" Padding="10" Margin="5"
        Background="Aquamarine" Height="45">
            <Label>1) DockPanel.Dock="Top"</Label>
        </Border>
        <Border DockPanel.Dock="Top" Padding="10" Margin="5"
        Background="PaleVioletRed" Height="45" Width="200">
            <Label>2) DockPanel.Dock="Top"</Label>
        </Border>
        <Border DockPanel.Dock="Left" Padding="10" Margin="5"
        Background="Bisque" Width="200">
            <Label>3) DockPanel.Dock="Left"</Label>
        </Border>
        <Border DockPanel.Dock="Bottom" Padding="10" Margin="5"
        Background="Ivory" Width="200" HorizontalAlignment="Right">
            <Label>4) DockPanel.Dock="Bottom"</Label>
        </Border>
        <Border Padding="10" Margin="5" Background="BlueViolet">
            <Label Foreground="White">5) Last control</Label>
        </Border>
    </DockPanel>
</Window>

```

This code uses the `Border` control introduced earlier to clearly mark out the docked control regions in the example layout, along with `Label` controls to output simple informative text. To understand the layout, you must read it from top to bottom, looking at each control in turn:

1. The first `Border` control is docked to the top of the `DockPanel`. The total area taken up in the `DockPanel` is the top 55 pixels (`Height + 2 × Margin`). Note that the `Padding` property does not affect this layout, as it is inside the edge of the `Border`, but this property does control the positioning of the embedded `Label` control. The `Border` control fills any available space along the edge it is docked to if not constrained by `Height` or `Width` properties, which is why it stretches across the `DockPanel`.
2. The second `Border` control is also docked to the top of the `DockPanel`, and takes up another 55 pixels from the top of the display area. This `Border` control also includes a `Width` property, which causes the border to take up only a portion of the width of the `DockPanel`. It is positioned centrally, as the default value for `HorizontalAlignment` in a `DockPanel` is `Center`.

3. The third `Border` control is docked to the left of the `DockPanel` and takes up 210 pixels of the left of the display.
4. The fourth `Border` control is docked to the bottom of the `DockPanel` and takes up 30 pixels plus the height of the `Label` control it contains (whatever that is). This height is determined by the `Margin`, `Padding`, and contents of the `Border` control, as it is not specified explicitly. The `Border` control is locked to the bottom-right corner of the `DockPanel`, as it has a `HorizontalAlignment` of `Right`.
5. The fifth and final `Border` control fills the remaining space.

Run this example and experiment with resizing content. Note that the further up the stacking order a control is, the more priority is given to its space. By shrinking the window, the fifth `Border` control can be completely obscured by controls further up the stacking order. Be careful when using `DockPanel` control layout to avoid this, perhaps by setting minimum dimensions for the window.

## StackPanel

You can think of `StackPanel` as being a slimmed down version of `DockPanel`, where the edge to which child controls are docked is fixed for those controls. The other difference between these controls is that the last child control of a `StackPanel` doesn't fill the remaining space. However, controls will, by default, stretch to the edges of the `StackPanel` control.

The direction in which controls are stacked is determined by three properties. `Orientation` can be set to `Horizontal` or `Vertical`, and `HorizontalAlignment` and `VerticalAlignment` can be used to determine whether control stacks are positioned next to the top, bottom, left, or right edge of the `StackPanel`. You can even make the stacked controls stack at the center of the `StackPanel` using the `Center` value for the alignment you use.

Figure 14-5 shows two `StackPanel` controls, each of which contains three buttons. The top `StackPanel` has its `Orientation` property set to `Horizontal` and the bottom one has `Orientation` set to `Vertical`.




FIGURE 14-5

The code used here is as follows (you can find it in the downloaded code at `LayoutExamples\StackPanels.xaml`):

```
<Window x:Class="LayoutExamples.StackPanels"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:LayoutExamples"
    mc:Ignorable="d"
    Title="StackPanels" Height="300" Width="300">
    <Grid>
        <StackPanel HorizontalAlignment="Left" Height="128" VerticalAlignment="Top"
            Width="284" Orientation="Horizontal">
            <Button Content="Button" Height="128" VerticalAlignment="Top"
                Width="75"/>
            <Button Content="Button" Height="128" VerticalAlignment="Top"
                Width="75"/>
            <Button Content="Button" Height="128" VerticalAlignment="Top"
                Width="75"/>
        </StackPanel>
        <StackPanel HorizontalAlignment="Left" Height="128" VerticalAlignment="Top"
            Width="284" Margin="0,128,0,0" Orientation="Vertical">
            <Button Content="Button" HorizontalAlignment="Left" Width="284"/>
            <Button Content="Button" HorizontalAlignment="Left" Width="284"/>
            <Button Content="Button" HorizontalAlignment="Left" Width="284"/>
        </StackPanel>
    </Grid>
</Window>
```

## WrapPanel

`WrapPanel` is essentially an extended version of `StackPanel`; controls that “don’t fit” are moved to additional rows (or columns). Figure 14-6 shows a `WrapPanel` control containing multiple shapes, with the window resized to two sizes.



FIGURE 14-6

The code to achieve this effect is shown here (you can find it in the downloaded code at `LayoutExamples\WrapPanel.xaml`):

```
<Window x:Class="LayoutExamples.WrapPanel"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
```

```
xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
xmlns:local="clr-namespace:LayoutExamples"
mc:Ignorable="d"
Title="WrapPanel" Height="92" Width="260">
<WrapPanel Background="AliceBlue">
    <Rectangle Fill="#FF000000" Height="50" Width="50" Stroke="Black"
RadiusX="10" RadiusY="10" />
    <Rectangle Fill="#FF111111" Height="50" Width="50" Stroke="Black"
RadiusX="10" RadiusY="10" />
    <Rectangle Fill="#FF222222" Height="50" Width="50" Stroke="Black"
RadiusX="10" RadiusY="10" />
    <Rectangle Fill="#FFFFFFFF" Height="50" Width="50" Stroke="Black"
RadiusX="10" RadiusY="10" />
</WrapPanel>
</Window>
```

WrapPanel controls are a great way to create a dynamic layout that enables users to control exactly how content should be viewed.

## Grid

Grid controls can have multiple rows and columns that you can use to lay out child controls. You have used Grid controls several times already in this chapter, but in all cases you used a Grid with a single row and a single column. To add more rows and columns, you must use the RowDefinitions and ColumnDefinitions properties, which are collections of RowDefinition and ColumnDefinition objects, respectively, and are specified using property element syntax:

```
<Grid>
    <Grid.RowDefinitions>
        <RowDefinition />
        <RowDefinition />
    </Grid.RowDefinitions>
    <Grid.ColumnDefinitions>
        <ColumnDefinition />
        <ColumnDefinition />
    </Grid.ColumnDefinitions>
    ...
</Grid>
```

This code defines a Grid control with two rows and two columns. Note that no extra information is required here; with this code, each row and column is dynamically resized automatically as the Grid control resizes. Each row will be a third of the height of the Grid, and each column will be half the width. You can display lines between cells in a Grid by setting the Grid>ShowGridlines property to true.

**NOTE** You can also define rows and columns in the grid by clicking the edges of the grid in the Design View. If you move the mouse pointer to the edge of the grid, a yellow line is drawn across the Design View; if you click the edge, the necessary XAML is inserted. When you do this, the Width and Height properties of the rows and columns are always set by the designer, but you can delete them or drag the lines to suit your needs.

You can control the resizing with the `Width`, `Height`, `MinWidth`, `MaxWidth`, `MinHeight`, and `MaxHeight` properties. For example, setting the `Width` property of a column ensures that the column stays at that width. You can also set the `width` property of a column to `*`, which means “fill the remaining space after calculating the width of all other columns.” This is actually the default. When you have multiple columns with a `Width` of `*`, then the remaining space is divided between them equally. The `*` value can also be used with the `Height` property of rows. The other possible value for `Height` and `width` is `Auto`, which sizes the row or column according to its content. You can also use `GridSplitter` controls to enable users to customize the dimensions of rows and columns by clicking and dragging.

Child controls of a `Grid` control can use the attached `Grid.Column` and `Grid.Row` properties to specify which cell they are contained in. Both these properties default to `0`, so if you omit them, then the child control is placed in the top-left cell. Child controls can also use `Grid.ColumnSpan` and `Grid.RowSpan` to be positioned over multiple cells in a table, where the upper-left cell is specified by `Grid.Column` and `Grid.Row`.

### TRY IT OUT Using Rows and Columns: Ch14Ex01\MainWindow.xaml

Return to the example from the beginning of the chapter with the two buttons and follow these steps.

1. Select the `Grid` control by clicking in the XAML View.
2. Move the mouse pointer to the top edge of the grid in Design View; you’ll see an orange line appear across the surface of the grid. Allow room for a button and click to create two columns.
3. Repeat Step 2 on the left edge of the window, creating two rows.
4. Select the first of the two buttons. Note that the action of adding the rows and columns automatically added the `Grid.Row` and `Grid.Column` properties to the button. Change the `Grid.Row` and `Grid.Column` attached properties to `0`.
5. Adjust the `Margin` property to make the button fully visible in the cell.
6. The second button has also been adjusted. For example, a `Margin` has been added. Now delete the `Margin` property from the second button.

7. Add a `GridSplitter` control to the XAML View just before the closing tag of the `Grid` control and set its properties like this:

```
<GridSplitter Grid.RowSpan="2" Width="3" BorderThickness="2" BorderBrush="Black" />
```

8. Run the application. The complete XAML should look like this:

```
<Window x:Class="Ch14Ex01.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:local="clr-namespace:Ch14Ex01"
    mc:Ignorable="d"
    Title="MainWindow" Height="350" Width="525" KeyDown="Window_KeyDown"
    PreviewKeyDown="Window_PreviewKeyDown">
    <Grid KeyDown="Grid_KeyDown" PreviewKeyDown="Grid_PreviewKeyDown">
        <Grid.RowDefinitions>
            <RowDefinition Height="109*"/>
            <RowDefinition Height="210*"/>
        </Grid.RowDefinitions>
        <Grid.ColumnDefinitions>
            <ColumnDefinition Width="191*"/>
            <ColumnDefinition Width="326*"/>
        </Grid.ColumnDefinitions>
        <Button x:Name="button" Content="Button" HorizontalAlignment="Left"
            Margin="27,4,0,0" VerticalAlignment="Top" Width="75" Grid.Column="0"
            Grid.Row="0"/>
        <Button x:Name="rotatedButton" Content="2nd Button" Width="75" Height="22"
            FontWeight="Bold" RenderTransformOrigin="0.5,0.5"
            KeyDown="rotatedButton_KeyDown"
            PreviewKeyDown="rotatedButton_PreviewKeyDown" Grid.Column="1"
            Grid.Row="1" >
            <Button.RenderTransform>
                <TransformGroup>
                    <ScaleTransform/>
                    <SkewTransform/>
                    <RotateTransform Angle="-23.896"/>
                    <TranslateTransform/>
                </TransformGroup>
            </Button.RenderTransform>
        </Button>
        <GridSplitter Grid.RowSpan="2" Width="3" BorderThickness="2"
            BorderBrush="Black" />
    </Grid>
</Window>
```

Figure 14-7 shows the application running with the splitter pushed to two positions.




FIGURE 14-7

### How It Works

By dividing the grid into two columns and two rows, you have changed how the controls can be positioned in the grid. When you set the `Grid.Row` and `Grid.Column` to 0 for the first button, you move it from its previous position on the form to the top-left section.

The second button more or less stays put, but when you drag the `GridSplitter` slider, you see that the margin of the button is now relative to the left edge of the column in which it is placed, meaning that it slides across the window as you move the slider.

---

## THE GAME CLIENT

Now that you know the basics of what it means to work with WPF and Visual Studio, it is time to start working with the controls to create something useful. The remaining sections of this chapter and Chapter 15 are dedicated to writing a game client for the card game you have been developing over the previous chapters. You are going to use a lot of controls to write the game client, and you are even going to write one yourself.

In this chapter you are going to write the supporting dialog boxes of the game — this includes the About, Options, and New Game windows.

### The About Window

An About window, or About box as it's sometimes called, is used to display information about the developer of the application and the application itself. Some About windows are quite complex, like the one found in Microsoft Office applications and Visual Studio, and display version and licensing

information. By convention, the About window can be accessed from the Help menu where it is usually the last item on the list.

Figure 14-8 shows a screenshot of the finished dialog box that you are about to create.

## Designing the User Interface

An About window is not something that the user is going to see very often. In fact, the reason that it is usually located on the Help menu is that it is very often only used when the user needs to find information about the version of the application or who to contact when something is wrong. But this also means that it is something the user has a specific purpose for visiting and if you include such a window in your application, you should treat it as important.



**FIGURE 14-8**

Whenever you are designing an application, you should strive to keep the look and feel as consistent as possible. This means that you should stick to a few select colors and use the same styling of controls everywhere in the application. In the case of Karli Cards, you are going to work with three main colors — red, black, and white.

If you look at Figure 14-8 you will see that the top-left corner of the window is occupied by a Wrox Press logo. You have not used images before, but adding a few select images to your applications can make the user interface look more professional.

## The Image Control

Image is a very simple control that can be used to great effect. It allows you to display a single image and to resize this image as you see fit. The control exposes two properties, as shown in Table 14-2.

**TABLE 14-2:** Image Control

PROPERTY	DESCRIPTION
Source	Use this property to specify the location of the image. This can be a location on disk or somewhere on the web. As you will see in Chapter 15, it is also possible to create a static resource and use it as the source.

PROPERTY	DESCRIPTION
Stretch	<p>It's actually pretty rare to have an image that is exactly the right size for your purpose, and sometimes the size of the image must change as the application window is resized. You can use this property to control how the image behaves. There are four possibilities:</p> <ul style="list-style-type: none"> <li>None — The image doesn't resize.</li> <li>Fill — The image resizes to fill the entire space. This may distort the image.</li> <li>Uniform — The image keeps its aspect ratio and doesn't fill the available space if this would change the aspect ratio.</li> <li>UniformToFill — The image keeps its aspect ratio and fills the available space. If keeping the ratio means that some of the image is too large for the space available, the image is clipped to fit.</li> </ul>

## The Label Control

You have already seen this most simple of controls used in some of the previous examples. It displays simple text information to the user and in some cases relays information about shortcut keys. The control uses the `Content` property to display its text. The `Label` control displays text on a single line. If you prefix a letter with an underscore “\_” character, the letter will become underlined and it will then be possible to access the control directly by using the prefixed letter and Alt. For example, `_Name` assigns the shortcut Alt+N to any control directly following the label.

## The TextBlock Control

Like `Label`, this control displays simple text without any complicated formatting. Unlike the `Label`, the `TextBlock` control is capable of displaying multiple lines of text. It is not possible to format individual parts of the text.

The `TextBlock` displays the text even if it will not fit in the space granted to the control. The control itself does not provide any scrollbars in this case, but it can be wrapped in a handy view control when needed: the `ScrollViewer`.

## The Button Control

Like the `Label` control, you have already seen quite a bit of the `Button` control. This control is used everywhere and is easily recognized on a user interface. Your users will expect that they can left-click it to perform an action — no more and no less. Altering this behavior will most likely lead to bad interface design and frustrated users.

By default, the button displays itself with a single short line of text or an image that describes what happens when you click on it.

The button does not contain any properties to display images or text, but you can use the `Content` property to display simple text or embed an `Image` control in the content to display an image. You can find this code in the downloaded code at `Ch14Ex01\ImageButton.xaml`:

```
<Button HorizontalAlignment="Left" VerticalAlignment="Top" Width="75" Margin="10" >
    <StackPanel Orientation="Horizontal">
        <Image Source=".\\Images\\Delete_black_32x32.png" Stretch="UniformToFill" Width="16" Height="16" />
        <TextBlock>Delete</TextBlock>
    </StackPanel>
</Button>
```

**NOTE** The image for the button is included in the code download in Ch14Ex01\Images.

Figure 14-9 shows the Delete button with text and an image.




FIGURE 14-9

**NOTE** To complete the following example, you need an image for a banner. This image is included in the download for this chapter in KarliCards Gui\Images\Banner.png.

### TRY IT OUT Creating the About Window: KarliCards Gui\About.xaml

Before you can start the About window, you need a project to work on. This is just one of many windows you are going to make in this and the next chapter, so go ahead and create a new WPF application project and name it KarliCards Gui. Name the solution KarliCards.

1. In the Solution Explorer, right-click the KarliCards Gui project and select Add ➔ Window. Name the window About.xaml.
2. Resize the window by clicking and dragging it or by setting these properties:

```
Height="300" Width="434" MinWidth="434" MinHeight="300"
ResizeMode="CanResizeWithGrip"
```

3. Select the Grid and create four rows by clicking at the edges of the grid. Don't worry too much about the exact positioning of the rows; instead change the values like this:

```
<Grid.RowDefinitions>
    <RowDefinition Height="58"/>
    <RowDefinition Height="20"/>
    <RowDefinition />
    <RowDefinition Height="42"/>
</Grid.RowDefinitions>
```

4. Drag a Canvas control from the Toolbox into the top-most row. Remove any properties inserted by Visual Studio and add this:

```
Grid.Row="0" Background="#C40D42"
```

5. Select the new canvas and drag an image control onto it. Change its properties like so:

```
Height="56" Canvas.Left="0" Canvas.Top="0" Stretch="UniformToFill"  
Source=".\\Images\\Banner.png"
```

6. Right-click the project and Select Add  $\Rightarrow$  New Folder. Create a directory called `Images`.  
7. Right-click the new directory in the Solution Explorer and select Add  $\Rightarrow$  Existing Item. Browse to the images of this chapter. Select them all and click Add. The banner is now displayed in Design.  
8. Select `Canvas` and drag a `Label` control onto it. Change its properties like this:

```
Canvas.Right="10" Canvas.Top="25" Content="Karli Cards" Foreground="#FFF7EFEF"  
FontFamily="Times New Roman"
```

9. Select `Grid` and drag a new `canvas` control onto it. Change its properties to:

```
Grid.Row="1" Background="Black"
```

10. Select the new `Canvas` control and drag a `Label` onto it. Change its properties like this:

```
Canvas.Left="5" Canvas.Top="0" FontWeight="Bold" FontFamily="Arial"  
Foreground="White"  
Content="Karli Cards (c) Copyright 2012 by Wrox Press and all readers"
```

11. Select `Grid` again, and drag the last `Canvas` into the bottom-most row. Change its properties like this:

```
Grid.Row="3"
```

12. Select the new `Canvas` control and drag a `Button` onto it. Change its properties to this:

```
Content="_OK" Canvas.Right="12" Canvas.Bottom="10" Width="75"
```

13. Select `Grid` again, and drag a `StackPanel` into the last center row. Change its properties to:

```
Grid.Row="2"
```

14. Select `StackPanel` and drag two `Label` controls and one `TextBlock` into it, in that order.

15. Change the top-most `Label` like this:

```
Content="CardLib and Idea developed by Karli Watson" HorizontalAlignment="Left"  
VerticalAlignment="Top" Padding="20,20,0,0" FontWeight="Bold"  
Foreground="#FF8B6F6F"
```

16. Change the next `Label` like this:

```
Content="Graphical User Interface developed by Jacob Hammer"  
HorizontalAlignment="Left" Padding="20, 0,0,0" VerticalAlignment="Top"  
FontWeight="Bold" Foreground="#FF8B6F6F"
```

17. Change `TextBlock` like this:

```
Text="Karli Cards developed with Visual C# 6 for Wrox Press.  
You can visit Wrox Press at http://www.wrox.com."
```

```
Margin="0, 10, 0, 0" Padding="20, 0, 0, 0" TextWrapping="Wrap"
HorizontalAlignment="Left" VerticalAlignment="Top" Height="39"
```

18. Double-click the button and, in the event handler, add this code:

```
private void Button_Click(object sender, RoutedEventArgs e)
{
    this.Close();
}
```

19. In the Solution Explorer, double-click the `App.xaml` file and change the name `MainWindow.xaml` to `About.xaml`.

20. Run the application.

### How It Works

You begin by setting some properties on the window. By setting `MinWidth` and `MinHeight`, you prevent the user from resizing the window to a point where it obscures the content. The `ResizeMode` is set to `CanResizeWithGrip`, which displays a small grip section in the bottom-right corner of the window that indicates to the user that the window can be resized.

Next you add four rows to the grid. By doing this, you define the basic structure of the window. By setting rows 1, 2, and 4 to fixed heights, you ensure that only the third row can change height; this is the row that holds the content.

Then you add the first `Canvas` control. This provides you with a handy place to set the background color of the first row. By ensuring that the canvas has no specific size, you force the canvas to fill the top row in the grid.

The `Image` control that is added to the canvas is fixed to the left and top edges of the canvas. This ensures that as the window resizes, the image stays put. You also gave the image a fixed height, but left the width open. With the `Stretch` property set to `UniformToFill`, this allows the `Image` control to use the height as a guide for the aspect ratio. The control simply changes its width to match the scale specified by the height and aspect ratio.

For the final part of the first row you add a single `Label` control and bind it to the top-right edge of the canvas, ensuring that when the window resizes, the `Label` moves with the right edge.

Then you start on the second row, which is filled by another `Canvas` control that has a `Label` added to it.

The bottom `Canvas` is more of the same, but this time you add a button to it and bind that button to the bottom-right side of the canvas. This ensures that when the window is resized, the button sticks to the bottom-right side of the window. The underscore “`_`” before the text `OK` creates a `Alt+O` shortcut for the button.

Finally, you add a `StackPanel` to the third row and add `Labels` and a `TextBlock` control to it. By setting the `Padding` of the first label to `20, 20, 0, 0`, you push the content down from the row above by 20 pixels and out from the left edge, also by 20 pixels.

The padding of the next label is set to `20,0,0,0`, which pushes the content out from the edge because the space between the two labels is fine and doesn't need any extra space.

The `TextBlock` was then introduced. The property `TextWrapping` is set to `Wrap`, which causes the text to wrap if it can't fit on a single line. As the window resizes and the line becomes longer, the text is automatically fitted into as few lines as needed. Both the `Margin` and `Padding` properties are used here. The `Margin` property is set so it pushes the entire control down 10 pixels from the labels above, and the `Padding` is set so it pushes the content of the control in by 20 pixels from the left edge.

The code in the event handler closes the window. In this case, this is the same as closing the entire application, because in Step 19 you changed the startup window to be the `About` window, so closing it is the same as closing the application.

## The Options Window

The next window you are going to create is the Options window. This window will allow the players to set a number of parameters that will alter the game play. It will also allow you to use some controls that you haven't used yet: the `CheckBox`, `RadioButton`, `ComboBox`, `TextBox`, and `TabControl` controls.

Figure 14-10 shows the window with the first tab selected. At first glance the window looks much like the `About` window, but there is a lot more to do on this window.

### The `TextBox` Control

Previously in this chapter you used the `Label` and `TextBlock` controls. These controls are designed exclusively for displaying text to the user. The `TextBox` control allows the user to type text into the application. Although it can just display text as well, you should not use it for this purpose unless the user is allowed to edit the displayed text. If you decide that you want to display text using a text box, be sure to set its `IsEnabled` property to `false` to prevent users from being able to edit it.




FIGURE 14-10

You control how the text is displayed and can be entered into the `TextBox` using a number of properties shown in Table 14-3.

**TABLE 14-3:** TextBox Properties

PROPERTY	DESCRIPTION
Text	The text currently displayed in the TextBox control.
IsEnabled	When this is set to <code>true</code> , the user can edit the text in the TextBox. When it is <code>false</code> , the text is grayed out and the user cannot give focus to the control.
TextWrapping	Sometimes you want the TextBox to display only a single line of text. In this case, you can set this property to <code>NoWrap</code> . This is the default. If you want your text to be displayed on multiple lines, you can set it to either <code>Wrap</code> or <code>WrapWithOverflow</code> . <code>Wrap</code> will cause the text that extends beyond the edge of the box to be moved to the line below. <code>WrapWithOverflow</code> will in some cases allow very long words to extend beyond the edge if no suitable breakpoint can be determined.
VerticalScrollBarVisibility	If your TextBox allows the user to enter multiple lines of text, then the user can potentially type text that will disappear below the lower edge of the box. In that case, it's a good idea to display a scrollbar. Set this to <code>Auto</code> if you want the scrollbar to appear only if the text is too long to be displayed. Set it to <code>Visible</code> to always display it, and <code>Hidden</code> or <code>Disabled</code> to never display a scrollbar.
AcceptsReturn	This property controls how text can be entered into the control. If you set this to <code>false</code> , which is the default, then the user can't break the line with a Return.

## The CheckBox Control

Checkboxes present the users with options that they can select or clear. You should use a CheckBox if you want to present an option to the users that can be turned on or off, or want the users to answer yes or no to a question. For example, in the Options dialog box, you want the user to answer to decide whether they should play against the computer. To this end a CheckBox with the text “Play Against Computer” is used.

A CheckBox is designed to be used as a single entity that is unaffected by other Checkboxes on the view. You will sometimes see Checkboxes used in a way that links them together so that selecting one causes another to become cleared, but this is not the intended use for this control. If you want this functionality, you should use a RadioButton, described in the next section.

Checkboxes can also display a third state, which is known as “indeterminate” and is supposed to indicate that the yes/no answer could not be answered. This state is commonly used when a CheckBox is used to show information about something else. For example, Checkboxes are sometimes used to indicate whether all child nodes in a Tree View are selected. In this case, the CheckBox

will be selected if all nodes are selected, cleared if none are, and indeterminate if some, but not all, are selected.

Table 14-4 lists the properties commonly used to control the CheckBox control.

**TABLE 14-4:** CheckBox Properties

PROPERTY	DESCRIPTION
Content	The CheckBox is a Content control and its display can therefore be heavily customized. Adding a text to the Content property yields the default view.
IsThreeState	Used to indicate if the control can have two or three states. The default is false, meaning that only two possible values exist.
IsChecked	This is either true or false. By default, setting it to true displays a check-mark. If IsThreeState is true, null is possible and indicates that the state is indeterminate.

## The RadioButton Control

RadioButtons are used with other RadioButtons to allow users to choose between multiple options where only one can be selected at any time. You should use RadioButtons when you want the users to answer a question that has a very limited number of possible values. If there are more than four or five possible values, you should consider using a ListBox or a ComboBox instead. In the Options window you will create shortly, the user can choose the skill level of the computer player. There are three options: Dumb, Good, and Cheats. Only one should ever be selected at any given time.

When more than one RadioButton is displayed in the same view they will by default know about each other and as soon as any one of them is selected, all the others are cleared. If you have multiple unrelated RadioButtons on the same view, they can be grouped together to avoid controls clearing the values of unrelated controls.

You can control RadioButtons with the properties listed in Table 14-5.

**TABLE 14-5:** RadioButton Properties

PROPERTY	DESCRIPTION
Content	RadioButtons are Content controls and can therefore have their display modified. By default, you enter a text in the Content.
IsChecked	This is either true or false. If IsThreeState is true, null is possible and indicates that the state is indeterminate.
GroupName	The name of the group the control belongs to. By default this is empty and any RadioButtons without a GroupName is considered in the same group.

## The ComboBox Control

Like the `RadioButton` and `CheckBox` controls, `ComboBoxes` allow users to select exactly one option. However, `ComboBoxes` are fundamentally different from the other two in two ways:

- `ComboBoxes` display the possible choices in a drop-down list.
- It is possible to allow the users to type new values.

`ComboBoxes` are commonly used to display long lists of values, such as country or state names, but they can be used for many purposes. In the Options dialog box, a `ComboBox` is used to display a list from which the user can choose the number of players. Although this could just as well have been done using `RadioButtons`, the use of a `ComboBox` saves space in the view.

A `ComboBox` can be changed to display itself with a `TextBox` at the top that allows the users to type any values that they feel are missing. One of the exercises of this chapter asks you to add a `ComboBox` to the Options dialog box from which the users can either type their name or select it from a list.

The two properties — `IsReadOnly` and `IsEditable` — are very important for the behavior of the control and work together to provide four possible ways for the user to select the value of the `ComboBox` using the keyboard (see Table 14-6):

TABLE 14-6: `IsReadOnly` and `IsEditable` Combinations

	ISREADONLY IS TRUE	ISREADONLY IS FALSE
<code>IsEditable</code> is true	The <code>TextBox</code> is displayed but the control does not react to key presses. If a selection is made in the list, the text can be selected in the <code>TextBox</code> .	The <code>TextBox</code> is displayed and the user can type anything she wishes. If something is typed that is in the list, it is selected. The control will display the best possible match as the user is typing.
<code>IsEditable</code> is false	When <code>IsEditable</code> is false, <code>IsReadOnly</code> no longer has any effect because the <code>TextBox</code> is not displayed. When the control is selected, the user can select a value from the list by typing but it is not possible to type a value that isn't in the list.	

A `ComboBox` is an Items control, which means that you can add multiple items to it. Table 14-7 shows additional properties for the `ComboBox` control.

**TABLE 14-7:** Other ComboBox Properties

COMBOBOX PROPERTY	DESCRIPTION
Text	The Text property represents the text displayed at the head of the ComboBox. It is either an element of the list or a new text typed by the user.
SelectedIndex	Represents the index of the selected item in the list. If this is -1 then no selection is made. This is also the case if the user has typed something that was not in the list.
SelectedItem	Represents the actual item of the list, not just the index or the text. If nothing is selected or the user has typed something new, this returns null.

## The TabControl

The TabControl is radically different than the other controls presented this section. It is a layout control that is used to group controls on pages that can be selected by clicking on them.

Tab controls are used when you want to display a lot of information in a single window but don't want to clutter the view too much. In this case, you should divide the information into groups of related items and create a single page for each group. Generally speaking, you should never allow controls on one page to affect controls on another page. If you do so anyway, the user will not realize that something has changed on another page and will be confused when settings change behind her back.

By default each page is constructed of `TabItems` that, by default, are populated by a single `Grid` control, but you can change the `Grid` to any other control as you see fit. On each tab, you can lay out your UI and, by selecting the `TabItems`, you can change between the tabs. Each `TabItem` has a `Header` that can be used to display the tab itself. This can be used as a `Content` control, meaning that you can customize how the header is displayed so that it can be more than just a text.

### TRY IT OUT Designing the Options Window: KarliCards Gui \Options.xaml

The first thing that you probably notice when you see the Options window is that it looks remarkably like the About window, and that is true. Because of that, it is possible to reuse at least some of the code from the previous example.

1. Right-click the project in the Solution Explorer and chose Add  $\Rightarrow$  Window. Name the window `Options.xaml`.
2. Delete the `Grid` control that is inserted by default.
3. Open the `About.xaml` window described earlier, copy the `Grid` control and all its content, and paste it into the new `Options.xaml` file.

4. Change the window properties like this:

```
Title="Options" Height="345" Width="434" ResizeMode="NoResize"
```

5. Delete the StackPanel and all of its content.

6. Delete the Canvas control with the Grid.Row property set to 3 and all of its content.

7. Delete the Label control from the Canvas control with the Grid.Row property set to 1.

8. Change the Label control in the Canvas with the Grid.Row property set to 0 like this:

```
<Label Canvas.Right="10" Canvas.Top="13" Content="Options" Foreground="#FFF7EFEF"  
FontFamily="Times New Roman" FontSize="24" FontWeight="Bold" />
```

9. Drag a StackPanel into the bottom row and set its properties to this:

```
Grid.Row="3" Orientation="Horizontal" FlowDirection="RightToLeft"
```

10. Add two buttons to the StackPanel like this:

```
<Button Content="_Cancel" Height="22" Width="75" Margin="10,0,0,0"  
Name="cancelButton" />  
<Button Content="_OK" Height="22" Width="75" Margin="10,0,0,0"  
Name="okButton" />
```

11. Drag a TabControl into the second row and set its properties like this:

```
Grid.RowSpan="2" Canvas.Left="10" Canvas.Top="2" Width="408" Height="208"  
Grid.Row="1"
```

12. Change the Header property of each of the two TabItem controls to Game and Computer Player, respectively.

Your window now looks like Figure 14-11 and it is time to insert some content into the tab items.




FIGURE 14-11

13. Select the Game TabItem and drag a CheckBox control onto it. Set its properties like this:

```
Content="Play against computer" HorizontalAlignment="Left" Margin="11,33,0,0"
VerticalAlignment="Top" Name="playAgainstComputerCheck"
```

14. Drag a Label control and then a ComboBox control into the TabItem and set their properties like this:

```
<Label Content="Number of players" HorizontalAlignment="Left"
Margin="10,54,0,0" VerticalAlignment="Top" />
<ComboBox HorizontalAlignment="Left" Margin="196,58,0,0"
VerticalAlignment="Top" Width="86" Name="numberOfPlayersComboBox"
SelectedIndex="0" >
    <ComboBoxItem>2</ComboBoxItem>
    <ComboBoxItem>3</ComboBoxItem>
    <ComboBoxItem>4</ComboBoxItem>
</ComboBox>
```

15. Select the second TabItem with the header Computer Player. Drag a Label and three RadioButtons onto the Grid and set their properties like this:

```
<Label Content="Skill Level" HorizontalAlignment="Left"
Margin="10,10,0,0" VerticalAlignment="Top" />
<RadioButton Content="Dumb" HorizontalAlignment="Left"
Margin="37,41,0,0" VerticalAlignment="Top" IsChecked="True"
Name="dumbAIRadioButton"/>
<RadioButton Content="Good" HorizontalAlignment="Left"
Margin="37,62,0,0" VerticalAlignment="Top" Name="goodAIRadioButton"/>
<RadioButton Content="Cheats" HorizontalAlignment="Left"
Margin="37,83,0,0" VerticalAlignment="Top"
Name="cheatingAIRadioButton"/>
```

16. The layout of the window is now complete. Open the App.xaml file and change StartupUri to Options.xaml.

17. Run the application.

### How It Works

The window's ResizeMode is set to NoResize. You can therefore position the controls without regard to what happens if the window changes size, because the user can no longer resize the window.

The StackPanel in Step 9 has a new property, FlowDirection, which is set to RightToLeft. This causes the two buttons that are added to it to cling to the right edge of the dialog box rather than the left edge that is the default. Interestingly, this also changes the meaning of the Margin property of the two buttons, causing Left and Right to be swapped.

The RadioButtons on the second tab are set up without specifying a GroupName, which causes them to be grouped together. You set the IsChecked property to true on the first one, which makes this the default selection.

## Handling Events in the Options Window

The window looks fine at this point, and there are even a few things users can do with it, although nothing happens when a setting is changed. Users expect that the options they choose are stored and used by the application. You could do this by storing the values of the controls in the window, but this is not very flexible and mixes the data of the application with the GUI, which is not a good idea. Instead, you should create a class to hold the selections made by the users.

### TRY IT OUT Handling Events: KarliCards Gui\Options.xaml

In this example, you will add a new class to the project that will contain the selections made by the user and handle events that happen as the user changes selections.

1. Add a new class to the project and name it GameOptions.cs.
2. Enter this code:

```
using System;
namespace KarliCards_Gui
{
    [Serializable]
    public class GameOptions
    {
        public bool PlayAgainstComputer { get; set; }
        public int NumberOfPlayers { get; set; }
        public int MinutesBeforeLoss { get; set; }
        public ComputerSkillLevel ComputerSkill { get; set; }
    }
    [Serializable]
    public enum ComputerSkillLevel
    {
        Dumb,
        Good,
        Cheats
    }
}
```

3. Return to the Options.xaml.cs code-behind file and add a private field to hold the GameOptions instance:

```
private GameOptions _gameOptions;
```

4. Add this code to the constructor:

```
using System.IO;
using System.Windows;
using System.Xml.Serialization;

namespace KarliCards_Gui
{
    /// <summary>
    /// Interaction logic for Options.xaml
    /// </summary>
    public partial class Options : Window
```

```

    {
        private GameOptions _gameOptions;

        public Options()
        {
            if (_gameOptions == null)
            {
                if (File.Exists("GameOptions.xml"))
                {
                    using (var stream = File.OpenRead("GameOptions.xml"))
                    {
                        var serializer = new XmlSerializer(typeof(GameOptions));
                        _gameOptions = serializer.Deserialize(stream) as GameOptions;
                    }
                }
                else
                    _gameOptions = new GameOptions();
            }

            InitializeComponent();
        }
    }

```

5. Go to Design View and double-click each of the three RadioButtons to add the Checked event handler to the code-behind file. Change the handlers like this:

```

private void dumbAIRadioButton_CheckedChanged(object sender, RoutedEventArgs e)
{
    _gameOptions.ComputerSkill = ComputerSkillLevel.Dumb;
}
private void goodAIRadioButton_CheckedChanged(object sender, RoutedEventArgs e)
{
    _gameOptions.ComputerSkill = ComputerSkillLevel.Good;
}
private void cheatingAIRadioButton_CheckedChanged(object sender, RoutedEventArgs e)
{
    _gameOptions.ComputerSkill = ComputerSkillLevel.Cheats;
}

```

6. Return to Design View and select the TextBox on the Game tab. Click the lightning icon on the Properties panel and double-click the GotFocus event to add the handler to the code-behind file.
7. Enter this code:

```

private void timeAllowedTextBox_GotFocus(object sender, RoutedEventArgs e)
{
    timeAllowedTextBox.SelectAll();
}

```

8. Select the TextBox again in Design View and add the PreviewMouseLeftButtonDown event handler to the code-behind file.
9. Enter this code:

```

private void timeAllowedTextBox_PreviewMouseLeftButtonDown(object sender,
    MouseButtonEventArgs e)

```

```
{  
    var control = sender as TextBox;  
    if (control == null)  
        return;  
    Keyboard.Focus(control);  
    e.Handled = true;  
}
```

10. Run the application.

### How It Works

The new class is currently just a number of properties that store the values from the Options window. It is marked as `Serializable` to make it possible to save it to a file.

The `Checked` event of a `RadioButton` is raised whenever the user selects it. You handle this event in order to set the value of the `ComputerSkillLevel` property of the `GameOptions` instance.

---

## Data Binding

Data binding is a way of declaratively connecting controls with data. In the Options window, you handled the `Checked` event of the `RadioButtons` in order to set the value of the `ComputerSkillLevel` property in the `GameOptions` class. This works well, and you can use code and event handling to set all the values you have in a window, but very often it is better to bind the properties of your controls directly to the data.

A binding consists of four components:

- The binding target, which specifies the object on which the binding is used
- The target property, which specifies the property to set
- The binding source, which specifies the object used by the binding
- The source property, which specifies which property holds the data

You don't always set all of these elements explicitly; particularly the binding target is very often implicitly specified by the fact that you are setting a binding to a property on a control.

The binding source is always set in order to make a binding work, but it can be set in several ways. In the following sections and in Chapter 15, you are going to see several ways of binding data from sources.

### The DataContext

A `DataContext` control defines a data source that can be used for data binding on all child elements of an element. You will often have a single instance of a class that holds most of the data that is used in a view. If this is the case you can set the `DataContext` of the window to the instance of that object, which makes you able to bind properties from that class in your view. This is demonstrated in the "Dynamic Binding to External Objects" section.

## Binding to Local Objects

You can bind to any .NET object that has the data you need as long as the compiler can locate the object. If the object is found in the same context, that is the same XAML block, as the control using the object, you can specify the binding source by setting the `ElementName` property of the binding. Take a look at this changed `ComboBox` from the Options window:

```
<ComboBox HorizontalAlignment="Left" Margin="196,58,0,0" VerticalAlignment="Top"
Width="86" Name="numberOfPlayersComboBox" SelectedIndex="0"
IsEnabled="{Binding ElementName=playAgainstComputerCheck, Path=IsChecked}" >
```

Notice the `.IsEnabled` property. Instead of specifying `true` or `false`, there is now lengthy text within a couple of curly brackets. This way of specifying property values is called *markup extension syntax*, and is a shorthand for specifying properties. The same could have been written like this:

```
<ComboBox HorizontalAlignment="Left" Margin="196,58,0,0"
VerticalAlignment="Top" Width="86" Name="numberOfPlayersComboBox"
SelectedIndex="0" >
    <ComboBox.IsEnabled>
        <Binding ElementName="playAgainstComputerCheck"
Path="IsChecked" />
    </ComboBox.IsEnabled>
```

Both examples set the binding source to the `playAgainstComputerCheck CheckBox`. The source property is specified in the `Path` to be the `.IsChecked` property.

The binding target is set to the `.IsEnabled` property. Both examples do this by specifying the binding as the content of the property — they just do it using different syntax. Finally, the binding target is implicitly specified by the fact that the binding is done on the `ComboBox`.

The binding in this example causes the `.IsEnabled` property of the `ComboBox` to be set or cleared depending on the value of the `.IsChecked` property of the `CheckBox`. The result is that without any code, the `ComboBox` is enabled and disabled when the user changes the value of the `CheckBox`.

## Static Binding to External Objects

It is possible to create object instances on the fly by specifying that a class is used as a resource in the XAML. This is done by adding a namespace to the XAML to allow the class to be located, and then declaring the class as a resource on an element in the XAML.

You can create resource references on parent elements of the object that you want to data bind.

### TRY IT OUT Creating a Static Data Binding: KarliCards Gui\NumberOfPlayers.cs

In this example you create a new class to hold the data for the `ComboBox` in the Options window and bind it to the control.

1. Add a new class to the project and name it `NumberOfPlayers.cs`.
2. Add this code:

```
using System.Collections.ObjectModel;
namespace KarliCards_Gui
```

```
{  
    public class NumberOfPlayers : ObservableCollection<int>  
    {  
        public NumberOfPlayers()  
            : base()  
        {  
            Add(2);  
            Add(3);  
            Add(4);  
        }  
    }  
}
```

3. Return to the Options.xaml file's Design View and select the Window root element.
4. Select the Canvas element that contains the ComboBox and add this code below it, and above the TabControl declaration.

```
<Canvas.Resources>  
    <local:NumberOfPlayers x:Key="numberOfPlayersData" />  
</Canvas.Resources>
```

5. Select the ComboBox and remove the three ComboBoxItems from it.
6. Add this property to it:

```
ItemsSource="{Binding Source={StaticResource numberOfPlayersData}}"
```

## How It Works

There is a lot happening in this example. The class `NumberOfPlayers` derives from a special collection named `ObservableCollection`. This base class is a collection that has been extended to make it work better with WPF. In the constructor of the class, you add the values to the collection.

Next you create a new resource on the `Canvas`. You could have created this resource on any parent element of the `ComboBox`. When a resource is specified on an element, all child elements can use it.

Finally you set the `ItemsSource` to a binding. The `ItemsSource` property is specifically designed to allow you to specify a binding for the collection of items on an `Items` control. In the binding you just need to specify the binding source. The binding target, target property, and source property settings are handled by the `ItemsSource` property.

---

## Dynamic Binding to External Objects

Now you can bind to objects that are created on the fly as they are needed in order to provide some data. What if you already have an instantiated object that you want to use for data binding? In that case, you need to do a little plumbing in the code.

In the case of the Options window, you don't want the options to be cleared every time the window is opened, and you want the selections the user made to persist and be used in the rest of the application.

You can do this in code by setting the value of the `DataContext` property to the instance.

**TRY IT OUT** Creating Dynamic Bindings: KarliCards Gui\GameOptions.cs

In this example you bind the remaining controls to the GameOptions instance in the Options window.

1. Go to the Options.xaml.cs code-behind file.
  2. At the bottom of the constructor, but above InitializeComponent(), add this line:
- ```
DataContext = _gameOptions;
```
3. Go to the GameOptions class and change it like this:

```
using System;
using System.ComponentModel;
namespace KarliCards_Gui
{
    [Serializable]
    public class GameOptions : INotifyPropertyChanged
    {
        private bool _playAgainstComputer = true;
        private int _numberOfPlayers = 2;
        private ComputerSkillLevel _computerSkill = ComputerSkillLevel.Dumb;
        public int NumberOfPlayers
        {
            get { return _numberOfPlayers; }
            set
            {
                _numberOfPlayers = value;
                OnPropertyChanged(nameof(NumberOfPlayers));
            }
        }
        public bool PlayAgainstComputer
        {
            get { return _playAgainstComputer; }
            set
            {
                _playAgainstComputer = value;
                OnPropertyChanged(nameof(PlayAgainstComputer));
            }
        }
        public ComputerSkillLevel ComputerSkill
        {
            get { return _computerSkill; }
            set
            {
                _computerSkill = value;
                OnPropertyChanged(nameof(ComputerSkill));
            }
        }
        public event PropertyChangedEventHandler PropertyChanged;
        private void OnPropertyChanged(string propertyName)
        {
            PropertyChanged?.Invoke(this, new PropertyChangedEventArgs(propertyName));
        }
    }
}
```

```
public enum ComputerSkillLevel
{
    Dumb,
    Good,
    Cheats
}
```

4. Return Options.xaml and select the CheckBox. Add the IsChecked property like this:  
`IsChecked="{Binding Path=PlayAgainstComputer}"`
5. Select the ComboBox and change it like this, removing the SelectedIndex property and changing the ItemsSource and SelectedValue properties:  
`<ComboBox HorizontalAlignment="Left" Margin="196,58,0,0" VerticalAlignment="Top"
Width="86" Name="numberOfPlayersComboBox"
ItemsSource="{Binding Source={StaticResource numberOfPlayersData}}"
SelectedValue="{Binding Path=NumberOfPlayers}" />`
6. Select and double-click the OK button to add the Click event handler to the code-behind file. Do the same with the Cancel button and add this code to the handlers:

```
private void okButton_Click(object sender, RoutedEventArgs e)
{
    using (var stream = File.Open("GameOptions.xml", FileMode.Create))
    {
        var serializer = new XmlSerializer(typeof(GameOptions));
        serializer.Serialize(stream, _gameOptions);
    }
    Close();
}
private void cancelButton_Click(object sender, RoutedEventArgs e)
{
    _gameOptions = null;
    Close();
}
```

7. Run the application.

### How It Works

Setting the `DataContext` of the window to an instance of `GameOptions` allows you to bind to this instance simply by specifying the property to use in the binding. This is done in Steps 4 and 5. Note that the `ComboBox` is filled with items from a static resource, but the selected value is set in the `GameOptions` instance.

The `GameOptions` class is changed quite a bit. It now implements the `INotifyPropertyChanged` interface, which means that the class is now able to inform WPF that a property has changed. In order for this notification to work, you have to call the subscribers to the `PropertyChanged` event defined by the interface. For this to happen, the property setters have to actively call them, which is done using the helper method `OnPropertyChanged`.

When the `OnPropertyChanged` method is called, we use a new expression introduced by C# 6: `nameof`. When we call `nameof(...)` with an expression, it will retrieve the name of the final identifier. This is particularly useful in the case of the `OnPropertyChanged` method, because it takes the name of the property that is being changed as a string.

The OK button event handler saves the settings to disk using an `XmlSerializer`. The Cancel event handler sets the game options field to null, ensuring that the selections made by the user are cleared. Both event handlers close the window.

## Starting a Game with the `ListBox` Control

You are now only one window short of having created all the supporting windows in the game. The last window before creating the game board is a window where the player can add new players and select the players who will be participating in a new game. This window will use a `ListBox` to display the names of the players.

`ListBoxes` and `ComboBoxes` can often be used for the same purpose, but where a `ComboBox` normally allows you to select only a single entry, `ListBoxes` often allows the user to select multiple items. Another key difference is that a `ListBox` will display its content in a list that is always expanded. This means that it takes up more real estate on the window, but it allows the user to see the options available right away.

Table 14-8 lists a few particularly interesting properties for the `ListBox` control.

**TABLE 14-8:** Interesting `ListBox` Properties

| PROPERTY                   | DESCRIPTION                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>SelectionMode</code> | This property controls how the user can select items from the list. There are three possible values: <code>Single</code> , which allows the user to select only one item, <code>Multiple</code> , which allows the user to select multiple items without holding down the <code>Ctrl</code> key, and <code>Extended</code> , which allows the user to select multiple consecutive items by holding down the <code>Shift</code> key, and non-consecutive items by holding down the <code>Ctrl</code> key. |
| <code>SelectedItem</code>  | Gets or sets the first selected item or null if nothing is selected. Even if multiple items are selected, only the first item is returned.                                                                                                                                                                                                                                                                                                                                                               |
| <code>SelectedItems</code> | Gets a list containing the items that are currently selected.                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| <code>SelectedIndex</code> | Works like <code>SelectedItem</code> , but returns the index instead of the item itself and -1 instead of null if nothing is selected.                                                                                                                                                                                                                                                                                                                                                                   |

**TRY IT OUT** Creating the Start Game Window: KarliCards Gui\StartGame.xaml

This window is displayed to the players when a new game starts. It will allow the players to enter their names and select them from a list of known players.

1. Create a new window and name it `StartGame.xaml`.
2. Delete the `Grid` element from the window and copy the main `Grid` and its content from the `Options.xaml` window instead.
3. Remove all the content from the `Canvas` control that has its `Grid.Row` property set to 1.
4. Change the window title to “Start New Game” and set these properties:

```
Height="345" Width="445" ResizeMode="NoResize"
```

5. Change the content of the label in grid row 0 to “New Game.”
6. Open the `GameOptions.cs` file and add these fields at the top of the class:

```
private ObservableCollection<string> _playerNames =  
new ObservableCollection<string>();  
public List<string> SelectedPlayers { get; set; }
```

7. The previous code used `System.Collections.Generic` and the `System.Collections.ObjectModel` namespaces, so include these:

```
using System.Collections.Generic;  
using System.Collections.ObjectModel;
```

8. Add a constructor to initialize the `SelectedPlayers` collection:

```
public GameOptions()  
{  
    SelectedPlayers = new List<string>();  
}
```

9. Add a property and two methods to the class like this:

```
public ObservableCollection<string> PlayerNames  
{  
    get  
    {  
        return _playerNames;  
    }  
    set  
    {  
        _playerNames = value;  
        OnPropertyChanged("PlayerNames");  
    }  
}  
public void AddPlayer(string playerName)  
{  
    if (_playerNames.Contains(playerName))  
        return;
```

```

    _playerNames.Add(playerName);
    OnPropertyChanged("PlayerNames");
}

```

10. Return to the StartGame.xaml window.
11. Add a ListBox, two Labels, a TextBox, and a Button to the grid below the Canvas in grid row 1 and change the controls to look like those shown in Figure 14-12.




FIGURE 14-12

12. Set the Name property of the controls as shown in Table 14-9.

TABLE 14-9: The Name Property

| CONTROL | NAME               |
|---------|--------------------|
| TextBox | newPlayerTextBox   |
| Button  | addNewPlayerButton |
| ListBox | playerNamesListBox |

13. Set the ItemsSource of the ListBox like this:

```
ItemsSource="{Binding Path=PlayerNames}"
```

14. Add the ListBox's SelectionChanged event handler to the code-behind file and add this code:

```

private void playerNamesListBox_SelectionChanged(object sender,
SelectionChangedEventArgs e)
{
    if (_gameOptions.PlayAgainstComputer)
        okButton.IsEnabled = (playerNamesListBox.SelectedItems.Count == 1);
    else
        okButton.IsEnabled = (playerNamesListBox.SelectedItems.Count ==
_gameOptions.NumberOfPlayers);
}

```

15. Set the `IsEnabled` property of the OK button to `false`.
16. Open the code-behind file and add this field to the top of the class:

```
private GameOptions _gameOptions;
```
17. Copy the constructor from the `Options.xaml.cs` code-behind (though not the name) and add these lines to the end after `InitializeComponent` (Note: You will need to add using declarations for `System.IO` and `System.Xml.Serialization`):

```
if (_gameOptions.PlayAgainstComputer)
    playerNamesListBox.SelectionMode = SelectionMode.Single;
else
    playerNamesListBox.SelectionMode = SelectionMode.Extended;
```

18. Select the Add button and add the `click` event handler. Add this code:

```
private void addNewPlayerButton_Click(object sender, RoutedEventArgs e)
{
    if (!string.IsNullOrWhiteSpace(newPlayerTextBox.Text))
        _gameOptions.AddPlayer(newPlayerTextBox.Text);
    newPlayerTextBox.Text = string.Empty;
}
```

19. Copy the event handler for the OK and Cancel buttons from the `options.xaml.cs` code-behind files to this code-behind.
20. Add these lines to the top of the OK button handler:

```
foreach (string item in playerNamesListBox.SelectedItems)
{
    _gameOptions.SelectedPlayers.Add(item);
}
```

21. Go to the `App.xaml` file and change the `StartupUri` to `StartGame.xaml`.
22. Run the application.

### How It Works

You started by adding code to the `GameOptions` class that holds information about all the known players and the current selection made in the `StartGame` window.

The `ListBox`'s `ItemsSource` property is the same as you saw on the `ComboBox` earlier. But where you were able to bind the selected value of the `ComboBox` directly to a value, it is more complicated with a `ListBox`. If you try to bind the `SelectedValues` property you will find that it is read-only and therefore can't be used for data binding. The work-around used here is to use the OK button to store the values through code. Note that the cast to `IList<string>` works here because the content of the `ListBox` is strings at the moment, but if you decided to change the default behavior and display something else, then this selection of items must be changed as well.

The `ListBox`'s `SelectionChanged` event is raised whenever something happens that changes the selection. In this case you want to handle this event to check if the number of items selected is correct. If the

game is to be played against a computer, then there can only be one human player; otherwise the correct number of human players must be selected.

**NOTE** Chapter 15 discusses the Styles, Control, and Item templates and shows why you can't always know what type the content of a control is.

## EXERCISES

- 14.1 A `TextBlock` control can be used to display large amounts of text, but the control does not provide any way to scroll the text itself if the text extends beyond the viewport. By combining the `TextBlock` with another control, create a window that contains a `TextBlock` with a lot of text that can be scrolled and where the scrollbar appears only if the text extends beyond the viewport.
- 14.2 The `Slider` and `ProgressBar` controls have a few things in common, such as a minimum, maximum, and current value. Using only data binding on the `ProgressBar`, create a window with a slider and a progress bar, where the `Slider` control controls the minimum, maximum, and current value of the progress bar.
- 14.3 Change the `ProgressBar` in the previous question to display itself diagonally from the bottom-left corner to the top-right corner of the window.
- 14.4 Create a new class with the name `PersistentSlider` and three properties: `MinValue`, `MaxValue`, and `CurrentValue`. The class must be able to participate in data binding and all the properties must be able to notify bound controls of changes.
  - a. In the code-behind of the window you created in the two previous exercises, create a new field of type `PersistentSlider` and initialize it with some default values.
  - b. In the constructor, bind the instance to the windows data source.
  - c. Bind the slider's `Minimum`, `Maximum`, and `Value` properties to the data source.

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

| KEY CONCEPT                   | DESCRIPTION                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|-------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>XAML</b>                   | XAML is a language that uses XML syntax and enables controls to be added to a user interface in a declarative, hierarchical way.                                                                                                                                                                                                                                                                                                                                                          |
| <b>Data binding</b>           | You can use data binding to connect properties of controls to the value of other controls. You can also define resources and use code defined in classes outside your views as a data source for both values of properties and as content for controls. <code>DataContexts</code> can be used to specify the binding source of existing object instances and thereby allow you to bind to instances that are created in other parts of your application.                                  |
| <b>Routed events</b>          | Routed events are special events used in WPF. They come in two flavors: bubbling and tunneling. Bubbling events are first called on the control on which they are activated and then bubble up through the view tree to the root element. Tunneling events move the other way, from the root element to the control that was activated by the user. Both bubbling and tunneling can be stopped by setting the <code>Handled</code> property of the event arguments to <code>true</code> . |
| <b>INotifyPropertyChanged</b> | The <code>INotifyPropertyChanged</code> interface is implemented by a class that will be used from a WPF view. When property setters of the class are called, they raise the event <code>PropertyChanged</code> with the name of the property that changed its value. Any control property that is bound to the property that raised the event will be notified of the change and can update itself accordingly.                                                                          |
| <b>ObservableCollections</b>  | An <code>ObservableCollection</code> is a collection that, among others, implement the <code>INotifyPropertyChanged</code> interface. You use this specialized collection when you want to provide properties or values that are lists to a WPF view for data binding.                                                                                                                                                                                                                    |
| <b>Content controls</b>       | Content controls can contain a single control in their content. An example of such a control is <code>Button</code> . This control can be <code>Grid</code> or <code>StackPanel</code> ; they allow you to create complex customizations.                                                                                                                                                                                                                                                 |
| <b>Items controls</b>         | Items controls can contain a list of controls in their content. An example of such a control is the <code>ListBox</code> . Each control in the list can be customized.                                                                                                                                                                                                                                                                                                                    |

| KEY CONCEPT            | DESCRIPTION                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Layout controls</b> | <p>You learned to use a number of controls that are used to help you create the view:</p> <ol style="list-style-type: none"><li>1. <code>Canvas</code> allows for explicit positioning of controls but little else.</li><li>2. <code>StackPanel</code> stacks controls horizontally or vertically.</li><li>3. <code>WrapPanel</code> stacks controls and wraps them to the next line or column depending on the orientation of the panel.</li><li>4. <code>DockPanel</code> allows you to dock controls to the edges of the control or fill the entire content.</li><li>5. <code>Grid</code> allows you to define rows and columns and use these to position the controls.</li></ol>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| <b>UI controls</b>     | <p>UI controls display themselves on the view, often using the layout controls to guide their positions. These controls were used:</p> <ol style="list-style-type: none"><li>1. <code>Label</code> controls display short text.</li><li>2. <code>TextBlock</code> controls display text that can need multiple lines to display.</li><li>3. <code>TextBox</code> controls allow the users to provide text input.</li><li>4. <code>Button</code> controls allow the users to perform a single action.</li><li>5. <code>Image</code> controls are used to display an image.</li><li>6. <code>Checkboxes</code> let the users answer yes/no questions such as "Play Against Computer?"</li><li>7. <code>RadioButtons</code> let the users select exactly one from multiple options.</li><li>8. <code>ComboBoxes</code> display a drop-down list of items from which the user can select a single item. The control can also display a <code>TextBox</code>, letting the user enter new options.</li><li>9. <code>ListBox</code> controls display a list of items. Unlike the <code>ComboBox</code> the list is always expanded. The control allows for multiple items being selected.</li><li>10. <code>TabControl</code> allows you to group controls on pages.</li></ol> |



# 15

## Advanced Desktop Programming

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Using routed commands instead of events
- Creating menus using the Menu control and routed commands
- Working with styling controls and applications using XAML styles
- Creating value converters
- Using timelines to create animations
- Defining and referencing static and dynamic resources
- Creating user controls when the common controls are not enough

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 15 download and individually named according to the names throughout the chapter.

Until this point you have used Windows Presentation Foundation (WPF) in much the same way that you use the other major technology for creating windows applications in Visual Studio: Windows Forms. But that is about to change. WPF can style any control and use templates to change existing controls to look nothing like they do out-of-the-box. In addition to that, you are going to start working more and more by typing XAML. Although this might seem like a burden at first, the ability to move and fine-tune the display by setting properties will quickly become second nature, and you will find that there is quite a bit in XAML that cannot be done in the designer, such as creating animations.

Now it is time to continue where you left off in Chapter 14 and continue with the game client.

## THE MAIN WINDOW

The main window of the application is where the game is played, and it therefore doesn't have many controls on it. You'll construct the game in this chapter, but before you start, there are three things that you must do. You need to add the main window to the project, add menus to the window, and bind the windows you already constructed to the menu items.

### The Menu Control

Most applications include menus and toolbars of some kind. Both are a means to the same end: to provide easy navigation of the application's content. Toolbars generally contain a subset of the same entries that the menus provide and can be thought of as shortcuts to the menu items.

Visual Studio ships with both a `Menu` and a `Toolbar` control. The example here shows the use of the `Menu` control but using the `Toolbar` is very similar.

By default, the menu item appears as a horizontal bar from which you can drop down lists of items. The control is an `Items` control, so it is possible to change the default items contained in the content; however, you would normally use `MenuItem`s in some form, as shown in the following example. Each `MenuItem` can contain other menu items, and you can build complex menus by nesting `MenuItem`s within each other, but you should try to keep the menu structure as simple as possible.

You can control how the `MenuItem` displays using a number of properties (see Table 15-1).

TABLE 15-1: Displaying MenuItem Properties

| PROPERTY                 | DESCRIPTION                                                                  |
|--------------------------|------------------------------------------------------------------------------|
| <code>Icon</code>        | Displays an icon by the left edge of the control                             |
| <code>IsCheckable</code> | Displays a <code>CheckBox</code> by the left edge of the control             |
| <code>IsChecked</code>   | Gets or sets the value of a <code>CheckBox</code> on a <code>MenuItem</code> |

### Routed Commands with Menus

*Routed commands* were briefly discussed in Chapter 14, but now you are going to see them in action for the first time. Recall that these commands are akin to events in that they execute code when a user performs an action, and they can return a state indicating whether they can be executed at any given time.

There are at least three reasons why you would want to use routed commands instead of events:

1. The action that will cause an event to occur can be triggered from multiple locations in your application.
2. The UI element should be accessible only under certain conditions, such as a `Save` button being disabled if there's nothing to save.
3. You want to disconnect the code that handles the event from the code-behind file.

If any of these scenarios matches yours, consider using routed commands. In the case of the game, some of the items in the menu should also potentially be available from a toolbar. In addition, the Save action should be available only when a game is in progress and it should potentially be available from both a menu and the toolbar.

**NOTE** It is important to have the correct default namespace set in the KarliCards GUI project in order to make the examples work. If you get compiler errors stating that a class or resource isn't a member of a namespace, you probably used a different namespace than the one that is being used in the book. The KarliCards solution uses two root namespaces: Ch13CardLib for the Ch13CardLib project and KarliCards\_Gui for the KarliCards GUI project. If you experience problems, try changing the namespaces throughout the projects to match those used in the book.

## TRY IT OUT Creating the Main Window: KarliCards Gui\MainWindow.xaml

In this example you create the main window for the game. Because this window is the main window of the application, it will use the windows you have already created.

1. Add a new window to the project and name it GameClient.xaml.
2. Change the title to “Karli Cards Game Client” and remove the Height and Width properties.
3. Set the WindowState property to Maximized.
4. Add this namespace:  
`xmlns:src="clr-namespace:KarliCards_Gui"`
5. Remove the grid from the window and copy the grid and all its content from the StartGame window.
6. Delete everything inside the `<Grid>` tags except the canvas that is positioned in `Grid.Row = 0` and the `<Grid.RowDefinitions>`.
7. Drag a DockPanel control into grid row 1. Set its properties like this:  
`Grid.Row="1" Margin="0"`
8. Select DockPanel and drag a Menu control onto it. Note that the control expands to fill the entire DockPanel — this is what you want.
9. Change the menu's properties to give it a black background, bold text weight, and white foreground color:  
`Background="Black" FontWeight="Bold" Foreground="White"`
10. Right-click the menu in the design view and choose Add MenuItem.
11. Change the Header property to `_File`. Note the leading underscore. Also, if it isn't already, go ahead and set the foreground to white.

12. Add another `MenuItem` inside the `_File` item by right-clicking the `_File` item and selecting Add `MenuItem`. Set the `Height`, `Width`, `Header`, and `Foreground` properties like this:

```
<MenuItem Header="_File" Foreground="White">
    <MenuItem Header="_New Game..." Height="22"
        Width="200" Foreground="Black" />
</MenuItem>
```

13. Add the following `MenuItem`s to the File menu:

```
<MenuItem Header="_Open" Width="200" Foreground="Black"/>
<MenuItem Header="_Save" Width="200" Foreground="Black" Command="Save">
    <MenuItem.Icon>
        <Image Source="Images\base_floppydisk_32.png" Width="20" />
    </MenuItem.Icon>
</MenuItem>
<Separator Width="145" Foreground="Black"/>
<MenuItem Header="_Close" Width="200" Foreground="Black" Command="Close"/>
```

14. Add these `MenuItem`s to the menu on the same level as the `File` `MenuItem`.

```
<MenuItem Header="_Game" Background="Black" Foreground="White">
    <MenuItem Header="_Undo" HorizontalAlignment="Left"
        Width="145" Foreground="Black"/>
</MenuItem>
<MenuItem Header="_Tools" Background="Black" Foreground="White">
    <MenuItem Header="_Options" HorizontalAlignment="Left"
        Width="145" Foreground="Black"/>
</MenuItem>
<MenuItem Header="Help" Background="Black" Foreground="White">
    <MenuItem Header="_About" HorizontalAlignment="Left"
        Width="145" Foreground="Black"/>
</MenuItem>
```

15. Change the background color of the main grid control to green. You can set the background color to a standard color by clicking the color box to the right of the `Background` box and selecting Custom Expression. Then type the name of the color you would like, in this case “Green.”

16. Above the first grid control, add this command binding to the window:


```
<Window.CommandBindings>
    <CommandBinding Command="ApplicationCommands.Close"
        CanExecute="CommandCanExecute" Executed="CommandExecuted" />
    <CommandBinding Command="ApplicationCommands.Save"
        CanExecute="CommandCanExecute" Executed="CommandExecuted" />
</Window.CommandBindings>
```

17. Change the content of the label in `Grid` row 0 from “New Game” to “Karli Cards.”

18. Add a new grid to the grid in row 2 and name it `contentGrid`:

```
<Grid Grid.Row="2" x:Name="contentGrid" />
```

Your window should now look like Figure 15-1.



**FIGURE 15-1**

19. Go to the `GameClient.xaml.cs` code-behind file and add these two methods:

```
private void CommandCanExecute(object sender, CanExecuteRoutedEventArgs e)
{
    if (e.Command == ApplicationCommands.Close)
        e.CanExecute = true;
    if (e.Command == ApplicationCommands.Save)
        e.CanExecute = false;
    e.Handled = true;
}
private void CommandExecuted(object sender, ExecutedRoutedEventArgs e)
{
    if (e.Command == ApplicationCommands.Close)
        this.Close();
    e.Handled = true;
}
```

20. Change the `StartupUri` in the `App.xaml` file to `GameClient.xaml` and run the application.

### How It Works

When you run the application you will notice that the Game Client window is initially displayed as maximized, but you can resize the window as you like. When you hold down the Alt key, the File menu gets focus and the F in File is underlined, indicating that you can expand the menu by pressing F.

When you expand the menu you can see that the Save menu is disabled, but it displays a disk icon as well as the text “Ctrl-S” to the right of the element title. This means that you can access it by pressing Ctrl-S (when it is enabled). You might wonder why this is displayed, as you haven’t set any shortcut keys anywhere. However, you did set a command for the menu item:

```
<MenuItem Header="_Save" Width="200" Foreground="Black" Command="Save">
```

The Save command is defined by WPF. Save and Close, which are used in the File menu, are defined in the `ApplicationCommands` class, which also defines Cut, Copy, Paste, and Print. When you specify the Save command for a `MenuItem`, the shortcut key Ctrl-S is assigned to the menu item because it's the standard key combination used to access that function in most Windows applications.

In the code-behind file, you added two methods used to determine the state and action taken by the commands. In the XAML, you created two command bindings that used the methods like this:

```
<Window.CommandBindings>
    <CommandBinding Command="ApplicationCommands.Close"
        CanExecute="CommandCanExecute" Executed="CommandExecuted" />
    <CommandBinding Command="ApplicationCommands.Save"
        CanExecute="CommandCanExecute" Executed="CommandExecuted" />
</Window.CommandBindings>
private void CommandCanExecute(object sender, CanExecuteRoutedEventArgs e)
{
    if (e.Command == ApplicationCommands.Close)
        e.CanExecute = true;
    if (e.Command == ApplicationCommands.Save)
        e.CanExecute = false;
    e.Handled = true;
}
private void CommandExecuted(object sender, ExecutedRoutedEventArgs e)
{
    if (e.Command == ApplicationCommands.Close)
        this.Close();
    e.Handled = true;
}
```

The `CanExecute` part of the command binding specifies a method that is called to determine whether the command should be available to the user at the moment. The `Executed` part specifies a method that should be called when the user activates the command. Note that it doesn't matter from where the command is activated. If a menu item and a button both include the Save command, the binding works for both.

The current implementation of `CommandCanExecute` is too simple for real life, where you would do some calculation to determine whether the application is ready to save anything. Since you don't have a game to save yet, just returning `false` for the Save command is appropriate. You do this by setting the `e.CanExecute` property on the `CanExecuteRoutedEventArgs` class. The Close command, on the other hand, can be executed just fine, so you return `true` for that one.

`CommandExecuted` performs the same test as `CommandCanExecute`. If it determines that the command to execute is the Close command, then it closes the current window.

---

## CREATING AND STYLING CONTROLS

It's time to step away from the client implementation of the game and start looking more at the game itself. One key feature of a graphical card game is...the cards. Obviously, you are not going to find a "Playing Card" control in the standard controls that ship with WPF, so you have to create it yourself.

One of the best features of WPF is the complete control it provides designers over the look and feel of user interfaces. Central to this is the capability to style controls however you want, in two or three dimensions. Until now, you have been using the basic styling for controls that is supplied with .NET, but the actual possibilities are endless.

This section describes two basic techniques:

- **Styles** — Sets of properties that are applied to a control as a batch
- **Templates** — The controls that are used to build the display for a control

There is some overlap here, as styles can contain templates.

## Styles

WPF controls have a property called `Style` (inherited from `FrameworkElement`) that can be set to an instance of the `Style` class. The `Style` class is quite complex and is capable of advanced styling functionality, but at its heart it is essentially a set of `Setter` objects. Each `Setter` object is responsible for setting the value of a property according to its `Property` property (the name of the property to set) and its `Value` property (the value to set the property to). You can either fully qualify the name you use in `Property` to the control type (for example, `Button.Foreground`), or you can set the `TargetType` property of the `Style` object (for example, `Button`), so that it is capable of resolving property names.

The following code shows how to use a `Style` object to set the `Foreground` property of a `Button` control:

```
<Button>
    Click me!
    <Button.Style>
        <Style TargetType="Button">
            <Setter Property="Foreground">
                <Setter.Value>
                    <SolidColorBrush Color="Purple" />
                </Setter.Value>
            </Setter>
        </Style>
    </Button.Style>
</Button>
```

Obviously, in this case it would be far easier simply to set the `Foreground` property of the button in the usual way. Styles become much more useful when you turn them into resources, because resources can be reused. You will learn how to do this in the “WPF User Controls” section later in the chapter.

## Templates

Controls are constructed using templates, which you can customize. A template consists of a hierarchy of controls used to build the display of a control, which may include a content presenter for controls such as buttons that display content.

The template of a control is stored in its `Template` property, which is an instance of the `ControlTemplate` class. The `ControlTemplate` class includes a `TargetType` property that you can set to the type of control for which you are defining a template, and it can contain a single control. This control can be a container such as `Grid`, so this doesn't exactly limit what you can do.

Typically, you set the template for a class by using a style. This simply involves providing controls to use for the `Template` property in the following way:

```
<Button>
    Click me!
    <Button.Style>
        <Style TargetType="Button">
            <Setter Property="Template">
                <Setter.Value>
                    <ControlTemplate TargetType="Button">
                        ...
                    </ControlTemplate>
                </Setter.Value>
            </Setter>
        </Style>
    </Button.Style>
</Button>
```

Some controls may require more than one template. For example, `CheckBox` controls use one template for a check box (`CheckBox.Template`) and one template to output text next to the check box (`CheckBox.ContentTemplate`).

Templates that require content presenters can include a `ContentPresenter` control at the location where you want to output content. Some controls — especially those that output collections of items — use alternative techniques, which aren't covered in this chapter.

Again, replacing templates is most useful when combined with resources. However, as control styling is a very common technique, it is worth looking at how to do it in a Try It Out.

## TRY IT OUT Using Styles and Templates: ControlStyling\MainWindow.xaml

1. Create a new WPF application called `ControlStyling`.
2. Modify the code in `MainWindow.xaml` as follows:

```
<Grid Background="Black">
    <Button Margin="20" Click="Button_Click">
        Would anyone use a button like this?
        <Button.Style>
            <Style TargetType="Button">
                <Setter Property="FontSize" Value="18" />
                <Setter Property="FontFamily" Value="arial" />
                <Setter Property="FontWeight" Value="bold" />
                <Setter Property="Foreground">
                    <Setter.Value>
                        <LinearGradientBrush StartPoint="0.5,0" EndPoint="0.5,1">
                            <LinearGradientBrush.GradientStops>
                                <GradientStop Offset="0.0" Color="Purple" />
```

```
<GradientStop Offset="0.5" Color="Azure" />
    <GradientStop Offset="1.0" Color="Purple" />
</LinearGradientBrush.GradientStops>
</LinearGradientBrush>
</Setter.Value>
</Setter>
<Setter Property="Template">
    <Setter.Value>
        <ControlTemplate TargetType="Button">
            <Grid>
                <Grid.ColumnDefinitions>
                    <ColumnDefinition Width="50" />
                    <ColumnDefinition />
                    <ColumnDefinition Width="50" />
                </Grid.ColumnDefinitions>
                <Grid.RowDefinitions>
                    <RowDefinition MinHeight="50" />
                </Grid.RowDefinitions>
                <Ellipse Grid.Column="0" Height="50">
                    <Ellipse.Fill>
                        <RadialGradientBrush>
                            <RadialGradientBrush.GradientStops>
                                <GradientStop Offset="0.0" Color="Yellow" />
                                <GradientStop Offset="1.0" Color="Red" />
                            </RadialGradientBrush.GradientStops>
                        </RadialGradientBrush>
                    </Ellipse.Fill>
                </Ellipse>
                <Grid Grid.Column="1">
                    <Rectangle RadiusX="10" RadiusY="10">
                        <Rectangle.Fill>
                            <RadialGradientBrush>
                                <RadialGradientBrush.GradientStops>
                                    <GradientStop Offset="0.0" Color="Yellow" />
                                    <GradientStop Offset="1.0" Color="Red" />
                                </RadialGradientBrush.GradientStops>
                            </RadialGradientBrush>
                        </Rectangle.Fill>
                    </Rectangle>
                    <ContentPresenter Margin="20,0,20,0"
                        HorizontalAlignment="Center"
                        VerticalAlignment="Center" />
                </Grid>
                <Ellipse Grid.Column="2" Height="50">
                    <Ellipse.Fill>
                        <RadialGradientBrush>
                            <RadialGradientBrush.GradientStops>
                                <GradientStop Offset="0.0" Color="Yellow" />
                                <GradientStop Offset="1.0" Color="Red" />
                            </RadialGradientBrush.GradientStops>
                        </RadialGradientBrush>
                    </Ellipse.Fill>
                </Ellipse>
            </Grid>
        </ControlTemplate>
```

```

        </Setter.Value>
    </Setter>
</Style>
</Button.Style>
</Button>
</Grid>

```

3. Modify the code in `MainWindow.xaml.cs` as follows:

```

public partial class MainWindow : Window
{
    ...
    private void Button_Click(object sender, RoutedEventArgs e)
    {
        MessageBox.Show("Button clicked.");
    }
}

```

4. Run the application and click the button once. Figure 15-2 shows the result.



FIGURE 15-2

## How It Works

First, let me apologize for the truly nasty-looking button shown in this example. However, aesthetic considerations aside, this example does show that you can completely change how a button looks in WPF without a lot of effort. Note that changing the button template does not change the functionality of the button. That is, you can still click on the button and respond to that click in an event handler.

You probably noticed that certain things you associate with Windows buttons aren't implemented in the template used here. In particular, there is no visual feedback when you roll over the button or when you click it. This button also looks exactly the same whether it has focus or not. To achieve these missing effects, you need to learn about *triggers*, which are the subject of the next section.

Before doing that, though, consider the example code in a little more detail, focusing on styles and templates and looking at how the template was created.

The example starts with ordinary code that you would use to display a `Button` control:

```

<Button Margin="20" Click="Button_Click">
    Would anyone use a button like this?

```

This provides basic properties and content for the button. The `Style` property is set to a `Style` object, which begins by setting three simple font properties of the `Button` control:

```

<Button.Style>
<Style TargetType="Button">

```

```

<Setter Property="FontSize" Value="18" />
<Setter Property="FontFamily" Value="arial" />
<Setter Property="FontWeight" Value="bold" />

```

Next, the `Button.Foreground` property is set using property element syntax because a brush is used:

```

<Setter Property="Foreground">
    <Setter.Value>
        <LinearGradientBrush StartPoint="0.5,0" EndPoint="0.5,1">
            <LinearGradientBrush.GradientStops>
                <GradientStop Offset="0.0" Color="Purple" />
                <GradientStop Offset="0.5" Color="Azure" />
                <GradientStop Offset="1.0" Color="Purple" />
            </LinearGradientBrush.GradientStops>
        </LinearGradientBrush>
    </Setter.Value>
</Setter>

```

The remainder of the code for the `Style` object sets the `Button.Template` property to a `ControlTemplate` object:

```

<Setter Property="Template">
    <Setter.Value>
        <ControlTemplate TargetType="Button">
            ...
        </ControlTemplate>
    </Setter.Value>
</Setter>
</Style>
</Button.Style>
</Button>

```

The template code can be summarized as a `Grid` control that contains three cells in a single row. In turn, these cells contain an `Ellipse`, a `Rectangle`, the `ContentPresenter` for the template, and another `Ellipse`:

```

<Grid>
    <Ellipse Grid.Column="0" Height="50">
        ...
    </Ellipse>
    <Grid Grid.Column="1">
        <Rectangle RadiusX="10" RadiusY="10">
            ...
        </Rectangle>
        <ContentPresenter Margin="20,0,20,0"
            HorizontalAlignment="Center"
            VerticalAlignment="Center" />
    </Grid>
    <Ellipse Grid.Column="2" Height="50">
        ...
    </Ellipse>
</Grid>

```

## Value Converters

You may have wondered at some of the assignments that you have used in the examples so far: How, for example, can you assign the string value “true” to a Boolean property? You have learned that C# is type-safe and the compiler should not allow that kind of thing to happen! Happily, the reality is that it doesn’t. XAML and WPF make extensive use of something called *value converters*, which can convert from one type to another behind the scenes.

WPF ships with converters for just about all the standard scenarios you can think of, so you can always convert from `int` to `string` or `bool` and `integer`. But what happens when you want to convert something that is not included? Then you have to implement the converter yourself.

Let’s look at one example that is very common: the inverted `bool` converter. Imagine that you have a check box on a dialog box. Depending on whether it’s checked, another part of the dialog box will be disabled or enabled. Quite often, the answer must be reversed for this to make sense. Take a look at the Options dialog box, for example. It has a check box with the question, “Play against computer?” Selecting this option should disable the `ComboBox` and `TextBoxes` on the dialog box. The value of `.IsChecked` would be `true`, so binding that to `.IsEnabled` of the other two controls will not work. Enter the `InversedBoolConverter`. This converter will simply inverse the `bool` value.

### The `IValueConverter` Interface

In order to create a `ValueConverter`, you must implement the `IValueConverter` interface. This interface has two methods: `Convert` and `ConvertBack`. These might seem self-explanatory, but they are actually a bit complicated.

```
object Convert(object value, Type targetType,
    object parameter, CultureInfo culture);
object ConvertBack(object value, Type targetType,
    object parameter, CultureInfo culture);
```

You use the `Convert` method when converting to a target type and use the `ConvertBack` method for the reverse operation. The `value` parameter denotes the value to convert and the `targetType` is the type it should be converted to. The `parameter` can be used to set a helper. Exercise 15.1 at the end of this chapter requires you to use this parameter to create a specific value converter.

### `ValueConversionAttribute`

In addition to implementing the interface, you can set an attribute on the class that implements the converter. This is not needed, but it is a great help both for tools and users of your converter. The `ValueConversionAttribute` takes two parameters, both of which are `Type` objects. This means that you explicitly set the types that the converter will convert to and from.

#### TRY IT OUT Create a Value Converter: KarliCards Gui

This example builds on the `KarliCards Gui` project you created earlier.

1. Create a new class and name it `InversedBoolConverter`.
2. Modify the class like this. You must also include the `System.Windows.Data` namespace:

```
[ValueConversion(typeof(bool), typeof(bool))]
```

```

public class InverseBoolConverter : IValueConverter
{
    public object Convert(object value, Type targetType, object parameter,
        System.Globalization.CultureInfo culture)
    {
        return !(bool)value;
    }
    public object ConvertBack(object value, Type targetType, object parameter,
        System.Globalization.CultureInfo culture)
    {
        return !(bool)value;
    }
}

```

3. Go to the Options.xaml file and create a new static resource for the window:

```

<Window.Resources>
    <local:InverseBoolConverter x:Key="inverseBool" />
</Window.Resources>

```

4. Set the `.IsEnabled` property of the combo box to this binding:

```

.IsEnabled="{Binding ElementName=playAgainstComputerCheck,
    Path=IsChecked, Converter={StaticResource inverseBool}}"

```

### *How It Works*

The conversion is obviously very simple in this case — it simply returns `false` if the value is true and vice-versa. You are going to see more complicated converters later in the chapter.

In the XAML, you create a resource for the converter to be able to reference it from the binding that needs it. In the binding, you set the `Converter` property to use the converter.

## Triggers

Events in WPF can include all manner of things, including button clicks, application startup and shutdown events, and so on. There are, in fact, several types of triggers in WPF, all of which inherit from a base `TriggerBase` class. One such trigger is the `EventTrigger` class, which contains a collection of actions, each of which is an object that derives from the base `TriggerAction` class. These actions are executed when the trigger is activated.

Not a lot of classes inherit from `TriggerAction` in WPF, but you can, of course, define your own. You can use `EventTrigger` to trigger animations using the `BeginStoryboard` action, manipulate storyboards using `ControllableStoryboardAction`, and trigger sound effects with `SoundPlayerAction`. As this latter trigger is mostly used in animations, you'll look at it in the next section.

Every control has a `Triggers` property that you can use to define triggers directly on that control. You can also define triggers further up the hierarchy — for example, on a `Window` object as shown earlier. The type of trigger you will use most often when you are styling controls is `Trigger` (although you will still use `EventTrigger` to trigger control animations). The `Trigger` class is used

to set properties in response to changes to other properties, and is particularly useful when used in `Style` objects.

Trigger objects are configured as follows:

- To define what property a `Trigger` object monitors, you use the `Trigger.Property` property.
- To define when the `Trigger` object activates, you set the `Trigger.Value` property.
- To define the actions taken by a `Trigger`, you set the `Trigger.Setters` property to a collection of `Setter` objects.

The `Setter` objects referred to here are exactly the same objects that you saw in the “Styles” section earlier.

For example, the following trigger examines the value of a property called `MyBooleanValue`, and when that property is `true` it sets the value of the `Opacity` property to 0.5:

```
<Trigger Property="MyBooleanValue" Value="true">
    <Setter Property="Opacity" Value="0.5" />
</Trigger>
```

On its own, this code doesn’t tell you very much, as it is not associated with any control or style. The following code is much more explanatory; it shows a `Trigger` as you would use it in a `Style` object:

```
<Style TargetType="Button">
    <Style.Triggers>
        <Trigger Property="IsMouseOver" Value="true">
            <Setter Property="Foreground" Value="Yellow" />
        </Trigger>
    </Style.Triggers>
</Style>
```

This code changes the `Foreground` property of a `Button` control to `Yellow` when the `Button` `.IsMouseOver` property is `true`. `IsMouseOver` is one of several extremely useful properties that you can use as a shortcut to find out information about controls and control state. As its name suggests, it is `true` if the mouse is over the control. This enables you to code for mouse rollovers. Other properties like this include `IsFocused`, to determine whether a control has focus; `IsHitTestVisible`, which indicates whether it is possible to click on a control (that is, it is not obscured by controls further up the stacking order); and `IsPressed`, which indicates whether a button is pressed. The last of these only applies to buttons that inherit from `ButtonBase`, whereas the others are available on all controls.

You can also achieve a great deal by using the `ControlTemplate.Triggers` property, which enables you to create templates for controls that include triggers. This is how the default `Button` template is able to respond to mouse rollovers, clicks, and focus changes with its template. This is also what you must modify to implement this functionality for yourself.

## Animations

Animations are created by using storyboards. The absolute best way to define animations is, without a doubt, to use a designer such as Expression Blend. However, you can also define them by editing XAML code directly, and by implication from code-behind (as XAML is simply a way to build a WPF object model).

A storyboard is defined using a `Storyboard` object, which contains one or more timelines. You can define timelines by using key frames or by using one of several simpler objects that encapsulate entire animations. Complex storyboards may even contain nested storyboards.

A `Storyboard` is contained in a resource dictionary, so you must identify it with an `x:Key` property.

Within the timeline of a storyboard, you can animate properties of any element in your application that is of type `double`, `Point`, or `Color`. This covers most of the things that you may want to change, so it's quite flexible. There are some things that you can't do, such as replace one brush with another, but there are ways to achieve pretty much any effect you can imagine given these three types.

Each of these three types has two associated timeline controls that you can use as children of `Storyboard`. These six controls are `DoubleAnimation`, `DoubleAnimationUsingKeyFrames`, `PointAnimation`, `PointAnimationUsingKeyFrames`, `ColorAnimation`, and `ColorAnimationUsingKeyFrames`. Every timeline control can be associated with a specific property of a specific control by using the attached properties `Storyboard.TargetName` and `Storyboard.TargetProperty`. For example, you would set these properties to `MyRectangle` and `Width` if you wanted to animate the `Width` property of a `Rectangle` control with a `Name` property of `MyRectangle`. You would use either `DoubleAnimation` or `DoubleAnimationUsingKeyFrames` to animate this property. You will see examples of using storyboards as this chapter progresses.

Next, you'll look at the simple, animation timelines without key frames, and then move on to look at the timelines that use key frames.

## Timelines without Key Frames

The timelines without key frames are `DoubleAnimation`, `PointAnimation`, and `ColorAnimation`. These timelines have identical property names, although the types of these properties vary according to the type of the timeline (note that all duration properties are specified in the form [days.] hours:minutes:seconds in XAML code). Table 15-2 describes these properties.

TABLE 15-2: The Timeline Properties

PROPERTY	DESCRIPTION
<code>Name</code>	The name of the timeline, so that you can refer to it from other places.
<code>BeginTime</code>	How long after the storyboard is triggered before the timeline starts.
<code>Duration</code>	How long the timeline lasts.

*continues*

**TABLE 15-2** (continued)

PROPERTY	DESCRIPTION
AutoReverse	Whether the timeline reverses when it completes and returns properties to their original values. This property is a Boolean value.
RepeatBehavior	Set this to a specified duration to make the timeline repeat as indicated — an integer followed by x (for example, 5x) to repeat the timeline a set number of times; or use Forever to make the timeline repeat until the storyboard is paused or stopped.
FillBehavior	How the timeline behaves if it completes while the storyboard is still continuing. You can use HoldEnd to leave properties at the values they are at when the timeline completes (the default), or Stop to return them to their original values.
SpeedRatio	Controls the speed of the animation relative to the values specified in other properties. The default value is 1, but you can change it from other code to speed up or slow down animations.
From	The initial value to set the property to at the start of the animation. You can omit this value to use the current value of the property.
To	The final value for the property at the end of the animation. You can omit this value to use the current value of the property.
By	Use this value to animate from the current value of a property to the sum of the current value and the value you specify. You can use this property on its own or in combination with From.

For example, the following timeline will animate the `Width` property of a `Rectangle` control with a `Name` property of `MyRectangle` between 100 and 200 over five seconds:

```
<Storyboard x:Key="RectangleExpander">
<DoubleAnimation Storyboard.TargetName="MyRectangle"
    Storyboard.TargetProperty="Width" Duration="00:00:05"
    From="100" To="200" />
</Storyboard>
```

## Timelines with Key Frames

The timelines with key frames are `DoubleAnimationUsingKeyFrames`, `PointAnimationUsingKeyFrames`, and `ColorAnimationUsingKeyFrames`. These timeline classes use the same properties as the timeline classes in the previous section, except that they don't have `From`, `To`, or `By` properties. Instead, they have a `KeyFrames` property that is a collection of key frame objects.

These timelines can contain any number of key frames, each of which can cause the value being animated to behave in a different way. There are three types of key frames for each type of timeline:

- **Discrete** — A discrete key frame causes the value being animated to jump to a specified value with no transition.
- **Linear** — A linear key frame causes the value being animated to animate to a specified value in a linear transition.
- **Spline** — A spline key frame causes the value being animated to animate to a specified value in a nonlinear transition defined by a cubic Bezier curve function.

There are therefore nine types of key frame objects: `DiscreteDoubleKeyFrame`, `LinearDoubleKeyFrame`, `SplineDoubleKeyFrame`, `DiscreteColorKeyFrame`, `LinearColorKeyFrame`, `SplineColorKeyFrame`, `DiscretePointKeyFrame`, `LinearPointKeyFrame`, and `SplinePointKeyFrame`.

The key frame classes have the same three properties as the timeline classes examined in the previous section. The four spline key frame classes add one additional property: `KeySpline` (see Table 15-3).

**TABLE 15-3:** Properties of the Spline Key Frame Classes

PROPERTY	USAGE
Name	The name of the key frame, so that you can refer to it from other places.
KeyTime	The location of the key frame expressed as an amount of time after the timeline starts.
Value	The value that the property will reach or be set to when the key frame is reached.
KeySpline	Two sets of two numbers in the form <code>cp1x, cp1y, cp2x, cp2y</code> that define the cubic Bezier function to use to animate the property. ( <i>Spline key frames only</i> .)

For example, you could animate the position of an `Ellipse` in a square by animating its `Center` property, which is of type `Point`, as follows:

```
<Storyboard x:Key="EllipseMover">
    <PointAnimationUsingKeyFrames Storyboard.TargetName="MyEllipse"
        Storyboard.TargetProperty="Center" RepeatBehavior="Forever">
        <LinearPointKeyFrame KeyTime="00:00:00" Value="50,50" />
        <LinearPointKeyFrame KeyTime="00:00:01" Value="100,50" />
        <LinearPointKeyFrame KeyTime="00:00:02" Value="100,100" />
        <LinearPointKeyFrame KeyTime="00:00:03" Value="50,100" />
        <LinearPointKeyFrame KeyTime="00:00:04" Value="50,50" />
    </PointAnimationUsingKeyFrames>
</Storyboard>
```

Point values are specified in `x,y` form in XAML code.

## WPF USER CONTROLS

WPF provides a set of controls that are useful in many situations. However, as with all the .NET development frameworks, it also enables you to extend this functionality. Specifically, you can create your own controls by deriving your classes from classes in the WPF class hierarchy.

One of the most useful controls you can derive from is `UserControl`. This class gives you all the basic functionality that you are likely to require from a WPF control, and it enables your control to snap in beside the existing WPF control suite seamlessly. Everything you might hope to achieve with a WPF control — such as animation, styling, and templating — can be achieved with user controls.

You can add user controls to your project by using the Project  $\Rightarrow$  Add User Control menu item. This gives you a blank canvas (well, actually a blank `Grid`) to work from. User controls are defined using the top-level `UserControl` element in XAML, and the class in the code-behind derives from the `System.Windows.Controls.UserControl` class.

Once you have added a user control to your project, you can add controls to lay out the control and code-behind to configure the control. When you have finished doing that, you can use it throughout your application, and even reuse it in other applications.

One of the crucial things you need to know when creating user controls is how to implement dependency properties. Chapter 14 briefly discussed dependency properties, and now that you are getting closer to writing your own controls, it is time to take a closer look at them.

## Implementing Dependency Properties

You can add dependency properties to any class that inherits from `System.Windows.DependencyObject`. This class is in the inheritance hierarchy for many classes in WPF, including all the controls and `UserControl`.

To implement a dependency property to a class, you add a public, static member to your class definition of type `System.Windows.DependencyProperty`. The name of this member is up to you, but best practice is to follow the naming convention `<PropertyName>Property`:

```
public static DependencyProperty MyStringProperty;
```

It might seem odd that this property is defined as static, as you end up with a property that can be uniquely defined for each instance of your class. The WPF property framework keeps track of things for you, so you don't have to worry about this for the moment.

The member you add must be configured by using the static `DependencyProperty.Register()` method:

```
public static DependencyProperty MyStringProperty =
    DependencyProperty.Register(...);
```

This method takes between three and five parameters, as shown in the Table 15-4 (these are shown in order, with the first three parameters being the mandatory ones).

**TABLE 15-4:** The Register( ) Method's Parameters

PARAMETER	USAGE
string name	The name of the property
Type propertyType	The type of the property
Type ownerType	The type of the class containing the property
PropertyMetadata typeMetadata	Additional property settings: the default value of the property and callback methods to use for property change notifications and coercion
ValidateValueCallback validateValueCallback	The callback method to use to validate property values

There are other methods that you can use to register dependency properties, such as `RegisterAttached()`, which you can use to implement an attached property. You won't look at these other methods in this chapter, but it's worth reading up on them.

For example, you could register the `MyStringProperty` dependency property using three parameters as follows:

```
public class MyClass : DependencyObject
{
    public static DependencyProperty MyStringProperty = DependencyProperty.Register(
        "MyString",
        typeof(string),
        typeof(MyClass));
}
```

You can also include a .NET property that can be used to access dependency properties directly (although this isn't mandatory, as you will see shortly). However, because dependency properties are defined as static members, you cannot use the same syntax you would use with ordinary properties. To access the value of a dependency property, you have to use methods that are inherited from `DependencyObject`, as follows:

```
public string MyString
{
    get { return (string)GetValue(MyStringProperty); }
    set { SetValue(MyStringProperty, value); }
}
```

Here, the `GetValue()` and `SetValue()` methods get and set, respectively, the value of the `MyStringProperty`, dependency property for the current instance. These two methods are public, so client code can use them directly to manipulate dependency property values. This is why adding a .NET property to access a dependency property is not mandatory.

If you want to set metadata for a property, then you must use an object that derives from `PropertyMetadata`, such as `FrameworkPropertyMetadata`, and pass this instance as the fourth parameter to `Register()`. There are 11 overloads of the `FrameworkPropertyMetadata` constructor, and they take one or more of the parameters shown in Table 15-5.

TABLE 15-5: Overloads for the `FrameworkPropertyMetadata` Constructor

PARAMETER TYPE	USAGE
<code>object defaultValue</code>	The default value for the property.
<code>FrameworkPropertyMetadataOptions flags</code>	A combination of the flags (from the <code>FrameworkPropertyMetadataOptions</code> enum) that you can use to specify additional metadata for a property. For example, you might use <code>AffectsArrange</code> to declare that changes to the property might affect control layout. This would cause the layout engine for a window to recalculate control layout if the property changed. See the MSDN documentation for a full list of the options available here.
<code>PropertyChangedCallback propertyChangedCallback</code>	The callback method to use when the property value changes.
<code>CoerceValueCallback coerceValueCallback</code>	The callback method to use if the property value is coerced.
<code>bool isAnimationProhibited</code>	Specifies whether this property can be changed by an animation.
<code>UpdateSourceTrigger defaultUpdateSourceTrigger</code>	When property values are data-bound, this property determines when the data source is updated, according to values in the <code>UpdateSourceTrigger</code> enum. The default value is <code>PropertyChanged</code> , which means that the binding source is updated as soon as the property changes. This is not always appropriate — for example, the <code>TextBox</code> <code>.Text</code> property uses a value of <code>LostFocus</code> for this property. This ensures that the binding source is not updated prematurely. You can also use the value <code>Explicit</code> to specify that the binding source should be updated only when requested (by calling the <code>UpdateSource()</code> method of a class derived from <code>DependencyObject</code> ).

A simple example of using `FrameworkPropertyMetadata` is to use it to set the default value of a property:

```

public static DependencyProperty MyStringProperty =
    DependencyProperty.Register(
        "MyString",
        typeof(string),
        typeof(MyClass),
        new FrameworkPropertyMetadata("Default value"));

```

You have so far learned about three callback methods that you can specify, for property change notification, property coercion, and property value validation. These callbacks, like the dependency property itself, must all be implemented as public, static methods. Each callback has a specific return type and parameter list that you must use on your callback method.

Now it is time to get back on track and continue with the game client for Karli Cards. In the following Try It Out, you create a user control that can represent a playing card in the application.

## TRY IT OUT User Controls: KarliCards Gui.CardControl.xaml

Return to the KarliCards Gui project from the previous Try It Out.

1. This example uses the CardLib project that you created in Chapter 13, so you have to add this to the solution. Begin by right-clicking the solution name in the Solution Explorer and choosing Add  $\Rightarrow$  Existing Project. Browse to and select the Ch13CardLib.csproj file from the Chapter 13 code examples.
2. Add a reference to the Ch13CardLib project by right-clicking References and choosing Add Reference in the KarliCards Gui project. Click Projects  $\Rightarrow$  Solution from the tree on the left and select Ch13CardLib. Click OK.
3. Add a new user control called CardControl to the KarliCards Gui project, and modify the code in CardControl.xaml as follows:

```

<UserControl x:Class="KarliCards_Gui.CardControl"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:local="clr-namespace:KarliCards_Gui"
    mc:Ignorable="d"
    Height="154" Width="100" Name ="UserControl">
    <UserControl.Resources>
        <local:RankNameConverter x:Key="rankConverter"/>
        <DataTemplate x:Key="SuitTemplate">
            <TextBlock Text="{Binding}" />
        </DataTemplate>
        <Style TargetType="Image" x:Key="SuitImage">
            <Style.Triggers>
                <DataTrigger Binding="{Binding ElementName=UserControl, Path=Suit}"
Value="Club">
                    <Setter Property="Source" Value="Images\Clubs.png" />
                </DataTrigger>
                <DataTrigger Binding="{Binding ElementName=UserControl, Path=Suit}"
Value="Heart">
                    <Setter Property="Source" Value="Images\Hearts.png" />
                </DataTrigger>
            </Style.Triggers>
        </Style>
    </UserControl.Resources>

```

```
        </DataTrigger>
        <DataTrigger Binding="{Binding ElementName= UserControl, Path= Suit}" Value="Diamond">
            <Setter Property="Source" Value="Images\Diamonds.png" />
        </DataTrigger>
        <DataTrigger Binding="{Binding ElementName= UserControl, Path= Suit}" Value="Spade">
            <Setter Property="Source" Value="Images\Spades.png" />
        </DataTrigger>
    </Style.Triggers>
</Style>
</UserControl.Resources>
<Grid>
    <Rectangle Stroke="{x:Null}" RadiusX="12.5" RadiusY="12.5">
        <Rectangle.Fill>
            <LinearGradientBrush EndPoint="0.47,-0.167" StartPoint="0.86,0.92">
                <GradientStop Color="#FFD1C78F" Offset="0"/>
                <GradientStop Color="#FFFFFF" Offset="1"/>
            </LinearGradientBrush>
        </Rectangle.Fill>
        <Rectangle.Effect>
            <DropShadowEffect Direction="145" BlurRadius="10" ShadowDepth="0" />
        </Rectangle.Effect>
    </Rectangle>
    <Label x:Name="SuitLabel" Content="{Binding Path= Suit, ElementName= UserControl, Mode= Default}" ContentTemplate="{DynamicResource SuitTemplate}" HorizontalAlignment="Center" VerticalAlignment="Center" Margin="8,51,8,60" />
    <Label x:Name="RankLabel" Grid.ZIndex="1" Content="{Binding Path= Rank, ElementName= UserControl, Mode= Default, Converter={StaticResource ResourceKey= rankConverter}}" ContentTemplate="{DynamicResource SuitTemplate}" HorizontalAlignment="Left" VerticalAlignment="Top" Margin="8,8,0,0" />
    <Label x:Name="RankLabelInverted" Content="{Binding Path= Rank, ElementName= UserControl, Mode= Default, Converter={StaticResource ResourceKey= rankConverter}}" ContentTemplate="{DynamicResource SuitTemplate}" HorizontalAlignment="Right" VerticalAlignment="Bottom" Margin="0,0,8,8" RenderTransformOrigin="0.5,0.5">
        <Label.RenderTransform>
            <RotateTransform Angle="180" />
        </Label.RenderTransform>
    </Label>
    <Image Name="TopRightImage" Style="{StaticResource ResourceKey= SuitImage}" Margin="12,12,8,0" HorizontalAlignment="Right" VerticalAlignment="Top" Width="18.5" Height="18.5" Stretch="UniformToFill" />
    <Image Name="BottomLeftImage" Style="{StaticResource ResourceKey= SuitImage}" Margin="12,0,8,12" HorizontalAlignment="Left" VerticalAlignment="Bottom" Width="18.5" Height="18.5" Stretch="UniformToFill" RenderTransformOrigin="0.5,0.5">
        <Image.RenderTransform>
            <RotateTransform Angle="180" />
        </Image.RenderTransform>
    </Image>

```

```

</Image>
<Path Fill="#FFFFFF" Stretch="Fill" Stroke="{x:Null}"
Margin="0,0,35.218,-0.077" Data="F1 M110.5,51 L145.16457,51 C116.5986,
76.731148 115.63518,132.69684 121.63533,149.34013 133.45299,
182.12018 152.15821,195.69803 161.79765,200.07669 L110.5,200 C103.59644,
200 98,194.40356 98,187.5 L98,63.5 C98,56.596439 103.59644,51 110.5,51 z">
<Path.OpacityMask>
<LinearGradientBrush EndPoint="0.957,1.127" StartPoint="0,-0.06">
<GradientStop Color="#FF000000" Offset="0"/>
<GradientStop Color="#00FFFFFF" Offset="1"/>
</LinearGradientBrush>
</Path.OpacityMask>
</Path>
</Grid>
</UserControl>

```

**4.** Add three dependency properties to the class:

```

public static DependencyProperty SuitProperty = DependencyProperty.Register(
    "Suit",
    typeof(Ch13CardLib.Suit),
    typeof(CardControl),
    new PropertyMetadata(Ch13CardLib.Suit.Club,
new PropertyChangedCallback(OnSuitChanged)));
public static DependencyProperty RankProperty = DependencyProperty.Register(
    "Rank",
    typeof(Ch13CardLib.Rank),
    typeof(CardControl),
    new PropertyMetadata(Ch13CardLib.Rank.Ace));
public static DependencyProperty IsFaceUpProperty = DependencyProperty.Register(
    "IsFaceUp",
    typeof(bool),
    typeof(CardControl),
    new PropertyMetadata(true, new PropertyChangedCallback(OnIsFaceUpChanged)));
public bool IsFaceUp
{
    get { return (bool)GetValue(IsFaceUpProperty); }
    set { SetValue(IsFaceUpProperty, value); }
}
public Ch13CardLib.Suit Suit
{
    get { return (Ch13CardLib.Suit)GetValue(SuitProperty); }
    set { SetValue(SuitProperty, value); }
}
public Ch13CardLib.Rank Rank
{
    get { return (Ch13CardLib.Rank)GetValue(RankProperty); }
    set { SetValue(RankProperty, value); }
}

```

**5.** Add the change event handlers to the class:

```

public static void OnSuitChanged(DependencyObject source,
    DependencyPropertyChangedEventArgs args)
{

```

```
        var control = source as CardControl;
        control.SetTextColor();
    }
    private static void OnIsFaceUpChanged(DependencyObject source,
   DependencyPropertyChangedEventArgs args)
    {
        var control = source as CardControl;
        control.RankLabel.Visibility = control.SuitLabel.Visibility =
            control.RankLabelInverted.Visibility =
control.TopRightImage.Visibility =
control.BottomLeftImage.Visibility = control.IsFaceUp ?
Visibility.Visible : Visibility.Hidden;
    }
}
```

6. Add a property to the class:

```
private Ch13CardLib.Card _card;
public Ch13CardLib.Card Card
{
    get { return _card; }
    private set { _card = value; Suit = _card.suit; Rank = _card.rank; }
}
```

7. Add a helper method to set the text colors and overload the constructor to take a Card:

```
public CardControl(Ch13CardLib.Card card)
{
    InitializeComponent();
    Card = card;
}
private void SetTextColor()
{
    var color = (Suit == Ch13CardLib.Suit.Club || Suit == Ch13CardLib.Suit.Spade) ?
        new SolidColorBrush(Color.FromArgb(0, 0, 0)) :
        new SolidColorBrush(Color.FromArgb(255, 0, 0));
    RankLabel.Foreground = SuitLabel.Foreground = RankLabelInverted.Foreground =
        color;
}
```

8. Add a new value converter by adding a new class to the project. Name it `RankNameConverter.cs` and add this code:

```
using System;
using System.Windows;
using System.Windows.Data;
namespace KarliCards_Gui
{
    [ValueConversion(typeof(Ch13CardLib.Rank), typeof(string))]
    public class RankNameConverter : IValueConverter
    {
        public object Convert(object value, Type targetType,
object parameter, System.Globalization.CultureInfo culture)
    {
```

```
int source = (int)value;
if (source == 1 || source > 10)
{
    switch (source)
    {
        case 1:
            return "Ace";
        case 11:
            return "Jack";
        case 12:
            return "Queen";
        case 13:
            return "King";
        default:
            return DependencyProperty.UnsetValue;
    }
}
else
    return source.ToString();
}

public object ConvertBack(object value, Type targetType,
object parameter, System.Globalization.CultureInfo culture)
{
    return DependencyProperty.UnsetValue;
}
}
```

9. Go to the GameClient.xaml.cs code-behind file and change the constructor like this:

```
public GameClient()
{
    InitializeComponent();
    var position = new Point(15, 15);
    for (var i = 0; i < 4; i++)
    {
        var suit = (Ch13CardLib.Suit)i;
        position.Y = 15;
        for (int rank = 1; rank < 14; rank++)
        {
            position.Y += 30;
            var card = new CardControl(new Ch13CardLib.Card((Ch13CardLib.Suit)suit,
(Ch13CardLib.Rank)rank));
            card.VerticalAlignment = VerticalAlignment.Top;
            card.HorizontalAlignment = HorizontalAlignment.Left;
            card.Margin = new Thickness(position.X, position.Y, 0, 0);
            contentGrid.Children.Add(card);
        }
        position.X += 112;
    }
}
```

10. Run the application. The result is shown in Figure 15-3.



FIGURE 15-3

## How It Works

This example creates a user control with two dependent properties, and includes client code to use the control. This example covers plenty of ground, and the place to start looking at the code is with the `Card` control.

The `Card` control consists mostly of code that will be familiar to you from code you've seen earlier in this chapter. The layout code uses nothing new, although you might agree that the result is a bit prettier than the lurid button in the previous two examples.

The code in `Card` exposes three dependency properties, `Suit`, `Rank`, and `IsFaceUp`, to client code, and binds these properties to visual elements in the control layout. As a result, when you set `Suit` to `Club`, the word `Club` is displayed in the center of the card and the `Club` image is displayed in the top-right and bottom-left corners of the card. Similarly, the value of `Rank` is displayed in the other two corners of the card.

You'll look at the implementation of these properties in a moment. For now it is enough to know that they are enumerations originating from the CardLib project that you started in Chapter 10.

The three labels display the rank and suit of the card. Even though they are bound to different properties, they have a few things in common. They must display some text in red or black depending on the values of the bound properties. In this example, the color is set using the events raised when the Rank changes, but you can use triggers for this:

```
<Label x:Name="SuitLabel"
    Content="{Binding Path=Suit, ElementName=UserControl, Mode=Default}"
    ContentTemplate="{DynamicResource SuitTemplate}" HorizontalAlignment="Center"
    VerticalAlignment="Center" Margin="8,51,8,60" />
```

When you bind property values, you can also specify how to render the bound content, by using a data template. In this example, the *data template* is *SuitTemplate*, referenced as a dynamic resource (although in this case a static resource binding would also work fine). This template is defined in the user control resources section as follows:

```
<UserControl.Resources>
    <DataTemplate x:Key="SuitTemplate">
        <TextBlock Text="{Binding}" />
    </DataTemplate>
</UserControl.Resources>
```

The string value of *Suit* is therefore used as the *Text* property of a *TextBlock* control. This same *DataTemplate* definition is reused for the two rank labels. *Suit* is an enumeration, and the name of the value in the enumeration is automatically converted to a string to be displayed in the *Text* property.

The two Rank labels include a value converter in the binding.

```
<Label x:Name="RankLabel" Grid.ZIndex="1"
    Content="{Binding Path=Rank, ElementName=UserControl, Mode=Default,
    Converter={StaticResource ResourceKey=rankConverter}}"
    ContentTemplate="{DynamicResource SuitTemplate}"
    HorizontalAlignment="Left" VerticalAlignment="Top"
    Margin="8,8,0,0" />
```

The converter is included in the *UserControl* resources through this declaration:

```
<local:RankNameConverter x:Key="rankConverter" />
```

You will not break the control if you remove the value converter. Instead, you will see Ace, 2, 3, 4, and so on. You will also see the names of the enumeration values converted to string — Ace, Deuce, Three, Four, and so on. Although this is technically correct, it doesn't look quite right, so you convert the values to a combination of numbers and strings.

The final point to notice is the *Grid.ZIndex="1"* property assignment on the *RankLabel*. The *ZIndex* of a control on a *Grid* or *Canvas* determines the visual layer that holds the control. If two or more controls occupy the same space, then you can use the *zIndex* to force one of them to go to the front. Normally all controls have a *zIndex* of zero, so setting a single control to 1 means that it is moved to the front. This is necessary because the blur of the path would otherwise obscure the text.

For this data binding to work, you must define three dependency properties using techniques you learned in the previous section. These are defined in the code-behind for the user control as follows (they have simple .NET property wrappers, which there is no need to show here because of the simplicity of the code):

```
public static DependencyProperty SuitProperty = DependencyProperty.Register(
    "Suit",
    typeof(CardLib.Suit),
    typeof(CardControl),
    new PropertyMetadata(CardLib.Suit.Club,
new PropertyChangedCallback(OnSuitChanged)));
public static DependencyProperty RankProperty = DependencyProperty.Register(
    "Rank",
    typeof(CardLib.Rank),
    typeof(CardControl),
    new PropertyMetadata(CardLib.Rank.Ace));
public static DependencyProperty IsFaceUpProperty = DependencyProperty.Register(
    "IsFaceUp",typeof(bool),
    typeof(CardControl),
    new PropertyMetadata(true, new PropertyChangedCallback(OnIsFaceUpChanged)));
```

The dependency properties use a callback method to validate their values, and the `Suit` and `IsFaceUp` properties also have a callback method for when their values change.

When the value of `Suit` changes, the `OnSuitChanged()` callback method is called. This method is responsible for setting the text color to red (for hearts and diamonds) or black (for clubs and spades). It does this by calling a utility method on the source of the method call. This is necessary because the callback method is implemented as a static method, but it is passed the instance of the user control that raised the event as a parameter so that it can interact with it. The method called is `SetTextColor()`:

```
public static void OnSuitChanged(DependencyObject source,
    DependencyPropertyChangedEventArgs args)
{
    var control = source as CardControl;
    control.SetTextColor();
}
```

The `SetTextColor()` method is private but is obviously still accessible from `OnSuitChanged()`, as they are both members of the same class, despite being instance and static methods, respectively. `SetTextColor()` simply sets the `Foreground` property of the various labels of the control to a solid color brush that is either black or red, depending on the `Suit` value.

When `IsFaceUp` changes, the control displays or hides the images and labels that are used to display the current value of the control.

This is all you need to look at in the `Card` control. The code in the `GameClient.xaml.cs` code-behind file is included to display the cards and is only temporary. It generates one card for each of the 13 possible values and displays each suit in a column.

---

## PUTTING IT ALL TOGETHER

At this point in the development of the game, you have two independent dialog boxes, a card library, and a main window that provides a blank space for the game to be displayed on. That still leaves quite a lot of work, but with the foundation built, it's time to start building the game. The classes in the CardLib describe the game "domain model," that is, the objects that a game can be broken down into, which need to be refactored a bit to make it work better with a Windows application. Next you are going to write the game's "View Model," which is a class that is able to control the display of the game. Then you will create two additional user controls that use the Card user control to display the game visually. Finally, you will bind it all together in the game client.

**NOTE** *The term "View Model" comes from a much used design pattern in WPF: Model - View - ViewModel (MVVM). This design pattern describes how to separate code from the view and link it together. Although this book doesn't attempt to conform to this pattern, this example uses a lot of the elements from the pattern, such as separating the ViewModel from the views. In this context, the domain model described next is the "model" part of the MVVM name, and the Windows you have been creating are the views.*

## Refactoring the Domain Model

As stated, the domain model is the code that describes the objects of the game. At the moment, you have these classes in the CardLib project that describe objects of the game:

- Card
- Deck
- Rank
- Suit

In addition to these classes, the game needs a Player and a ComputerPlayer class, so you are going to add those. You also need to modify the Card and Deck classes a bit to make them work better in a Windows application.

There is a lot of work to do, so let's get started.

**NOTE** *This example does not use the CardClient class from the earlier chapters because the differences between console and Windows applications are so great that very little code can be reused.*

**TRY IT OUT** Finishing the Domain Model: KarliCards Gui

This example continues where the previous example left off.

1. Each player in the game can be in a number of “states” during the game. You can model this in a `PlayerState` enumeration. Go to the Ch13CardLib project and create a new `PlayerState` enumeration for the project. You can simply create a new class and replace code like this:

```
[Serializable]
public enum PlayerState
{
    Inactive,
    Active,
    MustDiscard,
    Winner,
    Loser
}
```

2. Next, you raise a few events when something happens on a player. For that, you need some custom event arguments, so add another class named `PlayerEventArgs`:

```
public class PlayerEventArgs : EventArgs
{
    public Player Player { get; set; }
    public PlayerState State { get; set; }
}
```

3. You also need to raise events when something happens to a card, so go ahead and create another class called `CardEventArgs`:

```
public class CardEventArgs : EventArgs
{
    public Card Card { get; set; }
}
```

4. The enumeration `ComputerSkillLevel` currently exists in the `GameOptions.cs` class (in the Karli Cards Gui project). Go ahead and cut it from there and move it to its own file in the Ch13CardLib project. This changes its namespace to `Ch13CardLib`, so you have to add the `Ch13CardLib` namespace to the `GameOptions.xaml.cs` and `Options.Xaml.cs` files:

```
using Ch13CardLib;
```

5. The `Deck` class should be changed. Rather than going back to this class multiple times over the course of this chapter, the following listing is the complete class.

```
using System;
using System.Collections.Generic;
using System.Linq;

namespace Ch13CardLib
{
    public delegate void LastCardDrawnHandler(Deck currentDeck);

    public class Deck : ICloneable
```

```
{  
    public event LastCardDrawnHandler LastCardDrawn;  
    private Cards cards = new Cards();  
  
    public Deck()  
    {  
        InsertAllCards();  
    }  
  
    protected Deck(Cards newCards)  
    {  
        cards = newCards;  
    }  
  
    public int CardsInDeck  
    {  
        get { return cards.Count; }  
    }  
  
    public Card GetCard(int cardNum)  
    {  
        if (cardNum >= 0 && cardNum <= 51)  
        {  
            if ((cardNum == 51) && (LastCardDrawn != null)) LastCardDrawn(this);  
            return cards[cardNum];  
        }  
        else  
            throw new CardOutOfRangeException(cards.Clone() as Cards);  
    }  
  
    public void Shuffle()  
    {  
        Cards newDeck = new Cards();  
        bool[] assigned = new bool[cards.Count];  
        Random sourceGen = new Random();  
        for (int i = 0; i < cards.Count; i++)  
        {  
            int sourceCard = 0;  
            bool foundCard = false;  
            while (foundCard == false)  
            {  
                sourceCard = sourceGen.Next(cards.Count);  
                if (assigned[sourceCard] == false)  
                    foundCard = true;  
            }  
            assigned[sourceCard] = true;  
            newDeck.Add(cards[sourceCard]);  
        }  
        newDeck.CopyTo(cards);  
    }  
  
    public void ReshuffleDiscarded(List<Card> cardsInPlay)  
    {  
        InsertAllCards(cardsInPlay);  
        Shuffle();  
    }  
}
```

```
public Card Draw()
{
    if (cards.Count == 0) return null;

    var card = cards[0];
    cards.RemoveAt(0);
    return card;
}

public Card SelectCardOfSpecificSuit(Suit suit)
{
    Card selectedCard = cards.FirstOrDefault(card => card?.suit == suit);
    if (selectedCard == null) return Draw();
    cards.Remove(selectedCard);
    return selectedCard;
}

public object Clone()
{
    Deck newDeck = new Deck(cards.Clone() as Cards);
    return newDeck;
}

private void InsertAllCards()
{
    for (int suitVal = 0; suitVal < 4; suitVal++)
    {
        for (int rankVal = 1; rankVal < 14; rankVal++)
        {
            cards.Add(new Card((Suit)suitVal, (Rank)rankVal));
        }
    }
}

private void InsertAllCards(List<Card> except)
{
    for (int suitVal = 0; suitVal < 4; suitVal++)
    {
        for (int rankVal = 1; rankVal < 14; rankVal++)
        {
            var card = new Card((Suit)suitVal, (Rank)rankVal);
            if (except?.Contains(card))
                continue;
            cards.Add(card);
        }
    }
}
```

6. There will be two types of players in the game: a Player, which is controlled by a real person; and a ComputerPlayer, which is controlled by the game. These two classes are omitted from the code listings here, but you can find them in the code download for the chapter as `Player.cs` and `ComputerPlayer.cs`. Go ahead and add them to the `CardLib` project.

## How It Works

That was a lot of code and a lot of changes! However, when you run the application, nothing seems to have changed, but a lot of plumbing has been put in to make the game work.

The Deck class has been extended with a few new methods. Whenever the deck is emptied, the discarded cards should be put back in play. In order to do this, an overload of the `InsertAllCards` method that takes a list of the cards that are in play has been added. The property `CardsInDeck` will be used to tell how many cards are left in the deck. If the players draw every card in the deck, you want to shuffle all the discarded cards back into the deck, and so the `Shuffle` method now allows the deck to contain fewer than 52 cards and the `ReshuffleDiscarded` method allows you to perform the reshuffle. `Draw` and `SelectCardOfSpecificSuit` are both used to draw a card. Most of the code in the `Player` and `ComputerPlayer` classes that you added to the project from the downloaded code is pretty easy to understand. The `Player` class can draw and discard cards. This is shared with the `ComputerPlayer`, but the computer is also equipped with the ability to decide which cards to draw and discard without user interaction. The `ComputerPlayer` class can also cheat:

```
public void PerformDraw(Deck deck, Card availableCard)
{
    switch (Skill)
    {
        case ComputerSkillLevel.Dumb:
            DrawCard(deck);
            break;
        default:
            DrawBestCard(deck, availableCard, (Skill == ComputerSkillLevel.Cheats));
            break;
    }
}
public void PerformDiscard(Deck deck)
{
    switch (Skill)
    {
        case ComputerSkillLevel.Dumb:
            int discardIndex = _random.Next(Hand.Count);
            DiscardCard(Hand[discardIndex]);
            break;
        default:
            DiscardWorstCard();
            break;
    }
}
private void DrawBestCard(Deck deck, Card availableCard, bool cheat = false)
{
    var bestSuit = CalculateBestSuit();
    if (availableCard.suit == bestSuit)
        AddCard(availableCard);
    else if (cheat == false)
        DrawCard(deck);
    else
        AddCard(deck.SelectCardOfSpecificSuit(bestSuit));
}
```

Cheating is assisted by a deck that allows the computer to select a card of a specific suit. If you allow the computer to cheat, you are going to have a hard time winning any games!

You will also notice that the `Player` class implements the `IPropertyChanged` interface and the properties `PlayerName` and `State` use this to notify any observers of changes. Particularly, the `State` property is important later as changes to this property will drive the game forward.

---

## The View Models

The purpose of a view model is to hold the state of the view that displays it. In the case of the Karli Cards, this means that you already have a view model class: the `GameOptions` class. This class holds the state of the `Options` and `StartGame` windows. At the moment, you can't get the selected players from the options, so you have to add that ability. The view model of the Game Client window is missing, so that is the next task to do.

The view model for the execution of the game must reflect all the parts of the game as it is running. The parts of the game are:

- The deck from which the current player draw a card
- A card that can be taken by the current player instead of drawing a card
- A current player
- A number of participating players

The view model should also be able to notify observers of changes, and that means implementing `IPropertyChanged` again.

In addition to these abilities, the view model should also provide a way of starting a new game. You will do this by creating a new routed command for the menu. The command is created in the view model, but is called from the view.

### TRY IT OUT The View Model: KarliCards Gui

This example continues with the KarliCards Gui project.

1. Add the following namespaces to the `GameOptions` class using statements:

```
using System.Windows.Input;
using System.IO;
using System.Xml.Serialization;
```

2. Add a new command to the `GameOptions` class:

```
public static RoutedCommand OptionsCommand = new RoutedCommand("Show Options",
typeof(GameOptions), new InputGestureCollection(new List<InputGesture>
{ new KeyGesture(Key.O, ModifierKeys.Control) }));
```

3. Add two new methods to the class:

```
public void Save()
{
    using (var stream = File.Open("GameOptions.xml", FileMode.Create))
    {
        var serializer = new XmlSerializer(typeof(GameOptions));
        serializer.Serialize(stream, this);
    }
}
public static GameOptions Create()
{
    if (File.Exists("GameOptions.xml"))
    {
        using (var stream = File.OpenRead("GameOptions.xml"))
        {
            var serializer = new XmlSerializer(typeof(GameOptions));
            return serializer.Deserialize(stream) as GameOptions;
        }
    }
    else
        return new GameOptions();
}
```

4. Change the OK click event handler of the Options.xaml.cs code-behind file like this:

```
private void okButton_Click(object sender, RoutedEventArgs e)
{
    this.DialogResult = true;
    _gameOptions.Save();
    this.Close();
}
```

5. Delete everything except the InitializeComponent call from the constructor and hook the DataContextChanged event like this:

```
public Options()
{
    _gameOptions = GameOptions.Create();
    DataContext = _gameOptions;
    InitializeComponent();
```

6. Open the StartGame.xaml.cs code-behind file and select the last four lines of the code in the constructor. Extract a new method called ChangeListBoxOptions by right-clicking the selected code and selecting Quick Actions ⇔ ExtractMethod:

```
private void ChangeListBoxOptions()
{
    if (_gameOptions.PlayAgainstComputer)
        playersListBox.SelectionMode = SelectionMode.Single;
    else
        playersListBox.SelectionMode = SelectionMode.Extended;
}
```

7. Delete everything except the `InitializeComponent` call from the constructor and hook the `DataContextChanged` event like this:

```
public StartGame()
{
    InitializeComponent();
    DataContextChanged += StartGame_DataContextChanged;
}
```

8. Add the `StartGame_DataContextChanged` event handler:

```
void StartGame_DataContextChanged(object sender,
DependencyPropertyChangedEventArgs e)
{
    _gameOptions = DataContext as GameOptions;
    ChangeListBoxOptions();
}
```

9. Change the OK click event handler like this:

```
private void okButton_Click(object sender, RoutedEventArgs e)
{
    var gameOptions = DataContext as GameOptions;
    gameOptions.SelectedPlayers = new List<string>();
    foreach (string item in playerNamesListBox.SelectedItems)
    {
        gameOptions.SelectedPlayers.Add(item);
    }
    this.DialogResult = true;
    this.Close();
}
```

10. Create a new class and name it `GameViewModel`. Start by implementing the `INotifyPropertyChanged` interface:

```
using Ch13CardLib;
using System.Collections.Generic;
using System.ComponentModel;
using System.Linq;
using System.Windows.Input;
namespace KarliCards_Gui
{
    public class GameViewModel : INotifyPropertyChanged
    {
        public event PropertyChangedEventHandler PropertyChanged;
        private void OnPropertyChanged(string propertyName) =>
        PropertyChanged?.Invoke(this, new PropertyChangedEventArgs(propertyName));
    }
}
```

11. Add a property to hold the current player. This property should use the `OnPropertyChanged` event:

```
private Player _currentPlayer;
public Player CurrentPlayer
```

```

    {
        get { return _currentPlayer; }
        set
        {
            _currentPlayer = value;
            OnPropertyChanged(nameof(CurrentPlayer));
        }
    }
}

```

12. Add four more properties and their related fields to the class, just as you did with the `CurrentPlayer` property. The property names and field names are shown in Table 15-6.

TABLE 15-6: Property and Field Names

TYPE	PROPERTY NAME	FIELD NAME
<code>List&lt;Player&gt;</code>	<code>Players</code>	<code>_players</code>
<code>Card</code>	<code>CurrentAvailableCard</code>	<code>_availableCard</code>
<code>Deck</code>	<code>GameDeck</code>	<code>_deck</code>
<code>bool</code>	<code>GameStarted</code>	<code>_gameStarted</code>

13. Add this private field to hold the game options:

```
private GameOptions _gameOptions;
```

14. Add two Routed commands:

```

public static RoutedCommand StartGameCommand =
new RoutedCommand("Start New Game", typeof(GameViewModel),
new InputGestureCollection(new List<InputGesture>
{ new KeyGesture(Key.N, ModifierKeys.Control) }));
public static RoutedCommand ShowAboutCommand =
new RoutedCommand("Show About Dialog", typeof(GameViewModel));

```

15. Add a new default constructor:

```

public GameViewModel()
{
    _players = new List<Player>();
    _gameOptions = GameOptions.Create();
}

```

16. When a game is started, the players and deck must be initialized. Add this code to the class:

```

public void StartNewGame()
{
    if (_gameOptions.SelectedPlayers.Count < 1 ||
(_gameOptions.SelectedPlayers.Count == 1
&& !_gameOptions.PlayAgainstComputer))
        return;
    CreateGameDeck();
}

```

```
        CreatePlayers();
        InitializeGame();
        GameStarted = true;
    }
    private void InitializeGame()
    {
        AssignCurrentPlayer(0);
        CurrentAvailableCard = GameDeck.Draw();
    }
    private void AssignCurrentPlayer(int index)
    {
        CurrentPlayer = Players[index];
        if (!Players.Any(x => x.State == PlayerState.Winner))
            Players.ForEach(x => x.State = (x == Players[index] ? PlayerState.Active :
PlayerState.Inactive));
    }
    private void InitializePlayer(Player player)
    {
        player.DrawNewHand(GameDeck);
        player.OnCardDiscarded += player_OnCardDiscarded;
        player.OnPlayerHasWon += player_OnPlayerHasWon;
        Players.Add(player);
    }
    private void CreateGameDeck()
    {
        GameDeck = new Deck();
        GameDeck.Shuffle();
    }
    private void CreatePlayers()
    {
        Players.Clear();
        for (var i = 0; i < _gameOptions.NumberOfPlayers; i++)
        {
            if (i < _gameOptions.SelectedPlayers.Count)
                InitializePlayer(new Player { Index = i, PlayerName =
_gameOptions.SelectedPlayers[i] });
            else
                InitializePlayer(new ComputerPlayer { Index = i, Skill =
_gameOptions.ComputerSkill });
        }
    }
```

17. Finally, add the two event handlers for the events generated by the players:

```
void player_OnPlayerHasWon(object sender, PlayerEventArgs e)
{
    Players.ForEach(x => x.State = (x == e.Player ? PlayerState.Winner :
PlayerState.Loser));
}
void player_OnCardDiscarded(object sender, CardEventArgs e)
{
    CurrentAvailableCard = e.Card;
    var nextIndex = CurrentPlayer.Index + 1 >= _gameOptions.NumberOfPlayers ? 0 :
CurrentPlayer.Index + 1;
    if (GameDeck.CardsInDeck == 0)
```

```

    {
        var cardsInPlay = new List<Card>();
        foreach (var player in Players)
            cardsInPlay.AddRange(player.GetCards());
        cardsInPlay.Add(CurrentAvailableCard);
        GameDeck.ReshuffleDiscarded(cardsInPlay);
    }
    AssignCurrentPlayer(nextIndex);
}

```

18. Go to the GameClient.xaml file and add a new namespace to the Window declaration:

```
xmlns:vm="clr-namespace:KarliCards_Gui.ViewModel"
```

19. Below the Window declaration, add a DataContext declaration:

```
<Window.DataContext >
    <local:GameViewModel />
</Window.DataContext>
```

20. Add three command bindings to the CommandBindings declarations:

```

<CommandBinding Command="local:GameViewModel.StartGameCommand"
    CanExecute="CommandCanExecute" Executed="CommandExecuted" />
<CommandBinding Command="local:GameViewModel.ShowAboutCommand"
    CanExecute="CommandCanExecute" Executed="CommandExecuted" />
    <CommandBinding Command="local:GameOptions.OptionsCommand"
    CanExecute="CommandCanExecute" Executed="CommandExecuted" />

```

21. Add a command to the New Game menu item like this:

```
<MenuItem Header="_New Game..." Foreground="Black" Width="200"
    Command="local:GameViewModel.StartGameCommand" />
```

22. Add a command to the Options menu item and set the width to 200 like this:

```
<MenuItem Header="_Options" HorizontalAlignment="Left" Width="200"
    Foreground="Black" Command="local:GameOptions.OptionsCommand"/>
```

23. Add a command to the About menu item like this:

```
<MenuItem Header="_About" HorizontalAlignment="Left" Width="145"
    Foreground="Black" Command="local:GameViewModel.ShowAboutCommand" />
```

24. Go to the code-behind file and change the CommandCanExecute and CommandExecuted methods like this:

```

private void CommandCanExecute(object sender, CanExecuteRoutedEventArgs e)
{
    if (e.Command == ApplicationCommands.Close)
        e.CanExecute = true;
    if (e.Command == ApplicationCommands.Save)
        e.CanExecute = false;
    if (e.Command == GameViewModel.StartGameCommand)
        e.CanExecute = true;
}

```

```
        if (e.Command == GameOptions.OptionsCommand)
            e.CanExecute = true;
        if (e.Command == GameViewModel.ShowAboutCommand)
            e.CanExecute = true;
        e.Handled = true;
    }

    private void CommandExecuted(object sender, ExecutedRoutedEventArgs e)
    {
        if (e.Command == ApplicationCommands.Close)
            this.Close();
        if (e.Command == GameViewModel.StartGameCommand)
        {
            var model = new GameViewModel();
            StartGame startGameDialog = new StartGame();
            var options = GameOptions.Create();
            startGameDialog.DataContext = options;
            var result = startGameDialog.ShowDialog();
            if (result.HasValue && result.Value == true)
            {
                options.Save();
                model.StartNewGame();
                DataContext = model;
            }
        }
        if (e.Command == GameOptions.OptionsCommand)
        {
            var dialog = new Options();
            var result = dialog.ShowDialog();
            if (result.HasValue && result.Value == true)
                DataContext = new GameViewModel(); // Clear current game
        }
        if (e.Command == GameViewModel.ShowAboutCommand)
        {
            var dialog = new About();
            dialog.ShowDialog();
        }
        e.Handled = true;
    }
```

25. Delete everything from the constructor except the `InitializeComponent()` call.

### How It Works

Once again you have done a lot of work with very little to show for it when you run the application, but the menus have changed. The Options and New Game menu items have been given shortcut keys and can now be accessed using Ctrl-O and Ctrl-N. This is displayed when you drop down the menus. This has happened because you created two new commands for the menu. You did this in `GameOptions.cs` and `GameViewModel.cs`, respectively:

```
public static RoutedCommand OptionsCommand = new RoutedCommand("Show Options",
    typeof(GameOptions), new InputGestureCollection(new List<InputGesture>
    { new KeyGesture(Key.O, ModifierKeys.Control) }));
public static RoutedCommand StartGameCommand =
```

---

```

new RoutedCommand("Start New Game", typeof(GameViewModel),
new InputGestureCollection(new List<InputGesture>
{ new KeyGesture(Key.N, ModifierKeys.Control) }));

```

When you assign a list of `InputGestures` to the command, the shortcuts are automatically associated with the menus.

In the code-behind for the game client, you also added code to display the two windows as dialog boxes.

```

if (e.Command == GameViewModel.StartGameCommand)
{
    var model = new GameViewModel();
    StartGame startGameDialog = new StartGame();
    startGameDialog.DataContext = model.GameOptions;
    var result = startGameDialog.ShowDialog();
    if (result.HasValue && result.Value == true)
    {
        model.GameOptions.Save();
        model.StartNewGame();
        DataContext = model;
    }
}

```

By showing the windows as dialog boxes, you can return a value that indicates whether the result of the dialog box should be used. You can't return a value directly from the window; instead, you set the window's `DialogResult` property to either true or false to indicate success or failure:

```

private void okButton_Click(object sender, RoutedEventArgs e)
{
    this.DialogResult = true;
    this.Close();
}

```

In Chapter 14 you were told that if you want to set the `DataContext` to an existing object instance, you had to do so from code. This happens in the previous code, but the XAML in `GameClient.xaml` also instantiates a new instance when the application starts:

```

<Window.DataContext >
<vm:GameViewModel />
</Window.DataContext>

```

This instance ensures that there is a `DataContext` for the view, but it isn't used for much before it is exchanged for a new instance in the `StartGame` command.

The `GameViewModel` contains a lot of code but much of it is just properties and instantiation of the players and the Deck instances.

Once the game has started, the state of the players and `GameViewModel` drive the game forward as the computer or the players make choices. The `PlayerHasWon` event is handled in `GameViewModel` and ensures that the state of the other players changes to `Loser`.

```

void player_OnPlayerHasWon(object sender, PlayerEventArgs e)
{
}

```

```
    Players.ForEach(x => x.State = (x == e.Player ? PlayerState.Winner :  
PlayerState.Loser));  
}
```

The other event you created for the player is also handled here: `CardDiscarded` is used to indicate that a player has completed her turn. This causes the `CurrentPlayer` to be set to the next available player:

```
void player_OnCardDiscarded(object sender, CardEventArgs e)  
{  
    CurrentAvailableCard = e.Card;  
    var nextIndex = CurrentPlayer.Index + 1 >= _gameOptions.NumberOfPlayers ? 0 :  
        CurrentPlayer.Index + 1;  
    if (GameDeck.CardsInDeck == 0)  
    {  
        var cardsInPlay = new List<Card>();  
        foreach (var player in Players)  
            cardsInPlay.AddRange(player.GetCards());  
        cardsInPlay.Add(CurrentAvailableCard);  
        GameDeck.ReshuffleDiscarded(cardsInPlay);  
    }  
    AssignCurrentPlayer(nextIndex);  
}
```

This event handler also checks whether there are any more cards in the deck. If there are no more cards, the event handler collects a list of cards that are currently used in the game and makes the deck generate a new, shuffled deck containing only cards that have been discarded.

The `StartGame` method is called from the `CommandExecuted` method in the `GameClient.xaml.cs` code-behind file. This method uses three methods to create a new deck, to create and deal cards to the players, and finally to set the `CurrentPlayer` to start the game.

---

## Completing the Game

You now have a complete game that you can't play because nothing is being displayed in the game client. For the game to run, you need two additional user controls that will be positioned on the game client using a dock panel.

The two user controls are called `CardsInHand`, which displays a player's hand, and `GameDecks`, which displays the main deck and the available card.

**NOTE** You can add dependency properties by typing `propdp` and pressing Tab in the editor.

## TRY IT OUT Completing the Game: KarliCards Gui

Once again, this example continues with the KarliCards Gui project you have been working on.

1. Create a new user control in the GameClient project by right-clicking the project and selecting Add ➔ User Control. Name it CardsInHandControl.
2. Add a Label and a Canvas control to the Grid like this:

```
<Grid>
    <Label Name="PlayerNameLabel" Foreground="White" FontWeight="Bold"
FontSize="14" >
        <Label.Effect>
            <DropShadowEffect ShadowDepth="5" Opacity="0.5" Direction="145" />
        </Label.Effect>
    </Label>
    <Canvas Name="CardSurface">
    </Canvas>
</Grid>
```

3. Go to the code-behind file and use these using directives:

```
using Ch13CardLib;
using System;
using System.Threading;
using System.Windows;
using System.Windows.Controls;
using System.Windows.Input;
using System.Windows.Media;
using System.Windows.Threading;
```

4. There are four dependency properties. Type propdp and press the Tab key to insert the property template. Insert the Type, Name, OwnerClass, and default value. Use tab to switch from one value to the next. Set the values as shown in Table 15-7. Press the Return key after you finish editing the values to complete the template (see Table 15-7).

**TABLE 15-7:** Cards in Hand Dependency Properties

TYPE	NAME	OWNERCLASS	DEFAULT VALUE
Player	Owner	CardsInHandControl	null
GameViewModel	Game	CardsInHandControl	null
PlayerState	PlayerState	CardsInHandControl	PlayerState.Inactive
Orientation	PlayerOrientation	CardsInHandControl	Orientation.Horizontal

5. The `Owner` property requires a callback that should be called whenever the property changes. You can specify this as the second parameter of the constructor of the `PropertyMetadata` class that is used as the fourth parameter of the `register()` method. Change the registration like this:

```
public static readonly DependencyProperty OwnerProperty =
    DependencyProperty.Register(
        "Owner",
        typeof(Player),
        typeof(CardsInHandControl),
        new PropertyMetadata(null, new PropertyChangedCallback(OnOwnerChanged)));
```

6. Like the `Owner` property, the `PlayerState` and `PlayerOrientation` properties should also register a callback. Repeat Step 4 for these two properties using the names `OnPlayerStateChanged` and `OnPlayerOrientationChanged` for the callback methods.

7. Add the callback methods:

```
private static void OnOwnerChanged(DependencyObject source,
DependencyPropertyChangedEventArgs e)
{
    var control = source as CardsInHandControl;
    control.RedrawCards();
}
private static void OnPlayerStateChanged(DependencyObject source,
DependencyPropertyChangedEventArgs e)
{
    var control = source as CardsInHandControl;
    var computerPlayer = control.Owner as ComputerPlayer;
    if (computerPlayer != null)
    {
        if (computerPlayer.State == PlayerState.MustDiscard)
        {
            Thread delayedWorker = new Thread(control.DelayDiscard);
            delayedWorker.Start(new Payload { Deck = control.Game.GameDeck,
AvailableCard = control.Game.CurrentAvailableCard, Player = computerPlayer });
        }
        else if (computerPlayer.State == PlayerState.Active)
        {
            Thread delayedWorker = new Thread(control.DelayDraw);
            delayedWorker.Start(new Payload { Deck = control.Game.GameDeck,
AvailableCard = control.Game.CurrentAvailableCard, Player = computerPlayer });
        }
        control.RedrawCards();
    }
}
private static void OnPlayerOrientationChanged(DependencyObject source,
DependencyPropertyChangedEventArgs args)
{
    var control = source as CardsInHandControl;
    control.RedrawCards();
}
```

8. The callbacks require a number of helper methods. Start by adding the private class and two methods that are used by the `delayedWorker` threads in the `OnPlayerStateChanged` method:

```
private class Payload
{
    public Deck Deck { get; set; }
    public Card AvailableCard { get; set; }
    public ComputerPlayer Player { get; set; }
}
private void DelayDraw(object payload)
{
    Thread.Sleep(1250);
    var data = payload as Payload;
    Dispatcher.Invoke(DispatcherPriority.Normal,
new Action<Deck, Card>(data.Player.PerformDraw), data.Deck, data.AvailableCard);
}
private void DelayDiscard(object payload)
{
    Thread.Sleep(1250);
    var data = payload as Payload;
    Dispatcher.Invoke(DispatcherPriority.Normal,
new Action<Deck>(data.Player.PerformDiscard), data.Deck);
}
```

9. Add the methods used to draw the control:

```
private void RedrawCards()
{
    CardSurface.Children.Clear();
    if (Owner == null)
    {
        PlayerNameLabel.Content = string.Empty;
        return;
    }
    DrawPlayerName();
    DrawCards();
}
private void DrawCards()
{
    bool isFaceup = (Owner.State != PlayerState.Inactive);
    if (Owner is ComputerPlayer)
        isFaceup = (Owner.State == PlayerState.Loser ||
Owner.State == PlayerState.Winner);
    var cards = Owner.GetCards();
    if (cards == null || cards.Count == 0)
        return;
    for (var i = 0; i < cards.Count; i++)
    {
        var cardControl = new CardControl(cards[i]);
        if (PlayerOrientation == Orientation.Horizontal)
            cardControl.Margin = new Thickness(i * 35, 35, 0, 0);
        else
```

```
        cardControl.Margin = new Thickness(5, 35 + i * 30, 0, 0);
        cardControl.MouseDoubleClick += cardControl_MouseDoubleClick;
        cardControl.IsFaceUp = isFaceup;
        CardSurface.Children.Add(cardControl);
    }
}
private void DrawPlayerName()
{
    if (Owner.State == PlayerState.Winner || Owner.State == PlayerState.Loser)
        PlayerNameLabel.Content = Owner.PlayerName +
    (Owner.State == PlayerState.Winner ?
    " is the WINNER" : " has LOST");
    else
        PlayerNameLabel.Content = Owner.PlayerName;
    var isActivePlayer = (Owner.State == PlayerState.Active ||
    Owner.State == PlayerState.MustDiscard);
    PlayerNameLabel.FontSize = isActivePlayer ? 18 : 14;
    PlayerNameLabel.Foreground = isActivePlayer ?
    new SolidColorBrush(Colors.Gold) :
    new SolidColorBrush(Colors.White);
}
```

10. Finally, add the double-click handler that is called when the player double-clicks a card:

```
private void cardControl_MouseDoubleClick(object sender, MouseButtonEventArgs e)
{
    var selectedCard = sender as CardControl;
    if (Owner == null)
        return;
    if (Owner.State == PlayerState.MustDiscard)
        Owner.DiscardCard(selectedCard.Card);
    RedrawCards();
}
```

11. Create another user control like you did in Step 1 and name it GameDecksControl.

12. Remove the Grid and insert a Canvas control instead:

```
<Canvas Name="controlCanvas" Width="250" />
```

13. Go to the code-behind file use these namespaces:

```
using Ch13CardLib;
using System.Collections.Generic;
using System.Linq;
using System.Windows;
using System.Windows.Controls;
using System.Windows.Documents;
using System.Windows.Input;
```

14. As you did in Step 4, add four dependency properties with these values (see Table 15-8).

**TABLE 15-8:** Game Decks Dependency Properties

TYPE	NAME	OWNERCLASS	DEFAULT VALUE
bool	GameStarted	GameDecksControl	false
Player	CurrentPlayer	GameDecksControl	null
Deck	Deck	GameDecksControl	null
Card	AvailableCard	GameDecksControl	null

- 15.** All four properties require a callback method for when the property changes. Add these as you did in Step 4 with the names `OnGameStarted`, `OnPlayerChanged`, `OnDeckChanged`, and `OnAvailableCardChanged`.

- 16.** Add the callback methods:

```

private static void OnGameStarted(DependencyObject source,
DependencyPropertyChangedEventArgs e)
{
    var control = source as GameDecksControl;
    control.DrawDecks();
}
private static void OnPlayerChanged(DependencyObject source,
    DependencyPropertyChangedEventArgs e)
{
    var control = source as GameDecksControl;
    if (control.CurrentPlayer == null)
        return;
    control.CurrentPlayer.OnCardDiscarded +=
        control.CurrentPlayer_OnCardDiscarded;
    control.DrawDecks();
}
private void CurrentPlayer_OnCardDiscarded(object sender, CardEventArgs e)
{
    AvailableCard = e.Card;
    DrawDecks();
}
private static void OnDeckChanged(DependencyObject source,
DependencyPropertyChangedEventArgs e)
{
    var control = source as GameDecksControl;
    control.DrawDecks();
}
private static void OnAvailableCardChanged(DependencyObject source,
DependencyPropertyChangedEventArgs e)
{
    var control = source as GameDecksControl;
    control.DrawDecks();
}

```

**17.** Add the `DrawDecks` method:

```
private void DrawDecks()
{
    controlCanvas.Children.Clear();
    if (CurrentPlayer == null || Deck == null || !GameStarted)
        return;
    List<CardControl> stackedCards = new List<CardControl>();
    for (int i = 0; i < Deck.CardsInDeck; i++)
        stackedCards.Add(new CardControl(Deck.GetCard(i)) { Margin =
new Thickness(150 + (i * 1.25), 25 - (i * 1.25), 0, 0), IsFaceUp = false });
    if (stackedCards.Count > 0)
        stackedCards.Last().MouseDoubleClick += Deck_MouseDoubleClick;
    if (AvailableCard != null)
    {
        var availableCard = new CardControl(AvailableCard) { Margin =
new Thickness(0, 25, 0, 0) };
        availableCard.MouseDoubleClick += AvailableCard_MouseDoubleClick;
        controlCanvas.Children.Add(availableCard);
    }
    stackedCards.ForEach(x => controlCanvas.Children.Add(x));
}
```

**18.** Finally, add the event handlers for the cards:

```
void AvailableCard_MouseDoubleClick(object sender, MouseButtonEventArgs e)
{
    if (CurrentPlayer.State != PlayerState.Active)
        return;
    var control = sender as CardControl;
    CurrentPlayer.AddCard(control.Card);
    AvailableCard = null;
    DrawDecks();
}
void Deck_MouseDoubleClick(object sender, MouseButtonEventArgs e)
{
    if (CurrentPlayer.State != PlayerState.Active)
        return;
    CurrentPlayer.DrawCard(Deck);
    DrawDecks();
}
```

**19.** Return to the `GameClient.xaml` file and remove the `Grid` that is currently in Row 2. Instead, insert a new dock panel like this:

```
<DockPanel Grid.Row="2">
    <local:CardsInHandControl x:Name="Player2Hand" DockPanel.Dock="Right"
Height="380" Game="{Binding}">
        VerticalAlignment="Center" Width="180" PlayerOrientation="Vertical"
        Owner="{Binding Players[1]}" PlayerState="{Binding Players[1].State}" />
    <local:CardsInHandControl x:Name="Player4Hand" DockPanel.Dock="Left">
        HorizontalAlignment="Left" Height="380" VerticalAlignment="Center"
        PlayerOrientation="Vertical" Owner="{Binding Players[3]}" Width="180"
        PlayerState="{Binding Players[3].State}" Game="{Binding}"/>
```

```

<local:CardsInHandControl x:Name="Player1Hand" DockPanel.Dock="Top"
    HorizontalAlignment="Center" Height="154" VerticalAlignment="Top"
    PlayerOrientation="Horizontal" Owner="{Binding Players[0]}" Width="380"
    PlayerState="{Binding Players[0].State}" Game="{Binding}"/>
<local:CardsInHandControl x:Name="Player3Hand" DockPanel.Dock="Bottom"
    HorizontalAlignment="Center" Height="154" VerticalAlignment="Top"
    PlayerOrientation="Horizontal" Owner="{Binding Players[2]}" Width="380"
    PlayerState="{Binding Players[2].State}" Game="{Binding}"/>
<local:GameDecksControl Height="180" x:Name="GameDecks" Deck="{Binding
GameDeck}">
    AvailableCard="{Binding CurrentAvailableCard}"
    CurrentPlayer="{Binding CurrentPlayer}"
    GameStarted="{Binding GameStarted}"/>
</DockPanel>

```

20. Run the application. By default the ComputerPlayer class is enabled and the number of players is set to two. This means you select a single name in the Start Game dialog box. After that, you should be able to see something like Figure 15-4.




FIGURE 15-4

Double-click on the deck or available card to draw and then click a card from your hand to discard it.

## How It Works

Even though there is quite a bit of code in this example, most of it is the dependency properties, and the XAML is all about data binding these properties. The `CardsInHandControl` creates three properties that it uses to display itself and react to changes: `Game`, `Owner`, and `PlayerState`. `Game` and `Owner` are mostly used to draw, but the `PlayerState` is also used to control the `ComputerPlayer` actions.

```
private static void OnPlayerStateChanged(DependencyObject source,
   DependencyPropertyChangedEventArgs e)
{
    var control = source as CardsInHandControl;
    var computerPlayer = control.Owner as ComputerPlayer;
    if (computerPlayer != null)
    {
        if (computerPlayer.State == PlayerState.MustDiscard)
        {
            Thread delayedWorker = new Thread(control.DelayDiscard);
            delayedWorker.Start(new Payload
            {
                Deck = control.Game.GameDeck,
                AvailableCard = control.Game.CurrentAvailableCard,
                Player = computerPlayer
            });
        }
        else if (computerPlayer.State == PlayerState.Active)
        {
            Thread delayedWorker = new Thread(control.DelayDraw);
            delayedWorker.Start(new Payload
            {
                Deck = control.Game.GameDeck,
                AvailableCard = control.Game.CurrentAvailableCard,
                Player = computerPlayer
            });
        }
    }
    control.RedrawCards();
}
```

The `OnPlayerStateChanged` method, which is used to react to changes in the state of the player, determines if the current player is a `ComputerPlayer`. If it is, it checks to make sure that the computer player draws or discards a card. If this is the case, it creates a worker thread for this to happen and executes the methods on this thread. This allows the application to continue working while the computer is waiting:

```
private void DelayDraw(object payload)
{
    Thread.Sleep(1250);
    var data = payload as Payload;
    Dispatcher.Invoke(DispatcherPriority.Normal,
        new Action<Deck, Card>(data.Player.PerformDraw), data.Deck, data.AvailableCard);
}
```

The Dispatcher is used to invoke the call. This ensures that the calls are made on the GUI thread.

Drawing the cards is pretty straightforward. The program simply stacks them vertically or horizontally depending on the settings in PlayerOrientation.

The GameDecksControl is even simpler. It uses the CurrentPlayer class to be notified that the CurrentPlayer has changed. When this happens, it hooks the CardDiscarded event on the player, and uses this event to be notified that the card was discarded.

Finally, you add a dock panel to the game client with a CardsInHandControl on each side and with a GameDecksControl in the middle:

```
<local:CardsInHandControl x:Name="Player1Hand" DockPanel.Dock="Top"
    HorizontalAlignment="Center" Height="154" VerticalAlignment="Top"
    PlayerOrientation="Horizontal" Owner="{Binding Players[0]}" Width="380"
    PlayerState="{Binding Players[0].State}" Game="{Binding}" />
```

The binding for Game simply binds the DataContext of the game client directly to the Game property of the CardsInHandControl. The PlayerState is bound to the State property of a player. In this case, the player at index 0 is used to access the state.

## EXERCISES

- 15.1** The current game client has a problem. From the Options dialog box, you can set the skill level of the computer. The problem is that the radio buttons are not updated to reflect the choice the next time you open the Options dialog box. This is partly because there is nothing that tries to update them and partly because there is no value converter from ComputerSkillLevel. Fix this problem by creating a new value converter and setting the IsChecked binding instead of using the Checked event that is currently being used.

Hint: You must use the ConverterParameter part of the Converter binding.

- 15.2** The computer cheats, so you might want to allow the players to cheat as well. On the Options dialog box, create an option for the computer to play with open cards.
- 15.3** Create a status bar at the bottom of the game client that displays the current state of the game.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Styles	You can use styles to create styles for XAML elements that can be reused on many elements. Styles allow you to set the properties of an element. When you set the <code>Style</code> property of an element to point to a style you have defined, the properties of the element will use the values you specified in the <code>Style</code> property.
Templates	Templates are used to define the content of a control. Using templates you can change how standard controls are displayed. You can also build complex custom controls with them.
Value converters	Value converters are used to convert to and from two types. To create a value converter, you must implement the interface <code>IValueConverter</code> on a class.
User controls	User controls are used to create code and XAML that can be reused easily in your own project. This code and XAML can also be exported for use in other projects.

# PART III

# Cloud Programming

---

- **CHAPTER 16:** Basic Cloud Programming
- **CHAPTER 17:** Advanced Cloud Programming and Deployment



# 16

## Basic Cloud Programming

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Understanding the cloud, cloud programming, and the cloud optimized stack
- Programming for the cloud using cloud design patterns
- Using Microsoft Azure C# libraries to create a storage container
- Creating an ASP.NET 4.6 Web site that uses the storage container

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

The `wrox.com` code downloads for this chapter are found at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 16 download and individually named according to the names throughout the chapter.

In this book, the basics of C# programming are conveyed mostly using console applications, desktop applications with WPF, and Windows Universal Apps. Although these are viable and compelling development techniques, they are not good examples of programs to host and run in the cloud. These kinds of programs are classically deployed to and run on a user's computer, tablet, or mobile device. These programs are compiled into executables or dynamic linked libraries that have dependencies on preinstalled software like the .NET Framework, for example. These dependencies are generally assumed to be present on the location where they are installed, or they get included in the installation procedure. By contrast, an Internet application that is run in the cloud, based on ASP.NET, for instance, cannot rely on any such library or dependency being present on the computer or device from which the program is accessed. All dependencies instead are installed on the server hosting the Internet application and are accessed by a device using protocols such as HTTP, WS (web socket), FTP, or SMTP. Although console, desktop, and Windows Universal Apps can have dependencies in the cloud, like databases, storage containers, or web services, they themselves are generally not hosted there.

Programs that are accessed via web browsers and respond to REST API or WCF service requests are good candidates for running in the cloud. Development techniques used for creating those program types do not require any built-in dependency on the device from which they are called. In typical cases, those program types merely exchange information between themselves and render data in a legible and user-friendly fashion. Additionally, programs that receive and process large amounts of data are good candidates for running in the cloud, as utilizing the hyper-scalability of resources to accept and process the data is a fundamental feature of the cloud itself.

This chapter provides an overview of what the cloud is, some examples of patterns and techniques for successfully running a program in the cloud, and an example of creating and using cloud resources from an ASP.NET web site.

## THE CLOUD, CLOUD COMPUTING, AND THE CLOUD OPTIMIZED STACK

It is only a matter of time before you begin creating applications that run completely or partially in the cloud. It's no longer a question of "if" but "when." Deciding which components of your program will run in the cloud, the cloud type, and the cloud service model requires some investigation, understanding, and planning. For starters, you need to be clear on what the cloud is. The cloud is simply a large amount of commoditized computer hardware running inside a datacenter that can run programs and store large amounts of data. The differentiator is elasticity, which is the ability to scale up (for example, increase CPU and memory) and/or out (for example, increase number of virtual server instances) dynamically, then scale back down with seemingly minimal to no effort. This is an enormous difference from the current IT operational landscape where differentiated computer resources often go partially or completely unused in one area of the company, while in other areas there is a serious lack of computer resources. The cloud resolves this issue by providing access to computer resources as you need them, and when you don't need them those resources are given to someone else. For individual developers, the cloud is a place to deploy your program and expose it to the world. If by chance the program becomes a popular one, you can scale to meet your resources needs; if the program is a flop, then you are not out much money or time spent on setting up dedicated computer hardware and infrastructure.

Let's explore cloud type and cloud service models in more detail now. The common cloud types are public, private, and hybrid and are described in the following bullet points and illustrated in Figure 16-1.

- **Public cloud** is shared computer hardware and infrastructure owned and operated by a cloud provider like Microsoft Azure, Amazon AWS, Rackspace, or IBM Cloud. This cloud type is ideal for small and medium businesses that need to manage fluctuations in customer and user demands.
- **Private cloud** is dedicated computer hardware and infrastructure that exists onsite or in an outsourced data center. This cloud type is ideal for larger companies or those that must deliver a higher level of data security or government compliance.
- **Hybrid cloud** is a combination of both public and private cloud types whereby you choose which segments of your IT solution run on the private cloud and which run on the public cloud. The ideal solution is to run your businesses-critical programs that require a greater

level of security in the private cloud and run non-sensitive, possibly spiking tasks in the public cloud.




FIGURE 16-1

The number of cloud service models continues to increase, but the most common cloud service models are described in the following bullet points and illustrated in Figure 16-2.

- **Infrastructure as a Service (IaaS)** — You are responsible from the operating system upward. You are not responsible for the hardware or network infrastructure; however, you are responsible for operating system patches and third-party dependent libraries.
- **Platform as a Service (PaaS)** — You are responsible only for your program running on the chosen operating system and its dependencies. You are not responsible for operating system maintenance, hardware, or network infrastructure.
- **Software as a Service (SaaS)** — A software program or service used from a device that is accessed via the Internet. For example, O365, Salesforce, OneDrive or Box, all of which are accessible from anywhere with an Internet connection and do not require software to be installed on the client to function. You are only responsible for the software running on the platform and nothing else.




FIGURE 16-2

In summary, the cloud is an elastic structure of commoditized computer hardware for running programs. These programs run on IaaS, PaaS, or SaaS service models in a Hybrid, Public, or Private Cloud type.

Cloud programming is the development of code logic that runs on any of the cloud service models. The cloud program should incorporate portability, scalability, and resiliency patterns that improve the performance and stability of the program. Programs that do not implement these portability, scalability, and resiliency patterns would likely run in the cloud, but some circumstances such as a hardware failure or network latency issue may cause the program to execute an unexpected code path and terminate.

**NOTE** *Cloud programming patterns and best practices are discussed in the next section.*

Reflecting back to the elasticity of the cloud as being one of its most favorable benefits, it is important that not only the platform is able to scale, but the cloud program can as well. For example, does the code rely on backend resources, databases, read or open files, or parse through large data objects? These kinds of functional actions within a cloud program can reduce its ability to scale and therefore have a low support for throughput. Make sure your cloud program manages code paths that execute long running methods and perhaps place them into an offline process mechanism.

The cloud optimized stack is a concept used to refer to code that can handle high throughput, makes a small footprint, can run side-by-side with other applications on the same server, and is cross-platform enabled. Having a small footprint relates to packaging of only the components into your cloud program for which a dependency exists, making the deployment size as small as possible. Does the cloud program require the entire .NET Framework to function? Instead of requiring the entire .NET Framework, include only the libraries required to run your cloud program, and then compile your cloud program into a self-contained application to support side-by-side execution. The

cloud program can run alongside any other cloud program as it contains the dependencies within the binaries package itself. Finally, by using an open source version of C# named Mono, the cloud program can be packaged, compiled, and deployed onto operating systems that are not Microsoft, for example Mac OS X, Linux, or UNIX.

## CLOUD PATTERNS AND BEST PRACTICES

In the cloud, very brief moments of increased latency or downtime are expected and your code must be prepared for this and include logic to successfully recover from these platform exceptions. This is a significant mind shift if you have historically been coding for onsite or on-premise program executions. You need to unlearn a lot of what you know about managing exceptions and learn to embrace failure and create your code to recover from such failures.

In the previous section words like portability, scalability, and resiliency were touched upon in the context of integrating those concepts into your program slated to run in the cloud. But what does portability specifically mean here? A program is portable if it can be moved or executed on multiple platforms, for example Windows, Linux, and Mac OS X. Take for example some ASP.NET 4.6 features that sit on a new stack of open source technologies that provide the developer with options to compile code into binaries capable of running on any of those platforms. Traditionally, a developer who wrote a program using ASP.NET, with C# in the background, would run it on a Windows server using Internet Information Server (IIS). However, from a core cloud-centric perspective, the ability of your program and all its dependencies to move from one virtual machine to another, without manual or programmatic intervention, is the most applicable form of portability in this context. Remember that failures in the cloud are expected, and the virtual machine (VM) on which your program is running can be wiped out at any given time and then be rebuilt fresh on another VM. Therefore, your program needs to be portable and able to recover from such an event.

Scalability means that your code responds well when multiple customers use it. For example, if you have 1,500 requests per minute, that would be roughly 25 concurrent requests per second, if the request is completed and responded to in 1 second. However, if you have 15,000 request per minute, that would mean 250 concurrent requests per second. Will the cloud program respond in the same manner with 25 or 250 concurrent requests? How about 2,550? The following are a few cloud programming patterns that are useful for managing scalability.

- **Command and Query Responsibility Segregation (CQRS) pattern** — This pattern concerns the separation of operations that read data from operations that modify or update the data.
- **Materialized View pattern** — This modifies the storage structure to reflect the data query pattern. For example, creating views for specific highly used queries can make for more efficient querying.
- **Sharding pattern** — This breaks your data into multiple horizontal shards that contain a distinct subset of the data as opposed to vertical scaling via the addition of hardware capacity.
- **Valet Key pattern** — This gives clients direct access to the data store for streaming of or uploading of large files. Instead of having a web client manage the gatekeeping to the data store, it provides a client with a Valet Key and direct access to the data store.

**NOTE** These patterns cover some advanced C# coding techniques and therefore only descriptions of the patterns are provided. If you are interested in seeing the actual C# code to implement the patterns, they certainly exist and can be found by searching for them on the Internet.

Resiliency refers to how well your program responds and recovers from service faults and exceptions. Historically, IT infrastructures have been focused on failure prevention where the acceptability of downtime was minimal and 99.99% or 99.999% SLAs (service-level agreements) were the expectation. Running a program in the cloud, however, requires a reliability mind shift, one which embraces failure and is clearly oriented toward recovery and not prevention. Having multiple dependencies such as database, storage, network, and third-party services, some of which have no SLA, requires this shift in perspectives. User-friendly reactions in response to outages or situations that are not considered normal operation make your cloud program resilient. Here are a few cloud programming patterns that are useful for embedding resiliency into your cloud program:

- **Circuit Breaker pattern** — This is a code design that is aware of the state of remote services and will only attempt to make the connection if the service is available. This avoids attempting a request and wasting CPU cycles when it is already known the remote service is unavailable via previous failures.
- **Health Endpoint Monitoring pattern** — This checks that cloud-based applications are functionally available via the implementation of endpoint monitoring.
- **Retry pattern** — This retries the request after transient exceptions or failure. This pattern retries a number of times within a given timeframe and stops when the retry attempt threshold is breached.
- **Throttling pattern** — This manages the consumption of a cloud program so that SLAs can be met and the program remains functional under high load.

Using one or more of the patterns described in this section will help make your cloud migration more successful. The discussed patterns enhance usability of your program by improving the scalability and resiliency of it. This in turn makes for a more pleasant user or customer experience.

## USING MICROSOFT AZURE C# LIBRARIES TO CREATE A STORAGE CONTAINER

Although there are numerous cloud providers, the cloud provider used for the examples in this and the next chapter is Microsoft. The cloud platform provided by Microsoft is called Azure. Azure has many different kinds of features. For example the IaaS offering is called Azure VM, and the PaaS offering is called Azure Cloud Services. Additionally, Microsoft has SQL Azure for database, Azure Active Directory for user authentication, and Azure Storage for storing blobs, for example.

**NOTE** Exercise 1 requires that you have a Microsoft Azure subscription. If you do not have one, you can sign up for a free trial here: <http://azure.microsoft.com>.

The following two exercises will walk you through the creation of an Azure storage account and the creation of an Azure storage container using the Microsoft Azure Storage Client Library for .NET with C#. Then in the next section you will create an ASP.NET 4.6 Web Site, using Visual Studio, to access the images stored in the Azure storage container. Following along the flow of the book, the ASP.NET 4.6 Web Site will deal a hand of playing cards. The card images are the blobs stored in the Azure storage container.

### TRY IT OUT Try It Out: Create an Azure Storage Account

1. Access the Microsoft Azure Portal at <https://manage.windowsazure.com>.
2. Click on the STORAGE feature and then CREATE A STORAGE ACCOUNT as shown in Figure 16-3.
3. Enter the storage account name, location, and replication. For cost reasons, consider setting the redundancy to Locally Redundant. This avoids the small additional cost of a shadow copy of your storage account being created in another region. Your files, however, are copied three times within the same regional data center in which storage account exists.




FIGURE 16-3

4. Once the name, location, and replication data is entered, click CREATE STORAGE ACCOUNT and confirm you see something similar to what is shown in Figure 16-4. In this example, the Azure

storage account is `deckofcards`. Give your Azure storage account a different name. Remember the name, as it is used in the next Try It Out exercise.

The screenshot shows the Microsoft Azure portal interface. On the left, a navigation bar lists various cloud services: Cloud Services, SQL Databases, Storage (selected), HDInsight, Media Services, Service Bus, Visual Studio Online, Cache, BizTalk Services, Recovery Services, CDN, Automation, and Scheduler. The main content area is titled "storage" and displays a table with one row. The table has columns for NAME, STATUS, LOCATION, and SUBSCRIPTION. The row for the storage account "deckofcards" shows "Online" in the STATUS column, "West Europe" in the LOCATION column, and a QR code in the SUBSCRIPTION column. At the bottom of the page, a success message states "Storage account 'deckofcards001' created successfully." and includes "OK" and "CLOSE" buttons. The footer features a "NEW" button, "MANAGE ACCESS", "MANAGE DOMAIN", and "DELETE" buttons, along with a user icon and a help icon.

FIGURE 16-4

5. You have now successfully created an Azure storage account.

**NOTE** Storage accounts can be used for storing blobs, tables, queues, and files. Some examples of these include database backups, Azure Web App IIS logs, VM machine images, documents, or images, with a limit of 100TB per storage account.


## How It Works

The Microsoft Azure management console itself runs on the Microsoft Azure platform in the PaaS cloud service model, i.e. Azure Cloud Services. The management console is written by a product team within Microsoft and supported by additional Microsoft support staff. All the features you find on the left-hand side navigation bar can be created and utilized. Creating an Azure storage account with your subscriptions allocates storage space and a globally accessible URL for accessing the contents of the storage account (for example: <https://deckofcards.blob.core.windows.net>).


**TRY IT OUT****Create an Azure Storage Container Using the Microsoft Azure Storage Client Library**

You will create a console application using Visual Studio 2015 and the Microsoft Azure Storage Client libraries to create an Azure storage container and upload the 52 cards into it.

1. Create a new Console Application project by selecting File  $\Rightarrow$  New  $\Rightarrow$  Project within Visual Studio. In the New Project dialog box (see Figure 16-5), select the category Visual C# and the subcategory Windows, and then select the Console Application template. Name the project Ch16Ex01.
2. Add a directory named Cards to the project by right-clicking on Ch16Ex01  $\Rightarrow$  Add...  $\Rightarrow$  New Folder. Add the 52 card images to the directory, similar to what is shown in Figure 16-6. The images are available from the source code download site and are named from 0-1.PNG to 3-13.PNG.



**FIGURE 16-5**



**FIGURE 16-6**

3. Additionally, copy the Cards directory into C:\BegVCSharp\Chapter16\Ch16Ex01\Ch16Ex01\bin\Debug so that the compiled executable can find them when run.
4. Right-click again on the Ch16Ex01 project and select Manage NuGet Packages... from the popup menu.
5. In the search textbox, as shown in Figure 16-7, enter Windows Azure Storage and install the WindowsAzure.Storage client library. You can find more information about the Windows Azure Storage library here: <https://msdn.microsoft.com/library/azure/dn261237.aspx>.




FIGURE 16-7

6. Accept the user agreements and once the NuGet package and its dependencies are installed, you should see a ===== Finished===== message in the Output window of Visual Studio. Additionally, the References folder within Ch16Ex01 is expanded, and you can view the newly added binaries.
7. Open the App.config file and add the following <appSetting> settings into the <configuration> section. Notice that the AccountName is the name of the Azure storage account created in the previous Try It Out exercise (deckofcards). You would change this to the name of your Azure storage account. Refer to step 8 for instructions on how to get the AccountKey.

```
<appSettings>
<add key="StorageConnectionString"
    value="DefaultEndpointsProtocol=https;AccountName=<NAME>;
    AccountKey=<KEY>" />
</appSettings>
```
8. To get the Azure storage account key, access the Microsoft Azure management portal and navigate to your Azure storage account. As seen in Figure 16-8, at the bottom of the page there is an item

called Manage Access Keys. Select that and copy the PRIMARY ACCESS KEY to your clipboard and into the `App.config` file as the value for the `AccountKey`.




FIGURE 16-8

- Now add the code that creates the container, uploads the images, lists them, and if desired deletes them. First add the assembly references and the `try{...}catch{...}` C# framework to the `Main()` method, as shown here.

```
using static System.Console;
using System.IO;
using Microsoft.WindowsAzure;
using Microsoft.WindowsAzure.Storage;
using Microsoft.WindowsAzure.Storage.Auth;
using Microsoft.WindowsAzure.Storage.Blob;

static void Main(string[] args)
{
    try {
    }
    catch (StorageException ex)
    {
        WriteLine($"StorageException: {ex.Message}");
    }
    catch (Exception ex)
```

```
{  
    WriteLine($"Exception: {ex.Message}");  
}  
WriteLine("Press enter to exit.");  
ReadLine();  
}
```

10. Next, add the code within the `try{}` code block that creates the container, as shown here. Look at the parameter passed to the `blobClient.GetContainerReference("carddeck")`, `carddeck`. This is the name used for the Azure storage container. Content within this container can then be accessed via <https://deckofcards.blob.core.windows.net/carddeck/0-1.PNG>, for example. You can place any desired name as long as it meets the naming requirements (for example, it must be 3 to 63 characters long and must begin with a letter or number). If you provide a container name that does not meet the naming requirements, a 400 HTTP status error is returned.

```
CloudStorageAccount storageAccount =  
CloudStorageAccount.Parse(CloudConfigurationManager.GetSetting("StorageConnectionString"));  
CloudBlobClient blobClient = storageAccount.CreateCloudBlobClient();  
CloudBlobContainer container = blobClient.GetContainerReference("carddeck");  
if (container.CreateIfNotExists())  
{  
    WriteLine($"Created container '{container.Name}' " +  
        $"in storage account '{storageAccount.Credentials.AccountName}'.");  
}  
else  
{  
    WriteLine($"Container '{container.Name}' already exists " +  
        $"for storage account '{storageAccount.Credentials.AccountName}'.");  
}  
container.SetPermissions(new BlobContainerPermissions  
{ PublicAccess = BlobContainerPublicAccessType.Blob });  
WriteLine($"Permission for container '{container.Name}' is public.");
```

11. Add this code following the code that creates the container, which uploads the card images stored in the Cards folder.

```
int numberOfCards = 0;  
DirectoryInfo dir = new DirectoryInfo(@"Cards");  
foreach (FileInfo f in dir.GetFiles("*.*))  
{  
    CloudBlockBlob blockBlob = container.GetBlockBlobReference(f.Name);  
    using (var fileStream = System.IO.File.OpenRead(@"Cards\" + f.Name))  
    {  
        blockBlob.UploadFromStream(fileStream);  
        WriteLine($"Uploading: '{f.Name}' which " +  
            $"is {fileStream.Length} bytes.");  
    }  
    numberOfCards++;  
}  
WriteLine($"Uploaded {numberOfCards.ToString()} cards.");  
WriteLine();
```

12. Now that the images are uploaded, just to check that all went well, add this code to list the blobs stored in the newly created Azure storage container, named carddeck.

```
numberOfCards = 0;
foreach (IListBlobItem item in container.ListBlobs(null, false))
{
    if (item.GetType() == typeof(CloudBlockBlob))
    {
        CloudBlockBlob blob = (CloudBlockBlob)item;
        WriteLine($"Card image url '{blob.Uri}' with length " +
            $" of {blob.Properties.Length}");
    }
    numberOfCards++;
}
WriteLine($"Listed {numberOfCards.ToString()} cards.");
```

13. Now, if desired, you can delete the images that were just uploaded. This is really to show an example of how you can delete the blob files from your container programmatically.

```
WriteLine("Enter Y to delete listed cards, press enter to skip deletion:");
if (ReadLine() == "Y")
{
    numberOfCards = 0;
    foreach (IListBlobItem item in container.ListBlobs(null, false))
    {
        CloudBlockBlob blob = (CloudBlockBlob)item;
        CloudBlockBlob blockToDelete = container.GetBlockBlobReference(blob.Name);
        blockToDelete.Delete();
        WriteLine($"Deleted: '{blob.Name}' which was {blob.Name.Length} bytes.");
        numberOfCards++;
    }
    WriteLine($"Deleted {numberOfCards.ToString()} cards.");
}
```

14. Run the console application and review the output. You should see something similar to what is shown in Figure 16-9. Then access the Microsoft Azure management console and look on the container page for the newly created container named, for example, carddeck as shown in Figure 16-10. Click on the container to view its contents.

```
file:///C:/BegVCSharp/Chapter16/Ch16Ex01/Ch16Ex01/bin/Debug/Ch16Ex01.exe
Created container 'carddeck' in storage account 'deckofcards'.
Permission for container 'carddeck' is public.
Uploading: '0-1.PNG' which is 3925 bytes.
Uploading: '0-10.PNG' which is 9985 bytes.
Uploading: '0-11.PNG' which is 19657 bytes.
Uploading: '0-12.PNG' which is 19143 bytes.
Uploading: '0-13.PNG' which is 19091 bytes.
Uploading: '0-2.PNG' which is 4744 bytes.
Uploading: '0-3.PNG' which is 5621 bytes.
Uploading: '0-4.PNG' which is 5716 bytes.
Uploading: '0-5.PNG' which is 6872 bytes.
Uploading: '0-6.PNG' which is 7713 bytes.
Uploading: '0-7.PNG' which is 8486 bytes.
Uploading: '0-8.PNG' which is 9220 bytes.
Uploading: '0-9.PNG' which is 9722 bytes.
Uploading: '1-1.PNG' which is 3878 bytes.
Uploading: '1-10.PNG' which is 5269 bytes.
Uploading: '1-11.PNG' which is 18692 bytes.
Uploading: '1-12.PNG' which is 19142 bytes.
Uploading: '1-13.PNG' which is 19543 bytes.
Uploading: '1-2.PNG' which is 3109 bytes.
Uploading: '1-3.PNG' which is 3623 bytes.
Uploading: '1-4.PNG' which is 3376 bytes.
Uploading: '1-5.PNG' which is 3925 bytes.
Uploading: '1-6.PNG' which is 4494 bytes.
Uploading: '1-7.PNG' which is 5954 bytes.
Uploading: '1-8.PNG' which is 5555 bytes.
Uploading: '1-9.PNG' which is 5465 bytes.
```

FIGURE 16-9




FIGURE 16-10

## How It Works

It is programmatically possible to create a Microsoft Azure storage account, but the security aspect of that creation is relatively complex and that step is performed from within the Microsoft Azure Management console directly. Once an Azure storage account is created, you can then create multiple containers within the account. In this example, you created a container called `carddeck`. There is only a limit on the number of storage accounts per Microsoft Azure subscription and no limit on the number of containers within the storage account. You can create as much and as many as you want, but keep in mind that each comes with a cost.

The code is split into four sections (create the container, upload the images to the container, list the blobs in the container, and optionally delete the contents of the container). The first action taken was to set up the `try{ }...catch{ }` framework for the console application. This is a good practice because uncaught or unhandled exceptions typically crash the process (EXE), which is something that should always be avoided. The first `catch()` expression is the `StorageException` and captures exceptions thrown specifically from methods within the `Microsoft.WindowsAzure.Storage` namespace.

```
catch (StorageException ex)
```

Then there is a catch all exceptions expression that handles all other unexpected exceptions and writes the exception message for them to the console.

```
    catch (Exception ex)
```

The first line within the `try{}` code block creates the storage account using the details added to the `App.config` file.

```
CloudStorageAccount storageAccount =
    CloudStorageAccount.Parse(CloudConfigurationManager.GetSetting
        ("StorageConnectionString"));
```

The `App.config` file contains the storage account name and the secret storage account key that is needed for performing administrative actions on the Azure storage account. Next, you create a client that manages the interface with a specific blob container within the storage account. Then the code gets a reference to a specific container named `carddeck`.

```
CloudBlobClient blobClient = storageAccount.CreateCloudBlobClient();
CloudBlobContainer container =
    blobClient.GetContainerReference("carddeck");
```

Next the `container.CreateIfNotExists()` method is called. If the container is created, meaning it does not already exist, then the value `true` is returned and that information is written to the console. Otherwise, `false` is returned if the container does already exist.

```
if (container.CreateIfNotExists())
    {...} ...
```

Containers can be `Private` or `Public`. For this example the container is public, which means an access key is not required to access it. The container is set to be public by executing this code.

```
container.SetPermissions(new BlobContainerPermissions
    { PublicAccess = BlobContainerPublicAccessType.Blob });
```

At this point the container is created and publicly accessible, but it is empty. Using a `System.IO` method like `DirectoryInfo` and `FileInfo`, you created a `foreach` loop that added each of the card images to the `carddeck` storage container. The `GetBlockBlobReference()` method is used to set the reference to the specific image name to be added to the container. Then using the filename and path, the `System.IO.File.OpenRead()` method opens the actual file as a `FileStream`, and it is uploaded to the container via the `UploadFromStream()` method.

```
CloudBlockBlob blockBlob = container.GetBlockBlobReference(f.Name);
using (var fileStream = System.IO.File.OpenRead(@"Cards\" + f.Name))
{
    blockBlob.UploadFromStream(fileStream);
```

All of the files in the `Cards` directory are looped through and uploaded to the container. Using the same container object created during the initial creation of the `carddeck` container, by calling the

`ListBlob()` method, a list of existing blobs are returned as an `IEnumerable<IListBlobItems>`. You then loop through the list and write them to the console.

```
foreach (IListBlobItem item in container.ListBlobs(null, false))
{
    if (item.GetType() == typeof(CloudBlockBlob))
    {
        CloudBlockBlob blob = (CloudBlockBlob)item;
        WriteLine($"Card image url '{blob.Uri}' with length of " +
            $" {blob.Properties.Length}");
    }
    numberOfCards++;
}
```

As previously noted, there are numerous types of items that can be stored in a container, like blobs, tables, queues, and files. Therefore, prior to boxing the item as a `CloudBlockBlob`, it is important to confirm that the item is indeed a `CloudBlockBlob`. Other types to check for are `CloudPageBlob` and `CloudBlobDirectory`.

To delete the blobs in the container, first the list of blobs is retrieved in the same manner as previously performed when looping through them and writing them to the console. The difference when deleting them is that `GetBlockBlobReference(blob.Name)` is called to get a reference to the specific blob, then the `Delete()` method is called for that specific blob.

```
CloudBlockBlob blockBlobToDelete = container.GetBlockBlobReference(blob.Name);
blockBlobToDelete.Delete();
```

---

Now that the Microsoft Azure storage account and container are created and loaded with the images of a 52-card deck, you can create an ASP.NET Web Site to reference the Microsoft Azure storage container.

## CREATING AN ASP.NET 4.6 WEB SITE THAT USES THE STORAGE CONTAINER

Up to now there has not been any in-depth examination of what a web application is nor a discussion about the fundamental aspects of ASP.NET. This section provides some insight into these technical perspectives.

A web application causes a web server to send HTML code to a client. That code is displayed in a web browser such as Internet Explorer. When a user enters a URL string in the browser, an HTTP request is sent to the web server. The HTTP request contains the filename that is requested along with additional information such as a string identifying the client application, the languages that the client supports, and additional data belonging to the request. The web server returns an HTTP response that contains HTML code, which is interpreted by the web browser to display text boxes, buttons, and lists to the user.

ASP.NET is a technology for dynamically creating web pages with server-side code. These web pages can be developed with many similarities to client-side Windows programs. Instead of dealing

directly with the HTTP request and response and manually creating HTML code to send to the client, you can use controls such as `TextBox`, `Label`, `ComboBox`, and `Calendar`, which create HTML code.

Entire books exist on web applications and ASP.NET. Covering those topics in detail is beyond the scope of this book. However, this section will take a quick look at using ASP.NET runtime and creating an ASP.NET web site that uses the storage container.

Using ASP.NET for web applications on the client system requires only a simple web browser. You can use Internet Explorer, Chrome, Firefox, or any other web browser that supports HTML. The client system doesn't require .NET to be installed.

On the server system, the ASP.NET runtime is needed. If you have Internet Information Services (IIS) on the system, the ASP.NET runtime is configured with the server when the .NET Framework is installed. During development, there's no need to work with Internet Information Services because Visual Studio delivers its own ASP.NET Web Development server that you can use for testing and debugging the application.

To understand how the ASP.NET runtime goes into action, consider a typical web request from a browser (see Figure 16-11). The client requests a file, such as `default.cshtml`, from the server. ASP.NET web form pages usually have the file extension `.aspx`, (although ASP.NET MVC has no specific file extension), and `.cshtml` is used for Razor-based Web Sites. Because these file extensions are registered with IIS or known by the ASP.NET Web Development Server, the ASP.NET runtime and the ASP.NET worker process enter the picture. The IIS worker process is named `w3wp.exe` and is host to your application on the web server. With the first request to the `default.cshtml` file, the ASP.NET parser starts, and the compiler compiles the file together with the C# code, which is associated with the `.cshtml` file and creates an assembly. Then the assembly is compiled to native code by the JIT compiler of the .NET runtime. Then the Page object is destroyed. The assembly is kept for subsequent requests, though, so it is not necessary to compile the assembly again.




FIGURE 16-11

Now that you have a basic understanding of what web applications and ASP.NET are, perform the steps in the following Try It Out.

### TRY IT OUT Create an ASP.NET 4.6 Web Site That Deals Two Hands of Cards

Again, you will use Visual Studio 2015, but this time you create an ASP.NET Web Site that requests the names of two players, and then when the page is submitted, two hands of cards are dealt. Those cards are downloaded from the Microsoft Azure storage container created earlier, and the cards are displayed on the web page.

1. Create a new Web Site project by selecting File  $\Rightarrow$  New  $\Rightarrow$  Web Site... within Visual Studio. In the New Web Site dialog box (see Figure 16-12), select the category Visual C# and then select the ASP .NET Empty Web Site template. Name the Web Site Ch16Ex02.
2. Add an ASP.NET Folder named App\_Code by right-clicking on the Ch16Ex02 solution, and then select Add  $\Rightarrow$  Add ASP.NET Folder  $\Rightarrow$  App\_Code.




FIGURE 16-12

3. Download the sample code from the download site and place the following class files into the App\_Code folder you just created. Once downloaded, right-click on the App\_Code folder, select Add  $\Rightarrow$  Existing Item..., and select the seven classes from the downloaded example.
  - a. Card.cs
  - b. Cards.cs
  - c. Deck.cs
  - d. Game.cs
  - e. Player.cs
  - f. Rank.cs
  - g. Suit.cs

**NOTE** The classes in step 3 are very similar to those used in previous examples. Only a few modifications were implemented, like the removal of `WriteLine()` and `ReadLine()` methods and some unused methods. Look in the `Card.cs` class, and you will see a new constructor which contains the link to the card image.

- Add a `default.cshtml` Razor v3 file to the project by right-clicking on the CH16Ex02 solutions, and then select Add New Item...  $\Rightarrow$  Visual C#  $\Rightarrow$  Empty Page (Razor v3) as shown in Figure 16-13.




FIGURE 16-13

- Open the `default.cshtml` file and place the following code at the top of the page.

```
@{

    Player[] players = new Player[2];
    var player1 = Request["PlayerName1"];
    var player2 = Request["PlayerName2"];

    if(IsPost)
    {
        players[0] = new Player(player1);
        players[1] = new Player(player2);
        Game newGame = new Game();
        newGame.SetPlayers(players);
        newGame.DealHands();
    }
}
```

- Next, add this syntax under the code you added in step 5. Pay close attention to the `@card.image-Link`, which is the newly added parameter to the `Card` class.

```

<!DOCTYPE html>
<html lang="en">
<head>
<meta charset="utf-8" />
<style>
    body {font-family:Verdana; margin-left:50px; margin-top:50px; }
    div {border: 1px solid black; width:40%; margin:1.2em;padding:1em; }
</style>
<title>BensCards: a new and exciting card game. </title>
</head>
<body>
@if(IsPost)
    <label id="labelGoal">Which player has the best hand.</label>
    <br />
    <div>
        <p><label id="labelPlayer1">Player1: @player1</label></p>
        @foreach(Card card in players[0].PlayHand)
        {
            <img width="75px" height="100px" alt="cardImage"
            src=
                "https://deckofcards.blob.core.windows.net/carddeck/@card.imageLink" />
        }
    </div>
    <div>
        <p><label id="labelPlayer1">Player2: @player2</label></p>
        @foreach(Card card in players[1].PlayHand)
        {
            <img width="75px" height="100px" alt="cardImage"
            src=
                "https://deckofcards.blob.core.windows.net/carddeck/@card.imageLink" />
        }
    </div>
}
else
{
    <label id="labelGoal">
        Enter the players name and deal the cards.
    </label>
    <br /><br />
    <form method="post">
        <div>
            <p>Player 1: @Html.TextBox("PlayerName1")</p>
            <p>Player 2: @Html.TextBox("PlayerName2")</p>
            <p><input type="submit" value="Deal Cards" class="submit"></p>
        </div>
    </form>
}
</body>
</html>

```

7. Now, run the Web Site by pressing F5 or the Run button within Visual Studio. A browser will start up and you should see a page rendered similar to the one illustrated in Figure 16-14. First you are prompted to enter in the Player names. Enter any two names.




FIGURE 16-14

8. Press the Deal Cards button, and a hand of cards is dealt to each player. You would see something similar to what is shown in Figure 16-15.




FIGURE 16-15

You have now created a simple ASP.NET Web Site using Razor v3. The ASP.NET Web Site connects to the Azure storage account and container for displaying the images of the playing cards.

## How It Works

You certainly noticed a new technology named Razor that was used in the previous exercise. Razor is a view engine that was introduced with ASP.NET 3 MVC along with Visual Studio 2013. Razor, as you have seen, uses C#-like language (VB is supported too) that is placed within a @{ ... } code block and is compiled and executed when the page is requested from a browser. Take a look at this code:

```
@{  
  
    Player[] players = new Player[2];  
    var player1 = Request["PlayerName1"];  
    var player2 = Request["PlayerName2"];  
  
    if(IsPost)  
    {  
        players[0] = new Player(player1);  
        players[1] = new Player(player2);  
        Game newGame = new Game();  
        newGame.SetPlayers(players);  
        newGame.DealHands();  
    }  
}
```

The code is encapsulated within a @{ ... } code block and is compiled and executed by the Razor engine when accessed. When the page is accessed, an array of type `Player[]` is created and the contents of the query string are populated into the two variables called `player1` and `player2`. If the page is not a post back, which means the page was simply requested (GET) instead of a button click (POST), then the code within the `if(IsPost) {}` code block does not execute. If the request to the page is a POST, which happens when you click the Deal Cards button, the `Players` are instantiated, a new game is started, and the hands of cards get dealt.

The initial request to the `default.cshtml` file executes this code path because it is not a POST.

```
else  
{  
    <label id="labelGoal">  
        Enter the players name and deal the cards.  
    </label>  
    <br /><br />  
    <form method="post">  
        <div>  
            <p>Player 1: @Html.TextBox("PlayerName1")</p>  
            <p>Player 2: @Html.TextBox("PlayerName2")</p>  
            <p><input type="submit"  
                    value="Deal Cards"  
                    class="submit">  
            </p>  
        </div>  
    </form>  
}
```

The code renders two HTML `TextBox` controls that request the player names and a button. Once the information is entered, pressing the Deal Cards button executes a `POST` and the following code path is executed. The code loops through the cards dealt to each player of the game.

```

@if (IsPostBack)
{
    <label id="labelGoal">Which player has the best hand.</label>
    <br />
    <div>
        <p><label id="labelPlayer1">Player1: @player1</label></p>
        @foreach (Card card in players[0].PlayHand)
        {
            <img width="75"
                  height="100"
                  alt="cardImage"

            src=
            "https://deckofcards.blob.core.windows.net/carddeck/@card.imageLink" />
        }
    </div>
    <div>
        <p><label id="labelPlayer2">Player2: @player2</label></p>
        @foreach (Card card in players[1].PlayHand)
        {
            <img width="75"
                  height="100"
                  alt="cardImage"

            src=
            "https://deckofcards.blob.core.windows.net/carddeck/@card.imageLink" />
        }
    </div>
}

```

Notice that within both `foreach` loops there is a reference to the Azure storage account URL and the container created in the previous exercise.

**NOTE** The Azure storage account URL and container are for example only.  
 You should replace `deckofcards` with your Azure storage account and `carddeck` with your Azure storage container.

## EXERCISES

- 16.1 What information would you need to pass between the browser and the server to play the card game?
- 16.2 As web applications are stateless, describe some ways to store this information so it can be included with a web request.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
Defining the cloud	The cloud is an elastic structure of commoditized computer hardware for running programs. These programs run on IaaS, PaaS, or SaaS service models in a Hybrid, Public, or Private Cloud type.
Defining the cloud optimized stack	The cloud optimized stack is a concept used to refer to code that can handle high throughput, makes a small footprint, can run side-by-side with other applications on the same server, and is cross-platform enabled.
Creating a storage account	A storage account can contain an infinite number of containers. The storage account is the mechanism for controlling access to the containers created within it.
Creating a storage container with C#	A storage container exists within a storage account and contains the blobs, files, or data that are accessible from any place where an Internet connection exists.
Referencing the storage container from ASP.NET Razor	It is possible to reference a storage container from C# code. You use the storage account name, the container name, and the name of the blob, file, or data you need to access.

# 17

## Advanced Cloud Programming and Deployment

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Creating an ASP.NET Web API
- Deploying and consuming an ASP.NET Web API on Microsoft Azure
- Scaling an ASP.NET Web API on Microsoft Azure

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

The `wrox.com` code downloads for this chapter are found at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 17 download and individually named according to the names throughout the chapter.

Now that you have spent some time learning about the cloud and cloud programming, let's move forward and write some C# code that is a little more complex than what you did in the previous chapter. In this chapter, you continue exploring both ASP.NET and Microsoft Azure. You will modify the `CardLib` program so that it runs in the cloud as an ASP.NET Web API and, once deployed, you consume it from an ASP.NET Web Site.

**NOTE** To successfully complete the exercises in this chapter, you need a Microsoft Azure subscription. If you do not have one, you can sign up for a free trial here: <http://azure.microsoft.com>. It is quick and easy to do.

After creating, deploying, and consuming the ASP.NET Web API, you will learn how to scale it. The concept of scaling is important to grasp in the event that the cloud program you create

becomes popular. The example in this chapter uses free Microsoft Azure cloud resources. These free resources have low CPU, memory, and bandwidth thresholds and are easily breached under high usage. You will learn how to avoid any suspension of your cloud program due to resource threshold breaches by scaling when appropriate.

## CREATING AN ASP.NET WEB API

The Application Programming Interface (API) computer programming concept has been around for many decades and is generally described as a module that contains a set of functions useful for building software programs.

Originally, from a Windows client application perspective, these modules were dynamic linked libraries (.dll) and exposed programmatically accessible interfaces that exposed internal functions to other programs. In such a system, when a consuming program uses an API, it becomes dependent on the pattern of the interface. Changes to the interface cause exceptions and failures within the consuming program because the current procedure to access and execute the functions within the module is no longer valid. Once programs become dependent on an interface, it shouldn't be changed and when it is changed the event is commonly referred to as DLL Hell. For more information about DLL Hell, read this article: <http://www.desaware.com/tech/dllhell.aspx>.

As time moved on and the implementation of Internet and intranet solutions became mainstream, dependencies on technologies such as web services and Windows Communication Foundation (WCF) were made. Both web services and WCF exposed formal contractual interfaces that exposed the functions contained within them to other programs. As opposed to the previously mentioned DLL API where the module exists on the same computer as the one consuming it, the web service and WCF are hosted on a web server. As a result of being hosted on an Internet or intranet web server, access to the web interface is no longer confined to a single computer and is possible from any device, from any place with an Internet or networked intranet connection.

Recall from the previous chapter where the analysis of the cloud optimized stack took place. From the discussion you learned that in order to be considered cloud optimized, a program must have a small footprint, be able to handle high throughput, and be cross-platform enabled. An ASP.NET Web API is based on the ASP.NET MVC (Model, View, Controller) concept, which aligns directly with the new cloud optimized stack definition. If you have created and/or used web services or WCF in the past, you will see how much simpler and compact an ASP.NET Web API is in comparison. If you have never used either, take my word for it: It is.

In the following Try It Out, you will create an ASP.NET Web API that deals a hand of cards.

### TRY IT OUT Create an ASP.NET Web API

You will use Visual Studio 2015 to create an ASP.NET Web API that accepts a player's name and returns a hand of cards for that player.

1. Create a new ASP.NET Web API by selecting File  $\Rightarrow$  New  $\Rightarrow$  Project... within Visual Studio. In the New Project dialog box (see Figure 17-1), select the category Visual C#  $\Rightarrow$  Web and then select the ASP.NET Web Application template. Change the path to c:\BegVCSharp\Chapter17\, name the Web Application Ch17Ex01, and then click the OK button.




FIGURE 17-1

2. Next click on the Empty ASP.NET 4.6 Template and check the Web API checkbox so required folders and core references are added to the project. See Figure 17-2. Deselect the “Host in the cloud” check box for now. (Publishing of the Web API will be done later in this chapter.) Click the OK button.




FIGURE 17-2

3. Right-click on the Ch17Ex01 solution, and then select Add ➔ New Folder ➔ Rename to add a new folder named CardLib.
4. Download the sample code from the download site and place the following class files into the CardLib folder you just created. Once downloaded, right-click on the CardLib folder and select Add ➔ Existing Item... and select the seven classes from the downloaded example.
  - a. Card.cs
  - b. Cards.cs

- c. Deck.cs
- d. Game.cs
- e. Player.cs
- f. Rank.cs
- g. Suit.cs

**NOTE** The seven classes are the same as used in Ch16Ex02. If you have already downloaded the source code for that exercise, then you can reuse them here as well.

5. Next add a controller by right-clicking on the Controllers folder, selecting  $\Rightarrow$  Add  $\Rightarrow$  Controller..., and selecting Web API 2 Controller - Empty...  $\Rightarrow$  Add (Figure 17-3).




FIGURE 17-3

6. Name the controller HandOfCardsController.
7. Add this code to the HandOfCardsController class:

```
[Route("api/HandOfCards/{playerName}")]
public IEnumerable<Card> GetHandOfCards(string playerName)
{
    Player[] players = new Player[1];
    players[0] = new Player(playerName);
    Game newGame = new Game();
    newGame.SetPlayers(players);
    newGame.DealHands();
```

```
    var handOfCards = players[0].PlayHand;
    return handOfCards;
}
```

8. The ASP.NET Web API is now created and ready to be published to the cloud.

Congratulations! You have completed the creation of an ASP.NET Web API that returns a hand of cards.

### How It Works

When you want to create a new ASP.NET Web API, there are two options. The first is the method that this Try It Out described. The fact that you selected Empty from the Template selection window meant that the project would contain nothing other than the very basic necessities required to create an ASP .NET Web API. This resulted in very few configuration files and binaries being added to the solution; the footprint for this Web API is therefore very small and is just what is needed to run optimally in the cloud.

The other possible approach is to select the Web API template (see Figure 17-2) instead of the Empty one. This includes additional configuration files, many additional references, and a basic example of an ASP.NET MVC application. As this Try It Out is relatively small and did not require any MVC features, the Empty template was chosen. If additional functionalities and examples are needed in a future project of your own, consider selecting the Web API template because it constructs data pipelines and provides many proven coding patterns to build your solution on top of.

You add the same seven card classes used in the Chapter 16 example to a directory called CardLib. The contents of the GetHandOfCards() method are identical to those of the one in Chapter 16. The method accepts one parameter, the playerName, creates a new Game, sets the Players, deals the hand of cards, and returns the Cards class to the ASP.NET Web API consumer. The one additional line of code is this:

```
[Route("api/HandOfCards/{playerName}")]
```

The Route annotation is how ASP.NET decides which Web API method responds to which request. As you will come to realize, after publishing there is no specific file requested when you interface with an ASP.NET Web API. Unlike an ASP.NET Web Forms application where a request is sent to a file with an .aspx extension, the same is not true when calling a Web API (or an ASP.NET MVC application for that matter). A Web API request is sent to a web server, where the parameters are in the requested URL, separated by forward slashes. For example: “<http://contoso.com/api/{controllerName}/Parameter1/Parameter2/etc...>”.

**NOTE** Instead of using annotations for creating Route Maps, you can create them in a file called `WebApiConfig.cs` located in the `App_Start` directory.

Now that the ASP.NET Web API is created, move on to the next section to learn about deployment and then consumption of the Web API.

## DEPLOYING AND CONSUMING AN ASP.NET WEB API ON MICROSOFT AZURE

There are numerous options for deploying your Web App to the Microsoft Azure platform. One of the most popular methods is via a local Git repository or a public Git repository hosted on GitHub. Both the local and public Git repositories provide capabilities for version control, which is a very useful feature. Having version control lets the developer and release manager know what specific changes have been made, when they were made, and by whom. In the event that there are problems or unexpected exceptions when the binaries are compiled or the changes are deployed to the live environment, it is easy to find who to contact about it. Other deployment platforms that can be integrated into Microsoft Azure include Team Foundation Services, CodePlex, and Bitbucket, for example.

**NOTE** *There are numerous methods for making deployments of your code to the Microsoft Azure platform. Projects stored in a source code repository versus specific standalone code scenarios each have numerous and individual deployment options.*

As the code in the previous Try It Out is a standalone project that is not contained in a version control repository, the deployment is performed directly within the IDE, in this case Visual Studio 2015. Additional methods for deploying a solution not contained in a source code repository include, for example, Web Deploy (`msdeploy.exe`) and FTP.

Complete the following Try It Out to deploy an ASP.NET Web API to a Microsoft Azure Web App.

### TRY IT OUT Deploy an ASP.NET Web API to the Cloud

1. Right-click Ch17Ex01 project  $\Rightarrow$  Publish....
2. Select Microsoft Azure Web App  $\Rightarrow$  Manage subscriptions and import your Microsoft Azure subscription (if required).
3. Select the New... button, create a new Microsoft Azure Web App, enter the required values, and press the Create button (Figure 17-4).
  - **Web App Name:** Must be a unique name.
  - **Subscription:** If you have multiple Microsoft Azure Subscriptions, select the one that you want this Web App to be created in.
  - **Region:** Select the location where you would like the Web App created.
  - **Database server:** It is possible to create a database during the publishing of the Web App. In this example, a database is not required.




FIGURE 17-4

- Once created, the Publish Web window is rendered (Figure 17-5). Press the Publish button to deploy the ASP.NET Web API to the cloud. Before publishing you might consider pressing Validate Connection button to make sure the configuration and credentials are set up correctly.




FIGURE 17-5

5. After the ASP.NET Web API is successfully published, a browser will open notifying that the web app has been successfully created. You can also view the details of the Output window in Visual Studio to find more information about the publishing steps.
6. Check the response of the ASP.NET Web API. It is now globally accessible via, for example, `http://handofcards.azurewebsites.net/api/HandOfCards/Benjamin`, where handofcards is the name you provided when creating the Microsoft Azure Web App and Benjamin is the name of the player.

**NOTE** By default, different browsers render the results in different ways. For example, Internet Explorer prompts you to download a JSON file, while Chrome displays some XML data. The important aspect is that you get a response. Consuming the API is covered in the next section.

## How It Works

When you publish a Web App from within Visual Studio, it uses Web Deploy in the background to perform the actual deployment. Knowing this, if you have special requirements for the deployment, they can be set within the publish profile located in the `Properties\PublishProfiles` directory. The contents within `*.pubxml` contain the configuration items and dependencies for the given deployment.

Once the deployment completes, a browser is rendered to the main page of the Web App (illustrated by Figure 17-6), and not the ASP.NET Web API.




FIGURE 17-6

Unlike legacy APIs contained in a `.dll`, web service, or WCF service, it is uncommon in practice to access an ASP.NET Web API directly. Rather, in all cases, the call to the API comes from code contained in another (API consuming) project or solution.

Now that the ASP.NET Web API is deployed, it is consumable from any client with capabilities to make an HTTP request and parse a JSON file. The following Try It Out provides all the instructions you need to learn how to consume the just published ASP.NET Web API from an ASP.NET Web Page.

**NOTE** The following Try It Out modifies the Ch16Ex02 ASP.NET Web Site. The primary difference is that instead of retrieving the cards from classes contained in the Web Site itself, the cards are retrieved from the ASP.NET Web API created and deployed in this chapter.

## TRY IT OUT    Consume the Web API from a Web Site

You will use Visual Studio 2015 to modify the CH16Ex02 example so that it consumes an ASP.NET Web API. The Web API accepts a player's name and returns a hand of cards for that player.

1. Open Visual Studio 2015, select File  $\Rightarrow$  New  $\Rightarrow$  Web Site, and select ASP.NET Empty Web Site from the Visual C# list of installed Templates.
2. Change the Web Location to c:\BegVCSharp\Chapter17\Ch17Ex02 and then press the OK button to continue.

**NOTE** The output of an ASP.NET Web API is a JSON file, the format of which follows a standard format making it easily parsed. The most common means for parsing a JSON file is using the `Newtonsoft.Json` libraries.

3. To install the `Newtonsoft.Json` libraries used for parsing the JSON file, right-click on the Solution and select Manage NuGet Packages..., which opens a tab in Visual Studio similar to that shown in Figure 17-7.
4. Select `Newtonsoft.Json` from the Package list and press the Install button. A `Bin` directory is added to the ASP.NET Web Site that contains the `Newtonsoft.Json.dll` binary.
5. Add a `.cshtml` file to the solution by right-clicking on the Ch17Ex02 and selecting Add  $\Rightarrow$  Add New Item...  $\Rightarrow$  Empty Page (Razor v3), name it `default.cshtml`, and press the Add button.

**NOTE** The contents of the `default.cshtml` file here and the one previously created in CH16Ex02 are very similar, but some modifications are required. Consider copying the contents of `default.cshtml` from Chapter 16 instead of retyping the entire page.

6. Next, include the `Newtonsoft.Json` libraries into the Razor file by adding this statement at the very top of the page:

```
@using Newtonsoft.Json;
```




FIGURE 17-7

7. Add this code snippet directly below the line added in step 6:

```

@{
    List<string> cards = new List<string>();
    var playerName = Request["PlayerName"];

    if (IsPost)
    {
        string GetURL = "http://handofcards.azurewebsites.net/api/" +
            "HandOfCards/" + playerName;
        WebClient client = new WebClient();
        Stream dataStream = client.OpenRead(GetURL);
        StreamReader reader = new StreamReader(dataStream);
        var results =
            JsonConvert.DeserializeObject<dynamic>(reader.ReadLine());
        reader.Close();

        foreach (var item in results)
    }
}

```

```
        cards.Add((string)item.imageLink);  
    }  
}
```

8. Lastly, add the markup and Razor code to trigger the consumption of the ASP.NET Web API directly under the code added in step 7:

```

<html>
<head>
    <title>BensCards: Deal yourself a hand. </title>
</head>
<body>
    @if (IsPost)
    {
        <label id="labelGoal">Here is your hand of cards.</label>
        <br />
        <div>
            <p><label id="labelPlayer1">Player1: @playerName</label></p>
            @foreach (string card in cards)
            {
                <img width="75"
                    height="100"
                    alt="cardImage"
                    src=
                    "https://deckofcards.blob.core.windows.net/carddeck/@card" />
            }
        </div>
        <label id="errorMessageLabel" />
    }
    else
    {
        <label id="labelGoal">
            Enter the players name and deal the cards.
        </label>
        <br /><br />
        <form method="post">
            <div>
                <p>Player 1: @Html.TextBox("PlayerName")</p>
                <p><input type="submit" value="Deal Hand" class="submit"></p>
            </div>
        </form>
    }
</body>
</html>

```

9. Run the ASP.NET Web Site by pressing F5. Once rendered, enter a name and press the Deal Hand button. The ASP.NET Web Site consumes the ASP.NET Web API and renders a hand of cards, similar to that shown in Figure 17-8.




FIGURE 17-8

## How It Works

When the `default.cshtml` page is initially rendered, `IsPost` is `false` and therefore the calling of the ASP.NET Web API from the C# code contained in the Razor code block does not execute. Instead, only the portion of HTML code within the `else` code block gets displayed. The rendered portion contains a `TextBox` to capture the player name and a `Button` to trigger the posting of the page back to itself.

Once a player name is entered and the Deal Hand button is pressed, the `IsPost` property becomes `true` and the C# code within the Razor tag at the top of page is executed.

```
string GetURL = "http://handofcards.azurewebsites.net/api/HandOfCards/" +
    playerName;
 WebClient client = new WebClient();
 Stream dataStream = client.OpenRead(GetURL);
```

The web address stored in the `GetURL` string is the Internet or intranet location of the ASP.NET Web API and is used as a parameter for the `OpenRead()` method of the `WebClient` class. The `WebClient` contains the methods required to perform an HTTP request. The result of the `OpenRead()` method is stored in a `Stream` object.

```
StreamReader reader = new StreamReader(dataStream);
var results = JsonConvert.DeserializeObject<dynamic>(reader.ReadLine());
```

The `Stream` object is then passed as a parameter to the `StreamReader` constructor. Using the `ReadLine()` method of the `StreamReader` class as a parameter to deserialize the JSON file using the `Newtonsoft.Json` libraries, the `results` can then be enumerated through a `foreach` statement and added to a `List<string>` container named `cards`. The `cards` list can then be accessed for usage later in the page rendering process.

```

foreach (var item in results)
{
    cards.Add((string) item.imageLink);
}

```

**NOTE** Review the dynamic type discussed previously in Chapter 13. It is common practice to use the dynamic type with JSON files as the structure contained within it is not always castable to a strongly typed class.

Once the parsed results of the JSON file are loaded into the cards container, the markup code within the IsPost code block gets executed. The foreach loop within the Razor tags reads through the cards container and concatenates the image name with the link to the Microsoft Azure Blob Container created in Chapter 16.

```

@foreach (string card in cards)
{
    
}

```

You might consider deploying this ASP.NET Web Site to the Microsoft Azure platform using the acquired knowledge from the previous Try It Out. For example, simply right-click the Ch17Ex02 solution, select Publish Web App, and follow the publish wizard. Creating a Web App called "handofcards-consumer" would then be accessible from <http://handofcards-consumer.azurewebsites.net/>. As both the ASP.NET Web API and the Microsoft Azure Blob Container are accessible on the Internet from any place in the world, placing the ASP.NET Web Site on Azure would result in the same outcome (getting a hand of cards).

Over time, if either the consumer or the API become popular and begin receiving many requests, running the Web Apps in FREE mode would likely result in a resource threshold breach that renders the resources unusable. This would not be ideal. In the next section, you learn how to scale an ASP .NET Web API running as a Web App on the Microsoft Azure platform so that users and customers can access your responsive web resource when required.

## SCALING AN ASP.NET WEB API ON MICROSOFT AZURE

Scaling to meet the requirements of your users used to be a very tedious, time-consuming, and expensive activity. Historically, when a company wanted to increase server capacity to support more traffic, it required the acquisition, assembly, and configuration of physical hardware into a data center. Then, once the hardware was on the network, it was handed over to the application owners to install and configure the operating system, the required components, and the application source code. The time

required to perform such tasks resulted in companies installing a lot of physical capacity to manage peak time usage; however, during times of non-peak usage, the extra capacity simply went unused and sat idle, which is a very expensive and non-optimal way to allocate resources.

A better approach is to use cloud platforms, like Microsoft Azure, that provide the ability to optimally utilize physical resources to scale up, down, and out during the times when the resources are required. When you need physical resources like CPU, disk space or memory, you scale up or out to meet the demands, and when the demand for your cloud-hosted services reduces, you can scale back down and save your financial resources for use with other projects and services.

**NOTE** To successfully complete the exercises in this chapter, you need a Microsoft Azure subscription. If you do not have one, you can sign up for a free trial here: <http://azure.microsoft.com>. It is quick and easy to do.

The remainder of this chapter illustrates how to scale an ASP.NET Web API based on CPU demand and during a specific time frame.

### TRY IT OUT Scale an ASP.NET Web API Based on CPU Usage

1. Access the Microsoft Azure portal at <https://manage.windowsazure.com>.
2. Select the ASP.NET Web API you created earlier in this chapter, for example “handofcards.” As shown in Figure 17-9, notice that the Web App is in the FREE pricing tier. Auto Scaling is only available when the Web App is in STANDARD mode.




FIGURE 17-9

3. Scale your Web App up to STANDARD by clicking on the STANDARD Tier box and then the Save button at the bottom of the page.
4. Once the configuration is saved, scroll down and the options for capacity scaling are rendered (see Figure 17-10). Click on the CPU SCALE BY METRIC setting.




FIGURE 17-10

5. Change the max INSTANCE COUNT to 5 and the TARGET CPU to 50 and 80 respectively.
6. Save the configuration.

### How It Works

Running a Web App in FREE mode does not get you a whole lot. It is really for testing and learning how the Microsoft Azure platform works. Auto Scaling is only available in STANDARD mode and therefore you must scale to this tier to have access to this feature. Other modes like SHARED and BASIC have the capacity to scale but require a manual configuration to do so.

By default, the Auto Scale settings are to scale up to a maximum of 3 instances of this Web App when the CPU utilization averages between 60% and 80% for a given 60 minute time period examined every 5 minutes.

Notice the INSTANCE SIZE drop down in Figure 17-10. By default, the instance size is Small, which equates to 1 x 2.6 GHZ CPU and 1.75GB of memory. This means that when you scale to three instances, you receive three different virtual machines each with 1 x 2.6GHZ CPU and 1.75GB of memory. Had you set the INSTANCE SIZE to Large, you would instead get three virtual machines with 4 x 2.6GHZ CPUs and 7GB of memory each.

Finally, when the utilization of the CPU or CPUs on the virtual machines running your Web App breaches the lower threshold set as the TARGET CPU value, a new instance or virtual machine is added to your environment up to the maximum number set by the INSTANCE COUNT.

---

The auto scaling feature is very useful for managing unexpected peaks of usage and requests to your Web App. However, if you already know when your customers or users interact with your Web App, you can plan ahead and have the additional instances available slightly before they are actually needed. The benefit is that instead of a gradual increase or decrease of instances based on CPU usage, you can scale immediately to the number of CPUs and amount of memory required during only that specific timeframe. For example, if you know that your marketing department is running a campaign during the month of October, you can schedule additional resources to be available during that month. By having the required resources available and warmed up, you can avoid any delay in getting them allocated for use by your users or customers. Perform the steps described in the following Try It Out to see how.

#### TRY IT OUT Scale an ASP.NET Web API at a specific time

1. Access the Microsoft Azure portal at <https://manage.windowsazure.com>.
2. Select the ASP.NET Web API you created earlier in this chapter, for example "handofcards". As shown previously in Figure 17-9, notice that the Web App is in the FREE pricing tier. Auto Scaling is only available when the Web App is in STANDARD mode.
3. Scale your Web App up to STANDARD by clicking on the STANDARD Tier box and then the Save button at the bottom of the page.
4. Once the configuration is saved, scroll down and the options for scaling are rendered. Click on the "set up schedule times" link and you are presented with a pop-up, as illustrated by Figure 17-11.
5. Enter for example the name, dates, and times similar to those shown in Figure 17-12.
6. Press the check mark on the pop-up window, as shown in Figure 17-11.
7. Select the name of the scheduled scale profile from the EDIT SCALE SETTINGS FOR SCHEDULE dropdown. For example, select October, as shown in Figure 17-13. Then, select the number of INSTANCES for the selected schedule profile, for example 5.

8. Click the SAVE button at the bottom of the page, and when the configured time frame becomes current, the scale setting will take effect.




FIGURE 17-11

SPECIFIC DATES				
NAME	START AT	START TIME	END AT	END TIME
October	2015-10-01	09:00 AM	2015-10-31	06:00 PM
	NAME	YYYY-MM-DD	HH:MM AM/PM	YYYY-MM-DD

FIGURE 17-12




FIGURE 17-13

When you create a schedule for scaling your Web App, a name, start date, start time, end date, and end time are required. With this information, the Microsoft Azure platform manages the number of available instances, which are virtual machines that serve requests to your Web App during the configured time frame. It is possible to create numerous schedules, each having its own number of instances and scale settings. Simply create the schedule, save it, and when it's needed select it from the schedule drop down, and the resources will become available as expected.

## EXERCISES

- 17.1** Instead of consuming the ASP.NET Web API from an ASP.NET Web Site application, try consuming it from another program type like a console application or a Windows Universal App.
  
- 17.2** What is the maximum size and number of instances you can have for a Web App on the Microsoft Azure platform?

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>ASP.NET Web API</b>	An ASP.NET Web API is an Internet or intranet interface that exposes methods for consumption from external programs.
<b>Deploying to the cloud</b>	Use tools like Visual Studio, WebDeploy, Git, or FTP to deploy your program to the cloud.
<b>Consuming a Web API</b>	An ASP.NET Web API returns the output of the method in a JSON file. Use the <code>Newtonsoft.Json</code> class library to parse and use its content.
<b>Scaling in the cloud</b>	Microsoft Azure Web Apps let you auto scale based on a defined schedule or CPU usage. Being able to scale up when you need more of a resource and down when it is no longer needed is one of the most valuable benefits of the cloud.



# PART IV

# Data Access

---

- **CHAPTER 18:** Files
- **CHAPTER 19:** XML and JSON
- **CHAPTER 20:** LINQ
- **CHAPTER 21:** Databases



# 18

## Files

---

### WHAT YOU WILL LEARN IN THIS CHAPTER

---

- Discovering the `File` and `Directory` classes
- Understanding how .NET uses streams to access files
- Writing to and reading from a file
- Reading and writing compressed files
- Serializing and deserializing objects
- Monitoring files and directories for changes

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 18 download and program names match the names used in the examples throughout the chapter.

Files can be a great way to store data between instances of your application, or they can be used to transfer data between applications. User and application configuration settings can be stored to be retrieved the next time your application is run.

This chapter shows you how to use files effectively in your applications, touching on the major classes used to create, read from, and write to files, and the supporting classes used to manipulate the file system from C# code. Although you won't examine all of the classes in detail, this chapter goes into enough depth to give you a good idea of the concepts and fundamentals.

## FILE CLASSES FOR INPUT AND OUTPUT

Reading and writing files is an essential way to get data into your C# program (*input*) and send data out of your program (*output*). Because files are used for input and output, the file classes are contained in the `System.IO` namespace. (*IO* is a common abbreviation for *Input/Output*.)

`System.IO` contains the classes for reading and writing data to and from files, and you can reference this namespace in your C# application to gain access to these classes without fully qualifying type names.

The classes covered in this chapter are described in Table 18-1.

TABLE 18-1: File System Access Classes

CLASS	DESCRIPTION
<code>File</code>	A static utility class that exposes many static methods for moving, copying, and deleting files.
<code>Directory</code>	A static utility class that exposes many static methods for moving, copying, and deleting directories.
<code>Path</code>	A utility class used to manipulate path names.
<code>FileInfo</code>	Represents a physical file on disk, and has methods to manipulate this file. For any reading from and writing to the file, a <code>Stream</code> object must be created.
<code> DirectoryInfo</code>	Represents a physical directory on disk and has methods to manipulate this directory.
<code>FileSystemInfo</code>	Serves as the base class for both <code>FileInfo</code> and <code> DirectoryInfo</code> , making it possible to deal with files and directories at the same time using polymorphism.
<code> FileSystemWatcher</code>	The most advanced class you examine in this chapter. It is used to monitor files and directories, and it exposes events that your application can catch when changes occur in these locations.

You'll also look at the `System.IO.Compression` namespace, which enables you to read from and write to compressed files. In particular, you will look at the following two stream classes:

- `DeflateStream` — Represents a stream in which data is compressed automatically when writing, or uncompressed automatically when reading. Compression is achieved using the Deflate algorithm.
- `GZipStream` — Represents a stream in which data is compressed automatically when writing, or uncompressed automatically when reading. Compression is achieved using the GZIP (GNU Zip) algorithm.

## The File and Directory Classes

The `File` and `Directory` utility classes expose many static methods for manipulating, surprisingly enough, files and directories. These methods make it possible to move files, query and update attributes, and create `FileStream` objects. As you learned in Chapter 8, static methods can be called on classes without having to create instances of them.

Some of the most useful static methods of the `File` class are shown in the Table 18-2.

**TABLE 18-2:** Static Methods of the File Class

METHOD	DESCRIPTION
<code>Copy()</code>	Copies a file from a source location to a target location.
<code>Create()</code>	Creates a file in the specified path.
<code>Delete()</code>	Deletes a file.
<code>Open()</code>	Returns a <code>FileStream</code> object at the specified path.
<code>Move()</code>	Moves a specified file to a new location. You can specify a different name for the file in the new location.

Some useful static methods of the `Directory` class are shown in Table 18-3.

**TABLE 18-3:** Static Methods of the Directory Class

METHOD	DESCRIPTION
<code>CreateDirectory()</code>	Creates a directory with the specified path.
<code>Delete()</code>	Deletes the specified directory and all the files within it.
<code>GetDirectories()</code>	Returns an array of <code>string</code> objects that represent the names of the directories below the specified directory.
<code>EnumerateDirectories()</code>	Like <code>GetDirectories()</code> , but returns an <code>IEnumerable&lt;string&gt;</code> collection of directory names.
<code>.GetFiles()</code>	Returns an array of <code>string</code> objects that represent the names of the files in the specified directory.
<code>EnumerateFiles()</code>	Like <code>GetFiles()</code> , but returns an <code>IEnumerable&lt;string&gt;</code> collection of filenames.
<code>GetFileSystemEntries()</code>	Returns an array of <code>string</code> objects that represent the names of the files and directories in the specified directory.

*continues*

TABLE 18-3 (continued)

METHOD	DESCRIPTION
EnumerateFileSystemEntries()	Like GetFileSystemEntries(), but returns an <code>IEnumerable&lt;string&gt;</code> collection of file and directory names.
Move()	Moves the specified directory to a new location. You can specify a new name for the folder in the new location.

The three `EnumerateXxx()` methods provide better performance than their `GetXxx()` counterparts when a large amount of files or directories exist.

## The FileInfo Class

Unlike the `File` class, the `FileInfo` class is not static and does not have static methods. This class is useful only when instantiated. A `FileInfo` object represents a file on a disk or a network location, and you can create one by supplying a path to a file:

```
FileInfo aFile = new FileInfo(@"C:\Log.txt");
```

**NOTE** You will be working with strings representing the path of a file throughout this chapter, which means a lot of \ characters in your strings. Therefore, you should remember that you can precede a string value with @, which means that the string will be interpreted literally. Thus, \ will be interpreted as \, and not as an escape character. Without the @ prefix, you need to use \\ instead of \ to avoid having this character be interpreted as an escape character. In this chapter you'll stick to the @ prefix for your strings.

You can also pass the name of a directory to the `FileInfo` constructor, although in practical terms that isn't particularly useful. Doing this causes the base class of `FileInfo`, which is `FileSystemInfo`, to be initialized with all the directory information, but none of the `FileInfo` methods or properties relating specifically to files will work.

Many of the methods exposed by the `FileInfo` class are similar to those of the `File` class, but because `File` is a static class, it requires a string parameter that specifies the file location for every method call. Therefore, the following calls do the same thing:

```
FileInfo aFile = new FileInfo("Data.txt");
if (aFile.Exists)
    WriteLine("File Exists");
if (File.Exists("Data.txt"))
    WriteLine("File Exists");
```

In this code, a check is made to see whether the file `Data.txt` exists. Note that no directory information is specified here, which means that the current *working directory* is the only location

examined. This directory is the one containing the application that calls this code. You'll look at this in more detail a little later, in the section "Path Names and Relative Paths."

Most of the `FileInfo` methods mirror the `File` methods in this manner. In most cases it doesn't matter which technique you use, although the following criteria can help you to decide which is more appropriate:

- It makes sense to use methods on the static `File` class if you are making only a single method call — the single call will be faster because the .NET Framework won't have to go through the process of instantiating a new object and then calling the method.
- If your application is performing several operations on a file, then it makes more sense to instantiate a `FileInfo` object and use its methods — this saves time because the object will already be referencing the correct file on the file system, whereas the static class has to find it every time.

The `FileInfo` class also exposes properties relating to the underlying file, some of which can be manipulated to update the file. Many of these properties are inherited from `FileSystemInfo`, and thus apply to both the `FileInfo` and `DirectoryInfo` classes. The properties of `FileSystemInfo` are shown in Table 18-4.

TABLE 18-4: `FileSystemInfo` Properties

PROPERTY	DESCRIPTION
<code>Attributes</code>	Gets or sets the attributes of the current file or directory, using the <code>FileAttributes</code> enumeration.
<code>CreationTime</code> , <code>CreationTimeUtc</code>	Gets or sets the creation date and time of the current file, available in coordinated universal time (UTC) and non-UTC versions.
<code>Extension</code>	Retrieves the extension of the file. This property is read-only.
<code>Exists</code>	Determines whether a file exists. This is a read-only abstract property, and is overridden in <code>FileInfo</code> and <code>DirectoryInfo</code> .
<code>FullName</code>	Retrieves the full path of the file. This property is read-only.
<code>LastAccessTime</code> , <code>LastAccessTimeUtc</code>	Gets or sets the date and time that the current file was last accessed, available in UTC and non-UTC versions.
<code>LastWriteTime</code> , <code>LastWriteTimeUtc</code>	Gets or sets the date and time that the current file was last written to, available in UTC and non-UTC versions.
<code>Name</code>	Retrieves the full path of the file. This is a read-only abstract property, and is overridden in <code>FileInfo</code> and <code>DirectoryInfo</code> .

The properties specific to `FileInfo` are shown in Table 18-5.

TABLE 18-5: `FileInfo` Properties

PROPERTY	DESCRIPTION
<code>Directory</code>	Retrieves a <code> DirectoryInfo</code> object representing the directory containing the current file. This property is read-only.
<code>DirectoryName</code>	Returns the path to the file's directory. This property is read-only.
<code>IsReadOnly</code>	Shortcut to the read-only attribute of the file. This property is also accessible via <code>Attributes</code> .
<code>Length</code>	Gets the size of the file in bytes, returned as a <code>long</code> value. This property is read-only.

## The `DirectoryInfo` Class

The  `DirectoryInfo` class works exactly like the  `FileInfo` class. It is an instantiated object that represents a single directory on a machine. Like the  `FileInfo` class, many of the method calls are duplicated across  `Directory` and  `DirectoryInfo`. The guidelines for choosing whether to use the methods of  `File` or  `FileInfo` also apply to  `DirectoryInfo` methods:

- If you are making a single call, use the static  `Directory` class.
- If you are making a series of calls, use an instantiated  `DirectoryInfo` object.

The  `DirectoryInfo` class inherits most of its properties from  `FileSystemInfo`, as does  `FileInfo`, although these properties operate on directories instead of files. There are also two  `DirectoryInfo`-specific properties, shown in Table 18-6.

TABLE 18-6: Properties Unique to the  `DirectoryInfo` Class

PROPERTY	DESCRIPTION
<code>Parent</code>	Retrieves a <code> DirectoryInfo</code> object representing the directory containing the current directory. This property is read-only.
<code>Root</code>	Retrieves a <code> DirectoryInfo</code> object representing the root directory of the current volume — for example, the <code>C:\</code> directory. This property is read-only.

## Path Names and Relative Paths

When specifying a path name in .NET code, you can use absolute or relative path names. An *absolute* path name explicitly specifies a file or directory from a known location — such as the `C:` drive.

---

An example of this is C:\Work\LogFile.txt — this path defines exactly where the file is, with no ambiguity.

*Relative* path names are relative to a starting location. By using relative path names, no drive or known location needs to be specified. You saw this earlier, where the current working directory was the starting point, which is the default behavior for relative path names. For example, if your application is running in the C:\Development\FileDialog directory and uses the relative path LogFile.txt, the file references would be C:\Development\FileDialog\LogFile.txt. To move “up” a directory, the .. string is used. Thus, in the same application, the path ..\Log.txt points to the file C:\Development\Log.txt.

As shown earlier, the working directory is initially set to the directory in which your application is running. When you are developing with Visual Studio, this means the application is several directories beneath the project folder you created. It is usually located in *ProjectName\bin\Debug*. To access a file in the root folder of the project, then, you have to move up *two* directories with ..\..\..\. You will see this happen often throughout the chapter.

Should you need to, you can determine the working directory by using `Directory.GetCurrentDirectory()`, or you can set it to a new path by using `Directory.SetCurrentDirectory()`.

## STREAMS

All input and output in the .NET Framework involves the use of *streams*. A stream is an abstract representation of a *serial device*. A serial device is something that stores and/or accesses data in a linear manner, that is, one byte at a time, sequentially. This device can be a disk file, a network channel, a memory location, or any other object that supports linear reading, writing, or both. By keeping the device abstract, the underlying destination/source of the stream can be hidden. This level of abstraction enables code reuse, and enables you to write more generic routines because you don’t have to worry about the specifics of how data transfer actually occurs. Therefore, similar code can be transferred and reused when the application is reading from a file input stream, a network input stream, or any other kind of stream. Because you can ignore the physical mechanics of each device, you don’t need to worry about, for example, hard disk heads or memory allocation when dealing with a file stream.

A stream can represent almost any source such as a keyboard, a physical disk file, a network location, a printer, or even another program, but this chapter focuses on reading and writing disk files. The concepts applied to reading/writing disk files apply to most devices, so you’ll gain a basic understanding of streams and learn a proven approach that can be applied to many situations.

## Classes for Using Streams

The classes for using streams are contained in the same `System.IO` namespace along with the `File` and `Directory` classes. These classes are listed in Table 18-7.

**TABLE 18-7:** Stream Classes

CLASS	DESCRIPTION
FileStream	Represents a file that can be written to, read from, or both. This file can be written to and read from asynchronously or synchronously.
StreamReader	Reads character data from a stream and can be created by using a FileStream as a base.
StreamWriter	Writes character data to a stream and can be created by using a FileStream as a base.

Let's look now at how to use each of these classes.

## The FileStream Object

The `FileStream` object represents a stream pointing to a file on a disk or a network path. Although the class does expose methods for reading and writing bytes from and to the files, most often you will use a `StreamReader` or `StreamWriter` to perform these functions. That's because the `FileStream` class operates on bytes and byte arrays, whereas the `Stream` classes operate on character data. Character data is easier to work with, but certain operations, such as random file access (access to data at some point in the middle of a file), can be performed only by a `FileStream` object. You'll learn more about this later in the chapter.

There are several ways to create a `FileStream` object. The constructor has many different overloads, but the simplest takes just two arguments: the filename and a `FileMode` enumeration value:

```
FileStream aFile = new FileStream(filename, FileMode.<Member>);
```

The `FileMode` enumeration has several members that specify how the file is opened or created. You'll see the possibilities shortly. Another commonly used constructor is as follows:

```
FileStream aFile =
new FileStream(filename, FileMode.<Member>, FileAccess.<Member>);
```

The third parameter is a member of the `FileAccess` enumeration and is a way of specifying the purpose of the stream. The members of the `FileAccess` enumeration are shown in Table 18-8.

**TABLE 18-8:** FileAccess Enumeration Members

MEMBER	DESCRIPTION
Read	Opens the file for reading only
Write	Opens the file for writing only
ReadWrite	Opens the file for reading or writing

Attempting to perform an action other than that specified by the `FileAccess` enumeration member will result in an exception being thrown. This property is often used as a way to vary user access to the file based on the user's authorization level.

In the version of the `FileStream` constructor that doesn't use a `FileAccess` enumeration parameter, the default value is used, which is `FileAccess.ReadWrite`.

The  `FileMode` enumeration members are shown in Table 18-9. What actually happens when each of these values is used depends on whether the filename specified refers to an existing file. Note that the entries in this table refer to the position in the file that the stream points to when it is created, a topic you'll learn more about in the next section. Unless otherwise stated, the stream points to the beginning of a file.

TABLE 18-9: FileMode Enumeration Members

MEMBER	FILE EXISTS BEHAVIOR	NO FILE EXISTS BEHAVIOR
Append	The file is opened, with the stream positioned at the end of the file. Can be used only in conjunction with <code>FileAccess.Write</code> .	A new file is created. Can be used only in conjunction with <code>FileAccess.Write</code> .
Create	The file is destroyed, and a new file is created in its place.	A new file is created.
CreateNew	An exception is thrown.	A new file is created.
Open	The file is opened, with the stream positioned at the beginning of the file.	An exception is thrown.
OpenOrCreate	The file is opened, with the stream positioned at the beginning of the file.	A new file is created.
Truncate	The file is opened and erased. The stream is positioned at the beginning of the file. The original file creation date is retained.	An exception is thrown.

Both the `File` and `FileInfo` classes expose `OpenRead()` and `OpenWrite()` methods that make it easier to create `FileStream` objects. The first opens the file for read-only access, and the second allows write-only access. These methods provide shortcuts, so you do not have to provide all the information required in the form of parameters to the `FileStream` constructor. For example, the following line of code opens the `Data.txt` file for read-only access:

```
FileStream aFile = File.OpenRead("Data.txt");
```

The following code performs the same function:

```
FileInfo aFileInfo = new FileInfo("Data.txt");
FileStream aFile = aFileInfo.OpenRead();
```

## File Position

The `FileStream` class maintains an internal file pointer that points to the location within the file where the next read or write operation will occur. In most cases, when a file is opened, it points to the beginning of the file, but this pointer can be modified. This enables an application to read or write anywhere within the file, which in turn enables random access to a file and the capability to jump directly to a specific location in the file. This can save a lot of time when dealing with very large files because you can instantly move to the location you want.

The method that implements this functionality is the `Seek()` method, which takes two parameters. The first parameter specifies how far to move the file pointer, in bytes. The second parameter specifies where to start counting from, in the form of a value from the `SeekOrigin` enumeration. The `SeekOrigin` enumeration contains three values: `Begin`, `Current`, and `End`.

For example, the following line would move the file pointer to the eighth byte in the file, starting from the very first byte in the file:

```
aFile.Seek(8, SeekOrigin.Begin);
```

The following line would move the file pointer two bytes forward, starting from the current position. If this were executed directly after the previous line, then the file pointer would now point to the tenth byte in the file:

```
aFile.Seek(2, SeekOrigin.Current);
```

When you read from or write to a file, the file pointer changes as well. After you have read 10 bytes, the file pointer will point to the byte after the tenth byte read.

You can also specify negative seek positions, which could be combined with the `SeekOrigin.End` enumeration value to seek near the end of the file. The following seeks to the fifth byte from the end of the file:

```
aFile.Seek(-5, SeekOrigin.End);
```

Files accessed in this manner are sometimes referred to as *random access files* because an application can access any position within the file. The `StreamReader` and `StreamWriter` classes described later access files sequentially and do not allow you to manipulate the file pointer in this way.

## Reading Data

Reading data using the `FileStream` class is not as easy as using the `StreamReader` class, which you will look at later in this chapter. That's because the `FileStream` class deals exclusively with raw bytes. Working in raw bytes makes the `FileStream` class useful for any kind of data file, not just text files. By reading byte data, the `FileStream` object can be used to read files such as images or sound files. The cost of this flexibility is that you cannot use a `FileStream` to read data directly into

a string as you can with the `StreamReader` class. However, several conversion classes make it fairly easy to convert byte arrays into character arrays, and vice versa.

The `FileStream.Read()` method is the primary means to access data from a file that a `FileStream` object points to. This method reads the data from a file and then writes this data into a byte array. There are three parameters, the first being a byte array passed in to accept data from the `FileStream` object. The second parameter is the position in the byte array to begin writing data to — this is normally zero, to begin writing data from the file at the beginning of the array. The last parameter specifies how many bytes to read from the file.

The following Try It Out demonstrates reading data from a random access file. The file you will read from is actually the class file you create for the example.

### TRY IT OUT    Reading Data from Random Access Files: ReadFile\Program.cs

1. Create a new console application called `ReadFile` and save it in the directory `C:\BegVCSharp\Chapter18`.
2. Add the following using directives to the top of the `Program.cs` file:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.IO;
```

3. Add the following code to the `Main()` method:

```
static void Main(string[] args)
{
    byte[] byteData = new byte[200];
    char[] charData = new char[200];
    try
    {
        FileStream aFile = new FileStream("../Program.cs", FileMode.Open);
        aFile.Seek(174, SeekOrigin.Begin);
        aFile.Read(byteData, 0, 200);
    }
    catch(IOException e)
    {
        WriteLine("An IO exception has been thrown!");
        WriteLine(e.ToString());
        ReadKey();
        return;
    }
    Decoder d = Encoding.UTF8.GetDecoder();
    d.GetChars(byteData, 0, byteData.Length, charData, 0);
    WriteLine(charData);
    ReadKey();
}
```

4. Run the application. The result is shown in Figure 18-1.




FIGURE 18-1

## How It Works

This application opens its own .cs file to read from. It does so by navigating two directories up the file structure with the .. string in the following line:

```
FileStream aFile = new FileStream("../Program.cs", FileMode.Open);
```

The two lines that implement the actual seeking and reading from a specific point in the file are as follows:

```
aFile.Seek(174, SeekOrigin.Begin);
aFile.Read(byteData, 0, 200);
```

The first line moves the file pointer to byte number 174 in the file. This is the n of namespace in the Program.cs file; the 174 characters preceding it are the using directives. The second line reads the next 200 bytes into the byte array byteData.

Note that these two lines were enclosed in try...catch blocks to handle any exceptions that are thrown:

```
try
{
    aFile.Seek(113, SeekOrigin.Begin);
    aFile.Read(byteData, 0, 100);
}
catch(IOException e)
{
    WriteLine("An IO exception has been thrown!");
    WriteLine(e.ToString());
    ReadKey();
    return;
}
```

Almost all operations involving file I/O can throw an exception of type IOException. All production code should contain error handling, especially when dealing with the file system. The examples in this chapter all include a basic form of error handling.

Once you have the byte array from the file, you need to convert it into a character array so that you can display it to the Console. To do this, use the Decoder class from the System.Text namespace. This class is designed to convert raw bytes into more useful items, such as characters:

```
Decoder d = Encoding.UTF8.GetDecoder();
d.GetChars(byteData, 0, byteData.Length, charData, 0);
```

These lines create a `Decoder` object based on the UTF-8 encoding schema, which is the Unicode encoding schema. Then the `GetChars()` method is called, which takes an array of bytes and converts it to an array of characters. After that has been done, the character array can be written to the Console.

## Writing Data

The process for writing data to a random access file is very similar; a byte array must be created. The easiest way to do this is to first build the character array you want to write to the file. Next, use the `Encoder` object to convert it to a byte array, very much as you used the `Decoder` object. Last, call the `Write()` method to send the array to the file.

Here's a simple example to demonstrate how this is done.

### TRY IT OUT Writing Data to Random Access Files: WriteFile\Program.cs

1. Create a new console application called `WriteFile` and save it in the directory `C:\BegVCSharp\Chapter18`.
2. Add the following `using` directive to the top of the `Program.cs` file:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.IO;
```

3. Add the following code to the `Main()` method:

```
static void Main(string[] args)
{
    byte[] byteData;
    char[] charData;
    try
    {
        FileStream aFile = new FileStream("Temp.txt", FileMode.Create);
        charData = "My pink half of the drainpipe.".ToCharArray();
        byteData = new byte[charData.Length];
        Encoder e = Encoding.UTF8.GetEncoder();
        e.GetBytes(charData, 0, charData.Length, byteData, 0, true);
        // Move file pointer to beginning of file.
        aFile.Seek(0, SeekOrigin.Begin);
        aFile.Write(byteData, 0, byteData.Length);
    }
    catch (IOException ex)
    {
        WriteLine("An IO exception has been thrown!");
        WriteLine(ex.ToString());
        ReadKey();
        return;
    }
}
```

4. Run the application. It should run briefly and then close.
5. Navigate to the application directory — the file will have been saved there because you used a relative path. This is located in the `WriteFile\bin\Debug` folder. Open the `Temp.txt` file. You should see text in the file, as shown in Figure 18-2.




FIGURE 18-2

## How It Works

This application opens a file in its own directory and writes a simple string to it. In structure, this example is very similar to the previous example, except you use `Write()` instead of `Read()`, and `Encoder` instead of `Decoder`.

The following line creates a character array by using the `ToCharArray()` method of the `String` class. Because everything in C# is an object, the text "My pink half of the drainpipe." is actually a `string` object (albeit a slightly odd one), so these static methods can be called even on a string of characters:

```
CharData = "My pink half of the drainpipe.".ToCharArray();
```

The following lines show how to convert the character array to the correct byte array needed by the `FileStream` object:

```
Encoder e = Encoding.UTF8.GetEncoder();
e.GetBytes(charData, 0, charData.Length, byteData, 0, true);
```

This time, an `Encoder` object is created based on the UTF-8 encoding. You used Unicode for the decoding as well, and this time you need to encode the character data into the correct byte format before you can write to the stream. The `GetBytes()` method is where the magic happens. It converts the character array to the byte array. It accepts a character array as the first parameter (`charData` in this example), and the index to start in that array as the second parameter (0 for the start of the array). The third parameter is the number of characters to convert (`charData.Length` — the number of elements in the `charData` array). The fourth parameter is the byte array to place the data into (`byteData`), and the fifth parameter is the index to start writing from in the byte array (0 for the start of the `byteData` array).

The sixth, and final, parameter determines whether the `Encoder` object should flush its state after completion. This reflects the fact that the `Encoder` object retains an in-memory record of where it was in the byte array. This aids in subsequent calls to the `Encoder` object but is meaningless when only a single call is made. The final call to the `Encoder` must set this parameter to `true` to clear its memory and free the object for garbage collection.

After that, it is a simple matter of writing the byte array to the `FileStream` by using the `Write()` method:

```
aFile.Seek(0, SeekOrigin.Begin);
aFile.Write(byteData, 0, byteData.Length);
```

Like the `Read()` method, the `Write()` method has three parameters: a byte array containing the data to write to the file stream, the index in the array to start writing from, and the number of bytes to write.

---

## The StreamWriter Object

Working with arrays of bytes is not most people's idea of fun — having worked with the `FileStream` object, you might be wondering whether there is an easier way. Fear not, for once you have a `FileStream` object, you will usually create a `StreamWriter` or `StreamReader` and use its methods to manipulate the file. If you don't need the capability to change the file pointer to any arbitrary position, these classes make working with files much easier.

The `StreamWriter` class enables you to write characters and strings to a file, with the class handling the underlying conversions and writing to the `FileStream` object for you.

There are many ways to create a `StreamWriter` object. If you already have a `FileStream` object, then you can use it to create a `StreamWriter`:

```
FileStream aFile = new FileStream("Log.txt", FileMode.CreateNew);
StreamWriter sw = new StreamWriter(aFile);
```

A `StreamWriter` object can also be created directly from a file:

```
StreamWriter sw = new StreamWriter("Log.txt", true);
```

This constructor takes the filename and a Boolean value that specifies whether to append to the file or create a new one:

- If this is set to `false`, then a new file is created or the existing file is truncated and then opened.
- If it is set to `true`, then the file is opened and the data is retained. If there is no file, then a new one is created.

Unlike creating a `FileStream` object, creating a `StreamWriter` does not provide you with a similar range of options — other than the Boolean value to append or create a new file, you have no option for specifying the  `FileMode` property as you did with the `FileStream` class. Nor do you have an option to set the  `FileAccess` property, so you will always have read/write privileges to the file. To use any of the advanced parameters, you must first specify them in the `FileStream` constructor and then create a `StreamWriter` from the `FileStream` object, as you do in the following Try It Out.

### TRY IT OUT Writing Data to an Output Stream: StreamWrite\Program.cs

1. Create a new console application called `StreamWrite` and save it in the directory `C:\BegVCSharp\Chapter18`.
2. You will be using the `System.IO` namespace again, so add the following using directives near the top of the `Program.cs` file:

```
using System;
using System.Collections.Generic;
```

```
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.IO;
```

3. Add the following code to the Main() method:

```
static void Main(string[] args)
{
    try
    {
        FileStream aFile = new FileStream("Log.txt", FileMode.OpenOrCreate);
        StreamWriter sw = new StreamWriter(aFile);
        bool truth = true;
        // Write data to file.
        sw.WriteLine("Hello to you.");
        sw.Write($"It is now {DateTime.Now.ToString("MM/dd/yyyy")}");
        sw.Write("and things are looking good.");
        sw.Write("More than that,");
        sw.Write($" it's {truth} that C# is fun.");
        sw.Close();
    }
    catch(IOException e)
    {
        WriteLine("An IO exception has been thrown!");
        WriteLine(e.ToString());
        ReadLine();
        return;
    }
}
```

4. Build and run the project. If no errors are found, it should quickly run and close. Because you are not displaying anything on the console, it is not a very exciting program to watch.
5. Go to the application directory and find the Log.txt file. It is located in the StreamWrite\bin\Debug folder because you used a relative path.
6. Open the file. You should see the text shown in Figure 18-3.




FIGURE 18-3

## How It Works

This simple application demonstrates the two most important methods of the `StreamWriter` class, `Write()` and `WriteLine()`. Both of them have many overloaded versions for performing more advanced file output, but you used basic string output in this example.

The `WriteLine()` method writes the string passed to it, followed immediately by a newline character. You can see in the example that this causes the next write operation to begin on a new line:

```
sw.WriteLine("Hello to you.");
```

The `Write()` method simply writes the string passed to it to the file, without a newline character appended, enabling you to write a complete sentence or paragraph using more than one `Write()` statement. Just as you can write formatted data to the console, you can also write formatted data to files. For example, you can write out the value of variables to the file using interpolated string parameters:

```
sw.Write($"It is now {DateTime.Now.ToString()}");
```

`DateTime.Now` holds the current date; the `ToLongDateString()` method is used to convert this date into an easy-to-read form.

```
sw.Write("More than that,");
sw.Write(" it's {truth} that C# is fun.");
```

Again, you use interpolated string parameters, this time with `Write()` to display the Boolean value `truth` — you set this variable to `true` earlier, and its value is automatically converted into the string “True” for the formatting.

You can use `Write()` and format parameters to write comma-separated files:

```
[StreamWriter object].Write(${100}, {"A nice product"}, {10.50});
```

In a more sophisticated example, this data could come from a database or other data source.

## The StreamReader Object

Input streams are used to read data from an external source. Often, this will be a file on a disk or network location, but remember that this source could be almost anything that can send data, such as a network application or even the Console.

The `StreamReader` class is the one that you will be using to read data from files. Like the `StreamWriter` class, this is a generic class that can be used with any stream. In the next Try It Out, you again construct it around a `FileStream` object so that it points to the correct file.

`StreamReader` objects are created in much the same way as `StreamWriter` objects. The most common way to create one is to use a previously created `FileStream` object:

```
FileStream aFile = new FileStream("Log.txt", FileMode.Open);
StreamReader sr = new StreamReader(aFile);
```

Like `StreamWriter`, the `StreamReader` class can be created directly from a string containing the path to a particular file:

```
StreamReader sr = new StreamReader("Log.txt");
```

### TRY IT OUT    Reading Data from an Input Stream: StreamRead\Program.cs

1. Create a new console application called `StreamRead` and save it in the directory `C:\BegVCSharp\Chapter18`.

2. Import the `System.IO` and `System.Console` namespaces by placing the following lines of code near the top of `Program.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.IO;
using static System.Console;
```

3. Add the following code to the `Main()` method:

```
static void Main(string[] args)
{
    string line;
    try
    {
        FileStream aFile = new FileStream("Log.txt", FileMode.Open);
        StreamReader sr = new StreamReader(aFile);
        line = sr.ReadLine();
        // Read data in line by line.
        while(line != null)
        {
            WriteLine(line);
            line = sr.ReadLine();
        }
        sr.Close();
    }
    catch(IOException e)
    {
        WriteLine("An IO exception has been thrown!");
        WriteLine(e.ToString());
        return;
    }
    ReadKey();
}
```

4. Copy the `Log.txt` file, created in the previous example, into the `StreamRead\bin\Debug` directory. If you don't have a file named `Log.txt`, the `FileStream` constructor will throw an exception when it doesn't find it.
5. Run the application. You should see the text of the file written to the console, as shown in Figure 18-4.




FIGURE 18-4

## How It Works

This application is very similar to the previous one, with the obvious difference being that it is reading a file, rather than writing one. As before, you must import the `System.IO` namespace to be able to access the necessary classes.

You use the `ReadLine()` method to read text from the file. This method reads text until a new line is found, and returns the resulting text as a string. The method returns a `null` when the end of the file has been reached, which you use to test for the end of the file. Note that you use a `while` loop, which ensures that the line read isn't `null` before any code in the body of the loop is executed — that way, only the genuine contents of the file are displayed:

```
line = sr.ReadLine();
while(line != null)
{
    WriteLine(line);
    line = sr.ReadLine();
}
```

---

## Reading Data

The `ReadLine()` method is not the only way you can access data in a file. The `StreamReader` class has many methods for reading data.

The simplest of the reading methods is `Read()`. It returns the next character from the stream as a positive integer value or a `-1` if it has reached the end. This value can be converted into a character by using the `Convert` utility class. In the preceding example, the main parts of the program could be rewritten as follows:

```
StreamReader sr = new StreamReader(aFile);
int charCode;
charCode = sr.Read();
while(charCode != -1)
{
    Write(Convert.ToChar(charCode));
    charCode = sr.Read();
}
sr.Close();
```

A very convenient method to use with smaller files is the `ReadToEnd()` method. It reads the entire file and returns it as a string. In this case, the earlier application could be simplified to the following:

```
StreamReader sr = new StreamReader(aFile);
line = sr.ReadToEnd();
WriteLine(line);
sr.Close();
```

Although this might seem easy and convenient, be careful. By reading all the data into a string object, you are forcing the data in the file to exist in memory. Depending on the size of the data file, this can be prohibitive. If the data file is extremely large, then it is better to leave the data in the file and access it with the methods of the `StreamReader`.

Another way to deal with large files, which was introduced in .NET 4, is to use the static `File` `.ReadLines()` method. There are, in fact, several static methods of `File` that you can use to simplify reading and writing file data, but this one is particularly interesting in that it returns an `IEnumerable<string>` collection. You can iterate through the strings in this collection to read the file one line at a time. Using this method, you can rewrite the previous example as follows:

```
foreach (string alternativeLine in File.ReadLines("Log.txt"))
    WriteLine(alternativeLine);
```

There are, as you can see, several ways in .NET to achieve the same result — namely, reading data from a file. Choose the technique that suits you best.

## Asynchronous File Access

Sometimes — for example, when you are performing a lot of file access operations in one go or are working with very large files — reading and writing file system data can be slow. If this is the case, you might want to perform other operations while you wait. This is especially important with desktop applications, where you want your application to remain responsive to users while you are doing work in the background.

To facilitate this, .NET 4.5 introduced asynchronous ways to work with streams. This applies to the `FileStream` class, as well as to `StreamReader` and `StreamWriter`. If you have browsed through the definitions of these classes, you might have noticed some methods that end with the suffix `Async` — for example, `StreamReader` has a method called `ReadLineAsync()`, which is an asynchronous version of `ReadLine()`. These methods are designed to be used with the task-based asynchronous programming model.

Asynchronous programming is an advanced technique that isn't covered in detail in this book. However, if asynchronous file system access is something you are interested in doing then this is the place to start. You might also want to read *Professional C# 5.0 and .NET 4.5.1* by Christian Nagel, Jay Glynn, and Morgan Skinner (Wrox, 2014) for more details.

## Reading and Writing Compressed Files

Often when dealing with files, quite a lot of space is used up on the hard disk. This is particularly true for graphics and sound files. You've probably come across utilities that enable you to compress and decompress files, which are handy when you want to move them around or e-mail them. The `System.IO.Compression` namespace contains classes that enable you to compress files from your code, using either the GZIP or Deflate algorithm — both of which are publicly available and free for anyone to use.

There is a little bit more to compressing files than just compressing them, though. You've probably seen how commercial applications enable multiple files to be placed in a single compressed file, often called an *archive*. There are classes in the `System.IO.Compression` namespace that enable similar functionality. However, to keep things simple for this book you'll just look at one scenario: saving text data to a compressed file. You are unlikely to be able to access this file in an external utility, but the file will be much smaller than its uncompressed equivalent!

The two compression stream classes in the `System.IO.Compression` namespace that you'll look at here, `DeflateStream` and `GZipStream`, work very similarly. In both cases, you initialize them with

an existing stream, which, in the case of files, will be a `FileStream` object. After this you can use them with `StreamReader` and `StreamWriter` just like any other stream. All you need to specify in addition to that is whether the stream will be used for compression (saving files) or decompression (loading files) so that the class knows what to do with the data that passes through it. This is best illustrated with the following example.

### TRY IT OUT Reading and Writing Compressed Data: Compressor\Program.cs

1. Create a new console application called Compressor and save it in the directory C:\BegVCSharp\Chapter18.
2. Place the following lines of code near the top of `Program.cs`. You need to import the `System.Console`, `System.IO`, and `System.IO.Compression` namespaces to use the file and compression classes:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.IO;
using System.IO.Compression;
using static System.Console;
```

3. Add the following methods into the body of `Program.cs`, before the `Main()` method:


```
static void SaveCompressedFile(string filename, string data)
{
    FileStream fileStream =
        new FileStream(filename, FileMode.Create, FileAccess.Write);
    GZipStream compressionStream =
        new GZipStream(fileStream, CompressionMode.Compress);
    StreamWriter writer = new StreamWriter(compressionStream);
    writer.Write(data);
    writer.Close();
}
static string LoadCompressedFile(string filename)
{
    FileStream fileStream =
        new FileStream(filename, FileMode.Open, FileAccess.Read);
    GZipStream compressionStream =
        new GZipStream(fileStream, CompressionMode.Decompress);
    StreamReader reader = new StreamReader(compressionStream);
    string data = reader.ReadToEnd();
    reader.Close();
    return data;
}
```

4. Add the following code to the `Main()` method:

```
static void Main(string[] args)
{
    try
    {
        string filename = "compressedFile.txt";
```


```
WriteLine(
    "Enter a string to compress (will be repeated 100 times):");
string sourceString = ReadLine();
StringBuilder sourceStringMultiplier =
    new StringBuilder(sourceString.Length * 100);
for (int i = 0; i < 100; i++)
{
    sourceStringMultiplier.Append(sourceString);
}
sourceString = sourceStringMultiplier.ToString();
WriteLine($"Source data is {sourceString.Length} bytes long.");
SaveCompressedFile(filename, sourceString);
WriteLine($"\\nData saved to {filename}.");
FileInfo compressedFileData = new FileInfo(filename);
Write($"Compressed file is {compressedFileData.Length}");
WriteLine(" bytes long.");
string recoveredString = LoadCompressedFile(filename);
recoveredString = recoveredString.Substring(
    0, recoveredString.Length / 100);
WriteLine($"\\nRecovered data: {recoveredString},");
ReadKey();
}
catch (IOException ex)
{
    WriteLine("An IO exception has been thrown!");
    WriteLine(ex.ToString());
    ReadKey();
}
```

5. Run the application and enter a suitably long string. An example result is shown in Figure 18-5.



**FIGURE 18-5**

6. Open compressedFile.txt in Notepad. The text is shown in Figure 18-6.



**FIGURE 18-6**

## How It Works

In this example, you define two methods for saving and loading a compressed text file. The first of these, `SaveCompressedFile()`, is as follows:

```
static void SaveCompressedFile(string filename, string data)
{
    FileStream fileStream =
        new FileStream(filename, FileMode.Create, FileAccess.Write);
    GZipStream compressionStream =
        new GZipStream(fileStream, CompressionMode.Compress);
    StreamWriter writer = new StreamWriter(compressionStream);
    writer.Write(data);
    writer.Close();
}
```

The code starts by creating a `FileStream` object, and then uses it to create a `GZipStream` object. Note that you could replace all occurrences of `GZipStream` in this code with `DeflateStream` — the classes work in the same way. You use the `CompressionMode.Compress` enumeration value to specify that data is to be compressed, and then use a `StreamWriter` to write data to the file.

`LoadCompressedFile()` mirrors the `SaveCompressedFile()` method. Instead of saving to a filename, it loads a compressed file into a string:

```
static string LoadCompressedFile(string filename)
{
    FileStream fileStream =
        new FileStream(filename, FileMode.Open, FileAccess.Read);
    GZipStream compressionStream =
        new GZipStream(fileStream, CompressionMode.Decompress);
    StreamReader reader = new StreamReader(compressionStream);
    string data = reader.ReadToEnd();
    reader.Close();
    return data;
}
```

The differences are as you would expect — different `FileMode`, `FileAccess`, and `CompressionMode` enumeration values to load and uncompress data, and the use of a `StreamReader` to get the uncompressed text out of the file.

The code in `Main()` is a simple test of these methods. It simply asks for a string, duplicates the string 100 times to make things interesting, compresses it to a file, and then retrieves it. In the example, the first sentence of book XI of *The Iliad* repeated 100 times is 19,400 characters long, but when compressed, it takes up only 225 bytes — that's a compression ratio of more than 80:1. Admittedly, this is a bit of a cheat — the GZIP algorithm works particularly well with repetitive data, but it does illustrate compression in action.

You also looked at the text stored in the compressed file. Obviously, it isn't easily readable, which has implications should you want to share data between applications, for example. However, because the file was compressed with a known algorithm, at least you know that it is possible for applications to uncompress it.

## MONITORING THE FILE SYSTEM

Sometimes an application must do more than just read and write files to the file system. For example, it might be important to know when files or directories are being modified. The .NET Framework has made it easy to create custom applications that do just that.

The class that helps you to do this is the `FileSystemWatcher` class. It exposes several events that your application can catch. This enables your application to respond to file system events.

The basic procedure for using the `FileSystemWatcher` is simple. First, you must set a handful of properties, which specify where to monitor, what to monitor, and when it should raise the event that your application will handle. Then you give it the addresses of your custom event handlers, so that it can call these when significant events occur. Finally, you turn it on and wait for the events.

The properties that must be set before a `FileSystemWatcher` object is enabled are shown in Table 18-10.

TABLE 18-10: `FileSystemWatcher` Properties

PROPERTY	DESCRIPTION
<code>Path</code>	Must be set to the file location or directory to monitor.
<code>NotifyFilter</code>	A combination of <code>NotifyFilters</code> enumeration values that specify what to watch for within the monitored files. These represent properties of the file or folders being monitored. If any of the specified properties change, then an event is raised. The possible enumeration values are <code>Attributes</code> , <code>CreationTime</code> , <code>DirectoryName</code> , <code>FileName</code> , <code>LastAccess</code> , <code>LastWrite</code> , <code>Security</code> , and <code>Size</code> . Note that these can be combined using the binary OR operator.
<code>Filter</code>	A filter specifying which files to monitor — for example, <code>*.txt</code> .

Once these are set, you must write event handlers for four events: `Changed`, `Created`, `Deleted`, and `Renamed`. As shown in Chapter 13, this is simply a matter of creating your own method and assigning it to the object's event. By assigning your own event handler to these methods, your method will be called when the event is fired. Each event will fire when a file or directory matching the `Path`, `NotifyFilter`, and `Filter` property is modified.

Once you have set the properties and the events, set the `EnableRaisingEvents` property to `true` to begin the monitoring. In the following Try It Out, you use `FileSystemWatcher` in a simple client application to keep tabs on a directory of your choice.

### TRY IT OUT Monitoring the File System: FileWatch

Here's a more sophisticated example using much of what you have learned in this chapter.

1. Create a new WPF application called `FileWatch` and save it in the directory `C:\BegVCSharp\Chapter18`.

2. Modify MainWindow.xaml as follows (the resulting window is shown in Figure 18-7):

```
<Window x:Class="FileWatch.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    Title="File Monitor" Height="160" Width="300">
    <Grid>
        <Grid.RowDefinitions>
            <RowDefinition Height="Auto" />
            <RowDefinition Height="Auto" />
            <RowDefinition />
        </Grid.RowDefinitions>
        <Grid Margin="4">
            <Grid.ColumnDefinitions>
                <ColumnDefinition />
                <ColumnDefinition Width="Auto" />
            </Grid.ColumnDefinitions>
            <TextBox Name="LocationBox"TextChanged="LocationBox_TextChanged" />
            <Button Name="BrowseButton" Grid.Column="1" Margin="4,0,0,0"
                Content="Browse..." Click="BrowseButton_Click" />
        </Grid>
        <Button Name="WatchButton" Content="Watch!" Margin="4" Grid.Row="1"
            Click="WatchButton_Click" IsEnabled="False" />
        <ListBox Name="WatchOutput" Margin="4" Grid.Row="2" />
    </Grid>
</Window>
```

3. Add the following using directives to MainWindow.xaml.cs:

```
using System.IO;
using Microsoft.Win32;
```



FIGURE 18-7

4. Add a field of type FileSystemWatcher class to the MainWindow class:

```
namespace FileWatch
{
    /// <summary>
    /// Interaction logic for MainWindow.xaml
    /// </summary>
    public partial class MainWindow : Window
    {
        // File System Watcher object.
        private FileSystemWatcher watcher;
```

5. Add the following utility method to the class to allow messages to be added to the output from a background thread:

```
private void AddMessage(string message
{
    Dispatcher.BeginInvoke(new Action(
        () => WatchOutput.Items.Insert(
            0, message)));
}
```

6. Just after the `InitializeComponent()` method call in the window constructor, add the following code. This code is needed to initialize the `FileSystemWatcher` object and associate the events to calls to `AddMessage()`:

```
public MainWindow()
{
    InitializeComponent();
    watcher = new FileSystemWatcher();
    watcher.Deleted += (s, e) =>
        AddMessage($"File: {e.FullPath} Deleted");
    watcher.Renamed += (s, e) =>
        AddMessage($"File renamed from {e.OldName} to {e.FullPath}");
    watcher.Changed += (s, e) =>
        AddMessage($"File: {e.FullPath} {e.ChangeType.ToString()}");
    watcher.Created += (s, e) =>
        AddMessage($"File: {e.FullPath} Created");
}
```

7. Add the `Click` event handler for the `Browse` button. The code in this event handler opens the Open File dialog box, enabling the user to select a file to monitor:

```
private void BrowseButton_Click(object sender, RoutedEventArgs e)
{
    OpenFileDialog dialog = new OpenFileDialog();
    if (dialog.ShowDialog(this) == true)
    {
        LocationBox.Text = dialog.FileName;
    }
}
```

The `ShowDialog()` method returns a `bool?` value reflecting how the user exited the File Open dialog box (the user could have clicked OK or pressed the Cancel button). You need to confirm that the user did not click the Cancel button, so you compare the result from the method call to `true` before saving the user's file selection to the `TextBox`.

8. Add the `TextChanged` event handler for the `TextBox` to ensure the `Watch!` button is enabled when the `TextBox` contains text:

```
private void LocationBox_TextChanged(object sender, TextChangedEventArgs e)
{
    WatchButton.IsEnabled = !string.IsNullOrEmpty(LocationBox.Text);
}
```

9. Add the following code to the Click event handler for the Watch! button, which starts the `FileSystemWatcher`:

```
private void WatchButton_Click(object sender, RoutedEventArgs e)
{
    watcher.Path = System.IO.Path.GetDirectoryName(LocationBox.Text);
    watcher.Filter = System.IO.Path.GetFileName(LocationBox.Text);
    watcher.NotifyFilter = NotifyFilters.LastWrite |
        NotifyFilters.FileName | NotifyFilters.Size;
    AddMessage("Watching " + LocationBox.Text);
    // Begin watching.
    watcher.EnableRaisingEvents = true;
}
```

10. Create a directory called `C:\TempWatch` and a file in this directory called `temp.txt`.
11. Run the application. If everything builds successfully, click the Browse button and select `C:\TempWatch\temp.txt`.
12. Click the Watch! button to begin monitoring the file. The only change you will see in your application is a message confirming that the file is being watched.
13. Using Windows Explorer, navigate to `C:\TempWatch`. Open `temp.txt` in Notepad, add some text to the file, and save it.
14. Rename the file.
15. You should see a description of the changes to the file you selected to watch, as shown in Figure 18-8.




FIGURE 18-8

## How It Works

This application is fairly simple, but it demonstrates how the `FileSystemWatcher` works. Try playing with the string you put into the monitor text box. If you specify `*.*` in a directory, it will monitor all changes in the directory.

Most of the code in the application is related to setting up the `FileSystemWatcher` object to watch the correct location:

```
watcher.Path = System.IO.Path.GetDirectoryName(LocationBox.Text);
watcher.Filter = System.IO.Path.GetFileName(LocationBox.Text);
```

```
watcher.NotifyFilter = NotifyFilters.LastWrite |
    NotifyFilters.FileName | NotifyFilters.Size;
AddMessage("Watching " + LocationBox.Text);
// Begin watching.
watcher.EnableRaisingEvents = true;
```

The code first sets the path to the directory to monitor. It uses a new object you have not looked at yet: `System.IO.Path`. This is a static class, much like the static `File` object. It exposes many static methods to manipulate and extract information out of file location strings. You first use it to extract the directory name the user typed in the text box, using the `GetDirectoryName()` method.

The next line sets the filter for the object. This can be an actual file, in which case it would only monitor the file, or it could be something like `*.txt`, in which case it would monitor all the `.txt` files in the directory specified. Again, you use the `Path` static object to extract the information from the supplied file location.

The `NotifyFilter` is a combination of `NotifyFilters` enumeration values that specify what constitutes a change. In this example, you have indicated that if the last write time stamp, the filename, or the size of the file changes, your application should be notified of the change. After updating the UI, you set the `EnableRaisingEvents` property to `true` to begin monitoring.

Before that, however, you have to create the object and set the event handlers:

```
watcher = new FileSystemWatcher();
watcher.Deleted += (s, e) =>
    AddMessage($"File: {e.FullPath} Deleted");
watcher.Renamed += (s, e) =>
    AddMessage($"File renamed from {e.OldName} to {e.FullPath}");
watcher.Changed += (s, e) =>
    AddMessage($"File: {e.FullPath} {e.ChangeType.ToString()}");
watcher.Created += (s, e) =>
    AddMessage($"File: {e.FullPath} Created");
```

This code uses lambda expressions to create anonymous event handler methods for the events raised by the `watcher` object when a file is deleted, renamed, changed, or created. These event handlers simply call the `AddMessage()` method with an informative message. Obviously, you could implement a more sophisticated response, depending on your application. When a file is added to a directory, you could move it somewhere else or read the contents and fire off a new process using the information. The possibilities are endless!

## EXERCISES

- 18.1 Which namespace enables an application to work with files?
- 18.2 When would you use a `FileStream` object to write to a file instead of using a `StreamWriter` object?
- 18.3 Which methods of the `StreamReader` class enable you to read data from files and what does each one do?
- 18.4 Which class would you use to compress a stream by using the Deflate algorithm?

- 18.5** Which events does the `FileSystemWatcher` class expose and what are they for?
- 
- 18.6** Modify the FileWatch application you built in this chapter by adding the capability to turn the file system monitoring on and off without exiting the application.
- 

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>Streams</b>	A stream is an abstract representation of a serial device that you can read from or write to a byte at a time. Files are an example of such a device. There are two types of streams — input and output — for reading from and writing to devices, respectively.
<b>File classes</b>	There are numerous classes in the .NET Framework that abstract file system access, including <code>File</code> and <code>Directory</code> for dealing with files and directories through static methods, and <code>FileInfo</code> and <code> DirectoryInfo</code> , which can be instantiated to represent specific files and directories. The latter pair of classes is useful when you perform multiple operations on files and directories, as those classes don't require a path for every method call. Typical operations that you can perform on files and directories include interrogating and changing properties, creating, deleting, and copying.
<b>File paths</b>	File and directory paths can be absolute or relative. An absolute path gives a complete description of a location starting from the root of the drive that contains it; all parent directories are separated from child directories with backslashes. Relative directories are similar, but start from a defined point in the file system, such as the directory where an application is executing (the working directory). To navigate the file system, you often use the <code>..</code> parent directory alias.
<b>The <code>FileStream</code> object</b>	The <code>FileStream</code> object provides access to the contents of a file, for reading and writing purposes. It accesses file data at the byte level, and so is not always the best choice for accessing file data. A <code>FileStream</code> instance maintains a position byte index within a file so that you can navigate through the contents of a file. Accessing a file at any point in this way is known as <i>random access</i> .
<b>Reading and writing to streams</b>	An easier way to read and write file data is to use the <code>StreamReader</code> and <code>StreamWriter</code> classes in combination with a <code>FileStream</code> . These enable you to read and write character and string data rather than working with bytes. These types expose familiar methods for working with strings, including <code>ReadLine()</code> and <code>WriteLine()</code> . Because they work with string data, these classes make it easy to work with comma-delimited files, which are a common way to represent structured data.

---

<b>Compressed files</b>	You can use the <code>DeflateStream</code> and <code>GZipStream</code> compressed stream classes to read and write compressed data from and to files. These classes work with byte data much like <code>FileStream</code> , but as with <code>FileStream</code> you can access data through <code>StreamReader</code> and <code>StreamWriter</code> classes to simplify your code.
<b>Monitoring the file system</b>	You can use the <code>FileSystemWatcher</code> class to monitor changes to file system data. You can monitor both files and directories, and provide a filter, if required, to modify only those files that have a specific file extension. <code>FileSystemWatcher</code> instances notify you of changes by raising events that you can handle in your code.

---



# 19

## XML and JSON

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- XML basics
- JSON basics
- XML schemas
- XML Document Object Model
- Converting XML to JSON
- Searching XML documents using XPath

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 19 download and individually named according to the names throughout the chapter.

Just as programming languages like C# describe computer logic in a format that is readable by both machines and humans, XML and JSON are both *data languages*, which are used storing data in a simple text format that can be read by both humans and nearly any computer.

Most C# .NET applications use XML in some form for storing data, such as .config files for storing configuration details and XAML files used in WPF and Windows Store applications. Because of this important fact, we'll spend the most time in this chapter on XML, with just a short look at JSON on the side.

During this chapter you will learn the basics of XML and JSON and then learn how to create XML documents and schemas. You will learn the basics of the XmlDocument class, how to read and write XML, how to insert and delete nodes, how to convert XML to JSON format, and finally how to search for data in XML documents using XPath.

## XML BASICS

Extensible Markup Language (XML) is a *data language*, which is a way of storing data in a simple text format that can be read by both humans and nearly any computer. It is a W3C standard format like HTML ([www.w3.org/XML](http://www.w3.org/XML)). It has been fully adopted by Microsoft in the .NET Framework and other Microsoft products. Even the document formats introduced with the newer versions of Microsoft Office are based on XML, although the Office applications themselves are not .NET applications.

The ins and outs of XML can be very complicated, so you won't look at every single detail here. Luckily, most tasks don't require a detailed knowledge of XML because Visual Studio typically takes care of most of the work — you will rarely have to write an XML document by hand. If you want to learn about XML in more depth, read a book such as *Beginning XML* by Joe Fawcett, Danny Ayers, and Liam Quin (Wrox, 2012) or one of the many online tutorials such as [www.xmlnews.org/docs/xml-basics.html](http://www.xmlnews.org/docs/xml-basics.html) or <http://www.w3schools.com/xml/>.

The basic format is very simple, as you can see in the following example that shows an XML format for sharing data about books.

```
<book>
  <title>Beginning Visual C# 2015</title>
  <author>Benjamin Perkins et al</author>
  <code>096689</code>
</book>
```

In this example each book has a title, an author, and a unique code identifying the book. Each book's data is contained in a book *element* beginning with a `<book>` tag and ending with the `</book>` end tag. The title, author, and code values are stored in nested elements inside the book element.

Optionally, an element may also have *attributes* inside the tag itself. If the book code were an attribute of the book element instead of its own element, you'd see the book element beginning with something like this: `<book code=096689>`. To keep it simple we'll stick with elements in this chapter's examples. Generically both attributes and elements are called *nodes*, like the nodes of a graph.

## JSON BASICS

Another data language you may encounter when developing C# applications is JSON. JSON stands for JavaScript Object Notation. Like XML, it is also a standard ([www.json.org](http://www.json.org)), though as you can tell from the name it is derived from the JavaScript language rather than C#. While not used throughout .NET like XML, it is a common format for transferring data from web services and web browsers.

JSON also has a very simple format. The same book data we showed previously in XML is presented here in JSON:

```
{"book": [{"title": "Beginning Visual C# 2015",
  "author": "Benamin Perkins et al",
  "code": "096689"}]}
```

As with the previous XML example, we see the same book with title, author, and a unique code. JSON uses curly braces ({} ) to delimit blocks of data and square brackets ([]) to delimit arrays similar to the way C#, JavaScript, and other C-like languages use curly braces for blocks of code and square brackets for arrays.

JSON is a more compact format than XML, but it is much harder for humans to read, especially as the curly braces and brackets become deeply nested in complex data.

## XML SCHEMAS

An XML document may be described by a *schema*, which is another XML file describing what elements and attributes are allowed in a particular document. You can *validate* an XML document against a schema, ensuring that your program doesn't encounter data it isn't prepared to handle. The standard schema XML format used with C# is XSD (for XML Schema Definition).

Figure 19-1 includes a long list of schemas recognized by Visual Studio, but it will not automatically remember schemas you've used. If you are using a schema repeatedly and don't want to browse for it every time you need it, you can copy it to the following location: C:\Program Files\Microsoft visual Studio 14.0\Xml\Schemas. Any schema copied to that location will show up on the XML Schemas dialog box.

### TRY IT OUT

### Creating an XML Document in Visual Studio: Chapter19\XML and Schema\GhostStories.xml

Follow these steps to create an XML document:

1. Open Visual Studio and select File  $\Rightarrow$  New  $\Rightarrow$  File from the menu. If you don't see this option, create a new project, right-click the project in the Solution Explorer, and choose to add a new item. Then select XML File from the dialog box.
2. In the New File dialog box, select XML File and click Open. Visual Studio creates a new XML document for you. As Figure 19-2 shows, Visual Studio adds an XML declaration, complete with an encoding attribute. (It also colors the attributes and elements.)
3. Save the file by pressing Ctrl+S or by selecting File  $\Rightarrow$  Save XMLfile1.xml from the menu. Visual Studio asks you where to save the file and what to call the file; save it in the BegVCSharp\Chapter19\XML and Schemas folder as GhostStories.xml.
4. Move the cursor to the line underneath the XML declaration, and type the text `<stories>`. Notice how Visual Studio automatically puts the end tag in as soon as you type the greater than sign to close the opening tag.
5. Type this XML file and then click Save:

```
<stories>
  <story>
    <title>A House in Aungier Street</title>
```

```
<author>
    <name>Sheridan Le Fanu</name>
    <nationality>Irish</nationality>
</author>
<rating>eerie</rating>
</story>
<story>
    <title>The Signalman</title>
    <author>
        <name>Charles Dickens</name>
        <nationality>English</nationality>
    </author>
    <rating>atmospheric</rating>
</story>
<story>
    <title>The Turn of the Screw</title>
    <author>
        <name>Henry James</name>
        <nationality>American</nationality>
    </author>
    <rating>a bit dull</rating>
</story>
</stories>
```

6. It is now possible to let Visual Studio create a schema that fits the XML you have written. Do this by selecting the Create Schema menu option from the XML menu. Save the resulting XSD file by clicking Save as GhostStories.xsd.
7. Return to the XML file and type the following XML before the ending `</stories>` tag:

```
<story>
    <title>Number 13</title>
    <author>
        <name>M.R. James</name>
        <nationality>English</nationality>
    </author>
    <rating>mysterious</rating>
</story>
```

You are now getting IntelliSense hints when you begin typing the starting tags. That's because Visual Studio knows to connect the newly created XSD schema to the XML file you are typing.

8. It is possible to create this link between XML and one or more schemas in Visual Studio. Select `XML ⇔ Schemas`. That brings up the XML Schemas dialog box shown in Figure 19-1. At the top of the long list of schemas that Visual Studio recognizes, you will see `GhostStories.xsd`. To the left of it is a checkmark, which indicates that this schema is being used on the current XML document.




FIGURE 19-1




FIGURE 19-2

## XML DOCUMENT OBJECT MODEL

The XML Document Object Model (XML DOM) is a set of classes used to access and manipulate XML in a very intuitive way. The DOM is perhaps not the quickest way to read XML data, but as soon as you understand the relationship between the classes and the elements of an XML document, you will find it very easy to use.

The classes that make up the DOM can be found in the namespace `System.Xml`. There are several classes and namespaces in this namespace, but this chapter focuses on only a few of the classes that enable you to easily manipulate XML. These classes are described in Table 19-1.

TABLE 19-1: Common DOM Classes

CLASS	DESCRIPTION
<code>XmlNode</code>	Represents a single node in a document tree. It is the base of many of the classes shown in this chapter. If this node represents the root of an XML document, you can navigate to any position in the document from it.
<code> XmlDocument</code>	Extends the <code>XmlNode</code> class, but is often the first object you use when using XML. That's because this class is used to load and save data from disk or elsewhere.

*continues*

**TABLE 19-1** (continued)

CLASS	DESCRIPTION
XmlElement	Represents a single element in the XML document. <code>XmlElement</code> is derived from <code>XmlNode</code> , which in turn is derived from <code>XmlNode</code> .
XmlAttribute	Represents a single attribute. Like the <code> XmlDocument</code> class, it is derived from the <code> XmlNode</code> class.
XmlText	Represents the text between a starting tag and a closing tag.
XmlComment	Represents a special kind of node that is not regarded as part of the document other than to provide information to the reader about parts of the document.
XmlNodeList	Represents a collection of nodes.

## The `XmlDocument` Class

Usually, the first thing your application will want to do with XML is read it from disk. As described in Table 19-1, this is the domain of the  `XmlDocument` class. You can think of the  `XmlDocument` as an in-memory representation of the file on disk. Once you have used the  `XmlDocument` class to load a file into memory, you can obtain the root node of the document from it and start reading and manipulating the XML:

```
using System.Xml;
.
.
.
 XmlDocument document = new XmlDocument();
document.Load(@"C:\BegVCSharp\Chapter19\XML and Schema\books.xml");
```

The two lines of code create a new instance of the  `XmlDocument` class and load the file `books.xml` into it.

**NOTE** Note the folder name is an absolute path; your folder structure may differ and if so you should adjust the path following `document.Load` to reflect the actual folder path on your computer.

Remember that the  `XmlDocument` class is located in the `System.Xml` namespace, and you should insert a `using System.Xml;` in the `using` section at the beginning of the code.

In addition to loading and saving the XML, the  `XmlDocument` class is also responsible for maintaining the XML structure itself. Therefore, you will find numerous methods on this class that are used to create, alter, and delete nodes in the tree. You will look at some of those methods shortly, but to present the methods properly, you need to know a bit more about another class: `XmlElement`.

## The `XmlElement` Class

Now that the document has been loaded into memory, you want to do something with it. The  `DocumentElement` property of the  `XmlDocument` instance you created in the preceding code returns

an instance of an `XmlElement` that represents the root element of the  `XmlDocument`. This element is important because it gives you access to every bit of information in the document:

```
 XmlDocument document = new XmlDocument();
document.Load(@"C:\BegVCSharp\Chapter19\
XML and Schema\books.xml");
XmlElement element = document.DocumentElement;
```

After you have the root element of the document, you are ready to use the information. The `XmlElement` class contains methods and properties for manipulating the nodes and attributes of the tree. Let's examine the properties for navigating the XML elements first, shown in Table 19-2.

**TABLE 19-2: XmlElement Properties**

PROPERTY	DESCRIPTION
FirstChild	<p>Returns the first child element after this one. If you recall the <code>books.xml</code> file from earlier in the chapter, the root node of the document was called "books" and the next node after that was "book." In that document, then, the first child of the root node "books" is "book."</p> <pre>&lt;books&gt;      Root node &lt;book&gt;      FirstChild</pre> <p><code>FirstChild</code> returns an <code>XmlNode</code> object, and you should test for the type of the returned node because it is unlikely to always be an <code>XmlElement</code> instance. In the books example, the child of the <code>Title</code> element is, in fact, an <code>XmlText</code> node that represents the text Beginning Visual C#.</p>
LastChild	<p>Operates exactly like the <code>FirstChild</code> property except that it returns the last child of the current node. In the case of the books example, the last child of the "books" node will still be a "book" node, but it will be the node representing the "Beginning XML" book.</p> <pre>&lt;books&gt; Root node &lt;book&gt; FirstChild &lt;title&gt;Beginning Visual C# 2015&lt;/title&gt; &lt;author&gt;Benjamin Perkins et al&lt;/author&gt; &lt;code&gt;096689&lt;/code&gt; &lt;/book&gt; &lt;book&gt; LastChild &lt;title&gt;Beginning XML&lt;/title&gt; &lt;author&gt;Joe Fawcett et al&lt;/author&gt; &lt;code&gt;162132&lt;/code&gt; &lt;/book&gt; &lt;/books&gt;</pre>

*continues*

**TABLE 19-2** (continued)

PROPERTY	DESCRIPTION
ParentNode	Returns the parent of the current node. In the books example, the “books” node is the parent of both of the “book” nodes.
NextSibling	Where <code>FirstChild</code> and <code>LastChild</code> properties return the leaf node of the current node, the <code>NextSibling</code> node returns the next node that has the same parent node. In the case of the books example, that means getting the <code>NextSibling</code> of the <code>title</code> element will return the <code>author</code> element, and calling <code>NextSibling</code> on that will return the <code>code</code> element.
HasChildNodes	Enables you to check whether the current element has child elements without actually getting the value from <code>FirstChild</code> and examining that against <code>null</code> .

Using the five properties from Table 19-2, it is possible to run through an entire  `XmlDocument`, as shown in the following Try It Out.

### TRY IT OUT Looping through All Nodes in an XML Document: Chapter19\LoopThroughXml Document\MainWindows.xaml.cs

In this example, you are going to create a small WPF application that loops through all the nodes of an XML document and prints out the name of the element or the text contained in the element in the case of an  `XmlText` element. This code uses `Books.xml`, which you saw in the “Schemas” section earlier; if you didn’t create that file as you worked through that section, you can find it in `Chapter19\XML and Schemas\` in this chapter’s downloadable code.

1. Begin by creating a new WPF project by selecting `File`  $\Rightarrow$  `New`  $\Rightarrow$  `Project`. In the dialog box that appears, select `Windows`  $\Rightarrow$  `WPF Application`. Name the project `LoopThroughXmlDocument` and press Enter.
2. Design the form as shown in Figure 19-3 by dragging a `TextBlock` control and a `Button` control onto the form.
3. Name the `TextBlock` control `textBlockResults` and name the button `buttonLoop`. Allow the `TextBlock` to fill all the space not used by the button.
4. Add the event handler for the `Click` event for the button and enter the code that follows. Don’t forget to add `using System.Xml;` to the `using` section at the top of the file:

```
private void buttonLoop_Click(object sender, RoutedEventArgs e)
{
    XmlDocument document = new XmlDocument();
    document.Load(booksFile);
    textBlockResults.Text =
        FormatText(document.DocumentElement as XmlNode, "", "");
}
private string FormatText(XmlNode node, string text, string indent)
{
```

```

        if (node is XmlText)
        {
            text += node.Value;
            return text;
        }
        if (string.IsNullOrEmpty(indent))
            indent = "";
        else
        {
            text += "\r\n" + indent;
        }
        if (node is XmlComment)
        {
            text += node.OuterXml;
            return text;
        }
        text += "<" + node.Name;
        if (node.Attributes.Count > 0)
        {
            AddAttributes(node, ref text);
        }
        if (node.HasChildNodes)
        {
            text += ">";
            foreach (XmlNode child in node.ChildNodes)
            {
                text = FormatText(child, text, indent + "  ");
            }
            if (node.ChildNodes.Count == 1 &&
                (node.FirstChild is XmlText || node.FirstChild is XmlComment))
                text += "</" + node.Name + ">";
            else
                text += "\r\n" + indent + "</" + node.Name + ">";
        }
        else
            text += " />";
        return text;
    }
    private void AddAttributes(XmlNode node, ref string text)
    {
        foreach (XmlAttribute xa in node.Attributes)
        {
            text += " " + xa.Name + "=" + xa.Value + "'";
        }
    }
}

```

5. Add the private const that holds the location of the file that is loaded. You can change the location to reflect the location you put the file on your local system:

```
private const string booksFile =
@"C:\BegVCSharp\Chapter19\XML and Schema\Books.xml";
```

6. Run the application and click Loop. You should get a result like the one shown in Figure 19-4.




FIGURE 19-3




FIGURE 19-4

## How It Works

When you click the button, the  `XmlDocument` method `Load` is called. This method loads the XML from a file into the  `XmlDocument` instance, which can then be used to access the elements of the XML. Then you call a method that enables you to loop through the XML recursively, passing the root node of the XML document to the method. The root element is obtained with the property  `DocumentElement` of the  `XmlDocument` class. Aside from the check for `null` on the `root` parameter that is passed into the `FormatText` method, the first line to note is the `if` sentence:

```
if (node is XmlText)
{
  ...
}
```

Recall that the `is` operator enables you to examine the type of an object, and it returns `true` if the instance is of the specified type. Even though the root node is declared as an `XmlNode`, that is merely the base type of the objects you are going to work with. By using the `is` operator to test the type of the objects, you are able to determine the type of the object at runtime and select the action to perform based on that.

Inside the `FormatText` method you generate the text for the textbox. You have to know the type of the current instance of `root` because the information you want to display is obtained differently for different elements: You want to display the name of `XmlElement`s and the value of `XmlText` elements.

## Changing the Values of Nodes

Before you examine how to change the value of a node, it is important to realize that very rarely is the value of a node a simple thing. In fact, you will find that although all of the classes that derive from `XmlNode` include a property called `value`, it very rarely returns anything useful to you. Although this can feel like a bit of a letdown at first, you'll find it is actually quite logical. Examine the books example from earlier:

```
<books>
  <book>
    <title>Beginning Visual C# 2015</title>
    <author>Benjamin Perkins et al</author>
    <code>096689</code>
  </book>
  <book>
</books>
```

Every single tag pair in the document resolves into a node in the DOM. Remember that when you looped through all the nodes in the document, you encountered a number of `XmlElement` nodes and three `XmlText` nodes. The `XmlElement` nodes in this XML are `<books>`, `<book>`, `<title>`, `<author>`, and `<code>`. The `XmlText` nodes are the text between the starting and closing tags of title, author, and code. Although it could be argued that the value of title, author, and code is the text between the tags, that text is itself a node; and it is that node that actually holds the value. The other tags clearly have no value associated with them other than other nodes.

The following line is in the `if` block near the top of the code in the earlier `FormatText` method. It executes when the current node is an `XmlText` node.

```
text += node.Value;
```

You can see that the `Value` property of the `XmlText` node instance is used to get the value of the node.

Nodes of the type `XmlElement` return `null` if you use their `Value` property, but it is possible to get the information between the starting and closing tags of an `XmlElement` if you use one of two other methods: `InnerText` and `InnerXml`. That means you are able to manipulate the value of nodes using two methods and a property, as described in Table 19-3.

TABLE 19-3: Three Ways to Get the Value of a Node

PROPERTY	DESCRIPTION
InnerText	Gets the text of all the child nodes of the current node and returns it as a single concatenated string. This means if you get the value of <code>InnerText</code> from the <code>book</code> node in the preceding XML, the string <code>Beginning Visual C# 2015#Benjamin Perkins et al096689</code> is returned. If you get the <code>InnerText</code> of the <code>title</code> node, only <code>"Beginning Visual C# 2015"</code> is returned. You can set the text using this method, but be careful if you do so because if you set the text of a wrong node you may overwrite information you did not want to change.
InnerXml	Returns the text like <code>InnerText</code> , but it also returns all of the tags. Therefore, if you get the value of <code>InnerXml</code> on the <code>book</code> node, the result is the following string: <code>&lt;title&gt;Beginning Visual C# 2015&lt;/title&gt;&lt;author&gt;Benjamin Perkins et al&lt;/author&gt;&lt;code&gt;096689&lt;/code&gt;</code> As you can see, this can be quite useful if you have a string containing XML that you want to inject directly into your XML document. However, you are entirely responsible for the string yourself, and if you insert badly formed XML, the application will generate an exception.
Value	The “cleanest” way to manipulate information in the document, but as mentioned earlier, only a few of the classes actually return anything useful when you get the value. The classes that will return the desired text are as follows: <code>XmlText</code> <code>XmlComment</code> <code>XmlAttribute</code>

## Inserting New Nodes

Now that you’ve seen that you can move around in the XML document and even get the values of the elements, let’s examine how to change the structure of the document by adding nodes to the books document you’ve been using.

To insert new elements in the list, you need to examine the new methods that are placed on the  `XmlDocument` and  `XmlNode` classes, shown in Table 19-4. The  `XmlDocument` class has methods that enable you to create new  `XmlNode` and  `XmlElement` instances, which is nice because both of these classes have only a protected constructor, which means you cannot create an instance of either directly with `new`.

**TABLE 19-4:** Methods for Creating Nodes

METHOD	DESCRIPTION
CreateNode	Creates any kind of node. There are three overloads of the method, two of which enable you to create nodes of the type found in the <code>XmlNodeType</code> enumeration and one that enables you to specify the type of node to use as a string. Unless you are quite sure about specifying a node type other than those in the enumeration, use the two overloads that use the enumeration. The method returns an instance of <code>XmlNode</code> that can then be cast to the appropriate type explicitly.
CreateElement	A version of <code>CreateNode</code> that creates only nodes of the <code>XmlElement</code> variety.
CreateAttribute	A version of <code>CreateNode</code> that creates only nodes of the <code>XmlAttribute</code> variety.
CreateTextNode	Creates — yes, you guessed it — nodes of the type <code>XmlText</code> .
CreateComment	This method is included here to highlight the diversity of node types that can be created. This method doesn't create a node that is actually part of the data represented by the XML document, but rather is a comment meant for any human eyes that might have to read the data. You can pick up comments when reading the document in your applications as well.

The methods in Table 19-4 are all used to create the nodes themselves, but after calling any of them you have to do something with them before they become interesting. Immediately after creation, the nodes contain no additional information, and they are not yet inserted into the document. To do either, you should use methods that are found on any class derived from `XmlNode` (including  `XmlDocument` and `XmlElement`), described in Table 19-5.

**TABLE 19-5:** Methods for Inserting Nodes

METHOD	DESCRIPTION
AppendChild	Appends a child node to a node of type <code>XmlNode</code> or a derived type. Remember that the node you append appears at the bottom of the list of children of the node on which the method is called. If you don't care about the order of the children, there's no problem; if you do care, remember to append the nodes in the correct sequence.
InsertAfter	Controls exactly where you want to insert the new node. The method takes two parameters — the first is the new node and the second is the node after which the new node should be inserted.
InsertBefore	Works exactly like <code>InsertAfter</code> , except that the new node is inserted before the node you supply as a reference.

In the following Try It Out, you build on the previous example and insert a book node in the `books.xml` document. There is no code in the example to clean up the document (yet), so if you run it several times you will probably end up with a lot of identical nodes.


**TRY IT OUT****Creating Nodes: Chapter19\LoopThrough XmlDocument\MainWindow.xaml.cs**

This example builds on the `LoopThrough XmlDocument` project you created earlier. Follow these steps to add a node to the `books.xml` document:

1. Wrap the `TextBlock` in a `ScrollViewer` and set its `VerticalScrollBarVisibility` property to `Auto`.
2. Add a button beneath the existing button on the form and name it `buttonCreateNode`. Change its `Content` property to `Create`.
3. Add the Click event handler to the new button and enter the following code:

```
private void buttonCreateNode_Click(object sender, RoutedEventArgs e)
{
    // Load the XML document.
    XmlDocument document = new XmlDocument();
    document.Load(booksFile);
    // Get the root element.
    XmlElement root = document.DocumentElement;
    // Create the new nodes.
    XmlElement newBook = document.CreateElement("book");
    XmlElement newTitle = document.CreateElement("title");
    XmlElement newAuthor = document.CreateElement("author");
    XmlElement newCode = document.CreateElement("code");
    XmlText title = document.CreateTextNode("Beginning Visual C# 2015");
    XmlText author = document.CreateTextNode("Karli Watson et al");
    XmlText code = document.CreateTextNode("314418");
    XmlComment comment = document.CreateComment("The previous edition");
    // Insert the elements.
    newBook.AppendChild(comment);
    newBook.AppendChild(newTitle);
    newBook.AppendChild(newAuthor);
    newBook.AppendChild(newCode);
    newTitle.AppendChild(title);
    newAuthor.AppendChild(author);
    newCode.AppendChild(code);
    root.InsertAfter(newBook, root.LastChild);
    document.Save(booksFile);
}
```

4. Run the application and click Create. Then click Loop, and you should see the dialog box shown in Figure 19-5.



**FIGURE 19-5**

There is one important type of node that you didn't create in the preceding example: the `XmlAttribute`. That is left as an exercise at the end of the chapter.

### How It Works

The code in the `buttonCreateNode_Click` method is where all the creation of nodes happens. It creates eight new nodes, four of which are of type `XmlElement`, three of type `XmlText`, and one of type `XmlComment`.

All of the nodes are created with the method of the encapsulating  `XmlDocument` instance. The `XmlElement` nodes are created with the `CreateElement` method, the `XmlText` nodes are created with the `CreateTextNode` method, and the `XmlComment` node is created with the `CreateComment` method.

After the nodes have been created, they still need to be inserted into the XML tree. This is done with the `AppendChild` method on the element to which the new node should become a child. The only exception to this is the `book` node, which is the root node of all of the new nodes. This node is inserted into the tree using the `InsertAfter` method of the root object. Whereas all of the nodes that are inserted using `AppendChild` always become the last node in the list of child nodes, `InsertAfter` enables you to position the node where you want it.

### Deleting Nodes

Now that you've seen how to create new nodes, all that is left is to learn how to delete them again. All classes derived from `XmlNode` include two methods, shown in Table 19-6, that enable you to remove nodes from the document.

**TABLE 19-6:** Methods for Removing Nodes

METHOD	DESCRIPTION
RemoveAll	Removes all child nodes in the node on which it is called. What is slightly less obvious is that it also removes all attributes on the node because they are regarded as child nodes as well.
RemoveChild	Removes a single child in the node on which it is called. The method returns the node that has been removed from the document, but you can reinsert it if you change your mind.

The following short Try It Out extends the application you've been creating over the past two examples to include the capability to delete nodes. For now, it finds only the last instance of the book node and removes it.

### TRY IT OUT Removing Nodes: Chapter19\LoopThroughXmlDocument\MainWindow.xaml.cs

This example builds on the LoopThrough XmlDocument project you created earlier. The following steps enable you to find and remove the final instance of the book node:

1. Add a new button below the two that already exist and name it **buttonDeleteNode**. Set its Content property to Delete.
2. Double-click the new button and enter the following code:

```
private void buttonDeleteNode_Click(object sender, RoutedEventArgs e)
{
    // Load the XML document.
    XmlDocument document = new XmlDocument();
    document.Load(booksFile);
    // Get the root element.
    XmlElement root = document.DocumentElement;
    // Find the node. root is the <books> tag, so its last child
    // which will be the last <book> node.
    if (root.HasChildNodes)
    {
        XmlNode book = root.LastChild;
        // Delete the child.
        root.RemoveChild(book);
        // Save the document back to disk.
        document.Save(booksFile);
    }
}
```

3. Run the application. When you click the Delete Node button and then the Loop button, the last node in the tree will disappear.

### How It Works

After the initial steps to load the XML into the `XmlDocument` object, you examine the root element to see whether there are any child elements in the XML you loaded. If there are, you use the `LastChild`

property of the `XmlElement` class to get the last child. After that, removing the element is as simple as calling `RemoveChild`, which passes in the instance of the element you want to remove — in this case, the last child of the root element.

## Selecting Nodes

You now know how to move back and forth in an XML document, how to manipulate the values of the document, how to create new nodes, and how to delete them again. Only one thing remains in this section: how to select nodes without having to traverse the entire tree.

The `XmlNode` class includes two methods, described in Table 19-7, commonly used to select nodes from the document without running through every node in it: `SelectSingleNode` and `SelectNodes`, both of which use a special query language, called XPath, to select the nodes. You learn about that shortly.

**TABLE 19-7:** Methods for Selecting Nodes

METHOD	DESCRIPTION
<code>SelectSingleNode</code>	Selects a single node. If you create a query that fetches more than one node, only the first node will be returned.
<code>SelectNodes</code>	Returns a node collection in the form of an <code>XmlNodeList</code> class.

## CONVERTING XML TO JSON

We mentioned the JSON data language in the introduction to this chapter. There is limited support for JSON in the C# system libraries, but you can use a free third-party JSON library to work with JSON to convert XML to JSON and vice versa, and to do other manipulations with JSON similar to the .NET classes for XML. One such library available via the NuGet Package Manager in Visual Studio is the Newtonsoft JSON.NET package. Help and a full tutorial for this package are available at [www.json.net](http://www.json.net).

The following short Try It Out extends the application you've been creating over the previous examples in the chapter to include the capability to convert XML to JSON.

### TRY IT OUT

Convert: Chapter19\LoopThrough XmlDocument\MainWindow.xaml.cs

This example builds on the `LoopThrough XmlDocument` project you created earlier. The following steps enable you to find and remove the final instance of the book node:

1. In the Visual Studio menu, go to Tools → NuGet Package Manager → Manage NuGet Packages for Solution. Choose the `Newtonsoft.Json` package as shown in Figure 19-6. Click the Install button, and click OK on the Review Changes dialog to complete the installation.




FIGURE 19-6

2. Add a new button below the three that already exist and name it **buttonXMLtoJSON**. Set its Content property to XML>JSON.
3. Double-click the new button and enter the following code:

```
private void buttonXMLtoJSON_Click(object sender, RoutedEventArgs e)
{
    // Load the XML document.
    XmlDocument document = new XmlDocument();
    document.Load(booksFile);

    string json = Newtonsoft.Json.JsonConvert.SerializeXmlNode(document);

    textBlockResults.Text = json;
}
```

4. Run the application. Click the XML > JSON button. The JSON version of the book's data will appear in the main window as shown in Figure 19-7.




FIGURE 19-7

## How It Works

After the initial steps to load the XML into the  `XmlDocument`  object, you call the  `Newtonsoft JSON`  package method  `JsonConvert.SerializeXmlNode`  to convert your XML document to a text string in JSON format. Then you show the JSON text in any child elements in the XML you loaded. If there are any child elements, you use the  `textBlockResults`  window. As you can see, the JSON version of the book's data is more compact than the XML but a bit harder to read. That is why JSON is more often used for data transfer across the network rather than for storage in files that might be directly read by humans.

## SEARCHING XML WITH XPATH

XPath is a query language for XML documents, much as SQL is for relational databases. It is used by the two methods described in Table 19-7 that enable you to avoid the hassle of walking the entire tree of an XML document. It does take a little getting used to, however, because the syntax is nothing like SQL or C#.

**NOTE** *XPath is quite extensive, and only a small part of it is covered here so you can start selecting nodes. If you are interested in learning more, take a look at [www.w3.org/TR/xpath](http://www.w3.org/TR/xpath) and the Visual Studio help pages.*

To properly see XPath in action, you are going to use an XML file called `Elements.xml`, which contains a partial list of the chemical elements of the periodic table. You will find a subset of that XML listed in the “Selecting Nodes” Try It Out example later in the chapter, and it can be found in the download code for this chapter on this book’s website as `Elements.xml`.

Table 19-8 lists some of the most common operations you can perform with XPath. If nothing else is stated, the XPath query example makes a selection that is relative to the node on which it is performed. Where it is necessary to have a node name, you can assume the current node is the `<element>` node in the XML document.

TABLE 19-8: Common XPath Operations

PURPOSE	XPATH QUERY EXAMPLE
Select the current node.	.
Select the parent of the current node.	..
Select all child nodes of the current node.	*
Select all child nodes with a specific name — in this case, <code>title</code> .	Title
Select an attribute of the current node.	@Type
Select all attributes of the current node.	@*
Select a child node by index — in this case, the second element node.	element [2]
Select all the text nodes of the current node.	text ()
Select one or more grandchildren of the current node.	element/text ()
Select all nodes in the document with a particular name — in this case, all <code>mass</code> nodes.	//mass

*continues*

**TABLE 19-8** (continued)

PURPOSE	XPATH QUERY EXAMPLE
Select all nodes in the document with a particular name and a particular parent name — in this case, the parent name is element and the node name is name.	//element/name
Select a node where a value criterion is met — in this case, the element for which the name of the element is Hydrogen.	//element [name= 'Hydrogen']
Select a node where an attribute value criterion is met — in this case, the Type attribute is Noble_Gas.	//element [@Type= 'Noble_Gas']

In the following Try It Out, you'll create a small application that enables you to execute and see the results of a number of predefined queries, as well as enter your own queries.

### TRY IT OUT Selecting Nodes: Chapter19\XpathQuery\Elements.xml

As previously mentioned, this example uses an XML file called `Elements.xml`. You can download the file from the book's website or type part of it in from here:

```
<?xml version="1.0"?>
<elements>
  <!--First Non-Metal-->
  <element Type="Non-Metal">
    <name>Hydrogen</name>
    <symbol>H</symbol>
    <number>1</number>
    <specification>
      <mass>1.007825</mass>
      <density>0.0899 g/cm3</density>
    </specification>
  </element>
  <!--First Noble Gas-->
  <element Type="Noble Gas">
    <name>Helium</name>
    <symbol>He</symbol>
    <number>2</number>
    <specification>
      <mass>4.002602</mass>
      <density>0.1785 g/cm3</density>
    </specification>
  </element>
  <!--First Halogen-->
  <element Type="Halogen">
    <name>Fluorine</name>
    <symbol>F</symbol>
    <number>9</number>
    <specification>
```

```

<mass>18.998404</mass>
<density>1.696 g/cm3</density>
</specification>
</element>
<element Type="Noble Gas">
<name>Neon</name>
<symbol>Ne</symbol>
<number>10</number>
<specification>
<mass>20.1797</mass>
<density>0.901 g/cm3</density>
</specification>
</element>
</elements>

```

Save the XML file as `Elements.xml`. Remember to change the path to the file in the code that follows. This example is a small query tool that you can use to test different queries on the XML provided with the code.

Follow these steps to create a WPF application with querying capability:

1. Create a new WPF application and name it **XPath Query**.
2. Create the dialog box shown in Figure 19-8. Name the controls as shown in the figure, except for the button, which should be named **buttonExecute**. Wrap the `TextBlock` in a `ScrollViewer` control and set its `VerticalScrollBarVisibility` property to `Auto`.
3. Go to the Code view and include the `using` directive.
4. Add a private field to hold the document, and initialize it in the constructor:

```

private XmlDocument document;

public MainWindow()
{
    InitializeComponent();
    document = new XmlDocument();
    document.Load(@"C:\BegVCSharp\Chapter19\XML and Schema\Elements.xml");
}

```

5. You need a few helper methods to display the result of the queries in the `textBlockResult` `TextBlock`:

```

private void Update(XmlNodeList nodes)
{
    if (nodes == null || nodes.Count == 0)
    {
        textBlockResult.Text = "The query yielded no results";
        return;
    }
    string text = "";
    foreach (XmlNode node in nodes)
    {
        text = FormatText(node, text, "") + "\r\n";
    }
    textBlockResult.Text = text;
}

```

6. Update the constructor to display the entire contents of the XML file when the application starts:

```
public MainWindow()
{
    InitializeComponent();
    document = new XmlDocument();
    document.Load(@"C:\BegVCSharp\Chapter19\XML and Schema\Elements.xml");
    Update(document.DocumentElement.SelectNodes("."));
}
```

7. Copy and paste the two methods `FormatText` and `AddAttributes` from the previous Try It Out sections to the new project.
8. Finally, insert the code that executes whatever the user enters in the text box:

```
private void buttonExecute_Click(object sender, RoutedEventArgs e)
{
    try
    {
        XmlNodeList nodes = document.DocumentElement.SelectNodes(textBoxQuery.Text);
        Update(nodes);
    }
    catch (Exception err)
    {
        textBlockResult.Text = err.Message;
    }
}
```

9. Run the application and type the following query into the `textBoxQuery` textbox to select the element node that contains a node with the text Hydrogen:

```
element [name='Hydrogen']
```




FIGURE 19-8

### How It Works

The `buttonExecute_Click` method performs the queries. Because you can't know in advance if the queries typed into the `textBoxQuery` are going to yield a single node or multiple nodes, you must use

the `SelectNodes` method. This will either return an `XmlNodeList` object or throw one of the exceptions regarding XPath if the query used is illegal.

The `Update` method is responsible for looping through the content of the `XmlNodeList` selected by `SelectNodes`. It calls `FormatText` from the earlier examples with each of the nodes, and `FormatText` is responsible for recursively traversing the node tree and creating readable text you can use in the `textBoxResult` control.

In the exercises at the end of the chapter, you will find a number of additional XPath queries to try. Before you enter them into the `XPathQuery` application to see the result, try to determine for yourself the query's outcome.

## EXERCISES

**19.1** Change the Insert example in the “Creating Nodes” Try It Out section to insert an attribute called `Pages` with the value `1000+` on the `book` node.

**19.2** Determine the outcome of the following XPath queries and then verify your results by typing the queries into the `XPathQuery` application from the “Selecting Nodes” Try It Out. Remember that all of your queries are being executed on the `DocumentElement`, which is the `elements` node.

```
//elements
element
element[@Type='Noble Gas']
//mass
//mass..
element/specification[mass='20.1797']
element/name[text()='Neon']
```

Solution:

**19.3** On many Windows systems the default viewer of XML is a web browser. If you are using Internet Explorer you will see a nicely formatted view of the XML when you load the `Elements.xml` file into it. Why would it not be ideal to display the XML from our queries in a browser control instead of a text box?

**19.4** Use the Newtonsoft library to convert JSON to XML button as well (the reverse of the example shown in the chapter).

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
XML basics	XML documents are created from an XML declaration, XML namespaces, XML elements, and attributes. The XML declaration defines the XML version. XML namespaces are used to define vocabularies and XML elements and attributes are used to define the XML document content.
JSON basics	JSON is a data language used when transferring JavaScript and web services. JSON is more compact than the XML but harder to read.
XML schema	XML schemas are used to define the structure of XML documents. Schemas are especially useful when you need to exchange information with third parties. By agreeing on a schema for the data that is exchanged, you and the third party will be able to check that the documents are valid.
XML DOM	The Document Object Model (XML DOM) is the basis for .NET Framework classes provided for creating and manipulating XML.
JSON packages	You can use a JSON package such as Newtonsoft to convert XML to JSON and vice versa, and do other manipulations with JSON similar to the .NET classes for XML.
XPath	XPath is one of the possible ways to query data in XML documents. To use XPath, you must be familiar with the structure of the XML document in order to be able to select individual elements from it. Although XPath can be used on any well-formed XML document, the fact that you must know the structure of the document when you create the query means that ensuring that the document is valid also ensures that the query will work from document to document, as long as the documents are valid against the same schema.

# 20

## LINQ

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- LINQ to XML
- LINQ providers
- LINQ query syntax
- LINQ method syntax
- Lambda expressions
- Ordering query results
- Aggregates (Count, Sum, Min, Max, Average)
- SelectDistinctQuery
- Group queries
- Joins

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 20 download and individually named according to the names throughout the chapter.

This chapter introduces Language INtegrated Query (LINQ). LINQ is an extension to the C# language that integrates data query directly into the programming language itself.

Before LINQ this sort of work required writing a lot of looping code, and additional processing such as sorting or grouping the found objects required even more code that would differ depending on the data source. LINQ provides a portable, consistent way of querying, sorting, and grouping many different kinds of data (XML, JSON, SQL databases, collections of objects, web services, corporate directories, and more).

First you'll build on the previous chapter by learning the additional capabilities that the `System.Xml.Linq` namespace adds for creating XML. Then you'll get into the heart of LINQ by using query syntax, method syntax, lambda expressions, sorting, grouping, and joining related results.

LINQ is large enough that complete coverage of all its facilities and methods is beyond the scope of a beginning book. However, you will see examples of each of the different types of statements and operators you are likely to need as a user of LINQ, and you will be pointed to resources for more in-depth coverage as appropriate.

## LINQ TO XML

LINQ to XML is an alternate set of classes for XML that enables the use of LINQ for XML data and also makes certain operations with XML easier even if you are not using LINQ. We will look at a couple of specific cases where LINQ to XML has advantages over the XML DOM (Document Object Model) introduced in the previous chapter.

### LINQ to XML Functional Constructors

While you can create XML documents in code with the XML DOM, LINQ to XML provides an easier way to create XML documents called *functional construction*. In formal construction the constructor calls can be nested in a way that naturally reflects the structure of the XML document. In the following Try It Out, you use functional constructors to make a simple XML document containing customers and orders.

#### TRY IT OUT LINQ to XML: BegVCSharp\_20\_1\_LinqtoXmlConstructors

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called `BegVCSharp_20_1_LinqToXmlConstructors` in the directory `C:\BegVCSharp\Chapter20`.
2. Open the main source file `Program.cs`.
3. Add a reference to the `System.Xml.Linq` namespace to the beginning of `Program.cs`, as shown here:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Xml.Linq;
using System.Text;
using static System.Console;
```

4. Add the following code to the `Main()` method in `Program.cs`:

```
static void Main(string[] args)
{
    XDocument xdoc = new XDocument(
```

```

        new XElement("customers",
            new XElement("customer",
                new XAttribute("ID", "A"),
                new XAttribute("City", "New York"),
                new XAttribute("Region", "North America"),
                new XElement("order",
                    new XAttribute("Item", "Widget"),
                    new XAttribute("Price", 100)
                ),
                new XElement("order",
                    new XAttribute("Item", "Tire"),
                    new XAttribute("Price", 200)
                )
            ),
            new XElement("customer",
                new XAttribute("ID", "B"),
                new XAttribute("City", "Mumbai"),
                new XAttribute("Region", "Asia"),
                new XElement("order",
                    new XAttribute("Item", "Oven"),
                    new XAttribute("Price", 501)
                )
            )
        );
    WriteLine(xdoc);
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}

```

5. Compile and execute the program (you can just press F5 for Start Debugging). You will see the output shown here:

```

<customers>
    <customer ID="A" City="New York" Region="North America">
        <order Item="Widget" Price="100" />
        <order Item="Tire" Price="200" />
    </customer>
    <customer ID="B" City="Mumbai" Region="Asia">
        <order Item="Oven" Price="501" />
    </customer>
</customers>
Program finished, press Enter/Return to continue:

```

The XML document shown on the output screen contains a very simplified set of customer/order data. Note that the root element of the XML document is `<customers>`, which contains two nested `<customer>` elements. These in turn contain a number of nested `<order>` elements. The `<customer>` elements have two attributes, `City` and `Region`, and the `<order>` elements have `Item` and `Price` attributes.

Press Enter/Return to exit the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you might need to press Enter/Return twice.

## How It Works

The first step is to reference the `System.Xml.Linq` namespace. All of the XML examples in this chapter require that you add this line to your program:

```
using System.Xml.Linq;
```

Although the `System.Linq` namespace is included by default when you create a project, the `System.Xml.Linq` namespace is not included; you must add this line explicitly.

Next are the calls to the LINQ to XML constructors `XDocument()`, `XElement()`, and `XAttribute()`, which are nested inside one another as shown here:

```
XDocument xdoc = new XDocument(
    new XElement("customers",
        new XElement("customer",
            new XAttribute("ID", "A"),
            ...
        )
    )
);
```

Note that the code here looks like the XML itself, where the document contains elements and each element contains attributes and other elements. Take a look at each of these constructors in turn:

- `XDocument()` — The highest-level object in the LINQ to XML constructor hierarchy is `XDocument()`, which represents the complete XML document. It appears in your code here:

```
static void Main(string[] args)
{
    XDocument xdoc = new XDocument(
    ...
);
```

The parameter list for `XDocument()` is omitted in the previous code fragment so you can see where the `XDocument()` call begins and ends. Like all the LINQ to XML constructors, `XDocument()` takes an array of objects (`object[]`) as one of its parameters so that a number of other objects created by other constructors can be passed to it. All the other constructors you call in this program are parameters in the one call to the `XDocument()` constructor. The first (and only) parameter you pass in this program is the `XElement()` constructor.

- `XElement()` — An XML document must have a root element, so in most cases the parameter list of `XDocument()` will begin with an `XElement` object. The `XElement()` constructor takes the name of the element as a string, followed by a list of the XML objects contained within that element. Here, the root element is `"customers"`, which in turn contains a list of `"customer"` elements:

```
new XElement("customers",
    new XElement("customer",
    ...
),
...
)
```

The `"customer"` element does not contain any other XML elements. Instead, it contains three XML attributes, which are constructed with the `XAttribute()` constructor.

- `XAttribute()` — Here you add three XML attributes to the `"customer"` element, named `"ID"`, `"City"`, and `"Region"`:

```
new XAttribute("ID", "A"),
```

```

        new XAttribute("City", "New York"),
        new XAttribute("Region", "North America"),
    
```

Because an XML attribute is by definition a leaf XML node containing no other XML nodes, the `XAttribute()` constructor takes only the name of the attribute and its value as parameters. In this case, the three attributes generated are `ID="A"`, `City="New York"`, and `Region="North America"`.

- Other LINQ to XML constructors — Although you do not call them in this program, there are other LINQ to XML constructors for all the XML node types, such as `XDeclaration()` for the XML declaration at the start of an XML document, `XComment()` for an XML comment, and so on. These other constructors are not used often but are available if you need them for precise control over formatting an XML document.

Finishing up the explanation of the first example, you add two child `"order"` elements to the `"customer"` element following the `"ID"`, `"City"`, and `"Region"` attributes:

```

new XElement("order",
    new XAttribute("Item", "Widget"),
    new XAttribute("Price", 100)
),
new XElement("order",
    new XAttribute("Item", "Tire"),
    new XAttribute("Price", 200)
)
    
```

These order elements have `"Item"` and `"Price"` attributes but no other children.

Next, you display the contents of the `XDocument` to the console screen:

```
WriteLine(xdoc);
```

This prints the text of the XML document using the default `Tostring()` method of `XDocument()`.

Finally, you pause the screen so you can see the console output, and then wait until the user presses Enter:

```
Write("Program finished, press Enter/Return to continue:");
ReadLine();
```

After that your program exits the `Main()` method, which ends the program.

## Working with XML Fragments

Unlike the XML DOM, LINQ to XML works with XML fragments (partial or incomplete XML documents) in very much the same way as complete XML documents. When working with a fragment, you simply work with `XElement` as the top-level XML object instead of `XDocument`.

**NOTE** *The only restriction on working with XML fragments is that you cannot add some of the more esoteric XML node types that apply only to XML documents or XML fragments, such as `XComment` for XML comments, `XDeclaration` for the XML document declaration, and `XProcessingInstruction` for XML processing instructions.*

In the following Try It Out, you load, save, and manipulate an XML element and its child nodes, just as you did for an XML document.

### TRY IT OUT Working with XML Fragments: BegVCSharp\_20\_2\_XMLFragments

Follow these steps to create the example in Visual Studio 2015:

1. Either modify the previous example or create a new console application called BegVCSharp\_20\_2\_XMLFragments in the directory C:\BegVCSharp\Chapter20.
2. Open the main source file Program.cs.
3. Add a reference to the System.Xml.Linq namespace to the beginning of Program.cs, as shown here:

```
using System;
using System.Collections.Generic;
using System.Xml.Linq;
using System.Text;
using static System.Console;
```

This will already be present if you are modifying the previous example.

4. Add the XML element without the containing XML document constructor used in the previous examples to the Main() method in Program.cs:

```
static void Main(string[] args)
{
    XElement xcust =
        new XElement("customers",
            new XElement("customer",
                new XAttribute("ID", "A"),
                new XAttribute("City", "New York"),
                new XAttribute("Region", "North America"),
                new XElement("order",
                    new XAttribute("Item", "Widget"),
                    new XAttribute("Price", 100)
                ),
                new XElement("order",
                    new XAttribute("Item", "Tire"),
                    new XAttribute("Price", 200)
                )
            ),
            new XElement("customer",
                new XAttribute("ID", "B"),
                new XAttribute("City", "Mumbai"),
                new XAttribute("Region", "Asia"),
                new XElement("order",
                    new XAttribute("Item", "Oven"),
                    new XAttribute("Price", 501)
                )
            )
        );
}
```

5. After the XML element constructor code you added in the previous step, add the following code to save, load, and display the XML element:

```

        string xmlFileName =
    @"C:\BegVCSharp\Chapter20\BegVCSharp_20_2_XMLFragments\fragment.xml";
    xcust.Save(xmlFileName);
    XElement xcust2 = XElement.Load(xmlFileName);
    WriteLine("Contents of xcust:");
    WriteLine(xcust);
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}

```

**NOTE** Note the `xmlFileName` is an absolute path; your folder structure may differ and if so you should adjust the path to reflect the actual folder path on your computer.

6. Compile and execute the program (you can just press F5 for Start Debugging). You should see the following output in the console window:

```

Contents of XElement xcust:
<customers>
  <customer ID="A" City="New York" Region="North America">
    <order Item="Widget" Price="100" />
    <order Item="Tire" Price="200" />
  </customer>
  <customer ID="B" City="Mumbai" Region="Asia">
    <order Item="Oven" Price="501" />
  </customer>
</customers>
Program finished, press Enter/Return to continue:

```

Press Enter/Return to finish the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you might need to press Enter/Return twice.

### How It Works

Both `XElement` and `XDocument` inherit from the LINQ to XML `XContainer` class, which implements an XML node that can contain other XML nodes. Both classes also implement `Load()` and `Save()`, so most operations that can be performed on an `XDocument` in LINQ to XML can also be performed on an `XElement` instance and its children.

You simply create an `XElement` instance that has the same structure as the `XDocument` used in previous examples but omits the containing `XDocument`. All the operations for this particular program work the same with the `XElement` fragment.

`XElement` also supports the `Load()` and `Parse()` methods for loading XML from files and strings, respectively.

## LINQ PROVIDERS

LINQ to XML is just one example of a *LINQ provider*. Visual Studio 2015 and the .NET Framework 4.5 come with a number of built-in LINQ providers that provide query solutions for different types of data:

- **LINQ to Objects** — Provides queries on any kind of C# in-memory object, such as arrays, lists, and other collection types. All of the examples in the previous chapter use LINQ to Objects. However, you can use the techniques you learn in this chapter with all of the varieties of LINQ.
- **LINQ to XML** — As you have just seen, this provides creation and manipulation of XML documents using the same syntax and general query mechanism as the other LINQ varieties.
- **LINQ to Entities** — The Entity Framework is the newest set of data interface classes in .NET 4, recommended by Microsoft for new development. In this chapter you will add an ADO.NET Entity Framework data source to your Visual C# project, then query it using LINQ to Entities.
- **LINQ to Data Set** — The `DataSet` object was introduced in the first version of the .NET Framework. This variety of LINQ enables legacy .NET data to be queried easily with LINQ.
- **LINQ to SQL** — This is an alternative LINQ interface that has been superseded by LINQ to Entities.
- **PLINQ** — PLINQ, or Parallel LINQ, extends LINQ to Objects with a parallel programming library that can split up a query to execute simultaneously on a multicore processor.
- **LINQ to JSON** — Included in the Newtonsoft package you used in the previous chapter, this library supports creation and manipulation of JSON documents using the same syntax and general query mechanism as the other LINQ varieties.

With so many varieties of LINQ, it is impossible to cover them all in a beginning book, but the syntax and methods you will see apply to all. Let's next look at the LINQ query syntax using the LINQ to Objects provider.

## LINQ QUERY SYNTAX

In the following Try It Out, you use LINQ to create a query to find some data in a simple in-memory array of objects and print it to the console.

### TRY IT OUT First LINQ Program: BegVCSharp\_20\_3\_QuerySyntax\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called `BegVCSharp_20_3_QuerySyntax` in the directory `C:\BegVCSharp\Chapter20`, and then open the main source file `Program.cs`.
2. Notice that Visual Studio 2015 includes the `System.Linq` namespace by default in `Program.cs`:

```
using System;
using System.Collections.Generic;
```

```
using System.Linq;
using System.Text;
using static System.Console;
using System.Threading.Text;
```

3. Add the following code to the Main() method in Program.cs:

```
static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
"Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
    var queryResults =
        from n in names
        where n.StartsWith("S")
        select n;
    WriteLine("Names beginning with S:");
    foreach (var item in queryResults) {
        WriteLine(item);
    }
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}
```

4. Compile and execute the program (you can just press F5 for Start Debugging). You will see the names in the list beginning with S in the order they were declared in the array, as shown here:

```
Names beginning with S:
Smith
Smythe
Small
Singh
Samba
Program finished, press Enter/Return to continue:
```

Simply press Enter/Return to finish the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you may need to press Enter/Return twice. That finishes the program run.

## How It Works

The first step is to reference the System.Linq namespace, which is done automatically by Visual Studio 2015 when you create a project:

```
using System.Linq;
```

All the underlying base system support classes for LINQ reside in the System.Linq namespace. If you create a C# source file outside of Visual Studio 2015 or edit a project created from a previous version, you may have to add the using System.Linq directive manually.

The next step is to create some data, which is done in this example by declaring and initializing the array of names:

```
string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe", "Small",
"Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
```

This is a trivial set of data, but it is good to start with an example for which the result of the query is obvious. The actual LINQ query statement is the next part of the program:

```
var queryResults =  
    from n in names  
    where n.StartsWith("S")  
    select n;
```

That is an odd-looking statement, isn't it? It almost looks like something from a language other than C#, and the `from...where...select` syntax is deliberately similar to that of the SQL database query language. However, this statement is not SQL; it is indeed C#, as you saw when you typed in the code in Visual Studio 2015 — the `from`, `where`, and `select` were highlighted as keywords, and the odd-looking syntax is perfectly fine to the compiler.

The LINQ query statement in this program uses the LINQ declarative query syntax:

```
var queryResults =  
    from n in names  
    where n.StartsWith("S")  
    select n;
```

The statement has four parts: the result variable declaration beginning with `var`, which is assigned using a *query expression* consisting of the `from` clause; the `where` clause; and the `select` clause. Let's look at each of these parts in turn.

---

## Declaring a Variable for Results Using the `var` Keyword

The LINQ query starts by declaring a variable to hold the results of the query, which is usually done by declaring a variable with the `var` keyword:

```
var queryResult =
```

`var` is a keyword in C# created to declare a general variable type that is ideal for holding the results of LINQ queries. The `var` keyword tells the C# compiler to infer the type of the result based on the query. That way, you don't have to declare ahead of time what type of objects will be returned from the LINQ query — the compiler takes care of it for you. If the query can return multiple items, then it acts like a collection of the objects in the query data source (technically, it is not a collection; it just looks that way).

**NOTE** If you want to know the details, the query result will be a type that implements the `IEnumerable<T>` interface. The angle brackets with `T(<T>)` following `IEnumerable` indicate that it is a generic type. Generics are described in Chapter 12.

In this particular case, the compiler creates a special LINQ data type that provides an ordered list of strings (strings because the data source is a collection of strings).

By the way, the name `queryResult` is arbitrary — you can name the result anything you want. It could be `namesBeginningWithS` or anything else that makes sense in your program.

## Specifying the Data Source: `from` Clause

The next part of the LINQ query is the `from` clause, which specifies the data you are querying:

```
from n in names
```

Your data source in this case is `names`, the array of strings declared earlier. The variable `n` is just a stand-in for an individual element in the data source, similar to the variable name following a `foreach` statement. By specifying `from`, you are indicating that you are going to *query* a subset of the collection, rather than iterate through all the elements.

Speaking of iteration, a LINQ data source must be *enumerable* — that is, it must be an array or collection of items from which you can pick one or more elements to iterate through.

**NOTE** *Enumerable means the data source must support the `IEnumerable<T>` interface, which is supported for any C# array or collection of items.*

The data source cannot be a single value or object, such as a single `int` variable. You already have such a single item, so there is no point in querying it!

## Specify Condition: `where` Clause

In the next part of the LINQ query, you specify the condition for your query using the `where` clause, which looks like this:

```
where n.StartsWith("S")
```

Any Boolean (true or false) expression that can be applied to the items in the data source can be specified in the `where` clause. Actually, the `where` clause is optional and can even be omitted, but in almost all cases you will want to specify a `where` condition to limit the results to only the data you want. The `where` clause is called a *restriction operator* in LINQ because it restricts the results of the query.

Here, you specify that the name string starts with the letter S, but you could specify anything else about the string instead — for example, a length greater than 10 (`where n.Length > 10`) or containing a Q (`where n.Contains("Q")`).

## Selecting Items: `select` Clause

Finally, the `select` clause specifies which items appear in the result set. The `select` clause looks like this:

```
select n
```

The `select` clause is required because you must specify which items from your query appear in the result set. For this set of data, it is not very interesting because you have only one item, the name, in each element of the result set. You'll look at some examples with more complex objects in the result set where the usefulness of the `select` clause will be more apparent, but first, you need to finish the example.

## Finishing Up: Using the `foreach` Loop

Now you print out the results of the query. Like the array used as the data source, the results of a LINQ query like this are *enumerable*, meaning you can iterate through the results with a `foreach` statement:

```
WriteLine("Names beginning with S:");
foreach (var item in queryResults) {
    WriteLine(item);
}
```

In this case, you matched five names — Smith, Smythe, Small, Singh, and Samba — so that is what you display in the `foreach` loop.

## Deferred Query Execution

You may be thinking that the `foreach` loop really isn't part of LINQ itself — it's only looping through your results. While it's true that the `foreach` construct is not itself part of LINQ, nevertheless, it is the part of your code that actually executes the LINQ query! The assignment of the `queryResults` variable only saves a plan for executing the query; with LINQ, the data itself is not retrieved until the results are accessed. This is called *deferred query execution* or *lazy evaluation* of queries. Execution will be deferred for any query that produces a sequence — that is, a list — of results.

Now, back to the code. You've printed out the results; it's time to finish the program:

```
Write("Program finished, press Enter/Return to continue:");
ReadLine();
```

These lines just ensure that the results of the console program stay on the screen until you press a key, even if you press F5 instead of Ctrl+F5. You'll use this construct in most of the other LINQ examples as well.

## LINQ METHOD SYNTAX

There are multiple ways of doing the same thing with LINQ, as is often the case in programming. As noted, the previous example was written using the LINQ *query syntax*; in the next example, you will write the same program using LINQ's *method syntax* (also called *explicit syntax*, but the term *method syntax* is used here).

## LINQ Extension Methods

LINQ is implemented as a series of extension methods to collections, arrays, query results, and any other object that implements the `IEnumerable<T>` interface. You can see these methods with the Visual Studio IntelliSense feature. For example, in Visual Studio 2015, open the `Program.cs` file in the `FirstLINQquery` program you just completed and type in a new reference to the `names` array just below it:

```
string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe", "Small",
    "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
names.
```

Just as you type the period following `names`, you will see the methods available for `names` listed by the Visual Studio IntelliSense feature.

The `Where<T>` method and most of the other available methods are extension methods (as shown in the documentation appearing to the right of the `Where<T>` method, it begins with `extension`). You can see that they are LINQ extensions by commenting out the `using System.Linq` directive at the top; you will find that `Where<T>`, `Union<T>`, `Take<T>`, and most of the other methods in the list no longer appear. The `from...where...select` query expression you used in the previous example is translated by the C# compiler into a series of calls to these methods. When using the LINQ method syntax, you call these methods directly.

## Query Syntax versus Method Syntax

The query syntax is the preferred way of programming queries in LINQ, as it is generally easier to read and is simpler to use for the most common queries. However, it is important to have a basic understanding of the method syntax because some LINQ capabilities either are not available in the query syntax, or are just easier to use in the method syntax.

**NOTE** As the Visual Studio 2015 online help recommends, use query syntax whenever possible, and method syntax whenever necessary.

In this chapter, you will mostly use the query syntax, but the method syntax is pointed out in situations where it is needed, and you'll learn how to use the method syntax to solve the problem.

Most of the LINQ methods that use the method syntax require that you pass a method or function to evaluate the query expression. The method/function parameter is passed in the form of a delegate, which typically references an anonymous method.

Luckily, LINQ makes doing this much easier than it sounds! You create the method/function by using a *lambda expression*, which encapsulates the delegate in an elegant manner.

## Lambda Expressions

A lambda expression is a simple way to create a method on-the-fly for use in your LINQ query. It uses the `=>` operator, which declares the parameters for your method followed by the method logic all on a single line!

**NOTE** The term “lambda expression” comes from lambda calculus, which is a mathematical field important in programming language theory. Look it up if you’re mathematically inclined. Luckily you don’t need the math in order to use lambdas in C#!

For example, consider the lambda expression:

```
n => n < 0
```

This declares a method with a single parameter named `n`. The method returns `true` if `n` is less than zero, otherwise `false`. It’s dead simple. You don’t have to come up with a method name, put in a return statement, or wrap any code with curly braces.

Returning a true/false value like this is typical for methods used in LINQ lambdas, but it doesn’t have to be done. For example, here is a lambda that creates a method that returns the sum of two variables. This lambda uses multiple parameters:

```
(a, b) => a + b
```

This declares a method with two parameters named `a` and `b`. The method logic returns the sum of `a` and `b`. You don’t have to declare what type `a` and `b` are. They can be `int` or `double` or `string`. The C# compiler infers the types.

Finally, consider this lambda expression:

```
n => n.StartsWith("S")
```

This method returns `true` if `n` starts with the letter S, otherwise `false`. Try this out in an actual program to see this more clearly.

### TRY IT OUT Using LINQ Method Syntax and Lambda Expressions: BegVCSharp\_20\_4\_MethodSyntax\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. You can either modify the previous example or create a new console application called `BegVCSharp_20_4_MethodSyntax` in the directory `C:\BegVCSharp\Chapter20`. Open the main source file `Program.cs`.
2. Again, Visual Studio 2015 includes the `Linq` namespace automatically in `Program.cs`:  
`using System.Linq;`

3. Add the following code to the `Main()` method in `Program.cs`:

```
static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
    "Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
    var queryResults = names.Where(n => n.StartsWith("S"));
    WriteLine("Names beginning with S:");
    foreach (var item in queryResults) {
        WriteLine(item);
    }
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}
```

4. Compile and execute the program (you can just press F5). You will see the same output of names in the list beginning with S, in the order they were declared in the array, as shown here:

```
Names beginning with S:
Smith
Smythe
Small
Singh
Samba
Program finished, press Enter/Return to continue:
```

## How It Works

As before, the `System.Linq` namespace is referenced automatically by Visual Studio 2015:

```
using System.Linq;
```

The same source data as before is created again by declaring and initializing the array of names:

```
string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe", "Small", "Ruiz",
    "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
```

The part that is different is the LINQ query, which is now a call to the `Where()` method instead of a query expression:

```
var queryResults = names.Where(n => n.StartsWith("S"));
```

The C# compiler compiles the lambda expression `n => n.StartsWith("S")` into an anonymous method that is executed by `Where()` on each item in the `names` array. If the lambda expression returns true for an item, that item is included in the result set returned by `Where()`. The C# compiler infers that the `Where()` method should accept `string` as the input type for each item from the definition of the input source (the `names` array, in this case).

Well, a lot is going on in that one line, isn't it? For the simplest type of query like this, the method syntax is actually shorter than the query syntax because you do not need the `from` or `select` clauses; however, most queries are more complex than this.

The rest of the example is the same as the previous one — you print out the results of the query in a `foreach` loop and pause the output so you can see it before the program finishes execution:

```
foreach (var item in queryResults) {
    WriteLine(item);
}
```

```
Write("Program finished, press Enter/Return to continue:");
ReadLine();
```

An explanation of these lines isn't repeated here because that was covered in the "How It Works" section following the first example in the chapter. Let's move on to explore how to use more of LINQ's capabilities.

## ORDERING QUERY RESULTS

Once you have located some data of interest with a `where` clause (or `Where()` method invocation), LINQ makes it easy to perform further processing — such as reordering the results — on the resulting data. In the following Try It Out, you put the results from your first query in alphabetical order.

### TRY IT OUT Ordering Query Results: BegVCSharp\_20\_5\_OrderQueryResults\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. You can either modify the `QuerySyntax` example or create a new console application project called `BegVCSharp_20_5_OrderQueryResults` in the directory `C:\BegVCSharp\Chapter20`.
2. Open the main source file `Program.cs`. As before, Visual Studio 2015 includes the `using System.Linq;` namespace directive automatically in `Program.cs`.
3. Add the following code to the `Main()` method in `Program.cs`:

```
static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
"Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
    var queryResults =
        from n in names
        where n.StartsWith("S")
        orderby n
        select n;
    WriteLine("Names beginning with S ordered alphabetically:");
    foreach (var item in queryResults) {
        WriteLine(item);
    }
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}
```

4. Compile and execute the program. You will see the names in the list beginning with S in alphabetical order, as shown here:

```
Names beginning with S:
Samba
Singh
Small
```

---

```

Smith
Smythe
Program finished, press Enter/Return to continue:

```

## How It Works

This program is nearly identical to the previous example, except for one additional line added to the query statement:

```

var queryResults =
    from n in names
    where n.StartsWith("S")
    orderby n
    select n;

```

---

## UNDERSTANDING THE ORDERBY CLAUSE

The orderby clause looks like this:

```
orderby n
```

Like the where clause, the orderby clause is optional. Just by adding one line, you can order the results of any arbitrary query, which would otherwise require at least several lines of additional code and probably additional methods or collections to store the results of the reordered result, depending on the sorting algorithm you chose to implement. If multiple types needed to be sorted, you would have to implement a set of ordering methods for each one. With LINQ, you don't need to worry about any of that; just add one additional clause in the query statement and you're done.

By default, orderby orders in ascending order (A to Z), but you can specify descending order (from Z to A) simply by adding the descending keyword:

```
orderby n descending
```

This orders the example results as follows:

```

Smythe
Smith
Small
Singh
Samba

```

Plus, you can order by any arbitrary expression without having to rewrite the query; for example, to order by the last letter in the name instead of normal alphabetical order, you just change the orderby clause to the following:

```
orderby n.Substring(n.Length - 1)
```

This results in the following output:

```

Samba
Smythe
Smith
Singh
Small

```

**NOTE** The last letters are in alphabetical order (a, e, h, h, l). However, you will notice that the execution is implementation-dependent, meaning there's no guarantee of order beyond what is specified in the `orderby` clause. The last letter is the only letter considered, so, in this case, Smith came before Singh.

## QUERYING A LARGE DATA SET

All this LINQ syntax is well and good, you may be saying, but what is the point? You can see the expected results clearly just by looking at the source array, so why go to all this trouble to query something that is obvious by just looking? As mentioned earlier, sometimes the results of a query are not so obvious. In the following Try It Out, you create a very large array of numbers and query it using LINQ.

### TRY IT OUT Querying a Large Data Set: BegVCSharp\_20\_6\_LargeNumberQuery\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called BegVCSharp\_20\_6\_LargeNumberQuery in the directory C:\BegVCSharp\Chapter20. As before, when you create the project, Visual Studio 2015 already includes the `Linq` namespace method in `Program.cs`:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using static System.Console;
```

2. Add the following code to the `Main()` method:

```
static void Main(string[] args)
{
    int[] numbers = GenerateLotsOfNumbers(12045678);
    var queryResults =
        from n in numbers
        where n < 1000
        select n
    ;
    WriteLine("Numbers less than 1000:");
    foreach (var item in queryResults)
    {
        WriteLine(item);
    }
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}
```

3. Add the following method to generate the list of random numbers:

```
private static int[] GenerateLotsOfNumbers(int count)
{
    Random generator = new Random(0);
    int[] result = new int[count];
    for (int i = 0; i < count; i++)
    {
        result[i] = generator.Next();
    }
    return result;
}
```

4. Compile and execute the program. You will see a list of numbers less than 1,000, as shown here:

Numbers less than 1000:

```
714
24
677
350
257
719
584
```

Program finished, press Enter/Return to continue:

## How It Works

As before, the first step is to reference the `System.Linq` namespace, which is done automatically by Visual Studio 2015 when you create the project:

```
using System.Linq;
```

The next step is to create some data, which is done in this example by creating and calling the `GenerateLotsOfNumbers()` method:

```
int[] numbers = GenerateLotsOfNumbers(12345678);
private static int[] GenerateLotsOfNumbers(int count)
{
    Random generator = new Random(0);
    int[] result = new int[count];
    for (int i = 0; i < count; i++)
    {
        result[i] = generator.Next();
    }
    return result;
}
```

This is not a trivial set of data — there are more than 12 million numbers in the array! In one of the exercises at the end of the chapter, you will change the `size` parameter passed to the `GenerateLotsOfNumbers()` method to generate variously sized sets of random numbers and see how this affects the query results. As you will see when doing the exercises, the size shown here of 12,345,678 is just large enough for the program to generate some random numbers less than 1,000, in order to have results to show for this first query.

The values should be randomly distributed over the range of a signed integer (from zero to more than two billion). By creating the random number generator with a seed of 0, you ensure that the same set

of random numbers is created each time and is repeatable, so you get the same query results as shown here, but what those query results are is unknown until you try some queries. Luckily, LINQ makes those queries easy!

The query statement itself is similar to what you did with the names before, selecting some numbers that meet a condition (in this case, numbers less than 1,000):

```
var queryResults =
    from n in numbers
    where n < 1000
    select n
```

The `orderby` clause isn't needed here and would add extra processing time (not noticeably for this query, but more so as you vary the conditions in the next example).

You print out the results of the query with a `foreach` statement, just as in the previous example:

```
WriteLine("Numbers less than 1000:");
foreach (var item in queryResults) {
    WriteLine(item);
}
```

Again, output to the console and read a character to pause the output:

```
Write("Program finished, press Enter/Return to continue:");
ReadLine();
```

The pause code appears in all the following examples but isn't shown again because it is the same for each one.

It is very easy with LINQ to change the query conditions to explore different characteristics of the data set. However, depending on how many results the query returns, it may not make sense to print all the results each time. In the next section you'll see how LINQ provides aggregate operators to deal with that issue.

---

## USING AGGREGATE OPERATORS

Often, a query returns more results than you might expect. For example, if you were to change the condition of the large-number query program you just created to list the numbers greater than 1,000, rather than the numbers less than 1,000, there would be so many query results that the numbers would not stop printing!

Luckily, LINQ provides a set of aggregate operators that enable you to analyze the results of a query without having to loop through them all. Table 20-1 shows the most commonly used aggregate operators for a set of numeric results such as those from the large-number query. These may be familiar to you if you have used a database query language such as SQL.

**TABLE 20-1:** Aggregate Operators for Numeric Results

OPERATOR	DESCRIPTION
Count ()	Count of results
Min ()	Minimum value in results
Max ()	Maximum value in results
Average ()	Average value of numeric results
Sum ()	Total of all of numeric results

There are more aggregate operators, such as `Aggregate()`, for executing arbitrary code in a manner that enables you to code your own aggregate function. However, those are for advanced users and therefore beyond the scope of this book.

**NOTE** Because the aggregate operators return a simple scalar type instead of a sequence for their results, their use forces immediate execution of query results with no deferred execution.

In the following Try It Out, you modify the large-number query and use aggregate operators to explore the result set from the greater-than version of the large-number query using LINQ.

### TRY IT OUT      Numeric Aggregate Operators: BegVCSharp\_20\_7\_NumericAggregates\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. For this example, you can either modify the `LargeNumberQuery` example you just made or create a new console project named `BegVCSharp_20_7_NumericAggregates` in the directory `C:\BegVCSharp\Chapter20`.
2. As before, when you create the project, Visual Studio 2015 includes the `Linq` namespace method in `Program.cs`. You just need to modify the `Main()` method as shown in the following code and in the rest of this Try It Out. As with the previous example, the `orderby` clause is not used in this query. However, the condition on the `where` clause is the opposite of the previous example (the numbers are greater than 1,000 (`n > 1000`), instead of less than 1,000):

```
static void Main(string[] args)
{
    int[] numbers = GenerateLotsOfNumbers(12345678);
    WriteLine("Numeric Aggregates");
    var queryResults =
```

```
        from n in numbers
        where n > 1000
        select n
    ;
    WriteLine("Count of Numbers > 1000");
    WriteLine(queryResults.Count());
    WriteLine("Max of Numbers > 1000");
    WriteLine(queryResults.Max());
    WriteLine("Min of Numbers > 1000");
    WriteLine(queryResults.Min());
    WriteLine("Average of Numbers > 1000");
    WriteLine(queryResults.Average());
    WriteLine("Sum of Numbers > 1000");
    WriteLine(queryResults.Sum(n => (long) n));
    Write("Program finished, press Enter/Return to continue:");
    ReadLine();
}
```

3. If it is not already present, add the same `GenerateLotsOfNumbers()` method used in the previous example:

```
private static int[] GenerateLotsOfNumbers(int count)
{
    Random generator = new Random(0);
    int[] result = new int[count];
    for (int i = 0; i < count; i++)
    {
        result[i] = generator.Next();
    }
    return result;
}
```

4. Compile and execute. You will see the count, minimum, maximum, and average values as shown here:

```
Numeric Aggregates
Count of Numbers > 1000
12345671
Maximum of Numbers > 1000
2147483591
Minimum of Numbers > 1000
1034
Average of Numbers > 1000
1073643807.50298
Sum of Numbers > 1000
13254853218619179
Program finished, press Enter/Return to continue:
```

This query produces many more results than the previous example (more than 12 million). Using `orderby` on this result set would definitely have a noticeable impact on performance! The largest number (maximum) in the result set is over two billion and the smallest (minimum) is just over one thousand, as expected. The average is around one billion, near the middle of the range of possible values. Looks like the `Random()` function generates a good distribution of numbers!

## How It Works

The first part of the program is exactly the same as the previous example, with the reference to the `System.Linq` namespace, and the use of the `GenerateLotsOfNumbers()` method to generate the source data:

```
int[] numbers = GenerateLotsOfNumbers(12345678);
```

The query is the same as the previous example, except for changing the `where` condition from less than to greater than:

```
var queryResults =
    from n in numbers
    where n > 1000
    select n;
```

As noted before, this query using the greater-than condition produces many more results than the less-than query (with this particular data set). By using the aggregate operators, you are able to explore the results of the query without having to print out each result or do a comparison in a `foreach` loop. Each one appears as a method that can be called on the result set, similar to methods on a collection type.

Look at the use of each aggregate operator:

► `Count()`:

```
WriteLine("Count of Numbers > 1000");
WriteLine(queryResults.Count());
```

`Count()` returns the number of rows in the query results — in this case, 12,345,671 rows.

► `Max()`:

```
WriteLine("Max of Numbers > 1000");
WriteLine(queryResults.Max());
```

`Max()` returns the maximum value in the query results — in this case, a number larger than two billion: 2,147,483,591, which is very close to the maximum value of an `int` (`int.MaxValue` or 2,147,483,647).

► `Min()`:

```
WriteLine("Min of Numbers > 1000");
WriteLine(queryResults.Min());
```

`min()` returns the minimum value in the query results — in this case, 1,034.

► `Average()`:

```
WriteLine("Average of Numbers > 1000");
WriteLine(queryResults.Average());
```

`Average()` returns the average value of the query results, which in this case is 1,073,643,807.50298, a value very close to the middle of the range of possible values from 1,000 to more than two billion. This is rather meaningless with an arbitrary set of large numbers, but it shows the kind of query result analysis that is possible. You'll look at a more practical use of these operators with some business-oriented data in the last part of the chapter.

```
>     Sum():
        WriteLine("Sum of Numbers > 1000");
        WriteLine(queryResults.Sum(n => (long) n));
```

You passed the lambda expression `n => (long) n` to the `Sum()` method call to get the sum of all the numbers. Although `Sum()` has a no-parameter overload, like `Count()`, `Min()`, `Max()`, and so on, using that version of the method call would cause an overflow error because there are so many large numbers in the data set that the sum of all of them would be too large to fit into a standard 32-bit `int`, which is what the no-parameter version of `Sum()` returns. The lambda expression enables you to convert the result of `Sum()` to a long 64-bit integer, which is what you need to hold the total of over 13 quadrillion without overflow — 13,254,853,218,619,179 lambda expressions enable you to perform this kind of fix-up easily.

**NOTE** In addition to `Count()`, which returns a 32-bit `int`, LINQ also provides a `LongCount()` method that returns the count of query results in a 64-bit integer. That is a special case, however — all the other operators require a lambda or a call to a conversion method if a 64-bit version of the number is needed.

## USING THE SELECT DISTINCT QUERY

Another type of query that those of you familiar with the SQL data query language will recognize is the `SELECT DISTINCT` query, in which you search for the unique values in your data — that is, the query removes any repeated values from the result set. This is a fairly common need when working with queries.

Suppose you need to find the distinct regions in the customer data used in the previous examples. There is no separate region list in the data you just used, so you need to find the unique, nonrepeating list of regions from the customer list itself. LINQ provides a `Distinct()` method that makes it easy to find this data. You'll use it in the following Try It Out.

### TRY IT OUT    Projection: Select Distinct Query: BegVCSharp\_20\_8\_SelectDistinctQuery\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called `BegVCSharp_20_8_SelectDistinctQuery` in the directory `C:\BegVCSharp\Chapter20`.
2. Enter this code to create the `Customer` class and the initialization of the `customers` list (`List<Customer> customers`):

```
class Customer
{
    public string ID { get; set; }
```

```

        public string City { get; set; }
        public string Country { get; set; }
        public string Region { get; set; }
        public decimal Sales { get; set; }

        public override string ToString()
        {
            return "ID: " + ID + " City: " + City +
                " Country: " + Country +
                " Region: " + Region +
                " Sales: " + Sales;
        }
    }
}

class Program
{
    static void Main(string[] args)
    {
        List<Customer> customers = new List<Customer> {
            new Customer { ID="A", City="New York", Country="USA",
Region="North America", Sales=9999 },
            new Customer { ID="B", City="Mumbai", Country="India",
Region="Asia", Sales=8888 },
            new Customer { ID="C", City="Karachi", Country="Pakistan",
Region="Asia", Sales=7777 },
            new Customer { ID="D", City="Delhi", Country="India",
Region="Asia", Sales=6666 },
            new Customer { ID="E", City="São Paulo", Country="Brazil",
Region="South America", Sales=5555 },
            new Customer { ID="F", City="Moscow", Country="Russia",
Region="Europe", Sales=4444 },
            new Customer { ID="G", City="Seoul", Country="Korea",
Region="Asia", Sales=3333 },
            new Customer { ID="H", City="Istanbul", Country="Turkey",
Region="Asia", Sales=2222 },
            new Customer { ID="I", City="Shanghai", Country="China",
Region="Asia", Sales=1111 },
            new Customer { ID="J", City="Lagos", Country="Nigeria",
Region="Africa", Sales=1000 },
            new Customer { ID="K", City="Mexico City", Country="Mexico",
Region="North America", Sales=2000 },
            new Customer { ID="L", City="Jakarta", Country="Indonesia",
Region="Asia", Sales=3000 },
            new Customer { ID="M", City="Tokyo", Country="Japan",
Region="Asia", Sales=4000 },
            new Customer { ID="N", City="Los Angeles", Country="USA",
Region="North America", Sales=5000 },
            new Customer { ID="O", City="Cairo", Country="Egypt",
Region="Africa", Sales=6000 },
            new Customer { ID="P", City="Tehran", Country="Iran",
Region="Asia", Sales=7000 },
            new Customer { ID="Q", City="London", Country="UK",
Region="Europe", Sales=8000 },
            new Customer { ID="R", City="Beijing", Country="China",
Region="Asia", Sales=9000 },
            new Customer { ID="S", City="Bogotá", Country="Colombia",
}

```

```
Region="South America", Sales=1001 },
    new Customer { ID="T", City="Lima", Country="Peru",
Region="South America", Sales=2002 }
};
```

3. In the `Main()` method, following the initialization of the `customers` list, enter (or modify) the query as shown here:

```
var queryResults = customers.Select(c => c.Region).Distinct();
```

4. Finish the remaining code in the `Main()` method as shown here.

```
foreach (var item in queryResults)
{
    WriteLine(item);
}
Write("Program finished, press Enter/Return to continue:");
ReadLine();
```

5. Compile and execute the program. You will see the unique regions where customers exist:

```
North America
Asia
South America
Europe
Africa
Program finished, press Enter/Return to continue:
```

## How It Works

The `Customer` class and `customers` list initialization are the same as in the previous example. In the query statement, you call the `Select()` method with a simple lambda expression to select the region from the `Customer` objects, and then call `Distinct()` to return only the unique results from `Select()`:

```
var queryResults = customers.Select(c => c.Region).Distinct();
```

Because `Distinct()` is available only in method syntax, you make the call to `Select()` using method syntax. However, you can call `Distinct()` to modify a query made in the query syntax as well:

```
var queryResults = (from c in customers select c.Region).Distinct();
```

Because query syntax is translated by the C# compiler into the same series of LINQ method calls as used in the method syntax, you can mix and match if it makes sense for readability and style.

---

## ORDERING BY MULTIPLE LEVELS

Now that you are dealing with objects with multiple properties, you might be able to envision a situation where ordering the query results by a single field is not enough. What if you wanted to query your customers and order the results alphabetically by region, but then order alphabetically by country or city name within a region? LINQ makes this very easy, as you will see in the following Try It Out.

**TRY IT OUT****Ordering By Multiple Levels: BegVCSharp\_20\_9\_MultiLevelOrdering\Program.cs**

Follow these steps to create the example in Visual Studio 2015:

1. Modify the previous example, BegVCSharp\_20\_8\_SelectDistinctQuery, or create a new console application called BegVCSharp\_20\_9\_MultiLevelOrdering in the directory C:\BegVCSharp\Chapter20.
2. Create the `Customer` class and the initialization of the `customers` list (`List<Customer> customers`) as shown in the BegVCSharp\_20\_8\_SelectDistinctQuery example; this code is exactly the same as in previous examples.
3. In the `Main()` method, following the initialization of the `customers` list, enter the following query:

```
var queryResults =
    from c in customers
    orderby c.Region, c.Country, c.City
    select new { c.ID, c.Region, c.Country, c.City }
;
```

4. The results processing loop and the remaining code in the `Main()` method are the same as in previous examples.
5. Compile and execute the program. You will see the selected properties from all customers ordered alphabetically by region first, then by country, and then by city, as shown here:

```
{
  ID = O, Region = Africa, Country = Egypt, City = Cairo
{
  ID = J, Region = Africa, Country = Nigeria, City = Lagos
{
  ID = R, Region = Asia, Country = China, City = Beijing
{
  ID = I, Region = Asia, Country = China, City = Shanghai
{
  ID = D, Region = Asia, Country = India, City = Delhi
{
  ID = B, Region = Asia, Country = India, City = Mumbai
{
  ID = L, Region = Asia, Country = Indonesia, City = Jakarta
{
  ID = P, Region = Asia, Country = Iran, City = Tehran
{
  ID = M, Region = Asia, Country = Japan, City = Tokyo
{
  ID = G, Region = Asia, Country = Korea, City = Seoul
{
  ID = C, Region = Asia, Country = Pakistan, City = Karachi
{
  ID = H, Region = Asia, Country = Turkey, City = Istanbul
{
  ID = F, Region = Europe, Country = Russia, City = Moscow
{
  ID = Q, Region = Europe, Country = UK, City = London
{
  ID = K, Region = North America, Country = Mexico, City = Mexico City
{
  ID = N, Region = North America, Country = USA, City = Los Angeles
{
  ID = A, Region = North America, Country = USA, City = New York
{
  ID = E, Region = South America, Country = Brazil, City = São Paulo
{
  ID = S, Region = South America, Country = Colombia, City = Bogotá
{
  ID = T, Region = South America, Country = Peru, City = Lima }
```

Program finished, press Enter/Return to continue:

## How It Works

The `Customer` class and `customers` list initialization are the same as in previous examples. In this query you have no `where` clause because you want to see all the customers, but you simply list the fields you want to sort by order in a comma-separated list in the `orderby` clause:

```
orderby c.Region, c.Country, c.City
```

Couldn't be easier, could it? It seems a bit counterintuitive that a simple list of fields is allowed in the `orderby` clause but not in the `select` clause, but that is how LINQ works. It makes sense if you realize that the `select` clause is creating a new object but the `orderby` clause, by definition, operates on a field-by-field basis.

You can add the `descending` keyword to any of the fields listed to reverse the sort order for that field. For example, to order this query by ascending region but descending country, simply add `descending` following `Country` in the list, like this:

```
orderby c.Region, c.Country descending, c.City
```

With `descending` added, you see following output:

```
{ ID = J, Region = Africa, Country = Nigeria, City = Lagos }
{ ID = O, Region = Africa, Country = Egypt, City = Cairo }
{ ID = H, Region = Asia, Country = Turkey, City = Istanbul }
{ ID = C, Region = Asia, Country = Pakistan, City = Karachi }
{ ID = G, Region = Asia, Country = Korea, City = Seoul }
{ ID = M, Region = Asia, Country = Japan, City = Tokyo }
{ ID = P, Region = Asia, Country = Iran, City = Tehran }
{ ID = L, Region = Asia, Country = Indonesia, City = Jakarta }
{ ID = D, Region = Asia, Country = India, City = Delhi }
{ ID = B, Region = Asia, Country = India, City = Mumbai }
{ ID = R, Region = Asia, Country = China, City = Beijing }
{ ID = I, Region = Asia, Country = China, City = Shanghai }
{ ID = Q, Region = Europe, Country = UK, City = London }
{ ID = F, Region = Europe, Country = Russia, City = Moscow }
{ ID = N, Region = North America, Country = USA, City = Los Angeles }
{ ID = A, Region = North America, Country = USA, City = New York }
{ ID = K, Region = North America, Country = Mexico, City = Mexico City }
{ ID = T, Region = South America, Country = Peru, City = Lima }
{ ID = S, Region = South America, Country = Colombia, City = Bogotá }
{ ID = E, Region = South America, Country = Brazil, City = São Paulo }
Program finished, press Enter/Return to continue:
```

Note that the cities in India and China are still in ascending order even though the country ordering has been reversed.

---

## USING GROUP QUERIES

A group query divides the data into groups and enables you to sort, calculate aggregates, and compare by group. These are often the most interesting queries in a business context (the ones that really drive decision-making). For example, you might want to compare sales by country or by region to decide where to open another store or hire more staff. You'll do that in the next Try It Out.

**TRY IT OUT** Using a Group Query: BegVCSharp\_20\_10\_GroupQuery\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called BegVCSharp\_20\_10\_GroupQuery in the directory C:\BegVCSharp\Chapter20.
2. Create the Customer class and the initialization of the customers list (List<Customer> customers), as shown in the BegVCSharp\_20\_8\_SelectDistinctQuery example; this code is exactly the same as previous examples.
3. In the Main() method, following the initialization of the customers list, enter two queries:

```
var queryResults =
    from c in customers
    group c by c.Region into cg
    select new { TotalSales = cg.Sum(c => c.Sales), Region = cg.Key }
;
var orderedResults =
    from cg in queryResults
    orderby cg.TotalSales descending
    select cg
;
```

4. Continuing in the Main() method, add the following print statement and foreach processing loop:

```
WriteLine("Total\t: By\nSales\t: Region\n-----\t -----");
foreach (var item in orderedResults)
{
    WriteLine($"{item.TotalSales}\t: {item.Region}");
}
```

5. The results processing loop and the remaining code in the Main() method are the same as in previous examples. Compile and execute the program. Here are the group results:

```
Total : By
Sales : Region
-----
52997 : Asia
16999 : North America
12444 : Europe
8558 : South America
7000 : Africa
```

### How It Works

The Customer class and customers list initialization are the same as in previous examples.

The data in a group query is grouped by a key field, the field for which all the members of each group share a value. In this example, the key field is the Region:

```
group c by c.Region
```

You want to calculate a total for each group, so you group into a new result set named cg:

```
group c by c.Region into cg
```

In the `select` clause, you project a new anonymous type whose properties are the total sales (calculated by referencing the `cg` result set) and the key value of the group, which you reference with the special group `Key`:

```
select new { TotalSales = cg.Sum(c => c.Sales), Region = cg.Key }
```

The group result set implements the LINQ `IGrouping` interface, which supports the `Key` property. You almost always want to reference the `Key` property in some way in processing group results, because it represents the criteria by which each group in your data was created.

You want to order the result in descending order by `TotalSales` field so you can see which region has the highest total sales, next highest, and so on. To do that, you create a second query to order the results from the group query:

```
var orderedResults =
    from cg in queryResults
    orderby cg.TotalSales descending
    select cg
;
```

The second query is a standard `select` query with an `orderby` clause, as you have seen in previous examples; it does not make use of any LINQ group capabilities except that the data source comes from the previous group query.

Next, you print out the results, with a little bit of formatting code to display the data with column headers and some separation between the totals and the group names:

```
WriteLine("Total\t: By\nSales\t: Region\n---\t ---");
foreach (var item in orderedResults)
{
    WriteLine($"{item.TotalSales}\t: {item.Region}");
}
```

This could be formatted in a more sophisticated way with field widths and by right-justifying the totals, but this is just an example so you don't need to bother — you can see the data clearly enough to understand what the code is doing.

## USING JOINS

A data set such as the `customers` and `orders` list you just created, with a shared key field (ID), enables a `join` query, whereby you can query related data in both lists with a single query, joining the results together with the key field. This is similar to the `JOIN` operation in the SQL data query language; and as you might expect, LINQ provides a `join` command in the query syntax, which you will use in the following Try It Out.

### TRY IT OUT Join Query: BegVCSharp\_20\_11\_JoinQuery\Program.cs

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called `BegVCSharp_20_11_JoinQuery` in the directory `C:\BegVCSharp\Chapter20`.

2. Copy the code to create the Customer class, the Order class, and the initialization of the customers list (`List<Customer> customers`) and orders list (`List<Order> orders`) from the previous example; this code is the same.
3. In the `Main()` method, following the initialization of the `customers` and `orders` list, enter this query:

```
var queryResults =
    from c in customers
    join o in orders on c.ID equals o.ID
    select new { c.ID, c.City, SalesBefore = c.Sales, NewOrder = o.Amount,
                SalesAfter = c.Sales+o.Amount };
```

4. Finish the program using the standard `foreach` query processing loop you used in earlier examples:

```
foreach (var item in queryResults)
{
    WriteLine(item);
}
```

5. Compile and execute the program. Here's the output:

```
{ ID = P, City = Tehran, SalesBefore = 7000, NewOrder = 100, SalesAfter = 7100 }
{ ID = Q, City = London, SalesBefore = 8000, NewOrder = 200, SalesAfter = 8200 }
{ ID = R, City = Beijing, SalesBefore = 9000, NewOrder = 300, SalesAfter = 9300 }
{ ID = S, City = Bogotá, SalesBefore = 1001, NewOrder = 400, SalesAfter = 1401 }
{ ID = T, City = Lima, SalesBefore = 2002, NewOrder = 500, SalesAfter = 2502 }
Program finished, press Enter/Return to continue:
```

## How It Works

The code declaring and initializing the `Customer` class, the `Order` class, and the `customers` and `orders` lists is the same as in the previous example.

The query uses the `join` keyword to unite the customers with their corresponding orders using the `ID` fields from the `Customer` and `Order` classes, respectively:

```
var queryResults =
    from c in customers
    join o in orders on c.ID equals o.ID
```

The `on` keyword is followed by the name of the key field (`ID`), and the `equals` keyword indicates the corresponding field in the other collection. The query result only includes the data for objects that have the same `ID` field value as the corresponding `ID` field in the other collection.

The `select` statement projects a new data type with properties named so that you can clearly see the original sales total, the new order, and the resulting new total:

```
select new { c.ID, c.City, SalesBefore = c.Sales, NewOrder = o.Amount,
            SalesAfter = c.Sales+o.Amount };
```

Although you do not increment the sales total in the `customer` object in this program, you could easily do so in the business logic of your program.

The logic of the `foreach` loop and the display of the values from the query are exactly the same as in previous programs in this chapter.

**EXERCISES**

- 
- 20.1** Modify the third example program (BegVCSharp\_20\_3\_QuerySyntax) to order the results in descending order.
- 
- 20.2** Modify the number passed to the `GenerateLotsOfNumbers()` method in the large number program example (BegVCSharp\_20\_6\_LargeNumberQuery) to create result sets of different sizes and see how query results are affected.
- 
- 20.3** Add an `orderby` clause to the query in the large number program example (BegVCSharp\_20\_6\_LargeNumberQuery) to see how this affects performance.
- 
- 20.4** Modify the query conditions in the large number program example (BegVCSharp\_20\_6\_LargeNumberQuery) to select larger and smaller subsets of the number list. How does this affect performance?
- 
- 20.5** Modify the method syntax example (BegVCSharp\_20\_4\_MethodSyntax) to eliminate the `where` clause entirely. How much output does it generate?
- 
- 20.6** Add aggregate operators to the third example program (BegVCSharp\_20\_3\_QuerySyntax). Which simple aggregate operators are available for this non-numeric result set?
- 

Answers to Exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>What LINQ is and when to use it</b>	LINQ is a query language built into C#. Use LINQ to query data from large collections of objects, XML, or databases.
<b>Parts of a LINQ query</b>	A LINQ query includes the <code>from</code> , <code>where</code> , <code>select</code> , and <code>orderby</code> clauses.
<b>How to get the results of a LINQ query</b>	Use the <code>foreach</code> statement to iterate through the results of a LINQ query.
<b>Deferred execution</b>	LINQ query execution is deferred until the <code>foreach</code> statement is executed.
<b>Method syntax and query syntax</b>	Use the query syntax for most LINQ queries and method queries when required. For any given query, the query syntax or the method syntax will give the same result.
<b>Lambda Expressions</b>	Lambda expressions let you declare a method on-the-fly for use in a LINQ query using the method syntax.
<b>Aggregate operators</b>	Use LINQ aggregate operators to obtain information about a large data set without having to iterate through every result.
<b>Group queries</b>	Use group queries to divide data into groups, then sort, calculate aggregates, and compare by group.
<b>Ordering</b>	Use the <code>orderby</code> operator to order the results of a query.
<b>Joins</b>	Use the <code>join</code> operator to query related data in multiple collections with a single query.



# 21

## Databases

### **WHAT YOU WILL LEARN IN THIS CHAPTER:**

---

- Using databases
- Understanding the Entity Framework
- Creating data with Code First
- Using LINQ with databases
- Navigating database relationships
- Creating and querying XML from databases

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

The wrox.com code downloads for this chapter are found at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 21 download and individually named according to the names throughout the chapter.

The previous chapter introduced LINQ (Language-Integrated Query) and showed how LINQ works with objects and XML. This chapter teaches you how to store your objects in a database and use LINQ to query the data.

### **USING DATABASES**

A *database* is a persistent, structured storehouse for data. There are many different kinds of databases, but the most common type you will encounter for storing and querying business data is *relational databases* such as Microsoft SQL Server and Oracle. Relational databases use the SQL database language (SQL stands for *Structured Query Language*) to query and manipulate their data. Traditionally, working with such a database required knowing at least

some SQL, either embedding SQL statements in your programming language or passing strings containing SQL statements to API calls or methods in a SQL-oriented database class library.

Sounds complicated, doesn't it? Well, the good news is that with Visual C# 2015 you can use a *Code First* approach to create objects in C#, store them in a database, and use LINQ to query the objects without having to use another language such as SQL.

## INSTALLING SQL SERVER EXPRESS

To run the examples shown in this chapter, you must install Microsoft SQL Server Express, the free lightweight version of Microsoft SQL Server. You will use the LocalDB option with SQL Server Express, which enables Visual Studio 2015 to create and open a database file directly without the need to connect to a separate server.

SQL Server Express with LocalDB supports the same SQL syntax as the full Microsoft SQL Server, so it is an appropriate version for beginners to learn on. Download SQL Server express from this link:

<http://www.microsoft.com/en-us/server-cloud/products/sql-server-editions/sql-server-express.aspx>

**NOTE** If you are familiar with SQL Server and have access to an instance of Microsoft SQL Server, you may skip this installation, although you will have to change the connection information to match your SQL Server instance. If you have never worked with SQL Server, then go ahead and install SQL Server Express.

## ENTITY FRAMEWORK

The class library in .NET that supports Code First is the newest version of the *Entity Framework*. The name comes from a database concept called the *entity-relationship* model, where an *entity* is the abstract concept of a data object such as a customer, which is related to other entities such as orders and products (for example, a customer places an order for products) in a relational database.

The Entity Framework maps the C# objects in your program to the entities in a relational database. This is called *object-relational mapping*. Object-relational mapping is code that maps your classes, objects, and properties in C# to the tables, rows, and columns that make up a relational database. Creating this mapping code by hand is tedious and time-consuming, but the Entity Framework makes it easy!

The Entity Framework is built on top of ADO.NET, the low-level data access library built into .NET. ADO.NET requires some knowledge of SQL, but luckily the Entity Framework also handles this for you and lets you concentrate on your C# code.

**NOTE** Technically the full name of the Entity Framework is the ADO.NET Entity Framework. You will see it referred to by its full name in some places in Visual Studio. In many blogs and articles, on the other hand, you will see the Entity Framework abbreviated to just EF.

Also with the Entity Framework you get LINQ to Entities, the LINQ provider for the Entity Framework that makes querying the database in C# easy. Now you'll get started by creating some objects in a database.

## A CODE FIRST DATABASE

In the following Try It Out, you create some objects in a database using Code First with the Entity Framework, then query the objects you created using LINQ to Entities.

### TRY IT OUT Code First Database: BegVCSharp\_21\_1\_CodeFirstDatabase

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application project called BegVCSharp\_21\_1\_CodeFirstDatabase in the directory C:\BegVCSharp\Chapter21.
2. Press OK to create the project.
3. To add the Entity Framework, use NuGet as you did in Chapter 19. Go to Tools ⇔ NuGet Package Manager ⇔ Manage NuGet Packages for Solution as shown in Figure 21-1.




FIGURE 21-1

4. Uncheck the Include Prerelease checkbox and get the Entity Framework latest stable release as shown in Figure 21-2. Click the Install button.




FIGURE 21-2

5. Click OK on the Preview dialog as shown in Figure 21-3.




FIGURE 21-3

6. Now the License Acceptance dialog for the Entity Framework appears as shown in Figure 21-3. Click the I Accept button.




FIGURE 21-4

7. Now the Entity Framework and its references are added to your project. You can see them in the References section of your project in Solution Explorer as shown in Figure 21-5.




FIGURE 21-5

8. Open the main source file `Program.cs` and add the following code. First add the Entity Framework namespace at the top of the file below the other `using` clauses:

```
using System.Data.Entity;
```

9. Next, add another `using` clause for data annotations. This enables you to give hints to the Entity Framework on how to set up the database. Finally, add the `System.Console` namespace as with previous examples:

```
using System.ComponentModel.DataAnnotations;
using static System.Console;
```

10. Next, you add a `Book` class with `Author`, `Title`, and `Code` similar to the example you used in Chapter 19. The `[Key]` attribute you see before the `Code` field is a data annotation, telling C# to use this field as the unique identifier for each object in the database.

```
namespace BegVCSharp_21_1_CodeFirstDatabase
{
    public class Book
    {
        public string Title { get; set; }
        public string Author { get; set; }
        [Key] public int Code { get; set; }
    }
}
```

11. Now add a `DbContext` class (Database Context) to manage create, update, and delete the table of books in the database:

```
public class BookContext : DbContext
{
    public DbSet<Book> Books { get; set; }
}
```

12. Next, add code in the `Main()` function to create a couple of `Book` objects, and save the book objects to the database:

```
class Program
{
    static void Main(string[] args)
    {
        using (var db = new BookContext())
        {
            Book book1 = new Book { Title = "Beginning Visual C# 2015",
                                  Author = "Perkins, Reid, and Hammer" };
            db.Books.Add(book1);

            Book book2 = new Book { Title = "Beginning XML",
                                  Author = "Fawcett, Quin, and Ayers" };
            db.Books.Add(book2);

            db.SaveChanges();
        }
    }
}
```

13. Finally, add the code for a simple LINQ query to list the books in the database after creation:

```
var query = from b in db.Books
            orderby b.Title
            select b;

WriteLine("All books in the database:");
foreach (var b in query)
{
    WriteLine($"{b.Title} by {b.Author}, code={b.Code}");
}

WriteLine("Press a key to exit...");
ReadKey();
}
```

The complete code for your program should now look like this:

```
using System.Data.Entity;
using System.Data.Annotations;
using static System.Console;

namespace BegVCSharp_21_1_CodeFirstDatabase
{
    public class Book
    {
        public string Title { get; set; }
        public string Author { get; set; }
        public int Code { get; set; }
    }

    public class BookContext : DbContext
    {
        public DbSet<Book> Books { get; set; }
    }
}

class Program
```

```

    {
        static void Main(string[] args)
        {
            using (var db = new BookContext())
            {
                Book book1 = new Book { Title = "Beginning Visual C# 2015",
                                      Author = "Perkins, Reid, and Hammer" };
                db.Books.Add(book1);

                Book book2 = new Book { Title = "Beginning XML",
                                      Author = "Fawcett, Quin, and Ayers" };
                db.Books.Add(book2);

                db.SaveChanges();

                var query = from b in db.Books
                           orderby b.Title
                           select b;

                WriteLine("All books in the database:");
                foreach (var b in query)
                {
                    WriteLine($"{b.Title} by {b.Author}, code={b.Code}");
                }

                WriteLine("Press a key to exit...");
                ReadKey();
            }
        }
    }
}

```

14. Compile and execute the program (you can just press F5 for Start Debugging). You will see the information for the books database appear as shown in Figure 21-6.




FIGURE 21-6

Press any key to finish the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you might need to press Enter/Return twice. That finishes the program run. Now look at how it works in detail.

### *How It Works*

As is shown in the previous chapter, this code uses extension classes from the `System.Linq` namespace, which is referenced by a `using` statement inserted automatically by Visual C# 2015 when you create the project:

```
using System.Linq;
```

Next you added the Entity Framework namespace at the top of the file below the other using clauses:

```
using System.Data.Entity;
```

Then you added the using clause for data annotations, so that you could add hints to tell the Entity Framework on how to set up the database, and the static System.Console namespace:

```
using System.ComponentModel.DataAnnotations;
using static System.Console;
```

Next, you added a Book class with Author, Title, and Code similar to the example used in Chapter 19. You used the [Key] attribute to identify the Code property as the unique identifier for each row in the database.

```
namespace BegVCSharp_21_1_CodeFirstDatabase
{
    public class Book
    {
        public string Title { get; set; }
        public string Author { get; set; }
        [Key] public int Code { get; set; }
    }
}
```

Next you created the BookContext class inheriting from the DbContext (Database Context) class in the Entity Framework for creating, updating, and deleting the book objects as needed in the database:

```
public class BookContext : DbContext
{
    public DbSet<Book> Books { get; set; }
}
```

The class member DbSet<Book> is a collection of all the Book entities in your database.

Next you added code to use the BookContext to create two Book objects and save them to the database:

```
using (var db = new BookContext())
{
    Book book1 = new Book { Title = "Beginning Visual C# 2015",
                           Author = "Perkins, Reid, and Hammer" };
    db.Books.Add(book1);

    Book book2 = new Book { Title = "Beginning XML",
                           Author = "Fawcett, Quin, and Ayers" };
    db.Books.Add(book2);

    db.SaveChanges();
}
```

The using(var db = new BookContext()) clause lets you create a new BookContext instance for use in all the following code between the curly braces. Besides being a convenient shorthand, the using() clause ensures that the database connection and other underlying plumbing objects associated with the connection are closed properly when your program is finished, even if there is an exception or other unexpected event.

The Book creation and assignment statements such as

```
Book book = new Book { Title = "Beginning Visual C# 2015",
                      Author = "Perkins, Reid, and Hammer" };
```

are fairly straightforward creation of Book objects; no database magic has occurred yet as these are simple objects in memory. You'll note that you did not assign any value for the Code property; at this point the unassigned Code property simply contains a default value.

Next you saved the changes to `BookContext db` to the database:

```
db.SaveChanges();
```

Now some magic has happened; because you used the `[Key]` attribute to identify Code as a key, a unique value was assigned to the Code field when each object was saved to the database. You don't have to use this value or even care what it is, because it is taken care of for you by the Entity Framework.

**NOTE** If you had not added the `[Key]` attribute to your object, you would have seen an exception like the one shown Figure 21-7 when running your program.




FIGURE 21-7

Finally, you execute the code for a simple LINQ query to list the books in the database after creation:

```
var query = from b in db.Books
            orderby b.Title
            select b;

            WriteLine("All books in the database:");
            foreach (var b in query)
            {
                WriteLine($"{b.Title} by {b.Author}, code={b.Code}");
            }

            WriteLine("Press a key to exit...");
            ReadKey();
        }
```

This LINQ query is very similar to the one you used in the previous chapter, but instead of querying objects in memory using the LINQ to Objects provider, you are querying the database with the LINQ to Entities provider. LINQ infers the correct provider based on the types referenced in the query; you don't have to make any changes in your logic.

Finally you just use the standard `ReadKey()` to pause the program before exiting so you can see the output.

That was easy, right? You created some objects, saved them to a database, and queried the database using LINQ.

## BUT WHERE IS MY DATABASE?

But wait, you say. Where is the database you created? You never specified a file name or a folder location—it was all magic! You can see the database in Visual Studio 2015 through the Server Explorer. Go to Tools  $\Rightarrow$  Connect to Database. The Entity Framework will create a database in the first local SQL Server instance it finds on your computer.

If you never had any databases on your computer previously, Visual C# 2015 creates a local SQL Server instance for you called `(localdb)\MSSQLLocalDB`. To connect to this database type `(localdb)\MSSQLLocalDB` into the Server Name field as shown in Figure 21-8.




FIGURE 21-8

**NOTE** If you had installed a previous version of Visual Studio before using Visual C# 2015, you might have to enter `(localdb)\v11.0` into the Server Name field, as this was the previous edition's local database name. Or if you have installed the SQL Server Express Edition, you might have to enter `.\sqlexpress`, as Entity Framework uses the first local SQL Server database it finds.

The database containing your data will be called `BegVCSharp_21_1_CodeFirstDatabase`.`BookContext` assuming you typed in the example name exactly as shown in the chapter. It will show up in the **Select or enter a database name** field after taking a moment to connect.

Now you can press OK and the database will appear in the Server Explorer Data Connections window in Visual C# 2015 as shown in Figure 21-9.




FIGURE 21-9

From here you can explore the database directly. For example you can right-click on the Books table and choose **Show Table Data** to see the data you entered as shown in Figure 21-10.




FIGURE 21-10

## NAVIGATING DATABASE RELATIONSHIPS

One of the most powerful aspects of the Entity Framework is its capability to automatically create LINQ objects to help you navigate relationships between related tables in the database.

In the following Try It Out, you add two new classes related to the Book class to make a simple bookstore inventory report. The new classes are called Store (to represent each bookstore) and Stock, to represent the inventory of books on hand (in the store on the shelf) and on order from the publisher. A diagram of these new classes and relationships is shown in Figure 21-11.

Each store has a name, address, and an Inventory collection consisting of one or more stock objects, one for each different book (title) carried by the store. The relationship between Store and Stock is one-to-many. Each stock record is related to exactly one book. The relationship between Stock and Book is one-to-one. You need the stock record because one store may have three copies of a particular book, but another store will have six copies of the same book.




FIGURE 21-11

You'll see how with Code First, all you have to do is create the C# objects and collections, and the Entity Framework will create the database structure for you and let you easily navigate the relationships between your database objects and then query the related objects in the database.

## TRY IT OUT Navigating Database Relationships: BegVCSharp\_21\_2\_DatabaseRelations

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application project called BegVCSharp\_21\_2\_DatabaseRelations in the directory C:\BegVCSharp\Chapter21.
2. Press OK to create the project.
3. Add the Entity Framework using NuGet as you did in the previous example. Go to Tools  $\leftrightarrow$  NuGet Package Manager  $\leftrightarrow$  Manage NuGet Packages for Solution.
4. In the NuGet Package Manager, choose the Entity Framework, uncheck the Include Prerelease checkbox and get the Entity Framework latest stable release. Click the Install button. It does not have to download because you already downloaded it in the previous step. Click OK on the Preview Changes and the I Accept button for the License Acceptance dialog.
5. Open the main source file Program.cs. As in the previous example, add the using statements for the System.Console, System.Data.Entity, and DataAnnotations namespaces, as well as the code to create the Book class:

```

using System.Data.Entity;
using System.ComponentModel.DataAnnotations;
using static System.Console;

namespace BegVCSharp_21_2_DatabaseRelations
{
    public class Book
    {
        public string Title { get; set; }
        public string Author { get; set; }
    }
}
    
```

```

    [Key]
    public int Code { get; set; }
}

```

6. Now declare the Store and Stock classes as shown below. Make sure to declare Inventory and Item as virtual. You'll see why in the How It Works section.

```

public class Store
{
    [Key]
    public int StoreId { get; set; }
    public string Name { get; set; }
    public string Address { get; set; }
    public virtual List<Stock> Inventory { get; set; }
}
public class Stock
{
    [Key]
    public int StockId { get; set; }
    public int OnHand { get; set; }
    public int OnOrder { get; set; }
    public virtual Book Item { get; set; }
}

```

7. Next add Stores and Stocks to the DbContext class:

```

public class BookContext : DbContext
{
    public DbSet<Book> Books { get; set; }
    public DbSet<Store> Stores { get; set; }
    public DbSet<Stock> Stocks { get; set; }
}

```

8. Now add code to the Main() method to use the BookContext and create the two instances of the Book class as in the previous example:

```

class Program
{
    static void Main(string[] args)
    {
        using (var db = new BookContext())
        {
            Book book1 = new Book
            {
                Title = "Beginning Visual C# 2015",
                Author = "Perkins, Reid, and Hammer"
            };
            db.Books.Add(book1);

            Book book2 = new Book
            {
                Title = "Beginning XML",
                Author = "Fawcett, Quin, and Ayers"
            };
            db.Books.Add(book2);
        }
    }
}

```

9. Now add an instance for the first store and its inventory, still inside the `using(var db = new BookContext())` clause:

```
var store1 = new Store
{
    Name = "Main St Books",
    Address = "123 Main St",
    Inventory = new List<Stock>()
};
db.Stores.Add(store1);

Stock store1book1 = new Stock
{ Item = book1, OnHand = 4, OnOrder = 6 };
store1.Inventory.Add(store1book1);

Stock store1book2 = new Stock
{ Item = book2, OnHand = 1, OnOrder = 9 };
store1.Inventory.Add(store1book2);
```

10. Now add an instance for the second store and its inventory:

```
var store2 = new Store
{
    Name = "Campus Books",
    Address = "321 College Ave",
    Inventory = new List<Stock>()
};

db.Stores.Add(store2);

Stock store2book1 = new Stock
{ Item = book1, OnHand = 7, OnOrder = 23 };
store2.Inventory.Add(store2book1);

Stock store2book2 = new Stock
{ Item = book2, OnHand = 2, OnOrder = 8 };
store2.Inventory.Add(store2book2);
```

11. Next save the database changes as in the previous example:

```
db.SaveChanges();
```

12. Now create a LINQ query on all the stores, and print out the results:

```
var query = from store in db.Stores
            orderby store.Name
            select store;
```


13. Finally add code to print out the results of the query and pause the output:

```
WriteLine("Bookstore Inventory Report:");
foreach (var store in query)
{
    WriteLine($"{store.Name} located at {store.Address}");
    foreach (Stock stock in store.Inventory)
    {
        WriteLine($"- Title: {stock.Item.Title}");
```

```
        WriteLine($"-- Copies in Store: {stock.OnHand}");
        WriteLine($"-- Copies on Order: {stock.OnOrder}");
    }
}

WriteLine("Press a key to exit...");
ReadKey();
}
}
}
```

14. Compile and execute the program (you can just press F5 for Start Debugging). You will see the information for the bookstore inventory appear as shown in Figure 21-12.



**FIGURE 21-12**

Press any key to finish the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you might need to press Enter/Return twice. That finishes the program run. Now look at how it works in detail.

## *How It Works*

The basics of the Entity Framework, `DbContext`, and data annotations were covered in the previous example, so here you'll concentrate on what is different.

The Store and Stock classes are similar to the original Book class but you added some new virtual properties for `Inventory` and `Item` as shown here:

```
public class Store
{
    [Key]
    public int StoreId { get; set; }
    public string Name { get; set; }
    public string Address { get; set; }
    public virtual List<Stock> Inventory { get; set; }
}
public class Stock
{
    [Key]
    public int StockId { get; set; }
```

```

        public int OnHand { get; set; }
        public int OnOrder { get; set; }
        public virtual Book Item{ get; set; }
    }
}

```

The `Inventory` property looks and behaves like a normal in-memory `List<Stock>` collection. However because it is declared as `virtual`, the Entity Framework can override its behavior when storing to and retrieving from the database.

The Entity Framework takes care of the database details such as adding a foreign key column to the `Stocks` table in the database to implement the `Inventory` relationship between a `Store` and its `Stock` records. Similarly the Entity Framework adds another foreign key column to the `Stock` table in the database to implement the `Item` relationship between `Stock` and `Book`. If you're curious you can see this in Server Explorer database design view of the `BegVCSharp_21_2_DatabaseRelations.BookContext` database as shown in Figure 21-13.

Name	Data Type	Allow Nulls	Default
StockId	int	<input type="checkbox"/>	
OnHand	int	<input type="checkbox"/>	
OnOrder	int	<input type="checkbox"/>	
Item_Code	int	<input checked="" type="checkbox"/>	
Store_StoreId	int	<input checked="" type="checkbox"/>	
		<input type="checkbox"/>	

Keys (1)

- PK\_dbo.Stocks (Primary Key, Clustered: StockId)
- Check Constraints (0)

Indexes (2)

- IX\_Item\_Code (Item\_Code)
- IX\_Store\_StoreId (Store\_StoreId)

Foreign Keys (2)

- FK\_dbo.Stocks\_dbo.Books\_Item\_Code (Code)
- FK\_dbo.Stocks\_dbo.Stores\_Store\_StoreId (StoreId)

Triggers (0)

FIGURE 21-13

In the past you would have had to decide how to map the collection in your program to foreign keys and columns in the database and keep that code up-to-date as your design changes. However, with the Entity Framework you do not need to know these details; with Code First you simply work with C# classes and collections and let the framework take care of the plumbing for you.

Next you added the `DbSet` classes for `Store` and `Stock` to the `BookContext`.

```

public class BookContext : DbContext
{
    public DbSet<Book> Books { get; set; }
    public DbSet<Store> Stores { get; set; }
    public DbSet<Stock> Stocks { get; set; }
}

```

Then you used those `DbSet` classes to create instances of two books, two stores, and two stock records for each book under each store:

```

class Program
{
    static void Main(string[] args)
    {
        using (var db = new BookContext())
        {
            Book book1 = new Book
            {
                Title = "Beginning Visual C# 2015",
                Author = "Perkins, Reid, and Hammer"
            };
            db.Books.Add(book1);

            Book book2 = new Book
            {
                Title = "Beginning XML",
                Author = "Fawcett, Quin, and Ayers"
            };
            db.Books.Add(book2);

            var store1 = new Store
            {
                Name = "Main St Books",
                Address = "123 Main St",
                Inventory = new List<Stock>()
            };
            db.Stores.Add(store1);

            Stock store1book1 = new Stock
            { Item = book1, OnHand = 4, OnOrder = 6 };
            store1.Inventory.Add(store1book1);

            Stock store1book2 = new Stock
            { Item = book2, OnHand = 1, OnOrder = 9 };
            store1.Inventory.Add(store1book2);
            var store2 = new Store
            {
                Name = "Campus Books",
                Address = "321 College Ave",
                Inventory = new List<Stock>()
            };

            db.Stores.Add(store2);

            Stock store2book1 = new Stock
            { Item = book1, OnHand = 7, OnOrder = 23 };
            store2.Inventory.Add(store2book1);

            Stock store2book2 = new Stock
            { Item = book2, OnHand = 2, OnOrder = 8 };
            store2.Inventory.Add(store2book2);
        }
    }
}

```

After creating the objects, you saved the changes to the database:

```
db.SaveChanges();
```

Then you made a simple LINQ query to list all the stores' information:

```
var query = from store in db.Stores
            orderby store.Name
            select store;
```

The code to print out the results of the query is very straightforward because it simply deals with objects and collections, no database-specific code:

```
WriteLine("Bookstore Inventory Report:");
foreach (var store in query)
{
    WriteLine($"{store.Name} located at {store.Address}");
    foreach (Stock stock in store.Inventory)
    {
        WriteLine($"- Title: {stock.Item.Title}");
        WriteLine($"-- Copies in Store: {stock.OnHand}");
        WriteLine($"-- Copies on Order: {stock.OnOrder}");
    }
}
```

To print the inventory under each store, you simply use a foreach loop like with any collection.

## HANDLING MIGRATIONS

Inevitably as you develop your code, you are going to change your mind. You will come up with a better name for one of your properties, or you will realize you need a new class or relationship. If you change the code in a class connected to a database, via the Entity Framework, you will encounter the Invalid Operation Exception shown in Figure 21-14 when you first run the changed program.




FIGURE 21-14

Keeping the database up to date with your changed classes is complicated, but again the Entity Framework steps in with a facility to make it relatively easy. As the error message suggests, you need to add the Code First Migrations package to your program.

To do this, go to Tools  $\Rightarrow$  NuGet Package Manager  $\Rightarrow$  Package Manager Console. This brings up a command window as shown in Figure 21-15.

To enable automatic migration of your database to your updated class structure, enter this command in the Package Manager Console at the PM> prompt:

```
Enable-Migrations -EnableAutomaticMigrations
```




FIGURE 21-15

This adds a Migrations class to your project, shown in Figure 21-16.




FIGURE 21-16

The Entity Framework will compare the timestamp of the database to your program and advise you when the database is out of sync with your classes. To update the database, simply enter this command in the Package Manager Console at the PM> prompt:

Update-Database

## CREATING AND QUERYING XML FROM AN EXISTING DATABASE

For the last example you will combine all you have learned about LINQ, databases, and XML.


XML is often used to communicate data between client and server machines or between “tiers” in a multitier application. It is quite common to query for some data in a database and then produce an XML document or fragment from that data to pass to another tier.

In the following Try It Out, you create a query to find some data in the previous example database, use LINQ to Entities to query the data, and then use LINQ to XML classes to convert the data to XML. This is an example of Database First as opposed to Code First programming where you take an existing database and generate C# objects from it.

**TRY IT OUT****Generating XML from Databases: BegVCSharp\_21\_3\_XMLfromDatabase**

Follow these steps to create the example in Visual Studio 2015:

1. Create a new console application called BegVCSharp\_21\_3\_XMLfromDatabase in the directory C:\BegVCSharp\Chapter21.
2. As described in the previous example, add the Entity Framework to the project.
3. Add a connection to the database used by the previous example by selecting Project  $\Rightarrow$  Add New Item. Choose ADO.NET Entity Data Model in the Add New Item dialog and change the name from Model1 to BookContext as shown in Figure 21-17.



**FIGURE 21-17**

4. In the Entity Data Model Wizard, choose the connection to BegVCSharp\_21\_2\_DatabaseRelations .BookContext database you created in the previous example as shown in Figure 21-18.
5. Open the main source file Program.cs.
6. Add a reference to the System.Xml.Linq namespace to the beginning of Program.cs, as shown:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Xml.Linq;
```

```
using System.Text;
using static System.Console;
```





FIGURE 21-18

- Add the following code to the Main() method in Program.cs:

```
static void Main(string[] args)
{
    using (var db = new BookContext())
    {
        var query = from store in db.Stores
                    orderby store.Name
                    select store;
        foreach (var s in query)
        {
            XElement storeElement = new XElement("store",
                new XAttribute("name", s.Name),
                new XAttribute("address", s.Address),
                from stock in s.Stocks
                select new XElement("stock",
                    new XAttribute("StockID", stock.StockId),
                    new XAttribute("onHand",
                        stock.OnHand),
                    new XAttribute("onOrder",
                        stock.OnOrder),
                    new XElement("book",
                        new XAttribute("title",
                            stock.Book.Title),
                        new XAttribute("author",
                            stock.Book.Author))
                ) // end book
            ) // end stock
        }
    }
}
```

```
        ); // end store
    WriteLine(storeElement);
}
Write("Program finished, press Enter/Return to continue:");
ReadLine();
}
```

8. Compile and execute the program (you can just press F5 for Start Debugging). You will see the output shown in Figure 21-19.



**FIGURE 21-19**

Simply press Enter/Return to exit the program and make the console screen disappear. If you used Ctrl+F5 (Start Without Debugging), you might need to press Enter/Return twice.

## *How It Works*

In Program.cs you added the reference to the System.Xml.Linq namespace in order to call the LINQ to XML constructor classes in addition to the Entity Framework classes.

When you added the Database First code by choosing ADO.NET Entity Data Model in the Add New Item dialog, Visual Studio generated a separate BookContext.cs class and added it to your project using the information from the existing BegVCSharp\_21\_2\_DatabaseRelations.BookContext database created in the previous example.

In the main program, you created an instance of the BooksContext database context class and the same LINQ to Entities query used in previous examples:

```
using (var db = new BookContext())
{
    var query = from store in db.Stores
                orderby store.Name
                select store;
```

When you processed the results of the query in a `foreach` loop, you used the LINQ to XML classes to transform the query results into XML using a nested set of LINQ to XML elements and attributes:

```
foreach (var s in query)
{
    XElement storeElement = new XElement("store",
```

```
        new XAttribute("name", s.Name),
        new XAttribute("address", s.Address),
        from stock in s.Stocks
        select new XElement("stock",
            new XAttribute("StockID", stock.StockId),
            new XAttribute("onHand",
                stock.OnHand),
            new XAttribute("onOrder",
                stock.OnOrder),
            new XElement("book",
                new XAttribute("title",
                    stock.Book.Title),
                new XAttribute("author",
                    stock.Book.Author)
            ) // end book
        ) // end stock
    ); // end store
    WriteLine(storeElement);
}
```

Congratulations! You have combined your data access knowledge from Chapters 19, 20, and 21 into a single program using the full power of LINQ and the Entity Framework!

## EXERCISES

- 21.1** Modify the first example BegVCSharp\_21\_1\_CodeFirstDatabase to prompt the user for title and author and store the user-entered data into the database.
- 21.2** The first example BegVCSharp\_21\_1\_CodeFirstDatabase will create duplicate records if run repeatedly. Modify the example to not create duplicates.
- 21.3** The generated BookContext class used in the last example BegVCSharp\_21\_3\_XMLfromDatabase does not use the same relationship names as the previous example BegVCSharp\_21\_2\_DatabaseRelations. Modify the BookContext class to use the same relationship names.
- 21.4** Create a database using Code First to store the data found in the GhostStories.xml file used in Chapter 19.

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>Using Databases</b>	A database is a persistent, structured storehouse for data. While there are many different kinds of databases, the most common type used for business data are relational databases.
<b>Entity Framework</b>	The Entity Framework is a set of .NET classes for object-relational mapping between C# objects and relational databases.
<b>How to Create Data with Code First</b>	By using the Code First classes in the Entity Framework you can create databases directly from C# classes and collections using object-relational mapping.
<b>How to use LINQ with Databases</b>	LINQ to Entities enables powerful queries on databases using the same Entity Framework classes to create the data.
<b>How to Navigate Database Relationships</b>	The Entity Framework enables creation and navigation of related entities in your database through the use of virtual properties and collections in your C# code.
<b>How to create and query XML from Databases</b>	You can construct XML from databases by combining LINQ to Entities, LINQ to Objects, and LINQ to XML in a single query.

# PART IV

## Additional Techniques

---

►CHAPTER 22: Windows Communication Foundation

►CHAPTER 23: Windows Store Apps



# 22

## Windows Communication Foundation

---

### WHAT YOU WILL LEARN IN THIS CHAPTER:

---

- Discovering WCF
- Mastering WCF concepts
- Understanding WCF programming

### WROX.COM CODE DOWNLOADS FOR THIS CHAPTER

The wrox.com code downloads for this chapter are found at [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) on the Download Code tab. The code is in the Chapter 22 download and individually named according to the names throughout the chapter.

In recent years, as use of the Internet has become more ubiquitous, there has been a rapid increase in *web services*. A web service is like a website that is used by a computer instead of a person. For example, instead of browsing to a website about your favorite TV program, you might instead use a desktop application that pulled in the same information via a web service. The advantage here is that the same web service might be used by all sorts of applications, and, indeed, by websites. Also, you can write your own application or website that uses third-party web services. Perhaps you might combine information about your favorite TV program with a mapping service to show filming locations.

The .NET Framework has supported web services for some time now. However, in the more recent versions of the framework, web services have been combined with another technology, called *remoting*, to create *Windows Communication Foundation* (WCF), which is a generic infrastructure for communication between applications.

Remoting makes it possible to create instances of objects in one process and use them from another process — even if the object is created on a computer other than the one that is using it. However, remoting on its own is limited, and isn't the easiest thing for a beginner programmer to learn.

WCF takes concepts such as services and platform-independent SOAP messaging from web services, and combines these with concepts such as host server applications and advanced binding capabilities from remoting. The result is a technology you can think of as a superset that includes both web services and remoting, but that is much more powerful than web services and much easier to use than remoting. Using WCF, you can move from simple applications to applications that use a *service-oriented architecture* (SOA). SOA means that you decentralize processing and make use of distributed processing by connecting to services and data as you need them across local networks and the Internet.

This chapter walks you through how to create and consume WCF services from your application code. But just as importantly, it also covers the principles behind WCF, so you understand why things work the way they do.

## WHAT IS WCF?

WCF is a technology that enables you to create services that you can access from other applications across process, machine, and network boundaries. You can use these services to share functionality across multiple applications, to expose data sources, or to abstract complicated processes.

The functionality that WCF services offer is encapsulated as individual methods that are exposed by the service. Each method — or, in WCF terminology, each *operation* — has an endpoint that you exchange data with in order to use it. This data exchange can be defined by one or more protocols, depending on the network that you use to connect to the service and your specific requirements.

In WCF, an endpoint can have multiple *bindings*, each of which specifies a means of communication. Bindings can also specify additional information, such as which security requirements must be met to communicate with the endpoint. A binding might require username and password authentication or a Windows user account token, for example. When you connect to an endpoint, the protocol that the binding uses affects the address that you use, as you will see shortly.

Once you have connected to an endpoint, you can communicate with it by using Simple Object Access Protocol (SOAP) messages. The form of the messages that you use depends on the operation you are using and the data structures that are required to send messages to (and receive messages from) that operation. WCF uses *contracts* to specify all of this. You can discover contracts through metadata exchange with a service. One commonly used format for service discovery is the Web Service Description Language (WSDL), which was originally used for web services, although WCF services can also be described in other ways.

**NOTE** WCF is something of a chameleon in how it can be used and set up. It is possible to create Representative State Transfer (REST) services using WCF. These services rely on simple HTTP requests to communicate between the client and the server, and because of this they can have a smaller footprint than the SOAP messages.

When you have identified a service and endpoint that you want to use, and after you know which binding you use and which contracts to adhere to, you can communicate with a WCF service as easily as with an object that you have defined locally. Communications with WCF services can be simple, one-way transactions, request/response messages, or full-duplex communications that can be initiated from either end of the communication channel. You can also use message payload optimization techniques, such as Message Transmission Optimization Mechanism (MTOM), to package data if required.

The WCF service itself might be running in one of a number of different processes on the computer where it is hosted. Unlike web services, which always run in IIS, you can choose a host process that is appropriate to your situation. You can use IIS to host WCF services, but you can also use Windows services or executables. If you are using TCP to communicate with a WCF service over a local network, there is no need even to have IIS installed on the PC that is hosting the service.

The WCF framework has been designed to enable you to customize nearly everything you have read about in this section. However, this is an advanced subject and you will only be using the techniques provided by default in .NET 4.5 in this chapter.

Now that you have covered the basics about WCF services, you will look in more detail at these concepts in the following sections.

## WCF CONCEPTS

This section describes the following aspects of WCF:

- WCF communication protocols
- Addresses, endpoints, and bindings
- Contracts
- Message patterns
- Behaviors
- Hosting

## WCF Communication Protocols

As described earlier, you can communicate with WCF services through a variety of transport protocols. In fact, five are defined in the .NET 4.5 Framework:

- **HTTP** — Enables you to communicate with WCF services from anywhere, including across the Internet. You can use HTTP communications to create WCF web services.
- **TCP** — Enables you to communicate with WCF services on your local network or across the Internet if you configure your firewall appropriately. TCP is more efficient than HTTP and has more capabilities, but it can be more complicated to configure.
- **UDP** — User Datagram Protocol is similar to TCP in that it enables communications via the local network or Internet, but it's implemented in a subtly different way. One of the

consequences of this implementation is that a service can broadcast messages to multiple clients simultaneously.

- **Named pipe** — Enables you to communicate with WCF services that are on the same machine as the calling code, but reside in a separate process.
- **MSMQ** — Microsoft Message Queuing is a queuing technology that enables messages sent by an application to be routed through a queue to arrive at a destination. MSMQ is a reliable messaging technology that ensures that a message sent to a queue will reach that queue. MSMQ is also inherently asynchronous, so a queued message will be processed only when messages ahead of it in the queue have been processed and a processing service is available.

These protocols often enable you to establish secure connections. For example, you can use the HTTPS protocol to establish an SSL connection across the Internet. TCP offers extensive possibilities for security in a local network by using the Windows security framework. UDP doesn't support security.

In order to connect to a WCF service, you must know where it is. In practice, this means knowing the address of an endpoint.

## Addresses, Endpoints, and Bindings

The type of address you use for a service depends on the protocol that you are using. Service addresses are formatted for the three protocols described in this chapter (MSMQ is not covered) as follows:

- **HTTP** — Addresses for the HTTP protocol are URLs of the familiar form `http://<server>:<port>/<service>`. For SSL connections, you can also use `https://<server>:<port>/<service>`. If you are hosting a service in IIS, `<service>` will be a file with a `.svc` extension. IIS addresses will probably include more subdirectories than this example — that is, more sections separated by `/` characters before the `.svc` file.
- **TCP** — Addresses for TCP are of the form `net.tcp://<server>:<port>/<service>`.
- **UDP** — Addresses for UDP are of the form `soap.udp://<server>:<port>/<service>`. Certain `<server>` values are required for multicast communications, but this is beyond the scope of this chapter.
- **Named pipe** — Addresses for named pipe connections are similar but have no port number. They are of the form `net.pipe://<server>/<service>`.

The address for a service is a base address that you can use to create addresses for endpoints representing operations. For example, you might have an operation at `net.tcp://<server>:<port>/<service>/operation1`.

For example, imagine you create a WCF service with a single operation that has bindings for all three of the protocols listed here. You might use the following base addresses:

```
http://www.mydomain.com/services/amazingservices/mygreatservice.svc  
net.tcp://myhugeserver:8080/mygreatservice  
net.pipe://localhost/mygreatservice
```

You could then use the following addresses for operations:

```
http://www.mydomain.com/services/amazingservices/mygreatservice.svc/greatop
net.tcp://myhugeserver:8080/mygreatservice/greatop
net.pipe://localhost/mygreatservice/greatop
```

Since .NET 4, it has been possible to use default endpoints for operations, without having to explicitly configure them. This simplifies configuration, especially in situations where you want to use standard endpoint addresses, as in the preceding examples.

Bindings, as mentioned earlier, specify more than just the transport protocol that will be used by an operation. You can also use them to specify the security requirements for communication over the transport protocol, transactional capabilities of the endpoint, message encoding, and much more.

Because bindings offer such a great degree of flexibility, the .NET Framework provides some pre-defined bindings that you can use. You can also use these bindings as starting points, tweaking them to obtain exactly the type of binding you want — up to a point. The predefined bindings have certain capabilities to which you must adhere. Each binding type is represented by a class in the `System.ServiceModel` namespace. Table 22-1 lists the most commonly used bindings along with some basic information about them.

**TABLE 22-1:** Binding Types

BINDING	DESCRIPTION
BasicHttpBinding	The simplest HTTP binding, and the default binding used by web services. It has limited security capabilities and no transactional support.
WSHttpBinding	A more advanced form of HTTP binding that is capable of using all the additional functionality that was introduced in WSE.
WSDualHttpBinding	Extends <code>WSHttpBinding</code> capabilities to include duplex communication capabilities. With duplex communication, the server can initiate communications with the client in addition to ordinary message exchange.
WSFederationHttpBinding	Extends <code>WSHttpBinding</code> capabilities to include federation capabilities. Federation enables third parties to implement single sign-on and other proprietary security measures. This is an advanced topic not covered in this chapter.
NetTcpBinding	Used for TCP communications, and enables you to configure security, transactions, and so on.
NetNamedPipeBinding	Used for named pipe communications, and enables you to configure security, transactions, and so on.
NetMsmqBinding	Used with MSMQ, which is not covered in this chapter.
NetPeerTcpBinding	Used for peer-to-peer binding, which is not covered in this chapter.

*continues*

**TABLE 22-1** (continued)

BINDING	DESCRIPTION
WebHttpBinding	Used for web services that use HTTP requests instead of SOAP messages.
UdpBinding	Allows binding to the UDP protocol.

Many of the binding classes have similar properties that you can use for additional configuration. For example, they have properties that you can use to configure timeout values. You'll learn more about this when you look at code later in this chapter.

Since .NET 4, endpoints have default bindings that vary according to the protocol used. These defaults are shown in Table 22-2.

**TABLE 22-2:** NET Default Bindings

PROTOCOL	DEFAULT BINDING
HTTP	BasicHttpBinding
TCP	NetTcpBinding
UDP	UdpBinding
Named pipe	NetNamedPipeBinding
MSMQ	NetMsmqBinding

## Contracts

Contracts define how WCF services can be used. Several types of contract can be defined:

- **Service contract** — Contains general information about a service and the operations exposed by a service. This includes, for example, the namespace used by service. Services have unique namespaces that are used when defining the schema for SOAP messages in order to avoid possible conflicts with other services.
- **Operation contract** — Defines how an operation is used. This includes the parameter and return types for an operation method along with additional information, such as whether a method will return a response message.
- **Message contract** — Enables you to customize how information is formatted inside SOAP messages — for example, whether data should be included in the SOAP header or SOAP message body. This can be useful when creating a WCF service that must integrate with legacy systems.
- **Fault contract** — Defines faults that an operation can return. When you use .NET clients, faults result in exceptions that you can catch and deal with in the normal way.

- **Data contract** — If you use complex types, such as user-defined structs and objects, as parameters or return types for operations, then you must define data contracts for these types. Data contracts define the types in terms of the data that they expose through properties.

You typically add contracts to service classes and methods by using attributes, as you will see later in this chapter.

## Message Patterns

In the previous section, you saw that an operation contract can define whether an operation returns a value. You've also read about duplex communications that are made possible by the `WSDualHttpBinding` binding. These are both forms of message patterns, of which there are three types:

- **Request/response messaging** — The “ordinary” way of exchanging messages, whereby every message sent to a service results in a response being sent back to the client. This doesn't necessarily mean that the client waits for a response, as you can call operations asynchronously in the usual way.
- **One-way, or simplex, messaging** — Messages are sent from the client to the WCF operation, but no response is sent.
- **Two-way, or duplex, messaging** — A more advanced scheme whereby the client effectively acts as a server as well as a client, and the server as a client as well as a server. Once set up, duplex messaging enables both the client and the server to send messages to each other, which might not have responses.

You'll see how these message patterns are used in practice later in this chapter.

## Behaviors

Behaviors are a way to apply additional configuration that is not directly exposed to a client to services and operations. By adding a behavior to a service, you can control how it is instantiated and used by its hosting process, how it participates in transactions, how multithreading issues are dealt with in the service, and so on. Operation behaviors can control whether impersonation is used in the operation execution, how the individual operation affects transactions, and more.

Since .NET 4, you can specify default behaviors at various levels, so that you don't have to specify every aspect of every behavior for every service and operation. Instead, you can provide defaults and override settings where necessary, which reduces the amount of configuration required.

## Hosting

In the introduction to this chapter, you learned that WCF services can be hosted in several different processes. These possibilities are as follows:

- **Web server** — IIS-hosted WCF services are the closest thing to web services that WCF offers. However, you can use advanced functionality and security features in WCF services that are

much more difficult to implement in web services. You can also integrate with IIS features such as IIS security.

- **Executable** — You can host a WCF service in any application type that you can create in .NET, such as console applications, Windows Forms applications, and WPF applications.
- **Windows service** — You can host a WCF service in a Windows service, which means that you can use the useful features that Windows services provide. This includes automatic startup and fault recovery.
- **Windows Activation Service (WAS)** — Designed specifically to host WCF services, WAS is basically a simple version of IIS that you can use where IIS is not available.

Two of the options in the preceding list — IIS and WAS — provide useful features for WCF services such as activation, process recycling, and object pooling. If you use either of the other two hosting options, the WCF service is said to be *self-hosted*. You will occasionally self-host services for testing purposes, but there can be very good reasons for creating self-hosted production-grade services. For example, you could be in a situation where you’re not allowed to install a web server on the computer on which your service should run. This might be the case if the service runs on a domain controller or if the local policy of your organization simply prohibits running IIS. In this case you can host the service in a Windows service and it will work every bit as well as it would otherwise.

## WCF PROGRAMMING

Now that you have covered all the basics, it is time to get started with some code. In this section you’ll start by looking at a simple web server–hosted WCF service and a console application client. After looking at the structure of the code created, you’ll learn about the basic structure of WCF services and client applications. Then you will look at some key topics in a bit more detail:

- Defining WCF service contracts
- Self-hosted WCF services

### TRY IT OUT A Simple WCF Service and Client: Ch22Ex01Client


1. Create a new WCF Service Application project called Ch22Ex01 in the directory C:\BegVCSharp\Chapter22.
2. Add a console application called Ch22Ex01Client to the solution.
3. On the Build menu, click Build Solution.
4. In the Ch22Ex01Client project, right click References in the Solution Explorer and select Add Service Reference.
5. In the Add Service Reference dialog box, click Discover.
6. When the development web server has started and information about the WCF service has been loaded, expand the reference to look at its details. Notice that there are two methods in the service: GetData and GetDataUsingDataContract.

7. Click OK to add the service reference.
8. Modify the code in Program.cs in the Ch22Ex01Client application as follows:

```
using Ch22Ex01Client.ServiceReference1;
using static System.Console;

namespace Ch22Ex01Client
{
    class Program
    {
        static void Main(string[] args)
        {
            Title = "Ch22Ex01Client";
            string numericInput = null;
            int intParam;
            do
            {
                WriteLine("Enter an integer and press enter to call the WCF service.");
                numericInput = ReadLine();
            }
            while (!int.TryParse(numericInput, out intParam));
            Service1Client client = new Service1Client();
            WriteLine(client.GetData(intParam));
            WriteLine("Press any key to exit.");
            ReadKey();
        }
    }
}
```

9. Right-click the Ch22Ex01Client project in the Solution Explorer and select Set as StartUp Project.
10. Run the application. Enter a number in the console application window and press Enter. The result is shown in Figure 22-1.



**FIGURE 22-1**

11. Exit the application, right-click the Service1.svc file in the Ch22Ex01 project in the Solution Explorer, and click View in Browser.
12. Review the information in the window.
13. Click the link at the top of the web page for the service to view the WSDL. Don't panic — you don't need to understand all the stuff in the WSDL file!

### **How It Works**

In this example you created a simple web server-hosted WCF service and console application client. You used the default Visual Studio template for a WCF service project, which meant that you didn't have to add any code. Instead, you used one of the operations defined in this default template,

`GetData()`. For the purposes of this example, the actual operation used isn't important; here, you are focusing on the structure of the code and the plumbing that makes things work.

First, look at the server project, Ch22Ex01. This consists of the following:

- A `Service1.svc` file that defines the hosting for the service
- A class definition, `CompositeType`, that defines a data contract used by the service (located in the `IService1.cs` code file)
- An interface definition, `IService1`, that defines the service contract and two operation contracts for the service
- A class definition, `Service1`, that implements `IService1` and defines the functionality of the service (located in the `Service1.svc.cs` code file)
- A `<system.serviceModel>` configuration section (in `Web.config`) that configures the service

The `Service1.svc` file contains the following line of code (to see this code, right-click the file in the Solution Explorer and select View Markup):

```
<%@ ServiceHost Language="C#" Debug="true" Service="Ch22Ex01.Service1"
CodeBehind="Service1.svc.cs" %>
```

This is a `ServiceHost` instruction that is used to tell the web server (the development web server in this case, although this also applies to IIS) what service is hosted at this address. The class that defines the service is declared in the `Service` attribute, and the code file that defines this class is declared in the `CodeBehind` attribute. This instruction is necessary in order to obtain the hosting features of the web server as defined in the previous sections.

Obviously, this file is not required for WCF services that aren't hosted in a web server. You'll learn how to self-host WCF services later in this chapter.

Next, the data contract `CompositeType` is defined in the `IService1.cs` file. You can see from the code that the data contract is simply a class definition that includes the `DataContract` attribute on the class definition and `DataMember` attributes on class members:

```
[DataContract]
public class CompositeType
{
    bool boolValue = true;
    string stringValue = "Hello ";
    [DataMember]
    public bool BoolValue
    {
        get { return boolValue; }
        set { boolValue = value; }
    }
    [DataMember]
    public string StringValue
    {
        get { return stringValue; }
        set { stringValue = value; }
    }
}
```

This data contract is exposed to the client application through metadata (if you looked through the WSDL file in the example you might have seen this). This enables client applications to define a type that can be serialized into a form that can be deserialized by the service into a `CompositeType` object. The client doesn't need to know the actual definition of this type; in fact, the class used by the client might have a different implementation. This simple way of defining data contracts is surprisingly powerful, and enables the exchange of complex data structures between the WCF service and its clients.

The `IService1.cs` file also contains the service contract for the service, which is defined as an interface with the `ServiceContract` attribute. Again, this interface is completely described in the metadata for the service, and can be recreated in client applications. The interface members constitute the operations exposed by the service, and each is used to create an operation contract by applying the `OperationContract` attribute. The example code includes two operations, one of which uses the data contract you looked at earlier:

```
[ServiceContract]
public interface IService1
{
    [OperationContract]
    string GetData(int value);
    [OperationContract]
    CompositeType GetDataUsingDataContract(CompositeType composite);
}
```

All four of the contract-defining attributes that you have seen so far can be further configured with attributes, as shown in the next section. The code that implements the service looks much like any other class definition:

```
public class Service1 : IService1
{
    public string GetData(int value)
    {
        return string.Format("You entered: {0}", value);
    }
    public CompositeType GetDataUsingDataContract(CompositeType composite)
    {
        ...
    }
}
```

Note that this class definition doesn't need to inherit from a particular type, and doesn't require any particular attributes. All it needs to do is implement the interface that defines the service contract. In fact, you can add attributes to this class and its members to specify behaviors, but these aren't mandatory.

The separation of the service contract (the interface) from the service implementation (the class) works extremely well. The client doesn't need to know anything about the class, which could include much more functionality than just the service implementation. A single class could even implement more than one service contract.

Finally, you come to the configuration in the `Web.config` file. Configuration of WCF services in config files is a feature that has been taken from .NET remoting, and it works with all types of WCF services (hosted or self-hosted) as well as clients of WCF services (as shown in a moment). The vocabulary of

this configuration is such that you can apply pretty much any configuration that you can think of to a service, and you can even extend this syntax.

WCF configuration code is contained in the `<system.serviceModel>` configuration section of `.config` or `app.config` files. In this example, there is not a lot of service configuration, as default values are used. In the `Web.config` file, the configuration section consists of a single subsection that supplies overrides to default values for the service behavior `<behaviors>`. The code for the `<system.serviceModel>` configuration section in `Web.config` (with comments removed for clarity) is as follows:

```
<system.serviceModel>
  <behaviors>
    <serviceBehaviors>
      <behavior>
        <serviceMetadata httpGetEnabled="true" httpsGetEnabled="true" />
        <serviceDebug includeExceptionDetailInFaults="false" />
      </behavior>
    </serviceBehaviors>
  </behaviors>
</system.serviceModel>
```

This section can define one or more behaviors in `<behavior>` child sections, which can be reused on multiple other elements. A `<behavior>` section can be given a name to facilitate this reuse (so that it can be referenced from elsewhere), or can be used without a name (as in this example) to specify overrides to default behavior settings.

**NOTE** If nondefault configuration were being used, you would expect to see a `<services>` section inside `<system.serviceModel>`, containing one or more `<services>` child sections. In turn, the `<service>` sections can contain child `<endpoint>` sections, each of which (you guessed it) defines an endpoint for the service. In fact, the endpoints defined are base endpoints for the service. Endpoints for operations are inferred from these.

One of the default behavior overrides in `Web.config` is as follows:

```
<serviceDebug includeExceptionDetailInFaults="false"/>
```

This setting can be set to `true` to expose exception details in any faults that are transmitted to the client, which is something you would usually allow only in development.

The other default behavior override in `Web.config` relates to metadata. Metadata is used to enable clients to obtain descriptions of WCF services. The default configuration defines two default endpoints for services. One is the endpoint that clients use to access the service; the other is an endpoint used to obtain metadata from the service. This can be disabled in the `Web.config` file as follows:

```
<serviceMetadata httpGetEnabled="false" httpsGetEnabled="false" />
```

Alternatively, you could remove this line of configuration code entirely, as the default behavior does not enable metadata exchange.

If you try disabling this in the example it won't stop your client from being able to access the service, because it has already obtained the metadata it needed when you added the service reference. However,

disabling metadata will prevent other clients from using the Add Service Reference tool for this service. Typically, web services in a production environment will not need to expose metadata, so you should disable this functionality after the development phase is complete.

Without metadata, another common way to access a WCF service is to define its contracts in a separate assembly, which is referenced by both the hosting project and the client project. The client can then generate a proxy by using these contracts directly, rather than through exposed metadata.

Now that you've looked at the WCF service code, it's time to look at the client, and in particular at what using the Add Service Reference tool actually did. You will notice in the Solution Explorer that the client includes a folder called Service References, and if you expand that you will see an item called ServiceReference1, which was the name you chose when you added the reference.

The Add Service Reference tool creates all the classes you require to access the service. This includes a proxy class for the service that includes methods for all the operations exposed by the service (Service1Client), and a client-side class generated from the data contract (CompositeType).

**NOTE** You can browse through the code that is generated by the Add Service Reference tool if you want (by displaying all files in the project, including the hidden ones), although at this point it's probably best not to, because it contains quite a lot of confusing code.

The tool also adds a configuration file to the project, app.config. This configuration defines two things:

- Binding information for the service endpoint
- The address and contract for the endpoint

The binding information is taken from the service description:

```
<configuration>
  <system.serviceModel>
    <bindings>
      <basicHttpBinding>
        <binding name="BasicHttpBinding_IService1" />
      </basicHttpBinding>
    </bindings>
  </system.serviceModel>
</configuration>
```

This binding is used in the endpoint configuration, along with the base address of the service (which is the address of the .svc file for web server-hosted services) and the client-side version of the contract IServicel:

```
<client>
  <endpoint address="http://localhost:49227/Service1.svc"
    binding="basicHttpBinding"
    bindingConfiguration="BasicHttpBinding_IService1"
    contract="ServiceReference1.IService1"
    name="BasicHttpBinding_IService1" />
</client>
</system.serviceModel>
</configuration>
```

If you remove the `<bindings>` section as well as the `bindingConfiguration` attribute from the `<endpoint>` element, then the client will use the default binding configuration.

The `<binding>` element, which has the name `BasicHttpBinding_IService1`, is included so that you can use it to customize the configuration of the binding. There are a number of configuration settings that you might use here, ranging from timeout settings to message size limits and security settings. If these had been specified in the service project to be nondefault values, then you would have seen them in the `app.config` file, since they would have been copied across. In order for the client to communicate with the service, the binding configurations must match. You won't look at WCF service configuration in great depth in this chapter.

This example has covered a lot of ground, and it is worth summarizing what you have learned before moving on:

- WCF service definitions:
  - Services are defined by a service contract interface that includes operation contract members.
  - Services are implemented in a class that implements the service contract interface.
  - Data contracts are simply type definitions that use data contract attributes.
- WCF service configuration:
  - You can use configuration files (`web.config` or `app.config`) to configure WCF services.
- WCF web server hosting:
  - Web server hosting uses `.svc` files as service base addresses.
- WCF client configuration:
  - You can use configuration files (`web.config` or `app.config`) to configure WCF service clients.

The following section explores contracts in more detail.

---

## The WCF Test Client

In the previous Try It Out, you created both a service and a client in order to look at how the basic WCF architecture works and how configuration of WCF services is achieved. In practice, though, the client application you want to use might be complex, and it can be tricky to test services properly.

To ease the development of WCF services, Visual Studio provides a test tool you can use to ensure that your WCF operations work correctly. This tool is automatically configured to work with your WCF service projects, so if you run your project the tool will appear. All you need to do is ensure that the service you want to test (that is, the `.svc` file) is set to be the startup page for the WCF service project. Alternatively, you can run the test client as a standalone application. You can find the test client on 64-bit operating systems at `C:\Program Files (x86)\Microsoft Visual Studio 14.0\Common7\IDE\WcfTestClient.exe`.

If you are using a 32-bit operating system, the path is the same except the root folder is Program Files.

The tool enables you to invoke service operations and inspect the service in some other ways. The following Try It Out illustrates this.

### TRY IT OUT Using the WCF Test Client: Ch22Ex01\Web.config

1. Open the WCF Service Application project from the previous Try It Out, Ch22Ex01.
2. Right-click the Service1.svc service in Solution Explorer and click Set As Start Page.
3. Right-click the Ch22Ex01 project in Solution Explorer and click Set As StartUp Project.
4. In Web.config, ensure that metadata is enabled:
 

```
<serviceMetadata httpGetEnabled="true" httpsGetEnabled="true" />
```
5. Run the application. The WCF test client appears.
6. In the left pane of the test client, double-click Config File. The config file used to access the service is displayed in the right pane.
7. In the left pane, double-click the GetDataUsingDataContract() operation.
8. In the pane that appears on the right, change the value of BoolValue to True and StringValue to Test String, and then click Invoke.
9. If a security prompt dialog box appears, click OK to confirm that you are happy to send information to the service.
10. The operation result appears, as shown in Figure 22-2.




FIGURE 22-2

11. Click the XML tab to view the request and response XML.
12. Close the WCF Test Client. This will stop debugging in Visual Studio.

### How It Works

In this example you used the WCF test client to inspect and invoke an operation on the service you created in the previous Try It Out. The first thing you probably noticed is a slight delay while the service is loaded. This is because the test client has to inspect the service to determine its capabilities. This discovery uses the same metadata as the Add Service Reference tool, which is why you must ensure that metadata is available (it's possible you experimented with disabling it in the previous Try It Out). Once discovery is complete, you can view the service and its operations in the left pane of the tool.

Next, you looked at the configuration used to access the service. As with the client application from the previous Try It Out, this is generated automatically from the service metadata, and contains exactly the same code. You can edit this configuration file through the tool if you need to, by right-clicking on the Config File item and clicking Edit WCF Configuration. An example of this configuration is shown in Figure 22-3, which includes the binding configuration options mentioned earlier in this chapter.




FIGURE 22-3

Finally, you invoked an operation. The test client allows you to enter the parameters to use and invoke the method, then displays the result, all without you writing any client code. You also saw how to view the actual XML that is sent and received to obtain the result. This information is quite technical, but it can be absolutely critical when debugging more complex services.

## Defining WCF Service Contracts

The previous examples showed how the WCF infrastructure makes it easy for you to define contracts for WCF services with a combination of classes, interfaces, and attributes. This section takes a deeper look at this technique.

### Data Contracts

To define a data contract for a service, you apply the `DataContractAttribute` attribute to a class definition. This attribute is found in the `System.Runtime.Serialization` namespace. You can configure this attribute with the properties shown in Table 22-3.

**TABLE 22-3:** DataContractAttribute Properties

PROPERTY	DESCRIPTION
Name	Names the data contract with a different name than the one you use for the class definition. This name will be used in SOAP messages and client-side data objects that are defined from service metadata.
Namespace	Defines the namespace that the data contract uses in SOAP messages.
IsReference	Affects the way that objects are serialized. If this is set to <code>true</code> , then an object instance is serialized only once even if it is referenced several times, which can be important in some situations. The default is <code>false</code> .

The `Name` and `Namespace` properties are useful when you need interoperability with existing SOAP message formats (as are the similarly named properties for other contracts), but otherwise you will probably not require them.

Each class member that is part of a data contract must use the `DataMemberAttribute` attribute, which is also found in the `System.Runtime.Serialization` namespace. Table 22-4 lists this attribute's properties.

**TABLE 22-4:** DataMemberAttribute Properties

PROPERTY	DESCRIPTION
Name	Specifies the name of the data member when serialized (the default is the member name).
IsRequired	Specifies whether the member must be present in a SOAP message.
Order	An <code>int</code> value specifying the order of serializing or deserializing the member, which might be required if one member must be present before another can be understood. Lower <code>Order</code> members are processed first.
EmitDefaultValue	Set this to <code>false</code> to prevent members from being included in SOAP messages if their value is the default value for the member.

## Service Contracts

Service contracts are defined by applying the `System.ServiceModel.ServiceContractAttribute` attribute to an interface definition. You can customize the service contract with the properties shown in Table 22-5.

TABLE 22-5: ServiceContractAttribute Properties

PROPERTY	DESCRIPTION
Name	Specifies the name of the service contract as defined in the <code>&lt;portType&gt;</code> element in WSDL.
Namespace	Defines the namespace of the service contract used by the <code>&lt;portType&gt;</code> element in WSDL.
ConfigurationName	The name of the service contract as used in the configuration file.
HasProtectionLevel	Determines whether messages used by the service have explicitly defined protection levels. Protection levels enable you to sign, or sign and encrypt, messages.
ProtectionLevel	The protection level to use for message protection.
SessionMode	Determines whether sessions are enabled for messages. If you use sessions, then you can ensure that messages sent to different endpoints of a service are correlated — that is, they use the same service instance and so can share state, and so on.
CallbackContract	For duplex messaging the client exposes a contract as well as the service. This is because, as discussed earlier, the client in duplex communications also acts as a server. This property enables you to specify which contract the client uses.

## Operation Contracts

Within interfaces that define service contracts, you define members as operations by applying the `System.ServiceModel.OperationContractAttribute` attribute. This attribute has the properties described in Table 22-6.

TABLE 22-6: OperationContractAttribute Properties

Property	Description
Name	Specifies the name of the service operation. The default is the member name.
IsOneWay	Specifies whether the operation returns a response. If you set this to true, then clients won't wait for the operation to complete before continuing.

Property	Description
AsyncPattern	If set to true, the operation is implemented as two methods that you can use to call the operation asynchronously: <code>Begin&lt;methodName&gt;()</code> and <code>End&lt;methodName&gt;()</code> .
HasProtectionLevel	See the previous section.
ProtectionLevel	See the previous section.
IsInitiating	If sessions are used, then this property determines whether calling this operation can start a new session.
IsTerminating	If sessions are used, then this property determines whether calling this operation terminates the current session.
Action	If you are using addressing (an advanced capability of WCF services), then an operation has an associated action name, which you can specify with this property.
ReplyAction	As with <code>Action</code> , but specifies the action name for the operation response.

**NOTE** In the .NET 4.5 Framework, when you add a service reference, Visual Studio also generates asynchronous proxy methods to call the service, regardless of whether `AsyncPattern` is set to true. These methods, which have the suffix `Async`, use the new asynchronous techniques that are included in .NET 4.5, and are asynchronous only from the point of view of the calling code. Internally, they call the synchronous WCF operations.

## Message Contracts

The earlier example didn't use message contract specifications. If you use these, then you do so by defining a class that represents the message and applying the `MessageContractAttribute` attribute to the class. You then apply `MessageBodyMemberAttribute`, `MessageHeaderAttribute`, or `MessageHeaderArrayAttribute` attributes to members of this class. All these attributes are in the `System.ServiceModel` namespace. You are unlikely to want to do this unless you need a very high degree of control over the SOAP messages used by WCF services, so details are not provided here.

## Fault Contracts

If you have a particular exception type — for example, a custom exception — that you want to make available to client applications, then you can apply the `System.ServiceModel.FaultContractAttribute` attribute to the operation that might generate this exception.

**TRY IT OUT** WCF Contracts: Ch22Ex02Contracts

1. Create a new WCF Service Application project called Ch22Ex02 in the directory C:\BegVCSharp\Chapter22.
2. Add a class library project called Ch22Ex02Contracts to the solution and remove the Class1.cs file.
3. Add references to the System.Runtime.Serialization.dll and System.ServiceModel.dll assemblies to the Ch22Ex02Contracts project.
4. Add a class called Person to the Ch22Ex02Contracts project and modify the code in Person.cs as follows:

```
using System.Runtime.Serialization;
namespace Ch22Ex02Contracts
{
    [DataContract]
    public class Person
    {
        [DataMember]
        public string Name { get; set; }
        [DataMember]
        public int Mark { get; set; }
    }
}
```

5. Add an interface called IAwardService to the Ch22Ex02Contracts project and modify the code in IAwardService.cs as follows:

```
using System.ServiceModel;
namespace Ch22Ex02Contracts
{
    [ServiceContract(SessionMode = SessionMode.Required)]
    public interface IAwardService
    {
        [OperationContract(IsOneWay = true, IsInitiating = true)]
        void SetPassMark(int passMark);
        [OperationContract]
        Person[] GetAwardedPeople(Person[] peopleToTest);
    }
}
```

6. In the Ch22Ex02 project, add a reference to the Ch22Ex02Contracts project.
7. Remove IService1.cs and Service1.svc from the Ch22Ex02 project.
8. Add a new WCF service called AwardService to Ch22Ex02.
9. Remove the IAwardService.cs file from the Ch22Ex02 project.
10. Modify the code in AwardService.svc.cs as follows:

```
using System.Collections.Generic;
using Ch22Ex02Contracts;
namespace Ch22Ex02
{
    public class AwardService : IAwardService
```

```

    {
        private int passMark;
        public void SetPassMark(int passMark)
        {
            this.passMark = passMark;
        }
        public Person[] GetAwardedPeople(Person[] peopleToTest)
        {
            List<Person> result = new List<Person>();
            foreach (Person person in peopleToTest)
            {
                if (person.Mark > passMark)
                {
                    result.Add(person);
                }
            }
            return result.ToArray();
        }
    }
}

```

11. Modify the service configuration section in `Web.config` as follows:

```

<system.serviceModel>
    <protocolMapping>
        <add scheme="http" binding="wsHttpBinding" />
    </protocolMapping>
    ...
</system.serviceModel>

```

12. Open the project properties for Ch22Ex02. In the Web section, make a note of the port used in the hosting settings. If you don't have IIS installed, you can set a specific port for use in the Visual Studio Development Server instead.
13. Add a new console project called Ch22Ex02Client to the solution and set it as the startup project.
14. Add references to the `System.ServiceModel.dll` assembly and the Ch22Ex02Contracts project to the Ch22Ex02Client project.
15. Modify the code in `Program.cs` in Ch22Ex02Client as follows (ensure that you use the port number you obtained earlier in the `EndpointAddress` constructor, the example code uses port 49284):

```

using System;
using static System.Console;
using System.ServiceModel;
using Ch22Ex02Contracts;
namespace Ch22Ex02Client
{
    class Program
    {
        static void Main(string[] args)
        {
            Person[] people = new Person[]
            {
                new Person { Mark = 46, Name="Jim" },
                new Person { Mark = 73, Name="Mike" },
                new Person { Mark = 92, Name="Stefan" },
            }
        }
    }
}


```

```
    new Person { Mark = 24, Name="Arthur" }
};

WriteLine("People:");
OutputPeople(people);
IAwardService client = ChannelFactory<IAwardService>.CreateChannel(
    new WSHttpBinding(),
    new EndpointAddress("http://localhost:38831/AwardService.svc"));
client.SetPassMark(70);
Person[] awardedPeople = client.GetAwardedPeople(people);
WriteLine();
WriteLine("Awarded people:");
OutputPeople(awardedPeople);
.ReadKey();
}

static void OutputPeople(Person[] people)
{
    foreach (Person person in people)
        WriteLine("{0}, mark: {1}", person.Name, person.Mark);
}
}
```

16. If you are using IIS, simply run the application. If you are using the development server, you must ensure the development server is running for the service, so run the service project first. You can do this by setting the Ch22Ex02 project as the startup project and then pressing Ctrl+F5. This will start the service without debugging. Then set the startup project to the Ch22Ex02Client project again and press F5. The result is shown in Figure 22-4.



**FIGURE 22-4**

## *How It Works*

In this example, you created a set of contracts in a class library project and used that class library in both a WCF service and a client. The service, as in the previous example, is hosted in a web server. The configuration for this service is reduced to the bare minimum.

The main difference in this example is that no metadata is required by the client, as the client has access to the contract assembly. Instead of generating a proxy class from metadata, the client obtains a reference to the service contract interface through an alternative method. Another point to note about this example is the use of a session to maintain state in the service, which requires the `WSHttpBinding` binding instead of the `BasicHttpBinding` binding.

The data contract used in this example is for a simple class called `Person`, which has a string property called `Name` and an `int` property called `Mark`. You used the `DataContractAttribute` and

`DataMemberAttribute` attributes with no customization, and there is no need to reiterate the code for this contract here.

The service contract is defined by applying the `ServiceContractAttribute` attribute to the `IAwardService` interface. The `SessionMode` property of this attribute is set to `SessionMode.Required`, as this service requires state:

```
[ServiceContract(SessionMode=SessionMode.Required)]
public interface IAwardService
{
```

The first operation contract, `SetPassMark()`, is the one that sets state, and therefore has the `IsInitiating` property of `OperationContractAttribute` set to `true`. This operation doesn't return anything, so it is defined as a one-way operation by setting `IsOneWay` to `true`:

```
[OperationContract(IsOneWay=true, IsInitiating=true)]
void SetPassMark(int passMark);
```

The other operation contract, `GetAwardedPeople()`, does not require any customization and uses the data contract defined earlier:

```
[OperationContract]
Person[] GetAwardedPeople(Person[] peopleToTest);
}
```

Remember that these two types, `Person` and `IAwardService`, are available to both the service and the client. The service implements the `IAwardService` contract in a type called `AwardService`, which doesn't contain any remarkable code. The only difference between this class and the service class you saw earlier is that it is stateful. This is permissible, as a session is defined to correlate messages from a client.

To ensure that the service uses the `WSHttpBinding` binding, you added the following to `Web.config` for the service:

```
<protocolMapping>
  <add scheme="http" binding="wsHttpBinding" />
</protocolMapping>
```

This overrides the default mapping for HTTP binding. Alternatively, you could configure the service manually and keep the existing default, but this override is much simpler. However, be aware that this type of override is applied to all services in a project. If you have more than one service in a project, then you would have to ensure that this binding is acceptable to each of them.

The client is more interesting, primarily because of this code:

```
IAwardService client = ChannelFactory<IAwardService>.CreateChannel(
  new WSHttpBinding(),
  new EndpointAddress("http://localhost:38831/AwardService.svc"));
```

The client application has no `app.config` file to configure communications with the service, and no proxy class defined from metadata to communicate with the service. Instead, a proxy class is created through the `ChannelFactory<T>.CreateChannel()` method. This method creates a proxy class that implements the `IAwardService` client, although behind the scenes the generated class communicates with the service just like the metadata-generated proxy shown earlier.

**NOTE** If you create a proxy class with `ChannelFactory<T>.CreateChannel()`, the communication channel will, by default, time out after a minute, which can lead to communication errors. There are ways to keep connections alive, but they are beyond the scope of this chapter.

Creating proxy classes in this way is an extremely useful technique that you can use to quickly generate a client application on-the-fly.

## Self-Hosted WCF Services

So far in this chapter you have seen WCF services that are hosted in web servers. This enables you to communicate across the Internet, but for local network communications it is not the most efficient way of doing things. For one thing, you need a web server on the computer that hosts the service. In addition, the architecture of your applications might be such that having an independent WCF service isn't desirable.

Instead, you might want to use a self-hosted WCF service. A *self-hosted* WCF service exists in a process that you create, rather than in the process of a specially made hosting application such as a web server. This means, for example, that you can use a console application or Windows application to host your service.

To self-host a WCF service, you use the `System.ServiceModel.ServiceHost` class. You instantiate this class with either the type of the service you want to host or an instance of the service class. You can configure a service host through properties or methods, or (and this is the clever part) through a configuration file. In fact, host processes, such as web servers, use a `ServiceHost` instance to do their hosting. The difference when self-hosting is that you interact with this class directly. However, the configuration you place in the `<system.serviceModel>` section of the `app.config` file for your host application uses exactly the same syntax as the configuration sections you've already seen in this chapter.

You can expose a self-hosted service through any protocol that you like, although typically you will use TCP or named pipe binding in this type of application. Services accessed through HTTP are more likely to live inside web server processes, because you get the additional functionality that web servers offer, such as security and other features.

If you want to host a service called `MyService`, you could use code such as the following to create an instance of `ServiceHost`:

```
ServiceHost host = new ServiceHost(typeof(MyService));
```

If you want to host an instance of `MyService` called `myServiceObject`, you could code as follows to create an instance of `ServiceHost`:

```
MyService myServiceObject = new MyService();
ServiceHost host = new ServiceHost(myServiceObject);
```

**WARNING** Hosting a service instance in a ServiceHost works only if you configure the service so that calls are always routed to the same object instance. To do this, you must apply a ServiceBehaviorAttribute attribute to the service class and set the InstanceContextMode property of this attribute to InstanceContextMode.Single.

After creating a ServiceHost instance you can configure the service and its endpoints and binding through properties. Alternatively, if you put your configuration in a .config file, the ServiceHost instance will be configured automatically.

To start hosting a service once you have a configured ServiceHost instance, you use the ServiceHost.Open() method. Similarly, you stop hosting the service through the ServiceHost.Close() method. When you first start hosting a TCP-bound service, you might, if you have it enabled, receive a warning from the Windows Firewall service, as it will block the TCP port by default. You must open the TCP port for the service to begin listening on the port.

In the following Try it Out you use self-hosting techniques to expose some functionality of a WPF application through a WCF service.

### TRY IT OUT Self-Hosted WCF Services: Ch22Ex03

1. Create a new WPF application called Ch22Ex03 in the directory C:\BegVCSharp\Chapter22.
2. Add a new WCF service to the project called AppControlService by using the Add New Item Wizard.
3. Modify the code in MainWindow.xaml as follows:

```
<Window x:Class="Ch22Ex03.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    Title="Stellar Evolution" Height="450" Width="430"
    Loaded="Window_Loaded" Closing="Window_Closing">
    <Grid Height="400" Width="400" HorizontalAlignment="Center"
        VerticalAlignment="Center">
        <Rectangle Fill="Black" RadiusX="20" RadiusY="20"
            StrokeThickness="10">
            <Rectangle.Stroke>
                <LinearGradientBrush EndPoint="0.358,0.02"
                    StartPoint="0.642,0.98">
                    <GradientStop Color="#FF121A5D" Offset="0" />
                    <GradientStop Color="#FFB1B9FF" Offset="1" />
                </LinearGradientBrush>
            </Rectangle.Stroke>
        </Rectangle>
        <Ellipse Name="AnimatableEllipse" Stroke="{x:Null}" Height="0"
            Width="0" HorizontalAlignment="Center">
    
```

```
    VerticalAlignment="Center">
    <Ellipse.Fill>
        <RadialGradientBrush>
            <GradientStop Color="#FFFFFF" Offset="0" />
            <GradientStop Color="#FFFFFF" Offset="1" />
        </RadialGradientBrush>
    </Ellipse.Fill>
    <Ellipse.Effect>
        <DropShadowEffect ShadowDepth="0" Color="#FFFFFF"
            BlurRadius="50" />
    </Ellipse.Effect>
</Ellipse>
</Grid>
</Window>
```

4. Modify the code in MainWindow.xaml.cs as follows:

```
using System.Windows.Shapes;
using System.ServiceModel;
using System.Windows.Media.Animation;
namespace Ch22Ex03
{
    /// <summary>
    /// Interaction logic for MainWindow.xaml
    /// </summary>
    public partial class MainWindow : Window
    {
        private AppControlService service;
        private ServiceHost host;
        public MainWindow()
        {
            InitializeComponent();
        }
        private void Window_Loaded(object sender, RoutedEventArgs e)
        {
            service = new AppControlService(this);
            host = new ServiceHost(service);
            host.Open();
        }
        private void Window_Closing(object sender,
            System.ComponentModel.CancelEventArgs e)
        {
            host.Close();
        }
        internal void SetRadius(double radius, string foreTo,
            TimeSpan duration)
        {
            if (radius > 200)
            {
                radius = 200;
            }
            Color foreToColor = Colors.Red;
            try
            {
                foreToColor = (Color)ColorConverter.ConvertFromString(foreTo);
            }
```

```
        catch
        {
            // Ignore color conversion failure.
        }
        Duration animationLength = new Duration(duration);
        DoubleAnimation radiusAnimation = new DoubleAnimation(
            radius * 2, animationLength);
        ColorAnimation colorAnimation = new ColorAnimation(
            foreToColor, animationLength);
        AnimatableEllipse.BeginAnimation(Ellipse.HeightProperty,
            radiusAnimation);
        AnimatableEllipse.BeginAnimation(Ellipse.WidthProperty,
            radiusAnimation);
        ((RadialGradientBrush)AnimatableEllipse.Fill).GradientStops[1]
            .BeginAnimation(GradientStop.ColorProperty, colorAnimation);
    }
}
```

5. Modify the code in `IAppControlService.cs` as follows:

```
[ServiceContract]
public interface IAppControlService
{
    [OperationContract]
    void SetRadius(int radius, string foreTo, int seconds);
}
```

6. Modify the code in AppControlService.cs as follows:

```
[ServiceBehavior(InstanceContextMode=InstanceContextMode.Single)]
public class AppControlService : IAppControlService
{
    private MainWindow hostApp;
    public AppControlService(MainWindow hostApp)
    {
        this.hostApp = hostApp;
    }
    public void SetRadius(int radius, string foreTo, int seconds)
    {
        hostApp.SetRadius(radius, foreTo, new TimeSpan(0, 0, seconds));
    }
}
```

7. Modify the code in app.config as follows:

```
<configuration>
  <system.serviceModel>
    <services>
      <service name="Ch22Ex03.AppControlService">
        <endpoint address="net.tcp://localhost:8081/AppControlService"
          binding="netTcpBinding"
          contract="Ch22Ex03.IAppControlService" />
      </service>
    </services>
  </system.serviceModel>
</configuration>
```

8. Add a new console application to the project called Ch22Ex03Client.
9. Right-click the solution in the Solution Explorer and click Set StartUp Projects.
10. Configure the solution to have multiple startup projects, with both projects being started simultaneously.
11. Add references to System.ServiceModel.dll and Ch22Ex03 to the Ch22Ex03Client project.
12. Modify the code in Program.cs as follows:

```
using Ch22Ex03;
using System.ServiceModel;
using static System.Console;
namespace Ch22Ex03Client
{
    class Program
    {
        static void Main(string[] args)
        {
            WriteLine("Press enter to begin.");
            ReadLine();
            WriteLine("Opening channel.");
            IAppControlService client =
                ChannelFactory<IAppControlService>.CreateChannel(
                    new NetTcpBinding(),
                    new EndpointAddress(
                        "net.tcp://localhost:8081/AppControlService"));
            WriteLine("Creating sun.");
            client.SetRadius(100, "yellow", 3);
            WriteLine("Press enter to continue.");
            ReadLine();
            WriteLine("Growing sun to red giant.");
            client.SetRadius(200, "Red", 5);
            WriteLine("Press enter to continue.");
            ReadLine();
            WriteLine("Collapsing sun to neutron star.");
            client.SetRadius(50, "AliceBlue", 2);
            WriteLine("Finished. Press enter to exit.");
            ReadLine();
        }
    }
}
```

13. Run the solution. If prompted, unblock the Windows Firewall TCP port so that the WCF can listen for connections.
14. When both the Stellar Evolution window and the console application window are displayed, press Enter in the console window. The result is shown in Figure 22-5.




FIGURE 22-5

15. Continue pressing Enter in the console window to continue the stellar evolution cycle.
16. Close the Stellar Evolution window to stop debugging.

### How It Works

In this example you have added a WCF service to a WPF application and used it to control the animation of an `Ellipse` control. You have created a simple client application to test the service. Don't worry too much about the XAML code in this example if you are not familiar with WPF yet; it's the WCF plumbing that is of interest here.

The WCF service, `AppControlService`, exposes a single operation, `SetRadius()`, which clients call to control the animation. This method communicates with an identically named method defined in the `Window1` class for the WPF application. For this to work, the service needs a reference to the application, so you must host an object instance of the service. As discussed previously, this means that the service must use a behavior attribute:

```
[ServiceBehavior(InstanceContextMode=InstanceContextMode.Single)]
public class AppControlService : IAppControlService
{
    ...
}
```

In `Window1.xaml.cs`, the service instance is created in the `Windows_Loaded()` event handler. This method also begins hosting by creating a `ServiceHost` object for the service and calling its `Open()` method:

```
public partial class Window1 : Window
{
    private AppControlService service;
    private ServiceHost host;
    ...
    private void Window_Loaded(object sender, RoutedEventArgs e)
    {
        service = new AppControlService(this);
        host = new ServiceHost(service);
        host.Open();
    }
}
```

When the application closes, hosting is terminated in the `Window_Closing()` event handler.

The configuration file is again about as simple as it can be. It defines a single endpoint for the WCF service that listens at a `net.tcp` address, on port 8081, and uses the default `NetTcpBinding` binding:

```
<service name="Ch22Ex03.AppControlService">
    <endpoint address="net.tcp://localhost:8081/AppControlService"
        binding="netTcpBinding"
        contract="Ch22Ex03.IAppControlService" />
</service>
```

This matches up with code in the client app:

```
IAppControlService client =
    ChannelFactory<IAppControlService>.CreateChannel(
        new NetTcpBinding(),
        new EndpointAddress(
            "net.tcp://localhost:8081/AppControlService"));
```

When the client has created a client proxy class, it can call the `SetRadius()` method with radius, color, and animation duration parameters, and these are forwarded to the WPF application through the service. Simple code in the WPF application then defines and uses animations to change the size and color of the ellipse.

This code would work across a network if you used a machine name, rather than `localhost`, and if the network permitted traffic on the specified port. Alternatively, you could separate the client and host application further, and connect across the Internet. Either way, WCF services provide an excellent means of communication that doesn't take much effort to set up.

## EXERCISES

### 22.1 Which of the following applications can host WCF services?

- a. Web applications
- b. Windows Forms applications
- c. Windows services

- d. COM+ applications
  - e. Console applications
- 

**22.2** Which type of contract would you implement if you wanted to exchange parameters of type `MyClass` with a WCF service? Which attributes would you require?

---

**22.3** If you host a WCF service in a web application, what extension will the base endpoint for the service use?

---

**22.4** When self-hosting WCF services, you must configure the service by setting properties and calling methods of the `ServiceHost` class. True or false?

---

**22.5** Provide the code for a service contract, `IMusicPlayer`, with operations defined for `Play()`, `Stop()`, and `GetTrackInformation()`. Use one-way methods where appropriate. What other contracts might you define for this service to work?

---

Answers to the exercises can be found in Appendix A.

## ► WHAT YOU LEARNED IN THIS CHAPTER

TOPIC	KEY CONCEPTS
<b>WCF fundamentals</b>	WCF provides a framework for creating and communicating with remote services. It combines elements of the web service and remoting architectures along with new technologies to achieve this.
<b>Communication protocols</b>	You can communicate with a WCF service by any one of several protocols, including HTTP and TCP. This means that you can use services that are local to your client application, or that are separated by machine or network boundaries. To do this, you access a specific endpoint for the service through a binding corresponding to the protocol and features that you require. You can control these features, such as using session state or exposing metadata, through behaviors. .NET 4.5 includes many default settings to make it very easy to define a simple service.
<b>Communication payload</b>	Typically, calls to responses from WCF services are encoded as SOAP messages. However, there are alternatives, such as plain HTTP messages, and you can define your own payload types from scratch if you need to.
<b>Hosting</b>	WCF services might be hosted in IIS or in a Windows service, or they can be self-hosted. Using a host such as IIS enables you to make use of the host's built-in capabilities, including security and application pooling. Self-hosting is more flexible, but it can require more configuration and coding.
<b>Contracts</b>	You define the interface between a WCF service and client code through contracts. Services themselves, along with any operations they expose, are defined with service and operation contracts. Data types are defined with data contracts. Further customization of communications is achieved with message and fault contracts.
<b>Client applications</b>	Client applications communicate with WCF services by means of a proxy class. Proxy classes implement the service contract interface for the service, and any calls to operation methods of this interface are redirected to the service. You can generate a proxy by using the Add Service Reference tool, or you can create one programmatically through channel factory methods. In order for communications to succeed, the client must be configured to match the service configuration.

# 23

## Universal Apps

### **WHAT YOU WILL LEARN IN THIS CHAPTER**

---

- Enabling your Windows 10 device for development
- Developing Windows Universal apps using XAML and C#
- Using common Windows Universal Apps
- Packaging and deploying an app

### **WROX.COM CODE DOWNLOADS FOR THIS CHAPTER**

You can find the wrox.com code downloads for this chapter at [www.wrox.com/go/beginningvisualcsharp2015programming](http://www.wrox.com/go/beginningvisualcsharp2015programming) on the Download Code tab. The code is in the Chapter 23 download and individually named according to the names throughout the chapter.

Windows Universal apps is a hot topic for Windows developers all over the world. With the release of Windows 8, Microsoft took a huge leap from targeting the desktop and laptop computers almost exclusively toward becoming a real player on the market for tablet PCs and smart phones. Windows 8 shipped with a new API for developing apps and a Windows Store that allows users to download apps in a secure and predictable way. With Windows 10 and the Universal Windows Platform (UWP), Microsoft has taken app productivity to the next level by introducing Universal apps. These apps can target all Windows platforms, from phones over Xbox to the Windows desktop.

## **GETTING STARTED**

Writing Universal apps requires a few initial steps before you can get going. In the previous version of Visual Studio, you were required to get a Windows 8 Developer License that should be renewed quite often. For Windows 10, this is no longer needed for development, though you still need a store account to be able to publish the app. While developing the app, you can simply register your Windows 10 device for development.

Before you can start working on Windows Universal apps, you must enable development on your device and, unless they're already installed, you must install Universal Windows App Development Tools.

If you are using Visual Studio Express for Windows 10, or if you open a solution to create a Windows Universal app in another version of Visual Studio, you will get prompted to enable Developer Mode with the dialog shown in Figure 23.1. When you see this dialog, click the link "settings for developers," select the "Developer Mode" option, and then click yes to the warning that you are selecting a less secure option.




FIGURE 23-1

**NOTE** Developer mode has two options: Sideloaded apps and Developer mode. Sideloaded apps is a more secure option because in this mode you cannot install apps that are not trusted on the device. Developer mode, however, allows you to debug your apps on the device, so this is what you need for this chapter.

You may not have the Universal Windows App Development Tools installed. This comes automatically with some versions of Visual Studio, but if you don't have it, then simply open the New projects dialog  $\Rightarrow$  Visual C#  $\Rightarrow$  Windows  $\Rightarrow$  Universal, and you should see a link to the installer. Click this link to install the tools.

## UNIVERSAL APPS

Windows Universal Apps are apps that can target multiple device types. Traditional applications, like the WPF desktop game you wrote earlier in this book, target a single device type, such as a PC. With the introduction of the Universal Windows Platform, Microsoft has made it possible to write a single app that is able to run on multiple devices, and much effort has been put into making development of this kind of app a pleasant experience for the developer.

The primary challenges of developing apps that work on a large set of heterogeneous devices are that you can't know in advance how large the screen is or how the user will interact with the device.

If you simply scrunch the Karli Card WPF application from earlier in the book to the screen size of a phone, it will look terrible on even the largest of phones. Another aspect of this is that phone users will expect your app to be able to adjust its orientation on the screen. In this chapter we will introduce the concepts of responsive UI and adaptive triggers to solve these problems.

Universal Apps are deployed through the Windows Store, and this presents its own set of challenges for packaging the app. In order to get your app onto the store, you must undergo a fairly rigorous testing process and pass a number of requirements set by Microsoft. In the final part of this chapter we will examine this process so that you are ready to publish your own apps.

## APP CONCEPTS AND DESIGN

There are great differences in how applications display themselves on a phone and on the Windows Desktop. The design of applications running on the Windows Desktop is largely unchanged, albeit with much better graphics, since the introduction of Windows 95. The design features a window with a caption bar, three buttons in the top-right to maximize, minimize, and close the application and buttons, radio-buttons, check-boxes, and so on to display content. The generation of apps that was introduced with Windows 8 does things a little differently. They are designed to work with touch rather than mouse and keyboard, may or may not have a caption bar, and can swivel to fit the orientation of the device they are running on, just to mention a few differences.

When Microsoft launched Windows 8, they also released a fairly substantial design guide for apps, and you should be aware of this guide, even if you don't have to stick to it at all times. Even though apps will run on a diverse set of devices, they have a number of common traits that you should be aware of, so let's take a look at some of them and compare how Windows Store apps match up against desktop applications.

**NOTE** You can download the design guide for Windows 8 apps here: <http://go.microsoft.com/fwlink/p/?linkid=258743>.

## Screen Orientation

All Windows applications should be able to resize themselves gracefully. One aspect that is particularly important is the fact that handheld devices can move in three dimensions. Your users will expect your app to move with the orientation of the screen. So, if the user flips her tablet around, your app should follow the movement.

## Menus and Toolbars

Classic desktop apps use menus and toolbars for navigation between views. Universal apps can do so as well, but they are more likely to use toolbars than menus. Desktop apps usually display the visual components of the menu and toolbar all the time, but Universal apps will often choose not to do so to save precious real-estate on the smaller screen.

Rather than forcing your users to look at the complexity of your app through the menu, the app style presents the application to the users, and they can activate the menu when needed. When the menu is displayed, it should be simple, containing only the main options. It is up to you to decide where and when to display the menu.

## Tiles and Badges

Windows uses something called *live tiles* to display the apps on the Start menu and page. The “live” part of the name springs from the fact that the tiles can change based on the current content or state of the app. For example, you will see photo apps rotating through your pictures on the Start page, mail clients displaying the number of unread mails, games displaying screenshots from the last save, and so on. The possibilities are virtually endless.

Providing a good tile for your application is more important than providing a good icon for a desktop application, and that’s pretty important as well. Tiles are embedded in the manifest for the application, and, as you will see later in the chapter, they are easy to include using Visual Studio.

A *badge* is a small version of the tile that Windows can use on the Lock Screen and in other situations. You don’t have to provide a badge for your app unless it will show notifications on the Lock Screen.

## App Lifetime

Classic Windows Desktop applications can be closed by clicking a button in the top-right corner of the caption bar, but Universal apps don’t normally display a caption bar, so how do you close them? Generally speaking, you don’t need to close an app. Whenever a Universal app loses focus, it is suspended and will stop using processor resources entirely. This allows many apps to appear to be running at the same time, when in fact they are just suspended. The suspension happens automatically in Windows as soon as an app loses focus. It’s not really something that you notice as a user, but it is a very important fact to know and handle as an app developer.

## Lock Screen Apps

Some apps should keep running when they lose focus. Examples of this kind of app include GPS navigation and audio-streaming apps. Users expect these types of apps to continue running even if they start driving or begin using other apps. If your app needs to keep running in the background, you must declare it as a Lock Screen app and provide information to display notifications on the Lock Screen.

## APP DEVELOPMENT

When you start developing Windows Universal apps, you have a number of options regarding programming and UI language. This book uses C# and XAML, but other possibilities include using JavaScript and HTML5, C++ and DirectX, or Visual Basic and XAML.

The XAML that is used to create the user interfaces of the Universal apps is not entirely identical to the XAML used by WPF, but it is close enough that you should feel comfortable working with it. Many of the controls you are familiar with exist for Universal apps as well, though they tend to look slightly differently than their Windows Desktop counterparts. There are also a number of controls that are optimized to touch.

**NOTE** As was the case with Windows 8 apps, Microsoft has released a design guide for Universal Apps. You can find it here: <https://msdn.microsoft.com/library/windows/apps/hh465424.aspx>

## Adaptive Displays

Adaptive displays are displays that are able to change in response to user actions such as a phone being flipped on its side or the window changing size. Your app should be able to gracefully switch from portrait to landscape mode when the user flips her phone on the side and should work and look good regardless of whether it is deployed on a laptop or on a phone.

The first thing you will notice when you create a new Windows Universal app project is that the page displayed in the designer looks rather small. This is because this project defaults to a view that is optimized for a 5" phone display. You can change this using the Device Preview panel shown in Figure 23-2. You can also use this panel to change the layout from portrait to landscape.



FIGURE 23-2

A well-behaved app is able to display itself well in many if not all of the form-factors shown in the Device Preview panel. Considering that the range in this list is anything from a 569×320 pixels Internet of Things (IoT) device to a 3840×2160 pixels Surface Hub, this is a daunting task. Happily,

you will be aided by Visual Studio and the Universal Windows Platform framework. When you change the resolution (or screen size) from the drop-down, Visual Studio will resize your application, and you will immediately be able to see what the page looks like. In addition to that, controls that assist in creating an adaptive design for the application are included in the toolbox, and you can take advantage of them to easily create UIs that will transform nicely.

## Relative Panel

In Chapters 14 and 15 you used Grid and StackPanels controls to create a UI that worked well on a static display. But in a world where you must target many display sizes, you want something that will be better able to move the controls around for you. Enter the RelativePanel control.

The relative panel allows you to specify how controls should be positioned relative to one another. As you would expect, you can position controls to the left, right, above, or below other controls, but you can also do a few other nice tricks. It is possible to place a control in relation to the left, right, or center of another, both horizontally and vertically, and align the edges of the controls with the edges of the panel. This means no more fiddling with pixels to get two controls to line up perfectly on the display.

## Adaptive Triggers

Adaptive triggers are new to the Visual State Manager. Using these triggers you can change the layout of your application based on the size of the display. When combined with a relative panel, this is a very potent feature that in a fairly straightforward manner lets you build what the web-world refers to as responsive UIs and Microsoft calls adaptive displays.

### TRY IT OUT Adaptive Displays: Ch23Ex01

1. Create a new Windows Universal app project by selecting File  $\Rightarrow$  New  $\Rightarrow$  Project and expanding the Installed  $\Rightarrow$  Visual C#  $\Rightarrow$  Windows  $\Rightarrow$  Universal. Select the Blank App (Universal Windows) project and name it AdaptiveDisplay.
2. Add a RelativePanel control to the Grid. Set its margin to 20 and HorizontalAlignment to Stretch.
3. Add a textBlock and TextBox to the panel:

```
<RelativePanel HorizontalAlignment="Stretch" Margin="20" >
    <TextBlock x:Name="textBlockFirstName" Text="First name" Margin="0, 10, 10, 5" />
    <TextBox x:Name="textBoxFirstName" Text="" Width="400" RelativePanel.
        RightOf="textBlockFirstName" RelativePanel.AlignVerticalCenterWith="textBlockFirstN
        ame" />
</RelativePanel>
```

4. Add a Visual State Manager in the grid. It is critically important that it is the first child of the grid!

```
<VisualStateManager.VisualStateGroups>
    <VisualStateGroup>
        <VisualState x:Name="narrowView">
            <VisualState.StateTriggers>
                <AdaptiveTrigger MinWindowWidth="0" />
            </VisualState.StateTriggers>
            <VisualState.Setters>
```

```

        <Setter Target="textBoxFirstName.(RelativePanel.Below)"
               Value="textBlockFirstName" />
    <Setter
        Target="textBoxFirstName.(RelativePanel.AlignVerticalCenterWith)"
        Value="" />
    <Setter
        Target="textBoxFirstName.(RelativePanel.AlignLeftWith)"
        Value="textBlockFirstName" />
    </VisualState.Setters>
</VisualState>
<VisualState x:Name="wideView">
    <VisualState.StateTriggers>
        <AdaptiveTrigger MinWindowWidth="720" />
    </VisualState.StateTriggers>
</VisualState>
</VisualStateGroup>
</VisualStateManager.VisualStateGroups>

```

5. Change the target display in the Device Preview drop-down. When you select one of the smaller phone displays, the TextBox will pop down below the TextBlock. If you pick a tablet or another larger display, the TextBox will pop back up to the right of the TextBlock.

### *How It Works*

It is important that the Visual State Manager in the example is placed as the first child of the root grid. This allows the interpreter to find the controls that are referenced. You will not get any errors if you place it in another position, but you will not get the expected result.

The Visual State Manager uses the `AdaptiveTrigger` with the property `MinWindowWidth` to change the behavior of the display:

```
<AdaptiveTrigger MinWindowWidth="0" />
```

We define two states, one that is activated if the view is at least 0 pixels wide, and another that activates if the view is at least 720 pixels wide. You might expect that both will be active when the view is wider than 720 pixels, but that's not how it works. Rather, only one of the states will be active at any time, and the one that matches best will be selected. So, when the view is 1024 pixels wide, only the wide state is selected.

In the `narrowView`, we set three properties:

```

<VisualState.Setters>
    <Setter Target="textBoxFirstName.(RelativePanel.Below)"
           Value="textBlockFirstName" />
    <Setter
        Target="textBoxFirstName.(RelativePanel.
AlignVerticalCenterWith)"
        Value="" />
    <Setter
        Target="textBoxFirstName.(RelativePanel.AlignLeftWith)"
        Value="textBlockFirstName" />

```

First we ensure that the textbox should be moved below the `TextBlock`. Second, we clear the `AlignVerticalCenterWith` property. If we didn't change this, it will overrule the instruction to move

the control below the `textBlock`. This is because the `AlignVerticalCenterWith` property is set directly on the control, so if we leave it, it will take precedence over the `Below` directive of the View State. Another approach would be to refrain from setting any of the properties directly on the controls and only use view states. Finally, we align the left edges of the control.

The `wideView` state is in fact left empty. This means that no modifications to the properties defined directly on the controls should be made, making this the default state.

**NOTE** *The current version of Visual Studio sometimes fails to move the controls based on the selection in the Device Preview panel. If this happens, select another view size from the drop-down and the view should adjust correctly.*

## FlipView

The FlipView is a nice little control that works very well with handheld devices. It allows the user to swipe left or right to display some content. It is often used to display images one at a time and allows the user to use the swipe gesture to move between the images.

By default, the FlipView allows the user to swipe left or right to move the content in view, but this can be changed to move up or down. When a mouse is used, the scroll button will work as well.

### TRY IT OUT FlipView: Ch23Ex02

1. Create a new Windows Universal app project by selecting File  $\Rightarrow$  New  $\Rightarrow$  Project and expanding Installed  $\Rightarrow$  Visual C#  $\Rightarrow$  Windows  $\Rightarrow$  Universal. Select the Blank App (Universal Windows) project and name it `PictureViewer`.
2. Add three RelativePanels within the Grid tag on the MainPage.

```
<RelativePanel Margin="20">
    <RelativePanel x:Name="LeftPanel" Margin="0,10,0,0" >
        </RelativePanel>
        <RelativePanel x:Name="RightPanel" Margin="20,10,0,0">
            </RelativePanel>
    </RelativePanel>
```

3. Add a FlipView to the panel named LeftPanel like this:

```
<FlipView x:Name="flipView" VerticalAlignment="Stretch"
    HorizontalAlignment="Stretch" >
    <FlipView.ItemTemplate>
        <DataTemplate>
            <Image x:Name="image" Source="{Binding}" Stretch="Uniform"/>
        </DataTemplate>
    </FlipView.ItemTemplate>
</FlipView>
```

4. Add three TextBlocks to the panel named RightPanel like this:

```

<TextBlock x:Name="textBlockCurrentImageDisplayName"
    Margin="10,10,10,0" FontSize="24" FontWeight="Bold"
    RelativePanel.AlignLeftWithPanel="True"
    RelativePanel.AlignRightWithPanel="True" />
<TextBlock x:Name="textBlockCurrentImageImageHeight" Margin="10,10,10,0"
    FontSize="24" FontWeight="Bold"
    RelativePanel.AlignLeftWithPanel="True"
    RelativePanel.AlignRightWithPanel="True"
    RelativePanel.Below="textBlockCurrentImageDisplayName" />
<TextBlock x:Name="textBlockCurrentImageImageWidth" Margin="10,10,10,0"
    FontSize="24" FontWeight="Bold"
    RelativePanel.AlignLeftWithPanel="True"
    RelativePanel.AlignRightWithPanel="True"
    RelativePanel.Below="textBlockCurrentImageImageHeight" />

```

5. Add the following Visual State Manager to control the appearance when the app is resizing. Add it as the first child to the Grid tag:

```

<VisualStateManager.VisualStateGroups>
    <VisualStateGroup>
        <VisualState x:Name="narrowView">
            <VisualState.StateTriggers>
                <AdaptiveTrigger MinWindowWidth="0" />
            </VisualState.StateTriggers>
            <VisualState.Setters>
                <Setter Target="RightPanel.(RelativePanel.Below)" Value="LeftPanel" />
                <Setter Target="RightPanel.(RelativePanel.AlignLeftWithPanel)"
Value="True" />
                <Setter Target="RightPanel.(RelativePanel.AlignRightWithPanel)"
Value="True" />
                <Setter Target="RightPanel.Margin" Value="0,10,0,0" />
                <Setter Target="LeftPanel.(RelativePanel.AlignTopWithPanel)" Value="True" />
                <Setter Target="LeftPanel.(RelativePanel.AlignLeftWithPanel)" Value="True" />
                <Setter Target="LeftPanel.(RelativePanel.AlignRightWithPanel)"
Value="True" />
                <Setter Target="LeftPanel.Height" Value="560" />
            </VisualState.Setters>
        </VisualState>
        <VisualState x:Name="wideView">
            <VisualState.StateTriggers>
                <AdaptiveTrigger MinWindowWidth="720" />
            </VisualState.StateTriggers>
            <VisualState.Setters>
                <Setter Target="RightPanel.(RelativePanel.AlignBottomWithPanel)"
Value="True" />
                <Setter Target="RightPanel.(RelativePanel.AlignRightWithPanel)"
Value="True" />
                <Setter Target="RightPanel.(RelativePanel.AlignTopWithPanel)" Value="True"
/>
                <Setter Target="RightPanel.Width" Value="200" />
                <Setter Target="LeftPanel.(RelativePanel.LeftOf)" Value="RightPanel" />
                <Setter Target="LeftPanel.(RelativePanel.AlignBottomWithPanel)"
Value="True" />
            </VisualState.Setters>
        </VisualState>
    </VisualStateGroup>

```

```
<Setter Target="LeftPanel.(RelativePanel.AlignTopWithPanel)" Value="True" />
<Setter Target="LeftPanel.(RelativePanel.AlignLeftWithPanel)" Value="True" />
</VisualState.Setters>
</VisualState>
</VisualStateGroup>
</VisualStateManager.VisualStateGroups>
```

6. Create a new class and name it ImageProperties. Add three properties to it like this:

```
namespace PictureViewer
{
    class ImageProperties
    {
        public string FileName { get; set; }
        public int Width { get; set; }
        public int Height { get; set; }
    }
}
```

7. Go to the code-behind for the main page and add these using statements:

```
using System;
using System.Collections.Generic;
using Windows.Storage;
using Windows.UI.Xaml;
using Windows.UI.Xaml.Controls;
using Windows.Storage.Search;
using Windows.UI.Xaml.Media.Imaging;
using Windows.UI.Popups;
using System.Linq;
```

8. Create a private field to hold some information about the pictures that are being displayed:

```
private IList<ImageProperties> imageProperties = new List<ImageProperties>();
```

9. Add a method to load the files:

```
private async void GetFiles()
{
    try
    {
        StorageFolder picturesFolder = KnownFolders.PicturesLibrary;
        IReadOnlyList<StorageFile> sortedItems = await picturesFolder.
GetFilesAsync(CommonFileQuery.OrderByDate);
        var images = new List<BitmapImage>();

        if (sortedItems.Any())
        {
            foreach (StorageFile file in sortedItems)
            {
                if (file.FileType.ToUpper() == ".JPG")
                {
                    using (Windows.Storage.Streams.IRandomAccessStream fileStream = await
file.OpenAsync(FileAccessMode.Read))
                    {
                        BitmapImage bitmapImage = new BitmapImage();

```

```
        await bitmapImage.SetSourceAsync(fileStream);
        images.Add(bitmapImage);
        imageProperties.Add(new ImageProperties
    {
        FileName = file.DisplayName,
        Height = bitmapImage.PixelHeight,
        Width = bitmapImage.PixelWidth
    });
    if (imageProperties.Count > 10)
        break;
}
}
}
else
{
    var message = new MessageDialog("There are no images in the Pictures
Library");
    await message.ShowAsync();
}
flipView.ItemsSource = images;
}
catch (UnauthorizedAccessException)
{
    var message = new MessageDialog("The app does not have access to the Pictures
Library on this device.");
    await message.ShowAsync();
}
}
```

10. Select the Page tag in the XAML of the main page and add the Loading event to it. Then implement this event handler for it:

```
private void Page_Loaded(object sender, RoutedEventArgs e)
{
    GetFiles();
}
```

11. Select the FlipView in the XAML of the Main Page and implement the SelectionChanged event:

```
private void flipView_SelectionChanged(object sender, SelectionChangedEventArgs e)
{
    if (flipView.SelectedIndex >= 0)
    {
        textBlockCurrentImageDisplayName.Text = imageProperties[flipView.
SelectedIndex].FileName;
        textBlockCurrentImageHeight.Text = imageProperties[flipView.
SelectedIndex].Height.ToString();
        textBlockCurrentImageWidth.Text = imageProperties[flipView.
SelectedIndex].Width.ToString();
    }
}
```

12. Finally, double-click the Package.appxmanifest file to open the manifest file designer.

13. Select the Capabilities tab and ensure that the Pictures Library capability is checked.

14. Run the app.

### How It Works

The code uses three `RelativePanels` to move its content about. None of the panels have any positioning instructions directly on them; instead the entire layout is defined in the Visual State Manager. In this case, we are using two adaptive triggers, one that activates if the view is wider than 720 pixels and one that activates if the view is wider than 0 pixels.

The `FlipView` itself is almost the least amount of code in the example.

```
<FlipView x:Name="flipView" VerticalAlignment="Stretch"
HorizontalAlignment="Stretch" >
    <FlipView.ItemTemplate>
        <DataTemplate>
            <Image x:Name="image" Source="{Binding}" Stretch="Uniform"/>
        </DataTemplate>
    </FlipView.ItemTemplate>
</FlipView>
```

In this code, we tell the `FlipView` that it should use the `ItemTemplate` we define here, and this just includes a single `Image` control. From this it is apparent that you can use the `FlipView` to display anything, not just images.

The code in the `GetFiles` method demonstrates several of the interfaces that you can use to access files and resources in an app. Later in this chapter we will discuss the concept of sandboxed apps and what limitations they put on your code, but you have already seen some of this in action in this example. The following code gets a `StorageFolder` object if the app has access to it and throws a `UnauthorizedAccessException` exception if it does not.

```
StorageFolder picturesFolder = KnownFolders.PicturesLibrary;
```

In normal .NET you don't have this class, and the determination of whether access is granted is based on the user's permissions in the file system. For apps, this is very different. Here you have to declare up front which resources the app will need access to, and the user must accept these for the application to run. In step 13 you declared that the app will include the capability to access to the Pictures Library. If you didn't do that, you will get an exception when you run the app.

Next you used the `GetFilesAsync` method of the `storageFolder` to retrieve the files, ordered by date.

Once we have the files, we open them by calling `OpenAsync` on the `StorageFile` objects

```
using (Windows.Storage.Streams.IRandomAccessStream fileStream = await file.
OpenAsync(FileAccessMode.Read))
```

This returns a file stream we can use to access the content of the files. In this case, we don't want to write to it, so we specify Read access only.

Finally, we set the `ItemsSource` of the `FlipView` to the list of images we have loaded from the Pictures Library.

---

## Sandboxed Apps

At this point, it is worth taking a step back and looking at some of the limitations of the .NET framework for the Windows Universal Platform. Apps running on mobile devices have limited access to the OS on which they run, and this means that there are types of applications that you simply cannot write. If you require direct access to the file system to access Windows system files, for instance, you must write a classic Windows desktop application.

When you are writing Universal apps in C#, you will find that the limiting factor is in the .NET Framework that is referenced from your application, where common namespaces and classes are missing entirely or have fewer methods available than before. If you open Visual Studio, create a new Blank app, and then expand the References node, you will see that the references are very different from those for Windows Desktop Apps. There are three ApplicationInsights references, each of which allows you to monitor various aspects of your app, and two references to .NET and Windows. The first of the latter is a changed version of .NET and the second is the Windows Core API. At this point you might expect that you could simply change the references to use the normal .NET Framework, and indeed this will work. That is, it will work right up to the point when you try to publish your app to the Windows Store, at which point it will be rejected for non-compliance with the specifications.

The sandboxed nature of the Windows Universal Apps, and the process they must go through before they are admitted into the Windows Store, means that the users should rarely have to fear downloading malicious apps through the store. Obviously, there are people who will try to circumvent this, and users should never let their guard down; however, it is considerably harder to place malicious programs on Windows computers through Windows Store apps than it is through normal download and installation.

## Disk Access

Desktop applications can access the disk pretty much as they like, with a few exceptions. One such exception is that they are normally prohibited from writing to the Program Files folder and other system folders. Windows Universal Apps can access only a few very specific locations on disk directly. These locations include the folder in which the app is installed, the AppData folder associated with the app, and a few special folders such as the Documents folder. Access to the files and folders have also been moved in the .NET Framework for Universal apps to make sure that the developer can't accidentally write to a forbidden location.

In order to allow the user control over where files should be stored and read from in your app, Windows provides you with three File Picker contracts: `FolderOpenPicker`, `FileOpenPicker`, and `FileSavePicker`. These picker classes can be used from your app to gain secure access to the local disk.

As you saw earlier, you can also use the `KnownFolders` class to access resources on a device. You should use the `KnownFolders` class when you want to read or write to locations that the user must grant access to for the app to be able to open them.

## Serialization, Streams, and Async Programming

In Chapter 14, you used the `[Serializable]` attribute to allow classes to be serialized. .NET for Universal apps do not include this attribute, but you can use a similar attribute called `[DataContract]` instead. The `DataContract` attribute works with the `DataContractSerializer` class to serialize the content of a class. In order to get the serialized content to or from disk, you need to use some file access types, but unlike with normal .NET, you can't create these directly. Instead, you use file pickers to create the stream objects, which you can use with `DataContractSerializer` to save and load your files.

**NOTE** The projects you can download for this chapter from [www.wrox.com/go/beginningvisualc#2015programming](http://www.wrox.com/go/beginningvisualc#2015programming) include a certificate file that you may not be able to use, but can generate yourself. Follow these steps to do so:

1. With the project open, double-click the file `Package.appxmanifest`.
2. Select the Packaging tab.
3. Click Choose Certificate.
4. Select Create test certificate from the Configure Certificate.
5. Click OK.

The next Try It Out demonstrates using `DataContractSerializer` with streams created by `FileOpenPicker` and `FileSavePicker` to load and save XML representations of a data model.

### TRY IT OUT Disk Access: Ch23Ex03

1. Create a new project in Visual Studio by selecting Blank App (Universal Windows) and name it `DataSerialization`.
2. Create a new class in the project named `AppData`.
3. Mark the class with the `[DataContract]` attribute and add the `System.Runtime.Serialization` namespace to the using section:

```
using System.Runtime.Serialization;

namespace DataSerialization
{
    [DataContract]
    class AppData
    {
    }
}
```

4. Add a property of type `int` to the class and mark it with the `[DataMember]` attribute:

```
[DataMember]
public int TheAnswer { get; set; }
```

5. Add a new enum to the project called `AppStates`. Mark it with the `[DataContract]` attribute:

```
using System.Runtime.Serialization;
namespace DataSerialization
{
    [DataContract]
    public enum AppStates
    {
        [EnumMember]
        Started,
        [EnumMember]
        Suspended,
        [EnumMember]
        Closing
    }
}
```

6. Add three values to `AppStates`, taking care to mark each one with the `[EnumMember]` attribute:

```
[EnumMember]
Started,
[EnumMember]
Suspended,
[EnumMember]
Closing
```

7. Add two new properties to the `AppData` class:

```
[DataMember]
public AppStates State { get; set; }
[DataMember]
public object StateData { get; set; }
```

8. Add a new class with the name `AppStateData` and mark it with the `[DataContract]` attribute:

```
using System.Runtime.Serialization;
namespace DataSerialization
{
    [DataContract]
    public class AppStateData
    {
        [DataMember]
        public string Data { get; set; }
    }
}
```

9. Add a `[KnownType]` attribute to the `AppData` class like this:

```
[DataContract]
[KnownType(typeof(AppStateData))]
public class AppData
{
```

10. Double-click the `MainPage.xaml` file in the Solution Explorer and drag two buttons onto the page. Set their content and name properties to `Save` and `Load`.

11. Create a `click` event handler for the `Save` button and navigate to it in the code-behind file. Add this code (note the `async` keyword in the method declaration):

```
private async void Save_Click(object sender, RoutedEventArgs e)
{
    var data = new AppData
    {
        State = AppStates.Started,
```

```

        TheAnswer = 42,
        StateData = new AppStateData { Data = "The data is being saved" }
    };
    var fileSavePicker = new FileSavePicker
    {
        SuggestedStartLocation = PickerLocationId.DocumentsLibrary,
        DefaultFileExtension = ".xml",
    };
    fileSavePicker.FileTypeChoices.Add("XML file", new[] { ".xml" });
    var file = await fileSavePicker.PickSaveFileAsync();
    if (file != null)
    {
        var stream = await file.OpenStreamForWriteAsync();
        var serializer = new DataContractSerializer(typeof(AppData));
        serializer.WriteObject(stream, data);
    }
}

```

12. Create the `click` event handler for the Load button and add this code (note the `async` keyword again):

```

private async void Load_Click(object sender, RoutedEventArgs e)
{
    var fileOpenPicker = new FileOpenPicker
    {
        SuggestedStartLocation = PickerLocationId.DocumentsLibrary,
        ViewMode = PickerViewMode.Thumbnail
    };
    fileOpenPicker.FileTypeFilter.Add(".xml");
    var file = await fileOpenPicker.PickSingleFileAsync();
    if (file != null)
    {
        var stream = await file.OpenStreamForReadAsync();
        var serializer = new DataContractSerializer(typeof(AppData));
        var data = serializer.ReadObject(stream);
    }
}

```

13. You will need to add these two namespaces to the code-behind file:

```

using System.Runtime.Serialization;
using Windows.Storage.Pickers;

```

14. Run the app.

## How It Works

In Steps 1 through 9, you create the data model of the app. All classes and enumerations are marked with the `[DataContract]` attribute, but notice the difference in how members are marked. Properties and fields in classes can be marked with the `[DataMember]` attribute, but members of an enumeration must be marked with `[EnumMember]`:

```

[DataContract]
public class AppStateData
{
    [DataMember]

```

```

        public string Data { get; set; }
    }
    [DataContract]
    public enum AppStates
    {
        [EnumMember]
        Started,
        [EnumMember]
        Suspended,
        [EnumMember]
        Closing
    }
}

```

There is another attribute that is not shown here that can be of interest: `CollectionDataContract`. It can be set on custom collections.

You also add a property with a `Type` object. In order for the serializer to be able to serialize this property, you must tell it what types it could be. You do this by setting the `[KnownTypes]` attribute on the class that contains the property.

The `Save` and `Load` methods demonstrate some of the new file pickers. After displaying the pickers, you get a `StorageFile` instance back:

```

var file = await fileOpenPicker.PickSingleFileAsync();
if (file != null)
{
    var stream = await file.OpenStreamForReadAsync();
    var serializer = new DataContractSerializer(typeof(AppData));
    var data = serializer.ReadObject(stream);
}

```

This object can be used to open a stream for read or write operations. It is not shown directly here, but you can also use it directly with the `FileIO` class, which provides some simple methods for writing and reading data.

## Navigation between Pages

Navigating between pages within an app is similar to how web applications navigate. You can call the method `Navigate` to go from one page to another; you can go back by calling the `Back` method. The following Try It Out demonstrates how to move between pages in an app using three basic pages.

### TRY IT OUT    Navigation: Ch23Ex04

1. Create a new project in Visual Studio by selecting Blank App (Universal Windows) and name it `BasicNavigation`.
2. Select and delete the `MainPage.xaml` file.
3. Right-click the project and select Add ➔ New item. Add a new page using the Blank Page template and name it `BlankPage1`.

4. Repeat Step 3 twice so you have a total of three pages in the project, naming the pages `BlankPage2` and `BlankPage3` respectively.
5. Open the `App.xaml.cs` code-behind file and locate the `OnLaunched` method. This method uses the `MainPage` that you just deleted, so change the reference to `BlankPage1` instead.
6. On the `BlankPage1`, insert a stack panel, a `TextBlock`, and three buttons into the grid:

```
<Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
<TextBlock x:Name="textBlockCaption" Text="Page 1" HorizontalAlignment="Center"
Margin="10" VerticalAlignment="Top"/>
<StackPanel Orientation="Horizontal" Grid.Row="1" HorizontalAlignment="Center">
    <Button Content="Page 2" Click="buttonGoto2_Click" />
    <Button Content="Page 3" Click="buttonGoto3_Click" />
    <Button Content="Back" Click="buttonGoBack_Click" />
</StackPanel>
</Grid>
```

7. Add the event handlers for the click events like this:

```
private void buttonGoto2_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage2));
}
private void buttonGoto3_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage3));
}
private void buttonGoBack_Click(object sender, RoutedEventArgs e)
{
    if (Frame.CanGoBack) this.Frame.GoBack();
```

8. Open the second page (`BlankPage2`) and add a similar stack panel to it:

```
<TextBlock x:Name="textBlockCaption" Text="Page 2" HorizontalAlignment="Center"
Margin="10" VerticalAlignment="Top"/>
<StackPanel Orientation="Horizontal" Grid.Row="1" HorizontalAlignment="Center">
    <Button Content="Page 1" Click="buttonGoto1_Click" />
    <Button Content="Page 3" Click="buttonGoto3_Click" />
    <Button Content="Back" Click="buttonGoBack_Click" />
</StackPanel>
```

9. Add the navigation to the event handlers:

```
private void buttonGoto1_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage1));
}
private void buttonGoto3_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage3));
}
private void buttonGoBack_Click(object sender, RoutedEventArgs e)
{
    if (Frame.CanGoBack) this.Frame.GoBack();
```

10. Open the third page and add another stack panel that includes a Home button:

```
<TextBlock x:Name="textBlockCaption" Text="Page 3" HorizontalAlignment="Center" Margin="10" VerticalAlignment="Top"/>
<StackPanel Orientation="Horizontal" Grid.Row="1" HorizontalAlignment="Center">
    <Button Content="Page 1" Click="buttonGoto1_Click" />
    <Button Content="Page 2" Click="buttonGoto2_Click" />
    <Button Content="Back" Click="buttonGoBack_Click" />
</StackPanel>
```

11. Add the event handlers:

```
private void buttonGoto1_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage1));
}
private void buttonGoto2_Click(object sender, RoutedEventArgs e)
{
    Frame.Navigate(typeof(BlankPage2));
}
private void buttonGoBack_Click(object sender, RoutedEventArgs e)
{
    if (Frame.CanGoBack) this.Frame.GoBack();
}
```

12. Run the app. The app displays the front page with three buttons.

### *How It Works*

When you run the application, it displays a splash screen when loading and then displays the first page. The first time you click one of the buttons, the `Navigate` method is called using the type of the page you want to navigate to.

```
Frame.Navigate(typeof(BlankPage2));
```

It is not shown in this example, but the `Navigate` method includes an overload that allows you to send parameters to the page that is being navigated to. When you navigate between the pages, you will notice that if you go back to Page 1 using one of the buttons, the Back button remains active.

On each page, you use the `GoBack` event implementation to go back to the previous page. Before the `GoBack` method is called, the `CanGoBack` property is checked. If you fail to do so and call `GoBack` on the first page displayed, you will get an exception.

```
if (Frame.CanGoBack) this.Frame.GoBack();
```

Each time you navigate to a page, a new instance is created. You can change this behavior by enabling the property `NavigationCacheMode` in the constructor of your pages; for example, like this:

```
public BasicPage1()
{
    this.InitializeComponent();
    NavigationCacheMode = Windows.UI.Xaml.Navigation.NavigationCacheMode.Enabled;
}
```

This will cause the page to become cached.

## The CommandBar Control

A `CommandBar` provides the users with much the same functionality that a tool bar provides in desktop applications, but you should keep them much simpler, usually limiting the available options to fewer than eight items in a bar.

You can display more than one `CommandBar` at a time, but keep in mind that this clutters up the user interface, and you should not display more than one bar just to show more options. On the other hand, if you want to provide more than one kind of navigation, it is sometimes beneficial to show a top and bottom bar at the same time.

Visual Studio ships with the `CommandBar` control, which makes it very easy to create this kind of control. The following Try It Out creates an App Bar with a number of standard items on it.

### TRY IT OUT Creating CommandBars: Ch23Ex05

1. Return to the `BasicNavigation` example from earlier.
2. Add a `CommandBar` to all three pages. Place it as a child of the grid control on each of the pages:

```
<CommandBar>
    <AppBarToggleButton x:Name="toggleButtonBold" Icon="Bold" Label="Bold"
Click="AppBarToggleButtonBold_Click" />
    <AppBarSeparator />
    <AppBarButton Icon="Back" Label="Back" Click="buttonGoBack_Click"/>
    <AppBarButton Icon="Forward" Label="Forward" Click="AppBarButtonForward_Click"/>

    <CommandBar.SecondaryCommands>
        <AppBarButton Icon="Camera" Label="Take picture" />
        <AppBarButton Icon="Help" Label="Help" />
    </CommandBar.SecondaryCommands>
</CommandBar>
```

3. Add this event handler to all three pages:

```
private void AppBarButtonForward_Click(object sender, RoutedEventArgs e)
{
    if (Frame.CanGoForward) this.Frame.GoForward();
}
private void AppBarToggleButtonBold_Click(object sender, RoutedEventArgs e)
{
    AppBarToggleButton toggleButton = sender as AppBarToggleButton;
    bool isChecked = toggleButton.IsChecked.HasValue ?
        (bool)toggleButton.IsChecked.Value : false;
    textBlockCaption.FontWeight = isChecked ? FontWeights.Bold : FontWeights.Normal;
}
```

4. Add this using statement to all pages:

```
using Windows.UI.Text;
```

5. On all three pages, change the `margin` of the `TextBox` to `10,50,10,10`.
6. Run the app.

## How It Works

When you run the app, you can now use the command bar buttons to move back and forth in the list of pages that you have visited. The command bar itself is very easy to work with.

The command bar is built using three types. The first one is the `AppBarToggleButton`.

```
<AppBarToggleButton x:Name="toggleButtonBold" Icon="Bold" Label="Bold"  
Click="AppBarToggleButtonBold_Click" />
```

This type of button can be used to display a state that can be toggled on or off.

The second type is the `AppBarButton`, which works like any other button, and in fact you can see that the click event of the `AppBarButtonBack` button is handled by the same event handler as the `ButtonBack` from the previous example.

```
<AppBarButton Icon="Back" Label="Back" Click="buttonGoBack_Click"/>
```

The third type that is used in the command bar is the `AppBarSeparator`. This control simply displays a separator on the bar.

Finally, two buttons are located inside a `CommandBar.SecondaryCommands` tag:

```
<CommandBar.SecondaryCommands>  
    <AppBarButton Icon="Camera" Label="Take picture" />  
    <AppBarButton Icon="Help" Label="Help" />  
</CommandBar.SecondaryCommands>  
</CommandBar>
```

These commands are not displayed directly on the command bar. Instead they are displayed as a drop-down when you click the three dots that are displayed.

## Managing State

Unlike a desktop application, an app must expect to be suspended at any time. This happens when the user switches to another app or to the desktop, so it's a very common scenario that must be handled by all apps. When an app is suspended, Windows will save the values of your variables and data structures and restore them when the app resumes. However, your app may have been suspended for an extended period of time, so if you have data that changes over time, such as a news feed, then you should refresh this when the app is restored.

When the app is suspended, you should also consider saving any data that should persist between invocations of the app, as you will not get a chance to do so if the app is subsequently terminated by Windows or the user.

When your app is about to be suspended, a `Suspending` event is sent, which you should handle. When the app is returned to life, it will receive a `Resuming` event. By handling these two events and saving the state of the application, you can return the app to the state it was in before the suspension, and the user shouldn't notice anything.

**TRY IT OUT** Resume from Suspension: Ch23Ex06

1. Return to the previous example and create a new class named AppState:

```
using System.Collections.Generic;

namespace BasicNavigation
{
    public static class AppState
    {
        private static Dictionary<string, bool> state = new Dictionary<string, bool>();

        public static bool GetState(string pageName) => state.ContainsKey(pageName) ? state[pageName] : false;

        public static void SetState(string pageName, bool isBold)
        {
            if (state.ContainsKey(pageName))
                state[pageName] = isBold;
            else
                state.Add(pageName, isBold);
        }

        public static void Save()
        {
            var settings = Windows.Storage.ApplicationData.Current.RoamingSettings;
            foreach (var key in state.Keys)
            {
                settings.Values[key] = state[key];
            }
        }

        public static void Load(string pageName)
        {
            if (!state.ContainsKey(pageName) && Windows.Storage.ApplicationData.Current.RoamingSettings.Values.ContainsKey(pageName))
                state.Add(pageName, (bool)Windows.Storage.ApplicationData.Current.RoamingSettings.Values[pageName]);
        }
    }
}
```

2. Open the code-behind of the app.xaml file, and locate the OnSuspending method at the bottom. Add AppState.Save(); like this:

```
private void OnSuspending(object sender, SuspendingEventArgs e)
{
    var deferral = e.SuspendingOperation.GetDeferral();
    //TODO: Save application state and stop any background activity
    AppState.Save();
    deferral.Complete();
}
```

3. Add these lines to the bottom of the OnLaunched method, just above Window.Current.Activate();:

```
AppState.Load(typeof(BlankPage1).Name);
AppState.Load(typeof(BlankPage2).Name);
AppState.Load(typeof(BlankPage3).Name);
```

4. Go to the BlankPage1 and add the loaded event on Page class like this:

```
<Page
    x:Class="BasicNavigation.BlankPage1"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:BasicNavigation"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d" Loaded="Page_Loaded">
```

5. Implement the event:

```
private void Page_Loaded(object sender, RoutedEventArgs e)
{
    toggleButtonBold.IsChecked = AppState.GetState(GetType().Name);
    AppBarToggleButtonBold_Click(toggleButtonBold, new RoutedEventArgs());
}
```

6. Change the Click event handler for the toggle button to save the state of the page when the button is pressed:

```
private void AppBarToggleButtonBold_Click(object sender, RoutedEventArgs e)
{
    AppState.SetState(GetType().Name, (bool)toggleButtonBold.IsChecked);
    ...
}
```

7. Repeat steps 4 through 6 for BlankPage2 and BlankPage3.
8. Set a break point in the OnSuspending method in the app.xaml code-behind file.
9. Run the app.
10. Once the app is running, click the Bold button on one or two of the pages. Then, while the app is still running, return to Visual Studio. Notice that there is a Debug Location toolbar being displayed, on which you will see a drop-down with the text Lifecycle Events. Expand this and click Suspend.
11. Once you step through the OnSuspended method, the app is suspended. Expand the drop-down again and click Resume.

## How It Works

The `AppState` class uses the `Windows.Storage.ApplicationData` class to save the application settings. This class allows you to access the application data store and lets you easily set some simple values. You should only store simple types in this store, so if you need to save very complex state for the app you should consider some other mechanism, like a database or XML files.

The app already handles the Suspending event in the `app.xaml` code-behind file, so you can simply use this. If you had to handle the suspension differently for individual pages, you should handle this event on the pages themselves as well.

In the `OnSuspending` event, we save the state of the entire app so that it can be retrieved when the app restarts. Since there's no data that must be updated on any of the pages when the app resumes from suspension, we don't handle the `Resuming` event.

We restore the state when the app loads by loading it in the `OnLaunched` method, also in the `app.xaml` code-behind.

---

## COMMON ELEMENTS OF WINDOWS STORE APPS

All Windows Store apps should provide their own Tiles and Badges. Tiles give your app presence on the Start page in Windows and allow you to display information about the app. Badges allow Windows to display a small image that represents your app on the Lock Screen.

Tiles are important because users tend to be fickle and are inclined to make decisions based on how well an app presents itself. Also, a tile should be easily identifiable; if you make your users search for a tile that disappears in the other tiles, they're unlikely to be in a good mood by the time that they finally locate it.

There are many possible tile sizes in Windows Store apps and if your app is targeting many different display sizes, you should supply tailored tiles for all the suggested sizes, or at the very least provide tiles that scale well. If you don't provide a tile of the right size, Windows will scale what you do supply to the correct size and this will often look bad. So, for a professional app, make sure you tailor make tiles for every expected tile size.

Badges are smaller than the tiles ( $24 \times 24$  pixels) and are used when Windows displays the app on the Lock Screen. If you set a badge image for your app, you must also enable Lock Screen notifications. Badges can also be scaled, so provide all the appropriate sizes.

Splash screens are displayed while the app loads; and since that should take only a moment or two, they should not be too complex or provide any kind of information to the users, except for clearly identifying which app is currently starting. Splash screens are exactly  $620 \times 300$  pixels, but you can make them smaller by making parts of the image transparent. And once again, scaled versions should be supplied.

Finally, you should supply a "Store Logo" of exactly  $50 \times 50$  pixels and of course the scaled versions of it.

Tiles, badges, and logos are embedded in the apps package manifest, which can be edited easily in the Visual Studio Manifest Package editor. If you have downloaded the code for this book, you can use the tiles and badge supplied with the code (in the `Assets` folder), but otherwise you can quickly create the images in Paint or in a similar application.

### TRY IT OUT Adding Tiles and Badges

1. Use an image editor like Paint to create PNG images with these sizes:
  - $620 \times 300$

- 310 × 150
- 310 × 310
- 150 × 150
- 71 × 71
- 50 × 50
- 44 × 44
- 24 × 24

Name the images so you can recognize them without opening them.

2. Open the project from the previous example.
3. Double-click the file `Package.appxmanifest` in the Solution Explorer to open the package editor.
4. Below the Visual Assets heading, you will find a menu on the left where you can change the tiles, logos, and splash screens. Add the images here by clicking the buttons for the scale 100 images and browsing to them.
5. Right-click the project in the Solution Explorer and select Deploy.

### ***How It Works***

Go to the Start menu and find the app. You will probably have to click All Apps or search for the name. Notice that the small tile is displayed in the list. If you right-click it and pin the app to the Start menu, one of the larger tiles is used. You can right-click the tile and select Resize to change the size.

When the app runs, the splash screen briefly appears.

Right-click the app in the menu and select Uninstall to remove it again.

---

## **THE WINDOWS STORE**

After you create your app, you will probably want to distribute it to the public, and the way to do this is to use the Windows Store. Microsoft has gone to great lengths to create a store that is secure and lets Windows users download apps from it without too much fear of downloading malicious code. Unfortunately, this means you must endure a lengthy process to get your app in the store.

## **Packaging an App**

You have already seen some of the contents of the `package.appxmanifest` file when you had to specify the Picture Viewer required access to the Pictures Library and when adding Tiles to the app. When you are ready to package your app for the App Store, you must return to this file and set a number of other values.

Before you package your app, you should go through each of the six tabs for configuring the `package.appxmanifest` and consider every option you have:

- **Application:** Name your app well! Along with the store logo, this is probably the very first thing your potential users see about your app, so naming it something generic is not effective. Try to pick an interesting name that also indicates the purpose of the app.
- **Visual Assets:** In the last example you added tiles to the app. You should ensure that there is at least one image for every category on the Visual Assets tab.
- **Capabilities:** On this tab you specify which capabilities your app requires. Be warned that users will view your app suspiciously if it requires capabilities that don't appear reasonable. For instance, if you require access to the chat messages on the device, there had better be a good reason; otherwise it is likely that this will be seen as a potential breach of privacy. Most apps shouldn't require more than a few capabilities, but you must pick all that you use. If you don't accurately specify what you need, then the app will receive an access denied exception when it tries to access the resource.
- **Declarations:** On the Declarations tab, you can register the app as a provider of services. For instance, if your app works as a search provider, then you can add this declaration to the app and specify the required properties.
- **Content URIs:** If your app navigates to a remote page, it will have limited access to the system. You can use Content URIs to give a web page access to geo-location devices and the clipboard.
- **Packaging:** On this tab you can set the properties of the package, including the name of the developer/publisher, the version of the app, and the certificate used to sign the package.

## Creating the Package

Once you have specified all you need in the `appxmanifest`, you are ready to package your app. You can do this directly from Visual Studio by selecting **Store** ➔ **Create App Packages**. This will launch the Create App Packages wizard. A few steps into the wizard you will be required to log in with a store account. If you don't have one, you must create one. You must have a store account to be able to publish to the app store and to be able to get paid for your app.

At some point during the wizard, you will be shown the Select and Configure Packages page. On this, it is important to select all three of the target architectures (x86, x64, and ARM) to allow the app to be deployed to the widest range of devices.

On the final page you will be given options on how to validate that your app can be submitted to the app store. Launch the Windows App Certification Kit and learn if your app is ready to be submitted. If any problems are detected, you must fix them and go through the Create App Packages wizard again. If your app passes inspection, you can upload the package.

**EXERCISES**

- 
- 23.1 Extend the Ch23Ex06 example by adding a WebView control to the page BlankPage1 and use the navigate method to show a web page of your choice. Add an event handler to the page that will navigate the webView to another web page when the app resumes from suspension.
- 
- 23.2 If you want your app to work as a voice recorder, you must ensure that the app has access to the microphone on the device. How do you ensure that the app will not get an UnauthorizedAccessException when it tries to use the microphone on the device?
- 
- 23.3 Many apps running on Windows Phone use a style of navigation known as Pivot. You can create Universal apps that use this style as well. Create an app that uses the Pivot control to display three views, one displaying a web page, another displaying the text "Hello Pivot!," and the third showing the Wrox logo. You can find the logo here: [http://media.wiley.com/assets/253/59/wrox\\_logo.gif](http://media.wiley.com/assets/253/59/wrox_logo.gif)
-

## ► WHAT YOU LEARNED IN THIS CHAPTER

KEY CONCEPT	DESCRIPTION
<b>Windows Universal App XAML</b>	Windows Universal app XAML is used with C# to create the GUI for Windows Universal apps. It includes many of the same controls that you know from WPF, but some have changed, others are missing, and new controls have been introduced.
<b>Visual State manager</b>	You saw how to use a Visual State manager to change the look of your controls and pages simply by changing the visual state of the control. This leads to a lot less code in exchange for slightly more complex XAML.
<b>App State</b>	Windows Universal apps are suspended when the user switches to another app or to the desktop, so it's important to handle this suspension and save the app state when it happens.
<b>App store account</b>	This account is used for deploying apps to the Windows Store.
<b>Navigation</b>	Navigation in Windows Universal apps is done in much the same way that it is in web applications, using method calls to move back and forth in the page structure.

# APPENDIX

# Exercise Solutions

There are no exercises in Chapters 1 and 2.

## CHAPTER 3 SOLUTIONS

### Exercise 1

```
super.smashing.great
```

### Exercise 2

b), as it starts with a number, and e), as it contains a full stop.

### Exercise 3

No, there is no theoretical limit to the size of a string that may be contained in a `string` variable.

### Exercise 4

The `*` and `/` operators have the highest precedence here, followed by `+`, `%`, and finally `+=`. The precedence in the exercise can be illustrated using parentheses as follows:

```
resultVar += (((var1 * var2) + var3) % (var4 / var5));
```

### Exercise 5

```
using static System.Console;
using static System.Convert;
static void Main(string[] args)
{
    int firstNumber, secondNumber, thirdNumber, fourthNumber;
    WriteLine("Give me a number:");
    firstNumber =ToInt32(ReadLine());
```

```
        WriteLine("Give me another number:");
        secondNumber =ToInt32(Console.ReadLine());
        WriteLine("Give me another number:");
        thirdNumber =ToInt32(ReadLine());
        WriteLine("Give me another number:");
        fourthNumber =ToInt32(ReadLine());
        WriteLine($"The product of {firstNumber}, {secondNumber}, " +
            $"{thirdNumber}, and {fourthNumber} is " +
            $"{firstNumber * secondNumber * thirdNumber * fourthNumber}.");
    }
}
```

Note that `Convert.ToInt32()` is used here, which isn't covered in the chapter.

## CHAPTER 4 SOLUTIONS

### Exercise 1

```
(var1 > 10) ^ (var2 > 10)
```

### Exercise 2

```
using static System.Console;
using static System.Convert;
static void Main(string[] args)
{
    bool numbersOK = false;
    double var1, var2;
    var1 = 0;
    var2 = 0;
    while (!numbersOK)
    {
        WriteLine("Give me a number:");
        var1 = ToDouble(ReadLine());
        WriteLine("Give me another number:");
        var2 = ToDouble(ReadLine());
        if ((var1 > 10) && (var2 > 10))
        {
            numbersOK = true;
        }
        else
        {
            if ((var1 <= 10) && (var2 <= 10))
            {
                numbersOK = true;
            }
            else
            {
                WriteLine("Only one number may be greater than 10.");
            }
        }
    }
    WriteLine($"You entered {var1} and {var2}.");
}
```

Note that this can be performed better using different logic, for example:

```
static void Main(string[] args)
{
    bool numbersOK = false;
    double var1, var2;
    var1 = 0;
    var2 = 0;
    while (!numbersOK)
    {
        WriteLine("Give me a number:");
        var1 = ToDouble(ReadLine());
        WriteLine("Give me another number:");
        var2 = Convert.ToDouble(ReadLine());
        if ((var1 > 10) && (var2 > 10))
        {
            WriteLine("Only one number may be greater than 10.");
        }
        else
        {
            numbersOK = true;
        }
    }
    WriteLine($"You entered {var1} and {var2}.");
}
```

## Exercise 3

The code should read:

```
int i;
for (i = 1; i <= 10; i++)
{
    if ((i % 2) == 0)
        continue;
    WriteLine(i);
}
```

Using the = assignment operator instead of the Boolean == operator is a very common mistake.

# CHAPTER 5 SOLUTIONS

## Exercise 1

Conversions a and c can't be performed implicitly.

## Exercise 2

```
enum color : short
{
    Red, Orange, Yellow, Green, Blue, Indigo, Violet, Black, White
}
```

Yes, as the `byte` type can hold numbers between 0 and 255, so byte-based enumerations can hold 256 entries with individual values, or more if duplicate values are used for entries.

## Exercise 3

The code will not compile, for the following reasons:

- End of statement semicolons are missing.
- Second line attempts to access a nonexistent sixth element of `blab`.
- Second line attempts to assign a string that isn't enclosed in double quotes.

## Exercise 4

```
using static System.Console;
static void Main(string[] args)
{
    WriteLine("Enter a string:");
    string myString = ReadLine();
    string reversedString = "";
    for (int index = myString.Length - 1; index >= 0; index--)
    {
        reversedString += myString[index];
    }
    WriteLine($"Reversed: {reversedString}");
}
```

## Exercise 5

```
using static System.Console;
static void Main(string[] args)
{
    WriteLine("Enter a string:");
    string myString = ReadLine();
    myString = myString.Replace("no", "yes");
    WriteLine($"Replaced \"no\" with \"yes\": {myString}");
}
```

## Exercise 6

```
using static System.Console;
static void Main(string[] args)
{
    WriteLine("Enter a string:");
    string myString = ReadLine();
    myString = "\"" + myString.Replace(" ", "\\" + "\"") + "\"";
    WriteLine($"Added double quotes around words: {myString}");
}
```

Or using `String.Split()`:

```
using static System.Console;
static void Main(string[] args)
```

```

    {
        WriteLine("Enter a string:");
        string myString = ReadLine();
        string[] myWords = myString.Split(' ');
        WriteLine("Adding double quotes around words:");
        foreach (string myWord in myWords)
        {
            Write($"\"{myWord}\\" ");
        }
    }
}

```

## CHAPTER 6 SOLUTIONS

### Exercise 1

The first function has a return type of `bool`, but doesn't return a `bool` value.

The second function has a `params` argument, but this argument isn't at the end of the argument list.

### Exercise 2

```

using static System.Console;
static void Main(string[] args)
{
    if (args.Length != 2)
    {
        WriteLine("Two arguments required.");
        return;
    }
    string param1 = args[0];
    int param2 =ToInt32(args[1]);
    WriteLine($"String parameter: {param1},");
    WriteLine($"Integer parameter: {param2},");
}

```

Note that this answer contains code that checks that two arguments have been supplied, which wasn't part of the question but seems logical in this situation.

### Exercise 3

```

class Program
{
    using static System.Console;
    delegate string ReadLineDelegate();
    static void Main(string[] args)
    {
        ReadLineDelegate readLine = new ReadLineDelegate(ReadLine);
        WriteLine("Type a string:");
        string userInput = readLine();
        WriteLine($"You typed: {userInput}");
    }
}

```

## Exercise 4

```
struct order
{
    public string itemName;
    public int unitCount;
    public double unitCost;
    public double TotalCost() => unitCount * unitCost;}
```

## Exercise 5

```
struct order
{
    public string itemName;
    public int unitCount;
    public double unitCost;
    public double TotalCost() => unitCount * unitCost;
    public string Info() => "Order information: " + unitCount.ToString() +
        " " + itemName + " items at $" + unitCost.ToString() +
        " each, total cost $" + TotalCost().ToString();
}
```

# CHAPTER 7 SOLUTIONS

## Exercise 1

This statement is true only for information that you want to make available in all builds. More often, you will want debugging information to be written out only when debug builds are used. In this situation, the `Debug.WriteLine()` version is preferable.

Using the `Debug.WriteLine()` version also has the advantage that it will not be compiled into release builds, thus reducing the size of the resultant code.

## Exercise 2

```
static void Main(string[] args)
{
    for (int i = 1; i < 10000; i++)
    {
        WriteLine($"Loop cycle {i}");
        if (i == 5000)
        {
            WriteLine(args[999]);
        }
    }
}
```

In VS, you can place a breakpoint on the following line:

```
WriteLine("Loop cycle {0}", i);
```

The properties of the breakpoint should be modified such that the hit count criterion is “break when hit count is equal to 5000”.

## Exercise 3

False. `finally` blocks always execute. This may occur after a `catch` block has been processed.

## Exercise 4

```
static void Main(string[] args)
{
    Orientation myDirection;
    for (byte myByte = 2; myByte < 10; myByte++)
    {
        try
        {
            myDirection = checked((Orientation)myByte);
            if ((myDirection < Orientation.North) ||
                (myDirection > Orientation.West))
            {
                throw new ArgumentOutOfRangeException("myByte",
                    "Value must be between 1 and 4");
            }
        }
        catch (ArgumentOutOfRangeException e)
        {
            // If this section is reached then myByte < 1 or myByte > 4.
            WriteLine(e.Message);
            WriteLine("Assigning default value, Orientation.North.");
            myDirection = Orientation.North;
        }
        WriteLine($"myDirection = {myDirection}");
    }
}
```

Note that this is a bit of a trick question. Because the enumeration is based on the `byte` type, any `byte` value may be assigned to it, even if that value isn't assigned a name in the enumeration. In the previous code, you can generate your own exception if necessary.

# CHAPTER 8 SOLUTIONS

## Exercise 1

B, d, and e. Public, private, and protected are all real levels of accessibility.

## Exercise 2

False. You should never call the destructor of an object manually; the .NET runtime environment will do this for you during garbage collection.

## Exercise 3

No, you can call static methods without any class instances.

## Exercise 4




FIGURE A-1

## Exercise 5

```

static void ManipulateDrink(HotDrink drink)
{
    drink.AddMilk();
    drink.Drink();
    ICup cupInterface = (ICup)drink;
    cupInterface.Wash();
}
  
```

Note the explicit cast to `ICup`. This is necessary as `HotDrink` doesn't support the `ICup` interface, but you know that the two cup objects that might be passed to this function do. However, this is dangerous, as other classes deriving from `HotDrink` are possible, which might not support `ICup`, but could be passed to this function. To correct this, you should check to see if the interface is supported:

```

static void ManipulateDrink(HotDrink drink)
{
    drink.AddMilk();
    drink.Drink();
    if (drink is ICup)
    {
        ICup cupInterface = drink as ICup;
        cupInterface.Wash();
    }
}
  
```

The `is` and `as` operators used here are covered in Chapter 11.

## CHAPTER 9 SOLUTIONS

### Exercise 1

myDerivedClass derives from MyClass, but MyClass is sealed and can't be derived from.

### Exercise 2

You can define a noncreatable class by defining it as a static class or by defining all of its constructors as private.

### Exercise 3

Noncreatable classes can be useful through the static members they possess. In fact, you can even get instances of these classes through these members, as shown here:

```
class CreateMe
{
    private CreateMe()
    {
    }
    static public CreateMe GetCreateMe()
    {
        return new CreateMe();
    }
}
```

Here, the public constructor has access to the private constructor, as it is part of the same class definition.

### Exercise 4

For simplicity, the following class definitions are shown as part of a single code file, rather than listing a separate code file for each:

```
namespace Vehicles
{
    public abstract class Vehicle
    {
    }
    public abstract class Car : Vehicle
    {
    }
    public abstract class Train : Vehicle
    {
    }
}
```

```
public interface IPassengerCarrier
{
}
public interface IHeavyLoadCarrier
{
}
public class SUV : Car, IPassengerCarrier
{
}
public class Pickup : Car, IPassengerCarrier, IHeavyLoadCarrier
{
}
public class Compact : Car, IPassengerCarrier
{
}
public class PassengerTrain : Train, IPassengerCarrier
{
}
public class FreightTrain : Train, IHeavyLoadCarrier
{
}
public class T424DoubleBogey : Train, IHeavyLoadCarrier
{
}
}
```

## Exercise 5

```
using System;
using static System.Console;
using Vehicles;
namespace Traffic
{
    class Program
    {
        static void Main(string[] args)
        {
            AddPassenger(new Compact());
            AddPassenger(new SUV());
            AddPassenger(new Pickup());
            AddPassenger(new PassengerTrain());
            ReadKey();
        }
        static void AddPassenger(IPassengerCarrier Vehicle)
        {
            WriteLine(Vehicle.ToString());
        }
    }
}
```

## CHAPTER 10 SOLUTIONS

### Exercise 1

```
class MyClass
{
    protected string myString;
    public string ContainedString
    {
        set
        {
            myString = value;
        }
    }
    public virtual string GetString() => myString;
}
```

### Exercise 2

```
class MyDerivedClass : MyClass
{
    public override string GetString() => base.GetString() +
        " (output from derived class)";
}
```

### Exercise 3

If a method has a return type, then it is possible to use it as part of an expression:

```
x = Manipulate(y, z);
```

If no implementation is provided for a partial method, then it will be removed by the compiler along with all places where it is used. In the preceding code this would leave the result of x unclear because no replacement for the `Manipulate()` method is available. It might be the case that without this method you would simply want to ignore the entire line of code, but the compiler cannot decide whether this is what you want.

Methods with no return types are not called as part of expressions, so it is safe for the compiler to remove all references to the partial method calls.

Similarly, `out` parameters are forbidden since variables used as an `out` parameter must be undefined before the method call and will be defined after the method call. Removing the method call would break this behavior.

### Exercise 4

```
class MyCopyableClass
{
    protected int myInt;
```

```
public int ContainedInt
{
    get
    {
        return myInt;
    }
    set
    {
        myInt = value;
    }
}
public MyCopyableClass GetCopy() => (MyCopyableClass)MemberwiseClone();
```

The client code:

```
class Program
{
    using static System.Console;
    static void Main(string[] args)
    {
        MyCopyableClass obj1 = new MyCopyableClass();
        obj1.ContainedInt = 5;
        MyCopyableClass obj2 = obj1.GetCopy();
        obj1.ContainedInt = 9;
        WriteLine(obj2.ContainedInt);
    }
}
```

This code displays 5, showing that the copied object has its own version of the `myInt` field.

## Exercise 5

```
using System;
using static System.Console;
using Ch10CardLib;
namespace Exercise_Answers
{
    class Class1
    {
        static void Main(string[] args)
        {
            while(true)
            {
                Deck playDeck = new Deck();
                playDeck.Shuffle();
                bool isFlush = false;
                int flushHandIndex = 0;
                for (int hand = 0; hand < 10; hand++)
                {
                    isFlush = true;
                    Suit flushSuit = playDeck.GetCard(hand * 5).suit;
                    for (int card = 1; card < 5; card++)
                    {
```

```
        if (playDeck.GetCard(hand * 5 + card).suit != flushSuit)
        {
            isFlush = false;
        }
    }
    if (isFlush)
    {
        flushHandIndex = hand * 5;
        break;
    }
}
if (isFlush)
{
    WriteLine("Flush!");
    for (int card = 0; card < 5; card++)
    {
        WriteLine(playDeck.GetCard(flushHandIndex + card));
    }
}
else
{
    WriteLine("No flush.");
}
ReadLine();
}
}
```

This code is looped as flushes are uncommon. You might need to press Return several times before a flush is found in a shuffled deck. To verify that everything is working as it should, try commenting out the line that shuffles the deck.

## CHAPTER 11 SOLUTIONS

## Exercise 1

```
using System;
using System.Collections;
namespace Exercise_Answers
{
    public class People : DictionaryBase
    {
        public void Add(Person newPerson) =>
            Dictionary.Add(newPerson.Name, newPerson);

        public void Remove(string name) => Dictionary.Remove(name);

        public Person this[string name]
        {
            get
            {
                return (Person)Dictionary[name];
            }
        }
    }
}
```

```
        }
        set
        {
            Dictionary[name] = value;
        }
    }
}
```

## Exercise 2

```
public class Person
{
    private string name;
    private int age;
    public string Name
    {
        get
        {
            return name;
        }
        set
        {
            name = value;
        }
    }
    public int Age
    {
        get
        {
            return age;
        }
        set
        {
            age = value;
        }
    }
    public static bool operator >(Person p1, Person p2) =>
        p1.Age > p2.Age;
    public static bool operator <(Person p1, Person p2) =>
        p1.Age < p2.Age;
    public static bool operator >=(Person p1, Person p2) =>
        !(p1 < p2);
    public static bool operator <=(Person p1, Person p2) =>
        !(p1 > p2);
}
```

## Exercise 3

```
public Person[] GetOldest()
{
    Person oldestPerson = null;
    People oldestPeople = new People();
    Person currentPerson;
    foreach (DictionaryEntry p in Dictionary)
```

```

{
    currentPerson = p.Value as Person;
    if (oldestPerson == null)
    {
        oldestPerson = currentPerson;
        oldestPeople.Add(oldestPerson);
    }
    else
    {
        if (currentPerson > oldestPerson)
        {
            oldestPeople.Clear();
            oldestPeople.Add(currentPerson);
            oldestPerson = currentPerson;
        }
        else
        {
            if (currentPerson >= oldestPerson)
            {
                oldestPeople.Add(currentPerson);
            }
        }
    }
}
Person[] oldestPeopleArray = new Person[oldestPeople.Count];
int copyIndex = 0;
foreach (DictionaryEntry p in oldestPeople)
{
    oldestPeopleArray[copyIndex] = p.Value as Person;
    copyIndex++;
}
return oldestPeopleArray;
}

```

This function is made more complex by the fact that no `==` operator has been defined for `Person`, but the logic can still be constructed without this. In addition, returning a `People` instance would be simpler, as it is easier to manipulate this class during processing. As a compromise, a `People` instance is used throughout the function, and then converted into an array of `Person` instances at the end.

## Exercise 4

```

public class People : DictionaryBase, ICloneable
{
    public object Clone()
    {
        People clonedPeople = new People();
        Person currentPerson, newPerson;
        foreach (DictionaryEntry p in Dictionary)
        {
            currentPerson = p.Value as Person;
            newPerson = new Person();
            newPerson.Name = currentPerson.Name;
            newPerson.Age = currentPerson.Age;
        }
    }
}

```

```
        clonedPeople.Add(newPerson);
    }
    return clonedPeople;
}
...
}
```

You could simplify this by implementing `ICloneable` on the `Person` class.

## Exercise 5

```
public IEnumerable Ages
{
    get
    {
        foreach (object person in Dictionary.Values)
            yield return (person as Person).Age;
    }
}
```

# CHAPTER 12 SOLUTIONS

## Exercise 1

a, b, and e: Yes

c and d: No, although they can use generic type parameters supplied by the class containing them.

f: No

## Exercise 2

```
public static double? operator *(Vector op1, Vector op2)
{
    try
    {
        double angleDiff = (double)(op2.ThetaRadians.Value -
            op1.ThetaRadians.Value);
        return op1.R.Value * op2.R.Value * Math.Cos(angleDiff);
    }
    catch
    {
        return null;
    }
}
```

## Exercise 3

You can't instantiate `T` without enforcing the `new()` constraint on it, which ensures that a public default constructor is available:

```
public class Instantiator<T>
    where T : new()
```

```

    {
        public T instance;
        public Instantiator()
        {
            instance = new T();
        }
    }
}

```

## Exercise 4

The same generic type parameter, `T`, is used on both the generic class and the generic method. You need to rename one or both. For example:

```

public class StringGetter<U>
{
    public string GetString<T>(T item) => item.ToString();
}

```

## Exercise 5

One way of doing this is as follows:

```

public class ShortList<T> : IList<T>
{
    protected IList<T> innerCollection;
    protected int maxSize = 10;
    public ShortList()
        : this(10)
    {
    }
    public ShortList(int size)
    {
        maxSize = size;
        innerCollection = new List<T>();
    }
    public ShortList(IEnumerable<T> list)
        : this(10, list)
    {
    }
    public ShortList(int size, IEnumerable<T> list)
    {
        maxSize = size;
        innerCollection = new List<T>(list);
        if (Count > maxSize)
        {
            ThrowTooManyItemsException();
        }
    }
    protected void ThrowTooManyItemsException()
    {
        throw new IndexOutOfRangeException(
            "Unable to add any more items, maximum size is " + maxSize.ToString()
            + " items.");
    }
}
#region IList<T> Members

```

```
public int IndexOf(T item) => innerCollection.IndexOf(item);
public void Insert(int index, T item)
{
    if (Count < maxSize)
    {
        innerCollection.Insert(index, item);
    }
    else
    {
        ThrowTooManyItemsException();
    }
}
public void RemoveAt(int index)
{
    innerCollection.RemoveAt(index);
}
public T this[int index]
{
    get
    {
        return innerCollection[index];
    }
    set
    {
        innerCollection[index] = value;
    }
}
#endregion
#region ICollection<T> Members
public void Add(T item)
{
    if (Count < maxSize)
    {
        innerCollection.Add(item);
    }
    else
    {
        ThrowTooManyItemsException();
    }
}
public void Clear()
{
    innerCollection.Clear();
}
public bool Contains(T item) => innerCollection.Contains(item);
public void CopyTo(T[] array, int arrayIndex)
{
    innerCollection.CopyTo(array, arrayIndex);
}
public int Count
{
    get
    {
        return innerCollection.Count;
    }
}
```

```

public bool IsReadOnly
{
    get
    {
        return innerCollection.IsReadOnly;
    }
}
public bool Remove(T item) => innerCollection.Remove(item);
#endregion
#region IEnumerable<T> Members
public IEnumerator<T> GetEnumerator() =>
    innerCollection.GetEnumerator();
#endregion
#region IEnumerable Members
IEnumerator IEnumerable.GetEnumerator() => GetEnumerator();
#endregion
}

```

## Exercise 6

No, it won't. The type parameter `T` is defined as being covariant. However, covariant type parameters can be used only as return values of methods, not as method arguments. If you try this out you will get the following compiler error (assuming you use the namespace `VarianceDemo`):

```
Invalid variance: The type parameter 'T' must be contravariantly valid on
'VarianceDemo.IMethaneProducer<T>.BelchAt(T)'. 'T' is covariant.
```

# CHAPTER 13 SOLUTIONS

## Exercise 1

```

using static System.Console;
public void ProcessEvent(object source, EventArgs e)
{
    if (e is MessageArrivedEventArgs)
    {
        WriteLine("Connection.MessageArrived event received.");
        WriteLine($"Message: {(e as MessageArrivedEventArgs).Message}");
    }
    if (e is ElapsedEventArgs)
    {
        WriteLine("Timer.Elapsed event received.");
        WriteLine($"SignalTime: {(e as ElapsedEventArgs).SignalTime}");
    }
}

```

## Exercise 2

Modify `Player.cs` as follows (one modified method, two new ones—comments in the code explain the changes):

```

public bool HasWon()
{
    // get temporary copy of hand, which may get modified.

```

```
Cards tempHand = (Cards)PlayHand.Clone();
// find three and four of a kind sets
bool fourOfAKind = false;
bool threeOfAKind = false;
int fourRank = -1;
int threeRank = -1;
int cardsOfRank;
for (int matchRank = 0; matchRank < 13; matchRank++)
{
    cardsOfRank = 0;
    foreach (Card c in tempHand)
    {
        if (c.rank == (Rank)matchRank)
        {
            cardsOfRank++;
        }
    }
    if (cardsOfRank == 4)
    {
        // mark set of four
        fourRank = matchRank;
        fourOfAKind      if (cardsOfRank == 3)
    {
        // two threes means no win possible
        // (threeOfAKind will be true only if this code
        // has already executed)
        if (threeOfAKind == true)
        {
            return false;
        }
        // mark set of three
        threeRank = matchRank;
        threeOfAKind = true;
    }
}
// check simple win condition
if (threeOfAKind && fourOfAKind)
{
    return true;
}
// simplify hand if three or four of a kind is found,
// by removing used cards
if (fourOfAKind || threeOfAKind)
{
    for (int cardIndex = tempHand.Count - 1; cardIndex >= 0; cardIndex--)
    {
        if ((tempHand[cardIndex].rank == (Rank)fourRank)
            || (tempHand[cardIndex].rank == (Rank)threeRank))
        {
            tempHand.RemoveAt(cardIndex);
        }
    }
}
// at this point the method may have returned, because:
// - a set of four and a set of three has been found, winning.
```

```
// - two sets of three have been found, losing.  
// if the method hasn't returned then:  
// - no sets have been found, and tempHand contains 7 cards.  
// - a set of three has been found, and tempHand contains 4 cards.  
// - a set of four has been found, and tempHand contains 3 cards.  
// find run of four sets, start by looking for cards of same suit  
// in the same way as before  
bool fourOfASuit = false;  
bool threeOfASuit = false;  
int fourSuit = -1;  
int threeSuit = -1;  
int cardsOfSuit;  
for (int matchSuit = 0; matchSuit < 4; matchSuit++)  
{  
    cardsOfSuit = 0;  
    foreach (Card c in tempHand)  
    {  
        if (c.suit == (Suit)matchSuit)  
        {  
            cardsOfSuit++;  
        }  
    }  
    if (cardsOfSuit == 7)  
    {  
        // if all cards are the same suit then two runs  
        // are possible, but not definite.  
        threeOfASuit = true;  
        threeSuit = matchSuit;  
        fourOfASuit = true;  
        fourSuit = matchSuit;  
    }  
    if (cardsOfSuit == 4)  
    {  
        // mark four card suit.  
        fourOfASuit = true;  
        fourSuit = matchSuit;  
    }  
    if (cardsOfSuit == 3)  
    {  
        // mark three card suit.  
        threeOfASuit = true;  
        threeSuit = matchSuit;  
    }  
}  
if (!(threeOfASuit || fourOfASuit))  
{  
    // need at least one run possibility to continue.  
    return false;  
}  
if (tempHand.Count == 7)  
{  
    if (!(threeOfASuit && fourOfASuit))  
    {  
        // need a three and a four card suit.  
        return false;  
    }
```

```
// create two temporary sets for checking.
Cards set1 = new Cards();
Cards set2 = new Cards();
// if all 7 cards are the same suit...
if (threeSuit == fourSuit)
{
    // get min and max cards
    int maxVal, minVal;
    GetLimits(tempHand, out maxVal, out minVal);
    for (int cardIndex = tempHand.Count - 1; cardIndex >= 0; cardIndex--)
    {
        if (((int)tempHand[cardIndex].rank < (minVal + 3))
            || ((int)tempHand[cardIndex].rank > (maxVal - 3)))
        {
            // remove all cards in a three card set that
            // starts at minVal or ends at maxVal.
            tempHand.RemoveAt(cardIndex);
        }
    }
    if (tempHand.Count != 1)
    {
        // if more then one card is left then there aren't two runs.
        return false;
    }
    if ((tempHand[0].rank == (Rank)(minVal + 3))
        || (tempHand[0].rank == (Rank)(maxVal - 3)))
    {
        // if spare card can make one of the three card sets into a
        // four card set then there are two sets.
        return true;
    }
    else
    {
        // if spare card doesn't fit then there are two sets of three
        // cards but no set of four cards.
        return false;
    }
}
// if three card and four card suits are different...
foreach (Card card in tempHand)
{
    // split cards into sets.
    if (card.suit == (Suit)threeSuit)
    {
        set1.Add(card);
    }
    else
    {
        set2.Add(card);
    }
}
// check if sets are sequential.
if (isSequential(set1) && isSequential(set2))
{
    return true;
```

```
        }
    else
    {
        return false;
    }
}
// if four cards remain (three of a kind found)
if (tempHand.Count == 4)
{
    // if four cards remain then they must be the same suit.
    if (!fourOfASuit)
    {
        return false;
    }
    // won if cards are sequential.
    if (isSequential(tempHand))
    {
        return true;
    }
}
// if three cards remain (four of a kind found)
if (tempHand.Count == 3)
{
    // if three cards remain then they must be the same suit.
    if (!threeOfASuit)
    {
        return false;
    }
    // won if cards are sequential.
    if (isSequential(tempHand))
    {
        return true;
    }
}
// return false if two valid sets don't exist.
return false;
}
// utility method to get max and min ranks of cards
// (same suit assumed)
private void GetLimits(Cards cards, out int maxVal, out int minVal)
{
    maxVal = 0;
    minVal = 14;
    foreach (Card card in cards)
    {
        if ((int)card.rank > maxVal)
        {
            maxVal = (int)card.rank;
        }
        if ((int)card.rank < minVal)
        {
            minVal = (int)card.rank;
        }
    }
}
```

```
// utility method to see if cards are in a run
// (same suit assumed)
private bool isSequential(Cards cards)
{
    int maxVal, minVal;
    GetLimits(cards, out maxVal, out minVal);
    if ((maxVal - minVal) == (cards.Count - 1))
    {
        return true;
    }
    else
    {
        return false;
    }
}
```

## Exercise 3

In order to use an object initializer with a class, you must include a default, parameter-less constructor. You could either add one to this class or remove the nondefault constructor that is there already. Once you have done this, you can use the following code to instantiate and initialize this class in one step:

```
Giraffe myPetGiraffe = new Giraffe
{
    NeckLength = "3.14",
    Name = "Gerald"
};
```

## Exercise 4

False. When you use the `var` keyword to declare a variable, the variable is still strongly typed; the compiler determines the type of the variable.

## Exercise 5

You can use the `Equals()` method that is implemented for you. Note that you cannot use the `==` operator to do this, as this compares variables to determine if they both refer to the same object.

## Exercise 6

The extension method must be static:

```
public static string ToAcronym(this string inputString)
```

## Exercise 7

You must include the extension method in a static class that is accessible from the namespace that contains your client code. You could do this either by including the code in the same namespace or by importing the namespace containing the class.

## Exercise 8

One way to do this is as follows:

```
public static string ToAcronym(this string inputString) =>
    inputString.Trim().Split(' ').Aggregate<string, string>("", 
        (a, b) => a + (b.Length > 0 ? 
            b.ToUpper()[0].ToString() : ""));
```

Here the tertiary operator prevents multiple spaces from causing errors. Note also that the version of `Aggregate()` with two generic type parameters is required, as a seed value is necessary.

# CHAPTER 14 SOLUTIONS

## Exercise 1

Wrap the `TextBlock` control in a `ScrollViewer` panel. Set the `VerticalScrollBarVisibility` property to `Auto` to make the scrollbar appear when the text extends beyond the bottom edge of the control.

```
<Window x:Class="Answers.MainWindow"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    Title="14.1 Solution" Height="350" Width="525">
    <Grid>
        <Grid.RowDefinitions>
            <RowDefinition Height="75"/>
            <RowDefinition />
        </Grid.RowDefinitions>
        <Label Content="Enter text" HorizontalAlignment="Left" Margin="10,10,0,0"
            VerticalAlignment="Top"/>
        <TextBox HorizontalAlignment="Left" Margin="76,12,0,0" TextWrapping="Wrap"
            VerticalAlignment="Top" Height="53" Width="423" AcceptsReturn="True"
            Name="textTextBox">
        </TextBox>
        <ScrollViewer HorizontalAlignment="Left" Height="217" Margin="10,10,0,0"
            Grid.Row="1" VerticalAlignment="Top" Width="489"
            VerticalScrollBarVisibility="Auto">
            <TextBlock TextWrapping="Wrap" Text="{Binding ElementName=textTextBox,
                Path=Text}"/>
        </ScrollViewer>
    </Grid>
</Window>
```

## Exercise 2

After dragging a `Slider` and `ProgressBar` control into the view, set the minimum and maximum values of the slider to 1 and 100 and the `Value` property to 1. Bind the same values of the `ProgressBar` to the `Slider`.

```
<Window x:Class="Answers.Ch14Solution2"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
```

```
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    Title="14.2 Solution" Height="300" Width="300">
<Grid>
    <Slider HorizontalAlignment="Left" Margin="10,10,0,0" VerticalAlignment="Top"
Width="264" Minimum="1" Maximum="100" Name="valueSlider"/>
    <ProgressBar HorizontalAlignment="Left" Height="24" Margin="10,77,0,0"
VerticalAlignment="Top" Width="264"
Minimum="{Binding ElementName=valueSlider, Path=Minimum}"
Maximum="{Binding ElementName=valueSlider, Path=Maximum}"
Value="{Binding ElementName=valueSlider, Path=Value}"/>
</Grid>
</Window>
```

## Exercise 3

You can use a `RenderTransform` to do this. In Design View, you can position the cursor over the edge of the control and when you see a quarter circle icon for the mouse pointer, click and drag the control to the desired position.

```
<Window x:Class="Answers.Ch14Solution3"
        xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
        xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
        Title="14.3 Solution" Height="300" Width="300">
    <Grid>
        <Slider HorizontalAlignment="Left" Margin="10,10,0,0" VerticalAlignment="Top"
Width="264" Minimum="1" Maximum="100" Name="valueSlider"/>
        <ProgressBar HorizontalAlignment="Left" Height="24" Margin="-17,125,-10,0"
VerticalAlignment="Top" Width="311"
Minimum="{Binding ElementName=valueSlider, Path=Minimum}" Maximum="{Binding
ElementName=valueSlider, Path=Maximum}"
Value="{Binding ElementName=valueSlider, Path=Value}"
RenderTransformOrigin="0.5,0.5">
            <ProgressBar.RenderTransform>
                <TransformGroup>
                    <ScaleTransform/>
                    <SkewTransform/>
                    <RotateTransform Angle="-36.973"/>
                    <TranslateTransform/>
                </TransformGroup>
            </ProgressBar.RenderTransform>
        </ProgressBar>
    </Grid>
</Window>
```

## Exercise 4

The `PersistentSlider` class must implement the `INotifyPropertyChanged` interface.

Create a field to hold the value of each of the three properties.

In each of the setters of the properties, implement a call to any subscribers of the `PropertyChanged` event. You are advised to create a helper method, called `OnPropertyChanged`, for this purpose.

`PersistentSlider.cs`:

```
using System.ComponentModel;
namespace Answers
```

```

{
    public class PersistentSlider : INotifyPropertyChanged
    {
        private int _minValue;
        private int _maxValue;
        private int _currentValue;
        public int MinValue
        {
            get { return _minValue; }
            set { _minValue = value; OnPropertyChanged(nameof(MinValue)); }
        }
        public int MaxValue
        {
            get { return _maxValue; }
            set { _maxValue = value; OnPropertyChanged(nameof(MaxValue)); }
        }
        public int CurrentValue
        {
            get { return _currentValue; }
            set { _currentValue = value; OnPropertyChanged(nameof(CurrentValue)); }
        }
        public event PropertyChangedEventHandler PropertyChanged;
        protected void OnPropertyChanged(string propertyName) => PropertyChanged?.Invoke(this, new PropertyChangedEventArgs(propertyName));
    }
}

```

1. In the code-behind file, add a field like this:

```
private PersistentSlider _sliderData = new PersistentSlider { MinValue = 1, MaxValue = 200, CurrentValue = 100 };
```

2. In the constructor, set the `DataContext` property of the current instance to the field you just created:

```
this.DataContext = _sliderData;
InitializeComponent();
```

3. In the XAML, change the `Slider` control to use the data context. Only the `Path` needs to be set:

```
<Window x:Class="Answers.Ch14Solution4"
        xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
        xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
        Title="14.4 Solution" Height="300" Width="300">
    <Grid>
        <Slider HorizontalAlignment="Left" Margin="10,10,0,0"
VerticalAlignment="Top"
Width="264" Minimum="{Binding Path=MinValue}"
Maximum="{Binding Path=MaxValue}" Value="{Binding Path=CurrentValue}"
Name="valueSlider"/>
        <ProgressBar HorizontalAlignment="Left" Height="24" Margin="-17,125,-10,0"
VerticalAlignment="Top" Width="311"
Minimum="{Binding ElementName=valueSlider, Path=Minimum}"
Maximum="{Binding ElementName=valueSlider, Path=Maximum}"
Value="{Binding ElementName=valueSlider, Path=Value}"
RenderTransformOrigin="0.5,0.5">
    
```

```
<ProgressBar.RenderTransform>
<TransformGroup>
<ScaleTransform/>
<SkewTransform/>
<RotateTransform Angle="-36.973"/>
<TranslateTransform/>
</TransformGroup>
</ProgressBar.RenderTransform>
</ProgressBar>
</Grid>
</Window>
```

## CHAPTER 15 SOLUTIONS

### Exercise 1

Solution:

1. Create a new class with the name ComputerSkillValueConverter like this:

```
using Ch13CardLib;
using System;
using System.Windows.Data;

namespace KarliCards_Gui
{
    [ValueConversion(typeof(ComputerSkillLevel), typeof(bool))]
    public class ComputerSkillValueConverter : IValueConverter
    {
        public object Convert(object value, Type targetType, object parameter,
System.Globalization.CultureInfo culture)
        {
            string helper = parameter as string;
            if (string.IsNullOrWhiteSpace(helper))
                return false;

            ComputerSkillLevel skillLevel = (ComputerSkillLevel)value;
            return (skillLevel.ToString() == helper);
        }

        public object ConvertBack(object value, Type targetType, object parameter,
System.Globalization.CultureInfo culture)
        {
            string parameterString = parameter as string;
            if (parameterString == null)
                return ComputerSkillLevel.Dumb;

            return Enum.Parse(targetType, parameterString);
        }
    }
}
```

2. Add a static resource declaration to the Options.xaml:

```
<Window.Resources>
    <src:ComputerSkillValueConverter x:Key="skillConverter" />
</Window.Resources>
```

3. Change the radio buttons like this:

```
<RadioButton Content="Dumb" HorizontalAlignment="Left"
Margin="37,41,0,0" VerticalAlignment="Top" Name="dumbAIRadioButton"
IsChecked="{Binding ComputerSkill, Converter={StaticResource skillConverter},
ConverterParameter=Dumb}" />
<RadioButton Content="Good" HorizontalAlignment="Left"
Margin="37,62,0,0" VerticalAlignment="Top" Name="goodAIRadioButton"
IsChecked="{Binding ComputerSkill, Converter={StaticResource skillConverter},
ConverterParameter=Good}" />
<RadioButton Content="Cheats" HorizontalAlignment="Left"
Margin="37,83,0,0" VerticalAlignment="Top" Name="cheatingAIRadioButton"
IsChecked="{Binding ComputerSkill, Converter={StaticResource skillConverter},
ConverterParameter=Cheats}" />
```

4. Delete the events from the code-behind file.

## Exercise 2

Solution:

1. Add a new check box to the Options.xaml dialog box:

```
<CheckBox Content="Plays with open cards" HorizontalAlignment="Left"
Margin="10,100, 0,0" VerticalAlignment="Top"
IsChecked="{Binding ComputerPlaysWithOpenHand}" />
```

2. Add a new property to the GameOptions.cs class:

```
private bool _computerPlaysWithOpenHand;
public bool ComputerPlaysWithOpenHand
{
    get { return _computerPlaysWithOpenHand; }
    set
    {
        _computerPlaysWithOpenHand = value;
        OnPropertyChanged(nameof(ComputerPlaysWithOpenHand));
    }
}
```

3. Add a new dependency property to the CardsInHandControl:

```
public bool ComputerPlaysWithOpenHand
{
    get { return (bool)GetValue(ComputerPlaysWithOpenHandProperty); }
    set { SetValue(ComputerPlaysWithOpenHandProperty, value); }
}

public static readonly DependencyProperty ComputerPlaysWithOpenHandProperty =
    DependencyProperty.Register("ComputerPlaysWithOpenHand", typeof(bool),
    typeof(CardsInHandControl), new PropertyMetadata(false));
```

4. In the DrawCards method of the CardsInHandControl, change the test for isFaceUp:

```
if (Owner is ComputerPlayer)
    isFaceup = (Owner.State == CardLib.PlayerState.Loser ||
Owner.State == CardLib.PlayerState.Winner || ComputerPlaysWithOpenHand);
```

5. Add a new property to the GameViewModel class:

```
public bool ComputerPlaysWithOpenHand
{
    get { return _gameOptions.ComputerPlaysWithOpenHand; }
}
```

6. Bind the new property to the CardsInHandControls on the game client to all four players:

```
ComputerPlaysWithOpenHand="{Binding GameOptions.ComputerPlaysWithOpenHand}"
```

## Exercise 3

Solution:

1. Add a new property to the GameViewModel like this:

```
private string _currentStatusText = "Game is not started";
public string CurrentStatusText
{
    get { return _currentStatusText; }
    set
    {
        _currentStatusText = value;
        OnPropertyChanged(nameof(CurrentStatusText));
    }
}
```

2. Change the CurrentPlayer property like this:

```
public Player CurrentPlayer
{
    get { return _currentPlayer; }
    set
    {
        _currentPlayer = value;
        OnPropertyChanged("CurrentPlayer");
        if (!Players.Any(x => x.State == PlayerState.Winner))
        {
            Players.ForEach(x => x.State = (x == value ? PlayerState.Active :
PlayerState.Inactive));
            CurrentStatusText = $"Player {CurrentPlayer.PlayerName} ready";
        }
        else
        {
            var winner = Players.Where(x => x.HasWon).FirstOrDefault();
            if (winner != null)
```

```
        CurrentStatusText = $"Player {winner.PlayerName} has WON!"  
    }  
}
```

3. Add this line at the end of the StartNewGame method:

```
CurrentStatusText = string.Format("New game stated. Player {0} to start",  
CurrentPlayer.PlayerName);
```

4. Add a status bar to the game client XAML and set the binding to the new property:

```
<StatusBar Grid.Row="3" HorizontalAlignment="Center" Margin="0,0,0,15"
VerticalAlignment="Center" Background="Green" Foreground="White"
FontWeight="Bold">
    <StatusBarItem VerticalAlignment="Center">
        <TextBlock Text="{Binding CurrentStatusText}" />
    </StatusBarItem>
</StatusBar>
```

## CHAPTER 16 SOLUTIONS

## Exercise 1

To find the answer to this question, you should have a look at the `PlayGame()` method in the `Game.cs` file. Have a look through the method and list the variables it references while within the main `do...while` loop. This information would need to be sent back and forth between the client and server for the game to work via a web site:

- How many people are playing and what are their names?
  - Who is the current player?
  - The player's hand of cards.
  - The current card in play.
  - The player's action, for example taking, drawing or discarding.
  - A list of discarded cards.
  - The status of the game, such as whether somebody won.

## Exercise 2

You can store the information in a database and then retrieve the required data with each call, and you can pass the required information back and forth between the client and server using the ASP .NET Session Object or VIEWSTATE.

For information about the ASP.NET Session Object, read this article: <https://msdn.microsoft.com/en-us/library/ms178581.aspx>

For information about VIEWSTATE, read this article: <https://msdn.microsoft.com/en-us/library/ms972976.aspx>

## CHAPTER 17 SOLUTIONS

### Exercise 1

```
...
using System.Net;
using System.IO;
using Newtonsoft.Json;
using static System.Console;
namespace handofcards
{
    class Program
    {
        static void Main(string[] args)
        {
            List<string> cards = new List<string>();
            var playerName = "Benjamin";
            string GetURL =
                "http://handofcards.azurewebsites.net/api/HandOfCards/" +
                playerName;
            WebClient client = new WebClient();
            Stream dataStream = client.OpenRead(GetURL);
            StreamReader reader = new StreamReader(dataStream);
            var results =
                JsonConvert.DeserializeObject<dynamic>(reader.ReadLine());
            reader.Close();
            foreach (var item in results)
            {
                WriteLine((string)item.imageLink);
            }
            ReadLine();
        }
    }
}
```

### Exercise 2

The maximum size of a Web App VM is 4 CPU/Cores (~2.6Ghz) with 7GB of RAM.

The maximum number of VMs that you can have in Standard mode is 10. The maximum number of VMs you can have in Premium mode is 50. That translates into a maximum 200 x 2.6Ghz cores with 350GB of memory loaded across 50 virtual machines.

Note that this is for Web Apps. You can utilize Azure VMs or Azure Cloud Services to get even more cores and memory.

## CHAPTER 18 SOLUTIONS

### Exercise 1

System.IO

### Exercise 2

You use a `FileStream` object to write to a file when you need random access to files, or when you are not dealing with string data.

### Exercise 3

- `Peek()`: Gets the value of the next character in the file but does not advance the file position
- `Read()`: Gets the value of the next character in the file and advances the file position
- `Read(char[] buffer, int index, int count)`: Reads `count` characters into `buffer`, starting at `buffer[index]`
- `ReadLine()`: Gets a line of text
- `ReadToEnd()`: Gets all text in a file

### Exercise 4

DeflateStream

### Exercise 5

- `Changed`: Occurs when a file is modified
- `Created`: Occurs when a file is created
- `Deleted`: Occurs when a file is deleted
- `Renamed`: Occurs when a file is renamed

### Exercise 6

Add a button that toggles the value of the `FileSystemWatcher.EnableRaisingEvents` property.

## CHAPTER 19 SOLUTIONS

### Exercise 1

1. Double-click the Create Node button to go to the event handler doing the work.
2. Below the creation of the `xmlComment`, insert the following three lines:

```
XmlAttribute newPages = document.CreateAttribute("pages");
newPages.Value = "1000";
newBook.Attributes.Append(newPages);
```

## Exercise 2

1. //elements — Returns all nodes in the document.
2. element — Returns every element node in the document but leaves the element root node out.
3. element[@Type='Noble Gas'] — Returns every element that includes an attribute with the name Type, which has a value of Noble Gas.
4. //mass — Returns all nodes with the name mass.
5. //mass/.. — The .. causes the XPath to move one up from the selected node, which means that this query selects all the nodes that include a mass node.
6. element/specification[mass='20.1797'] — Selects the specification element that contains a mass node with the value 20.1797.
7. element/name[text()='Neon'] — To select the node whose contents you are testing, you can use the text() function. This selects the name node with the text Neon.

## Exercise 3

Recall that XML can be valid, well-formed, or invalid. Whenever you select part of an XML document, you are left with a fragment of the whole. This means that there is a good chance that the XML you've selected is in fact invalid XML on its own. Most XML viewers will refuse to display XML that isn't well-formed, so it is not possible to display the results of many queries directly in a standard XML viewer.

## Exercise 4

Add a new button JSON>XML to MainWindow.xaml and then add the following code to MainWindow.xaml.cs:

```
private void buttonConvertXMLtoJSON_Click(object sender, RoutedEventArgs e)
{
    // Load the XML document.
    XmlDocument document = new XmlDocument();

    document.Load(@"C:\BegVCSharp\Chapter19\XML and Schema\Books.xml");

    string json = Newtonsoft.Json.JsonConvert.SerializeXmlNode(document);

    textBlockResults.Text = json;

    System.IO.File.AppendAllText
        (@"C:\BegVCSharp\Chapter19\XML and Schema\Books.json", json);
}
```

```

private void buttonConvertJSONtoXML_Click(object sender, RoutedEventArgs e)
{
    // Load the json document.
    string json = System.IO.File.ReadAllText
        ("@C:\BegVCSharp\Chapter19\XML and Schema\Books.json");

    XmlDocument document =
        Newtonsoft.Json.JsonConvert.DeserializeXmlNode(json);

    textBlockResults.Text =
        FormatText(document.DocumentElement as XmlNode, "", "");
}

```

## CHAPTER 20 SOLUTIONS

### Exercise 1

```

static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
        "Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };

    var queryResults =
        from n in names
        where n.StartsWith("S")
        orderby n descending
        select n;

    Console.WriteLine("Names beginning with S:");

    foreach (var item in queryResults) {
        Console.WriteLine(item);
    }

    Console.Write("Program finished, press Enter/Return to continue:");
    Console.ReadLine();
}

```

### Exercise 2

Sets smaller than 5,000,000 have no numbers < 1000:

```

static void Main(string[] args)
{
    int[] arraySizes = {      100,      1000,     10000,   100000,
                           1000000,  5000000, 10000000, 50000000 };

    foreach (int i in arraySizes) {
        int[] numbers = generateLotsOfNumbers(i);
        var queryResults = from n in numbers
                           where n < 1000
                           select n;
        Console.WriteLine("number array size = {0}: Count(n < 1000) = {1}",

```

```
        numbers.Length, queryResults.Count()
    );
}

Console.WriteLine("Program finished, press Enter/Return to continue:");
Console.ReadLine();
}
```

## Exercise 3

Does not affect performance noticeably for n < 1000:

```
static void Main(string[] args)
{
    int[] numbers = generateLotsOfNumbers(12345678);

    var queryResults =
        from n in numbers
        where n < 1000
        orderby n
        select n
    ;

    Console.WriteLine("Numbers less than 1000:");
    foreach (var item in queryResults)
    {
        Console.WriteLine(item);
    }

    Console.WriteLine("Program finished, press Enter/Return to continue:");
    Console.ReadLine();
}
```

## Exercise 4

Very large subsets such as n > 1000 instead of n < 1000 are very slow:

```
static void Main(string[] args)
{
    int[] numbers = generateLotsOfNumbers(12345678);

    var queryResults =
        from n in numbers
        where n > 1000
        select n
    ;

    Console.WriteLine("Numbers less than 1000:");
    foreach (var item in queryResults)
    {
        Console.WriteLine(item);
    }
}
```

```

    }

    Console.WriteLine("Program finished, press Enter/Return to continue:");
    Console.ReadLine();
}

```

## Exercise 5

All the names are output because there is no query.

```

static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
"Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };

    var queryResults = names;

    foreach (var item in queryResults) {
        Console.WriteLine(item);
    }

    Console.WriteLine("Program finished, press Enter/Return to continue:");
    Console.ReadLine();
}

```

## Exercise 6

```

static void Main(string[] args)
{
    string[] names = { "Alonso", "Zheng", "Smith", "Jones", "Smythe",
"Small", "Ruiz", "Hsieh", "Jorgenson", "Ilyich", "Singh", "Samba", "Fatimah" };
    // only Min() and Max() are available (if no lambda is used)
    // for a result set like this consisting only of strings
    Console.WriteLine("Min(names) = " + names.Min());
    Console.WriteLine("Max(names) = " + names.Max());
    var queryResults =
        from n in names
        where n.StartsWith("S")
        select n;

    Console.WriteLine("Query result: names starting with S");
    foreach (var item in queryResults)
    {
        Console.WriteLine(item);
    }

    Console.WriteLine("Min(queryResults) = " + queryResults.Min());
    Console.WriteLine("Max(queryResults) = " + queryResults.Max());

    Console.WriteLine("Program finished, press Enter/Return to continue:");
    Console.ReadLine();
}

```

## CHAPTER 21 SOLUTIONS

### Exercise 1

Comment out the explicit creation of the two books and replace with code to prompt for a new title and author such as shown in this code:

```
//Book book = new Book { Title = "Beginning Visual C# 2015",
//                        Author = "Perkins, Reid, and Hammer" };
//db.Books.Add(book);
//book = new Book { Title = "Beginning XML", Author = "Fawcett, Quin, and Ayers" };

        string title;
        string author;
        Book book;

        do
        {
            Console.WriteLine("Title: "); title = Console.ReadLine();
            Console.WriteLine("Author: "); author = Console.ReadLine();

            if (!string.IsNullOrEmpty(author))
            {
                book = new Book { Title = title, Author = author };

                db.Books.Add(book);

                db.SaveChanges();
            }
        } while (!string.IsNullOrEmpty(author));
```

### Exercise 2

Add a test LINQ query to see if a book with same title and author already exists before adding to database. Use code like this:

```
Book book = new Book { Title = "Beginning Visual C# 2015",
                      Author = "Perkins, Reid, and Hammer" };

var testQuery = from b in db.Books
               where b.Title == book.Title && b.Author == book.Author
               select b;

if (testQuery.Count() < 1)
{
    db.Books.Add(book);
    db.SaveChanges();
}
```

## Exercise 3

Modify the generated classes Stock.cs, Store.cs, and BookContext.cs to use the Inventory and Item names, then change the references to these in Program.cs:

```
public partial class Stock
{
    ...
    public virtual Store Store { get; set; }
}
public partial class Store
{
    ...
    public Store()
    {
        Inventory = new HashSet<Stock>();
    }
    ...
    public virtual ICollection<Stock> Inventory { get; set; }
}

public partial class BookContext : DbContext
{
    ...
    protected override void OnModelCreating(DbModelBuilder modelBuilder)
    {
        modelBuilder.Entity<Book>()
            .HasMany(e => e.Inventory)
            .WithOptional(e => e.Item)
            .HasForeignKey(e => e.Item_Code);

        modelBuilder.Entity<Store>()
            .HasMany(e => e.Inventory)
            .WithOptional(e => e.Store)
            .HasForeignKey(e => e.Store_StoreId);
    }
}

class Program
{
    static void Main(string[] args)
    {
        using (var db = new BookContext())
        {
            var query = from store in db.Stores
                        orderby store.Name
                        select store;
```

```
foreach (var s in query)
{
    XElement storeElement = new XElement("store",
        new XAttribute("name", s.Name),
        new XAttribute("address", s.Address),
        from stock in s.Inventory
        select new XElement("stock",
            new XAttribute("StockID", stock.StockId),
            new XAttribute("onHand",
                stock.OnHand),
            new XAttribute("onOrder",
                stock.OnOrder),
            new XElement("book",
                new XAttribute("title",
                    stock.Item.Title),
                new XAttribute("author",
                    stock.Item.Author)
            )// end book
        ) // end stock
    ); // end store
    Console.WriteLine(storeElement);
}
```

## Exercise 4

Use the following code:

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.Data.Entity;
using System.ComponentModel.DataAnnotations;

namespace BegVCSharp_21_Exercise4_GhostStories
{
    public class Story
    {
        [Key]
        public int StoryID { get; set; }
        public string Title { get; set; }
        public Author Author { get; set; }
        public string Rating { get; set; }
    }

    public class Author
    {
        [Key]
        public int AuthorId { get; set; }
        public string Name { get; set; }
        public string Nationality { get; set; }
    }

    public class StoryContext : DbContext
    {
        public DbSet<Author> Authors { get; set; }
```

```
        public DbSet<Story> Stories { get; set; }
    }

    class Program
    {
        static void Main(string[] args)
        {
            using (var db = new StoryContext())
            {
                Author author1 = new Author
                {
                    Name = "Henry James",
                    Nationality = "American"
                };
                Story story1 = new Story
                {
                    Title = "The Turn of the Screw",
                    Author = author1,
                    Rating = "a bit dull"
                };
                db.Stories.Add(story1);

                db.SaveChanges();

                var query = from story in db.Stories
                           orderby story.Title
                           select story;

                Console.WriteLine("Ghost Stories:");
                Console.WriteLine();
                foreach (var story in query)
                {
                    Console.WriteLine(story.Title);
                    Console.WriteLine();
                }

                Console.WriteLine("Press a key to exit...");
                Console.ReadKey();
            }
        }
    }
}
```

## CHAPTER 22 SOLUTIONS

### Exercise 1

All of the above.

### Exercise 2

You would implement a data contract, with the `DataContractAttribute` and `DataMemberAttribute` attributes.

## Exercise 3

Use the .svc extension.

## Exercise 4

That is one way of doing things, but it is usually easier to put all your WCF configuration in a separate configuration file, either web.config or app.config.

## Exercise 5

```
[ServiceContract]
public interface IMusicPlayer
{
    [OperationContract(IsOneWay=true)]
    void Play();
    [OperationContract(IsOneWay=true)]
    void Stop();
    [OperationContract]
    TrackInformation GetCurrentTrackInformation();
}
```

You would also want a data contract to encapsulate track information; TrackInformation in the preceding code.

# CHAPTER 23 SOLUTIONS

## Exercise 1

1. Modify the XAML of the page BlankPage1 like this:

```
<Page
    x:Class="BasicNavigation.BlankPage1"
    xmlns="http://schemas.microsoft.com/winfx/2006/xaml/presentation"
    xmlns:x="http://schemas.microsoft.com/winfx/2006/xaml"
    xmlns:local="using:BasicNavigation"
    xmlns:d="http://schemas.microsoft.com/expression/blend/2008"
    xmlns:mc="http://schemas.openxmlformats.org/markup-compatibility/2006"
    mc:Ignorable="d" Loaded="Page_Loaded">

    <Grid Background="{ThemeResource ApplicationPageBackgroundThemeBrush}">
        <CommandBar>
            <AppBarToggleButton x:Name="toggleButtonBold" Icon="Bold" Label="Bold"
                Click="AppBarToggleButtonBold_Click" />
            <AppBarSeparator />
            <AppBarButton Icon="Back" Label="Back" Click="buttonGoBack_Click"/>
            <AppBarButton Icon="Forward" Label="Forward" Click="AppBarButtonForward_Click"/>
        </CommandBar>
        <SecondaryCommands>
            <AppBarButton Icon="Camera" Label="Take picture" />
        </SecondaryCommands>
    </Grid>

```

```

        <AppBarButton Icon="Help" Label="Help" />
    </CommandBar.SecondaryCommands>
</CommandBar>

        <TextBlock x:Name="textBlockCaption" Text="Page 1"
HorizontalAlignment="Center" Margin="10,50,10,10" VerticalAlignment="Top"/>
        <StackPanel Orientation="Horizontal" Grid.Row="1"
HorizontalAlignment="Center" VerticalAlignment="Bottom">
            <Button Content="Page 2" Click="buttonGoto2_Click" />
            <Button Content="Page 3" Click="buttonGoto3_Click" />
            <Button Content="Back" Click="buttonGoBack_Click" />
        </StackPanel>
        <WebView x:Name="webViewControl" HorizontalAlignment="Stretch"
Margin="0,75,0,40" VerticalAlignment="Stretch" />
    </Grid>

</Page>

```

2. Go to the code-behind and add these lines:

```

webViewControl.Navigate(new Uri("http://www.wrox.com"));
Application.Current.Resuming += (sender, o) => webViewControl.Navigate(new
Uri("http://www.amazon.com/BEGINNING-VISUAL-C-2015-PROGRAMMING/dp/1119096685/ref
=sr_1_1?ie=UTF8&qid=1444947234&sr=8-1&keywords=beginning+visual+c%23+2015"));

```

## Exercise 2

You specify which capabilities the app has in the `Package.appxmanifest` file on the Capabilities tab. In order to avoid getting an `UnauthorizedAccessException` when you access the microphone, you must ensure that the Microphone capability is checked.

## Exercise 3

1. Create a new Universal app project.
2. Drag a Pivot control onto the design view.
3. Change the first PivotItem like this:

```

<PivotItem Header="Wrox Homepage">
    <Grid>
        <WebView Name="WebViewControl" />
    </Grid>
</PivotItem>

```

4. Change the second PivotItem like this:

```

<PivotItem Header="Hello Pivot!">
    <Grid>
        <TextBlock Text="Hello Pivot!" HorizontalAlignment="Center"
VerticalAlignment="Center" />
    </Grid>
</PivotItem>

```

5. Add a third PivotItem like this:

```
<PivotItem Header="Wrox Logo">
  <Grid>
    <Image Source="http://media.wiley.com/assets/253/59/wrox_logo.gif" />
  </Grid>
</PivotItem>
```

6. Finally, navigate the web view control to a page you choose by calling Navigate in the constructor of the page:

```
WebViewControl.Navigate(new Uri("http://www.wrox.com"));
```

# INDEX

## Symbols

\ (backslash), strings, 564  
. (period), 48  
, (commas), 96, 112–113  
“ “ (double quotation marks), 38  
( ) parentheses, 48–49, 110  
?: operator, 68  
-- operator, 44  
- operator, 42  
/ operator, 43  
- operator, 43  
? operator, 306–307  
:: operator, 342–343  
\\" (backslashes), 38, 41, 564  
|| operator, 55–56  
| = operator, 56  
+ operator, 42–43  
++ operator, 44, 67  
+= operator, 47  
-= operator, 47  
/= operator, 48  
!= operator, 54  
= operator, 38, 47, 86  
== operator, 54, 379  
% operator, 43  
%= operator, 48  
&& operator, 55–56  
&= operator, 56  
/\* characters, comments, 31–32  
/\* characters, comments, 31–32  
\* operator, 42

\*= operator, 48  
@ (at symbol), 39, 564  
^ = operator, 56  
{ } (curly braces), JSON, 595  
< operator, 54  
<= operator, 54  
<> (angle brackets), generic types, 302  
> operator, 54  
>= operator, 54

## A

About windows, 433–434, 436–439  
absolute path names, 566–567, 590, 598  
abstract classes, 174, 188–190  
abstract classes  
declaring, 188–189  
inheritance and, 174  
vs. interfaces, 209–212  
abstract keyword

accumulators, as lambda expressions, 399–401  
adaptive displays, 714–717  
adaptive triggers, 715  
Add New Item Wizard, adding classes, 203–204  
Add Service Reference tool, 690  
addresses, using WCF, 681, 690  
ADO.NET (Active Data Objects .NET)  
    Entity Data Model, 673  
    Entity Framework built on, 654  
    writing applications with C#, 10  
Advanced Build, overflow settings, 82  
advanced method parameters  
    example, 387–390  
    named parameters, 386–387  
    optional parameters, 385–386  
    overview of, 384–385  
aggregate operators, LINQ, 637–641, 650  
aliases, namespaces and, 342–343  
alignment, control property, 423–424  
Amazon AWS, cloud options, 516–518  
angle brackets (<>), generic types, 302  
animations  
    controlling with WCF service, 706  
    overview of, 475  
    timelines with key frames, 476–477  
    timelines without key frames, 475–476  
anonymous methods  
    creating, 357  
    defined, 403  
    lambda expressions for, 392–393  
    overview of, 391–392  
anonymous types  
    defined, 404  
    example using, 378–380  
    overview of, 376–378  
API (application programming interface). *See* web API  
AppBarButton, 730  
AppBarButtonBack, 730  
AppBarToggleButton, 730  
AppControlService, 706  
AppendChild() method, inserting nodes, 605–607  
application programming interface. *See* web API  
ApplicationException class, 343  
applications  
    writing with C#, 9–10  
    writing with .NET Framework, 5  
apps  
    developing, 713–714  
    sandboxed apps, 722  
    universal. *See* Universal Apps  
AppState class, 732  
args parameter, Main(), 125  
arguments  
    Main(), and command-line, 125–127  
    parameters vs., 112  
ArrayList class, 255–258, 292, 294–295  
arrays. *See also* collections  
    of arrays, 98–99  
    declaring, 93–94  
    defined, 77  
    foreach loops and, 95–96  
    how it works, 95  
    iteration of, 256–257  
    multidimensional, 96–98  
    overview of, 92–93  
    parameter, 114–116  
    review, 105  
    using, 94  
    vs. collections, 253–255  
as operator, 297–298, 300  
ASP.NET  
    consuming web API from web site, 547–551  
    creating site that deals two hands of cards, 532–537  
    creating site that uses a storage container, 530–532  
    creating web API, 540–543  
    deploying web API, 544–546

scaling web API at specific time, 554–556  
 scaling web API based on CPU usage, 552–554  
 scaling web API to user requirements, 551–552  
 assemblies, CIL code stored in, 6  
 assertions, entering break mode with, 146–147  
 assignment operators  
   Boolean, 56  
   inability to overload, 283  
   types of, 47–48  
 async programming, Universal Apps, 722  
 asynchronous file access, 581  
 attached property, WPF controls, 416–417  
 attributes  
   creating, 367–368  
   DataContract, 723  
   defined, 403  
   OptionalAttribute, 386  
   overview of, 365–366  
   reading, 366–367  
   Serializable, 723  
   ValueConversionAttribute, 472–473  
   WCF contract, 688–689  
   XML, 594  
 Auto Scaling, 552–554  
 auto-completion, of statements, 102–104  
 automatic properties, class members, 226–227  
 await keyword, 154

**B**

backslash (\), 564  
 backslashes (\\\), 38, 41  
 badges  
   adding, 733–734  
   Universal apps, 713  
   Windows Store apps, 733  
 base arrays  
   declaring arrays, 93–94  
   defined, 93  
   using foreach loops, 95–96

base keyword  
 constructor execution sequence, 198–199  
 member hiding and overrides, 229, 250  
*Beginning XML* (Fawcett), 594  
 behaviors, WCF, 684, 689  
 best practices, cloud computing, 519–520  
 binary operators  
   Boolean assignment/bitwise operators as, 56  
   Boolean comparison operators as, 54–55  
   defined, 42  
   mathematical operators as, 42  
   overloading, 281–283  
 bindings  
   addresses, endpoints, and, 681–682  
   WCF, 679  
   WCF contract, 700  
   WCF service, 690–691  
 block-structured language, basic C# syntax, 30–31  
 bool type  
   as Boolean type, 36  
   no implicit conversion of, 79  
   overview of, 36–37  
   storing result of comparison, 54  
 Boole, George, 54  
 Boolean  
   bitwise and assignment operators, 56–58  
   bool type. *See* bool type  
   comparison operators, 54–55  
   conditional Boolean operators, 55–56  
   overview of, 54  
   review, 75  
 Border control, 424, 427–428  
 Box, SaaS, 517  
 boxing, value types, 275–277  
 branching  
   defined, 53  
   with if statement, 59–63  
   overview of, 59  
   review, 75

- branching (*continued*)  
with `switch` statement, 63–66  
with ternary operator, 59
- break mode  
breakpoints in, 145–146  
Call Stack window and, 152  
debugging in, 144–145  
Immediate and Command windows, 151–152  
monitoring variable content, 148–150  
options for entering, 146–147  
stepping through code, 150–151
- break statements  
exiting infinite loops with, 73  
exiting loops with, 72  
terminating `switch` statement with, 64–65
- breakpoints  
adding, 145–146  
viewing tracepoints with, 143
- bubbling events, 419
- Button control  
as content controls, 458  
in game client example, 435–436  
looping through all nodes in XML document, 600–603  
name property, 455  
styling, 470–471  
as UI control, 459
- `buttonCreateNode_Click()`, 606–607
- `buttonExecute_Click()`, 614
- buttons  
creating desktop applications, 22–26  
example of, 183–184
- `buttonXMLtoJSON`, 610
- byte arrays  
`FileStream` class operating on, 568  
LINQ to XML constructors, 622  
reading data, 571–573  
writing data, 574–575
- byte type  
in enumerations, 86, 88–89  
explicit conversions of, 80–81
- implicit numeric conversions of, 79  
as integer type, 35  
using structs, 91–92

## C

- C#  
basic syntax, 30–34  
creating storage container using Azure, 520–521  
.NET. *See .NET Framework*  
review, 12  
in this book, 10  
understanding, 8–9  
unmanaged code and, 6  
Visual Studio 2015 and, 10–12  
writing applications in, 9–10  
[wrox.com](#) downloads for, 3
- C#, writing program in  
console applications, 17–20  
desktop applications, 22–26  
Error list window, 22  
overview of, 13–14  
Properties window, 21–22  
review, 27  
Solution Explorer, 20–21  
Visual Studio 2015 development environment, 14–17  
[wrox.com](#) downloads for, 13
- C++ language, 8–9
- Call Hierarchy view, 202
- Call Hierarchy window, 248
- Call Stack window, error handling, 152
- callback methods, games, 504–505, 507–508
- Canvas control  
completing game application, 503  
content layout controls, 422  
creating About window, 436–438  
designing Options window, 444  
as layout control, 459  
overview of, 424–425
- case sensitivity, C# syntax, 32

case statements, 64–65  
casting conversions, 80  
catch keyword  
    error handling with, 153–155  
    exception handling example, 156–160  
char array, *FileStream*, 571–572, 574  
char type  
    implicit numeric conversions of, 79  
    string manipulation, 99–101  
    as text type, 36  
CheckBox control  
    creating Options window, 440–441  
    designing Options window, 445  
    as UI control, 459  
checked keyword, 81–82  
CIL (Common Intermediate Language) code, 5–6, 7  
Circuit Breaker pattern, cloud programming, 520  
Class Details window, class diagram, 206  
class diagrams, 204, 205–206  
class families, 216  
class keyword, 188–189  
class libraries, 206–209  
class library example  
    adding Card class, 242–243  
    adding client console, 246–247  
    adding Deck class, 243–246  
    adding Suit and Rank enumerations, 240–242  
    planning Card and Deck classes, 238–239  
    writing, 239–240  
class members  
    accessibility modifiers and, 235  
    adding Card class to class library, 242–243  
    adding client console to class library, 246–247  
    adding Deck class to class library, 243–246  
    adding Suit and Rank enumerations in class library, 240–242  
    applying fields, methods, and properties, 223–225  
    automatic properties, 226–227  
Call Hierarchy window, 248  
calling overridden/hidden base class methods, 229  
defining fields, 218–219  
defining methods, 219–220  
defining nested types, 230–232  
defining partial methods, 237–238  
defining properties, 220–222  
definition, 218  
explicit interface members, 234–235  
hiding base class methods, 227–229  
interfaces, 232–234  
overview of, 217  
partial definitions, 235–237  
planning Card and Deck classes in class library, 238–239  
refactoring, 225–226  
review, 248–250  
this keyword, 230  
writing class library application, 239–240  
Class View window, 20–21, 200–203  
classes  
    adding, 203–204  
    class diagrams, 204–206  
    class libraries, 206–209  
    Class View window, 200–202  
    collections, 178  
    common DOM, 597–598  
    constructor execution sequence, 196–200  
    constructors/destructors, 168–169  
    containment, 177–178  
    defining constructors and destructors, 195–196  
    defining in C#, 188–190  
    defining interfaces, 190–191  
    exercise defining, 191–192  
    implementing interfaces, 233–234  
    inheritance, 172–175  
    inheriting from `System.Object`, 193–195  
    interfaces vs. abstract, 209–212  
    Object Browser working with, 202–203  
    object types by, 165  
    operator overloading, 179

classes (*continued*)  
overview of, 187  
polymorphism, 175–177  
reference types and, 180  
review, 216  
shallow copying vs. deep copying, 214–215  
structs vs., 212–214  
classes, file  
file system access, 567–568  
input/output, 562–567  
monitoring file system, 584–588  
review, 590  
streams. *See streams*  
click events, event handlers for, 727–728  
`Click Me` button, WPF, 25–26  
clients  
adding client console, 246–247  
proxy class and, 701, 707, 709  
WCF test, 691–693  
cloud optimized stack, 538  
cloud programming, advanced  
consuming web API from web site, 547–551  
creating web API, 540–543  
deploying web API, 544–546  
overview of, 539  
review, 556–557  
scaling web API at specific time, 554–556  
scaling web API based on CPU usage, 552–554  
scaling web API to user requirements, 551–552  
cloud programming, basic  
best practices, 519–520  
creating storage accounts, 521–522  
creating storage container, 520–521, 523–530  
creating web site that deals two hands of cards, 532–537  
creating web site that uses storage container, 530–532  
overview of, 515–519  
review, 538  
writing applications with C#, 9

CLR (Common Language Runtime), .NET  
defined, 5  
managed code controlled by, 6  
writing .NET application, 8

code  
decorating with attributes, 365  
separation of concerns, 409  
stepping through, 150–151

code blocks, variable scope, 122

Code First  
creating database objects, 654–661  
navigating database relationships, 663  
working with database objects, 653

code outlining, 236

code-behind files  
adding methods to, 465  
completing game application, 503, 506  
separation of concerns and, 409  
XAML, 411

coercion, enforcing properties, 416

collection classes, 252

`CollectionBase` class, 258–259

`CollectionDataContract`, 726

collections  
adding cards collection to `CardLib`, 262–264  
arrays vs., 253–255  
deep copying, 271–274  
defining, 258–259  
example animals, 260–262  
how they work, 255–258  
`IDictionary` for keyed, 264–266  
indexers and, 259–260  
initializers, 371–372  
iterators and, 266–271  
lambda expressions used with, 399–401  
`ObservableCollection`, 450, 454, 458  
overview of, 178–179, 251–253  
review, 298–300  
sorting, 292–295

using, 253  
 using generic collection class with CardLib, 321  
 columns, Grid control, 431–433  
 ComboBox control  
   Options window, 442–443, 445  
   as UI control, 459  
 Command and Query Responsibility Segregation (CQRS), 519  
 command bars, 729–730  
 Command window, manual Visual Studio operations in, 151–152  
 command-line applications, console, 17–20, 27  
 command-line parameters, 125–127, 134  
 commands  
   InputGestures, 501  
   routed commands, 419–422  
 commas (,), 96, 112–113  
 comments, 31, 33–34  
 Common Intermediate Language (CIL) code, 5–6, 7  
 Common Language Runtime. *See* CLR (Common Language Runtime), .NET  
 Common Type System (CTS), 5  
 communication payload, WCF, 680, 709  
 communication protocols, WCF, 680–681  
 comparison operators, 54–55, 283  
 comparisons  
   adding operator overloads to CardLib, 284–289  
   bool type storing results of, 54  
   `IComparable/IComparer` interfaces, 290–291  
   with `if` statement, 60  
   `is` operator and, 277–279  
   of object references with `==` operator, 379  
   operator overloading and, 280–284  
   overview of, 275  
   review, 298–300  
   with `switch` statement, 63–66  
   with ternary operator, 59  
   type comparisons, 275–277  
   value comparisons, 279–280  
   what they do, 252

compiling code  
 creating console application, 19  
 writing .NET application, 5, 8  
 complex variable types  
   arrays. *See* arrays  
   enumerations, 85–89  
   structs, 89–92  
 compression, reading and writing files, 562, 581–583  
`CompressionMode.Compress` enumeration, 581, 583  
 condition  
   for LINQ queries, 628  
   `for` loop structure, 71  
 Condition option, Breakpoints window, 146  
 conditional operator  
   Boolean, 55–56  
   branching with, 59  
   common usage of, 68  
   defined, 42  
 console applications  
   basic C# structure for, 33–34  
   creating simple, 17–20  
   defined, 14  
   using Solution Explorer, 20–21  
`Console.ReadLine()` command, 45–46  
`Console.WriteLine()` command, 45–47  
 const keyword  
   declaring arrays, 94  
   for global variables, 120–121  
   looping through all nodes in XML document, 601  
 constraints, generic type, 324–326  
 constructor initializer, 198–199  
 constructors  
   defining, 195–196  
   execution sequence, 196–200  
   LINQ to XML, 621  
   of objects, 168–169  
   static, 170

consuming web API, 547–551  
containment, 177–178  
Content controls, 422, 458  
content presenters, for templates, 468  
ContentPresenter control, for templates, 468  
continue command, interrupting loops, 72–73  
contracts, WCF  
    creating, 697–700  
    data contracts, 694  
    fault contracts, 696  
    function of, 679  
    message contracts, 696  
    operation contracts, 695–696  
    service contracts, 695  
    types of, 683–684  
contravariance, generic classes and, 336–337  
control event types, 417  
controllers, Web API 2 Controller, 542  
controls, Toolbox window UI, 24  
controls, user  
    adding to game application, 481–488  
    implementing dependency properties, 478–481  
    overview of, 478  
controls, WPF  
    adding to window, 413  
    layout, 422–423  
    overview of, 412  
    properties, 413–416  
    specifying positions for, 715  
    stack order of, 423  
    styles and templates applied to, 467–471  
    types of, 422  
    UI, 434–436  
ControlTemplate class, 468  
conversions  
    explicit, 80–83  
    how it works, 83–84  
    implicit, 78–80  
    as operator in, 297–298  
    overloading conversion operators, 295–297  
        overview of, 295  
        understanding, 252  
convert commands, explicit conversions, 83  
Convert.ToDouble(), 45–46, 51  
Convert.ToInt32(), 57  
Convert.ToString(), 89  
covariance, generic classes and, 336  
CPU usage, scaling web API, 551–554  
CQRS (Command and Query Responsibility Segregation), 519  
Create Schema menu option, XML, 596  
CreateAttribute() method, 595–596, 605  
CreateComment() method, nodes, 605–607  
CreateElement() method, nodes, 605–607  
CreateNode() method, 605–606  
CreateTextNode() method, 605–607  
.cs file extension, 21  
Ctrl+Shift+N, creating new project, 19  
CTS (Common Type System), 5  
curly braces ({}), JSON, 595  
custom exceptions  
    adding to CardLib, 343–345  
    defined, 403  
    overview of, 343

## D

data binding  
    defined, 458  
    dynamic, 450–453  
    to local objects, 449  
    overview of, 448  
    static, 449–450  
    with user control, 487–488  
data contracts, WCF  
    creating, 699  
    creating service contracts, 687–688  
    defined, 684  
    defining for service, 694–700  
data languages, XML and JSON as, 594

data sets, querying large, 635–637  
data sources  
defining with `DataContext` controls, 448  
specifying for LINQ queries, 628  
data templates. *See* templates  
data types  
immutable, 227  
in .NET Framework, 4–5  
object properties, fields and, 167  
reference vs. value, 180  
simple variable, 34–39  
databases  
connecting to, 661–662  
creating and querying XML from existing, 670–674  
creating using Code First, 654–661  
Entity Framework and, 653–654  
handling migrations, 669–670  
installing SQL Server Express, 653  
navigating relationships, 662–669  
overview of, 652  
review, 674–675  
using, 652–653  
`DataContext` control, data binding, 448  
`DataContract` attribute, 723  
.ddl file extension, libraries, 6  
debug builds  
breakpoints in, 145  
outputting debugging information, 138  
symbolic information in, 136  
`Debug` command, 144  
`Debug` toolbar, 144–145, 150–151  
`Debug` Windows menu, 151–152  
debugging, in Visual Studio  
assertions, 146–147  
in break mode, 144–145  
breakpoint use, 145–146  
Call Stack window, 152  
console applications, 19  
diagnostic output vs. tracepoints, 144  
in Error List window, 22  
Immediate and Command windows, 151–152  
monitoring variable content, 148–150  
in nonbreak (normal) mode, 136–137  
options for entering break mode, 146–147  
outputting debugging information, 137–138  
overview of, 136  
review, 162  
stepping through code, 150–151  
tracepoint use, 142–143  
writing text to output window, 138–142  
`Debug.WriteLine()`, 137–138  
decimal type, 36  
declared variables, 34, 37–39  
`Decoder` class, reading data using `FileStream`, 571–573  
decorating code, with attributes, 365  
decrement (--) operator, 44  
deep copying  
adding to `CardLib`, 273–274  
collections, 271–273  
using `Clone()`, 272–273  
vs. shallow copying, 214–215  
default keyword, generic classes, 324  
deferred execution, LINQ, 629, 650  
`DeflateStream` class  
defined, 562  
reading and writing compressed files, 581, 583  
`delegate` keyword, 130–133  
delegates  
calling functions through, 130–133  
defining for use with events, 352  
defining generic, 334  
lambda expressions as, 398  
multipurpose event handlers and, 353  
specifying event restrictions, 345–346  
deleting nodes, 607–609  
dependency properties  
adding to user control, 483, 488  
completing game application, 503, 506–507

- dependency properties (*continued*)
- defined, 412
  - features of, 416
  - implementing, 478–481
  - options for adding, 502
  - deploying web API, to cloud, 544–546
  - derived class
    - inheriting abstract classes and, 210
    - inheriting from parent class, 172–175
    - polymorphism and, 175–176
    - relationships between objects, 177–178
  - design, Universal Apps, 712–713
- Design View
- adding event handlers, 447
  - manipulating control properties, 413–416
  - WPF, 411–412
- desktop applications
- defined, 14
  - writing with C#, 9, 22–26
- destructors, 168–169, 195–196
- development, Windows Universal Apps, 711, 713–714
- Device Preview panel, 714–717
- diagnostic output
- outputting debugging information, 137–138
  - tracepoints vs., 144
  - writing text to output window, 138–142
- dictionaries
- `Dictionary<K, V>` interface, 311, 319–320
  - `IDictionary` interface, 264–266
  - review, 300
- dimension property, 423–424
- directories, monitoring file system, 588
- `Directory` class, 562, 563–564
- `DirectoryInfo` class, 562, 566
- discrete key frame, 477
- disk access, Universal Apps, 722–726
- disk space, scaling web API, 551–552
- displays
- adaptive, 714–717
- flip view for, 717–721
- .dll assemblies, class libraries, 206
- DLL Hell, 540
- DLR (Dynamic Language Runtime), 381
- do loops, 66–69
- `DockPanel` control
- completing game application, 508–509
  - content layout controls, 422
  - creating main window, 463
  - as layout control, 459
  - overview of, 426–428
- documentation, creating C#, 32
- domain model, refactoring, 489–494
- double quotation marks (“ “), 38
- double type, 36
- double values, 128–131
- dynamic data binding, 450–453
- dynamic keyword, variables, 381
- Dynamic Language Runtime (DLR), 381
- dynamic lookup
- defined, 404
  - dynamic types, 381–384
  - overview of, 380–381
- dynamic types, 380–381
- dynamic variables, 380
- E**
- editor features, WPF, 411–412
- elements
- array entries as, 93
  - XML, 594
- `Elements.xml`, 611–615
- Elvis operator (?), 306–307
- `Encoder` object, writing data using `FileStream`, 574
- endpoints, WCF
- creating WCF service, 690

default bindings for, 683  
for each operation, 679, 682  
`#endregion` keyword, 34  
Entity Data Model, ADO.NET, 673  
Entity Framework  
    adding with NuGet Package Manager, 663  
    automatic creation of LINQ objects, 662  
    creating database objects using Code First, 654–656  
    creating local server instance of database, 661  
    handling database migrations, 669–670  
    managing database details, 667  
    overview of, 653–654  
entity-relationship model, 653  
`enum` keyword, enumerations, 85–88  
enumerations  
    defined, 77  
    overview of, 85–89  
    review, 105  
`Enum.Parse()`, 89  
`Equals()` method, `System.Object` methods, 193  
error handling  
    Azure storage account, 527–528  
    exception handling example, 156–160  
    listing and configuring exceptions, 160–161  
    overview of, 153  
    review, 162  
    `try.catch.finally`, 153–156  
Error List window, Visual Studio 2015  
    defined, 17  
    displaying, 137  
    writing C# program, 22  
escape sequences  
    for `string` literals, 40–41  
    using, 38  
event handlers  
    adding to app pages, 729  
    adding to code, 180  
    anonymous methods and, 391–392  
    in CardLib, 357–365  
    for click events, 727–728  
    completing game application, 506, 508  
    defined, 403  
    EventHandler and Generic EventHandler<T>, 356  
    example using, 347–349  
    for game applications, 502  
    looping through all nodes in XML document, 600–601  
    monitoring file systems, 584, 586–587  
    multipurpose event handlers, 353–356  
    in Options window, 446–448  
    raising events and, 345–346  
    registering for, 352  
    return values and, 356–357  
    self-hosted WCF services and, 707  
    for user control, 483–484  
    WPF controls and, 418–419  
event-driven applications, 180  
events  
    anonymous methods, 357  
    in CardLib, 357–365  
    defined, 403  
    defining, 350–353  
    EventHandler and Generic EventHandler<T>, 356  
    handling, 347–349  
    managing state, 730  
    multipurpose event handlers, 353–356  
    objects raising (consuming), 180  
    overview of, 345–347  
    raising in game application, 490  
    reasons for using routed commands in place of, 462–463  
    return values and event handlers, 356–357  
    WPF controls, 417–419  
EventTrigger class, 473  
exception handling. *See* error handling

exceptions  
custom, 343–345  
Debug menu settings, 160  
.exe file extension  
creating console application, 19  
defined, 6  
hosting WCF services, 685  
explicit conversions, 80–83, 88  
explicit syntax. *See* method (explicit) syntax,  
LINQ  
expression trees, lambda expressions as, 398  
expression-bodied methods, functions, 112  
expressions  
assignment operators, 47–48  
basic C# syntax, 30–34  
evaluating/testing, 151–152  
lambda. *See* lambda expressions  
manipulating data with, 30  
mathematical operators, 42–47  
namespaces, 49–51  
operator precedence, 48–49  
overflow checking of, 81–82  
properties, 222  
review, 51–52  
understanding, 42  
Extensible Application Markup Language. *See*  
XAML (Extensible Application Markup  
Language)  
Extensible Markup Language (XML)  
creating/querying from existing database,  
670–674  
and JSON. *See* XML (Extensible Markup  
Language) and JSON  
extension methods, LINQ, 630  
extern keyword, methods, 219–220

defining for class members, 218–219  
example applying, 223–225  
of objects, 166–167

File class  
creating FileStream objects in, 568  
defined, 562  
static methods of, 563

file classes  
file system access, 567–568  
input/output, 562–567  
monitoring file system, 584–588  
review, 590  
streams. *See* streams

File menu, 465–466  
file picker contracts, 722  
file pointer, FileStream class, 570, 572  
FileAccess enumeration members, 568–569

FileInfo class  
creating FileStream objects, 568  
defined, 562  
FileSystemInfo properties of, 565–566  
overview of, 564–566

FileMode enumeration members, 569

filename  
creating FileStream object, 568  
verbatim string literals in, 41

files  
classes for input and output, 562–567  
creating StreamWriter object from, 575  
monitoring directories and, 584–588  
overview of, 561  
reading and writing compressed, 581–583  
review, 588–589  
streams. *See* streams

FileStream object  
asynchronous file access, 581  
creating StreamWriter object, 575  
file position, 570  
overview of, 568–570  
reading and writing compressed files, 581, 583  
reading data, 570–573

## F

fatal errors, 135  
fault contracts, WCF, 683, 696  
fields

reading data with `StreamReader`, 577–578  
 review, 590  
 writing data, 573–575  
`FileSystemInfo` class, 562, 565–566  
`FileSystemWatcher` class, 562, 584–588  
`finally` keyword  
 error handling, 153–155  
 exception handling example, 156–160  
`FirstChild` property, `XmlElement`, 599  
 flip view, handhelds, 717–721  
`float` type, 36, 79  
 floating-point types, 36  
 flow control  
   Boolean bitwise and assignment operators, 56–58  
   branching, 59–66  
   looping. *See* looping  
   operator precedence and, 58–59  
   overview of, 53  
   review, 73–75  
   using Boolean logic, 54–56  
`for` loops, 71–72, 95  
`foreach` loops  
   addressing elements in arrays, 95–96  
   iterating through arrays, 256–257  
   iterating through query results, 629  
   iterators and, 266–267  
   statement auto-completion in Visual Studio, 104  
   string manipulation, 100  
   using with jagged arrays, 98  
   using with multidimensional arrays, 97  
`FormatText` method, looping through all nodes in XML, 602–603  
`FrameworkPropertyMetadata` constructor, overloading, 480–481  
`from` clause, query syntax, 627, 628  
 fully qualified names, 121  
 functional construction, LINQ to XML, 619  
 functions  
   defining and using simple, 108–110  
`Main()`, 125–127

as members of `struct` types, 127–128  
 overview of, 107–108  
 parameters, 112–119  
 return values, 110–112  
 review, 133–134  
 using delegates, 130–133  
 using overloading of, 128–130  
 variable scope. *See* variable scope

## G

GAC (global assembly cache), placing code in, 6  
 garbage collection, .NET, 6–7  
 generic classes  
   constraining types, 324–326  
   contravariance and, 336–337  
   covariance and, 336  
   `default` keyword and, 324  
   defining, 322–323, 326–330  
   `Dictionary<K, V>` interface, 319–320  
   generic delegates, 334  
   generic interfaces, 332  
   generic methods, 333–334  
   generic operators, 331–332  
   generic structs, 332  
   inheriting from, 330–331  
   `List<T>` interface, 312, 314–319  
   null coalescing operator (`??`), 305–306  
   null condition operator (`?`), 306–307  
   nullable types, 303–304, 307–311  
   operators and nullable types, 304–305  
   overview of, 301–303  
   review, 337–339  
   sorting and searching generic lists, 313–316  
`System.Collections.Generic`, 311–312  
   using, 303  
   using generic collection class with `CardLib`, 321  
   variance and, 335  
   generic delegates, 334  
`Generic EventHandler<T>`, 356  
   generic interfaces, 332

generic methods, 333–334  
generic operators, 331–332  
generic structs, 332  
generic types, 180  
get keyword, accessor properties, 220, 222  
`GetBytes()`, `FileStream`, 574  
`GetChars()`, `FileStream`, 571–573  
`GetCopy()` method, shallow copying, 271–272  
`GetEnumerator()`, `IEnumerator`, 266–267  
`GetHashCode()`, `System.Object`, 194  
`GetType()`  
    `System.Object`, 194  
type comparisons, 275  
    using Boolean operators, 57  
global assembly cache (GAC), placing code in, 6  
global namespace  
    `::` operator, 342–343  
C# code contained in, 48  
global variables  
    local vs., 120–121  
    parameters and return values vs., 123–125  
    unsuitable for general purpose functions, 122  
GNU ZIP algorithm (`GZipStream` class), 562, 581–583  
`goto` statement, flow control for `case` statement, 64  
greater-than operator (`>`), overloading, 179  
Grid control  
    content layout controls, 422, 458–459  
    creating About window, 437–438  
    creating main window, 464  
    designing Options window, 443  
    overview of, 430–431  
    using rows and columns, 431–433  
group queries, LINQ, 645–647, 650  
`GZipStream` class (GNU ZIP algorithm), 562, 581–583

## H

handheld device, flip view for, 717–721  
`HasChildNodes` property, `XmlElement`, 600  
Health Endpoint Monitoring pattern, cloud, 520  
`Height`, alignment property, 423–424  
Help menus, About windows, 434  
hidden methods, 227–229  
Hit Count, Breakpoints window, 146  
`HorizontalAlignment`, control property, 423–424  
hosting WCF services, 684–685  
HTTP (HyperText Transport Protocol), 680–681, 683  
hybrid cloud, 516–518

## I

IaaS (Infrastructure as a Service), 517–518  
IBM Cloud, 516–518  
`ICloneable` interface, 214–215, 273–274  
`ICollection` interface, 253, 265  
`IComparable` interface, 290–291, 313–314  
`IComparer` interface  
    comparing with `IComparable`, 290–291  
    sorting and searching generic lists, 313–314  
    sorting collections, 295  
IDEs (integrated development environments), 10–11, 12  
`IDictionary` interface, 264–266  
`IDisposable` interface, 172  
`IEnumerable` interface  
    collections and, 253  
    defining generic interfaces, 332  
`DictionaryBase` class and, 265  
    looping through collections, 256–257  
`IEnumerator` interface, 266–267  
`if` statement, 59–63  
IIS (Internet Information Server)

creating ASP.NET web site, 531  
 hosting WCF services, 684–685  
 running ASP.NET with, 519

**IL** (Microsoft Intermediate Language). *See CIL* (Common Intermediate Language) code

**IList** interface

- ArrayList class and, 257
- collections and, 253
- defining collections, 259–260
- `IndexOf()` method, 258

**Image** control, 434–435, 459

**images**

- creating PNG, 733–734
- uploading to cloud storage account, 527

**Immediate window**, expressions, 151–152

**implicit conversions**, 43, 78–80

**increment operators**, 44

**indentation**, basic C# syntax, 30–31

**indexed lists**, arrays as, 93

**indexers**

- array-like access to collections, 259–260
- IDictionary** for keyed collections, 264–266

**infinite loops**, 73–74

**Infrastructure as a Service (IaaS)**, 517–518

**inheritance**

- from generic classes, 330–331
- hiding base class methods, 227–229
- review, 172–175
- from `System.Object`, 193–195

**initialization**

- array, 93–94
- arrays of arrays, 98
- `FileSystemWatcher`, 586
- `for` loop structure, 71
- multidimensional array, 96–98

**initializers**

- collection initializers, 371–372
- defined, 404
- example using, 372–374

object initializers, 368–371  
 overview of, 368

**INotifyPropertyChanged** interface, 452, 458

**input**

- errors in, 92
- file classes for, 562–563
- reading data from, 577–580

**InputGestures**, 501

**Input/Output** file classes. *See I/O (Input/Output)*

- file classes
- `InsertAfter()`, nodes, 605–607
- inserting nodes, 605–607

**instance members**

- class membership, 186
- members, 169–170

**instances**, class, 165

**instantiation**. *See also* initializers

- `FileInfo` class, 564–566
- of objects, 165

**int array**

- `Main()` returning, 125
- overloading functions, 128–129

**int value**

- implicit numeric conversions of, 79
- as integer type, 36
- returned when multiplying two `short` values, 85
- as underlying type in enumerations, 86, 105

**intArray**, 113–114

**integer types**, 35–36

**integer value**, `System.Convert` class, 57

**integrated development environments (IDEs)**, 10–11, 12

**IntelliSense**, Visual Studio

- creating XML document, 596
- showing available overloads, 128

**interface keyword**, 190–191

**interfaces**

- defining, 190–191
- defining generic interfaces, 332

**interfaces (*continued*)**

disposable objects and, 172  
implementing explicit interface members, 234–235  
implementing for class members, 232–233  
implementing in classes, 233–234  
interface members, 210  
interface polymorphism, 176–177  
`IValueConverter`, 472–473  
overview of, 171–172  
review, 171–172  
`System.Collections`, 252  
vs. abstract classes, 209–212  
internal classes, 188–189  
internal keyword, 218  
interoperability  
    dynamic lookup and, 380–381  
    between languages in .NET, 4–5  
interrupting loops, 72–73  
I/O (Input/Output) file classes  
    absolute vs. relative path names and, 566–567  
    `DirectoryInfo` class, 566  
    `File` and `Directory` classes, 563–564  
    `FileInfo` class, 564–566  
    overview, 562  
`IOException`, 571–572, 576  
`is` operator, 277–279  
Items controls, 422, 458  
iteration, through query results, 629  
iterators  
    collections and, 266–268, 270–271  
    implementing, 268–270  
    review, 300  
`IValueConverter` interface, 472

**J**

JavaScript Object Notation. *See* JSON (JavaScript Object Notation)  
JIT (just-in-time) compilers, 5, 8  
joins, LINQ, 647–648, 650

**JSON (JavaScript Object Notation)**

basics, 594–595  
converting XML to, 609–610  
packages, 616  
parsing file, 547–551  
just-in-time (JIT) compilers, 5, 8

**K**

key frames, and timelines, 475–477  
keys  
    accessing storage account with account, 524–525  
    for keyed collections, 264–266  
key-value pairs, defining collection of, 319–320  
keywords  
    beginning with #, 34  
    overflow checking of expressions, 81–82  
    variable naming and, 39

**L**

Label control  
    About window, 437–438  
    in game application, 435, 503  
    Options window, 444–445  
    as UI control, 459  
lambda expressions  
    anonymous methods and, 391–393  
    collections and, 399–400  
    defined, 404  
    as delegates and expression trees, 398  
    example using, 393–396  
    example with collections, 400–401  
    LINQ and, 631–633, 650  
    overview of, 391  
    parameters, 396  
    review, 402–404  
    statement bodies, 396–398  
landscape layout, changing from portrait to, 714–717

Language Integrated Query. *See LINQ (Language Integrated Query)*

language settings, Error List window, 22  
`LastChild` property, `XmlElement`, 599, 608–609  
layout  
    controls, 422–423, 459  
    Visual Studio environment, 15–16  
lazy evaluation, queries, 629  
libraries  
    class, 206–209  
    creating storage container using Azure C#, 520–521  
exception categories of .NET, 160  
overview of, 4–5  
standard types defined in, 35  
viewing list of, 21  
lifecycle of apps, 713  
line breaks, string variables and, 38  
line numbers, debugging, 22  
linear key frames, 477  
linking, 8  
LINQ (Language Integrated Query)  
    adding code for simple query, 657  
    aggregate operators used with, 637–641  
    creating and printing query results, 665, 669  
    declaring variables, 627–628  
    deferring query execution, 629  
    defined, 650  
    executing simple query, 660–661  
    extensions methods, 630  
    group queries, 645–647  
    iterating through query results, 629  
    joins, 647–648  
    lambda expressions, 631–633  
    LINQ to XML. *See LINQ to XML*  
    method syntax, 629  
    navigating database relationships, 662–663  
    `orderby` clause, 634  
    ordering by multiple levels, 643–645  
    ordering query results, 633–634  
    overview of, 618–619

providers for, 625  
query syntax, 625–627  
query syntax vs. method syntax, 630  
querying databases, 653  
querying large data sets, 635–637  
review, 649–650  
`SELECT DISTINCT` queries, 641–643  
selecting items, 628–629  
specifying condition, 628  
specifying data source, 628  
working with XML fragments, 622–624  
writing applications with C#, 10  
LINQ to Data Set, 625  
LINQ to Entities, 625, 654, 670  
LINQ to JSON, 625  
LINQ to Objects, 625  
LINQ to SQL, 625  
LINQ to XML  
    creating/querying XML from existing database, 670–674  
    overview of, 619–622  
    as type of LINQ provider, 625  
    working with XML fragments, 622–624  
`List<T>` interface  
    as generic collection type, 311  
    how it works, 316–319  
    overview of, 312  
    sorting and searching, 314–316  
    using, 312–313  
`ListBox` control  
    creating start game window, 454–457  
    as item control, 458  
    name property, 455  
    properties, 453  
    as UI control, 459  
lists  
    collection classes for, 252  
    sorting and searching generic, 313–314  
literal values  
    assigning to variables, 38  
    creating expressions, 42

- literal values (*continued*)  
specifying array, 93  
string literals, 40–41  
types of, 39–40  
live tiles, 713  
`LoadCompressedFile()`, 581–583  
local variables  
  global vs., 120–122  
  using identical names for, 124  
lock screen apps, 713  
logic (semantic) errors, 135  
`long` type  
  implicit numeric conversions of, 79  
  as integer type, 36  
  as underlying type in enumerations, 86  
looping  
  converting XML to JSON, 609–610  
  creating nodes, 606  
  defined, 53  
  do loops, 66–69  
  foreach loops, 95–96  
  infinite loops, 73–74  
  interrupting loops, 72–73  
  for loops, 71–72, 95  
  overview of, 66  
  review, 75  
  through all nodes in XML document, 600–603  
  variable scope and, 122–123  
  while loops, 69–71
- M**
- `Main()`  
  as entry point function for console application, 110  
  reading and writing compressed files, 581–583  
  reading data with `StreamReader`, 578  
  understanding variable scope, 119–121  
  using command-line arguments with, 125–127  
  writing data with `StreamWriter`, 576  
main window  
  adding for game application, 462  
  creating, 463–466  
`MainWindow.xaml`, 23–24, 585  
managed code  
  garbage collection and, 6–7  
  writing .NET application, 6, 8  
`Margin` property, 423–424  
markers, comments in C#, 31–32  
matching parameters, 114  
Materialized View pattern, cloud, 519  
mathematical operators  
  increment/decrement, 44  
  manipulating variables with, 45–47  
  simple, 42–43  
  string concatenation operator, 43  
`Maxima()`, 141  
`MaxValue()`, 113–114, 118–119, 128–129  
members  
  abstract class, 210  
  class members. *See* class members  
  object properties, fields and, 167  
  refactoring, 225–226  
`MemberwiseClone()`, 194  
memory, scaling web API, 551–552  
`MenuItem`, 462, 464  
menus  
  adding for game application, 462  
  creating main window, 464  
  developing Universal apps, 712  
  routed commands with, 462–463  
message contracts, WCF, 683, 696  
message patterns, WCF, 684  
Message Transmission Optimization Mechanism (MTOM), 680  
`MessageHandler`, 352  
metadata  
  in assemblies, 6  
  completing game application, 504  
  configuring WCF contracts, 689–690  
  `PropertyMetadata` class, 480

- 
- method (explicit) syntax, LINQ
    - combing with lambda expressions, 631–633
    - lambda expressions as extension of, 397
    - overview of, 629
    - review, 650
  - methods
    - adding to code-behind files, 465
    - advanced parameters of, 385–386
    - anonymous, 357, 391–393
    - for calling overridden or hidden base class, 229
    - for completing game application, 504–506
    - for creating nodes, 605
    - defining for class members, 219–220
    - `DirectoryInfo` class, 566
    - example applying, 223–225
    - exposed by objects, 167–168
    - `File` and `Directory` classes, 563
    - `FileInfo` class, 564–565
    - generic, 333–334
    - hiding base class methods, 227–229
    - for inserting nodes, 605–606
    - .NET functions vs., 108
    - partial, 237–238
    - using `this` keyword, 230
  - Microsoft Azure
    - cloud options, 516–518
    - consuming ASP.NET web API from web site, 547–551
    - creating storage accounts, 521–522
    - creating storage container using C# libraries, 520–521
    - creating storage container using storage client library, 523–530
    - deploying ASP.NET web API, 544–546
    - scaling web API at specific time, 554–556
    - scaling web API based on CPU usage, 552–554
    - scaling web API to user requirements, 551–552
  - Microsoft Intermediate Language (IL). *See CIL (Common Intermediate Language)* code
  - Microsoft Intermediate Language (MSIL). *See CIL (Common Intermediate Language)* code
  - Microsoft Message Queuing (MSMQ), 681, 683
  - migrations, handling database, 669–670
  - Model, View, Controller (MVC), 540
  - models
    - creating view model for game application, 494–502
    - refactoring domain, 489–494
  - Model-View-ViewModel (MVVM), 489
  - modules, .NET library, 4
  - monitoring
    - file system, 584–588
    - variable content, 148–150
  - Mono
    - C# version, 519
    - open-source .NET, 4
  - MSIL (Microsoft Intermediate Language). *See CIL (Common Intermediate Language)* code
  - MSMQ (Microsoft Message Queuing), 681, 683
  - MTOM (Message Transmission Optimization Mechanism), 680
  - multidimensional arrays, 96–98
  - multipurpose code, functions for, 108
  - multipurpose event handlers, 353–356
  - MVC (Model, View, Controller), 540
  - MVVM (Model-View-ViewModel), 489

## N

- name property, WPF, 455
- named method parameters, 404
- named parameters, 386–390
- Named Pipe, 680–681, 683
- namespace keyword, 48
- namespace qualifiers, 403
- namespaces
  - global namespace qualifier, 342–343
  - review, 51–52
  - XAML, 410–411

naming conventions, variable, 38  
native code, C#, 5, 8  
navigation  
  database relationships, 662–669  
  between pages of apps, 726–728  
  in Windows Universal apps, 737  
nested blocks  
  basic C# syntax, 31  
  variable scope and, 122  
nested object initializers, 370–371  
nested types, defining class members, 230–232  
.NET Framework  
  review, 12  
  understanding, 4  
  what it consists of, 4–5  
  writing applications with, 5–8  
new keyword  
  implementing interfaces, 233–234  
  initializing arrays, 93  
New Project, Visual Studio, 17–18  
`NextSibling` property, `XmlElement`, 600  
node values, changing  
  creating nodes, 606–607  
  deleting nodes, 607–609  
  inserting new nodes, 604–606  
  overview of, 603–604  
  selecting nodes, 609  
non-abstract members, in abstract classes, 210  
nonbreak (normal) mode  
  debugging in, 136–137  
  diagnostic output vs. tracepoints, 144  
  outputting debugging information, 137–138  
  tracepoint use, 142–143  
  writing text to Output window, 138–142  
normal mode. *See* nonbreak (normal) mode  
`NotifyFilter`, 587–588  
`NuGet Package Manager`, 609–610  
null coalescing operator (`??`), 305–306  
null condition operator (`?`), 306–307  
nullable types

example, 307–311  
null coalescing operator (`??`), 305–306  
null condition operator (`?`), 306–307  
operators and, 304–305  
reference types vs. value types, 180  
using generics, 303–304  
numeric aggregation, LINQ, 637–641  
numeric types  
  implicit conversions of, 79  
  overview of, 35–36

## O

`O365`, SaaS, 517  
`Object()`, `System.Object`, 193  
Object Browser, 202–203  
Object-oriented programming. *See* OOP (Object-oriented programming)  
object-relational mapping, Entity Framework for, 653  
objects  
  comparing object references, 379–380  
  constructors/destructors of, 168–169  
  disposable, 172  
  dynamic data binding to external, 450–453  
  event handlers and, 348  
  example of, 181–184  
  initializers, 368–374  
  lifecycle of, 168  
  methods exposed by, 167–168  
  properties and fields of, 166–167  
  relationships between, 177  
  review, 186  
  shallow copying vs. deep copying of, 214–215  
  static data binding to external, 449–450  
  understanding, 165  
`ObservableCollection`  
  creating start game window, 454  
  defined, 458  
  in static data binding, 450

OneDrive, SaaS, 517  
one-way (simplex), WCF message patterns, 684  
online resources  
list of languages using .NET Framework, 4  
Mono, 4  
Visual Studio Express products, 10  
XML tutorials, 594  
OOP (Object-oriented programming)  
collections, 178–179  
constructors/destructors of objects, 168–169  
containment, 177–178  
events, 180  
inheritance, 172–175  
interfaces, 171–172  
lifecycle of objects, 168  
methods exposed by objects, 167–168  
operator overloading, 179  
overview of, 4, 163–165  
polymorphism, 175–177  
properties and fields of objects, 166–167  
reference types vs. value types, 180  
relationships between objects, 177  
review, 186  
static and instance class members, 169–170  
techniques, 170–171  
what objects are, 165  
in WPF desktop applications, 180–185  
operands, 42  
operating systems, supporting .NET, 4  
operation contracts, WCF  
attribute properties, 695–696  
defined, 683  
example, 700  
operations  
common XPath, 611–612  
`for` loop, 71  
WCF, 679  
operator overloading  
adding overloads to `CardLib` class, 284–289  
conversion operators, 295–297

`FrameworkPropertyMetadata` constructor, 480–481  
in OOP, 179  
value comparisons and, 280–284  
operators  
`==` operator, 379  
assignment, 47–48  
Boolean, 54–58  
creating expressions with, 42  
decrement, 44  
generic, 331–332  
increment, 44  
LINQ aggregate, 637–641  
mathematical, 42–47  
null coalescing (?), 305–306  
null condition (?), 306–307  
nullable types and, 304–305  
precedence for, 48–49  
string concatenation and, 43  
optional parameters  
advanced methods, 385–386  
example of, 387–390  
methods, 404  
`OptionalAttribute`, 386  
Options window  
creating, 439–443  
designing, 443–445  
handling events, 446–448  
Oracle, as relational database, 652  
`orderby` clause, 634, 650  
ordering query results  
by multiple levels, 643–645  
`orderby` clause, 634, 650  
overview of, 633–634  
`orientation` enumeration, 87–89, 91–92  
`out` keyword, 118  
`out` parameters, 118–119  
outer variables, anonymous methods, 357  
output  
file classes. *See I/O (Input/Output) file classes*  
writing data to, 575–577

Output window  
diagnostic output vs. tracepoints, 144  
drop-down menu options, 137  
outputting debugging information, 137–138  
writing text to, 138–142  
overflow checking, 81–83  
overloading functions, 128–130  
overloading operators. *See* operator overloading  
overridden methods, 229  
**override** keyword  
defining methods, 219–220  
defining properties, 222  
**override** keyword  
hidden base class methods, 227–229  
overridden methods, 229

**P**

PaaS (Platform as a Service), 517–518  
Package Manager Console, database migrations, 669–670  
`package.appxmanifest` file, 734–735  
packages  
converting XML to JSON, 609–610  
creating, 735  
overview of, 734–735  
Padding property, controls, 423–424  
`PadLeft()`, string manipulation, 101–102  
`PadRight()`, string manipulation, 101–102  
Panel class, 422  
parameter arrays, 114–116  
parameters  
advanced method, 384–385  
creating anonymous methods, 357  
functions and, 98–99, 112–113  
global data vs. return values and, 123–125  
how it works, 113–114  
lambda expressions, 396  
matching, 114  
named, 386–390  
optional, 385–390  
out, 118–119  
parameter arrays, 114–116  
`Read()` method, 571  
reference and value, 116–118  
`Register()` method, 479  
review, 134  
using delegates to call functions, 130–131  
`Write()` method, 575  
**params** keyword, 114–116  
parent (base) class  
hidden base class methods, 227–229  
inheriting from, 172–175  
overridden or methods, 229  
polymorphism and, 175–176  
relationships between objects, 177–178  
parentheses (), 48–49, 110  
`ParentNode` property, `XmlElement`, 600  
partial class definitions, 235–237  
**partial** keyword, 235–237  
PascalCase, function names, 109  
Path class, 562  
paths  
absolute vs. relative, 566–567, 590  
adjusting folder structure, 598  
using \ character, 134  
period(.) character, 48  
Picture Viewer, 734  
Pictures Library, 734  
Platform as a Service (PaaS), 517–518  
PLINQ, 625  
pointers, to objects, 213  
polymorphism  
in collection example, 255  
interface, 176–177  
overview of, 175–176  
review of, 175–177  
variance vs., 335  
portrait layout, changing to landscape, 714–717  
precedence, operator

- with Boolean operators, 58–59, 63  
 overview of, 48–49  
 review, 51  
 preferences, Visual C# Development Settings, 14–15  
 prime numbers, enumerating collection of, 269–270  
 private, access properties for objects, 166  
 private cloud, 516–518  
**private keyword**  
     defining nested types, 231  
     defining properties, 221  
     member definitions, 218  
 procedural programming, 164  
 products, Visual Studio Express, 10  
`Program.cs`, viewing, 21  
 programming  
     C#. *See* C#, writing program in  
     .NET Framework support for, 4–5  
     OOP. *See* OOP (Object-oriented programming)  
     options for Universal apps, 713  
     WCF, 685–691  
 Progress controls, 457  
 prop code snippet, automatic properties, 226  
 properties  
     automatic, 226–227  
     CheckBox control, 441  
     class member, 220–222  
     ComboBox control, 442–443  
     DirectoryInfo class, 566  
     example applying, 223–225  
     FileSystemInfo, 565–566  
     FileSystemWatcher, 584  
     Image control, 434–435  
     ListBox control, 453  
     MenuItem, 462  
     object, 166–167  
     overflow, 82  
     property accessor, 235  
     RadioButton control, 441  
     read-only, 377–378  
     TextBox control, 440  
     timeline, 475–476  
     WCF data contract, 694  
     WCF operation contract, 695–696  
     WCF service contract, 695  
     WPF, manipulating control, 413–416  
     WPF alignment, margin, padding, dimension, 423–424  
     WPF attached, 416–417  
     WPF dependency, 416  
     XmlElement, 599  
 Properties window, Visual Studio 2015, 17, 21–22  
 Properties window, WPF, 25–27  
 PropertyMetadata class, 480  
 protected accessibility, 173  
 protected keyword, member definitions, 218  
 proxy class, WCF clients, 701, 707  
 public classes, 188–190  
 public cloud, 516–518  
 public interfaces, 190–191  
**public keyword**  
     access properties for objects, 166  
     defining fields, 218  
     defining methods, 219–220  
     defining properties, 220  
     defining structs, 90  
     member definitions, 218  
 Publish Web window, ASP.NET web API, 545

**Q**

- qualified names, 48  
 queries  
     LINQ. *See* LINQ (Language Integrated Query)  
     XML, from existing database, 670–674  
 query syntax, LINQ, 625–627, 630, 650

**R**

- Rackspace, cloud, 516–518  
 RadioButton control

**RadioButton control (*continued*)**

Options window, 441, 445  
as UI control, 459  
random access files  
defined, 570  
reading data from, 571–573  
writing data to, 573–575  
range checks, validating user input, 71  
raw bytes, *FileStream*, 570–573  
Razor v3, 536, 547–551  
*Read()* method, *FileStream*, 571–573  
*Read()* method, *StreamReader*, 579  
reading attribute values, 366–367  
reading file data  
review, 590  
with *StreamReader*, 577–580  
using *FileStream*, 570–573  
*ReadKey()*, console applications, 20  
*ReadLine()*, *StreamReader*, 578–579  
*ReadLines()*, *StreamReader*, 580  
readonly keyword, defining fields, 218  
read-only properties, anonymous types and, 377–378  
*ReadToEnd()*, *StreamReader*, 579  
rectangular arrays, 98  
ref keyword, reference parameters, 117–118  
refactoring, 225–226, 489–494  
reference parameters, 116–118  
reference types  
converting value types to, 297–298  
strings as, 41  
structs and, 212–213  
value types vs., 180  
*ReferenceEquals()*, *System.Object* methods, 193  
References, viewing in Solution Explorer, 21  
reflection  
dynamic lookup and, 380–381  
reading attributes, 366–367

#region keyword, 34  
*Register()*, parameters, 479  
relational databases, 652  
relational operators, Boolean, 54–55  
relationships, navigating database, 662–669  
relative path names, 566–567  
**RelativePanels** control  
adding, 717  
moving content, 721  
specifying control positions, 715  
release builds, breakpoints ignored in, 145  
remoting, WCF and, 678–679  
*RemoveAll()* method, nodes, 608–609  
*RemoveChild()* method, nodes, 608–609  
Representative State Transfer (REST), WCF, 679  
request/response, WCF message patterns, 684  
Reset All Settings, Visual Studio, 14–15  
resources, scaling web API to user, 551–554  
REST (Representative State Transfer), WCF, 679  
Retry pattern, cloud, 520  
return  
flow control for case statement, 64–65  
interrupting loops with, 72  
using return values with, 111  
return values  
event handlers and, 356–357  
exchanging data with functions, 110–112  
global data vs. parameters and, 123–125  
review, 134  
using delegates, 131  
reusable code, functions and, 108  
**Route** annotation, ASP.NET, 543  
routed commands  
example applying, 420–422  
with menus, 462–463  
overview of, 419–420  
routed events, 419, 458  
rows, *Grid* control, 431–433  
runtime, managed code as, 6

**S**

SaaS (Software as a Service), 517–518  
 sandboxed apps, 722  
`SaveCompressedFile()`, 581–583  
`sbyte` type  
 implicit numeric conversions of, 79  
 as integer type, 35  
 as underlying type in enumerations, 86  
 scaling web API, in cloud  
 based on CPU usage, 552–554  
 at specific time, 554–556  
 to user requirements, 551–552  
 schemas, XML, 595–597  
 scope. *See* variable scope  
 screen orientation, in Universal apps, 712  
 sealed classes, 175, 188–190  
 searching XML with XPath, 611–615  
`Seek()` method, `FileStream` file pointer, 570  
 SEH (structured exception handling)  
 C# syntax for, 153  
 as error handling. *See* error handling  
`select` clause, query syntax, 627–629  
`SELECT DISTINCT` queries, LINQ, 641–643  
 selections, LINQ queries, 628–629  
`SelectNodes()`, `XmlNode`, 609  
`SelectSingleNode()`, `XmlNode`, 598, 609  
 self-hosted services  
 creating, 702–707  
 defined, 685  
 overview of, 701–702  
 semantic (logic) errors, 135  
`Serializable` attribute, 723  
 serialization, 722–726  
 Server Explorer, accessing database from, 662  
 service contracts, WCF  
 attribute properties, 695  
 defined, 683  
 programming, 688, 700  
 service models, cloud, 517–518  
 set keyword, accessor properties, 220, 222  
 shallow copying, 214–215  
 Sharding pattern, cloud programming, 519  
 shared (static) members, classes, 169–170, 186  
 short type  
 explicit conversions of, 80–81  
 implicit numeric conversions of, 79–80  
 as integer type, 35  
 as underlying type in enumerations, 86  
`ShowDouble()`, 116–118, 124  
 signatures, function, 129–130  
 size, array, 93–95  
 Slider controls, 457  
 SOA (service-oriented architecture), WCF and, 679  
 SOAP (Simple Object Access Protocol), WCF and, 679  
 Software as a Service (SaaS), 517–518  
 Solution Explorer window, Visual Studio 2015, 16–17, 20–21  
 solutions, Visual Studio, 11  
 sorting collections, 292–295  
 spline, as key frame, 477  
`Split()`, statement auto-completion, 102–104  
 SQL (Structured Query Language), relational databases, 652–653  
 SQL Server  
 creating local server instance of database, 661  
 installing SQL Server Express, 653  
 as relational database, 652  
 stack order  
`DockPanel` control, 426  
`WPF` controls, 423  
 StackPanel control  
 as content control, 458  
 as content layout control, 423  
 creating About window, 437–438  
 designing Options window, 444  
 as layout control, 459  
 overview of, 428–429

Start Debugging, new project, 19  
Start Page, Visual Studio, 17–18  
Start Without Debugging, new project, 19  
state  
  app state, 737  
  managing, 730  
  managing app, 730  
  of objects, 166  
  resuming from suspension, 731–733  
statement bodies, lambda expressions, 396–398  
statements  
  auto-completion of, 102–104  
  basic C# syntax for, 30  
  looping. *See* looping  
static (shared) members, classes, 169–170, 186  
static classes, 170  
static constructors, 170  
static data binding, 449–450  
**static keyword**  
  accessing method of static class directly, 57  
  defining fields, 218–219  
  function definition with, 108–110  
  function of, 49  
  for global variables, 120–121  
  not required for struct functions, 127–128  
static methods, `File` and `Directory` classes, 563  
storage accounts, creating, 521–522  
storage containers, cloud computing  
  creating using Azure C# libraries, 520–521  
  creating web site using, 530–532  
  exercise, 523–530  
store accounts, Windows Store, 737  
storyboards  
  animations, 475  
  triggers used in, 473  
**StreamReader class**  
  asynchronous file access, 581  
  defined, 568  
  overview of, 577–580  
streams  
  asynchronous file access, 580  
  classes for using, 567–568  
  `FileStream` class, 568–575  
  reading and writing compressed files, 580–583  
  review, 590  
  `StreamReader` class, 577–580  
  `StreamWriter` class, 575–577  
  understanding, 567  
  Universal Apps, 722  
**StreamWriter class**  
  asynchronous file access, 581  
  defined, 568  
  overview of, 575–577  
  reading and writing compressed files, 581, 583  
stretch, as alignment property, 423–424  
string array, 95  
string literals, 40–41  
string variables, 43  
strings  
  \ character and @ prefix in, 564  
  binary + operator used with, 43  
  defined, 36–37  
  manipulating, 99–104  
  no implicit conversion of, 79  
  as text type, 36  
  using, 37–38  
  strongly typed, C# as, 374  
**struct keyword**, 90  
structs  
  adding functions to, 127–128  
  boxing, 275–276  
  defined, 77  
  defining, 90  
  example, 212–213  
  generic structs, 332  
  how it works, 90–92, 213–214  
  overview of, 89, 212  
  review, 105  
  unboxing, 276–277  
structured exception handling (SEH)

- C# syntax for, 153
  - as error handling. *See* error handling
  - Structured Query Language (SQL), relational
    - databases, 652–653
    - styles, applied to controls, 467–471
    - subscriptions to events, 345
    - suspension, app resuming after, 731–733
    - switch statement, branching with, 63–66
    - symbolic information, debug builds, 136
    - synchronization, Visual C# Developer Settings, 14–15
    - syntax, basic C#, 31–34
    - System namespace, 159–160
    - System.Array, 252–253
    - System.Attribute, 367
    - System.Collections, 252, 292
    - System.Collections.Generic, 311–312
    - System.Convert, 57
    - System.Diagnostics, 138–142
    - SystemException, 343
    - System.IO namespace, 562, 575–576
    - System.IO.Compression namespace, 562, 581–583
    - System.Linq namespace, 621
    - System.Nullable<T>, 303–304
    - System.Object
      - all classes inheriting from, 189, 193–195
      - shallow copying, 271
    - System.Reflection, 366–367
    - System.String, 227
    - System.Xml namespace, 598
- T**
- TabControl, 443–444, 459
  - tags, XML, 594
  - Task Manager, exiting infinite loops, 73
  - TCP (Transmission Control Protocol)
    - addresses, 681
    - bindings, 683
  - communication with WCF services, 680
  - Team Explorer window, Visual Studio 2015, 17
  - templates
    - applied to controls, 467–471
    - applying to user control, 487
    - creating web API, 543
  - ternary (or conditional) operator
    - branching with, 59
    - defined, 42
    - most common usage of, 68
  - text
    - types of, 36–37
    - writing to Output window, 138–142
  - TextBlock control
    - adding, 717
    - looping through all nodes in XML document, 600–603
    - as UI control, 459
  - TextBlockResult control, 610, 613–615
  - TextBox control
    - combining with other controls, 457
    - creating About window, 437–439
    - creating Options window, 439–440
    - in game client example, 435
    - name property, 455
  - this keyword, 230, 260
  - Throttling pattern, cloud programming, 520
  - ThrowException(), 159–160
  - throwing exceptions, 154–160
  - tiles
    - adding, 733–734
    - common in Windows Store apps, 733
    - developing Universal apps, 713
  - time, scaling web API at specific, 554–556
  - timelines, 475–477
  - ToCharArray(), 99, 574
  - ToLower(), string manipulation, 100
  - toolbars
    - creating new project, 19
    - developing Universal apps, 712–713

## Toolbox window

Visual Studio 2015, 16  
WPF, 23, 24  
`ToString()`, 192, 194  
`ToUpper()`, string manipulation, 100  
Trace command, 144  
`Trace.Assert()`, assertions, 146–147  
tracepoints, 142–144  
`Trace.WriteLine()`, 137–138, 160  
Transmission Control Protocol. *See* TCP (Transmission Control Protocol)  
`TriggerAction` class, 473  
`TriggerBase` class, 473  
triggers  
adaptive, 715  
in animation, 475  
overview of, 473–474  
`Trim()` command, string manipulation, 100–101  
`TrimEnd()` command, string manipulation, 101  
`TrimStart()` command, string manipulation, 101  
try keyword, 153–160  
`try.catch.finally`, 153–160  
tunneling events, 419  
two-way (duplex), WCF message patterns, 684  
type comparisons  
`is` operator and, 277–279  
overview of, 275–277  
review, 300  
type conversion  
explicit conversions, 80–83  
implicit conversions, 78–80  
with mathematical operators, 46  
overview of, 78  
in practice, 83–85  
review, 105  
type inference, 374–376, 404  
`typeof` operator, 89  
types. *See* data types  
typesafe language, C# as, 9

**U**

u characters, variable names, 36  
UDP (User Datagram Protocol)  
addresses, 681  
bindings, 683  
communication with WCF services, 680–681  
UI (user interface)  
controls, 459  
creating desktop applications, 22–26  
designing for game, 434  
languages for developing Universal apps, 713  
Visual Studio options for, 407  
`uint` type  
implicit numeric conversions of, 79  
as integer type, 36  
as underlying type in enumerations, 86  
`ulong` type  
implicit numeric conversions of, 79  
as integer type, 36  
as underlying type in enumerations, 86  
UML (Unified Modeling Language)  
class diagrams vs., 205  
method syntax, 167–168  
visualizing contained classes, 178  
working with classes and objects, 165  
unary operators  
defined, 42  
increment/decrement, 44  
mathematical operators as, 43  
overloading, 281–283  
unboxing, comparing objects, 275–277  
unchecked keyword, overflow checking, 81–82  
underlying type, enumerations, 86  
underscore character(\_), variable naming, 39  
Unicode escape sequences, 41  
Unified Modeling Language. *See* UML (Unified Modeling Language)  
Universal Apps  
adaptive displays, 714–717  
adding tiles and badges, 733–734

CommandBar control, 729  
 concepts and design, 712–713  
 creating command bars, 729–730  
 developing, 713–714  
 disk access, 722–726  
 elements of Windows Store apps, 733  
 flip view, 717–721  
 getting started, 710  
 navigating between pages, 726–728  
 overview of, 710–712  
 packaging for distribution to Windows Store, 734–735  
 resuming from suspension, 731–733  
 review, 736  
 sandboxed apps, 722  
 serialization, streams, and async programming, 722  
 state management, 730  
 unmanaged code, 6  
 user controls  
   adding to game application, 481–488  
   completing game application, 506  
   implementing dependency properties, 478–481  
   overview of, 478  
 User Datagram Protocol. *See* UDP (User Datagram Protocol)  
 user interface. *See* UI (user interface)  
 UserControl. *See* user controls  
 users  
   designing validation for input of, 71  
   scaling web API to requirements of, 551–552  
 ushort type  
   implicit numeric conversions of, 79  
   as integer type, 35  
   as underlying type in enumerations, 86  
 using keyword  
   controlling resources used by objects, 172  
   function of, 49  
   visualizing collections, 178–179  
 UWP (Universal Windows Platform), 710

**V**

val parameter, 116–117, 124–125  
 Valet Key pattern, cloud programming, 519  
 validation  
   of user input, 71  
   of XML document against schema, 595  
 value comparisons  
   adding operator overloads to CardLib, 284–289  
   operator overloading and, 280–284  
   overview of, 279–280  
   review, 300  
 value converters  
   overview of, 472–473  
   with user control, 484–485, 487  
 value parameters, 116–118  
 Value property. *See* node values, changing  
 value types  
   boxing and unboxing, 275–277  
   converting to reference types, 297–298  
   reference types vs., 180  
   structs as, 212–213  
 ValueConversionAttribute, 472–473  
 values  
   assigning to enumerations, 86–87  
   assigning to multidimensional arrays, 96–97  
   bool type, 54  
   node. *See* node values, changing  
   using return, 110–112  
 var keyword, 627–628  
 variable scope  
   assigning to multidimensional arrays, 97  
   how it works, 119–122  
   in other structures, 122–123  
   overview of, 119  
   parameters/return values vs. global data, 123–125  
   review, 134  
 variables  
   arrays as. *See* arrays  
   basic C# syntax, 30–34

variables (*continued*)  
    changing content of, 152  
    creating expressions with, 42  
    declaring, 86  
    declaring in LINQ queries, 627–628  
    dynamic, 380  
    dynamic keyword defining, 381  
    enumerations as, 85–89  
    as literal values, 39–41  
    manipulating with mathematical operators, 46  
    monitoring variable content, 148–150  
    naming, 39  
    outer, 357  
    overview of, 34, 77–78  
    reference types vs. value types, 180  
    review, 51–52, 104–105  
    shallow copying vs. deep copying, 214–215  
    simple types of, 34–39  
    statement auto-completion with, 102–104  
    storing data with, 30  
    string manipulation and, 99–104  
    strongly typed languages and, 374–376  
    structs as, 89–92  
    type conversion and, 78–85  
variance, 335–337  
vector class, 310–311, 316–319  
verbatim string literals, 41  
VerticalAlignment, 423–424  
view models  
    creating for game application, 494–502  
    MVVM design pattern, 489  
    purpose of, 494  
virtual classes, inheritance and, 173–174  
virtual keyword  
    defining methods, 219  
    defining properties, 222  
    implementing interfaces, 234  
Visual C# Developer Settings, 14–15, 19  
Visual State Manager, 718–719, 737  
Visual Studio 2015  
    consuming web API from web site, 547–551  
    creating Universal apps, 711  
    creating web API, 540–543  
    creating XML document in, 595–597  
    debugging in. *See* debugging, in Visual Studio  
    options for formatting code, 31  
    overview of, 10  
    review, 12  
    solutions, 11  
    statement auto-completion in, 102–104  
    testing WCF services, 691–693  
    Visual Studio Express products, 10  
    writing .NET application with, 5  
Visual Studio 2015 development environment  
    creating console application, 17–20  
    creating desktop application, 22–26  
    Error list window, 22  
    overview of, 14–17  
    Properties window, 21–22  
    Solution Explorer, 20–21  
    void keyword, 108–111, 125

## W

WAS (Windows Activation Service), WCF services, 685  
Watch window, monitoring variable content, 149–150  
WCF (Windows Communication Foundation)  
    addresses, endpoints, and bindings, 681–683  
    behaviors, 684  
    communication protocols, 680–681  
    concepts, 680  
    contracts, 683–684  
    creating contracts, 697–700  
    data contracts, 694  
    fault contracts, 696  
    hosting WCF services, 684–685  
    message contracts, 696  
    message patterns, 684

operation contracts, 695–696  
overview of, 678–680  
programming, 685–691  
review, 707–709  
self-hosted services, 701–707  
service contracts, 695  
WCF test client, 691–693  
web API  
    consuming from web site, 547–551  
    creating, 540–543  
    deploying, 544–546  
    scaling at specific time, 554–556  
    scaling based on CPU usage, 552–554  
    scaling to user requirements, 551–552  
    writing with C#, 9  
Web API 2 Controller, 542  
web pages, navigation between, 726–728  
web servers, hosting WCF services, 684–685  
Web Service Description Language (WSDL), 679  
web services  
    WCF and, 678–679, 684–685  
    WSDL, 679  
web sites  
    consuming web API from, 547–551  
    creating site that deals two hands of cards, 532–537  
    creating site that uses a storage container, 530–532  
What You See Is What You Get (WYSIWYG),  
    XAML view, 411  
where clause, query syntax, 627, 628  
while loops, 69–71  
whitespaces  
    basic C# syntax, 30  
    console application structure, 33–34  
width, alignment property, 423–424  
windows  
    adding to game application, 462  
    creating About, 436–439  
    creating Options, 439–443  
    Visual Studio, 16–17  
Windows, hosting WCF services, 685  
Windows 10, registering for app development, 710  
Windows Activation Service (WAS), WCF services, 685  
Windows Communication Foundation. *See* WCF  
    (Windows Communication Foundation)  
Windows Designer, 408  
Windows Forms  
    creating desktop applications, 26  
    creating user interfaces, 407  
    WPF compared to, 461  
    WPF replacing. *See* WPF (Windows Presentation Foundation)  
Windows Presentation Foundation. *See* WPF  
    (Windows Presentation Foundation)  
Windows Store  
    apps, 733  
    deploying Universal apps, 712  
    packaging apps for, 734–735  
    sandboxing apps and, 722  
    store accounts, 737  
    writing applications with C#, 9  
Windows Task Manager, exiting infinite loops, 73  
Windows Universal Apps. *See* Universal Apps  
WPF (Windows Presentation Foundation)  
    creating desktop applications, 22–26  
    OOP in desktop applications, 180–185  
WPF (Windows Presentation Foundation),  
    advanced desktop programming  
    adding main window and menus for game  
        application, 462  
        animations, 475–477  
    completing game application example, 502–511  
    creating main window, 463–466  
    implementing dependency properties, 478–481  
    overview of, 461  
    refactoring domain model, 489–494  
    routed commands with menus, 462–463  
    styles and templates applied to controls, 467–471

triggers, 473–474  
user controls, 478, 481–488  
value converters, 472–473  
view models, 494–502  
WPF (Windows Presentation Foundation), basic  
    desktop programming  
    accessing About window, 433–434  
    alignment, margin, padding, and dimension  
        properties, 423–424  
    attached property, 416–417  
    Border control, 424  
    Canvas control, 424–425  
    control layout, 422–423  
    controls, 412–413  
    controls used in game example, 434–436  
    creating About window, 436–439  
    creating Options window, 439–443  
    creating start game window using ListBox,  
        453–457  
    data binding, 448–449  
    dependency property, 416  
    designing Options window, 443–445  
    designing user interface, 434  
    DockPanel control, 426–428  
    dynamic data binding, 450–453  
    editor features, 411–412  
    event handling, 418–419  
    events, 417–418  
    Grid control, 430–433  
    handling events in Options window, 446–448  
    overview of, 407–408  
    properties, 413–416  
    review, 457–459  
    routed commands, 419–422  
    routed events, 419  
    stack order of controls, 423  
    StackPanel control, 428–429  
    static data binding, 449–450  
    types of controls, 422  
    WrapPanel control, 429–430  
    XAML and, 408–411

WrapPanel control  
    as content layout control, 423  
    as layout control, 459  
    overview of, 429–430  
Write()  
    defining and using, 108–110  
    in variable scope, 119–121  
    writing data using FileStream, 574–575  
    writing data with StreamWriter, 576–577  
WriteLine()  
    getting feedback about operations, 136–137  
    writing data with StreamWriter, 576–577  
writing data  
    review, 590  
    with StreamWriter, 575–577  
    using FileStream, 573–575  
WSDL (Web Service Description Language), 679  
WYSIWYG (What You See Is What You Get),  
    XAML view, 411

## X

XAML (Extensible Application Markup  
    Language)  
code-behind files, 411  
defined, 458  
defining user interfaces in WPF, 24–26  
developing Universal apps, 713–714  
example, 409–410  
manipulating control properties, 413–416  
namespace declarations, 410–411  
overview of, 408–409  
routed events, 420–421  
separation of concerns, 409  
Universal Apps. *See* Windows Universal Apps  
value converters, 472  
XAttribute, LINQ to XML constructors, 621  
XDeclaration, LINQ to XML constructors, 622  
XDocument, LINQ to XML constructors, 621  
XDocument, XML fragments, 624  
XElement, LINQ to XML constructors, 621

- 
- XElement, XML fragments, 624
  - XML (Extensible Markup Language), 670–674
  - XML (Extensible Markup Language) and JSON
    - changing values of nodes, 603–609
    - converting XML to JSON, 609–610
    - creating XML document in Visual Studio, 595–597
    - JSON basics, 594–595
    - overview of, 593
    - review, 615–616
    - searching XML with XPATH, 611–615
    - XML basics, 594
    - XML DOM, 597–603
    - XML schemas, 595–597
  - XmlComment class, 598, 606–607
  - XmlDocument class
    - converting XML to JSON, 609–610
    - creating nodes, 606–607
    - defined, 597–598
    - looping through all nodes, 600–603
    - overview of, 598
    - removing nodes, 608–609
    - searching XML with XPATH, 611–615
  - XmlElement class, 598–599, 605–607
  - XmlNode class
    - changing node values. *See* node values, changing defined, 597
    - searching XML with XPATH, 613–615
  - XmlNodeList class, 598
  - XmlText class, 598, 605–607
  - XPath, searching XML with, 611–615

# **WILEY END USER LICENSE AGREEMENT**

Go to [www.wiley.com/go/eula](http://www.wiley.com/go/eula) to access Wiley's ebook EULA.